

» Data Source: `azurerm_api_management`

Use this data source to access information about an existing API Management Service.

» Example Usage

```
data "azurerm_api_management" "test" {
  name                       = "search-api"
  resource_group_name       = "search-service"
}

output "api_management_id" {
  value = "${data.azurerm_api_management.test.id}"
}
```

» Argument Reference

- `name` - (Required) The name of the API Management service.
- `resource_group_name` - (Required) The Name of the Resource Group in which the API Management Service exists.

» Attributes Reference

- `id` - The ID of the API Management Service.
- `additional_location` - One or more `additional_location` blocks as defined below
- `location` - The Azure location where the API Management Service exists.
- `gateway_url` - The URL for the API Management Service's Gateway.
- `gateway_regional_url` - The URL for the Gateway in the Default Region.
- `hostname_configuration` - A `hostname_configuration` block as defined below.
- `management_api_url` - The URL for the Management API.
- `notification_sender_email` - The email address from which the notification will be sent.
- `portal_url` - The URL of the Publisher Portal.

- **public_ip_addresses** - The Public IP addresses of the API Management Service.
 - **publisher_name** - The name of the Publisher/Company of the API Management Service.
 - **publisher_email** - The email of Publisher/Company of the API Management Service.
 - **scm_url** - The SCM (Source Code Management) endpoint.
 - **sku** - A **sku** block as documented below.
 - **tags** - A mapping of tags assigned to the resource.
-

A **additional_location** block exports the following:

- **location** - The location name of the additional region among Azure Data center regions.
 - **gateway_regional_url** - Gateway URL of the API Management service in the Region.
 - **public_ip_addresses** - Public Static Load Balanced IP addresses of the API Management service in the additional location. Available only for Basic, Standard and Premium SKU.
-

A **hostname_configuration** block exports the following:

- **management** - One or more **management** blocks as documented below.
 - **portal** - One or more **portal** blocks as documented below.
 - **proxy** - One or more **proxy** blocks as documented below.
 - **scm** - One or more **scm** blocks as documented below.
-

A **management** block exports the following:

- **host_name** - The Hostname used for the Management API.
 - **key_vault_id** - The ID of the Key Vault Secret which contains the SSL Certificate.
 - **negotiate_client_certificate** - Is Client Certificate Negotiation enabled?
-

A **portal** block exports the following:

- **host_name** - The Hostname used for the Portal.

- `key_vault_id` - The ID of the Key Vault Secret which contains the SSL Certificate.
 - `negotiate_client_certificate` - Is Client Certificate Negotiation enabled?
-

A `proxy` block exports the following:

- `default_ssl_binding` - Is this the default SSL Binding?
 - `host_name` - The Hostname used for the Proxy.
 - `key_vault_id` - The ID of the Key Vault Secret which contains the SSL Certificate.
 - `negotiate_client_certificate` - Is Client Certificate Negotiation enabled?
-

A `scm` block exports the following:

- `host_name` - The Hostname used for the SCM URL.
 - `key_vault_id` - The ID of the Key Vault Secret which contains the SSL Certificate.
 - `negotiate_client_certificate` - Is Client Certificate Negotiation enabled?
-

A `sku` block exports the following:

- `name` - Specifies the plan's pricing tier.
- `capacity` - Specifies the number of units associated with this API Management service.

» Data Source: `azurerm_api_management_api`

Use this data source to access information about an existing API Management API.

» Example Usage

```
data "azurerm_api_management_api" "test" {
  name = "search-api"
  api_management_name = "search-api-management"
```

```

    resource_group_name = "search-service"
    revision             = "2"
  }

  output "api_management_api_id" {
    value = "${data.azurearm_api_management_api.test.id}"
  }

```

» Argument Reference

- **name** - (Required) The name of the API Management API.
- **api_management_name** - (Required) The name of the API Management Service in which the API Management API exists.
- **resource_group_name** - (Required) The Name of the Resource Group in which the API Management Service exists.
- **revision** - (Required) The Revision of the API Management API.

» Attributes Reference

- **id** - The ID of the API Management API.
- **description** - A description of the API Management API, which may include HTML formatting tags.
- **display_name** - The display name of the API.
- **is_current** - Is this the current API Revision?
- **is_online** - Is this API Revision online/accessible via the Gateway?
- **path** - The Path for this API Management API.
- **protocols** - A list of protocols the operations in this API can be invoked.
- **service_url** - Absolute URL of the backend service implementing this API.
- **soap_pass_through** - Should this API expose a SOAP frontend, rather than a HTTP frontend?
- **subscription_key_parameter_names** - A `subscription_key_parameter_names` block as documented below.
- **version** - The Version number of this API, if this API is versioned.
- **version_set_id** - The ID of the Version Set which this API is associated with.

A `subscription_key_parameter_names` block exports the following:

- `header` - The name of the HTTP Header which should be used for the Subscription Key.
 - `query` - The name of the QueryString parameter which should be used for the Subscription Key.
-

A `wsdl_selector` block exports the following:

- `service_name` - The name of service to import from WSDL.
- `endpoint_name` - The name of endpoint (port) to import from WSDL.

» Data Source: `azurerm_api_management_group`

Use this data source to access information about an existing API Management Group.

» Example Usage

```
data "azurerm_api_management_group" "test" {
  name                       = "my-group"
  api_management_name       = "example-apim"
  resource_group_name       = "search-service"
}

output "group_type" {
  value = "${data.azurerm_api_management_group.test.type}"
}
```

» Argument Reference

- `api_management_name` - (Required) The Name of the API Management Service in which this Group exists.
- `name` - (Required) The Name of the API Management Group.
- `resource_group_name` - (Required) The Name of the Resource Group in which the API Management Service exists.

» Attributes Reference

- `id` - The ID of the API Management Group.
- `display_name` - The display name of this API Management Group.
- `description` - The description of this API Management Group.
- `external_id` - The identifier of the external Group.
- `type` - The type of this API Management Group, such as `custom` or `external`.

» Data Source: `azurerm_api_management_product`

Use this data source to access information about an existing API Management Product.

» Example Usage

```
data "azurerm_api_management_product" "test" {
  product_id           = "my-product"
  api_management_name = "example-apim"
  resource_group_name = "search-service"
}

output "product_terms" {
  value = "${data.azurerm_api_management_product.test.terms}"
}
```

» Argument Reference

- `api_management_name` - (Required) The Name of the API Management Service in which this Product exists.
- `product_id` - (Required) The Identifier for the API Management Product.
- `resource_group_name` - (Required) The Name of the Resource Group in which the API Management Service exists.

» Attributes Reference

- `id` - The ID of the API Management Product.

- **approval_required** - Do subscribers need to be approved prior to being able to use the Product?
- **display_name** - The Display Name for this API Management Product.
- **published** - Is this Product Published?
- **subscription_required** - Is a Subscription required to access API's included in this Product?
- **description** - The description of this Product, which may include HTML formatting tags.
- **subscriptions_limit** - The number of subscriptions a user can have to this Product at the same time.
- **terms** - Any Terms and Conditions for this Product, which must be accepted by Developers before they can begin the Subscription process.

» Data Source: **azurerm_api_management_user**

Use this data source to access information about an existing API Management User.

» Example Usage

```
data "azurerm_api_management_user" "test" {
  user_id           = "my-user"
  api_management_name = "example-apim"
  resource_group_name = "search-service"
}

output "notes" {
  value = "${data.azurerm_api_management_user.test.notes}"
}
```

» Argument Reference

- **api_management_name** - (Required) The Name of the API Management Service in which this User exists.
- **resource_group_name** - (Required) The Name of the Resource Group in which the API Management Service exists.
- **user_id** - (Required) The Identifier for the User.

» Attributes Reference

- `id` - The ID of the API Management User.
- `first_name` - The First Name for the User.
- `last_name` - The Last Name for the User.
- `email` - The Email Address used for this User.
- `note` - Any notes about this User.
- `state` - The current state of this User, for example `active`, `blocked` or `pending`.

» Data Source: `azurerm_application_security_group`

Use this data source to access information about an existing Application Security Group.

» Example Usage

```
data "azurerm_application_security_group" "test" {
  name                = "tf-appsecuritygroup"
  resource_group_name = "my-resource-group"
}

output "application_security_group_id" {
  value = "${data.azurerm_application_security_group.test.id}"
}
```

» Argument Reference

The following arguments are supported:

- `name` - The name of the Application Security Group.
- `resource_group_name` - The name of the resource group in which the Application Security Group exists.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the Application Security Group.

- **location** - The supported Azure location where the Application Security Group exists.
- **tags** - A mapping of tags assigned to the resource.

» Data Source: `azurerm_app_service`

Use this data source to access information about an existing App Service.

» Example Usage

```
data "azurerm_app_service" "test" {
  name                = "search-app-service"
  resource_group_name = "search-service"
}

output "app_service_id" {
  value = "${data.azurerm_app_service.test.id}"
}
```

» Argument Reference

- **name** - (Required) The name of the App Service.
- **resource_group_name** - (Required) The Name of the Resource Group where the App Service exists.

» Attributes Reference

- **id** - The ID of the App Service.
- **location** - The Azure location where the App Service exists.
- **app_service_plan_id** - The ID of the App Service Plan within which the App Service exists.
- **app_settings** - A key-value pair of App Settings for the App Service.
- **connection_string** - An `connection_string` block as defined below.
- **client_affinity_enabled** - Does the App Service send session affinity cookies, which route client requests in the same session to the same instance?
- **enabled** - Is the App Service Enabled?

- `https_only` - Can the App Service only be accessed via HTTPS?
- `client_cert_enabled` - Does the App Service require client certificates for incoming requests?
- `site_config` - A `site_config` block as defined below.
- `tags` - A mapping of tags to assign to the resource.
- `outbound_ip_addresses` - A comma separated list of outbound IP addresses - such as `52.23.25.3,52.143.43.12`
- `possible_outbound_ip_addresses` - A comma separated list of outbound IP addresses - such as `52.23.25.3,52.143.43.12,52.143.43.17` - not all of which are necessarily in use. Superset of `outbound_ip_addresses`.

`connection_string` supports the following:

- `name` - The name of the Connection String.
- `type` - The type of the Connection String.
- `value` - The value for the Connection String.

A `cors` block exports the following:

- `allowed_origins` - A list of origins which are able to make cross-origin calls.
- `support_credentials` - Are credentials supported?

A `ip_restriction` block exports the following:

- `ip_address` - The IP Address used for this IP Restriction.
- `subnet_mask` - The Subnet mask used for this IP Restriction.

`site_config` supports the following:

- `always_on` - Is the app be loaded at all times?
- `app_command_line` - App command line to launch.
- `cors` - A `cors` block as defined above.
- `default_documents` - The ordering of default documents to load, if an address isn't specified.
- `dotnet_framework_version` - The version of the .net framework's CLR used in this App Service.

- `http2_enabled` - Is HTTP2 Enabled on this App Service?
- `ftps_state` - State of FTP / FTPS service for this AppService.
- `ip_restriction` - One or more `ip_restriction` blocks as defined above.
- `java_version` - The version of Java in use.
- `java_container` - The Java Container in use.
- `java_container_version` - The version of the Java Container in use.
- `linux_fx_version` - Linux App Framework and version for the AppService.
- `local_mysql_enabled` - Is "MySQL In App" Enabled? This runs a local MySQL instance with your app and shares resources from the App Service plan.
- `managed_pipeline_mode` - The Managed Pipeline Mode used in this App Service.
- `min_tls_version` - The minimum supported TLS version for this App Service.
- `php_version` - The version of PHP used in this App Service.
- `python_version` - The version of Python used in this App Service.
- `remote_debugging_enabled` - Is Remote Debugging Enabled in this App Service?
- `remote_debugging_version` - Which version of Visual Studio is the Remote Debugger compatible with?
- `scm_type` - The type of Source Control enabled for this App Service.
- `use_32_bit_worker_process` - Does the App Service run in 32 bit mode, rather than 64 bit mode?
- `websockets_enabled` - Are WebSockets enabled for this App Service?
- `virtual_network_name` - The name of the Virtual Network which this App Service is attached to.

» Data Source: `azurerm_app_service_plan`

Use this data source to access information about an existing App Service Plan (formerly known as a **Server Farm**).

» Example Usage

```
data "azurerm_app_service_plan" "test" {
  name                = "search-app-service-plan"
  resource_group_name = "search-service"
}

output "app_service_plan_id" {
  value = "${data.azurerm_app_service_plan.test.id}"
}
```

» Argument Reference

- **name** - (Required) The name of the App Service Plan.
- **resource_group_name** - (Required) The Name of the Resource Group where the App Service Plan exists.

» Attributes Reference

- **id** - The ID of the App Service Plan.
- **location** - The Azure location where the App Service Plan exists
- **kind** - The Operating System type of the App Service Plan
- **sku** - A sku block as documented below.
- **properties** - A **properties** block as documented below.
- **tags** - A mapping of tags assigned to the resource.
- **maximum_number_of_workers** - The maximum number of workers supported with the App Service Plan's sku.

A **sku** block supports the following:

- **tier** - Specifies the plan's pricing tier.
- **size** - Specifies the plan's instance size.
- **capacity** - Specifies the number of workers associated with this App Service Plan.

A **properties** block supports the following:

- **app_service_environment_id** - The ID of the App Service Environment where the App Service Plan is located.

- **maximum_number_of_workers** - Maximum number of instances that can be assigned to this App Service plan.
- **reserved** - Is this App Service Plan Reserved?
- **per_site_scaling** - Can Apps assigned to this App Service Plan be scaled independently?

» Data Source: **azurerm_application_insights**

Use this data source to access information about an existing Application Insights component.

» Example Usage

```
data "azurerm_application_insights" "test" {
  name                = "production"
  resource_group_name = "networking"
}

output "application_insights_instrumentation_key" {
  value = "${data.azurerm_application_insights.test.instrumentation_key}"
}
```

» Argument Reference

- **name** - (Required) Specifies the name of the Application Insights component.
- **resource_group_name** - (Required) Specifies the name of the resource group the Application Insights component is located in.

» Attributes Reference

- **id** - The ID of the Virtual Machine.
- **app_id** - The App ID associated with this Application Insights component.
- **application_type** - The type of the component.
- **instrumentation_key** - The instrumentation key of the Application Insights component.
- **location** - The Azure location where the component exists.
- **tags** - Tags applied to the component.

» Data Source: `azurerm_availability_set`

Use this data source to access information about an existing Availability Set.

» Example Usage

```
data "azurerm_availability_set" "test" {
  name                       = "tf-appsecuritygroup"
  resource_group_name       = "my-resource-group"
}

output "availability_set_id" {
  value = "${data.azurerm_availability_set.test.id}"
}
```

» Argument Reference

The following arguments are supported:

- `name` - The name of the Availability Set.
- `resource_group_name` - The name of the resource group in which the Availability Set exists.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the Availability Set.
- `location` - The supported Azure location where the Availability Set exists.
- `managed` - Whether the availability set is managed or not.
- `platform_fault_domain_count` - The number of fault domains that are used.
- `platform_update_domain_count` - The number of update domains that are used.
- `tags` - A mapping of tags assigned to the resource.

» Data Source: `azurerm_azuread_application`

Use this data source to access information about an existing Application within Azure Active Directory.

NOTE: The Azure Active Directory resources have been split out into a new AzureAD Provider - as such the AzureAD resources within the AzureRM Provider are deprecated and will be removed in the next major version (2.0). Information on how to migrate from the existing resources to the new AzureAD Provider can be found [here](#).

NOTE: If you're authenticating using a Service Principal then it must have permissions to both `Read and write all applications` and `Sign in and read user profile` within the Windows Azure Active Directory API.

» Example Usage

```
data "azurerm_azuread_application" "test" {
  name = "My First AzureAD Application"
}

output "azure_active_directory_object_id" {
  value = "${data.azurerm_azuread_application.test.id}"
}
```

» Argument Reference

- `object_id` - (Optional) Specifies the Object ID of the Application within Azure Active Directory.
- `name` - (Optional) Specifies the name of the Application within Azure Active Directory.

NOTE: Either an `object_id` or `name` must be specified.

» Attributes Reference

- `id` - the Object ID of the Azure Active Directory Application.
- `application_id` - the Application ID of the Azure Active Directory Application.
- `available_to_other_tenants` - Is this Azure AD Application available to other tenants?

- `identifier_uris` - A list of user-defined URI(s) that uniquely identify a Web application within its Azure AD tenant, or within a verified custom domain if the application is multi-tenant.
- `oauth2_allow_implicit_flow` - Does this Azure AD Application allow OAuth2.0 implicit flow tokens?
- `object_id` - the Object ID of the Azure Active Directory Application.
- `reply_urls` - A list of URLs that user tokens are sent to for sign in, or the redirect URIs that OAuth 2.0 authorization codes and access tokens are sent to.

» Data Source: `azurerm_azuread_service_principal`

Gets information about an existing Service Principal associated with an Application within Azure Active Directory.

NOTE: The Azure Active Directory resources have been split out into a new AzureAD Provider - as such the AzureAD resources within the AzureRM Provider are deprecated and will be removed in the next major version (2.0). Information on how to migrate from the existing resources to the new AzureAD Provider can be found [here](#).

NOTE: If you're authenticating using a Service Principal then it must have permissions to both `Read and write all applications` and `Sign in and read user profile` within the Windows Azure Active Directory API.

» Example Usage (by Application Display Name)

```
data "azurerm_azuread_service_principal" "test" {
  display_name = "my-awesome-application"
}
```

» Example Usage (by Application ID)

```
data "azurerm_azuread_service_principal" "test" {
  application_id = "00000000-0000-0000-0000-000000000000"
}
```

» Example Usage (by Object ID)

```
data "azurerm_azuread_service_principal" "test" {
  object_id = "00000000-0000-0000-0000-000000000000"
}
```


}

» Argument Reference

The following arguments are supported:

- `application_id` - (Optional) The ID of the Azure AD Application for which to create a Service Principal.
- `object_id` - (Optional) The ID of the Azure AD Service Principal.
- `display_name` - (Optional) The Display Name of the Azure AD Application associated with this Service Principal.

NOTE: At least one of `application_id`, `display_name` or `object_id` must be specified.

» Attributes Reference

The following attributes are exported:

- `id` - The Object ID for the Service Principal.

» Data Source: `azurerm_batch_account`

Use this data source to access information about an existing Batch Account.

» Example Usage

```
data "azurerm_batch_account" "test" {
  name                = "testbatchaccount"
  resource_group_name = "test"
}

output "pool_allocation_mode" {
  value = "${data.azurerm_batch_account.test.pool_allocation_mode}"
}
```

» Argument Reference

- `name` - (Required) The name of the Batch account.
- `resource_group_name` - (Required) The Name of the Resource Group where this Batch account exists.

» Attributes Reference

The following attributes are exported:

- `id` - The Batch account ID.
- `name` - The Batch account name.
- `location` - The Azure Region in which this Batch account exists.
- `pool_allocation_mode` - The pool allocation mode configured for this Batch account.
- `storage_account_id` - The ID of the Storage Account used for this Batch account.
- `primary_access_key` - The Batch account primary access key.
- `secondary_access_key` - The Batch account secondary access key.
- `account_endpoint` - The account endpoint used to interact with the Batch service.
- `tags` - A map of tags assigned to the Batch account.

NOTE: Primary and secondary access keys are only available when `pool_allocation_mode` is set to `BatchService`. See documentation for more information.

» Data Source: `azurerm_batch_certificate`

Use this data source to access information about an existing certificate in a Batch Account.

» Example Usage

```
data "azurerm_batch_certificate" "example" {
  name           = "SHA1-42C107874FD0E4A9583292A2F1098E8FE4B2EDDA"
  account_name   = "examplebatchaccount"
  resource_group_name = "example"
}

output "thumbprint" {
  value = "${data.azurerm_batch_certificate.example.thumbprint}"
}
```

» Argument Reference

- **name** - (Required) The name of the Batch certificate.
- **account_name** - (Required) The name of the Batch account.
- **resource_group_name** - (Required) The Name of the Resource Group where this Batch account exists.

» Attributes Reference

The following attributes are exported:

- **id** - The Batch certificate ID.
- **public_data** - The public key of the certificate.
- **format** - The format of the certificate, such as **Cer** or **Pfx**.
- **thumbprint** - The thumbprint of the certificate.
- **thumbprint_algorithm** - The algorithm of the certificate thumbprint.

» Data source: `azurerm_batch_pool`

Use this data source to access information about an existing Batch pool

» Example Usage

```
data "azurerm_batch_pool" "test" {
  name           = "testbatchpool"
  account_name   = "testbatchaccount"
  resource_group_name = "test"
}
```

» Attributes Reference

The following attributes are exported:

- **id** - The Batch pool ID.
- **name** - The name of the Batch pool.
- **account_name** - The name of the Batch account.
- **node_agent_sku_id** - The Sku of the node agents in the Batch pool.

- **vm_size** - The size of the VM created in the Batch pool.
- **fixed_scale** - A **fixed_scale** block that describes the scale settings when using fixed scale.
- **auto_scale** - A **auto_scale** block that describes the scale settings when using auto scale.
- **storage_image_reference** - The reference of the storage image used by the nodes in the Batch pool.
- **start_task** - A **start_task** block that describes the start task settings for the Batch pool.
- **max_tasks_per_node** - The maximum number of tasks that can run concurrently on a single compute node in the pool.
- **certificate** - One or more **certificate** blocks that describe the certificates installed on each compute node in the pool.

A **fixed_scale** block exports the following:

- **target_dedicated_nodes** - The number of nodes in the Batch pool.
- **target_low_priority_nodes** - The number of low priority nodes in the Batch pool.
- **resize_timeout** - The timeout for resize operations.

A **auto_scale** block exports the following:

- **evaluation_interval** - The interval to wait before evaluating if the pool needs to be scaled.
- **formula** - The autoscale formula that needs to be used for scaling the Batch pool.

A **start_task** block exports the following:

- **command_line** - The command line executed by the start task.
- **max_task_retry_count** - The number of retry count.
- **wait_for_success** - A flag that indicates if the Batch pool should wait for the start task to be completed.
- **environment** - A map of strings (key,value) that represents the environment variables to set in the start task.
- **user_identity** - A **user_identity** block that describes the user identity under which the start task runs.

- **resource_file** - (Optional) One or more **resource_file** blocks that describe the files to be downloaded to a compute node.
-

A **user_identity** block exports the following:

- **user_name** - The username to be used by the Batch pool start task.
 - **auto_user** - A **auto_user** block that describes the user identity under which the start task runs.
-

A **auto_user** block exports the following:

- **elevation_level** - The elevation level of the user identity under which the start task runs.
 - **scope** - The scope of the user identity under which the start task runs.
-

A **certificate** block exports the following:

- **id** - The fully qualified ID of the certificate installed on the pool.
- **store_location** - The location of the certificate store on the compute node into which the certificate is installed, either **CurrentUser** or **LocalMachine**.

NOTE: This property is applicable only for pools configured with Windows nodes (that is, created with `cloudServiceConfiguration`, or with `virtualMachineConfiguration` using a Windows image reference). For Linux compute nodes, the certificates are stored in a directory inside the task working directory and an environment variable `AZ_BATCH_CERTIFICATES_DIR` is supplied to the task to query for this location. For certificates with visibility of `'remoteUser'`, a `'certs'` directory is created in the user's home directory (e.g., `/home/{user-name}/certs`) and certificates are placed in that directory.

- **store_name** - The name of the certificate store on the compute node into which the certificate is installed.

NOTE: This property is applicable only for pools configured with Windows nodes (that is, created with `cloudServiceConfiguration`, or with `virtualMachineConfiguration` using a Windows image reference).

- **visibility** - Which user accounts on the compute node have access to the private data of the certificate.
-

A **resource_file** block exports the following:

- **auto_storage_container_name** - The storage container name in the auto storage account.
- **blob_prefix** - The blob prefix used when downloading blobs from an Azure Storage container.
- **file_mode** - The file permission mode attribute represented as a string in octal format (e.g. "0644").
- **file_path** - The location on the compute node to which to download the file, relative to the task's working directory. If the **http_url** property is specified, the **file_path** is required and describes the path which the file will be downloaded to, including the filename. Otherwise, if the **auto_storage_container_name** or **storage_container_url** property is specified.
- **http_url** - The URL of the file to download. If the URL is Azure Blob Storage, it must be readable using anonymous access.
- **storage_container_url** - The URL of the blob container within Azure Blob Storage.

» Data Source: **azurerm_builtin_role_definition**

Use this data source to access information about a built-in Role Definition. To access information about a custom Role Definition, please see the **azurerm_role_definition** data source instead.

NOTE: The this datasource has been deprecated in favour of **azurerm_role_definition** that now can look up role definitions by name. As such this data source will be removed in version 2.0 of the AzureRM Provider.

» Example Usage

```
data "azurerm_builtin_role_definition" "contributor" {
  name = "Contributor"
}

output "contributor_role_definition_id" {
  value = "${data.azurerm_builtin_role_definition.contributor.id}"
}
```

» Argument Reference

- **name** - (Required) Specifies the name of the built-in Role Definition. Possible values are: **Contributor**, **Owner**, **Reader** and **VirtualMachineContributor**.

» Attributes Reference

- **id** - the ID of the built-in Role Definition.
- **description** - the Description of the built-in Role.
- **type** - the Type of the Role.
- **permissions** - a **permissions** block as documented below.
- **assignable_scopes** - One or more assignable scopes for this Role Definition, such as `/subscriptions/0b1f6471-1bf0-4dda-aec3-111122223333`, `/subscriptions/0b1f6471-1bf0-4dda-aec3-111122223333/resourceGroups/myGroup`, or `/subscriptions/0b1f6471-1bf0-4dda-aec3-111122223333/resourceGroups/myGroup/providers/`

A **permissions** block contains:

- **actions** - a list of actions supported by this role
- **data_actions** - a list of data actions supported by this role
- **not_actions** - a list of actions which are denied by this role
- **not_data_actions** - a list of data actions which are denied by this role

» Data Source: `azurerm_cdn_profile`

Use this data source to access information about an existing CDN Profile.

» Example Usage

```
data "azurerm_cdn_profile" "test" {
  name                = "myfirstcdnprofile"
  resource_group_name = "example-resources"
}

output "cdn_profile_id" {
  value = "${data.azurerm_cdn_profile.test.id}"
}
```

» Argument Reference

- **name** - (Required) The name of the CDN Profile.

- **resource_group_name** - (Required) The name of the resource group in which the CDN Profile exists.

» Attributes Reference

- **location** - The Azure Region where the resource exists.
- **sku** - The pricing related information of current CDN profile.
- **tags** - A mapping of tags assigned to the resource.

» Data Source: `azurerm_client_config`

Use this data source to access the configuration of the AzureRM provider.

» Example Usage

```
data "azurerm_client_config" "current" {}

output "account_id" {
  value = "${data.azurerm_client_config.current.service_principal_application_id}"
}
```

» Argument Reference

There are no arguments available for this data source.

» Attributes Reference

- **client_id** is set to the Azure Client ID (Application Object ID).
- **tenant_id** is set to the Azure Tenant ID.
- **subscription_id** is set to the Azure Subscription ID.

Note: the following fields are only available when authenticating via a Service Principal (as opposed to using the Azure CLI):

- **service_principal_application_id** is the Service Principal Application ID.
- **service_principal_object_id** is the Service Principal Object ID.

Note: To better understand "application" and "service principal", please read Application and service principal objects in Azure Active Directory.

» Data Source: `azurerm_container_registry`

Use this data source to access information about an existing Container Registry.

» Example Usage

```
data "azurerm_container_registry" "test" {
  name                       = "testacr"
  resource_group_name = "test"
}

output "login_server" {
  value = "${data.azurerm_container_registry.test.login_server}"
}
```

» Argument Reference

- `name` - (Required) The name of the Container Registry.
- `resource_group_name` - (Required) The Name of the Resource Group where this Container Registry exists.

» Attributes Reference

The following attributes are exported:

- `id` - The Container Registry ID.
- `login_server` - The URL that can be used to log into the container registry.
- `admin_username` - The Username associated with the Container Registry Admin account - if the admin account is enabled.
- `admin_password` - The Password associated with the Container Registry Admin account - if the admin account is enabled.
- `location` - The Azure Region in which this Container Registry exists.
- `admin_enabled` - Is the Administrator account enabled for this Container Registry.
- `sku` - The SKU of this Container Registry, such as `Basic`.
- `storage_account_id` - The ID of the Storage Account used for this Container Registry. This is only returned for `Classic` SKU's.
- `tags` - A map of tags assigned to the Container Registry.

» Data Source: `azurerm_cosmosdb_account`

Use this data source to access information about an existing CosmosDB (formally DocumentDB) Account.

» Example Usage

```
data "azurerm_cosmosdb_account" "test" {
  name                       = "tfex-cosmosdb-account"
  resource_group_name = "tfex-cosmosdb-account-rg"
}

output "cosmosdb_account_endpoint" {
  value = "${data.azurerm_cosmosdb_account.jobs.endpoint}"
}
```

» Argument Reference

The following arguments are supported:

- `name` - (Required) Specifies the name of the CosmosDB Account.
- `resource_group_name` - (Required) Specifies the name of the resource group in which the CosmosDB Account resides.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the CosmosDB Account.
- `location` - The Azure location where the resource exists.
- `tags` - A mapping of tags assigned to the resource.
- `offer_type` - The Offer Type to used by this CosmosDB Account.
- `kind` - The Kind of the CosmosDB account.
- `ip_range_filter` - The current IP Filter for this CosmosDB account
- `enable_automatic_failover` - If automatic failover is enabled for this CosmosDB Account.
- `capabilities` - Capabilities enabled on this Cosmos DB account.
- `is_virtual_network_filter_enabled` - If virtual network filtering is enabled for this Cosmos DB account.

- **virtual_network_rule** - Subnets that are allowed to access this CosmosDB account.
- **enable_multiple_write_locations** - If multi-master is enabled for this Cosmos DB account.

consistency_policy The current consistency Settings for this CosmosDB account with the following properties:

- **consistency_level** - The Consistency Level used by this CosmosDB Account.
- **max_interval_in_seconds** - The amount of staleness (in seconds) tolerated when the consistency level is Bounded Staleness.
- **max_staleness_prefix** - The number of stale requests tolerated when the consistency level is Bounded Staleness.

geo_location The geographic locations data is replicated to with the following properties:

- **id** - The ID of the location.
- **location** - The name of the Azure region hosting replicated data.
- **priority** - The locations fail over priority.

virtual_network_rule The virtual network subnets allowed to access this Cosmos DB account with the following properties:

- **id** - The ID of the virtual network subnet.
- **endpoint** - The endpoint used to connect to the CosmosDB account.
- **read_endpoints** - A list of read endpoints available for this CosmosDB account.
- **write_endpoints** - A list of write endpoints available for this CosmosDB account.
- **primary_master_key** - The Primary master key for the CosmosDB Account.
- **secondary_master_key** - The Secondary master key for the CosmosDB Account.
- **primary_readonly_master_key** - The Primary read-only master Key for the CosmosDB Account.
- **secondary_readonly_master_key** - The Secondary read-only master key for the CosmosDB Account.

» Data Source: **azurerm_dns_zone**

Use this data source to access information about an existing DNS Zone.

» Example Usage

```
data "azurerm_dns_zone" "test" {
  name                = "search-eventhubns"
  resource_group_name = "search-service"
}

output "dns_zone_id" {
  value = "${data.azurerm_dns_zone.test.id}"
}
```

» Argument Reference

- **name** - (Required) The name of the DNS Zone.
- **resource_group_name** - (Optional) The Name of the Resource Group where the DNS Zone exists. If the Name of the Resource Group is not provided, the first DNS Zone from the list of DNS Zones in your subscription that matches **name** will be returned.

» Attributes Reference

- **id** - The ID of the DNS Zone.
- **max_number_of_record_sets** - Maximum number of Records in the zone.
- **number_of_record_sets** - The number of records already in the zone.
- **name_servers** - A list of values that make up the NS record for the zone.
- **zone_type** - The type of this DNS zone, such as **Public** or **Private**.
- **registration_virtual_network_ids** - A list of Virtual Network ID's that register hostnames in this DNS zone.
- **resolution_virtual_network_ids** - A list of Virtual Network ID's that resolve records in this DNS zone.
- **tags** - A mapping of tags to assign to the EventHub Namespace.

» Data Source: `azurerm__data__lake__store`

Use this data source to access information about an existing Data Lake Store.

» Example Usage

```
data "azurerm_data_lake_store" "test" {
  name                = "testdatalake"
  resource_group_name = "testdatalake"
}

output "data_lake_store_id" {
  value = "${data.azurerm_data_lake_store.test.id}"
}
```

» Argument Reference

- **name** - (Required) The name of the Data Lake Store.
- **resource_group_name** - (Required) The Name of the Resource Group where the Data Lake Store exists.

» Attributes Reference

- **id** - The ID of the Data Lake Store.
- **encryption_state** - the Encryption State of this Data Lake Store Account, such as **Enabled** or **Disabled**.
- **encryption_type** - the Encryption Type used for this Data Lake Store Account.
- **firewall_allow_azure_ips** - are Azure Service IP's allowed through the firewall?
- **firewall_state** - the state of the firewall, such as **Enabled** or **Disabled**.
- **tier** - Current monthly commitment tier for the account.
- **tags** - A mapping of tags to assign to the Data Lake Store.

» Data Source: `azurerm_dev_test_lab`

Use this data source to access information about an existing Dev Test Lab.

» Example Usage

```
data "azurerm_dev_test_lab" "test" {
  name = "example-lab"
}
```

```

    resource_group_name = "example-resources"
  }

  output "unique_identifier" {
    value = "${data.azurerm_dev_test_lab.test.unique_identifier}"
  }

```

» Argument Reference

- `name` - (Required) The name of the Dev Test Lab.
- `resource_group_name` - (Required) The Name of the Resource Group where the Dev Test Lab exists.

» Attributes Reference

- `id` - The ID of the Dev Test Lab.
- `artifacts_storage_account_id` - The ID of the Storage Account used for Artifact Storage.
- `default_storage_account_id` - The ID of the Default Storage Account for this Dev Test Lab.
- `default_premium_storage_account_id` - The ID of the Default Premium Storage Account for this Dev Test Lab.
- `key_vault_id` - The ID of the Key used for this Dev Test Lab.
- `location` - The Azure location where the Dev Test Lab exists.
- `premium_data_disk_storage_account_id` - The ID of the Storage Account used for Storage of Premium Data Disk.
- `storage_type` - The type of storage used by the Dev Test Lab.
- `tags` - A mapping of tags to assign to the resource.
- `unique_identifier` - The unique immutable identifier of the Dev Test Lab.

» Data Source: `azurerm_eventhub_namespace`

Use this data source to access information about an existing EventHub Namespace.

» Example Usage

```
data "azurerm_eventhub_namespace" "test" {
  name                       = "search-eventhubns"
  resource_group_name = "search-service"
}

output "eventhub_namespace_id" {
  value = "${data.azurerm_eventhub_namespace.test.id}"
}
```

» Argument Reference

- **name** - (Required) The name of the EventHub Namespace.
- **resource_group_name** - (Required) The Name of the Resource Group where the EventHub Namespace exists.

» Attributes Reference

- **id** - The ID of the EventHub Namespace.
- **location** - The Azure location where the EventHub Namespace exists
- **sku** - Defines which tier to use.
- **capacity** - The Capacity / Throughput Units for a **Standard** SKU namespace.
- **auto_inflate_enabled** - Is Auto Inflate enabled for the EventHub Namespace?
- **maximum_throughput_units** - Specifies the maximum number of throughput units when Auto Inflate is Enabled.
- **tags** - A mapping of tags to assign to the EventHub Namespace.

The following attributes are exported only if there is an authorization rule named **RootManageSharedAccessKey** which is created automatically by Azure.

- **default_primary_connection_string** - The primary connection string for the authorization rule **RootManageSharedAccessKey**.
- **default_secondary_connection_string** - The secondary connection string for the authorization rule **RootManageSharedAccessKey**.
- **default_primary_key** - The primary access key for the authorization rule **RootManageSharedAccessKey**.

- `default_secondary_key` - The secondary access key for the authorization rule `RootManageSharedAccessKey`.

» Data Source: `azurerm_express_route_circuit`

Use this data source to access information about an existing ExpressRoute circuit.

» Example Usage

```
data "azurerm_express_route_circuit" test {
  resource_group_name = "${azurerm_resource_group.test.name}"
  name                = "${azurerm_express_route_circuit.test.name}"
}

output "express_route_circuit_id" {
  value = "${data.azurerm_express_route_circuit.test.id}"
}

output "service_key" {
  value = "${data.azurerm_express_route_circuit.test.service_key}"
}
```

» Argument Reference

- `name` - (Required) The name of the ExpressRoute circuit.
- `resource_group_name` - (Required) The Name of the Resource Group where the ExpressRoute circuit exists.

» Attributes Reference

- `id` - The ID of the ExpressRoute circuit.
- `location` - The Azure location where the ExpressRoute circuit exists
- `peerings` - A `peerings` block for the ExpressRoute circuit as documented below
- `service_provider_provisioning_state` - The ExpressRoute circuit provisioning state from your chosen service provider. Possible values are "NotProvisioned", "Provisioning", "Provisioned", and "Deprovisioning".

- **service_key** - The string needed by the service provider to provision the ExpressRoute circuit.
- **service_provider_properties** - A **service_provider_properties** block for the ExpressRoute circuit as documented below
- **sku** - A **sku** block for the ExpressRoute circuit as documented below.

service_provider_properties supports the following:

- **service_provider_name** - The name of the ExpressRoute Service Provider.
- **peering_location** - The name of the peering location and **not** the Azure resource location.
- **bandwidth_in_mbps** - The bandwidth in Mbps of the ExpressRoute circuit.

peerings supports the following:

- **peering_type** - The type of the ExpressRoute Circuit Peering. Acceptable values include **AzurePrivatePeering**, **AzurePublicPeering** and **MicrosoftPeering**. Changing this forces a new resource to be created. ~> **NOTE:** only one Peering of each Type can be created per ExpressRoute circuit.
- **primary_peer_address_prefix** - A /30 subnet for the primary link.
- **secondary_peer_address_prefix** - A /30 subnet for the secondary link.
- **vlan_id** - A valid VLAN ID to establish this peering on.
- **shared_key** - The shared key. Can be a maximum of 25 characters.
- **azure_asn** - The Either a 16-bit or a 32-bit ASN for Azure.
- **peer_asn** - The Either a 16-bit or a 32-bit ASN. Can either be public or private.

sku supports the following:

- **tier** - The service tier. Possible values are **Standard** or **Premium**.
- **family** - The billing mode for bandwidth. Possible values are **MeteredData** or **UnlimitedData**.

» Data Source: **azurerm_firewall**

Use this data source to access information about an existing Azure Firewall.

» Example Usage

```
data "azurerm_firewall" "test" {
  name = "firewall1"
```

```

    resource_group_name = "firewall-RG"
  }

  output "firewall_private_ip" {
    value = "${data.azurerm_firewall.test.ip_configuration.0.private_ip_address}"
  }

```

» Argument Reference

- **name** - (Required) The name of the Azure Firewall.
- **resource_group_name** - (Required) The name of the Resource Group in which the Azure Firewall exists.

» Attributes Reference

The following attributes are exported:

- **id** - The Resource ID of the Azure Firewall.
- **ip_configuration** - A `ip_configuration` block as defined below.

A `ip_configuration` block exports the following:

- **subnet_id** - The Resource ID of the subnet where the Azure Firewall is deployed.
- **private_ip_address** - The private IP address of the Azure Firewall.
- **public_ip_address_id** - The Resource ID of the public IP address of the Azure Firewall.

» Data Source: `azurerm_hdinsight_cluster`

Use this data source to access information about an existing HDInsight Cluster.

» Example Usage

```

data "azurerm_hdinsight_cluster" "example" {
  name                = "example"
  resource_group_name = "example-resources"
}

```

```
output "https_endpoint" {
  value = "${data.azurearm_hdinsight_cluster.example.https_endpoint}"
}
```

» Argument Reference

- **name** - (Required) Specifies the name of this HDInsight Cluster.
- **resource_group_name** - (Required) Specifies the name of the Resource Group in which this HDInsight Cluster exists.

» Attributes Reference

- **location** - The Azure Region in which this HDInsight Cluster exists.
- **cluster_version** - The version of HDInsights which is used on this HDInsight Cluster.
- **component_versions** - A map of versions of software used on this HDInsights Cluster.
- **gateway** - A gateway block as defined below.
- **edge_ssh_endpoint** - The SSH Endpoint of the Edge Node for this HDInsight Cluster, if an Edge Node exists.
- **https_endpoint** - The HTTPS Endpoint for this HDInsight Cluster.
- **kind** - The kind of HDInsight Cluster this is, such as a Spark or Storm cluster.
- **tier** - The SKU / Tier of this HDInsight Cluster.
- **ssh_endpoint** - The SSH Endpoint for this HDInsight Cluster.
- **tags** - A map of tags assigned to the HDInsight Cluster.

A gateway block exports the following:

- **enabled** - Is the Ambari Portal enabled?
- **username** - The username used for the Ambari Portal.
- **password** - The password used for the Ambari Portal.

» Data Source: `azurerm_image`

Use this data source to access information about an existing Image.

» Example Usage

```
data "azurerm_image" "search" {
  name           = "search-api"
  resource_group_name = "packerimages"
}

output "image_id" {
  value = "${data.azurerm_image.search.id}"
}
```

» Argument Reference

- **name** - (Optional) The name of the Image.
- **name_regex** - (Optional) Regex pattern of the image to match.
- **sort_descending** - (Optional) By default when matching by regex, images are sorted by name in ascending order and the first match is chosen, to sort descending, set this flag.
- **resource_group_name** - (Required) The Name of the Resource Group where this Image exists.

» Attributes Reference

- **data_disk** - a collection of **data_disk** blocks as defined below.
- **name** - the name of the Image.
- **location** - the Azure Location where this Image exists.
- **os_disk** - a **os_disk** block as defined below.
- **tags** - a mapping of tags to assigned to the resource.
- **zone_resilient** - is zone resiliency enabled?

os_disk supports the following:

- **blob_uri** - the URI in Azure storage of the blob used to create the image.
- **caching** - the caching mode for the OS Disk, such as **ReadWrite**, **ReadOnly**, or **None**.
- **managed_disk_id** - the ID of the Managed Disk used as the OS Disk Image.
- **os_state** - the State of the OS used in the Image, such as **Generalized**.
- **os_type** - the type of Operating System used on the OS Disk. such as **Linux** or **Windows**.
- **size_gb** - the size of the OS Disk in GB.

data_disk supports the following:

- **blob_uri** - the URI in Azure storage of the blob used to create the image.

- `cached` - the caching mode for the Data Disk, such as `ReadWrite`, `ReadOnly`, or `None`.
- `lun` - the logical unit number of the data disk.
- `managed_disk_id` - the ID of the Managed Disk used as the Data Disk Image.
- `size_gb` - the size of this Data Disk in GB.

» Data Source: `azurerm_key_vault`

Use this data source to access information about an existing Key Vault.

» Example Usage

```
data "azurerm_key_vault" "test" {
  name                = "mykeyvault"
  resource_group_name = "some-resource-group"
}

output "vault_uri" {
  value = "${data.azurerm_key_vault.test.vault_uri}"
}
```

» Argument Reference

The following arguments are supported:

- `name` - (Required) Specifies the name of the Key Vault.
- `resource_group_name` - (Required) The name of the Resource Group in which the Key Vault exists.

» Attributes Reference

The following attributes are exported:

- `id` - The Vault ID.
- `vault_uri` - The URI of the vault for performing operations on keys and secrets.
- `location` - The Azure Region in which the Key Vault exists.
- `sku` - A `sku` block as described below.

- **tenant_id** - The Azure Active Directory Tenant ID used for authenticating requests to the Key Vault.
- **access_policy** - One or more **access_policy** blocks as defined below.
- **enabled_for_deployment** - Can Azure Virtual Machines retrieve certificates stored as secrets from the Key Vault?
- **enabled_for_disk_encryption** - Can Azure Disk Encryption retrieve secrets from the Key Vault?
- **enabled_for_template_deployment** - Can Azure Resource Manager retrieve secrets from the Key Vault?
- **tags** - A mapping of tags assigned to the Key Vault.

A **sku** block exports the following:

- **name** - The name of the SKU used for this Key Vault.

access_policy supports the following:

- **tenant_id** - The Azure Active Directory Tenant ID used to authenticate requests for this Key Vault.
- **object_id** - An Object ID of a User, Service Principal or Security Group.
- **application_id** - The Object ID of a Azure Active Directory Application.
- **certificate_permissions** - A list of certificate permissions applicable to this Access Policy.
- **key_permissions** - A list of key permissions applicable to this Access Policy.
- **secret_permissions** - A list of secret permissions applicable to this Access Policy.
- **storage_permissions** - A list of storage permissions applicable to this Access Policy.

» Data Source: **azurerm_key_vault_access_policy**

Use this data source to access information about the permissions from the Management Key Vault Templates.

» Example Usage

```
data "azurerm_key_vault_access_policy" "contributor" {
  name = "Key Management"
}
```

```
output "access_policy_key_permissions" {
  value = "${data.azurerm_key_vault_access_policy.key_permissions}"
}
```

» Argument Reference

- **name** - (Required) Specifies the name of the Management Template. Possible values are: Key Management, Secret Management, Certificate Management, Key & Secret Management, Key & Certificate Management, Secret & Certificate Management, Key, Secret, & Certificate Management

» Attributes Reference

- **id** - the ID of the Key Vault Access Policy
- **key_permissions** - the key permissions for the access policy
- **secret_permissions** - the secret permissions for the access policy
- **certificate_permissions** - the certificate permissions for the access policy

» Data Source: `azurerm_key_vault_key`

Use this data source to access information about an existing Key Vault Key.

Note: All arguments including the secret value will be stored in the raw state as plain-text. Read more about sensitive data in state.

» Example Usage

```
data "azurerm_key_vault_key" "test" {
  name      = "secret-sauce"
  vault_uri = "https://rickslab.vault.azure.net/"
}

output "key_type" {
  value = "${data.azurerm_key_vault_secret.test.key_type}"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Key Vault Key.
- **vault_uri** - (Required) Specifies the ID of the Key Vault Key Vault instance where the Key resides, available on the **azurerm_key_vault** Data Source / Resource.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Key Vault Key.
- **e** - The RSA public exponent of this Key Vault Key.
- **key_type** - Specifies the Key Type of this Key Vault Key
- **key_size** - Specifies the Size of this Key Vault Key.
- **key_opts** - A list of JSON web key operations assigned to this Key Vault Key
- **n** - The RSA modulus of this Key Vault Key.
- **tags** - A mapping of tags assigned to this Key Vault Key.
- **version** - The current version of the Key Vault Key.

» Data Source: **azurerm_key_vault_secret**

Use this data source to access information about an existing Key Vault Secret.

Note: All arguments including the secret value will be stored in the raw state as plain-text. Read more about sensitive data in state.

» Example Usage

```
data "azurerm_key_vault_secret" "test" {
  name          = "secret-sauce"
  key_vault_id = "${data.azurerm_key_vault.existing.id}"
}

output "secret_value" {
  value = "${data.azurerm_key_vault_secret.test.value}"
}
```


» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Key Vault Secret.
- **key_vault_id** - (Required) Specifies the ID of the Key Vault Key Vault instance where the Secret resides, available on the **azurerm_key_vault** Data Source / Resource.

» Attributes Reference

The following attributes are exported:

- **id** - The Key Vault Secret ID.
- **value** - The value of the Key Vault Secret.
- **version** - The current version of the Key Vault Secret.
- **content_type** - The content type for the Key Vault Secret.
- **tags** - Any tags assigned to this resource.

» Data Source: **azurerm_kubernetes_cluster**

Use this data source to access information about an existing Managed Kubernetes Cluster (AKS).

Note: All arguments including the client secret will be stored in the raw state as plain-text. Read more about sensitive data in state.

» Example Usage

```
data "azurerm_kubernetes_cluster" "test" {  
  name           = "myakscluster"  
  resource_group_name = "my-example-resource-group"  
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the managed Kubernetes Cluster.
- **resource_group_name** - (Required) The name of the Resource Group in which the managed Kubernetes Cluster exists.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the Kubernetes Managed Cluster.
- `addon_profile` - A `addon_profile` block as documented below.
- `agent_pool_profile` - An `agent_pool_profile` block as documented below.
- `dns_prefix` - The DNS Prefix of the managed Kubernetes cluster.
- `fqdn` - The FQDN of the Azure Kubernetes Managed Cluster.
- `kube_admin_config` - A `kube_admin_config` block as defined below. This is only available when Role Based Access Control with Azure Active Directory is enabled.
- `kube_admin_config_raw` - Raw Kubernetes config for the admin account to be used by kubectl and other compatible tools. This is only available when Role Based Access Control with Azure Active Directory is enabled.
- `kube_config` - A `kube_config` block as defined below.
- `kube_config_raw` - Base64 encoded Kubernetes configuration.
- `kubernetes_version` - The version of Kubernetes used on the managed Kubernetes Cluster.
- `location` - The Azure Region in which the managed Kubernetes Cluster exists.
- `linux_profile` - A `linux_profile` block as documented below.
- `network_profile` - A `network_profile` block as documented below.
- `node_resource_group` - Auto-generated Resource Group containing AKS Cluster resources.
- `role_based_access_control` - A `role_based_access_control` block as documented below.
- `service_principal` - A `service_principal` block as documented below.
- `tags` - A mapping of tags assigned to this resource.

A `addon_profile` block exports the following:

- `http_application_routing` - A `http_application_routing` block.
 - `oms_agent` - A `oms_agent` block.
-

A `agent_pool_profile` block exports the following:

- `count` - The number of Agents (VM's) in the Pool.
 - `max_pods` - The maximum number of pods that can run on each agent.
 - `name` - The name assigned to this pool of agents.
 - `os_disk_size_gb` - The size of the Agent VM's Operating System Disk in GB.
 - `os_type` - The Operating System used for the Agents.
 - `vm_size` - The size of each VM in the Agent Pool (e.g. `Standard_F1`).
 - `vnet_subnet_id` - The ID of the Subnet where the Agents in the Pool are provisioned.
-

A `azure_active_directory` block exports the following:

- `client_app_id` - The Client ID of an Azure Active Directory Application.
 - `server_app_id` - The Server ID of an Azure Active Directory Application.
 - `tenant_id` - The Tenant ID used for Azure Active Directory Application.
-

A `http_application_routing` block exports the following:

- `enabled` - Is HTTP Application Routing Enabled?
 - `http_application_routing_zone_name` - The Zone Name of the HTTP Application Routing.
-

The `kube_admin_config` and `kube_config` blocks exports the following:

- `client_key` - Base64 encoded private key used by clients to authenticate to the Kubernetes cluster.
- `client_certificate` - Base64 encoded public certificate used by clients to authenticate to the Kubernetes cluster.
- `cluster_ca_certificate` - Base64 encoded public CA certificate used as the root of trust for the Kubernetes cluster.
- `host` - The Kubernetes cluster server host.
- `username` - A username used to authenticate to the Kubernetes cluster.
- `password` - A password or token used to authenticate to the Kubernetes cluster.

NOTE: It's possible to use these credentials with the Kubernetes Provider like so:

```
provider "kubernetes" {
  host          = "${data.azurerm_kubernetes_cluster.main.kube_config.0.host}"
  username      = "${data.azurerm_kubernetes_cluster.main.kube_config.0.username}"
  password      = "${data.azurerm_kubernetes_cluster.main.kube_config.0.password}"
  client_certificate = "${base64decode(data.azurerm_kubernetes_cluster.main.kube_config.0.client_certificate)}"
  client_key      = "${base64decode(data.azurerm_kubernetes_cluster.main.kube_config.0.client_key)}"
  cluster_ca_certificate = "${base64decode(data.azurerm_kubernetes_cluster.main.kube_config.0.cluster_ca_certificate)}"
}
```

A `linux_profile` block exports the following:

- `admin_username` - The username associated with the administrator account of the managed Kubernetes Cluster.
- `ssh_key` - An `ssh_key` block as defined below.

A `network_profile` block exports the following:

- `docker_bridge_cidr` - IP address (in CIDR notation) used as the Docker bridge IP address on nodes.
- `dns_service_ip` - IP address within the Kubernetes service address range used by cluster service discovery (kube-dns).
- `network_plugin` - Network plugin used such as `azure` or `kubenet`.
- `network_policy` - Network policy to be used with Azure CNI. Eg: `calico` or `azure`
- `pod_cidr` - The CIDR used for pod IP addresses.
- `service_cidr` - Network range used by the Kubernetes service.

A `oms_agent` block exports the following:

- `enabled` - Is the OMS Agent Enabled?
- `log_analytics_workspace_id` - The ID of the Log Analytics Workspace which the OMS Agent should send data to.

A `role_based_access_control` block exports the following:

- `azure_active_directory` - A `azure_active_directory` block as documented above.

- `enabled` - Is Role Based Access Control enabled?
-

A `service_principal` block supports the following:

- `client_id` - The Client ID of the Service Principal used by this Managed Kubernetes Cluster.
-

A `ssh_key` block exports the following:

- `key_data` - The Public SSH Key used to access the cluster.

» Data Source: `azurerm_lb`

Use this data source to access information about an existing Load Balancer

» Example Usage

```
data "azurerm_lb" "test" {
  name                = "example-lb"
  resource_group_name = "example-resources"
}

output "loadbalancer_id" {
  value = "${data.azurerm_lb.test.id}"
}
```

» Argument Reference

- `name` - (Required) Specifies the name of the Load Balancer.
- `resource_group_name` - (Required) The name of the Resource Group in which the Load Balancer exists.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the Load Balancer.
- `frontend_ip_configuration` - (Optional) A `frontend_ip_configuration` block as documented below.

- `location` - The Azure location where the Load Balancer exists.
- `private_ip_address` - The first private IP address assigned to the load balancer in `frontend_ip_configuration` blocks, if any.
- `private_ip_addresses` - The list of private IP address assigned to the load balancer in `frontend_ip_configuration` blocks, if any.
- `sku` - The SKU of the Load Balancer.
- `tags` - A mapping of tags assigned to the resource.

A `frontend_ip_configuration` block exports the following:

- `name` - The name of the Frontend IP Configuration.
- `subnet_id` - The ID of the Subnet which is associated with the IP Configuration.
- `private_ip_address` - Private IP Address to assign to the Load Balancer.
- `private_ip_address_allocation` - The allocation method for the Private IP Address used by this Load Balancer.
- `public_ip_address_id` - The ID of a Public IP Address which is associated with this Load Balancer.
- `zones` - A list of Availability Zones which the Load Balancer's IP Addresses should be created in.

» Data Source: `azurerm_lb_backend_address_pool`

Use this data source to access information about an existing Load Balancer Backend Address Pool.

» Example Usage

```
data "azurerm_lb" "test" {
  name                       = "example-lb"
  resource_group_name       = "example-resources"
}

data "azurerm_lb_backend_address_pool" "test" {
  name                       = "first"
  loadbalancer_id           = "${data.azurerm_lb.test.id}"
}

output "backend_address_pool_id" {
  value = "${data.azurerm_lb_backend_address_pool.test.id}"
}
```

» Argument Reference

- **name** - (Required) Specifies the name of the Backend Address Pool.
- **loadbalancer_id** - (Required) The ID of the Load Balancer in which the Backend Address Pool exists.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Backend Address Pool.

» Data Source: `azurerm_log_analytics_workspace`

Use this data source to access information about an existing Log Analytics (formerly Operational Insights) Workspace.

» Example Usage

```
data "azurerm_log_analytics_workspace" "test" {
  name                = "acctest-01"
  resource_group_name = "acctest"
}

output "log_analytics_workspace_id" {
  value = "${data.azurerm_log_analytics_workspace.test.workspace_id}"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Log Analytics Workspace.
- **resource_group_name** - (Required) The name of the resource group in which the Log Analytics workspace is located in.

» Attributes Reference

The following attributes are exported:

- **id** - The Azure Resource ID of the Log Analytics Workspace.

- **primary_shared_key** - The Primary shared key for the Log Analytics Workspace.
- **secondary_shared_key** - The Secondary shared key for the Log Analytics Workspace.
- **workspace_id** - The Workspace (or Customer) ID for the Log Analytics Workspace.
- **portal_url** - The Portal URL for the Log Analytics Workspace.
- **sku** - The Sku of the Log Analytics Workspace.
- **retention_in_days** - The workspace data retention in days.
- **tags** - A mapping of tags assigned to the resource.

» Data Source: **azurerm_logic_app_workflow**

Use this data source to access information about an existing Logic App Workflow.

» Example Usage

```
data "azurerm_logic_app_workflow" "test" {
  name                = "workflow1"
  resource_group_name = "my-resource-group"
}

output "access_endpoint" {
  value = "${data.azurerm_logic_app_workflow.test.access_endpoint}"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Logic App Workflow.
- **resource_group_name** - (Required) The name of the Resource Group in which the Logic App Workflow exists.

» Attributes Reference

The following attributes are exported:

- `id` - The Logic App Workflow ID.
- `location` - The Azure location where the Logic App Workflow exists.
- `workflow_schema` - The Schema used for this Logic App Workflow.
- `workflow_version` - The version of the Schema used for this Logic App Workflow. Defaults to 1.0.0.0.
- `parameters` - A map of Key-Value pairs.
- `tags` - A mapping of tags assigned to the resource.
- `access_endpoint` - The Access Endpoint for the Logic App Workflow

» Data Source: `azurerm_managed_disk`

Use this data source to access information about an existing Managed Disk.

» Example Usage

```
data "azurerm_managed_disk" "datasourcemd" {
  name                = "testManagedDisk"
  resource_group_name = "acctestRG"
}

resource "azurerm_virtual_network" "test" {
  name                = "acctvn"
  address_space       = ["10.0.0.0/16"]
  location             = "West US 2"
  resource_group_name = "acctestRG"
}

resource "azurerm_subnet" "test" {
  name                = "acctsub"
  resource_group_name = "acctestRG"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix       = "10.0.2.0/24"
}

resource "azurerm_network_interface" "test" {
  name                = "acctni"
  location             = "West US 2"
  resource_group_name = "acctestRG"

  ip_configuration {
```

```

        name = "testconfiguration1"
        subnet_id = "${azurerm_subnet.test.id}"
        private_ip_address_allocation = "Dynamic"
    }
}

resource "azurerm_virtual_machine" "test" {
    name = "acctvm"
    location = "West US 2"
    resource_group_name = "acctestRG"
    network_interface_ids = ["${azurerm_network_interface.test.id}"]
    vm_size = "Standard_DS1_v2"

    storage_image_reference {
        publisher = "Canonical"
        offer = "UbuntuServer"
        sku = "16.04-LTS"
        version = "latest"
    }

    storage_os_disk {
        name = "myosdisk1"
        caching = "ReadWrite"
        create_option = "FromImage"
        managed_disk_type = "Standard_LRS"
    }

    storage_data_disk {
        name = "datadisk_new"
        managed_disk_type = "Standard_LRS"
        create_option = "Empty"
        lun = 0
        disk_size_gb = "1023"
    }

    storage_data_disk {
        name = "${data.azurerm_managed_disk.datasourcemd.name}"
        managed_disk_id = "${data.azurerm_managed_disk.datasourcemd.id}"
        create_option = "Attach"
        lun = 1
        disk_size_gb = "${data.azurerm_managed_disk.datasourcemd.disk_size_gb}"
    }

    os_profile {
        computer_name = "hostname"
        admin_username = "testadmin"
    }
}

```

```

    admin_password = "Password1234!"
  }

  os_profile_linux_config {
    disable_password_authentication = false
  }

  tags = {
    environment = "staging"
  }
}

```

» Argument Reference

- **name** - (Required) Specifies the name of the Managed Disk.
- **resource_group_name** - (Required) Specifies the name of the resource group.

» Attributes Reference

- **storage_account_type** - The storage account type for the managed disk.
- **source_uri** - The source URI for the managed disk
- **source_resource_id** - ID of an existing managed disk that the current resource was created from.
- **os_type** - The operating system for managed disk. Valid values are **Linux** or **Windows**
- **disk_size_gb** - The size of the managed disk in gigabytes.
- **tags** - A mapping of tags assigned to the resource.
- **zones** - A collection containing the availability zone the managed disk is allocated in.

» Data Source: `azurerm__management__group`

Use this data source to access information about an existing Management Group.

» Example Usage

```

data "azurerm_management_group" "test" {
  group_id = "00000000-0000-0000-0000-000000000000"
}

```

```
output "display_name" {
  value = "${data.azurerm_management_group.test.display_name}"
}
```

» Argument Reference

The following arguments are supported:

- `group_id` - (Required) Specifies the UUID of this Management Group.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the Management Group.
- `display_name` - A friendly name for the Management Group.
- `parent_management_group_id` - The ID of any Parent Management Group.
- `subscription_ids` - A list of Subscription ID's which are assigned to the Management Group.

» Data Source: `azurerm_monitor_action_group`

Use this data source to access the properties of an Action Group.

» Example Usage

```
data "azurerm_monitor_action_group" "example" {
  resource_group_name = "terraform-example-rg"
  name                = "tfex-actiongroup"
}

output "action_group_id" {
  value = "${data.azurerm_monitor_action_group.example.id}"
}
```

» Argument Reference

- **name** - (Required) Specifies the name of the Action Group.
- **resource_group_name** - (Required) Specifies the name of the resource group the Action Group is located in.

» Attributes Reference

- **id** - The ID of the Action Group.
- **short_name** - The short name of the action group.
- **enabled** - Whether this action group is enabled.
- **email_receiver** - One or more **email_receiver** blocks as defined below.
- **sms_receiver** - One or more **sms_receiver** blocks as defined below.
- **webhook_receiver** - One or more **webhook_receiver** blocks as defined below.

email_receiver supports the following:

- **name** - The name of the email receiver.
- **email_address** - The email address of this receiver.

sms_receiver supports the following:

- **name** - The name of the SMS receiver.
- **country_code** - The country code of the SMS receiver.
- **phone_number** - The phone number of the SMS receiver.

webhook_receiver supports the following:

- **name** - The name of the webhook receiver.
- **service_uri** - The URI where webhooks should be sent.

» Data Source: **azurerm__monitor__diagnostic__categories**

Use this data source to access information about the Monitor Diagnostics Categories supported by an existing Resource.

» Example Usage

```
data "azurerm_key_vault" "test" {
  name = "${azurerm_key_vault.test.name}"
}
```

```

    resource_group_name = "${azurerm_key_vault.test.resource_group_name}"
  }

  data "azurerm_monitor_diagnostic_categories" "test" {
    resource_id = "${azurerm_key_vault.test.id}"
  }

```

» Argument Reference

- `resource_id` - (Required) The ID of an existing Resource which Monitor Diagnostics Categories should be retrieved for.

» Attributes Reference

- `id` - The ID of the Resource.
- `logs` - A list of the Log Categories supported for this Resource.
- `metrics` - A list of the Metric Categories supported for this Resource.

» Data Source: `azurerm_monitor_log_profile`

Use this data source to access the properties of a Log Profile.

» Example Usage

```

data "azurerm_monitor_log_profile" "test" {
  name = "test-logprofile"
}

output "log_profile_storage_account_id" {
  value = "${data.azurerm_monitor_log_profile.test.storage_account_id}"
}

```

» Argument Reference

- `name` - (Required) Specifies the Name of the Log Profile.

» Attributes Reference

- `id` - The ID of the Log Profile.
 - `storage_account_id` - The resource id of the storage account in which the Activity Log is stored.
 - `servicebus_rule_id` - The service bus (or event hub) rule ID of the service bus (or event hub) namespace in which the Activity Log is streamed to.
 - `locations` - List of regions for which Activity Log events are stored or streamed.
 - `categories` - List of categories of the logs.
 - `retention_policy` - a `retention_policy` block as documented below.
-

The `retention_policy` block supports:

- `enabled` - A boolean value indicating whether the retention policy is enabled.
- `days` - The number of days for the retention policy.

» Data Source: `azurerm_network_interface`

Use this data source to access information about an existing Network Interface.

» Example Usage

```
data "azurerm_network_interface" "test" {
  name                = "acctest-nic"
  resource_group_name = "networking"
}

output "network_interface_id" {
  value = "${data.azurerm_network_interface.test.id}"
}
```

» Argument Reference

- `name` - (Required) Specifies the name of the Network Interface.
- `resource_group_name` - (Required) Specifies the name of the resource group the Network Interface is located in.

» Attributes Reference

- **id** - The ID of the Network Interface.
- **applied_dns_servers** - List of DNS servers applied to the specified Network Interface.
- **enable_accelerated_networking** - Indicates if accelerated networking is set on the specified Network Interface.
- **enable_ip_forwarding** - Indicate if IP forwarding is set on the specified Network Interface.
- **dns_servers** - The list of DNS servers used by the specified Network Interface.
- **internal_dns_name_label** - The internal dns name label of the specified Network Interface.
- **ip_configuration** - One or more **ip_configuration** blocks as defined below.
- **location** - The location of the specified Network Interface.
- **mac_address** - The MAC address used by the specified Network Interface.
- **network_security_group_id** - The ID of the network security group associated to the specified Network Interface.
- **private_ip_address** - The primary private ip address associated to the specified Network Interface.
- **private_ip_addresses** - The list of private ip addresses associates to the specified Network Interface.
- **tags** - List the tags associated to the specified Network Interface.
- **virtual_machine_id** - The ID of the virtual machine that the specified Network Interface is attached to.

A **ip_configuration** block contains:

- **name** - The name of the IP Configuration.
- **subnet_id** - The ID of the Subnet which the Network Interface is connected to.
- **private_ip_address** - The Private IP Address assigned to this Network Interface.
- **private_ip_address_allocation** - The IP Address allocation type for the Private address, such as **Dynamic** or **Static**.
- **public_ip_address_id** - The ID of the Public IP Address which is connected to this Network Interface.
- **application_gateway_backend_address_pools_ids** - A list of Backend Address Pool ID's within a Application Gateway that this Network Interface is connected to.
- **load_balancer_backend_address_pools_ids** - A list of Backend Address Pool ID's within a Load Balancer that this Network Interface is connected to.
- **load_balancer_inbound_nat_rules_ids** - A list of Inbound NAT Rule

ID's within a Load Balancer that this Network Interface is connected to.

- **primary** - is this the Primary IP Configuration for this Network Interface?

» Data Source: `azurerm__network__security__group`

Use this data source to access information about an existing Network Security Group.

» Example Usage

```
data "azurerm_network_security_group" "test" {
  name = "${azurerm_network_security_group.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

output "location" {
  value = "${data.azurerm_network_security_group.test.location}"
}
```

» Argument Reference

- **name** - (Required) Specifies the Name of the Network Security Group.
- **resource_group_name** - (Required) Specifies the Name of the Resource Group within which the Network Security Group exists

» Attributes Reference

- **id** - The ID of the Network Security Group.
- **location** - The supported Azure location where the resource exists.
- **security_rule** - One or more **security_rule** blocks as defined below.
- **tags** - A mapping of tags assigned to the resource.

The **security_rule** block supports:

- **name** - The name of the security rule.
- **description** - The description for this rule.
- **protocol** - The network protocol this rule applies to.
- **source_port_range** - The Source Port or Range.
- **destination_port_range** - The Destination Port or Range.

- `source_address_prefix` - CIDR or source IP range or * to match any IP.
- `destination_address_prefix` - CIDR or destination IP range or * to match any IP.
- `source_application_security_group_ids` - A List of source Application Security Group ID's
- `destination_application_security_group_ids` - A List of destination Application Security Group ID's
- `access` - Is network traffic is allowed or denied?
- `priority` - The priority of the rule
- `direction` - The direction specifies if rule will be evaluated on incoming or outgoing traffic.

» Data Source: `azurerm_network_watcher`

Use this data source to access information about an existing Network Watcher.

» Example Usage

```
data "azurerm_network_watcher" "test" {
  name           = "${azurerm_network_watcher.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

output "network_watcher_id" {
  value = "${data.azurerm_network_watcher.test.id}"
}
```

» Argument Reference

- `name` - (Required) Specifies the Name of the Network Watcher.
- `resource_group_name` - (Required) Specifies the Name of the Resource Group within which the Network Watcher exists.

» Attributes Reference

- `id` - The ID of the Network Watcher.
- `location` - The supported Azure location where the resource exists.

- **tags** - A mapping of tags assigned to the resource.

» Data Source: `azurerm_notification_hub_namespace`

Use this data source to access information about an existing Notification Hub Namespace.

» Example Usage

```
data "azurerm_notification_hub_namespace" "test" {
  name                = "my-namespace"
  resource_group_name = "my-resource-group"
}

output "servicebus_endpoint" {
  value = "${data.azurerm_notification_hub_namespace.test.servicebus_endpoint}"
}
```

» Argument Reference

- **name** - (Required) Specifies the Name of the Notification Hub Namespace.
- **resource_group_name** - (Required) Specifies the Name of the Resource Group within which the Notification Hub exists.

» Attributes Reference

- **id** - The ID of the Notification Hub Namespace.
- **location** - The Azure Region in which this Notification Hub Namespace exists.
- **namespace_type** - The Type of Namespace, such as `Messaging` or `NotificationHub`.
- **sku** - A `sku` block as defined below.
- **enabled** - Is this Notification Hub Namespace enabled?

A `sku` block exports the following:

- **name** - (Required) The name of the SKU to use for this Notification Hub Namespace. Possible values are `Free`, `Basic` or `Standard`.

» Data Source: `azurerm__platform__image`

Use this data source to access information about a Platform Image.

» Example Usage

```
data "azurerm_platform_image" "test" {
  location = "West Europe"
  publisher = "Canonical"
  offer     = "UbuntuServer"
  sku       = "16.04-LTS"
}

output "version" {
  value = "${data.azurerm_platform_image.test.version}"
}
```

» Argument Reference

- `location` - (Required) Specifies the Location to pull information about this Platform Image from.
- `publisher` - (Required) Specifies the Publisher associated with the Platform Image.
- `offer` - (Required) Specifies the Offer associated with the Platform Image.
- `sku` - (Required) Specifies the SKU of the Platform Image.

» Attributes Reference

- `id` - The ID of the Platform Image.
- `version` - The latest version of the Platform Image.

» Data Source: `azurerm__policy__definition`

Use this data source to access information about a Policy Definition, both custom and built in. Retrieves Policy Definitions from your current subscription by default.

» Example Usage

```
data "azurerm_policy_definition" "test" {
```

```

    display_name = "Allowed resource types"
}

output "id" {
  value = "${data.azurerm_policy_definition.test.id}"
}

```

» Argument Reference

- `display_name` - (Required) Specifies the name of the Policy Definition.
- `management_group_id` - (Optional) Only retrieve Policy Definitions from this Management Group.

» Attributes Reference

- `id` - The ID of the Policy Definition.
- `name` - The Name of the Policy Definition.
- `type` - The Type of Policy.
- `description` - The Description of the Policy.
- `policy_type` - The Type of the Policy, such as `Microsoft.Authorization/policyDefinitions`.
- `policy_rule` - The Rule as defined (in JSON) in the Policy.
- `parameters` - Any Parameters defined in the Policy.
- `metadata` - Any Metadata defined in the Policy.

» Data Source: `azurerm_public_ip`

Use this data source to access information about an existing Public IP Address.

» Example Usage (reference an existing)

```

data "azurerm_public_ip" "test" {
  name                = "name_of_public_ip"
  resource_group_name = "name_of_resource_group"
}

output "domain_name_label" {
  value = "${data.azurerm_public_ip.test.domain_name_label}"
}

output "public_ip_address" {
  value = "${data.azurerm_public_ip.test.ip_address}"
}

```

» Example Usage (Retrieve the Dynamic Public IP of a new VM)

```
resource "azurerm_resource_group" "test" {
  name      = "test-resources"
  location  = "West US 2"
}

resource "azurerm_virtual_network" "test" {
  name            = "test-network"
  address_space   = ["10.0.0.0/16"]
  location        = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_subnet" "test" {
  name                 = "acctsub"
  resource_group_name  = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix       = "10.0.2.0/24"
}

resource "azurerm_public_ip" "test" {
  name                = "test-pip"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  allocation_method   = "Dynamic"
  idle_timeout_in_minutes = 30

  tags = {
    environment = "test"
  }
}

resource "azurerm_network_interface" "test" {
  name            = "test-nic"
  location        = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  ip_configuration {
    name                       = "testconfiguration1"
    subnet_id                 = "${azurerm_subnet.test.id}"
    private_ip_address_allocation = "Static"
    private_ip_address        = "10.0.2.5"
    public_ip_address_id      = "${azurerm_public_ip.test.id}"
  }
}
```

```

    }
}

resource "azurerm_virtual_machine" "test" {
  name                = "test-vm"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  network_interface_ids = ["${azurerm_network_interface.test.id}"]

  # ...
}

data "azurerm_public_ip" "test" {
  name                = "${azurerm_public_ip.test.name}"
  resource_group_name = "${azurerm_virtual_machine.test.resource_group_name}"
}

output "public_ip_address" {
  value = "${data.azurerm_public_ip.test.ip_address}"
}

```

» Argument Reference

- **name** - (Required) Specifies the name of the public IP address.
- **resource_group_name** - (Required) Specifies the name of the resource group.

» Attributes Reference

- **domain_name_label** - The label for the Domain Name.
- **idle_timeout_in_minutes** - Specifies the timeout for the TCP idle connection.
- **fqdn** - Fully qualified domain name of the A DNS record associated with the public IP. This is the concatenation of the domainNameLabel and the regionalized DNS zone.
- **ip_address** - The IP address value that was allocated.
- **ip_version** - The IP version being used, for example IPv4 or IPv6.
- **tags** - A mapping of tags to assigned to the resource.

» Data Source: `azurerm_public_ips`

Use this data source to access information about a set of existing Public IP Addresses.

» Example Usage

```
data "azurerm_public_ips" "test" {
  resource_group_name = "pip-test"
  attached             = false
}
```

» Argument Reference

- `resource_group_name` - (Required) Specifies the name of the resource group.
- `attached` - (Optional) Filter to include IP Addresses which are attached to a device, such as a VM/LB (**true**) or unattached (**false**).
- `name_prefix` - (Optional) A prefix match used for the IP Addresses **name** field, case sensitive.
- `allocation_type` - (Optional) The Allocation Type for the Public IP Address. Possible values include **Static** or **Dynamic**.

» Attributes Reference

- `public_ips` - A List of `public_ips` blocks as defined below filtered by the criteria above.

A `public_ips` block contains:

- `id` - The ID of the Public IP Address
- `domain_name_label` - The Domain Name Label of the Public IP Address
- `fqdn` - The FQDN of the Public IP Address
- `name` - The Name of the Public IP Address

» Data Source: `azurerm_recovery_services_vault`

Use this data source to access information about an existing Recovery Services Vault.

» Example Usage

```
data "azurerm_recovery_services_vault" "vault" {  
  name                = "tfex-recovery_vault"  
  resource_group_name = "tfex-resource_group"  
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Recovery Services Vault.
- **resource_group_name** - (Required) The name of the resource group in which the Recovery Services Vault resides.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Recovery Services Vault.
- **location** - The Azure location where the resource resides.
- **tags** - A mapping of tags assigned to the resource.
- **sku** - The vault's current SKU.

» Data Source: `azurerm_recovery_services_protection_policy_vm`

Use this data source to access information about an existing Recovery Services VM Protection Policy.

» Example Usage

```
data "azurerm_recovery_services_protection_policy_vm" "policy" {  
  name                = "policy"  
  recovery_vault_name = "recovery_vault"  
  resource_group_name = "resource_group"  
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Recovery Services VM Protection Policy.
- **recovery_vault_name** - (Required) Specifies the name of the Recovery Services Vault.
- **resource_group_name** - (Required) The name of the resource group in which the Recovery Services VM Protection Policy resides.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Recovery Services VM Protection Policy.
- **tags** - A mapping of tags assigned to the resource.

» Data Source: `azurerm_resource_group`

Use this data source to access information about an existing Resource Group.

» Example Usage

```
data "azurerm_resource_group" "test" {
  name = "dsrg_test"
}

resource "azurerm_managed_disk" "test" {
  name                = "managed_disk_name"
  location            = "${data.azurerm_resource_group.test.location}"
  resource_group_name = "${data.azurerm_resource_group.test.name}"
  storage_account_type = "Standard_LRS"
  create_option       = "Empty"
  disk_size_gb        = "1"
}
```

» Argument Reference

- **name** - (Required) Specifies the name of the resource group.

NOTE: If the specified location doesn't match the actual resource group location, an error message with the actual location value will be shown.

» Attributes Reference

- `location` - The location of the resource group.
- `tags` - A mapping of tags assigned to the resource group.

» Data Source: `azurerm_role_definition`

Use this data source to access information about an existing Role Definition.

» Example Usage

```
data "azurerm_subscription" "primary" {}

resource "azurerm_role_definition" "custom" {
  role_definition_id = "00000000-0000-0000-0000-000000000000"
  name               = "CustomRoleDef"
  scope              = "${data.azurerm_subscription.primary.id}"

  #...
}

data "azurerm_role_definition" "custom" {
  role_definition_id = "${azurerm_role_definition.custom.role_definition_id}"
  scope              = "${data.azurerm_subscription.primary.id}" # /subscriptions/00000000-0000-0000-0000-000000000000
}

data "azurerm_role_definition" "custom-byname" {
  name = "${azurerm_role_definition.custom.name}"
  scope = "${data.azurerm_subscription.primary.id}"
}

data "azurerm_builtin_role_definition" "builtin" {
  name = "Contributor"
}

output "custom_role_definition_id" {
  value = "${data.azurerm_role_definition.custom.id}"
}

output "contributor_role_definition_id" {
```

```

    value = "${data.azurerm_role_definition.builtin.id}"
  }

```

» Argument Reference

- **name** - (Optional) Specifies the Name of either a built-in or custom Role Definition.

You can also use this for built-in roles such as **Contributor**, **Owner**, **Reader** and **Virtual Machine Contributor**

- **role_definition_id** - (Optional) Specifies the ID of the Role Definition as a UUID/GUID.
- **scope** - (Optional) Specifies the Scope at which the Custom Role Definition exists.

NOTE: One of **name** or **role_definition_id** must be specified.

» Attributes Reference

- **id** - the ID of the built-in Role Definition.
- **description** - the Description of the built-in Role.
- **type** - the Type of the Role.
- **permissions** - a **permissions** block as documented below.
- **assignable_scopes** - One or more assignable scopes for this Role Definition, such as `/subscriptions/0b1f6471-1bf0-4dda-aec3-111122223333`, `/subscriptions/0b1f6471-1bf0-4dda-aec3-111122223333/resourceGroups/myGroup`, or `/subscriptions/0b1f6471-1bf0-4dda-aec3-111122223333/resourceGroups/myGroup/providers/`

A **permissions** block contains:

- **actions** - a list of actions supported by this role
- **not_actions** - a list of actions which are denied by this role

» Data Source: `azurerm_route_table`

Use this data source to access information about an existing Route Table.

» Example Usage

```

data "azurerm_route_table" "test" {
  name                = "myroutetable"
  resource_group_name = "some-resource-group"
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Route Table.
- **resource_group_name** - (Required) The name of the Resource Group in which the Route Table exists.

» Attributes Reference

The following attributes are exported:

- **id** - The Route Table ID.
- **location** - The Azure Region in which the Route Table exists.
- **route** - One or more **route** blocks as documented below.
- **subnets** - The collection of Subnets associated with this route table.
- **tags** - A mapping of tags assigned to the Route Table.

The **route** block exports the following:

- **name** - The name of the Route.
- **address_prefix** - The destination CIDR to which the route applies.
- **next_hop_type** - The type of Azure hop the packet should be sent to.
- **next_hop_in_ip_address** - Contains the IP address packets should be forwarded to.

» Data Source: `azurerm_scheduler_job_collection`

Use this data source to access information about an existing Scheduler Job Collection.

NOTE: Support for Scheduler Job Collections has been deprecated by Microsoft in favour of Logic Apps (more information can be found at [this link](#)) - as such we plan to remove support for this data source as a part of version 2.0 of the AzureRM Provider.

» Example Usage

```
data "azurerm_scheduler_job_collection" "test" {
  name                       = "tfex-job-collection"
  resource_group_name       = "tfex-job-collection-rg"
```

```

}

output "job_collection_state" {
  value = "${data.azurearm_scheduler_job_collection.jobs.state}"
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Scheduler Job Collection.
- **resource_group_name** - (Required) Specifies the name of the resource group in which the Scheduler Job Collection resides.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Scheduler Job Collection.
- **location** - The Azure location where the resource exists.
- **tags** - A mapping of tags assigned to the resource.
- **sku** - The Job Collection's pricing level's SKU.
- **state** - The Job Collection's state.
- **quota** - The Job collection quotas as documented in the **quota** block below.

The **quota** block supports:

- **max_job_count** - Sets the maximum number of jobs in the collection.
- **max_recurrence_frequency** - The maximum frequency of recurrence.
- **max_retry_interval** - The maximum interval between retries.

» Data Source: `azurerm_servicebus_namespace`

Use this data source to access information about an existing ServiceBus Namespace.

» Example Usage

```
data "azurerm_servicebus_namespace" "test" {
  name                = "examplnamespace"
  resource_group_name = "example-resources"
}

output "location" {
  value = "${data.azurerm_servicebus_namespace.test.location}"
}
```

» Argument Reference

- **name** - (Required) Specifies the name of the ServiceBus Namespace.
- **resource_group_name** - (Required) Specifies the name of the Resource Group where the ServiceBus Namespace exists.

» Attributes Reference

- **location** - The location of the Resource Group in which the ServiceBus Namespace exists.
- **sku** - The Tier used for the ServiceBus Namespace.
- **capacity** - The capacity of the ServiceBus Namespace.
- **tags** - A mapping of tags assigned to the resource.

The following attributes are exported only if there is an authorization rule named `RootManageSharedAccessKey` which is created automatically by Azure.

- **default_primary_connection_string** - The primary connection string for the authorization rule `RootManageSharedAccessKey`.
- **default_secondary_connection_string** - The secondary connection string for the authorization rule `RootManageSharedAccessKey`.
- **default_primary_key** - The primary access key for the authorization rule `RootManageSharedAccessKey`.
- **default_secondary_key** - The secondary access key for the authorization rule `RootManageSharedAccessKey`.

» Data Source: `azurerm_shared_image`

Use this data source to access information about an existing Shared Image within a Shared Image Gallery.

NOTE Shared Image Galleries are currently in Public Preview. You can find more information, including how to register for the Public Preview [here](#).

» Example Usage

```
data "azurerm_shared_image" "test" {
  name                = "my-image"
  gallery_name        = "my-image-gallery"
  resource_group_name = "example-resources"
}
```

» Argument Reference

The following arguments are supported:

- `name` - (Required) The name of the Shared Image.
- `gallery_name` - (Required) The name of the Shared Image Gallery in which the Shared Image exists.
- `resource_group_name` - (Required) The name of the Resource Group in which the Shared Image Gallery exists.

» Attributes Reference

The following attributes are exported:

- `id` - The Resource ID of the Shared Image.
- `description` - The description of this Shared Image.
- `eula` - The End User Licence Agreement for the Shared Image.
- `location` - The supported Azure location where the Shared Image Gallery exists.
- `identifier` - An `identifier` block as defined below.
- `os_type` - The type of Operating System present in this Shared Image.
- `privacy_statement_uri` - The URI containing the Privacy Statement for this Shared Image.

- **release_note_uri** - The URI containing the Release Notes for this Shared Image.
 - **tags** - A mapping of tags assigned to the Shared Image.
-

A **identifier** block exports the following:

- **offer** - The Offer Name for this Shared Image.
- **publisher** - The Publisher Name for this Gallery Image.
- **sku** - The Name of the SKU for this Gallery Image.

» Data Source: **azurerm_shared_image_gallery**

Use this data source to access information about an existing Shared Image Gallery.

NOTE Shared Image Galleries are currently in Public Preview. You can find more information, including how to register for the Public Preview [here](#).

» Example Usage

```
data "azurerm_shared_image_gallery" "test" {
  name                = "my-image-gallery"
  resource_group_name = "example-resources"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Shared Image Gallery.
- **resource_group_name** - (Required) The name of the Resource Group in which the Shared Image Gallery exists.

» Attributes Reference

The following attributes are exported:

- **id** - The Resource ID of the Shared Image Gallery.
- **description** - A description for the Shared Image Gallery.

- **unique_name** - The unique name assigned to the Shared Image Gallery.
- **tags** - A mapping of tags which are assigned to the Shared Image Gallery.

» Data Source: **azurerm_shared_image_version**

Use this data source to access information about an existing Version of a Shared Image within a Shared Image Gallery.

NOTE Shared Image Galleries are currently in Public Preview. You can find more information, including how to register for the Public Preview [here](#).

» Example Usage

```
data "azurerm_shared_image_version" "test" {
  name           = "1.0.0"
  image_name     = "my-image"
  gallery_name   = "my-image-gallery"
  resource_group_name = "example-resources"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Image Version.
- **image_name** - (Required) The name of the Shared Image in which this Version exists.
- **gallery_name** - (Required) The name of the Shared Image in which the Shared Image exists.
- **resource_group_name** - (Required) The name of the Resource Group in which the Shared Image Gallery exists.

» Attributes Reference

The following attributes are exported:

- **id** - The Resource ID of the Shared Image.
- **exclude_from_latest** - Is this Image Version excluded from the **latest** filter?

- **location** - The supported Azure location where the Shared Image Gallery exists.
- **managed_image_id** - The ID of the Managed Image which was the source of this Shared Image Version.
- **target_region** - One or more **target_region** blocks as documented below.
- **tags** - A mapping of tags assigned to the Shared Image.

The **target_region** block exports the following:

- **name** - The Azure Region in which this Image Version exists.
- **regional_replica_count** - The number of replicas of the Image Version to be created per region.

» Data Source: **azurerm_stream_analytics_job**

Use this data source to access information about an existing Stream Analytics Job.

» Example Usage

```
data "azurerm_stream_analytics_job" "test" {
  name                = "example-job"
  resource_group_name = "example-resources"
}

output "job_id" {
  value = "${data.azurerm_stream_analytics_job.test.job_id}"
}
```

» Argument Reference

- **name** - (Required) Specifies the name of the Stream Analytics Job.
- **resource_group_name** - (Required) Specifies the name of the resource group the Stream Analytics Job is located in.

» Attributes Reference

- `id` - The ID of the Stream Analytics Job.
- `compatibility_level` - The compatibility level for this job.
- `data_locale` - The Data Locale of the Job.
- `events_late_arrival_max_delay_in_seconds` - The maximum tolerable delay in seconds where events arriving late could be included.
- `events_out_of_order_max_delay_in_seconds` - The maximum tolerable delay in seconds where out-of-order events can be adjusted to be back in order.
- `events_out_of_order_policy` - The policy which should be applied to events which arrive out of order in the input event stream.
- `job_id` - The Job ID assigned by the Stream Analytics Job.
- `location` - The Azure location where the Stream Analytics Job exists.
- `output_error_policy` - The policy which should be applied to events which arrive at the output and cannot be written to the external storage due to being malformed (such as missing column values, column values of wrong type or size).
- `streaming_units` - The number of streaming units that the streaming job uses.
- `transformation_query` - The query that will be run in the streaming job, written in Stream Analytics Query Language (SAQL).

» Data Source: `azurerm_storage_account`

Use this data source to access information about an existing Storage Account.

» Example Usage

```
data "azurerm_storage_account" "test" {
  name                = "packerimages"
  resource_group_name = "packer-storage"
}

output "storage_account_tier" {
  value = "${data.azurerm_storage_account.test.account_tier}"
}
```

» Argument Reference

- **name** - (Required) Specifies the name of the Storage Account
- **resource_group_name** - (Required) Specifies the name of the resource group the Storage Account is located in.

» Attributes Reference

- **id** - The ID of the Storage Account.
- **location** - The Azure location where the Storage Account exists
- **account_kind** - The Kind of account.
- **account_tier** - The Tier of this storage account.
- **account_replication_type** - The type of replication used for this storage account.
- **access_tier** - The access tier for **BlobStorage** accounts.
- **enable_blob_encryption** - Are Encryption Services are enabled for Blob storage? See here for more information.
- **enable_file_encryption** - Are Encryption Services are enabled for File storage? See here for more information.
- **enable_https_traffic_only** - Is traffic only allowed via HTTPS? See here for more information.
- **is_hns_enabled** - Is Hierarchical Namespace enabled?
- **account_encryption_source** - The Encryption Source for this Storage Account.
- **custom_domain** - A **custom_domain** block as documented below.
- **tags** - A mapping of tags to assigned to the resource.
- **primary_location** - The primary location of the Storage Account.
- **secondary_location** - The secondary location of the Storage Account.
- **primary_blob_endpoint** - The endpoint URL for blob storage in the primary location.
- **primary_blob_host** - The hostname with port if applicable for blob storage in the primary location.
- **secondary_blob_endpoint** - The endpoint URL for blob storage in the secondary location.
- **secondary_blob_host** - The hostname with port if applicable for blob storage in the secondary location.

- **primary_queue_endpoint** - The endpoint URL for queue storage in the primary location.
- **primary_queue_host** - The hostname with port if applicable for queue storage in the primary location.
- **secondary_queue_endpoint** - The endpoint URL for queue storage in the secondary location.
- **secondary_queue_host** - The hostname with port if applicable for queue storage in the secondary location.
- **primary_table_endpoint** - The endpoint URL for table storage in the primary location.
- **primary_table_host** - The hostname with port if applicable for table storage in the primary location.
- **secondary_table_endpoint** - The endpoint URL for table storage in the secondary location.
- **secondary_table_host** - The hostname with port if applicable for table storage in the secondary location.
- **primary_file_endpoint** - The endpoint URL for file storage in the primary location.
- **primary_file_host** - The hostname with port if applicable for file storage in the primary location.
- **secondary_file_endpoint** - The endpoint URL for file storage in the secondary location.
- **secondary_file_host** - The hostname with port if applicable for file storage in the secondary location.
- **primary_dfs_endpoint** - The endpoint URL for DFS storage in the primary location.
- **primary_dfs_host** - The hostname with port if applicable for DFS storage in the primary location.
- **secondary_dfs_endpoint** - The endpoint URL for DFS storage in the secondary location.
- **secondary_dfs_host** - The hostname with port if applicable for DFS storage in the secondary location.
- **primary_web_endpoint** - The endpoint URL for web storage in the primary location.
- **primary_web_host** - The hostname with port if applicable for web storage in the primary location.

- `secondary_web_endpoint` - The endpoint URL for web storage in the secondary location.
- `secondary_web_host` - The hostname with port if applicable for web storage in the secondary location.
- `primary_access_key` - The primary access key for the Storage Account.
- `secondary_access_key` - The secondary access key for the Storage Account.
- `primary_connection_string` - The connection string associated with the primary location
- `secondary_connection_string` - The connection string associated with the secondary location
- `primary_blob_connection_string` - The connection string associated with the primary blob location
- `secondary_blob_connection_string` - The connection string associated with the secondary blob location

-
- `custom_domain` supports the following:
 - `name` - The Custom Domain Name used for the Storage Account.

» Data Source: `azurerm_storage_account_sas`

Use this data source to obtain a Shared Access Signature (SAS Token) for an existing Storage Account.

Shared access signatures allow fine-grained, ephemeral access control to various aspects of an Azure Storage Account.

Note that this is an Account SAS and *not* a Service SAS.

» Example Usage

```
resource "azurerm_resource_group" "testrg" {
  name      = "resourceGroupName"
  location  = "westus"
}

resource "azurerm_storage_account" "testsa" {
  name                = "storageaccountname"
  resource_group_name = "${azurerm_resource_group.testrg.name}"
}
```

```

location                = "westus"
account_tier             = "Standard"
account_replication_type = "GRS"

tags = {
  environment = "staging"
}
}

data "azurerm_storage_account_sas" "test" {
  connection_string = "${azurerm_storage_account.testsa.primary_connection_string}"
  https_only        = true

  resource_types {
    service  = true
    container = false
    object   = false
  }

  services {
    blob  = true
    queue = false
    table = false
    file  = false
  }

  start  = "2018-03-21"
  expiry = "2020-03-21"

  permissions {
    read      = true
    write     = true
    delete   = false
    list      = false
    add       = true
    create    = true
    update    = false
    process   = false
  }
}

output "sas_url_query_string" {
  value = "${data.azurerm_storage_account_sas.test.sas}"
}

```


» Argument Reference

- **connection_string** - (Required) The connection string for the storage account to which this SAS applies. Typically directly from the **primary_connection_string** attribute of a terraform created **azurerm_storage_account** resource.
- **https_only** - (Optional) Only permit **https** access. If **false**, both **http** and **https** are permitted. Defaults to **true**.
- **resource_types** - (Required) A **resource_types** block as defined below.
- **services** - (Required) A **services** block as defined below.
- **start** - (Required) The starting time and date of validity of this SAS. Must be a valid ISO-8601 format time/date string.
- **expiry** - (Required) The expiration time and date of this SAS. Must be a valid ISO-8601 format time/date string.
- **permissions** - (Required) A **permissions** block as defined below.

resource_types is a set of **true/false** flags which define the storage account resource types that are granted access by this SAS. This can be thought of as the scope over which the permissions apply. A **service** will have larger scope (affecting all sub-resources) than **object**.

A **resource_types** block contains:

- **service** - (Required) Should permission be granted to the entire service?
- **container** - (Required) Should permission be granted to the container?
- **object** - (Required) Should permission be granted only to a specific object?

services is a set of **true/false** flags which define the storage account services that are granted access by this SAS.

A **services** block contains:

- **blob** - (Required) Should permission be granted to **blob** services within this storage account?
- **queue** - (Required) Should permission be granted to **queue** services within this storage account?
- **table** - (Required) Should permission be granted to **table** services within this storage account?
- **file** - (Required) Should permission be granted to **file** services within this storage account?

A **permissions** block contains:

- **read** - (Required) Should Read permissions be enabled for this SAS?

- **write** - (Required) Should Write permissions be enabled for this SAS?
- **delete** - (Required) Should Delete permissions be enabled for this SAS?
- **list** - (Required) Should List permissions be enabled for this SAS?
- **add** - (Required) Should Add permissions be enabled for this SAS?
- **create** - (Required) Should Create permissions be enabled for this SAS?
- **update** - (Required) Should Update permissions be enabled for this SAS?
- **process** - (Required) Should Process permissions be enabled for this SAS?

Refer to the SAS creation reference from Azure for additional details on the fields above.

» Attributes Reference

- **sas** - The computed Account Shared Access Signature (SAS).

» Data Source: `azurerm_subnet`

Use this data source to access information about an existing Subnet within a Virtual Network.

» Example Usage

```
data "azurerm_subnet" "test" {
  name                       = "backend"
  virtual_network_name      = "production"
  resource_group_name       = "networking"
}

output "subnet_id" {
  value = "${data.azurerm_subnet.test.id}"
}
```

» Argument Reference

- **name** - (Required) Specifies the name of the Subnet.
- **virtual_network_name** - (Required) Specifies the name of the Virtual Network this Subnet is located within.
- **resource_group_name** - (Required) Specifies the name of the resource group the Virtual Network is located in.

» Attributes Reference

- `id` - The ID of the Subnet.
- `address_prefix` - The address prefix used for the subnet.
- `network_security_group_id` - The ID of the Network Security Group associated with the subnet.
- `route_table_id` - The ID of the Route Table associated with this subnet.
- `ip_configurations` - The collection of IP Configurations with IPs within this subnet.
- `service_endpoints` - A list of Service Endpoints within this subnet.

» Data Source: `azurerm_subscription`

Use this data source to access information about an existing Subscription.

» Example Usage

```
data "azurerm_subscription" "current" {}

output "current_subscription_display_name" {
  value = "${data.azurerm_subscription.current.display_name}"
}
```

» Argument Reference

- `subscription_id` - (Optional) Specifies the ID of the subscription. If this argument is omitted, the subscription ID of the current Azure Resource Manager provider is used.

» Attributes Reference

- `id` - The ID of the Subscription.
- `display_name` - The subscription display name.
- `state` - The subscription state. Possible values are Enabled, Warned, PastDue, Disabled, and Deleted.
- `location_placement_id` - The subscription location placement ID.
- `quota_id` - The subscription quota ID.
- `spending_limit` - The subscription spending limit.

» Data Source: `azurerm__subscriptions`

Use this data source to access information about all the Subscriptions currently available.

» Example Usage

```
data "azurerm_subscriptions" "available" {}

output "available_subscriptions" {
  value = "${data.azurerm_subscriptions.available.subscriptions}"
}

output "first_available_subscription_display_name" {
  value = "${data.azurerm_subscriptions.available.subscriptions.0.display_name}"
}
```

» Argument Reference

- `display_name_prefix` - (Optional) A case-insensitive prefix which can be used to filter on the `display_name` field
- `display_name_contains` - (Optional) A case-insensitive value which must be contained within the `display_name` field, used to filter the results

» Attributes Reference

- `subscriptions` - One or more `subscription` blocks as defined below.

The `subscription` block contains:

- `display_name` - The subscription display name.
- `state` - The subscription state. Possible values are Enabled, Warned, PastDue, Disabled, and Deleted.
- `location_placement_id` - The subscription location placement ID.
- `quota_id` - The subscription quota ID.
- `spending_limit` - The subscription spending limit.

» Data Source: `azurerm__traffic__manager__geographical__location`

Use this data source to access the ID of a specified Traffic Manager Geographical Location within the Geographical Hierarchy.

» Example Usage (World)

```
data "azurerm_traffic_manager_geographical_location" "test" {
  name = "World"
}

output "location_code" {
  value = "${data.azurerm_traffic_manager_geographical_location.test.id}"
}
```

» Argument Reference

- **name** - (Required) Specifies the name of the Location, for example World, Europe or Germany.

» Attributes Reference

- **id** - The ID of this Location, also known as the Code of this Location.

» Data Source: `azurerm_virtual_machine`

Use this data source to access information about an existing Virtual Machine.

» Example Usage

```
data "azurerm_virtual_machine" "test" {
  name                = "production"
  resource_group_name = "networking"
}

output "virtual_machine_id" {
  value = "${data.azurerm_virtual_machine.test.id}"
}
```

» Argument Reference

- **name** - (Required) Specifies the name of the Virtual Machine.
- **resource_group_name** - (Required) Specifies the name of the resource group the Virtual Machine is located in.

» Attributes Reference

- `id` - The ID of the Virtual Machine.

» Data Source: `azurerm_virtual_network`

Use this data source to access information about an existing Virtual Network.

» Example Usage

```
data "azurerm_virtual_network" "test" {
  name                = "production"
  resource_group_name = "networking"
}

output "virtual_network_id" {
  value = "${data.azurerm_virtual_network.test.id}"
}
```

» Argument Reference

- `name` - (Required) Specifies the name of the Virtual Network.
- `resource_group_name` - (Required) Specifies the name of the resource group the Virtual Network is located in.

» Attributes Reference

- `id` - The ID of the virtual network.
- `address_spaces` - The list of address spaces used by the virtual network.
- `dns_servers` - The list of DNS servers used by the virtual network.
- `subnets` - The list of name of the subnets that are attached to this virtual network.
- `vnet_peerings` - A mapping of name - virtual network id of the virtual network peerings.

» Data Source: `azurerm_virtual_network_gateway`

Use this data source to access information about an existing Virtual Network Gateway.

» Example Usage

```
data "azurerm_virtual_network_gateway" "test" {
  name                = "production"
  resource_group_name = "networking"
}

output "virtual_network_gateway_id" {
  value = "${data.azurerm_virtual_network_gateway.test.id}"
}
```

» Argument Reference

- **name** - (Required) Specifies the name of the Virtual Network Gateway.
- **resource_group_name** - (Required) Specifies the name of the resource group the Virtual Network Gateway is located in.

» Attributes Reference

- **id** - The ID of the Virtual Network Gateway.
- **location** - The location/region where the Virtual Network Gateway is located.
- **type** - The type of the Virtual Network Gateway.
- **vpn_type** - The routing type of the Virtual Network Gateway.
- **enable_bgp** - Will BGP (Border Gateway Protocol) will be enabled for this Virtual Network Gateway.
- **active_active** - (Optional) Is this an Active-Active Gateway?
- **default_local_network_gateway_id** - The ID of the local network gateway through which outbound Internet traffic from the virtual network in which the gateway is created will be routed (*forced tunneling*). Refer to the Azure documentation on forced tunneling.
- **sku** - Configuration of the size and capacity of the Virtual Network Gateway.
- **ip_configuration** - One or two `ip_configuration` blocks documented below.
- **vpn_client_configuration** - A `vpn_client_configuration` block which is documented below.
- **tags** - A mapping of tags assigned to the resource.

The `ip_configuration` block supports:

- `name` - A user-defined name of the IP configuration.
- `private_ip_address_allocation` - Defines how the private IP address of the gateways virtual interface is assigned.
- `subnet_id` - The ID of the gateway subnet of a virtual network in which the virtual network gateway will be created. It is mandatory that the associated subnet is named `GatewaySubnet`. Therefore, each virtual network can contain at most a single Virtual Network Gateway.
- `public_ip_address_id` - The ID of the Public IP Address associated with the Virtual Network Gateway.

The `vpn_client_configuration` block supports:

- `address_space` - The address space out of which ip addresses for vpn clients will be taken. You can provide more than one address space, e.g. in CIDR notation.
- `root_certificate` - One or more `root_certificate` blocks which are defined below. These root certificates are used to sign the client certificate used by the VPN clients to connect to the gateway.
- `revoked_certificate` - One or more `revoked_certificate` blocks which are defined below.
- `radius_server_address` - (Optional) The address of the Radius server. This setting is incompatible with the use of `root_certificate` and `revoked_certificate`.
- `radius_server_secret` - (Optional) The secret used by the Radius server. This setting is incompatible with the use of `root_certificate` and `revoked_certificate`.
- `vpn_client_protocols` - (Optional) List of the protocols supported by the vpn client. The supported values are `SSTP`, `IkeV2` and `OpenVPN`.

The `bgp_settings` block supports:

- `asn` - The Autonomous System Number (ASN) to use as part of the BGP.
- `peering_address` - The BGP peer IP address of the virtual network gateway. This address is needed to configure the created gateway as a BGP Peer on the on-premises VPN devices.
- `peer_weight` - The weight added to routes which have been learned through BGP peering.

The `root_certificate` block supports:

- `name` - The user-defined name of the root certificate.

- `public_cert_data` - The public certificate of the root certificate authority. The certificate must be provided in Base-64 encoded X.509 format (PEM).

The `root_revoked_certificate` block supports:

- `name` - The user-defined name of the revoked certificate.
- `public_cert_data` - The SHA1 thumbprint of the certificate to be revoked.

» **azurerm_resource_group**

Manages a resource group on Azure.

» **Example Usage**

```
resource "azurerm_resource_group" "test" {
  name      = "testResourceGroup1"
  location  = "West US"

  tags = {
    environment = "Production"
  }
}
```

» **Argument Reference**

The following arguments are supported:

- `name` - (Required) The name of the resource group. Must be unique on your Azure subscription.
- `location` - (Required) The location where the resource group should be created. For a list of all Azure locations, please consult this link or run `az account list-locations --output table`.
- `tags` - (Optional) A mapping of tags to assign to the resource.

» **Attributes Reference**

The following attributes are exported:

- `id` - The resource group ID.

» Import

Resource Groups can be imported using the `resource id`, e.g.

```
terraform import azurerm_resource_group.mygroup /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_api_management

Manages an API Management Service.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "example-resources"
  location  = "West Europe"
}

resource "azurerm_api_management" "test" {
  name                = "example-apim"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  publisher_name       = "My Company"
  publisher_email      = "company@terraform.io"

  sku {
    name      = "Developer"
    capacity = 1
  }
}
```

» Argument Reference

The following arguments are supported:

- `name` - (Required) The name of the API Management Service. Changing this forces a new resource to be created.
- `location` - (Required) The Azure location where the API Management Service exists. Changing this forces a new resource to be created.
- `resource_group_name` - (Required) The name of the Resource Group in which the API Management Service should be exist. Changing this forces a new resource to be created.
- `publisher_name` - (Required) The name of publisher/company.

- **publisher_email** - (Required) The email of publisher/company.
- **sku** - (Required) A **sku** block as documented below.

-
- **additional_location** - (Optional) One or more **additional_location** blocks as defined below.
 - **certificate** - (Optional) One or more (up to 10) **certificate** blocks as defined below.
 - **identity** - (Optional) An **identity** block is documented below.
 - **hostname_configuration** - (Optional) A **hostname_configuration** block as defined below.
 - **notification_sender_email** - (Optional) Email address from which the notification will be sent.
 - **policy** - (Optional) A **policy** block as defined below.
 - **security** - (Optional) A **security** block as defined below.
 - **sign_in** - (Optional) A **sign_in** block as defined below.
 - **sign_up** - (Optional) A **sign_up** block as defined below.
 - **tags** - (Optional) A mapping of tags assigned to the resource.

A **additional_location** block supports the following:

- **location** - (Required) The name of the Azure Region in which the API Management Service should be expanded to.

A **certificate** block supports the following:

- **encoded_certificate** - (Required) The Base64 Encoded PFX Certificate.
- **certificate_password** - (Required) The password for the certificate.
- **store_name** - (Required) The name of the Certificate Store where this certificate should be stored. Possible values are **CertificateAuthority** and **Root**.

A **hostname_configuration** block supports the following:

- **management** - (Optional) One or more **management** blocks as documented below.
- **portal** - (Optional) One or more **portal** blocks as documented below.

- **proxy** - (Optional) One or more **proxy** blocks as documented below.
 - **scm** - (Optional) One or more **scm** blocks as documented below.
-

A **identity** block supports the following:

- **type** - (Required) Specifies the type of Managed Service Identity that should be configured on this API Management Service. At this time the only supported value is **SystemAssigned**.
-

A **management**, **portal** and **scm** block supports the following:

- **host_name** - (Required) The Hostname to use for the Management API.
- **key_vault_id** - (Optional) The ID of the Key Vault Secret containing the SSL Certificate, which must be of the type **application/x-pkcs12**.

NOTE: Setting this field requires the **identity** block to be specified, since this identity is used for to retrieve the Key Vault Certificate. Auto-updating the Certificate from the Key Vault requires the Secret version isn't specified.

- **certificate** - (Optional) The Base64 Encoded Certificate.
- **certificate_password** - (Optional) The password associated with the certificate provided above.

NOTE: Either **key_vault_id** or **certificate** and **certificate_password** must be specified.

- **negotiate_client_certificate** - (Optional) Should Client Certificate Negotiation be enabled for this Hostname? Defaults to **false**.
-

A **policy** block supports the following:

- **xml_content** - (Optional) The XML Content for this Policy.
 - **xml_link** - (Optional) A link to an API Management Policy XML Document, which must be publicly available.
-

A **proxy** block supports the following:

- **default_ssl_binding** - (Optional) Is the certificate associated with this Hostname the Default SSL Certificate? This is used when an SNI header isn't specified by a client. Defaults to **false**.
- **host_name** - (Required) The Hostname to use for the Management API.

- **key_vault_id** - (Optional) The ID of the Key Vault Secret containing the SSL Certificate, which must be of the type `application/x-pkcs12`.

NOTE: Setting this field requires the **identity** block to be specified, since this identity is used for to retrieve the Key Vault Certificate. Auto-updating the Certificate from the Key Vault requires the Secret version isn't specified.

- **certificate** - (Optional) The Base64 Encoded Certificate.
- **certificate_password** - (Optional) The password associated with the certificate provided above.

NOTE: Either **key_vault_id** or **certificate** and **certificate_password** must be specified.

- **negotiate_client_certificate** - (Optional) Should Client Certificate Negotiation be enabled for this Hostname? Defaults to **false**.

A **security** block supports the following:

- **disable_backend_ssl30** - (Optional) Should SSL 3.0 be disabled on the backend of the gateway? Defaults to **false**.

info: This maps to the `Microsoft.WindowsAzure.ApiManagement.Gateway.Security.Backend.Protocols.Ssl30` field

- **disable_backend_tls10** - (Optional) Should TLS 1.0 be disabled on the backend of the gateway? Defaults to **false**.

info: This maps to the `Microsoft.WindowsAzure.ApiManagement.Gateway.Security.Backend.Protocols.Tls10` field

- **disable_backend_tls11** - (Optional) Should TLS 1.1 be disabled on the backend of the gateway? Defaults to **false**.

info: This maps to the `Microsoft.WindowsAzure.ApiManagement.Gateway.Security.Backend.Protocols.Tls11` field

- **disable_frontend_ssl30** - (Optional) Should SSL 3.0 be disabled on the frontend of the gateway? Defaults to **false**.

info: This maps to the `Microsoft.WindowsAzure.ApiManagement.Gateway.Security.Protocols.Ssl30` field

- **disable_frontend_tls10** - (Optional) Should TLS 1.0 be disabled on the frontend of the gateway? Defaults to **false**.

info: This maps to the `Microsoft.WindowsAzure.ApiManagement.Gateway.Security.Protocols.Tls10` field

- **disable_frontend_tls11** - (Optional) Should TLS 1.1 be disabled on the frontend of the gateway? Defaults to **false**.

info: This maps to the `Microsoft.WindowsAzure.ApiManagement.Gateway.Security.Protocols.Tls11` field

- **disable_triple_des_chipers** - (Optional) Should the `TLS_RSA_WITH_3DES_EDE_CBC_SHA` cipher be disabled for all TLS versions (1.0, 1.1 and 1.2)? Defaults to `false`.

info: This maps to the `Microsoft.WindowsAzure.ApiManagement.Gateway.Security.Ciphers.TripleDes128` field

A **sku** block supports the following:

- **name** - (Required) Specifies the Pricing Tier for the API Management Service. Possible values include: `Developer`, `Basic`, `Standard` and `Premium`.
- **capacity** - (Required) Specifies the Pricing Capacity for the API Management Service.

A **sign_in** block supports the following:

- **enabled** - (Required) Should anonymous users be redirected to the sign in page?

A **sign_up** block supports the following:

- **enabled** - (Required) Can users sign up on the development portal?
- **terms_of_service** - (Optional) A **terms_of_service** block as defined below.

A **terms_of_service** block supports the following:

- **consent_required** - (Required) Should the user be asked for consent during sign up?
- **enabled** - (Required) Should Terms of Service be displayed during sign up?.
- **text** - (Required) The Terms of Service which users are required to agree to in order to sign up.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- **id** - The ID of the API Management Service.

- `additional_location` - One or more `additional_location` blocks as documented below.
- `gateway_url` - The URL of the Gateway for the API Management Service.
- `gateway_regional_url` - The Region URL for the Gateway of the API Management Service.
- `identity` - An `identity` block as defined below.
- `management_api_url` - The URL for the Management API associated with this API Management service.
- `portal_url` - The URL for the Publisher Portal associated with this API Management service.
- `public_ip_addresses` - The Public IP addresses of the API Management Service.
- `scm_url` - The URL for the SCM (Source Code Management) Endpoint associated with this API Management service.

An `additional_location` block exports the following:

- `gateway_regional_url` - The URL of the Regional Gateway for the API Management Service in the specified region.
- `public_ip_addresses` - Public Static Load Balanced IP addresses of the API Management service in the additional location. Available only for Basic, Standard and Premium SKU.

An `identity` block exports the following:

- `principal_id` - The Principal ID associated with this Managed Service Identity.
- `tenant_id` - The Tenant ID associated with this Managed Service Identity.

» Import

API Management Services can be imported using the `resource id`, e.g.

```
terraform import azurerm_api_management.test /subscriptions/00000000-0000-0000-0000-00000000
```

» `azurerm_api_management_api`

Manages an API within an API Management Service.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "example-resources"
  location  = "West Europe"
}

resource "azurerm_api_management" "test" {
  name                = "example-apim"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  publisher_name      = "My Company"
  publisher_email     = "company@terraform.io"

  sku {
    name      = "Developer"
    capacity = 1
  }
}

resource "azurerm_api_management_api" "test" {
  name                = "example-api"
  resource_group_name = "${azurerm_resource_group.test.name}"
  api_management_name = "${azurerm_api_management.test.name}"
  revision            = "1"
  display_name        = "Example API"
  path                = "example"
  protocols            = ["https"]

  import {
    content_format = "swagger-link-json"
    content_value  = "http://conferenceapi.azurewebsites.net/?format=json"
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the API Management API. Changing this forces a new resource to be created.
- **api_management_name** - (Required) The Name of the API Management Service where this API should be created. Changing this forces a new resource to be created.

- **resource_group_name** - (Required) The Name of the Resource Group where the API Management API exists. Changing this forces a new resource to be created.
- **revision** - (Required) The Revision which used for this API.
- **display_name** - (Required) The display name of the API.
- **path** - (Required) The Path for this API Management API, which is a relative URL which uniquely identifies this API and all of its resource paths within the API Management Service.
- **protocols** - (Required) A list of protocols the operations in this API can be invoked. Possible values are **http** and **https**.

-
- **description** - (Optional) A description of the API Management API, which may include HTML formatting tags.
 - **import** - (Optional) A **import** block as documented below.
 - **service_url** - (Optional) Absolute URL of the backend service implementing this API.
 - **soap_pass_through** - (Optional) Should this API expose a SOAP frontend, rather than a HTTP frontend? Defaults to **false**.
 - **subscription_key_parameter_names** - (Optional) A **subscription_key_parameter_names** block as documented below.

A **import** block supports the following:

- **content_format** - (Required) The format of the content from which the API Definition should be imported. Possible values are: **swagger-json**, **swagger-link-json**, **wadl-link-json**, **wadl-xml**, **wsdl** and **wsdl-link**.
- **content_value** - (Required) The Content from which the API Definition should be imported. When a **content_format** of ***-link-*** is specified this must be a URL, otherwise this must be defined inline.
- **wsdl_selector** - (Optional) A **wsdl_selector** block as defined below, which allows you to limit the import of a WSDL to only a subset of the document. This can only be specified when **content_format** is **wsdl** or **wsdl-link**.

A **subscription_key_parameter_names** block supports the following:

- **header** - (Required) The name of the HTTP Header which should be used for the Subscription Key.

- **query** - (Required) The name of the QueryString parameter which should be used for the Subscription Key.

A `wsdl_selector` block supports the following:

- **service_name** - (Required) The name of service to import from WSDL.
 - **endpoint_name** - (Required) The name of endpoint (port) to import from WSDL.
-

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- **id** - The ID of the API Management API.
- **is_current** - Is this the current API Revision?
- **is_online** - Is this API Revision online/accessible via the Gateway?
- **version** - The Version number of this API, if this API is versioned.
- **version_set_id** - The ID of the Version Set which this API is associated with.

» Import

API Management API's can be imported using the `resource_id`, e.g.

```
terraform import azurerm_api_management_api.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm__api__management__api__operation`

Manages an API Operation within an API Management Service.

» Example Usage

```
data "azurerm_api_management_api" "example" {
  name = "search-api"
  api_management_name = "search-api-management"
  resource_group_name = "search-service"
  revision = "2"
}
```

```

resource "azurerm_api_management_api_operation" "example" {
  operation_id      = "user-delete"
  api_name          = "${data.azurerm_api_management_api.example.name}"
  api_management_name = "${data.azurerm_api_management_api.example.api_management_name}"
  resource_group_name = "${data.azurerm_api_management_api.example.resource_group_name}"
  display_name      = "Delete User Operation"
  method            = "DELETE"
  url_template       = "/users/{id}/delete"
  description        = "This can only be done by the logged in user."

  response {
    status_code = 200
  }
}

```

» Argument Reference

The following arguments are supported:

- **operation_id** - (Required) A unique identifier for this API Operation. Changing this forces a new resource to be created.
 - **api_name** - (Required) The name of the API within the API Management Service where this API Operation should be created. Changing this forces a new resource to be created.
 - **api_management_name** - (Required) The Name of the API Management Service where the API exists. Changing this forces a new resource to be created.
 - **resource_group_name** - (Required) The Name of the Resource Group in which the API Management Service exists. Changing this forces a new resource to be created.
 - **display_name** - (Required) The Display Name for this API Management Operation.
 - **method** - (Required) The HTTP Method used for this API Management Operation, like GET, DELETE, PUT or POST - but not limited to these values.
 - **url_template** - (Required) The relative URL Template identifying the target resource for this operation, which may include parameters.
-
- **description** - (Optional) A description for this API Operation, which may include HTML formatting tags.
 - **request** - (Optional) A **request** block as defined below.

- **response** - (Optional) One or more **response** blocks as defined below.
 - **template_parameter** - (Optional) One or more **template_parameter** blocks as defined below.
-

A **form_parameter** block supports the following:

- **name** - (Required) The Name of this Form Parameter.
 - **required** - (Required) Is this Form Parameter Required?
 - **type** - (Required) The Type of this Form Parameter, such as a **string**.
 - **description** - (Optional) A description of this Form Parameter.
 - **default_value** - (Optional) The default value for this Form Parameter.
 - **values** - (Optional) One or more acceptable values for this Form Parameter.
-

A **header** block supports the following:

- **name** - (Required) The Name of this Header.
 - **required** - (Required) Is this Header Required?
 - **type** - (Required) The Type of this Header, such as a **string**.
 - **description** - (Optional) A description of this Header.
 - **default_value** - (Optional) The default value for this Header.
 - **values** - (Optional) One or more acceptable values for this Header.
-

A **query_parameter** block supports the following:

- **name** - (Required) The Name of this Query Parameter.
 - **required** - (Required) Is this Query Parameter Required?
 - **type** - (Required) The Type of this Query Parameter, such as a **string**.
 - **description** - (Optional) A description of this Query Parameter.
 - **default_value** - (Optional) The default value for this Query Parameter.
 - **values** - (Optional) One or more acceptable values for this Query Parameter.
-

A **request** block supports the following:

- **description** - (Required) A description of the HTTP Request, which may include HTML tags.
 - **header** - (Optional) One or more **header** blocks as defined above.
 - **query_parameter** - (Optional) One or more **query_parameter** blocks as defined above.
 - **representation** - (Optional) One or more **representation** blocks as defined below.
-

A **representation** block supports the following:

- **content_type** - (Required) The Content Type of this representation, such as **application/json**.
- **form_parameter** - (Optional) One or more **form_parameter** block as defined above.

NOTE: This is Required when **content_type** is set to **application/x-www-form-urlencoded** or **multipart/form-data**.

- **sample** - (Optional) An example of this representation.
- **schema_id** - (Optional) The ID of an API Management Schema which represents this Response.

NOTE: This can only be specified when **content_type** is not set to **application/x-www-form-urlencoded** or **multipart/form-data**.

- **type_name** - (Optional) The Type Name defined by the Schema.

NOTE: This can only be specified when **content_type** is not set to **application/x-www-form-urlencoded** or **multipart/form-data**.

A **response** block supports the following:

- **status_code** - (Required) The HTTP Status Code.
 - **description** - (Required) A description of the HTTP Response, which may include HTML tags.
 - **header** - (Optional) One or more **header** blocks as defined above.
 - **representation** - (Optional) One or more **representation** blocks as defined below.
-

A **template_parameter** block supports the following:

- **name** - (Required) The Name of this Template Parameter.

- **required** - (Required) Is this Template Parameter Required?
 - **type** - (Required) The Type of this Template Parameter, such as a **string**.
 - **description** - (Optional) A description of this Template Parameter.
 - **default_value** - (Optional) The default value for this Template Parameter.
 - **values** - (Optional) One or more acceptable values for this Template Parameter.
-

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- **id** - The ID of the API Management API Operation.

» Import

API Management API Operation's can be imported using the **resource id**, e.g.

```
terraform import azurerm_api_management_api_operation.test /subscriptions/00000000-0000-0000
```

» azurerm__api__management__version__set

Manages an API Version Set within a API Management Service.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "example-resources"
  location = "West US"
}

resource "azurerm_api_management" "example" {
  name                = "example-apim"
  location             = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
  publisher_name       = "pub1"
  publisher_email      = "pub1@email.com"

  sku {
```

```

        name      = "Developer"
        capacity = 1
    }
}

resource "azurerm_api_management_version_set" "example" {
    name                = "example-apimapi-1.0.0"
    resource_group_name = "${azurerm_resource_group.example.name}"
    api_management_name = "${azurerm_api_management.example.name}"
    display_name        = "ExampleAPIVersionSet"
    versioning_scheme   = "Segment"
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the API Version Set. Changing this forces a new resource to be created.
- **api_management_name** - (Required) The name of the API Management Service in which the API Version Set should exist. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the Resource Group in which the parent API Management Service exists. Changing this forces a new resource to be created.
- **display_name** - (Required) The display name of this API Version Set.
- **versioning_scheme** - (Required) Specifies where in an Inbound HTTP Request that the API Version should be read from. Possible values are **Header**, **Query** and **Segment**.

-
- **description** - (Optional) The description of API Version Set.
 - **version_header_name** - (Optional) The name of the Header which should be read from Inbound Requests which defines the API Version.

NOTE: This must be specified when **versioning_scheme** is set to **Header**.

- **version_query_name** - (Optional) The name of the Query String which should be read from Inbound Requests which defines the API Version.

NOTE: This must be specified when **versioning_scheme** is set to **Query**.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- `id` - The ID of the API Version Set.

» Import

API Version Set can be imported using the `resource id`, e.g.

```
terraform import azurerm_api_management_version_set.example /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_api_management_authorization_server`

Manages an Authorization Server within an API Management Service.

» Example Usage

```
data "azurerm_api_management_api" "example" {
  name                       = "search-api"
  api_management_name       = "search-api-management"
  resource_group_name       = "search-service"
  revision                  = "2"
}

resource "azurerm_api_management_authorization_server" "example" {
  name                       = "test-server"
  api_management_name       = "${data.azurerm_api_management.example.name}"
  resource_group_name       = "${data.azurerm_api_management.example.resource_group_name}"
  display_name              = "Test Server"
  authorization_endpoint     = "https://example.mydomain.com/client/authorize"
  client_id                 = "42424242-4242-4242-4242-424242424242"
  client_registration_endpoint = "https://example.mydomain.com/client/register"
  grant_types = [
    "authorizationCode",
  ]
}
```

» Argument Reference

The following arguments are supported:

- **api_management_name** - (Required) The name of the API Management Service in which this Authorization Server should be created. Changing this forces a new resource to be created.
- **authorization_methods** - (Required) The HTTP Verbs supported by the Authorization Endpoint. Possible values are **DELETE**, **GET**, **HEAD**, **OPTIONS**, **PATCH**, **POST**, **PUT** and **TRACE**.

NOTE: **GET** must always be present.

- **authorization_endpoint** - (Required) The OAUTH Authorization Endpoint.
- **client_id** - (Required) The Client/App ID registered with this Authorization Server.
- **client_registration_endpoint** - (Required) The URI of page where Client/App Registration is performed for this Authorization Server.
- **display_name** - (Required) The user-friendly name of this Authorization Server.
- **grant_types** - (Required) Form of Authorization Grants required when requesting an Access Token. Possible values are **authorizationCode**, **clientCredentials**, **implicit** and **resourceOwnerPassword**.
- **name** - (Required) The name of this Authorization Server. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the Resource Group in which the API Management Service exists. Changing this forces a new resource to be created.

-
- **bearer_token_sending_methods** - (Optional) The mechanism by which Access Tokens are passed to the API. Possible values are **authorizationHeader** and **query**.
 - **client_authentication_method** - (Optional) The Authentication Methods supported by the Token endpoint of this Authorization Server.. Possible values are **Basic** and **Body**.
 - **client_secret** - (Optional) The Client/App Secret registered with this Authorization Server.
 - **default_scope** - (Optional) The Default Scope used when requesting an Access Token, specified as a string containing space-delimited values.
 - **description** - (Optional) A description of the Authorization Server, which may contain HTML formatting tags.
 - **resource_owner_password** - (Optional) The password associated with the Resource Owner.

NOTE: This can only be specified when `grant_type` includes `resourceOwnerPassword`.

- `resource_owner_username` - (Optional) The username associated with the Resource Owner.

NOTE: This can only be specified when `grant_type` includes `resourceOwnerPassword`.

- `support_state` - (Optional) Does this Authorization Server support State? If this is set to `true` the client may use the state parameter to raise protocol security.
- `token_body_parameters` - (Optional) A `token_body_parameters` block as defined below.
- `token_endpoint` - (Optional) The OAUTH Token Endpoint.

A `token_body_parameter` block supports the following:

- `name` - (Required) The Name of the Parameter.
- `value` - (Required) The Value of the Parameter.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- `id` - The ID of the API Management Authorization Server.

» Import

API Management Authorization Servers can be imported using the `resource id`, e.g.

```
terraform import azurerm_api_management_authorization_server.example /subscriptions/00000000
```

» `azurerm__api__management__certificate`

Manages an Certificate within an API Management Service.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "example-resources"
  location = "West Europe"
}
```

```

resource "azurerm_api_management" "test" {
  name                = "example-apim"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  publisher_name       = "My Company"
  publisher_email      = "company@terraform.io"

  sku {
    name       = "Developer"
    capacity = 1
  }
}

resource "azurerm_api_management_certificate" "test" {
  name                = "example-cert"
  api_management_name = "${azurerm_api_management.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  data                = "${base64encode(file("example.pfx"))}"
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the API Management Certificate. Changing this forces a new resource to be created.
- **api_management_name** - (Required) The Name of the API Management Service where this Service should be created. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The Name of the Resource Group where the API Management Service exists. Changing this forces a new resource to be created.
- **data** - (Required) The base-64 encoded certificate data, which must be a PFX file. Changing this forces a new resource to be created.
- **password** - (Optional) The password used for this certificate. Changing this forces a new resource to be created.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- `id` - The ID of the API Management Certificate.
- `expiration` - The Expiration Date of this Certificate, formatted as an RFC3339 string.
- `subject` - The Subject of this Certificate.
- `thumbprint` - The Thumbprint of this Certificate.

» Import

API Management Certificates can be imported using the `resource id`, e.g.

```
terraform import azurerm_api_management_certificate.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_api_management_group`

Manages an API Management Group.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "example-resources"
  location = "West US"
}

resource "azurerm_api_management" "example" {
  name                = "example-apim"
  location             = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
  publisher_name      = "pub1"
  publisher_email     = "pub1@email.com"

  sku {
    name       = "Developer"
    capacity = 1
  }
}

resource "azurerm_api_management_group" "example" {
  name                = "example-aping"
  resource_group_name = "${azurerm_resource_group.example.name}"
  api_management_name = "${azurerm_api_management.example.name}"
  display_name       = "Example Group"
}
```

```

    description      = "This is an example API management group."
  }

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the API Management Group. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the Resource Group in which the API Management Group should exist. Changing this forces a new resource to be created.
- **api_management_name** - (Required) The name of the API Management Service in which the API Management Group should exist. Changing this forces a new resource to be created.
- **display_name** - (Required) The display name of this API Management Group.
- **description** - (Optional) The description of this API Management Group.
- **external_id** - (Optional) The identifier of the external Group. For example, an Azure Active Directory group `aad://<tenant>.onmicrosoft.com/groups/<group object id>`.
- **type** - (Optional) The type of this API Management Group. Possible values are `custom` and `external`. Default is `custom`.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- **id** - The ID of the API Management Group.

» Import

API Management Groups can be imported using the **resource id**, e.g.

```
terraform import azurerm_api_management_group.example /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_api_management_group_user

Manages an API Management User Assignment to a Group.

» Example Usage

```
data "azurerm_api_management_user" "example" {
  user_id            = "my-user"
  api_management_name = "example-apim"
  resource_group_name = "search-service"
}

resource "azurerm_api_management_group_user" "example" {
  user_id            = "${data.azurerm_api_management_user.example.id}"
  group_name         = "example-group"
  resource_group_name = "${data.azurerm_api_management_user.example.resource_group_name}"
  api_management_name = "${data.azurerm_api_management_user.example.api_management_name}"
}
```

» Argument Reference

The following arguments are supported:

- **user_id** - (Required) The ID of the API Management User which should be assigned to this API Management Group. Changing this forces a new resource to be created.
- **group_name** - (Required) The Name of the API Management Group within the API Management Service. Changing this forces a new resource to be created.
- **api_management_name** - (Required) The name of the API Management Service. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the Resource Group in which the API Management Service exists. Changing this forces a new resource to be created.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- **id** - The ID of the API Management Group User.

» Import

API Management Group Users can be imported using the **resource id**, e.g.

```
terraform import azurerm_api_management_group_user.example /subscriptions/00000000-0000-0000
```

» `azurerm_api_management_logger`

Manages a Logger within an API Management Service.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name      = "example-resources"
  location  = "West US"
}

resource "azurerm_application_insights" "example" {
  name                = "example-appinsights"
  location             = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
  application_type     = "Other"
}

resource "azurerm_api_management" "example" {
  name                = "example-apim"
  location             = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
  publisher_name      = "My Company"
  publisher_email     = "company@terraform.io"

  sku {
    name       = "Developer"
    capacity   = 1
  }
}

resource "azurerm_api_management_logger" "example" {
  name                = "example-logger"
  api_management_name = "${azurerm_api_management.example.name}"
  resource_group_name = "${azurerm_resource_group.example.name}"

  application_insights {
    instrumentation_key = "${azurerm_application_insights.example.instrumentation_key}"
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of this Logger, which must be unique within the API Management Service. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the Resource Group in which the API Management Service exists. Changing this forces a new resource to be created.
- **api_management_name** - (Required) The name of the API Management Service. Changing this forces a new resource to be created.
- **application_insights** - (Optional) An **application_insights** block as documented below.
- **buffered** - (Optional) Specifies whether records should be buffered in the Logger prior to publishing. Defaults to **true**.
- **description** - (Optional) A description of this Logger.
- **eventhub** - (Optional) An **eventhub** block as documented below.

An **application_insights** block supports the following:

- **instrumentation_key** - (Required) The instrumentation key used to push data to Application Insights.

An **eventhub** block supports the following:

- **name** - (Required) The name of an EventHub.
- **connection_string** - (Required) The connection string of an EventHub Namespace.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- **id** - The ID of the API Management Logger.

» Import

API Management Loggers can be imported using the **resource id**, e.g.

```
$ terraform import azurerm_api_management_logger.example /subscriptions/00000000-0000-0000-0000-000000000000/resourceGroups/default-rg/providers/Microsoft.ApiManagement/service/example/loggers/default-logger
```


» `azurerm_api_management_openid_connect_provider`

Manages an OpenID Connect Provider within a API Management Service.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "example-resources"
  location = "West Europe"
}

resource "azurerm_api_management" "test" {
  name                = "example-apim"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  publisher_name      = "My Company"
  publisher_email     = "company@terraform.io"

  sku {
    name       = "Developer"
    capacity = 1
  }
}

resource "azurerm_api_management_openid_connect_provider" "test" {
  name                = "example-provider"
  api_management_name = "${azurerm_api_management.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  client_id           = "00001111-2222-3333-4444-555566667777"
  display_name        = "Example Provider"
  metadata_endpoint   = "https://example.com/example"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) the Name of the OpenID Connect Provider which should be created within the API Management Service. Changing this forces a new resource to be created.
- **api_management_name** - (Required) The name of the API Management Service in which this OpenID Connect Provider should be created. Changing this forces a new resource to be created.

- **resource_group_name** - (Required) The name of the Resource Group where the API Management Service exists. Changing this forces a new resource to be created.
 - **client_id** - (Required) The Client ID used for the Client Application.
 - **client_secret** - (Required) The Client Secret used for the Client Application.
 - **display_name** - (Required) A user-friendly name for this OpenID Connect Provider.
 - **metadata_endpoint** - (Required) The URI of the Metadata endpoint.
-
- **description** - (Optional) A description of this OpenID Connect Provider.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- **id** - The ID of the API Management OpenID Connect Provider.

» Import

API Management OpenID Connect Providers can be imported using the resource id, e.g.

```
terraform import azurerm_api_management_openid_connect_provider.test /subscriptions/00000000
```

» azurerm_api_management_product

Manages an API Management Product.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "example-resources"
  location = "West Europe"
}

resource "azurerm_api_management" "test" {
  name          = "example-apim"
  location      = "${azurerm_resource_group.test.location}"
}
```

```

resource_group_name = "${azurerm_resource_group.test.name}"
publisher_name      = "My Company"
publisher_email     = "company@terraform.io"

sku {
  name      = "Developer"
  capacity = 1
}

resource "azurerm_api_management_product" "test" {
  product_id          = "test-product"
  api_management_name = "${azurerm_api_management.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  display_name        = "Test Product"
  subscription_required = true
  approval_required   = true
  published            = true
}

```

» Argument Reference

The following arguments are supported:

- **api_management_name** - (Required) The name of the API Management Service. Changing this forces a new resource to be created.
- **approval_required** - (Optional) Do subscribers need to be approved prior to being able to use the Product?

NOTE: **approval_required** can only be set when **subscription_required** is set to **true**.

- **display_name** - (Required) The Display Name for this API Management Product.
- **product_id** - (Required) The Identifier for this Product, which must be unique within the API Management Service. Changing this forces a new resource to be created.
- **published** - (Required) Is this Product Published?
- **resource_group_name** - (Required) The name of the Resource Group in which the API Management Service should be exist. Changing this forces a new resource to be created.
- **subscription_required** - (Required) Is a Subscription required to access API's included in this Product?

-
- **description** - (Optional) A description of this Product, which may include HTML formatting tags.
 - **subscriptions_limit** - (Optional) The number of subscriptions a user can have to this Product at the same time.

NOTE: **subscriptions_limit** can only be set when **subscription_required** is set to **true**.

- **terms** - (Optional) The Terms and Conditions for this Product, which must be accepted by Developers before they can begin the Subscription process.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- **id** - The ID of the API Management Product.

» Import

API Management Products can be imported using the **resource id**, e.g.

```
terraform import azurerm_api_management_product.test /subscriptions/00000000-0000-0000-0000-
```

» `azurerm_api_management_product_api`

Manages an API Management API Assignment to a Product.

» Example Usage

```
data "azurerm_api_management" "example" {
  name                = "example-api"
  resource_group_name = "example-resources"
}
```

```
data "azurerm_api_management_api" "example" {
  name                = "search-api"
  api_management_name = "${data.azurerm_api_management.example.name}"
  resource_group_name = "${data.azurerm_api_management.example.resource_group_name}"
  revision            = "2"
}
```

```

}
data "azurerm_api_management_product" "test" {
  product_id          = "my-product"
  api_management_name = "${data.azurerm_api_management.example.name}"
  resource_group_name = "${data.azurerm_api_management.example.resource_group_name}"
}

resource "azurerm_api_management_product_api" "example" {
  api_name          = "${data.azurerm_api_management_api.example.name}"
  product_id        = "${data.azurerm_api_management_product.example.product_id}"
  api_management_name = "${data.azurerm_api_management.example.name}"
  resource_group_name = "${data.azurerm_api_management.example.resource_group_name}"
}

```

» Argument Reference

The following arguments are supported:

- **api_name** - (Required) The Name of the API Management API within the API Management Service. Changing this forces a new resource to be created.
- **api_management_name** - (Required) The name of the API Management Service. Changing this forces a new resource to be created.
- **product_id** - (Required) The ID of the API Management Product within the API Management Service. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the Resource Group in which the API Management Service exists. Changing this forces a new resource to be created.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- **id** - The ID of the API Management Product API.

» Import

API Management Product API's can be imported using the **resource id**, e.g.

```
terraform import azurerm_api_management_product_api.example /subscriptions/00000000-0000-0000-0000-000000000000
```

» **azurerm_api_management_product_group**

Manages an API Management Product Assignment to a Group.

» **Example Usage**

```
data "azurerm_api_management" "example" {
  name                = "example-api"
  resource_group_name = "example-resources"
}

data "azurerm_api_management_product" "example" {
  product_id          = "my-product"
  api_management_name = "${data.azurerm_api_management.example.name}"
  resource_group_name = "${data.azurerm_api_management.example.resource_group_name}"
}

data "azurerm_api_management_group" "example" {
  name                = "my-group"
  api_management_name = "${data.azurerm_api_management.example.name}"
  resource_group_name = "${data.azurerm_api_management.example.resource_group_name}"
}

resource "azurerm_api_management_product_group" "example" {
  product_id          = "${data.azurerm_api_management_user.example.id}"
  group_name          = "${data.azurerm_api_management_group.example.name}"
  api_management_name = "${data.azurerm_api_management.example.name}"
  resource_group_name = "${data.azurerm_api_management.example.resource_group_name}"
}
```

» **Argument Reference**

The following arguments are supported:

- **product_id** - (Required) The ID of the API Management Product within the API Management Service. Changing this forces a new resource to be created.
- **group_name** - (Required) The Name of the API Management Group within the API Management Service. Changing this forces a new resource to be created.
- **api_management_name** - (Required) The name of the API Management Service. Changing this forces a new resource to be created.

- `resource_group_name` - (Required) The name of the Resource Group in which the API Management Service exists. Changing this forces a new resource to be created.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- `id` - The ID of the API Management Product Group.

» Import

API Management Product Groups can be imported using the `resource id`, e.g.

```
terraform import azurerm_api_management_product_group.example /subscriptions/000000000-0000-0000-0000-000000000000
```

» `azurerm_api_management_property`

Manages an API Management Property.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "example-resources"
  location = "West US"
}

resource "azurerm_api_management" "example" {
  name                = "example-apim"
  location            = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
  publisher_name      = "pub1"
  publisher_email     = "pub1@email.com"

  sku {
    name       = "Developer"
    capacity = 1
  }
}

resource "azurerm_api_management_property" "example" {
  name = "example-aping"
```

```

resource_group_name = "${azurerm_resource_group.example.name}"
api_management_name = "${azurerm_api_management.example.name}"
display_name        = "ExampleProperty"
value               = "Example Value"
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the API Management Property. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the Resource Group in which the API Management Property should exist. Changing this forces a new resource to be created.
- **api_management_name** - (Required) The name of the API Management Service in which the API Management Property should exist. Changing this forces a new resource to be created.
- **display_name** - (Required) The display name of this API Management Property.
- **value** - (Required) The value of this API Management Property.
- **secret** - (Optional) Specifies whether the API Management Property is secret. Valid values are **true** or **false**. The default value is **false**.

NOTE: setting the field **secret** to **true** doesn't make this field sensitive in Terraform, instead it marks the value as secret and encrypts the value in Azure.

- **tags** - (Optional) A list of tags to be applied to the API Management Property.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- **id** - The ID of the API Management Property.

» Import

API Management Properties can be imported using the **resource id**, e.g.

```
terraform import azurerm_api_management_property.example /subscriptions/00000000-0000-0000-0000-000000000000/resourceGroups/example/providers/Microsoft.ApiManagement/service/example/properties/example
```


» azurerm_app_service

Manages an App Service (within an App Service Plan).

Note: When using Slots - the `app_settings`, `connection_string` and `site_config` blocks on the `azurerm_app_service` resource will be overwritten when promoting a Slot using the `azurerm_app_service_active_slot` resource.

» Example Usage

This example provisions a Windows App Service. Other examples of the `azurerm_app_service` resource can be found in the `./examples/app-service` directory within the Github Repository

```
resource "azurerm_resource_group" "test" {
  name     = "example-resources"
  location = "West Europe"
}

resource "azurerm_app_service_plan" "test" {
  name                = "example-appserviceplan"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  sku {
    tier = "Standard"
    size = "S1"
  }
}

resource "azurerm_app_service" "test" {
  name                = "example-app-service"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  app_service_plan_id = "${azurerm_app_service_plan.test.id}"

  site_config {
    dotnet_framework_version = "v4.0"
    scm_type                  = "LocalGit"
  }

  app_settings = {
    "SOME_KEY" = "some-value"
  }
}
```

```

connection_string {
  name = "Database"
  type = "SQLServer"
  value = "Server=some-server.mydomain.com;Integrated Security=SSPI"
}
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the App Service. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the App Service.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **app_service_plan_id** - (Required) The ID of the App Service Plan within which to create this App Service.
- **app_settings** - (Optional) A key-value pair of App Settings.
- **connection_string** - (Optional) One or more **connection_string** blocks as defined below.
- **client_affinity_enabled** - (Optional) Should the App Service send session affinity cookies, which route client requests in the same session to the same instance?
- **client_cert_enabled** - (Optional) Does the App Service require client certificates for incoming requests? Defaults to **false**.
- **enabled** - (Optional) Is the App Service Enabled?
- **https_only** - (Optional) Can the App Service only be accessed via HTTPS? Defaults to **false**.
- **site_config** - (Optional) A **site_config** block as defined below.
- **tags** - (Optional) A mapping of tags to assign to the resource.
- **identity** - (Optional) A Managed Service Identity block as defined below.

A **connection_string** block supports the following:

- **name** - (Required) The name of the Connection String.

- **type** - (Required) The type of the Connection String. Possible values are **APIHub**, **Custom**, **DocDb**, **EventHub**, **MySQL**, **NotificationHub**, **PostgreSQL**, **RedisCache**, **ServiceBus**, **SQLAzure** and **SQLServer**.
 - **value** - (Required) The value for the Connection String.
-

A **identity** block supports the following:

- **type** - (Required) Specifies the identity type of the App Service. At this time the only allowed value is **SystemAssigned**.

The assigned **principal_id** and **tenant_id** can be retrieved after the App Service has been created. More details are available below.

A **site_config** block supports the following:

- **always_on** - (Optional) Should the app be loaded at all times? Defaults to **false**.
- **app_command_line** - (Optional) App command line to launch, e.g. **/sbin/myserver -b 0.0.0.0**.
- **cors** - (Optional) A **cors** block as defined below.
- **default_documents** - (Optional) The ordering of default documents to load, if an address isn't specified.
- **dotnet_framework_version** - (Optional) The version of the .net framework's CLR used in this App Service. Possible values are **v2.0** (which will use the latest version of the .net framework for the .net CLR v2 - currently **.net 3.5**) and **v4.0** (which corresponds to the latest version of the .net CLR v4 - which at the time of writing is **.net 4.7.1**). For more information on which .net CLR version to use based on the .net framework you're targeting - please see this table. Defaults to **v4.0**.
- **ftps_state** - (Optional) State of FTP / FTPS service for this App Service. Possible values include: **AllAllowed**, **FtpsOnly** and **Disabled**.
- **http2_enabled** - (Optional) Is HTTP2 Enabled on this App Service? Defaults to **false**.
- **ip_restriction** - (Optional) One or more **ip_restriction** blocks as defined below.
- **java_version** - (Optional) The version of Java to use. If specified **java_container** and **java_container_version** must also be specified. Possible values are **1.7** and **1.8**.

- `java_container` - (Optional) The Java Container to use. If specified `java_version` and `java_container_version` must also be specified. Possible values are JETTY and TOMCAT.
- `java_container_version` - (Optional) The version of the Java Container to use. If specified `java_version` and `java_container` must also be specified.
- `local_mysql_enabled` - (Optional) Is "MySQL In App" Enabled? This runs a local MySQL instance with your app and shares resources from the App Service plan.

NOTE: MySQL In App is not intended for production environments and will not scale beyond a single instance. Instead you may wish to use Azure Database for MySQL.

- `linux_fx_version` - (Optional) Linux App Framework and version for the App Service. Possible options are a Docker container (`DOCKER|<user/image:tag>`), a base-64 encoded Docker Compose file (`COMPOSE|${base64encode(file("compose.yml"))}`) or a base-64 encoded Kubernetes Manifest (`KUBE|${base64encode(file("kubernetes.yml"))}`).

Additional examples of how to run Containers via the `azurerm_app_service` resource can be found in the `./examples/app-service` directory within the Github Repository.

- `managed_pipeline_mode` - (Optional) The Managed Pipeline Mode. Possible values are `Integrated` and `Classic`. Defaults to `Integrated`.
- `min_tls_version` - (Optional) The minimum supported TLS version for the app service. Possible values are 1.0, 1.1, and 1.2. Defaults to 1.2 for new app services.
- `php_version` - (Optional) The version of PHP to use in this App Service. Possible values are 5.5, 5.6, 7.0, 7.1 and 7.2.
- `python_version` - (Optional) The version of Python to use in this App Service. Possible values are 2.7 and 3.4.
- `remote_debugging_enabled` - (Optional) Is Remote Debugging Enabled? Defaults to `false`.
- `remote_debugging_version` - (Optional) Which version of Visual Studio should the Remote Debugger be compatible with? Possible values are VS2012, VS2013, VS2015 and VS2017.
- `scm_type` - (Optional) The type of Source Control enabled for this App Service. Defaults to `None`. Possible values are: `BitbucketGit`, `BitbucketHg`, `CodePlexGit`, `CodePlexHg`, `Dropbox`, `ExternalGit`, `ExternalHg`, `GitHub`, `LocalGit`, `None`, `OneDrive`, `Tfs`, `VSO` and `VSTSRM`

- `use_32_bit_worker_process` - (Optional) Should the App Service run in 32 bit mode, rather than 64 bit mode?

NOTE: when using an App Service Plan in the **Free** or **Shared** Tiers `use_32_bit_worker_process` must be set to `true`.

- `virtual_network_name` - (Optional) The name of the Virtual Network which this App Service should be attached to.
- `websockets_enabled` - (Optional) Should WebSockets be enabled?

A `cors` block supports the following:

- `allowed_origins` - (Optional) A list of origins which should be able to make cross-origin calls. `*` can be used to allow all calls.
- `support_credentials` - (Optional) Are credentials supported?

A `ip_restriction` block supports the following:

- `ip_address` - (Required) The IP Address used for this IP Restriction.
- `subnet_mask` - (Optional) The Subnet mask used for this IP Restriction. Defaults to `255.255.255.255`.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the App Service.
- `default_site_hostname` - The Default Hostname associated with the App Service - such as `mysite.azurewebsites.net`
- `outbound_ip_addresses` - A comma separated list of outbound IP addresses - such as `52.23.25.3,52.143.43.12`
- `possible_outbound_ip_addresses` - A comma separated list of outbound IP addresses - such as `52.23.25.3,52.143.43.12,52.143.43.17` - not all of which are necessarily in use. Superset of `outbound_ip_addresses`.
- `source_control` - A `source_control` block as defined below, which contains the Source Control information when `scm_type` is set to `LocalGit`.
- `site_credential` - A `site_credential` block as defined below, which contains the site-level credentials used to publish to this App Service.
- `identity` - An `identity` block as defined below, which contains the Managed Service Identity information for this App Service.

`identity` exports the following:

- `principal_id` - The Principal ID for the Service Principal associated with the Managed Service Identity of this App Service.
- `tenant_id` - The Tenant ID for the Service Principal associated with the Managed Service Identity of this App Service.

You can access the Principal ID via `${azurerm_app_service.test.identity.0.principal_id}` and the Tenant ID via `${azurerm_app_service.test.identity.0.tenant_id}`

`site_credential` exports the following:

- `username` - The username which can be used to publish to this App Service
- `password` - The password associated with the username, which can be used to publish to this App Service.

NOTE: both `username` and `password` for the `site_credential` block are only exported when `scm_type` is set to `LocalGit`

`source_control` exports the following:

- `repo_url` - URL of the Git repository for this App Service.
- `branch` - Branch name of the Git repository for this App Service.

» Import

App Services can be imported using the `resource_id`, e.g.

```
terraform import azurerm_app_service.instance1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_app_service_active_slot`

Promotes an App Service Slot to Production within an App Service.

Note: When using Slots - the `app_settings`, `connection_string` and `site_config` blocks on the `azurerm_app_service` resource will be overwritten when promoting a Slot using the `azurerm_app_service_active_slot` resource.

» Example Usage

```
resource "random_id" "server" {  
  # ...  
}
```

```

}

resource "azurerm_resource_group" "test" {
  # ...
}

resource "azurerm_app_service_plan" "test" {
  # ...
}

resource "azurerm_app_service" "test" {
  # ...
}

resource "azurerm_app_service_slot" "test" {
  # ...
}

resource "azurerm_app_service_active_slot" "test" {
  resource_group_name = "${azurerm_resource_group.test.name}"
  app_service_name     = "${azurerm_app_service.test.name}"
  app_service_slot_name = "${azurerm_app_service_slot.test.name}"
}

```

» Argument Reference

The following arguments are supported:

- **resource_group_name** - (Required) The name of the resource group in which the App Service exists. Changing this forces a new resource to be created.
- **app_service_name** - (Required) The name of the App Service within which the Slot exists. Changing this forces a new resource to be created.
- **app_service_slot_name** - (Required) The name of the App Service Slot which should be promoted to the Production Slot within the App Service.

» azurerm__app__service__custom__hostname__binding

Manages a Hostname Binding within an App Service.

» Example Usage

```
resource "random_id" "server" {
  keepers = {
    azi_id = 1
  }

  byte_length = 8
}

resource "azurerm_resource_group" "test" {
  name      = "some-resource-group"
  location = "West Europe"
}

resource "azurerm_app_service_plan" "test" {
  name                = "some-app-service-plan"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  sku {
    tier = "Standard"
    size = "S1"
  }
}

resource "azurerm_app_service" "test" {
  name                = "${random_id.server.hex}"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  app_service_plan_id = "${azurerm_app_service_plan.test.id}"
}

resource "azurerm_app_service_custom_hostname_binding" "test" {
  hostname          = "www.mywebsite.com"
  app_service_name  = "${azurerm_app_service.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}
```

» Argument Reference

The following arguments are supported:

- **hostname** - (Required) Specifies the Custom Hostname to use for the App Service, example `www.example.com`. Changing this forces a new resource

to be created.

NOTE: A CNAME needs to be configured from this Hostname to the Azure Website - otherwise Azure will reject the Hostname Binding.

- **app_service_name** - (Required) The name of the App Service in which to add the Custom Hostname Binding. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the App Service exists. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the App Service Custom Hostname Binding

» Import

App Service Custom Hostname Bindings can be imported using the **resource id**, e.g.

```
terraform import azurerm_app_service_custom_hostname_binding.mywebsite /subscriptions/000000
```

» azurerm__app__service__plan

Manage an App Service Plan component.

» Example Usage (Dedicated)

```
resource "azurerm_resource_group" "test" {
  name      = "api-rg-pro"
  location  = "West Europe"
}

resource "azurerm_app_service_plan" "test" {
  name                = "api-appserviceplan-pro"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  sku {
    tier = "Standard"
  }
}
```

```

        size = "S1"
    }
}

```

» Example Usage (Shared / Consumption Plan)

```

resource "azurerm_resource_group" "test" {
  name      = "api-rg-pro"
  location  = "West Europe"
}

resource "azurerm_app_service_plan" "test" {
  name                  = "api-appserviceplan-pro"
  location              = "${azurerm_resource_group.test.location}"
  resource_group_name  = "${azurerm_resource_group.test.name}"
  kind                 = "FunctionApp"

  sku {
    tier = "Dynamic"
    size = "Y1"
  }
}

```

» Example Usage (Linux)

```

resource "azurerm_resource_group" "test" {
  name      = "api-rg-pro"
  location  = "West Europe"
}

resource "azurerm_app_service_plan" "test" {
  name                  = "api-appserviceplan-pro"
  location              = "${azurerm_resource_group.test.location}"
  resource_group_name  = "${azurerm_resource_group.test.name}"
  kind                 = "Linux"

  sku {
    tier = "Standard"
    size = "S1"
  }

  properties {
    reserved = true
  }
}

```

}

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the App Service Plan component. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the App Service Plan component.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **kind** - (Optional) The kind of the App Service Plan to create. Possible values are **Windows** (also available as **App**), **Linux** and **FunctionApp** (for a Consumption Plan). Defaults to **Windows**. Changing this forces a new resource to be created.

NOTE: When creating a **Linux** App Service Plan, the **reserved** field must be set to **true**.

- **sku** - (Required) A **sku** block as documented below.
- **app_service_environment_id** - (Optional) The ID of the App Service Environment where the App Service Plan should be located. Changing forces a new resource to be created.

NOTE: Attaching to an App Service Environment requires the App Service Plan use a **Premium** SKU (when using an ASEv1) and the **Isolated** SKU (for an ASEv2).

- **reserved** - (Optional) Is this App Service Plan **Reserved**. Defaults to **false**.
- **per_site_scaling** - (Optional) Can Apps assigned to this App Service Plan be scaled independently? If set to **false** apps assigned to this plan will scale to all instances of the plan. Defaults to **false**.
- **tags** - (Optional) A mapping of tags to assign to the resource.

sku supports the following:

- **tier** - (Required) Specifies the plan's pricing tier.
- **size** - (Required) Specifies the plan's instance size.
- **capacity** - (Optional) Specifies the number of workers associated with this App Service Plan.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the App Service Plan component.
- `maximum_number_of_workers` - The maximum number of workers supported with the App Service Plan's sku.

» Import

App Service Plan instances can be imported using the `resource id`, e.g.

```
terraform import azurerm_app_service_plan.instance1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_app_service_slot`

Manages an App Service Slot (within an App Service).

Note: When using Slots - the `app_settings`, `connection_string` and `site_config` blocks on the `azurerm_app_service` resource will be overwritten when promoting a Slot using the `azurerm_app_service_active_slot` resource.

» Example Usage (.net 4.x)

```
resource "random_id" "server" {
  keepers = {
    azi_id = 1
  }

  byte_length = 8
}

resource "azurerm_resource_group" "test" {
  name     = "some-resource-group"
  location = "West Europe"
}

resource "azurerm_app_service_plan" "test" {
  name                = "some-app-service-plan"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  sku {
    tier = "Standard"
  }
}
```

```

        size = "S1"
    }
}

resource "azurerm_app_service" "test" {
    name                = "${random_id.server.hex}"
    location            = "${azurerm_resource_group.test.location}"
    resource_group_name = "${azurerm_resource_group.test.name}"
    app_service_plan_id = "${azurerm_app_service_plan.test.id}"

    site_config {
        dotnet_framework_version = "v4.0"
    }

    app_settings = {
        "SOME_KEY" = "some-value"
    }

    connection_string {
        name = "Database"
        type = "SQLServer"
        value = "Server=some-server.mydomain.com;Integrated Security=SSPI"
    }
}

resource "azurerm_app_service_slot" "test" {
    name                = "${random_id.server.hex}"
    app_service_name    = "${azurerm_app_service.test.name}"
    location            = "${azurerm_resource_group.test.location}"
    resource_group_name = "${azurerm_resource_group.test.name}"
    app_service_plan_id = "${azurerm_app_service_plan.test.id}"

    site_config {
        dotnet_framework_version = "v4.0"
    }

    app_settings = {
        "SOME_KEY" = "some-value"
    }

    connection_string {
        name = "Database"
        type = "SQLServer"
        value = "Server=some-server.mydomain.com;Integrated Security=SSPI"
    }
}

```

» Example Usage (Java 1.8)

```
resource "random_id" "server" {
  keepers = {
    azi_id = 1
  }

  byte_length = 8
}

resource "azurerm_resource_group" "test" {
  name     = "some-resource-group"
  location = "West Europe"
}

resource "azurerm_app_service_plan" "test" {
  name                = "some-app-service-plan"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  sku {
    tier = "Standard"
    size = "S1"
  }
}

resource "azurerm_app_service" "test" {
  name                = "${random_id.server.hex}"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  app_service_plan_id = "${azurerm_app_service_plan.test.id}"

  site_config {
    java_version      = "1.8"
    java_container    = "JETTY"
    java_container_version = "9.3"
  }
}

resource "azurerm_app_service_slot" "test" {
  name                = "${random_id.server.hex}"
  app_service_name    = "${azurerm_app_service.test.name}"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  app_service_plan_id = "${azurerm_app_service_plan.test.id}"
}
```

```

site_config {
  java_version      = "1.8"
  java_container    = "JETTY"
  java_container_version = "9.3"
}
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the App Service Slot component. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the App Service Slot component.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **app_service_plan_id** - (Required) The ID of the App Service Plan within which to create this App Service Slot. Changing this forces a new resource to be created.
- **app_service_name** - (Required) The name of the App Service within which to create the App Service Slot. Changing this forces a new resource to be created.
- **app_settings** - (Optional) A key-value pair of App Settings.
- **connection_string** - (Optional) An `connection_string` block as defined below.
- **client_affinity_enabled** - (Optional) Should the App Service Slot send session affinity cookies, which route client requests in the same session to the same instance?
- **enabled** - (Optional) Is the App Service Slot Enabled?
- **https_only** - (Optional) Can the App Service Slot only be accessed via HTTPS? Defaults to `false`.
- **site_config** - (Optional) A `site_config` object as defined below.
- **identity** - (Optional) A Managed Service Identity block as defined below.
- **tags** - (Optional) A mapping of tags to assign to the resource.

`connection_string` supports the following:

- **name** - (Required) The name of the Connection String.
- **type** - (Required) The type of the Connection String. Possible values are `APIHub`, `Custom`, `DocDb`, `EventHub`, `MySQL`, `NotificationHub`, `PostgreSQL`, `RedisCache`, `ServiceBus`, `SQLAzure` and `SQLServer`.
- **value** - (Required) The value for the Connection String.

`site_config` supports the following:

- **app_command_line** - (Optional) App command line to launch, e.g. `/sbin/myserv -b 0.0.0.0`.
- **always_on** - (Optional) Should the app be loaded at all times? Defaults to `false`.
- **cors** - (Optional) A `cors` block as defined below.
- **default_documents** - (Optional) The ordering of default documents to load, if an address isn't specified.
- **dotnet_framework_version** - (Optional) The version of the .net framework's CLR used in this App Service Slot. Possible values are `v2.0` (which will use the latest version of the .net framework for the .net CLR v2 - currently `.net 3.5`) and `v4.0` (which corresponds to the latest version of the .net CLR v4 - which at the time of writing is `.net 4.7.1`). For more information on which .net CLR version to use based on the .net framework you're targeting - please see this table. Defaults to `v4.0`.
- **http2_enabled** - (Optional) Is HTTP2 Enabled on this App Service? Defaults to `false`.
- **ip_restriction** - (Optional) One or more `ip_restriction` blocks as defined below.
- **java_container** - (Optional) The Java Container to use. If specified **java_version** and **java_container_version** must also be specified. Possible values are `JETTY` and `TOMCAT`.
- **java_container_version** - (Optional) The version of the Java Container to use. If specified **java_version** and **java_container** must also be specified.
- **java_version** - (Optional) The version of Java to use. If specified **java_container** and **java_container_version** must also be specified. Possible values are `1.7` and `1.8`.
- **local_mysql_enabled** - (Optional) Is "MySQL In App" Enabled? This runs a local MySQL instance with your app and shares resources from the App Service plan.

NOTE: MySQL In App is not intended for production environments and will not scale beyond a single instance. Instead you may wish to use Azure Database for MySQL.

- **managed_pipeline_mode** - (Optional) The Managed Pipeline Mode. Possible values are **Integrated** and **Classic**. Defaults to **Integrated**.
- **min_tls_version** - (Optional) The minimum supported TLS version for the app service. Possible values are 1.0, 1.1, and 1.2. Defaults to 1.2 for new app services.
- **php_version** - (Optional) The version of PHP to use in this App Service Slot. Possible values are 5.5, 5.6, 7.0, 7.1 and 7.2.
- **python_version** - (Optional) The version of Python to use in this App Service Slot. Possible values are 2.7 and 3.4.
- **remote_debugging_enabled** - (Optional) Is Remote Debugging Enabled? Defaults to **false**.
- **remote_debugging_version** - (Optional) Which version of Visual Studio should the Remote Debugger be compatible with? Possible values are VS2012, VS2013, VS2015 and VS2017.
- **scm_type** - (Optional) The type of Source Control enabled for this App Service Slot. Defaults to **None**. Possible values are: **BitbucketGit**, **BitbucketHg**, **CodePlexGit**, **CodePlexHg**, **Dropbox**, **ExternalGit**, **ExternalHg**, **GitHub**, **LocalGit**, **None**, **OneDrive**, **Tfs**, **VSO** and **VSTSRM**
- **use_32_bit_worker_process** - (Optional) Should the App Service Slot run in 32 bit mode, rather than 64 bit mode?

Note: Deployment Slots are not supported in the **Free**, **Shared**, or **Basic** App Service Plans.

- **virtual_network_name** - (Optional) The name of the Virtual Network which this App Service Slot should be attached to.
- **websockets_enabled** - (Optional) Should WebSockets be enabled?

A **cors** block supports the following:

- **allowed_origins** - (Optional) A list of origins which should be able to make cross-origin calls. * can be used to allow all calls.
- **support_credentials** - (Optional) Are credentials supported?

A **identity** block supports the following:

- **type** - (Required) Specifies the identity type of the App Service. At this time the only allowed value is **SystemAssigned**.

The assigned `principal_id` and `tenant_id` can be retrieved after the App Service Slot has been created.

A `ip_restriction` block supports the following:

- `ip_address` - (Required) The IP Address used for this IP Restriction.
- `subnet_mask` - (Optional) The Subnet mask used for this IP Restriction. Defaults to `255.255.255.255`.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the App Service Slot.
- `default_site_hostname` - The Default Hostname associated with the App Service Slot - such as `mysite.azurewebsites.net`

» Import

App Service Slots can be imported using the `resource id`, e.g.

```
terraform import azurerm_app_service_slot.instance1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm__function__app`

Manages a Function App.

» Example Usage (with App Service Plan)

```
resource "azurerm_resource_group" "test" {
  name     = "azure-functions-test-rg"
  location = "westus2"
}

resource "azurerm_storage_account" "test" {
  name                        = "functionsapptestsa"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  location                   = "${azurerm_resource_group.test.location}"
  account_tier               = "Standard"
  account_replication_type   = "LRS"
}
```

```

resource "azurerm_app_service_plan" "test" {
  name                = "azure-functions-test-service-plan"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  sku {
    tier = "Standard"
    size = "S1"
  }
}

resource "azurerm_function_app" "test" {
  name                = "test-azure-functions"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  app_service_plan_id = "${azurerm_app_service_plan.test.id}"
  storage_connection_string = "${azurerm_storage_account.test.primary_connection_string}"
}

```

» Example Usage (in a Consumption Plan)

```

resource "azurerm_resource_group" "test" {
  name     = "azure-functions-cptest-rg"
  location = "westus2"
}

resource "azurerm_storage_account" "test" {
  name                = "functionsapptestsa"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location             = "${azurerm_resource_group.test.location}"
  account_tier         = "Standard"
  account_replication_type = "LRS"
}

resource "azurerm_app_service_plan" "test" {
  name                = "azure-functions-test-service-plan"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  kind                 = "FunctionApp"

  sku {
    tier = "Dynamic"
    size = "Y1"
  }
}

```

```

}

resource "azurerm_function_app" "test" {
  name                       = "test-azure-functions"
  location                  = "${azurerm_resource_group.test.location}"
  resource_group_name       = "${azurerm_resource_group.test.name}"
  app_service_plan_id       = "${azurerm_app_service_plan.test.id}"
  storage_connection_string = "${azurerm_storage_account.test.primary_connection_string}"
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Function App. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the Function App.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **app_service_plan_id** - (Required) The ID of the App Service Plan within which to create this Function App. Changing this forces a new resource to be created.
- **storage_connection_string** - (Required) The connection string of the backend storage account which will be used by this Function App (such as the dashboard, logs).
- **app_settings** - (Optional) A key-value pair of App Settings.
- **enable_built_in_logging** - (Optional) Should the built-in logging of this Function App be enabled? Defaults to **true**.
- **connection_string** - (Optional) An **connection_string** block as defined below.
- **client_affinity_enabled** - (Optional) Should the Function App send session affinity cookies, which route client requests in the same session to the same instance?
- **enabled** - (Optional) Is the Function App enabled?
- **https_only** - (Optional) Can the Function App only be accessed via HTTPS? Defaults to **false**.
- **version** - (Optional) The runtime version associated with the Function App. Defaults to **~1**.

- `site_config` - (Optional) A `site_config` object as defined below.
 - `identity` - (Optional) An `identity` block as defined below.
 - `tags` - (Optional) A mapping of tags to assign to the resource.
-

`connection_string` supports the following:

- `name` - (Required) The name of the Connection String.
 - `type` - (Required) The type of the Connection String. Possible values are `APIHub`, `Custom`, `DocDb`, `EventHub`, `MySQL`, `NotificationHub`, `PostgreSQL`, `RedisCache`, `ServiceBus`, `SQLAzure` and `SQLServer`.
 - `value` - (Required) The value for the Connection String.
-

`site_config` supports the following:

- `always_on` - (Optional) Should the Function App be loaded at all times? Defaults to `false`.
- `use_32_bit_worker_process` - (Optional) Should the Function App run in 32 bit mode, rather than 64 bit mode? Defaults to `true`.

Note: when using an App Service Plan in the **Free** or **Shared** Tiers `use_32_bit_worker_process` must be set to `true`.

- `websockets_enabled` - (Optional) Should WebSockets be enabled?
 - `linux_fx_version` - (Optional) Linux App Framework and version for the AppService, e.g. `DOCKER|golang:latest`.
-

`identity` supports the following:

- `type` - (Required) Specifies the identity type of the App Service. At this time the only allowed value is `SystemAssigned`.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the Function App
- `default_hostname` - The default hostname associated with the Function App - such as `mysite.azurewebsites.net`
- `outbound_ip_addresses` - A comma separated list of outbound IP addresses - such as `52.23.25.3,52.143.43.12`

- `possible_outbound_ip_addresses` - A comma separated list of outbound IP addresses - such as `52.23.25.3,52.143.43.12,52.143.43.17` - not all of which are necessarily in use. Superset of `outbound_ip_addresses`.
- `identity` - An `identity` block as defined below, which contains the Managed Service Identity information for this App Service.
- `site_credential` - A `site_credential` block as defined below, which contains the site-level credentials used to publish to this App Service.
- `kind` - The Function App kind - such as `functionapp,linux,container`

`identity` exports the following:

- `principal_id` - The Principal ID for the Service Principal associated with the Managed Service Identity of this App Service.
- `tenant_id` - The Tenant ID for the Service Principal associated with the Managed Service Identity of this App Service.

`site_credential` exports the following:

- `username` - The username which can be used to publish to this App Service
- `password` - The password associated with the username, which can be used to publish to this App Service.

» Import

Function Apps can be imported using the `resource id`, e.g.

```
terraform import azurerm_function_app.functionapp1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_role_assignment`

Assigns a given Principal (User or Application) to a given Role.

» Example Usage (using a built-in Role)

```
data "azurerm_subscription" "primary" {}

data "azurerm_client_config" "test" {}

resource "azurerm_role_assignment" "test" {
  scope                = "${data.azurerm_subscription.primary.id}"
  role_definition_name = "Reader"
```

```

    principal_id      = "${data.azurearm_client_config.test.service_principal_object_id}"
  }

```

» Example Usage (Custom Role & Service Principal)

```

data "azurerm_subscription" "primary" {}

data "azurerm_client_config" "test" {}

resource "azurerm_role_definition" "test" {
  role_definition_id = "00000000-0000-0000-0000-000000000000"
  name               = "my-custom-role-definition"
  scope              = "${data.azurerm_subscription.primary.id}"

  permissions {
    actions      = ["Microsoft.Resources/subscriptions/resourceGroups/read"]
    not_actions = []
  }

  assignable_scopes = [
    "${data.azurerm_subscription.primary.id}",
  ]
}

resource "azurerm_role_assignment" "test" {
  name                = "00000000-0000-0000-0000-000000000000"
  scope               = "${data.azurerm_subscription.primary.id}"
  role_definition_id = "${azurerm_role_definition.test.id}"
  principal_id        = "${data.azurearm_client_config.test.service_principal_object_id}"
}

```

» Example Usage (Custom Role & User)

```

data "azurerm_subscription" "primary" {}

data "azurerm_client_config" "test" {}

resource "azurerm_role_definition" "test" {
  role_definition_id = "00000000-0000-0000-0000-000000000000"
  name               = "my-custom-role-definition"
  scope              = "${data.azurerm_subscription.primary.id}"

  permissions {
    actions      = ["Microsoft.Resources/subscriptions/resourceGroups/read"]

```

```

    not_actions = []
  }

  assignable_scopes = [
    "${data.azurerm_subscription.primary.id}",
  ]
}

resource "azurerm_role_assignment" "test" {
  name           = "00000000-0000-0000-0000-000000000000"
  scope          = "${data.azurerm_subscription.primary.id}"
  role_definition_id = "${azurerm_role_definition.test.id}"
  principal_id    = "${data.azurerm_client_config.test.client_id}"
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Optional) A unique UUID/GUID for this Role Assignment - one will be generated if not specified. Changing this forces a new resource to be created.
- **scope** - (Required) The scope at which the Role Assignment applies too, such as `/subscriptions/0b1f6471-1bf0-4dda-aec3-111122223333`, `/subscriptions/0b1f6471-1bf0-4dda-aec3-111122223333/resourceGroups/myGroup`, or `/subscriptions/0b1f6471-1bf0-4dda-aec3-111122223333/resourceGroups/myGroup/providers/`. Changing this forces a new resource to be created.
- **role_definition_id** - (Optional) The Scoped-ID of the Role Definition. Changing this forces a new resource to be created. Conflicts with **role_definition_name**.
- **role_definition_name** - (Optional) The name of a built-in Role. Changing this forces a new resource to be created. Conflicts with **role_definition_id**.
- **principal_id** - (Required) The ID of the Principal (User or Application) to assign the Role Definition to. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The Role Assignment ID.

» Import

Role Assignments can be imported using the `resource id`, e.g.

```
terraform import azurerm_role_assignment.test /subscriptions/00000000-0000-0000-0000-00000000
```

» azurerm_role_definition

Manages a custom Role Definition, used to assign Roles to Users/Principals.
See 'Understand role definitions' in the Azure documentation for more details.

» Example Usage

```
data "azurerm_subscription" "primary" {}

resource "azurerm_role_definition" "test" {
  name           = "my-custom-role"
  scope          = "${data.azurerm_subscription.primary.id}"
  description    = "This is a custom role created via Terraform"

  permissions {
    actions      = ["*"]
    not_actions = []
  }

  assignable_scopes = [
    "${data.azurerm_subscription.primary.id}", # /subscriptions/00000000-0000-0000-0000-00000000
  ]
}
```

» Argument Reference

The following arguments are supported:

- `role_definition_id` - (Optional) A unique UUID/GUID which identifies this role - one will be generated if not specified. Changing this forces a new resource to be created.
- `name` - (Required) The name of the Role Definition. Changing this forces a new resource to be created.
- `scope` - (Required) The scope at which the Role Definition applies too, such as `/subscriptions/0b1f6471-1bf0-4dda-aec3-111122223333`, `/subscriptions/0b1f6471-1bf0-4dda-aec3-111122223333/resourceGroups/myGroup`,

or `/subscriptions/0b1f6471-1bf0-4dda-aec3-111122223333/resourceGroups/myGroup/providers/`
Changing this forces a new resource to be created.

- **description** - (Optional) A description of the Role Definition.
- **permissions** - (Required) A **permissions** block as defined below.
- **assignable_scopes** - (Required) One or more assignable scopes for this Role Definition, such as `/subscriptions/0b1f6471-1bf0-4dda-aec3-111122223333`, `/subscriptions/0b1f6471-1bf0-4dda-aec3-111122223333/resourceGroups/myGroup`, or `/subscriptions/0b1f6471-1bf0-4dda-aec3-111122223333/resourceGroups/myGroup/providers/`

A **permissions** block as the following properties:

- **actions** - (Optional) One or more Allowed Actions, such as `*`, `Microsoft.Resources/subscriptions/resourceGroups/read`. See 'Azure Resource Manager resource provider operations' for details.
- **data_actions** - (Optional) One or more Allowed Data Actions, such as `*`, `Microsoft.Storage/storageAccounts/blobServices/containers/blobs/read`. See 'Azure Resource Manager resource provider operations' for details.
- **not_actions** - (Optional) One or more Disallowed Actions, such as `*`, `Microsoft.Resources/subscriptions/resourceGroups/read`. See 'Azure Resource Manager resource provider operations' for details.
- **not_data_actions** - (Optional) One or more Disallowed Data Actions, such as `*`, `Microsoft.Resources/subscriptions/resourceGroups/read`. See 'Azure Resource Manager resource provider operations' for details.

» Attributes Reference

The following attributes are exported:

- **id** - The Role Definition ID.

» Import

Role Definitions can be imported using the **resource id**, e.g.

```
terraform import azurerm_role_definition.test /subscriptions/00000000-0000-0000-0000-00000000
```

» `azurerm_user_assigned_identity`

Manages a user assigned identity.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "acceptanceTestResourceGroup1"
  location  = "eastus"
}

resource "azurerm_user_assigned_identity" "testIdentity" {
  resource_group_name = "${azurerm_resource_group.test.name}"
  location             = "${azurerm_resource_group.test.location}"

  name = "search-api"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the user assigned identity. Changing this forces a new identity to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the user assigned identity.
- **location** - (Required) The location/region where the user assigned identity is created.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- **id** - The user assigned identity ID.
- **principal_id** - Service Principal ID associated with the user assigned identity.
- **client_id** - Client ID associated with the user assigned identity.

» Import

User Assigned Identities can be imported using the **resource id**, e.g.

```
terraform import azurerm_user_assigned_identity.testIdentity /subscriptions/00000000-0000-0000-0000-000000000000/resourceGroups/test-rg/providers/Microsoft.ManagedIdentity/userAssignedIdentities/testIdentity
```

» **azurerm_azuread_application**

Manages an Application within Azure Active Directory.

NOTE: The Azure Active Directory resources have been split out into a new AzureAD Provider - as such the AzureAD resources within the AzureRM Provider are deprecated and will be removed in the next major version (2.0). Information on how to migrate from the existing resources to the new AzureAD Provider can be found [here](#).

NOTE: If you're authenticating using a Service Principal then it must have permissions to both `Read` and `write` all applications and `Sign in and read user profile` within the Windows Azure Active Directory API.

» **Example Usage**

```
resource "azurerm_azuread_application" "test" {
  name                = "example"
  homepage            = "https://homepage"
  identifier_uris     = ["https://uri"]
  reply_urls         = ["https://replyurl"]
  available_to_other_tenants = false
  oauth2_allow_implicit_flow = true
}
```

» **Argument Reference**

The following arguments are supported:

- **name** - (Required) The display name for the application.
- **homepage** - (optional) The URL to the application's home page. If no homepage is specified this defaults to `https://{name}`.
- **identifier_uris** - (Optional) A list of user-defined URI(s) that uniquely identify a Web application within it's Azure AD tenant, or within a verified custom domain if the application is multi-tenant.
- **reply_urls** - (Optional) A list of URLs that user tokens are sent to for sign in, or the redirect URIs that OAuth 2.0 authorization codes and access tokens are sent to.
- **available_to_other_tenants** - (Optional) Is this Azure AD Application available to other tenants? Defaults to `false`.
- **oauth2_allow_implicit_flow** - (Optional) Does this Azure AD Application allow OAuth2.0 implicit flow tokens? Defaults to `false`.

» Attributes Reference

The following attributes are exported:

- `application_id` - The Application ID.

» Import

Azure Active Directory Applications can be imported using the `object id`, e.g.

```
terraform import azurerm_azuread_application.test 00000000-0000-0000-0000-000000000000
```

» `azurerm_azuread_service_principal`

Manages a Service Principal associated with an Application within Azure Active Directory.

NOTE: The Azure Active Directory resources have been split out into a new AzureAD Provider - as such the AzureAD resources within the AzureRM Provider are deprecated and will be removed in the next major version (2.0). Information on how to migrate from the existing resources to the new AzureAD Provider can be found [here](#).

NOTE: If you're authenticating using a Service Principal then it must have permissions to both `Read and write all applications` and `Sign in and read user profile` within the Windows Azure Active Directory API.

» Example Usage

```
resource "azurerm_azuread_application" "test" {
  name                = "example"
  homepage            = "http://homepage"
  identifier_uris     = ["http://uri"]
  reply_urls         = ["http://replyurl"]
  available_to_other_tenants = false
  oauth2_allow_implicit_flow = true
}

resource "azurerm_azuread_service_principal" "test" {
  application_id = "${azurerm_azuread_application.test.application_id}"
}
```

» Argument Reference

The following arguments are supported:

- **application_id** - (Required) The ID of the Azure AD Application for which to create a Service Principal.

» Attributes Reference

The following attributes are exported:

- **id** - The Object ID for the Service Principal.
- **display_name** - The Display Name of the Azure Active Directory Application associated with this Service Principal.

» Import

Azure Active Directory Service Principals can be imported using the **object id**, e.g.

```
terraform import azurerm_azuread_service_principal.test 00000000-0000-0000-0000-000000000000
```

» azurerm_azuread_service_principal_password

Manages a Password associated with a Service Principal within Azure Active Directory.

NOTE: The Azure Active Directory resources have been split out into a new AzureAD Provider - as such the AzureAD resources within the AzureRM Provider are deprecated and will be removed in the next major version (2.0). Information on how to migrate from the existing resources to the new AzureAD Provider can be found [here](#).

NOTE: If you're authenticating using a Service Principal then it must have permissions to both **Read** and **write** all applications and **Sign in and read user profile** within the Windows Azure Active Directory API.

» Example Usage

```
resource "azurerm_azuread_application" "test" {  
  name                = "example"  
  homepage            = "http://homepage"  
  identifier_uris     = ["http://uri"]  
}
```

```

    reply_urls                = ["http://replyurl"]
    available_to_other_tenants = false
    oauth2_allow_implicit_flow = true
}

resource "azurerm_azuread_service_principal" "test" {
  application_id = "${azurerm_azuread_application.test.application_id}"
}

resource "azurerm_azuread_service_principal_password" "test" {
  service_principal_id = "${azurerm_azuread_service_principal.test.id}"
  value                = "VT=uSgbTanZhyz@%nL9Hpd+Tfay_MRV#"
  end_date              = "2020-01-01T01:02:03Z"
}

```

» Argument Reference

The following arguments are supported:

- **service_principal_id** - (Required) The ID of the Service Principal for which this password should be created. Changing this field forces a new resource to be created.
- **value** - (Required) The Password for this Service Principal.
- **end_date** - (Required) The End Date which the Password is valid until, formatted as a RFC3339 date string (e.g. 2018-01-01T01:02:03Z). Changing this field forces a new resource to be created.
- **key_id** - (Optional) A GUID used to uniquely identify this Key. If not specified a GUID will be created. Changing this field forces a new resource to be created.
- **start_date** - (Optional) The Start Date which the Password is valid from, formatted as a RFC3339 date string (e.g. 2018-01-01T01:02:03Z). If this isn't specified, the current date is used. Changing this field forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The Key ID for the Service Principal Password.

» Import

Service Principal Passwords can be imported using the `object_id`, e.g.

```
terraform import azurerm_azuread_service_principal_password.test 00000000-0000-0000-0000-000000000000
```

NOTE: This ID format is unique to Terraform and is composed of the Service Principal's Object ID and the Service Principal Password's Key ID in the format `{ServicePrincipalObjectId}/{ServicePrincipalPasswordKeyId}`.

» `azurerm_automation_account`

Manages a Automation Account.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "resourceGroup1"
  location = "West Europe"
}

resource "azurerm_automation_account" "example" {
  name                = "automationAccount1"
  location             = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"

  sku {
    name = "Basic"
  }

  tags = {
    environment = "development"
  }
}
```

» Argument Reference

The following arguments are supported:

- `name` - (Required) Specifies the name of the Automation Account. Changing this forces a new resource to be created.

- **resource_group_name** - (Required) The name of the resource group in which the Automation Account is created. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **sku** - (Required) A **sku** block as defined below.
- **tags** - (Optional) A mapping of tags to assign to the resource.

sku supports the following:

- **name** - (Optional) The SKU name of the account - only **Basic** is supported at this time. Defaults to **Basic**.

» Attributes Reference

The following attributes are exported:

- **id** - The Automation Account ID.
- **dsc_server_endpoint** - The DSC Server Endpoint associated with this Automation Account.
- **dsc_primary_access_key** - The Primary Access Key for the DSC Endpoint associated with this Automation Account.
- **dsc_secondary_access_key** - The Secondary Access Key for the DSC Endpoint associated with this Automation Account.

» Import

Automation Accounts can be imported using the **resource id**, e.g.

```
terraform import azurerm_automation_account.account1 /subscriptions/00000000-0000-0000-0000-
```

» azurerm_automation_credential

Manages a Automation Credential.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "resourceGroup1"
  location = "West Europe"
```

```

}

resource "azurerm_automation_account" "example" {
  name           = "account1"
  location       = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"

  sku {
    name = "Basic"
  }
}

resource "azurerm_automation_credential" "example" {
  name           = "credential1"
  resource_group_name = "${azurerm_resource_group.example.name}"
  account_name   = "${azurerm_automation_account.example.name}"
  username       = "example_user"
  password       = "example_pwd"
  description    = "This is an example credential"
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Credential. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the Credential is created. Changing this forces a new resource to be created.
- **account_name** - (Required) The name of the automation account in which the Credential is created. Changing this forces a new resource to be created.
- **username** - (Required) The username associated with this Automation Credential.
- **password** - (Required) The password associated with this Automation Credential.
- **description** - (Optional) The description associated with this Automation Credential.

» Attributes Reference

The following attributes are exported:

- `id` - The Automation Credential ID.

» Import

Automation Credentials can be imported using the `resource id`, e.g.

```
terraform import azurerm_automation_credential.credential1 /subscriptions/00000000-0000-0000
```

» `azurerm_automation_dsc_configuration`

Manages a Automation DSC Configuration.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "resourceGroup1"
  location = "West Europe"
}

resource "azurerm_automation_account" "example" {
  name                = "account1"
  location             = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"

  sku {
    name = "Basic"
  }
}

resource "azurerm_automation_dsc_configuration" "example" {
  name                = "test"
  resource_group_name = "${azurerm_resource_group.example.name}"
  automation_account_name = "${azurerm_automation_account.example.name}"
  location             = "${azurerm_resource_group.example.location}"
  content_embedded     = "configuration test {}"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the DSC Configuration. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the DSC Configuration is created. Changing this forces a new resource to be created.
- **automation_account_name** - (Required) The name of the automation account in which the DSC Configuration is created. Changing this forces a new resource to be created.
- **content_embedded** - (Required) The PowerShell DSC Configuration script.
- **location** - (Required) Must be the same location as the Automation Account.
- **log_verbose** - (Optional) Verbose log option.
- **description** - (Optional) Description to go with DSC Configuration.

» Attributes Reference

The following attributes are exported:

- **id** - The DSC Configuration ID.

» **azurerm_automation_dsc_nodeconfiguration**

Manages a Automation DSC Node Configuration.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name      = "resourceGroup1"
  location = "West Europe"
}

resource "azurerm_automation_account" "example" {
  name              = "account1"
  location          = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
}
```

```

    sku {
        name = "Basic"
    }
}

resource "azurerm_automation_dsc_configuration" "example" {
    name = "test"
    resource_group_name = "${azurerm_resource_group.example.name}"
    automation_account_name = "${azurerm_automation_account.example.name}"
    location = "${azurerm_resource_group.example.location}"
    content_embedded = "configuration test {}"
}

resource "azurerm_automation_dsc_nodeconfiguration" "example" {
    name = "test.localhost"
    resource_group_name = "${azurerm_resource_group.example.name}"
    automation_account_name = "${azurerm_automation_account.example.name}"
    depends_on = ["azurerm_automation_dsc_configuration.example"]

    content_embedded = <<mofcontent
instance of MSFT_FileDirectoryConfiguration as $MSFT_FileDirectoryConfiguration1ref
{
    ResourceID = "[File]bla";
    Ensure = "Present";
    Contents = "bogus Content";
    DestinationPath = "c:\\\\bogus.txt";
    ModuleName = "PSDesiredStateConfiguration";
    SourceInfo = ".:3:9::file";
    ModuleVersion = "1.0";
    ConfigurationName = "bla";
};
instance of OMI_ConfigurationDocument
{
    Version="2.0.0";
    MinimumCompatibleVersion = "1.0.0";
    CompatibleVersionAdditionalProperties= {"Omi_BaseResource:ConfigurationName"};
    Author="bogusAuthor";
    GenerationDate="06/15/2018 14:06:24";
    GenerationHost="bogusComputer";
    Name="test";
};
mofcontent
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the DSC Node Configuration. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the DSC Node Configuration is created. Changing this forces a new resource to be created.
- **automation_account_name** - (Required) The name of the automation account in which the DSC Node Configuration is created. Changing this forces a new resource to be created.
- **content_embedded** - (Required) The PowerShell DSC Node Configuration (mof content).

» Attributes Reference

The following attributes are exported:

- **id** - The DSC Node Configuration ID.

» azurerm_automation_module

Manages a Automation Module.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name      = "resourceGroup1"
  location  = "West Europe"
}

resource "azurerm_automation_account" "example" {
  name                = "account1"
  location             = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"

  sku {
    name = "Basic"
  }
}
```

```

resource "azurerm_automation_module" "example" {
  name                = "xActiveDirectory"
  resource_group_name = "${azurerm_resource_group.example.name}"
  automation_account_name = "${azurerm_automation_account.example.name}"

  module_link {
    uri = "https://devopsgallerystorage.blob.core.windows.net/packages/xactivedirectory.2.19"
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Module. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the Module is created. Changing this forces a new resource to be created.
- **automation_account_name** - (Required) The name of the automation account in which the Module is created. Changing this forces a new resource to be created.
- **module_link** - (Required) The published Module link.

module_link supports the following:

- **uri** - (Required) The uri of the module content (zip or nupkg).

» Attributes Reference

The following attributes are exported:

- **id** - The Automation Module ID.

» azurerm__automation__runbook

Manages a Automation Runbook.

» Example Usage

```

resource "azurerm_resource_group" "example" {
  name = "resourceGroup1"
}

```

```

    location = "West Europe"
  }

  resource "azurerm_automation_account" "example" {
    name           = "account1"
    location       = "${azurerm_resource_group.example.location}"
    resource_group_name = "${azurerm_resource_group.example.name}"

    sku {
      name = "Basic"
    }
  }

  resource "azurerm_automation_runbook" "example" {
    name           = "Get-AzureVMTutorial"
    location       = "${azurerm_resource_group.example.location}"
    resource_group_name = "${azurerm_resource_group.example.name}"
    account_name   = "${azurerm_automation_account.example.name}"
    log_verbose    = "true"
    log_progress   = "true"
    description    = "This is an example runbook"
    runbook_type   = "PowerShellWorkflow"

    publish_content_link {
      uri = "https://raw.githubusercontent.com/Azure/azure-quickstart-templates/master/101-automate-vm/azure-quickstart-templates/101-automate-vm.json"
    }
  }
}

```

» Example Usage - custom content

```

resource "azurerm_resource_group" "example" {
  name     = "resourceGroup1"
  location = "West Europe"
}

resource "azurerm_automation_account" "example" {
  name           = "account1"
  location       = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"

  sku {
    name = "Basic"
  }
}

```



```

data "local_file" "example" {
  filename = "${path.module}/example.ps1"
}

resource "azurerm_automation_runbook" "example" {
  name                = "Get-AzureVMTutorial"
  location             = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
  account_name        = "${azurerm_automation_account.example.name}"
  log_verbose         = "true"
  log_progress        = "true"
  description         = "This is an example runbook"
  runbook_type        = "PowerShell"

  publish_content_link {
    uri = "https://raw.githubusercontent.com/Azure/azure-quickstart-templates/master/101-automate-vm/azure-powershell/Get-AzureVMTutorial.ps1"
  }

  content = "${data.local_file.example.content}"
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Runbook. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the Runbook is created. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **account_name** - (Required) The name of the automation account in which the Runbook is created. Changing this forces a new resource to be created.
- **runbook_type** - (Required) The type of the runbook - can be either **Graph**, **GraphPowerShell**, **GraphPowerShellWorkflow**, **PowerShellWorkflow**, **PowerShell** or **Script**.
- **log_progress** - (Required) Progress log option.
- **log_verbose** - (Required) Verbose log option.
- **publish_content_link** - (Required) The published runbook content link.
- **description** - (Optional) A description for this credential.

- **content** - (Optional) The desired content of the runbook.

NOTE The Azure API requires a **publish_content_link** to be supplied even when specifying your own **content**.

NOTE Setting **content** to an empty string will revert the runbook to the **publish_content_link**.

- **tags** - (Optional) A mapping of tags to assign to the resource.

publish_content_link supports the following:

- **uri** - (Required) The uri of the runbook content.

» Attributes Reference

The following attributes are exported:

- **id** - The Automation Runbook ID.

» Import

Automation Runbooks can be imported using the **resource id**, e.g.

```
terraform import azurerm_automation_runbook.Get-AzureVMTutorial /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_automation_schedule

Manages a Automation Schedule.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "tfex-automation-account"
  location = "West Europe"
}

resource "azurerm_automation_account" "example" {
  name                = "tfex-automation-account"
  location            = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"

  sku {
    name = "Basic"
  }
}
```

```

}

resource "azurerm_automation_schedule" "example" {
  name                        = "tfex-automation-schedule"
  resource_group_name        = "${azurerm_resource_group.example.name}"
  automation_account_name    = "${azurerm_automation_account.example.name}"
  frequency                  = "Week"
  interval                   = 1
  timezone                   = "Central Europe Standard Time"
  start_time                 = "2014-04-15T18:00:15+02:00"
  description                 = "This is an example schedule"

  advanced_schedule {
    week_days = ["Friday"]
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Schedule. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the Schedule is created. Changing this forces a new resource to be created.
- **automation_account_name** - (Required) The name of the automation account in which the Schedule is created. Changing this forces a new resource to be created.
- **frequency** - (Required) The frequency of the schedule. - can be either `OneTime`, `Day`, `Hour`, `Week`, or `Month`.
- **description** - (Optional) A description for this Schedule.
- **interval** - (Optional) The number of **frequencies** between runs. Only valid when frequency is `Day`, `Hour`, `Week`, or `Month` and defaults to 1.
- **start_time** - (Optional) Start time of the schedule. Must be at least five minutes in the future. Defaults to seven minutes in the future from the time the resource is created.
- **expiry_time** - (Optional) The end time of the schedule.
- **timezone** - (Optional) The timezone of the start time. Defaults to `UTC`. For possible values see: [https://msdn.microsoft.com/en-us/library/ms912391\(v=winembedded.11\).aspx](https://msdn.microsoft.com/en-us/library/ms912391(v=winembedded.11).aspx)

- **week_days** - (Optional) List of days of the week that the job should execute on. Only valid when frequency is **Week**.
- **month_days** - (Optional) List of days of the month that the job should execute on. Must be between 1 and 31. -1 for last day of the month. Only valid when frequency is **Month**.
- **monthly_occurrence** - (Optional) List of occurrences of days within a month. Only valid when frequency is **Month**. The **monthly_occurrence** block supports fields documented below.

The **monthly_occurrence** block supports:

- **day** - (Required) Day of the occurrence. Must be one of **Monday, Tuesday, Wednesday, Thursday, Friday, Saturday, Sunday**.
- **occurrence** - (Required) Occurrence of the week within the month. Must be between 1 and 5. -1 for last week within the month.

» Attributes Reference

The following attributes are exported:

- **id** - The Automation Schedule ID.

» Import

Automation Schedule can be imported using the **resource id**, e.g.

```
terraform import azurerm_automation_schedule.schedule1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_application_insights

Manage an Application Insights component.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "tf-test"
  location = "West Europe"
}

resource "azurerm_application_insights" "test" {
```

```

    name          = "tf-test-appinsights"
    location       = "West Europe"
    resource_group_name = "${azurerm_resource_group.test.name}"
    application_type = "Web"
  }

  output "instrumentation_key" {
    value = "${azurerm_application_insights.test.instrumentation_key}"
  }

  output "app_id" {
    value = "${azurerm_application_insights.test.app_id}"
  }

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Application Insights component. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the Application Insights component.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **application_type** - (Required) Specifies the type of Application Insights to create. Valid values are `ios` for *iOS*, `java` for *Java web*, `MobileCenter` for *App Center*, `Node.JS` for *Node.js*, `other` for *General*, `phone` for *Windows Phone*, `store` for *Windows Store* and `web` for *ASP.NET*. Please note these values are case sensitive; unmatched values are treated as *ASP.NET* by Azure. Changing this forces a new resource to be created.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Application Insights component.
- **app_id** - The App ID associated with this Application Insights component.
- **instrumentation_key** - The Instrumentation Key for this Application Insights component.

» Import

Application Insights instances can be imported using the `resource id`, e.g.

```
terraform import azurerm_application_insights.instance1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_application_insights_api_key`

Manages an Application Insights API key.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "tf-test"
  location = "West Europe"
}

resource "azurerm_application_insights" "test" {
  name                = "tf-test-appinsights"
  location            = "West Europe"
  resource_group_name = "${azurerm_resource_group.test.name}"
  application_type     = "web"
}

resource "azurerm_application_insights_api_key" "read_telemetry" {
  name                = "tf-test-appinsights-read-telemetry-api-key"
  application_insights_id = "${azurerm_application_insights.test.id}"
  read_permissions    = ["aggregate", "api", "draft", "extendqueries", "search"]
}

resource "azurerm_application_insights_api_key" "write_annotations" {
  name                = "tf-test-appinsights-write-annotations-api-key"
  application_insights_id = "${azurerm_application_insights.test.id}"
  write_permissions   = ["annotations"]
}

resource "azurerm_application_insights_api_key" "authenticate_sdk_control_channel" {
  name                = "tf-test-appinsights-authenticate-sdk-control-channel-api-key"
  application_insights_id = "${azurerm_application_insights.test.id}"
  read_permissions    = ["agentconfig"]
}

resource "azurerm_application_insights_api_key" "full_permissions" {
  name                = "tf-test-appinsights-full-permissions-api-key"
```

```

    application_insights_id = "${azurerm_application_insights.test.id}"
    read_permissions        = ["agentconfig", "aggregate", "api", "draft", "extendqueries", "search"]
    write_permissions       = ["annotations"]
  }

  output "read_telemetry_api_key" {
    value = "${azurerm_application_insights_api_key.read_telemetry.api_key}"
  }

  output "write_annotations_api_key" {
    value = "${azurerm_application_insights_api_key.write_annotations.api_key}"
  }

  output "authenticate_sdk_control_channel" {
    value = "${azurerm_application_insights_api_key.authenticate_sdk_control_channel.api_key}"
  }

  output "full_permissions_api_key" {
    value = "${azurerm_application_insights_api_key.full_permissions.api_key}"
  }

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Application Insights API key. Changing this forces a new resource to be created.
- **application_insights_id** - (Required) The ID of the Application Insights component on which the API key operates. Changing this forces a new resource to be created.
- **read_permissions** - (Optional) Specifies the list of read permissions granted to the API key. Valid values are `agentconfig`, `aggregate`, `api`, `draft`, `extendqueries`, `search`. Please note these values are case sensitive. Changing this forces a new resource to be created.
- **write_permissions** - (Optional) Specifies the list of write permissions granted to the API key. Valid values are `annotations`. Please note these values are case sensitive. Changing this forces a new resource to be created.

Note: At least one read or write permission must be defined.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the Application Insights API key.
- `api_key` - The API Key secret (Sensitive).

» Import

Application Insights API keys can be imported using the `resource id`, e.g.

```
terraform import azurerm_application_insights_api_key.my_key/subscriptions/00000000-0000-0000-0000-000000000000
```

Note: The secret `api_key` cannot be retrieved during an import. You will need to edit the state by hand to set the secret value if you happen to have it backed up somewhere.

» `azurerm_batch_account`

Manages an Azure Batch account.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "testbatch"
  location = "westeurope"
}

resource "azurerm_storage_account" "test" {
  name                        = "teststorage"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  location                   = "${azurerm_resource_group.test.location}"
  account_tier               = "Standard"
  account_replication_type   = "LRS"
}

resource "azurerm_batch_account" "test" {
  name                        = "testbatchaccount"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  location                   = "${azurerm_resource_group.test.location}"
  pool_allocation_mode       = "BatchService"
  storage_account_id         = "${azurerm_storage_account.test.id}"

  tags = {
    env = "test"
  }
}
```


» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Batch account. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the Batch account. Changing this forces a new resource to be created.

NOTE: To work around a bug in the Azure API this property is currently treated as case-insensitive. A future version of Terraform will require that the casing is correct.

- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **pool_allocation_mode** - (Optional) Specifies the mode to use for pool allocation. Possible values are **BatchService** or **UserSubscription**. Defaults to **BatchService**.
- **storage_account_id** - (Optional) Specifies the storage account to use for the Batch account. If not specified, Azure Batch will manage the storage.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- **id** - The Batch account ID.
- **primary_access_key** - The Batch account primary access key.
- **secondary_access_key** - The Batch account secondary access key.
- **account_endpoint** - The account endpoint used to interact with the Batch service.

NOTE: Primary and secondary access keys are only available when **pool_allocation_mode** is set to **BatchService**. See documentation for more information.

» `azurerm_batch_certificate`

Manages a certificate in an Azure Batch account.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "testbatch"
  location  = "westeurope"
}

resource "azurerm_storage_account" "test" {
  name                        = "teststorage"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  location                   = "${azurerm_resource_group.test.location}"
  account_tier                = "Standard"
  account_replication_type    = "LRS"
}

resource "azurerm_batch_account" "test" {
  name                      = "testbatchaccount"
  resource_group_name       = "${azurerm_resource_group.test.name}"
  location                  = "${azurerm_resource_group.test.location}"
  pool_allocation_mode      = "BatchService"
  storage_account_id        = "${azurerm_storage_account.test.id}"

  tags = {
    env = "test"
  }
}

resource "azurerm_batch_certificate" "test" {
  resource_group_name = "${azurerm_resource_group.test.name}"
  account_name         = "${azurerm_batch_account.test.name}"
  certificate           = "${base64encode(file("certificate.pfx"))}"
  format               = "Pfx"
  password             = "terraform"
  thumbprint           = "42C107874FD0E4A9583292A2F1098E8FE4B2EDDA"
  thumbprint_algorithm = "SHA1"
}
```

» Argument Reference

The following arguments are supported:

- **account_name** - (Required) Specifies the name of the Batch account. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the Batch account. Changing this forces a new resource

to be created.

- **certificate** - (Required) The base64-encoded contents of the certificate.
- **format** - (Required) The format of the certificate. Possible values are **Cer** or **Pfx**.
- **password** - (Optional) The password to access the certificate's private key. This must and can only be specified when **format** is **Pfx**.
- **thumbprint** - (Required) The thumbprint of the certificate. At this time the only supported value is 'SHA1'.

» Attributes Reference

The following attributes are exported:

- **id** - The Batch certificate ID.
- **name** - The generated name of the certificate.
- **public_data** - The public key of the certificate.

» **azurerm__batch__pool**

Manages an Azure Batch pool.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "testaccbatch"
  location  = "%s"
}

resource "azurerm_storage_account" "test" {
  name                = "testaccsa"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location            = "${azurerm_resource_group.test.location}"
  account_tier        = "Standard"
  account_replication_type = "LRS"
}

resource "azurerm_batch_account" "test" {
  name                = "testaccbatch"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location            = "${azurerm_resource_group.test.location}"
}
```

```

pool_allocation_mode = "BatchService"
storage_account_id   = "${azurerm_storage_account.test.id}"

tags = {
  env = "test"
}
}

resource "azurerm_batch_certificate" "testcer" {
  resource_group_name = "${azurerm_resource_group.test.name}"
  account_name        = "${azurerm_batch_account.test.name}"
  certificate          = "${base64encode(file("certificate.cer"))}"
  format              = "Cer"
  thumbprint          = "312d31a79fa0cef49c00f769afc2b73e9f4edf34"
  thumbprint_algorithm = "SHA1"
}

resource "azurerm_batch_pool" "test" {
  name                = "testaccpool"
  resource_group_name = "${azurerm_resource_group.test.name}"
  account_name        = "${azurerm_batch_account.test.name}"
  display_name        = "Test Acc Pool Auto"
  vm_size             = "Standard_A1"
  node_agent_sku_id   = "batch.node.ubuntu 16.04"

  auto_scale {
    evaluation_interval = "PT15M"
    formula             = <<EOF
      startingNumberOfVMs = 1;
      maxNumberOfVMs      = 25;
      pendingTaskSamplePercent = $PendingTasks.GetSamplePercent(180 * TimeInterval_Second);
      pendingTaskSamples = pendingTaskSamplePercent < 70 ? startingNumberOfVMs : avg($Pending
        $TargetDedicatedNodes=min(maxNumberOfVMs, pendingTaskSamples);
    EOF
  }

  storage_image_reference {
    publisher = "Canonical"
    offer     = "UbuntuServer"
    sku       = "16.04.0-LTS"
    version   = "latest"
  }

  start_task {
    command_line      = "echo 'Hello World from $env'"
    max_task_retry_count = 1
  }
}

```

```

wait_for_success      = true

environment = {
  env = "TEST"
}

user_identity {
  auto_user {
    elevation_level = "NonAdmin"
    scope           = "Task"
  }
}

certificate = {
  id           = "${azurerm_batch_certificate.testcer.id}"
  visibility = [ "StartTask" ]
}
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Batch pool. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the Batch pool. Changing this forces a new resource to be created.

NOTE: To work around a bug in the Azure API this property is currently treated as case-insensitive. A future version of Terraform will require that the casing is correct.

- **account_name** - (Required) Specifies the name of the Batch account in which the pool will be created. Changing this forces a new resource to be created.
- **node_agent_sku_id** - (Required) Specifies the Sku of the node agents that will be created in the Batch pool.
- **vm_size** - (Required) Specifies the size of the VM created in the Batch pool.
- **storage_image_reference** - (Required) A `storage_image_reference` for the virtual machines that will compose the Batch pool.
- **display_name** - (Optional) Specifies the display name of the Batch pool.

- **max_tasks_per_node** - (Optional) Specifies the maximum number of tasks that can run concurrently on a single compute node in the pool. Defaults to 1. Changing this forces a new resource to be created.
- **fixed_scale** - (Optional) A **fixed_scale** block that describes the scale settings when using fixed scale.
- **auto_scale** - (Optional) A **auto_scale** block that describes the scale settings when using auto scale.
- **start_task** - (Optional) A **start_task** block that describes the start task settings for the Batch pool.
- **certificate** - (Optional) One or more **certificate** blocks that describe the certificates to be installed on each compute node in the pool.

NOTE: For Windows compute nodes, the Batch service installs the certificates to the specified certificate store and location. For Linux compute nodes, the certificates are stored in a directory inside the task working directory and an environment variable **AZ_BATCH_CERTIFICATES_DIR** is supplied to the task to query for this location. For certificates with visibility of **remoteUser**, a **certs** directory is created in the user's home directory (e.g., **/home/{user-name}/certs**) and certificates are placed in that directory.

Please Note: **fixed_scale** and **auto_scale** blocks cannot be used both at the same time.

A **fixed_scale** block supports the following:

- **target_dedicated_nodes** - (Optional) The number of nodes in the Batch pool. Defaults to 1.
- **target_low_priority_nodes** - (Optional) The number of low priority nodes in the Batch pool. Defaults to 0.
- **resize_timeout** - (Optional) The timeout for resize operations. Defaults to PT15M.

A **auto_scale** block supports the following:

- **evaluation_interval** - (Optional) The interval to wait before evaluating if the pool needs to be scaled. Defaults to PT15M.
- **formula** - (Required) The autoscale formula that needs to be used for scaling the Batch pool.

A **start_task** block supports the following:

- **command_line** - (Required) The command line executed by the start task.

- **max_task_retry_count** - (Optional) The number of retry count. Defaults to 1.
- **wait_for_success** - (Optional) A flag that indicates if the Batch pool should wait for the start task to be completed. Default to **false**.
- **environment** - (Optional) A map of strings (key,value) that represents the environment variables to set in the start task.
- **user_identity** - (Required) A **user_identity** block that describes the user identity under which the start task runs.
- **resource_file** - (Optional) One or more **resource_file** blocks that describe the files to be downloaded to a compute node.

A **user_identity** block supports the following:

- **user_name** - (Optional) The username to be used by the Batch pool start task.
- **auto_user** - (Optional) A **auto_user** block that describes the user identity under which the start task runs.

Please Note: **user_name** and **auto_user** blocks cannot be used both at the same time, but you need to define one or the other.

A **auto_user** block supports the following:

- **elevation_level** - (Optional) The elevation level of the user identity under which the start task runs. Possible values are **Admin** or **NonAdmin**. Defaults to **NonAdmin**.
- **scope** - (Optional) The scope of the user identity under which the start task runs. Possible values are **Task** or **Pool**. Defaults to **Task**.

A **certificate** block supports the following:

- **id** - (Required) The ID of the Batch Certificate to install on the Batch Pool, which must be inside the same Batch Account.
- **store_location** - (Required) The location of the certificate store on the compute node into which to install the certificate. Possible values are **CurrentUser** or **LocalMachine**.

NOTE: This property is applicable only for pools configured with Windows nodes (that is, created with `cloudServiceConfiguration`, or with `virtualMachineConfiguration` using a Windows image reference). For Linux compute nodes, the certificates are stored in a directory inside the task working directory and an environment variable `AZ_BATCH_CERTIFICATES_DIR` is supplied

to the task to query for this location. For certificates with visibility of `remoteUser`, a 'certs' directory is created in the user's home directory (e.g., `/home/{user-name}/certs`) and certificates are placed in that directory.

- **store_name** - (Optional) The name of the certificate store on the compute node into which to install the certificate. This property is applicable only for pools configured with Windows nodes (that is, created with `cloudServiceConfiguration`, or with `virtualMachineConfiguration` using a Windows image reference). Common store names include: `My`, `Root`, `CA`, `Trust`, `Disallowed`, `TrustedPeople`, `TrustedPublisher`, `AuthRoot`, `AddressBook`, but any custom store name can also be used. The default value is `My`.
- **visibility** - (Optional) Which user accounts on the compute node should have access to the private data of the certificate.

A `resource_file` block supports the following:

- **auto_storage_container_name** - (Optional) The storage container name in the auto storage account.
- **blob_prefix** - (Optional) The blob prefix to use when downloading blobs from an Azure Storage container. Only the blobs whose names begin with the specified prefix will be downloaded. The property is valid only when **auto_storage_container_name** or **storage_container_url** is used. This prefix can be a partial filename or a subdirectory. If a prefix is not specified, all the files in the container will be downloaded.
- **file_mode** - (Optional) The file permission mode represented as a string in octal format (e.g. `"0644"`). This property applies only to files being downloaded to Linux compute nodes. It will be ignored if it is specified for a `resource_file` which will be downloaded to a Windows node. If this property is not specified for a Linux node, then a default value of `0770` is applied to the file.
- **file_path** - (Optional) The location on the compute node to which to download the file, relative to the task's working directory. If the **http_url** property is specified, the **file_path** is required and describes the path which the file will be downloaded to, including the filename. Otherwise, if the **auto_storage_container_name** or **storage_container_url** property is specified, **file_path** is optional and is the directory to download the files to. In the case where **file_path** is used as a directory, any directory structure already associated with the input data will be retained in full and appended to the specified `filePath` directory. The specified relative path cannot break out of the task's working directory (for example by using `..'`).
- **http_url** - (Optional) The URL of the file to download. If the URL is

Azure Blob Storage, it must be readable using anonymous access; that is, the Batch service does not present any credentials when downloading the blob. There are two ways to get such a URL for a blob in Azure storage: include a Shared Access Signature (SAS) granting read permissions on the blob, or set the ACL for the blob or its container to allow public access.

- **storage_container_url** - (Optional) The URL of the blob container within Azure Blob Storage. This URL must be readable and listable using anonymous access; that is, the Batch service does not present any credentials when downloading the blob. There are two ways to get such a URL for a blob in Azure storage: include a Shared Access Signature (SAS) granting read and list permissions on the blob, or set the ACL for the blob or its container to allow public access.

Please Note: Exactly one of `auto_storage_container_name`, `storage_container_url` and `auto_user` must be specified.

» Attributes Reference

The following attributes are exported:

- `id` - The Batch pool ID.

» `azurerm_cdn_endpoint`

A CDN Endpoint is the entity within a CDN Profile containing configuration information regarding caching behaviors and origins. The CDN Endpoint is exposed using the URL format `.azureedge.net`.

» Example Usage

```
resource "random_id" "server" {
  keepers = {
    azi_id = 1
  }

  byte_length = 8
}

resource "azurerm_resource_group" "test" {
  name     = "acceptanceTestResourceGroup1"
  location = "West US"
}
```

```

resource "azurerm_cdn_profile" "test" {
  name          = "exampleCdnProfile"
  location      = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  sku          = "Standard_Verizon"
}

resource "azurerm_cdn_endpoint" "test" {
  name          = "${random_id.server.hex}"
  profile_name  = "${azurerm_cdn_profile.test.name}"
  location      = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  origin {
    name      = "exampleCdnOrigin"
    host_name = "www.example.com"
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the CDN Endpoint. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the CDN Endpoint.
- **profile_name** - (Required) The CDN Profile to which to attach the CDN Endpoint.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **is_http_allowed** - (Optional) Defaults to `true`.
- **is_https_allowed** - (Optional) Defaults to `true`.
- **content_types_to_compress** - (Optional) An array of strings that indicates a content types on which compression will be applied. The value for the elements should be MIME types.
- **geo_filter** - (Optional) A set of Geo Filters for this CDN Endpoint. Each `geo_filter` block supports fields documented below.
- **is_compression_enabled** - (Optional) Indicates whether compression is to be enabled. Defaults to `false`.

- **querystring_caching_behaviour** - (Optional) Sets query string caching behavior. Allowed values are `IgnoreQueryString`, `BypassCaching` and `UseQueryString`. Defaults to `IgnoreQueryString`.
- **optimization_type** - (Optional) What types of optimization should this CDN Endpoint optimize for? Possible values include `DynamicSiteAcceleration`, `GeneralMediaStreaming`, `GeneralWebDelivery`, `LargeFileDownload` and `VideoOnDemandMediaStreaming`.
- **origin** - (Optional) The set of origins of the CDN endpoint. When multiple origins exist, the first origin will be used as primary and rest will be used as failover options. Each **origin** block supports fields documented below.
- **origin_host_header** - (Optional) The host header CDN provider will send along with content requests to origins. Defaults to the host name of the origin.
- **origin_path** - (Optional) The path used at for origin requests.
- **probe_path** - (Optional) the path to a file hosted on the origin which helps accelerate delivery of the dynamic content and calculate the most optimal routes for the CDN. This is relative to the **origin_path**.
- **tags** - (Optional) A mapping of tags to assign to the resource.

The **origin** block supports:

- **name** - (Required) The name of the origin. This is an arbitrary value. However, this value needs to be unique under the endpoint. Changing this forces a new resource to be created.
- **host_name** - (Required) A string that determines the hostname/IP address of the origin server. This string can be a domain name, Storage Account endpoint, Web App endpoint, IPv4 address or IPv6 address. Changing this forces a new resource to be created.
- **http_port** - (Optional) The HTTP port of the origin. Defaults to 80. Changing this forces a new resource to be created.
- **https_port** - (Optional) The HTTPS port of the origin. Defaults to 443. Changing this forces a new resource to be created.

The **geo_filter** block supports:

- **relative_path** - (Required) The relative path applicable to geo filter.
- **action** - (Required) The Action of the Geo Filter. Possible values include `Allow` and `Block`.
- **country_codes** - (Required) A List of two letter country codes (e.g. `US`, `GB`) to be associated with this Geo Filter.

» Attributes Reference

The following attributes are exported:

- `id` - The CDN Endpoint ID.

» Import

CDN Endpoints can be imported using the `resource id`, e.g.

```
terraform import azurem_cdn_endpoint.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_cdn_profile`

Manage a CDN Profile to create a collection of CDN Endpoints.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "resourceGroup1"
  location = "West US"
}

resource "azurerm_cdn_profile" "test" {
  name                = "exampleCdnProfile"
  location            = "West US"
  resource_group_name = "${azurerm_resource_group.test.name}"
  sku                 = "Standard_Verizon"

  tags = {
    environment = "Production"
    cost_center = "MSFT"
  }
}
```

» Argument Reference

The following arguments are supported:

- `name` - (Required) Specifies the name of the CDN Profile. Changing this forces a new resource to be created.

- **resource_group_name** - (Required) The name of the resource group in which to create the CDN Profile.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **sku** - (Required) The pricing related information of current CDN profile. Accepted values are `Standard_Akamai`, `Standard_ChinaCdn`, `Standard_Microsoft`, `Standard_Verizon` or `Premium_Verizon`.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- **id** - The CDN Profile ID.

» Import

CDN Profiles can be imported using the `resource id`, e.g.

```
terraform import azurerm_cdn_profile.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm__availability__set

Manages an availability set for virtual machines.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "resourceGroup1"
  location = "West US"
}

resource "azurerm_availability_set" "test" {
  name                        = "acceptanceTestAvailabilitySet1"
  location                   = "${azurerm_resource_group.test.location}"
  resource_group_name        = "${azurerm_resource_group.test.name}"

  tags = {
    environment = "Production"
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the availability set. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the availability set. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **platform_update_domain_count** - (Optional) Specifies the number of update domains that are used. Defaults to 5.

NOTE: The number of Update Domains varies depending on which Azure Region you're using - a list can be found [here](#).

- **platform_fault_domain_count** - (Optional) Specifies the number of fault domains that are used. Defaults to 3.

NOTE: The number of Fault Domains varies depending on which Azure Region you're using - a list can be found [here](#).

- **managed** - (Optional) Specifies whether the availability set is managed or not. Possible values are **true** (to specify aligned) or **false** (to specify classic). Default is **false**.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- **id** - The virtual Availability Set ID.

» Import

Availability Sets can be imported using the **resource id**, e.g.

```
terraform import azurerm_availability_set.group1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_image

Manage a custom virtual machine image that can be used to create virtual machines.

» Example Usage Creating from VHD

```
resource "azurerm_resource_group" "test" {
  name      = "acctest"
  location  = "West US"
}

resource "azurerm_image" "test" {
  name            = "acctest"
  location        = "West US"
  resource_group_name = "${azurerm_resource_group.test.name}"

  os_disk {
    os_type   = "Linux"
    os_state  = "Generalized"
    blob_uri  = "{blob_uri}"
    size_gb   = 30
  }
}
```

» Example Usage Creating from Virtual Machine (VM must be generalized beforehand)

```
resource "azurerm_resource_group" "test" {
  name      = "acctest"
  location  = "West US"
}

resource "azurerm_image" "test" {
  name            = "acctest"
  location        = "West US"
  resource_group_name = "${azurerm_resource_group.test.name}"
  source_virtual_machine_id = "{vm_id}"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the image. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the image. Changing this forces a new resource to be created.

- **location** - (Required) Specified the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **source_virtual_machine_id** - (Optional) The Virtual Machine ID from which to create the image.
- **os_disk** - (Optional) One or more **os_disk** elements as defined below.
- **data_disk** - (Optional) One or more **data_disk** elements as defined below.
- **tags** - (Optional) A mapping of tags to assign to the resource.
- **zone_resilient** - (Optional) Is zone resiliency enabled? Defaults to **false**. Changing this forces a new resource to be created.

Note: **zone_resilient** can only be set to **true** if the image is stored in a region that supports availability zones.

os_disk supports the following:

- **os_type** - (Required) Specifies the type of operating system contained in the the virtual machine image. Possible values are: Windows or Linux.
- **os_state** - (Required) Specifies the state of the operating system contained in the blob. Currently, the only value is Generalized.
- **managed_disk_id** - (Optional) Specifies the ID of the managed disk resource that you want to use to create the image.
- **blob_uri** - (Optional) Specifies the URI in Azure storage of the blob that you want to use to create the image.
- **caching** - (Optional) Specifies the caching mode as **ReadWrite**, **ReadOnly**, or **None**. The default is **None**.
- **size_gb** - (Optional) Specifies the size of the image to be created. The target size can't be smaller than the source size.

data_disk supports the following:

- **lun** - (Required) Specifies the logical unit number of the data disk.
- **managed_disk_id** - (Optional) Specifies the ID of the managed disk resource that you want to use to create the image.
- **blob_uri** - (Optional) Specifies the URI in Azure storage of the blob that you want to use to create the image.
- **caching** - (Optional) Specifies the caching mode as **ReadWrite**, **ReadOnly**, or **None**. The default is **None**.
- **size_gb** - (Optional) Specifies the size of the image to be created. The target size can't be smaller than the source size.

» Attributes Reference

The following attributes are exported:

- **id** - The managed image ID.

» Import

Image can be imported using the `resource id`, e.g.

```
terraform import azurerm_image.test /subscriptions/00000000-0000-0000-0000-000000000000/res
```

» azurerm__managed__disk

Manage a managed disk.

» Example Usage with Create Empty

```
resource "azurerm_resource_group" "test" {
  name     = "acctestRG"
  location = "West US 2"
}

resource "azurerm_managed_disk" "test" {
  name                = "acctestmd"
  location             = "West US 2"
  resource_group_name = "${azurerm_resource_group.test.name}"
  storage_account_type = "Standard_LRS"
  create_option       = "Empty"
  disk_size_gb        = "1"

  tags = {
    environment = "staging"
  }
}
```

» Example Usage with Create Copy

```
resource "azurerm_resource_group" "test" {
  name     = "acctestRG"
  location = "West US 2"
}

resource "azurerm_managed_disk" "source" {
  name                = "acctestmd1"
  location             = "West US 2"
  resource_group_name = "${azurerm_resource_group.test.name}"
  storage_account_type = "Standard_LRS"
  create_option       = "Empty"
}
```

```

    disk_size_gb          = "1"

    tags = {
      environment = "staging"
    }
  }

  resource "azurerm_managed_disk" "copy" {
    name                = "acctestmd2"
    location            = "West US 2"
    resource_group_name = "${azurerm_resource_group.test.name}"
    storage_account_type = "Standard_LRS"
    create_option       = "Copy"
    source_resource_id  = "${azurerm_managed_disk.source.id}"
    disk_size_gb       = "1"

    tags = {
      environment = "staging"
    }
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the managed disk. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the managed disk.
- **location** - (Required) Specified the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **storage_account_type** - (Required) The type of storage to use for the managed disk. Allowable values are `Standard_LRS`, `Premium_LRS`, `StandardSSD_LRS` or `UltraSSD_LRS`.
- **create_option** - (Required) The method to use when creating the managed disk. Possible values include:
 - `Import` - Import a VHD file in to the managed disk (VHD specified with `source_uri`).
 - `Empty` - Create an empty managed disk.
 - `Copy` - Copy an existing managed disk or snapshot (specified with `source_resource_id`).
 - `FromImage` - Copy a Platform Image (specified with `image_reference_id`)

- `source_uri` - (Optional) URI to a valid VHD file to be used when `create_option` is `Import`.
- `source_resource_id` - (Optional) ID of an existing managed disk to copy when `create_option` is `Copy`.
- `image_reference_id` - (Optional) ID of an existing platform/marketplace disk image to copy when `create_option` is `FromImage`.
- `os_type` - (Optional) Specify a value when the source of an `Import` or `Copy` operation targets a source that contains an operating system. Valid values are `Linux` or `Windows`
- `disk_size_gb` - (Optional, Required for a new managed disk) Specifies the size of the managed disk to create in gigabytes. If `create_option` is `Copy` or `FromImage`, then the value must be equal to or greater than the source's size.
- `encryption_settings` - (Optional) an `encryption_settings` block as defined below.
- `tags` - (Optional) A mapping of tags to assign to the resource.
- `zones` - (Optional) A collection containing the availability zone to allocate the Managed Disk in.

Please Note: Availability Zones are only supported in several regions at this time.

For more information on managed disks, such as sizing options and pricing, please check out the [azure documentation](#).

`encryption_settings` supports:

- `enabled` - (Required) Is Encryption enabled on this Managed Disk? Changing this forces a new resource to be created.
- `disk_encryption_key` - (Optional) A `disk_encryption_key` block as defined below.
- `key_encryption_key` - (Optional) A `key_encryption_key` block as defined below.

`disk_encryption_key` supports:

- `secret_url` - (Required) The URL to the Key Vault Secret used as the Disk Encryption Key. This can be found as `id` on the `azurerm_key_vault_secret` resource.
- `source_vault_id` - (Required) The URL of the Key Vault. This can be found as `vault_uri` on the `azurerm_key_vault` resource.

`key_encryption_key` supports:

- **key_url** - (Required) The URL to the Key Vault Key used as the Key Encryption Key. This can be found as **id** on the **azurerm_key_vault_secret** resource.
- **source_vault_id** - (Required) The URL of the Key Vault. This can be found as **vault_uri** on the **azurerm_key_vault** resource.

» Attributes Reference

The following attributes are exported:

- **id** - The managed disk ID.

» Import

Managed Disks can be imported using the **resource id**, e.g.

```
terraform import azurerm_managed_disk.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm__snapshot

Manages a Disk Snapshot.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "snapshot-rg"
  location = "West Europe"
}

resource "azurerm_managed_disk" "test" {
  name                = "managed-disk"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  storage_account_type = "Standard_LRS"
  create_option        = "Empty"
  disk_size_gb         = "10"
}

resource "azurerm_snapshot" "test" {
  name                = "snapshot"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}
```

```

    create_option      = "Copy"
    source_uri         = "${azurerm_managed_disk.test.id}"
  }

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Snapshot resource. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the Snapshot. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **create_option** - (Required) Indicates how the snapshot is to be created. Possible values are **Copy** or **Import**. Changing this forces a new resource to be created.

Note: One of **source_uri**, **source_resource_id** or **storage_account_id** must be specified.

- **source_uri** - (Optional) Specifies the URI to a Managed or Unmanaged Disk. Changing this forces a new resource to be created.
- **source_resource_id** - (Optional) Specifies a reference to an existing snapshot, when **create_option** is **Copy**. Changing this forces a new resource to be created.
- **storage_account_id** - (Optional) Specifies the ID of an storage account. Used with **source_uri** to allow authorization during import of unmanaged blobs from a different subscription. Changing this forces a new resource to be created.
- **disk_size_gb** - (Optional) The size of the Snapshotted Disk in GB.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- **id** - The Snapshot ID.
- **disk_size_gb** - The Size of the Snapshotted Disk in GB.

» Import

Snapshots can be imported using the `resource id`, e.g.

```
terraform import azurerm_snapshot.test /subscriptions/00000000-0000-0000-0000-000000000000/1
```

» azurerm__shared__image

Manages a Shared Image within a Shared Image Gallery.

NOTE Shared Image Galleries are currently in Public Preview. You can find more information, including how to register for the Public Preview [here](#).

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "example-resources"
  location  = "West Europe"
}

resource "azurerm_shared_image_gallery" "test" {
  name                = "example_image_gallery"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location            = "${azurerm_resource_group.test.location}"
  description         = "Shared images and things."

  tags = {
    Hello = "There"
    World = "Example"
  }
}

resource "azurerm_shared_image" "test" {
  name                = "my-image"
  gallery_name        = "${azurerm_shared_image_gallery.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location            = "${azurerm_resource_group.test.location}"
  os_type             = "Linux"

  identifier {
    publisher = "PublisherName"
    offer     = "OfferName"
    sku       = "ExampleSku"
  }
}
```

}

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Shared Image. Changing this forces a new resource to be created.
 - **gallery_name** - (Required) Specifies the name of the Shared Image Gallery in which this Shared Image should exist. Changing this forces a new resource to be created.
 - **resource_group_name** - (Required) The name of the resource group in which the Shared Image Gallery exists. Changing this forces a new resource to be created.
 - **location** - (Required) Specifies the supported Azure location where the Shared Image Gallery exists. Changing this forces a new resource to be created.
 - **identifier** - (Required) An **identifier** block as defined below.
 - **os_type** - (Required) The type of Operating System present in this Shared Image. Possible values are **Linux** and **Windows**.
-
- **description** - (Optional) A description of this Shared Image.
 - **eula** - (Optional) The End User Licence Agreement for the Shared Image.
 - **privacy_statement_uri** - (Optional) The URI containing the Privacy Statement associated with this Shared Image.
 - **release_note_uri** - (Optional) The URI containing the Release Notes associated with this Shared Image.
 - **tags** - (Optional) A mapping of tags to assign to the Shared Image.
-

A **identifier** block supports the following:

- **offer** - (Required) The Offer Name for this Shared Image.
- **publisher** - (Required) The Publisher Name for this Gallery Image.
- **sku** - (Required) The Name of the SKU for this Gallery Image.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the Shared Image.

» Import

Shared Images can be imported using the `resource id`, e.g.

```
terraform import azurerm_shared_image.image1 /subscriptions/00000000-0000-0000-0000-00000000
```

» `azurerm_shared_image_gallery`

Manages a Shared Image Gallery.

NOTE Shared Image Galleries are currently in Public Preview. You can find more information, including how to register for the Public Preview [here](#).

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "example-resources"
  location  = "West Europe"
}

resource "azurerm_shared_image_gallery" "test" {
  name                        = "example_image_gallery"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  location                   = "${azurerm_resource_group.test.location}"
  description                 = "Shared images and things."

  tags = {
    Hello = "There"
    World = "Example"
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Shared Image Gallery. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the Shared Image Gallery. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **description** - (Optional) A description for this Shared Image Gallery.
- **tags** - (Optional) A mapping of tags to assign to the Shared Image Gallery.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Shared Image Gallery.
- **unique_name** - The Unique Name for this Shared Image Gallery.

» Import

Shared Image Galleries can be imported using the `resource id`, e.g.

```
terraform import azurerm_shared_image_gallery.gallery1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm__shared__image__version

Manages a Version of a Shared Image within a Shared Image Gallery.

NOTE Shared Image Galleries are currently in Public Preview. You can find more information, including how to register for the Public Preview [here](#).

» Example Usage

```
data "azurerm_image" "existing" {
  name                = "search-api"
  resource_group_name = "packerimages"
}

data "azurerm_shared_image" "existing" {
  name            = "existing-image"
  gallery_name    = "existing_gallery"
}
```

```

    resource_group_name = "existing-resources"
  }

  resource "azurerm_shared_image_version" "test" {
    name                = "0.0.1"
    gallery_name        = "${data.azurerm_shared_image.existing.gallery_name}"
    image_name          = "${data.azurerm_shared_image.existing.name}"
    resource_group_name = "${data.azurerm_shared_image.existing.resource_group_name}"
    location            = "${data.azurerm_shared_image.existing.location}"
    managed_image_id    = "${data.azurerm_image.existing.id}"

    target_region {
      name                = "${data.azurerm_shared_image.existing.location}"
      regional_replica_count = "5"
    }
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The version number for this Image Version, such as 1.0.0. Changing this forces a new resource to be created.
- **gallery_name** - (Required) The name of the Shared Image Gallery in which the Shared Image exists. Changing this forces a new resource to be created.
- **image_name** - (Required) The name of the Shared Image within the Shared Image Gallery in which this Version should be created. Changing this forces a new resource to be created.
- **location** - (Required) The Azure Region in which the Shared Image Gallery exists. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the Resource Group in which the Shared Image Gallery exists. Changing this forces a new resource to be created.
- **managed_image_id** - (Required) The ID of the Managed Image which should be used for this Shared Image Version. Changing this forces a new resource to be created.

NOTE: The ID can be sourced from the `azurerm_image` Data Source or Resource.

- **target_region** - (Required) One or more `target_region` blocks as documented below.

- **exclude_from_latest** - (Optional) Should this Image Version be excluded from the **latest** filter? If set to **true** this Image Version won't be returned for the **latest** version. Defaults to **false**.
- **tags** - (Optional) A collection of tags which should be applied to this resource.

The **target_region** block exports the following:

- **name** - (Required) The Azure Region in which this Image Version should exist.
- **regional_replica_count** - (Required) The number of replicas of the Image Version to be created per region.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Shared Image Version.

» Import

Shared Image Versions can be imported using the **resource id**, e.g.

```
terraform import azurerm_shared_image_version.version1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_virtual_machine

Manages a Virtual Machine.

NOTE: Data Disks can be attached either directly on the **azurerm_virtual_machine** resource, or using the **azurerm_virtual_machine_data_disk_attachment** resource - but the two cannot be used together. If both are used against the same Virtual Machine, spurious changes will occur.

» Example Usage (from an Azure Platform Image)

This example provisions a Virtual Machine with Managed Disks. Other examples of the **azurerm_virtual_machine** resource can be found in the **./examples/virtual-machines** directory within the Github Repository

```

variable "prefix" {
  default = "tfvmex"
}

resource "azurerm_resource_group" "main" {
  name     = "${var.prefix}-resources"
  location = "West US 2"
}

resource "azurerm_virtual_network" "main" {
  name                = "${var.prefix}-network"
  address_space       = ["10.0.0.0/16"]
  location             = "${azurerm_resource_group.main.location}"
  resource_group_name = "${azurerm_resource_group.main.name}"
}

resource "azurerm_subnet" "internal" {
  name                 = "internal"
  resource_group_name = "${azurerm_resource_group.main.name}"
  virtual_network_name = "${azurerm_virtual_network.main.name}"
  address_prefix       = "10.0.2.0/24"
}

resource "azurerm_network_interface" "main" {
  name                = "${var.prefix}-nic"
  location             = "${azurerm_resource_group.main.location}"
  resource_group_name = "${azurerm_resource_group.main.name}"

  ip_configuration {
    name                       = "testconfiguration1"
    subnet_id                  = "${azurerm_subnet.internal.id}"
    private_ip_address_allocation = "Dynamic"
  }
}

resource "azurerm_virtual_machine" "main" {
  name                = "${var.prefix}-vm"
  location             = "${azurerm_resource_group.main.location}"
  resource_group_name = "${azurerm_resource_group.main.name}"
  network_interface_ids = ["${azurerm_network_interface.main.id}"]
  vm_size              = "Standard_DS1_v2"

  # Uncomment this line to delete the OS disk automatically when deleting the VM
  # delete_os_disk_on_termination = true

```

```

# Uncomment this line to delete the data disks automatically when deleting the VM
# delete_data_disks_on_termination = true

storage_image_reference {
  publisher = "Canonical"
  offer     = "UbuntuServer"
  sku       = "16.04-LTS"
  version   = "latest"
}
storage_os_disk {
  name              = "myosdisk1"
  caching           = "ReadWrite"
  create_option     = "FromImage"
  managed_disk_type = "Standard_LRS"
}
os_profile {
  computer_name = "hostname"
  admin_username = "testadmin"
  admin_password = "Password1234!"
}
os_profile_linux_config {
  disable_password_authentication = false
}
tags = {
  environment = "staging"
}
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Virtual Machine. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) Specifies the name of the Resource Group in which the Virtual Machine should exist. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the Azure Region where the Virtual Machine exists. Changing this forces a new resource to be created.
- **network_interface_ids** - (Required) A list of Network Interface ID's which should be associated with the Virtual Machine.
- **os_profile_linux_config** - (Required, when a Linux machine) A `os_profile_linux_config` block.

- `os_profile_windows_config` - (Required, when a Windows machine) A `os_profile_windows_config` block.
- `vm_size` - (Required) Specifies the size of the Virtual Machine.

-
- `availability_set_id` - (Optional) The ID of the Availability Set in which the Virtual Machine should exist. Changing this forces a new resource to be created.
 - `boot_diagnostics` - (Optional) A `boot_diagnostics` block.
 - `delete_os_disk_on_termination` - (Optional) Should the OS Disk (either the Managed Disk / VHD Blob) be deleted when the Virtual Machine is destroyed? Defaults to `false`.
 - `delete_data_disks_on_termination` - (Optional) Should the Data Disks (either the Managed Disks / VHD Blobs) be deleted when the Virtual Machine is destroyed? Defaults to `false`.
 - `identity` - (Optional) A `identity` block.
 - `license_type` - (Optional) Specifies the BYOL Type for this Virtual Machine. This is only applicable to Windows Virtual Machines. Possible values are `Windows_Client` and `Windows_Server`.
 - `os_profile` - (Optional) An `os_profile` block. Required when `create_option` in the `storage_os_disk` block is set to `FromImage`.
 - `os_profile_secrets` - (Optional) One or more `os_profile_secrets` blocks.
 - `plan` - (Optional) A `plan` block.
 - `primary_network_interface_id` - (Optional) The ID of the Network Interface (which must be attached to the Virtual Machine) which should be the Primary Network Interface for this Virtual Machine.
 - `storage_data_disk` - (Optional) One or more `storage_data_disk` blocks.

Please Note: Data Disks can also be attached either using this block or the `azurerm_virtual_machine_data_disk_attachment` resource - but not both.

- `storage_image_reference` - (Optional) A `storage_image_reference` block.
- `storage_os_disk` - (Required) A `storage_os_disk` block.
- `tags` - (Optional) A mapping of tags to assign to the Virtual Machine.
- `zones` - (Optional) A list of a single item of the Availability Zone which the Virtual Machine should be allocated in.

Please Note: Availability Zones are only supported in several regions at this time.

For more information on the different example configurations, please check out the Azure documentation

A `additional_unattend_config` block supports the following:

- `pass` - (Required) Specifies the name of the pass that the content applies to. The only allowable value is `oobeSystem`.
- `component` - (Required) Specifies the name of the component to configure with the added content. The only allowable value is `Microsoft-Windows-Shell-Setup`.
- `setting_name` - (Required) Specifies the name of the setting to which the content applies. Possible values are: `FirstLogonCommands` and `AutoLogon`.
- `content` - (Optional) Specifies the base-64 encoded XML formatted content that is added to the `unattend.xml` file for the specified path and component.

A `boot_diagnostics` block supports the following:

- `enabled` - (Required) Should Boot Diagnostics be enabled for this Virtual Machine?
- `storage_uri` - (Required) The Storage Account's Blob Endpoint which should hold the virtual machine's diagnostic files.

NOTE: This needs to be the root of a Storage Account and not a Storage Container.

A `identity` block supports the following:

- `type` - (Required) The Managed Service Identity Type of this Virtual Machine. Possible values are `SystemAssigned` (where Azure will generate a Service Principal for you), `UserAssigned` (where you can specify the Service Principal ID's) to be used by this Virtual Machine using the `identity_ids` field, and `SystemAssigned, UserAssigned` which assigns both a system managed identity as well as the specified user assigned identities.

NOTE: Managed Service Identity previously required the installation of a VM Extension, but this information is now available via the Azure Instance Metadata Service.

NOTE: When `type` is set to `SystemAssigned`, identity the Principal ID can be retrieved after the virtual machine has been created. See documentation for more information.

- `identity_ids` - (Optional) Specifies a list of user managed identity ids to be assigned to the VM. Required if `type` is `UserAssigned`.

A `os_profile` block supports the following:

- `computer_name` - (Required) Specifies the name of the Virtual Machine.
- `admin_username` - (Required) Specifies the name of the local administrator account.
- `admin_password` - (Required for Windows, Optional for Linux) The password associated with the local administrator account.

NOTE: If using Linux, it may be preferable to use SSH Key authentication (available in the `os_profile_linux_config` block) instead of password authentication.

NOTE: `admin_password` must be between 6-72 characters long and must satisfy at least 3 of password complexity requirements from the following: 1. Contains an uppercase character 2. Contains a lowercase character 3. Contains a numeric digit 4. Contains a special character

- `custom_data` - (Optional) Specifies custom data to supply to the machine. On Linux-based systems, this can be used as a cloud-init script. On other systems, this will be copied as a file on disk. Internally, Terraform will base64 encode this value before sending it to the API. The maximum length of the binary array is 65535 bytes.

A `os_profile_linux_config` block supports the following:

- `disable_password_authentication` - (Required) Specifies whether password authentication should be disabled. If set to `false`, an `admin_password` must be specified.
- `ssh_keys` - (Optional) One or more `ssh_keys` blocks. This field is required if `disable_password_authentication` is set to `true`.

A `os_profile_secrets` block supports the following:

- `source_vault_id` - (Required) Specifies the ID of the Key Vault to use.
- `vault_certificates` - (Required) One or more `vault_certificates` blocks.

A `os_profile_windows_config` block supports the following:

- `provision_vm_agent` - (Optional) Should the Azure Virtual Machine Guest Agent be installed on this Virtual Machine? Defaults to `false`.

NOTE: This is different from the Default value used for this field within Azure.

- `enable_automatic_upgrades` - (Optional) Are automatic updates enabled on this Virtual Machine? Defaults to `false`.
- `timezone` - (Optional) Specifies the time zone of the virtual machine, the possible values are defined here.
- `winrm` - (Optional) One or more `winrm` block.
- `additional_unattend_config` - (Optional) A `additional_unattend_config` block.

A `plan` block supports the following:

- `name` - (Required) Specifies the name of the image from the marketplace.
- `publisher` - (Required) Specifies the publisher of the image.
- `product` - (Required) Specifies the product of the image from the marketplace.

A `ssh_keys` block supports the following:

- `key_data` - (Required) The Public SSH Key which should be written to the `path` defined above.

NOTE: Rather than defining this in-line you can source this from a local file using the `file` function - for example `key_data = file("~/ssh/id_rsa.pub")`.

- `path` - (Required) The path of the destination file on the virtual machine

NOTE: Due to a limitation in the Azure VM Agent the only allowed `path` is `/home/{username}/.ssh/authorized_keys`.

A `storage_image_reference` block supports the following:

This block provisions the Virtual Machine from one of two sources: an Azure Platform Image (e.g. Ubuntu/Windows Server) or a Custom Image.

To provision from an Azure Platform Image, the following fields are applicable:

- `publisher` - (Required) Specifies the publisher of the image used to create the virtual machine. Changing this forces a new resource to be created.

- **offer** - (Required) Specifies the offer of the image used to create the virtual machine. Changing this forces a new resource to be created.
- **sku** - (Required) Specifies the SKU of the image used to create the virtual machine. Changing this forces a new resource to be created.
- **version** - (Optional) Specifies the version of the image used to create the virtual machine. Changing this forces a new resource to be created.

To provision a Custom Image, the following fields are applicable:

- **id** - (Required) Specifies the ID of the Custom Image which the Virtual Machine should be created from. Changing this forces a new resource to be created.

NOTE: An example of how to use this is available within the `./examples/virtual-machines/managed-disks` directory within the Github Repository

A `storage_data_disk` block supports the following:

NOTE: Data Disks can also be attached either using this block or the `azurerm_virtual_machine_data_disk_attachment` resource - but not both.

- **name** - (Required) The name of the Data Disk.
- **caching** - (Optional) Specifies the caching requirements for the Data Disk. Possible values include `None`, `ReadOnly` and `ReadWrite`.
- **create_option** - (Required) Specifies how the data disk should be created. Possible values are `Attach`, `FromImage` and `Empty`.
- **disk_size_gb** - (Required) Specifies the size of the data disk in gigabytes.
- **lun** - (Required) Specifies the logical unit number of the data disk. This needs to be unique within all the Data Disks on the Virtual Machine.
- **write_accelerator_enabled** - (Optional) Specifies if Write Accelerator is enabled on the disk. This can only be enabled on `Premium_LRS` managed disks with no caching and M-Series VMs. Defaults to `false`.

The following properties apply when using Managed Disks:

- **managed_disk_type** - (Optional) Specifies the type of managed disk to create. Possible values are either `Standard_LRS`, `StandardSSD_LRS` or `Premium_LRS`.
- **managed_disk_id** - (Optional) Specifies the ID of an Existing Managed Disk which should be attached to this Virtual Machine. When this field is set **create_option** must be set to `Attach`.

The following properties apply when using Unmanaged Disks:

- **vhd_uri** - (Optional) Specifies the URI of the VHD file backing this Unmanaged Data Disk. Changing this forces a new resource to be created.

A **storage_os_disk** block supports the following:

- **name** - (Required) Specifies the name of the OS Disk.
- **create_option** - (Required) Specifies how the OS Disk should be created. Possible values are **Attach** (managed disks only) and **FromImage**.
- **caching** - (Optional) Specifies the caching requirements for the OS Disk. Possible values include **None**, **ReadOnly** and **ReadWrite**.
- **disk_size_gb** - (Optional) Specifies the size of the OS Disk in gigabytes.
- **image_uri** - (Optional) Specifies the Image URI in the format **publisherName:offer:skus:version**. This field can also specify the VHD uri of a custom VM image to clone. When cloning a Custom (Unmanaged) Disk Image the **os_type** field must be set.
- **os_type** - (Optional) Specifies the Operating System on the OS Disk. Possible values are **Linux** and **Windows**.
- **write_accelerator_enabled** - (Optional) Specifies if Write Accelerator is enabled on the disk. This can only be enabled on **Premium_LRS** managed disks with no caching and M-Series VMs. Defaults to **false**.

The following properties apply when using Managed Disks:

- **managed_disk_id** - (Optional) Specifies the ID of an existing Managed Disk which should be attached as the OS Disk of this Virtual Machine. If this is set then the **create_option** must be set to **Attach**.
- **managed_disk_type** - (Optional) Specifies the type of Managed Disk which should be created. Possible values are **Standard_LRS**, **StandardSSD_LRS** or **Premium_LRS**.

The following properties apply when using Unmanaged Disks:

- **vhd_uri** - (Optional) Specifies the URI of the VHD file backing this Unmanaged OS Disk. Changing this forces a new resource to be created.

A **vault_certificates** block supports the following:

- **certificate_url** - (Required) The ID of the Key Vault Secret. Stored secret is the Base64 encoding of a JSON Object that which is encoded in UTF-8 of which the contents need to be:

```
{
  "data": "<Base64-encoded-certificate>",
  "dataType": "pfx",
}
```

```
    "password": "<pfx-file-password>"
  }
```

NOTE: If your certificate is stored in Azure Key Vault - this can be sourced from the `secret_id` property on the `azurerm_key_vault_certificate` resource.

- `certificate_store` - (Required, on windows machines) Specifies the certificate store on the Virtual Machine where the certificate should be added to, such as My.

A `winrm` block supports the following:

- `protocol` - (Required) Specifies the protocol of listener. Possible values are HTTP or HTTPS.
- `certificate_url` - (Optional) The ID of the Key Vault Secret which contains the encrypted Certificate which should be installed on the Virtual Machine. This certificate must also be specified in the `vault_certificates` block within the `os_profile_secrets` block.

NOTE: This can be sourced from the `secret_id` field on the `azurerm_key_vault_certificate` resource.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the Virtual Machine.

» Import

Virtual Machines can be imported using the `resource id`, e.g.

```
terraform import azurerm_virtual_machine.test /subscriptions/00000000-0000-0000-0000-00000000
```

» `azurerm_virtual_machine_data_disk_attachment`

Manages attaching a Disk to a Virtual Machine.

NOTE: Data Disks can be attached either directly on the `azurerm_virtual_machine` resource, or using the `azurerm_virtual_machine_data_disk_attachment` resource - but the two cannot be used together. If both are used against the same Virtual Machine, spurious changes will occur.

Please Note: only Managed Disks are supported via this separate resource, Unmanaged Disks can be attached using the `storage_data_disk` block in the `azurerm_virtual_machine` resource.

» Example Usage

```
variable "prefix" {
  default = "example"
}

locals {
  vm_name = "${var.prefix}-vm"
}

resource "azurerm_resource_group" "main" {
  name       = "${var.prefix}-resources"
  location   = "West Europe"
}

resource "azurerm_virtual_network" "main" {
  name                = "${var.prefix}-network"
  address_space       = ["10.0.0.0/16"]
  location             = "${azurerm_resource_group.main.location}"
  resource_group_name = "${azurerm_resource_group.main.name}"
}

resource "azurerm_subnet" "internal" {
  name                 = "internal"
  resource_group_name  = "${azurerm_resource_group.main.name}"
  virtual_network_name = "${azurerm_virtual_network.main.name}"
  address_prefix       = "10.0.2.0/24"
}

resource "azurerm_network_interface" "main" {
  name                 = "${var.prefix}-nic"
  location             = "${azurerm_resource_group.main.location}"
  resource_group_name  = "${azurerm_resource_group.main.name}"

  ip_configuration {
    name                       = "internal"
    subnet_id                 = "${azurerm_subnet.internal.id}"
    private_ip_address_allocation = "Dynamic"
  }
}
```

```

resource "azurerm_virtual_machine" "test" {
  name                = "${local.vm_name}"
  location            = "${azurerm_resource_group.main.location}"
  resource_group_name = "${azurerm_resource_group.main.name}"
  network_interface_ids = ["${azurerm_network_interface.main.id}"]
  vm_size             = "Standard_F2"

  storage_image_reference {
    publisher = "Canonical"
    offer     = "UbuntuServer"
    sku       = "16.04-LTS"
    version   = "latest"
  }

  storage_os_disk {
    name          = "myosdisk1"
    caching       = "ReadWrite"
    create_option = "FromImage"
    managed_disk_type = "Standard_LRS"
  }

  os_profile {
    computer_name = "${local.vm_name}"
    admin_username = "testadmin"
    admin_password = "Password1234!"
  }

  os_profile_linux_config {
    disable_password_authentication = false
  }
}

resource "azurerm_managed_disk" "test" {
  name                = "${local.vm_name}-disk1"
  location            = "${azurerm_resource_group.main.location}"
  resource_group_name = "${azurerm_resource_group.main.name}"
  storage_account_type = "Standard_LRS"
  create_option       = "Empty"
  disk_size_gb        = 10
}

resource "azurerm_virtual_machine_data_disk_attachment" "test" {
  managed_disk_id      = "${azurerm_managed_disk.test.id}"
  virtual_machine_id   = "${azurerm_virtual_machine.test.id}"
  lun                  = "10"
  caching               = "ReadWrite"
}

```

}

» Argument Reference

The following arguments are supported:

- **virtual_machine_id** - (Required) The ID of the Virtual Machine to which the Data Disk should be attached. Changing this forces a new resource to be created.
- **managed_disk_id** - (Required) The ID of an existing Managed Disk which should be attached. Changing this forces a new resource to be created.
- **lun** - (Required) The Logical Unit Number of the Data Disk, which needs to be unique within the Virtual Machine. Changing this forces a new resource to be created.
- **caching** - (Required) Specifies the caching requirements for this Data Disk. Possible values include **None**, **ReadOnly** and **ReadWrite**.
- **create_option** - (Optional) The Create Option of the Data Disk, such as **Empty** or **Attach**. Defaults to **Attach**. Changing this forces a new resource to be created.
- **write_accelerator_enabled** - (Optional) Specifies if Write Accelerator is enabled on the disk. This can only be enabled on **Premium_LRS** managed disks with no caching and M-Series VMs. Defaults to **false**.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Virtual Machine Data Disk attachment.

» Import

Virtual Machines Data Disk Attachments can be imported using the **resource id**, e.g.

```
terraform import azurerm_virtual_machine_data_disk_attachment.test /subscriptions/00000000-0
```

Please Note: This is a Terraform Unique ID matching the format: `{virtualMachineID}/dataDisks/{diskName}`

» azurerm_virtual_machine_extension

Manages a Virtual Machine Extension to provide post deployment configuration and run automated tasks.

NOTE: Custom Script Extensions for Linux & Windows require that the `commandToExecute` returns a 0 exit code to be classified as successfully deployed. You can achieve this by appending `exit 0` to the end of your `commandToExecute`.

NOTE: Custom Script Extensions require that the Azure Virtual Machine Guest Agent is running on the Virtual Machine.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "acctestRG"
  location = "West US"
}

resource "azurerm_virtual_network" "test" {
  name            = "acctvn"
  address_space   = ["10.0.0.0/16"]
  location        = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_subnet" "test" {
  name                 = "acctsub"
  resource_group_name = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix       = "10.0.2.0/24"
}

resource "azurerm_network_interface" "test" {
  name            = "acctni"
  location        = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  ip_configuration {
    name                          = "testconfiguration1"
    subnet_id                    = "${azurerm_subnet.test.id}"
    private_ip_address_allocation = "Dynamic"
  }
}
```



```

resource "azurerm_storage_account" "test" {
  name                = "accsa"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location            = "${azurerm_resource_group.test.location}"
  account_tier        = "Standard"
  account_replication_type = "LRS"

  tags = {
    environment = "staging"
  }
}

resource "azurerm_storage_container" "test" {
  name                = "vhds"
  resource_group_name = "${azurerm_resource_group.test.name}"
  storage_account_name = "${azurerm_storage_account.test.name}"
  container_access_type = "private"
}

resource "azurerm_virtual_machine" "test" {
  name                = "acctvm"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  network_interface_ids = ["${azurerm_network_interface.test.id}"]
  vm_size             = "Standard_F2"

  storage_image_reference {
    publisher = "Canonical"
    offer     = "UbuntuServer"
    sku       = "16.04-LTS"
    version   = "latest"
  }

  storage_os_disk {
    name                = "myosdisk1"
    vhd_uri             = "${azurerm_storage_account.test.primary_blob_endpoint}${azurerm_storage_
    caching             = "ReadWrite"
    create_option       = "FromImage"
  }

  os_profile {
    computer_name = "hostname"
    admin_username = "testadmin"
    admin_password = "Password1234!"
  }
}

```

```

os_profile_linux_config {
  disable_password_authentication = false
}

tags = {
  environment = "staging"
}
}

resource "azurerm_virtual_machine_extension" "test" {
  name                        = "hostname"
  location                  = "${azurerm_resource_group.test.location}"
  resource_group_name       = "${azurerm_resource_group.test.name}"
  virtual_machine_name      = "${azurerm_virtual_machine.test.name}"
  publisher                 = "Microsoft.Azure.Extensions"
  type                     = "CustomScript"
  type_handler_version      = "2.0"

  settings = <<SETTINGS
    {
      "commandToExecute": "hostname && uptime"
    }
  SETTINGS

  tags = {
    environment = "Production"
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the virtual machine extension peering. Changing this forces a new resource to be created.
- **location** - (Required) The location where the extension is created. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the virtual network. Changing this forces a new resource to be created.
- **virtual_machine_name** - (Required) The name of the virtual machine. Changing this forces a new resource to be created.

- **publisher** - (Required) The publisher of the extension, available publishers can be found by using the Azure CLI.
- **type** - (Required) The type of extension, available types for a publisher can be found using the Azure CLI.

Note: The **Publisher** and **Type** of Virtual Machine Extensions can be found using the Azure CLI, via: `shell $ az vm extension image list --location westus -o table`

- **type_handler_version** - (Required) Specifies the version of the extension to use, available versions can be found using the Azure CLI.
- **auto_upgrade_minor_version** - (Optional) Specifies if the platform deploys the latest minor version update to the **type_handler_version** specified.
- **settings** - (Required) The settings passed to the extension, these are specified as a JSON object in a string.

Please Note: Certain VM Extensions require that the keys in the **settings** block are case sensitive. If you're seeing unhelpful errors, please ensure the keys are consistent with how Azure is expecting them (for instance, for the **JsonAddDomainExtension** extension, the keys are expected to be in **TitleCase**.)

- **protected_settings** - (Optional) The **protected_settings** passed to the extension, like **settings**, these are specified as a JSON object in a string.

Please Note: Certain VM Extensions require that the keys in the **protected_settings** block are case sensitive. If you're seeing unhelpful errors, please ensure the keys are consistent with how Azure is expecting them (for instance, for the **JsonAddDomainExtension** extension, the keys are expected to be in **TitleCase**.)

- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- **id** - The Virtual Machine Extension ID.

» Import

Virtual Machine Extensions can be imported using the **resource id**, e.g.

```
terraform import azurerm_virtual_machine_extension.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_virtual_machine_scale_set

Manage a virtual machine scale set.

Note: All arguments including the administrator login and password will be stored in the raw state as plain-text. Read more about sensitive data in state.

» Example Usage with Managed Disks (Recommended)

```
resource "azurerm_resource_group" "test" {
  name     = "acctestRG"
  location = "West US 2"
}

resource "azurerm_virtual_network" "test" {
  name            = "acctvn"
  address_space   = ["10.0.0.0/16"]
  location        = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_subnet" "test" {
  name                 = "acctsub"
  resource_group_name = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix       = "10.0.2.0/24"
}

resource "azurerm_public_ip" "test" {
  name             = "test"
  location         = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  allocation_method = "Static"
  domain_name_label = "${azurerm_resource_group.test.name}"

  tags = {
    environment = "staging"
  }
}

resource "azurerm_lb" "test" {
  name             = "test"
  location         = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}
```

```

    frontend_ip_configuration {
      name                = "PublicIPAddress"
      public_ip_address_id = "${azurerm_public_ip.test.id}"
    }
  }

resource "azurerm_lb_backend_address_pool" "bpepool" {
  resource_group_name = "${azurerm_resource_group.test.name}"
  loadbalancer_id      = "${azurerm_lb.test.id}"
  name                 = "BackEndAddressPool"
}

resource "azurerm_lb_nat_pool" "lbnatpool" {
  resource_group_name      = "${azurerm_resource_group.test.name}"
  name                     = "ssh"
  loadbalancer_id          = "${azurerm_lb.test.id}"
  protocol                  = "Tcp"
  frontend_port_start      = 50000
  frontend_port_end        = 50119
  backend_port              = 22
  frontend_ip_configuration_name = "PublicIPAddress"
}

resource "azurerm_lb_probe" "test" {
  resource_group_name = "${azurerm_resource_group.test.name}"
  loadbalancer_id      = "${azurerm_lb.test.id}"
  name                 = "http-probe"
  protocol              = "Http"
  request_path          = "/health"
  port                  = 8080
}

resource "azurerm_virtual_machine_scale_set" "test" {
  name                = "mytestscaleset-1"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  # automatic rolling upgrade
  automatic_os_upgrade = true
  upgrade_policy_mode   = "Rolling"

  rolling_upgrade_policy {
    max_batch_instance_percent      = 20
    max_unhealthy_instance_percent = 20
    max_unhealthy_upgraded_instance_percent = 5
    pause_time_between_batches      = "PT0S"
  }
}

```

```

}

# required when using rolling upgrade policy
health_probe_id = "${azurerm_lb_probe.test.id}"

sku {
  name      = "Standard_F2"
  tier       = "Standard"
  capacity  = 2
}

storage_profile_image_reference {
  publisher = "Canonical"
  offer     = "UbuntuServer"
  sku       = "16.04-LTS"
  version   = "latest"
}

storage_profile_os_disk {
  name              = ""
  caching           = "ReadWrite"
  create_option     = "FromImage"
  managed_disk_type = "Standard_LRS"
}

storage_profile_data_disk {
  lun          = 0
  caching      = "ReadWrite"
  create_option = "Empty"
  disk_size_gb = 10
}

os_profile {
  computer_name_prefix = "testvm"
  admin_username       = "myadmin"
}

os_profile_linux_config {
  disable_password_authentication = true

  ssh_keys {
    path      = "/home/myadmin/.ssh/authorized_keys"
    key_data  = "${file("~/ssh/demo_key.pub")}"
  }
}

```

```

network_profile {
  name      = "terraformnetworkprofile"
  primary   = true

  ip_configuration {
    name                        = "TestIPConfiguration"
    primary                    = true
    subnet_id                  = "${azurerm_subnet.test.id}"
    load_balancer_backend_address_pool_ids = ["${azurerm_lb_backend_address_pool.bpepool.id}"]
    load_balancer_inbound_nat_rules_ids   = ["${element(azurerm_lb_nat_pool.lbnatpool.*.id, 0)}"]
  }
}

tags = {
  environment = "staging"
}
}

```

» Example Usage with Unmanaged Disks

```

resource "azurerm_resource_group" "test" {
  name     = "acctestRG"
  location = "West US"
}

resource "azurerm_virtual_network" "test" {
  name                = "acctvn"
  address_space       = ["10.0.0.0/16"]
  location             = "West US"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_subnet" "test" {
  name                 = "acctsub"
  resource_group_name = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix        = "10.0.2.0/24"
}

resource "azurerm_storage_account" "test" {
  name                = "accsa"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location             = "westus"
  account_tier         = "Standard"
  account_replication_type = "LRS"
}

```

```

tags = {
    environment = "staging"
}
}

resource "azurerm_storage_container" "test" {
    name                = "vhds"
    resource_group_name = "${azurerm_resource_group.test.name}"
    storage_account_name = "${azurerm_storage_account.test.name}"
    container_access_type = "private"
}

resource "azurerm_virtual_machine_scale_set" "test" {
    name                = "mytestsscaleset-1"
    location            = "West US"
    resource_group_name = "${azurerm_resource_group.test.name}"
    upgrade_policy_mode = "Manual"

    sku {
        name      = "Standard_F2"
        tier       = "Standard"
        capacity = 2
    }

    os_profile {
        computer_name_prefix = "testvm"
        admin_username       = "myadmin"
    }

    os_profile_linux_config {
        disable_password_authentication = true

        ssh_keys {
            path      = "/home/myadmin/.ssh/authorized_keys"
            key_data = "${file("~/ssh/demo_key.pub")}"
        }
    }

    network_profile {
        name = "TestNetworkProfile"
        primary = true

        ip_configuration {
            name      = "TestIPConfiguration"
            primary   = true
        }
    }
}

```



```

        subnet_id = "${azurerm_subnet.test.id}"
    }
}

storage_profile_os_disk {
    name          = "osDiskProfile"
    caching       = "ReadWrite"
    create_option = "FromImage"
    vhd_containers = ["${azurerm_storage_account.test.primary_blob_endpoint}${azurerm_storage_account.test.blob_name}"]
}

storage_profile_image_reference {
    publisher = "Canonical"
    offer     = "UbuntuServer"
    sku       = "16.04-LTS"
    version   = "latest"
}
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the virtual machine scale set resource. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the virtual machine scale set. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **network_profile** - (Required) A collection of network profile block as documented below.
- **os_profile** - (Required) A Virtual Machine OS Profile block as documented below.
- **os_profile_windows_config** - (Required, when a windows machine) A Windows config block as documented below.
- **os_profile_linux_config** - (Required, when a linux machine) A Linux config block as documented below.
- **sku** - (Required) A sku block as documented below.
- **storage_profile_os_disk** - (Required) A storage profile os disk block as documented below

- **upgrade_policy_mode** - (Required) Specifies the mode of an upgrade to virtual machines in the scale set. Possible values, **Rolling**, **Manual**, or **Automatic**. When choosing **Rolling**, you will need to set a health probe.

-
- **automatic_os_upgrade** - (Optional) Automatic OS patches can be applied by Azure to your scaleset. This is particularly useful when **upgrade_policy_mode** is set to **Rolling**. Defaults to **false**.
 - **boot_diagnostics** - (Optional) A boot diagnostics profile block as referenced below.
 - **extension** - (Optional) Can be specified multiple times to add extension profiles to the scale set. Each **extension** block supports the fields documented below.
 - **eviction_policy** - (Optional) Specifies the eviction policy for Virtual Machines in this Scale Set. Possible values are **Deallocate** and **Delete**.

NOTE: **eviction_policy** can only be set when **priority** is set to **Low**.

- **health_probe_id** - (Optional) Specifies the identifier for the load balancer health probe. Required when using **Rolling** as your **upgrade_policy_mode**.
- **license_type** - (Optional, when a Windows machine) Specifies the Windows OS license type. If supplied, the only allowed values are **Windows_Client** and **Windows_Server**.
- **os_profile_secrets** - (Optional) A collection of Secret blocks as documented below.
- **overprovision** - (Optional) Specifies whether the virtual machine scale set should be overprovisioned.
- **plan** - (Optional) A plan block as documented below.
- **priority** - (Optional) Specifies the priority for the Virtual Machines in the Scale Set. Defaults to **Regular**. Possible values are **Low** and **Regular**.
- **rolling_upgrade_policy** - (Optional) A **rolling_upgrade_policy** block as defined below. This is only applicable when the **upgrade_policy_mode** is **Rolling**.
- **single_placement_group** - (Optional) Specifies whether the scale set is limited to a single placement group with a maximum size of 100 virtual machines. If set to **false**, managed disks must be used. Default is **true**. Changing this forces a new resource to be created. See documentation for more information.
- **storage_profile_data_disk** - (Optional) A storage profile data disk block as documented below

- **storage_profile_image_reference** - (Optional) A storage profile image reference block as documented below.
- **tags** - (Optional) A mapping of tags to assign to the resource.
- **zones** - (Optional) A collection of availability zones to spread the Virtual Machines over.

Please Note: Availability Zones are only supported in several regions at this time.

sku supports the following:

- **name** - (Required) Specifies the size of virtual machines in a scale set.
- **tier** - (Optional) Specifies the tier of virtual machines in a scale set. Possible values, **standard** or **basic**.
- **capacity** - (Required) Specifies the number of virtual machines in the scale set.

rolling_upgrade_policy supports the following:

- **max_batch_instance_percent** - (Optional) The maximum percent of total virtual machine instances that will be upgraded simultaneously by the rolling upgrade in one batch. As this is a maximum, unhealthy instances in previous or future batches can cause the percentage of instances in a batch to decrease to ensure higher reliability. Defaults to 20.
- **max_unhealthy_instance_percent** - (Optional) The maximum percentage of the total virtual machine instances in the scale set that can be simultaneously unhealthy, either as a result of being upgraded, or by being found in an unhealthy state by the virtual machine health checks before the rolling upgrade aborts. This constraint will be checked prior to starting any batch. Defaults to 20.
- **max_unhealthy_upgraded_instance_percent** - (Optional) The maximum percentage of upgraded virtual machine instances that can be found to be in an unhealthy state. This check will happen after each batch is upgraded. If this percentage is ever exceeded, the rolling update aborts. Defaults to 20.
- **pause_time_between_batches** - (Optional) The wait time between completing the update for all virtual machines in one batch and starting the next batch. The time duration should be specified in ISO 8601 format for duration (https://en.wikipedia.org/wiki/ISO_8601#Durations). Defaults to 0 seconds represented as PT0S.

identity supports the following:

- **type** - (Required) Specifies the identity type to be assigned to the scale set. Allowable values are **SystemAssigned**, **UserAssigned**, and **SystemAssigned, UserAssigned**. For the **SystemAssigned** identity the

scale set's Service Principal ID (SPN) can be retrieved after the scale set has been created. See documentation for more information.

- **identity_ids** - (Optional) Specifies a list of user managed identity ids to be assigned to the VMSS. Required if **type** is **UserAssigned**.

```
resource "azurerm_virtual_machine_scale_set" "test" {
  name                = "vm-scaleset"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location            = "${azurerm_resource_group.test.location}"

  sku {
    name       = "${var.vm_sku}"
    tier        = "Standard"
    capacity   = "${var.instance_count}"
  }

  identity {
    type = "systemAssigned"
  }

  extension {
    name                = "MSILinuxExtension"
    publisher           = "Microsoft.ManagedIdentity"
    type                = "ManagedIdentityExtensionForLinux"
    type_handler_version = "1.0"
    settings            = "{\"port\": 50342}"
  }

  # ...
}

output "principal_id" {
  value = "${lookup(azurerm_virtual_machine_scale_set.test.identity[0], "principal_id")}"
}
```

os_profile supports the following:

- **computer_name_prefix** - (Required) Specifies the computer name prefix for all of the virtual machines in the scale set. Computer name prefixes must be 1 to 9 characters long for windows images and 1 - 58 for linux. Changing this forces a new resource to be created.
- **admin_username** - (Required) Specifies the administrator account name to use for all the instances of virtual machines in the scale set.
- **admin_password** - (Required) Specifies the administrator password to use for all the instances of virtual machines in a scale set.
- **custom_data** - (Optional) Specifies custom data to supply to the machine. On linux-based systems, this can be used as a cloud-init script. On other

systems, this will be copied as a file on disk. Internally, Terraform will base64 encode this value before sending it to the API. The maximum length of the binary array is 65535 bytes.

`os_profile_secrets` supports the following:

- `source_vault_id` - (Required) Specifies the key vault to use.
- `vault_certificates` - (Required, on windows machines) A collection of Vault Certificates as documented below

`vault_certificates` support the following:

- `certificate_url` - (Required) It is the Base64 encoding of a JSON Object that which is encoded in UTF-8 of which the contents need to be `data`, `dataType` and `password`.
- `certificate_store` - (Required, on windows machines) Specifies the certificate store on the Virtual Machine where the certificate should be added to.

`os_profile_windows_config` supports the following:

- `provision_vm_agent` - (Optional) Indicates whether virtual machine agent should be provisioned on the virtual machines in the scale set.
- `enable_automatic_upgrades` - (Optional) Indicates whether virtual machines in the scale set are enabled for automatic updates.
- `winrm` - (Optional) A collection of WinRM configuration blocks as documented below.
- `additional_unattend_config` - (Optional) An Additional Unattended Config block as documented below.

`winrm` supports the following:

- `protocol` - (Required) Specifies the protocol of listener
- `certificate_url` - (Optional) Specifies URL of the certificate with which new Virtual Machines is provisioned.

`additional_unattend_config` supports the following:

- `pass` - (Required) Specifies the name of the pass that the content applies to. The only allowable value is `oobeSystem`.
- `component` - (Required) Specifies the name of the component to configure with the added content. The only allowable value is `Microsoft-Windows-Shell-Setup`.
- `setting_name` - (Required) Specifies the name of the setting to which the content applies. Possible values are: `FirstLogonCommands` and `AutoLogon`.
- `content` - (Optional) Specifies the base-64 encoded XML formatted content that is added to the `unattend.xml` file for the specified path and component.

`os_profile_linux_config` supports the following:

- **disable_password_authentication** - (Optional) Specifies whether password authentication should be disabled. Defaults to **false**. Changing this forces a new resource to be created.
- **ssh_keys** - (Optional) Specifies a collection of **path** and **key_data** to be placed on the virtual machine.

***Note:** Please note that the only allowed **path** is `/home/<username>/.ssh/authorized_keys` due to a limitation of Azure*

network_profile supports the following:

- **name** - (Required) Specifies the name of the network interface configuration.
- **primary** - (Required) Indicates whether network interfaces created from the network interface configuration will be the primary NIC of the VM.
- **ip_configuration** - (Required) An **ip_configuration** block as documented below.
- **accelerated_networking** - (Optional) Specifies whether to enable accelerated networking or not. Defaults to **false**.
- **dns_settings** - (Optional) A **dns_settings** block as documented below.
- **ip_forwarding** - (Optional) Whether IP forwarding is enabled on this NIC. Defaults to **false**.
- **network_security_group_id** - (Optional) Specifies the identifier for the network security group.

dns_settings supports the following:

- **dns_servers** - (Required) Specifies an array of dns servers.

ip_configuration supports the following:

- **name** - (Required) Specifies name of the IP configuration.
- **subnet_id** - (Required) Specifies the identifier of the subnet.
- **application_gateway_backend_address_pool_ids** - (Optional) Specifies an array of references to backend address pools of application gateways. A scale set can reference backend address pools of one application gateway. Multiple scale sets cannot use the same application gateway.
- **load_balancer_backend_address_pool_ids** - (Optional) Specifies an array of references to backend address pools of load balancers. A scale set can reference backend address pools of one public and one internal load balancer. Multiple scale sets cannot use the same load balancer.
- **load_balancer_inbound_nat_rules_ids** - (Optional) Specifies an array of references to inbound NAT rules for load balancers.
- **primary** - (Required) Specifies if this **ip_configuration** is the primary one.
- **application_security_group_ids** - (Optional) Specifies up to 20 application security group IDs.
- **public_ip_address_configuration** - (Optional) Describes a virtual machines scale set IP Configuration's PublicIPAddress configuration. The **public_ip_address_configuration** is documented below.

`public_ip_address_configuration` supports the following:

- `name` - (Required) The name of the public ip address configuration
- `idle_timeout` - (Required) The idle timeout in minutes. This value must be between 4 and 32.
- `domain_name_label` - (Required) The domain name label for the dns settings.

`storage_profile_os_disk` supports the following:

- `name` - (Optional) Specifies the disk name. Must be specified when using unmanaged disk ('`managed_disk_type`' property not set).
- `vhd_containers` - (Optional) Specifies the vhd uri. Cannot be used when `image` or `managed_disk_type` is specified.
- `managed_disk_type` - (Optional) Specifies the type of managed disk to create. Value you must be either `Standard_LRS`, `StandardSSD_LRS` or `Premium_LRS`. Cannot be used when `vhd_containers` or `image` is specified.
- `create_option` - (Required) Specifies how the virtual machine should be created. The only possible option is `FromImage`.
- `caching` - (Optional) Specifies the caching requirements. Possible values include: `None` (default), `ReadOnly`, `ReadWrite`.
- `image` - (Optional) Specifies the blob uri for user image. A virtual machine scale set creates an os disk in the same container as the user image. Updating the osDisk image causes the existing disk to be deleted and a new one created with the new image. If the VM scale set is in Manual upgrade mode then the virtual machines are not updated until they have manualUpgrade applied to them. When setting this field `os_type` needs to be specified. Cannot be used when `vhd_containers`, `managed_disk_type` or `storage_profile_image_reference` are specified.
- `os_type` - (Optional) Specifies the operating system Type, valid values are windows, linux.

`storage_profile_data_disk` supports the following:

- `lun` - (Required) Specifies the Logical Unit Number of the disk in each virtual machine in the scale set.
- `create_option` - (Optional) Specifies how the data disk should be created. The only possible options are `FromImage` and `Empty`.
- `caching` - (Optional) Specifies the caching requirements. Possible values include: `None` (default), `ReadOnly`, `ReadWrite`.
- `disk_size_gb` - (Optional) Specifies the size of the disk in GB. This element is required when creating an empty disk.
- `managed_disk_type` - (Optional) Specifies the type of managed disk to create. Value must be either `Standard_LRS`, `StandardSSD_LRS` or `Premium_LRS`.

`storage_profile_image_reference` supports the following:

- **id** - (Optional) Specifies the ID of the (custom) image to use to create the virtual machine scale set, as in the example below.
- **publisher** - (Optional) Specifies the publisher of the image used to create the virtual machines.
- **offer** - (Optional) Specifies the offer of the image used to create the virtual machines.
- **sku** - (Optional) Specifies the SKU of the image used to create the virtual machines.
- **version** - (Optional) Specifies the version of the image used to create the virtual machines.

boot_diagnostics supports the following:

- **enabled**: (Required) Whether to enable boot diagnostics for the virtual machine.
- **storage_uri**: (Required) Blob endpoint for the storage account to hold the virtual machine's diagnostic files. This must be the root of a storage account, and not a storage container.

extension supports the following:

- **name** - (Required) Specifies the name of the extension.
- **publisher** - (Required) The publisher of the extension, available publishers can be found by using the Azure CLI.
- **type** - (Required) The type of extension, available types for a publisher can be found using the Azure CLI.
- **type_handler_version** - (Required) Specifies the version of the extension to use, available versions can be found using the Azure CLI.
- **auto_upgrade_minor_version** - (Optional) Specifies whether or not to use the latest minor version available.
- **settings** - (Required) The settings passed to the extension, these are specified as a JSON object in a string.
- **protected_settings** - (Optional) The protected_settings passed to the extension, like settings, these are specified as a JSON object in a string.

plan supports the following:

- **name** - (Required) Specifies the name of the image from the marketplace.
- **publisher** - (Required) Specifies the publisher of the image.
- **product** - (Required) Specifies the product of the image from the marketplace.

» Example of `storage_profile_image_reference` with `id`

```
resource "azurerm_image" "test" {
  name = "test"

  # ...
```



```

}

resource "azurerm_virtual_machine_scale_set" "test" {
  name = "test"

  # ...

  storage_profile_image_reference {
    id = "${azurerm_image.test.id}"
  }

  # ...
}

```

» Attributes Reference

The following attributes are exported:

- `id` - The virtual machine scale set ID.

» Import

Virtual Machine Scale Sets can be imported using the `resource id`, e.g.

```
terraform import azurerm_virtual_machine_scale_set.scaleset1 /subscriptions/00000000-0000-00
```

» `azurerm__container__group`

Manage as an Azure Container Group instance.

» Example Usage

```

resource "azurerm_resource_group" "aci-rg" {
  name     = "aci-test"
  location = "west us"
}

resource "azurerm_storage_account" "aci-sa" {
  name                = "acistorageacct"
  resource_group_name = "${azurerm_resource_group.aci-rg.name}"
  location            = "${azurerm_resource_group.aci-rg.location}"
  account_tier        = "Standard"
}

```

```

    account_replication_type = "LRS"
}

resource "azurerm_storage_share" "aci-share" {
    name = "aci-test-share"

    resource_group_name = "${azurerm_resource_group.aci-rg.name}"
    storage_account_name = "${azurerm_storage_account.aci-sa.name}"

    quota = 50
}

resource "azurerm_container_group" "aci-helloworld" {
    name = "aci-hw"
    location = "${azurerm_resource_group.aci-rg.location}"
    resource_group_name = "${azurerm_resource_group.aci-rg.name}"
    ip_address_type = "public"
    dns_name_label = "aci-label"
    os_type = "Linux"

    container {
        name = "hw"
        image = "seanmckenna/aci-hellofiles"
        cpu = "0.5"
        memory = "1.5"
        ports = {
            port = 80
            protocol = "TCP"
        }
        ports {
            port = 443
            protocol = "TCP"
        }
    }

    environment_variables = {
        "NODE_ENV" = "testing"
    }

    secure_environment_variables = {
        "ACCESS_KEY" = "secure_testing"
    }

    readiness_probe {
        exec = ["/bin/sh", "-c", "touch /tmp/healthy; sleep 30; rm -rf /tmp/healthy; sleep 600"]
    }
}

```

```

liveness_probe {
    exec = ["cat", "/tmp/healthy"]
}

commands = ["/bin/bash", "-c", "'/path to/myscript.sh'"]

volume {
    name          = "logs"
    mount_path    = "/aci/logs"
    read_only     = false
    share_name    = "${azurerm_storage_share.aci-share.name}"

    storage_account_name = "${azurerm_storage_account.aci-sa.name}"
    storage_account_key  = "${azurerm_storage_account.aci-sa.primary_access_key}"
}

container {
    name     = "sidecar"
    image    = "microsoft/aci-tutorial-sidecar"
    cpu      = "0.5"
    memory   = "1.5"
}

tags = {
    environment = "testing"
}
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Container Group. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the Container Group. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **container** - (Required) The definition of a container that is part of the group as documented in the `container` block below. Changing this forces a new resource to be created.

Note: if `os_type` is set to `Windows` currently only a single `container` block is supported.

- `os_type` - (Required) The OS for the container group. Allowed values are `Linux` and `Windows`. Changing this forces a new resource to be created.
-
- `diagnostics` - (Optional) A `diagnostics` block as documented below.
 - `dns_name_label` - (Optional) The DNS label/name for the container groups IP.
 - `ip_address_type` - (Optional) Specifies the ip address type of the container. `Public` is the only acceptable value at this time. Changing this forces a new resource to be created.
 - `image_registry_credential` - (Optional) A `image_registry_credential` block as documented below.
 - `restart_policy` - (Optional) Restart policy for the container group. Allowed values are `Always`, `Never`, `OnFailure`. Defaults to `Always`.
 - `tags` - (Optional) A mapping of tags to assign to the resource.
-

A `container` block supports:

- `name` - (Required) Specifies the name of the Container. Changing this forces a new resource to be created.
- `image` - (Required) The container image name. Changing this forces a new resource to be created.
- `cpu` - (Required) The required number of CPU cores of the containers. Changing this forces a new resource to be created.
- `memory` - (Required) The required memory of the containers in GB. Changing this forces a new resource to be created.
- `gpu` - (Optional) A `gpu` block as defined below.

Note: Gpu resources are currently only supported in Linux containers.

- `ports` - (Optional) A set of public ports for the container. Changing this forces a new resource to be created. Set as documented in the `ports` block below.
- `environment_variables` - (Optional) A list of environment variables to be set on the container. Specified as a map of name/value pairs. Changing this forces a new resource to be created.

- **secure_environment_variables** - (Optional) A list of sensitive environment variables to be set on the container. Specified as a map of name/value pairs. Changing this forces a new resource to be created.
- **readiness_probe** - (Optional) The definition of a readiness probe for this container as documented in the **readiness_probe** block below. Changing this forces a new resource to be created.
- **liveness_probe** - (Optional) The definition of a readiness probe for this container as documented in the **liveness_probe** block below. Changing this forces a new resource to be created.
- **command** - (Optional) A command line to be run on the container.

NOTE: The field **command** has been deprecated in favor of **commands** to better match the API.

- **commands** - (Optional) A list of commands which should be run on the container.
- **volume** - (Optional) The definition of a volume mount for this container as documented in the **volume** block below. Changing this forces a new resource to be created.

A **diagnostics** block supports:

- **log_analytics** - (Required) A **log_analytics** block as defined below.

A **image_registry_credential** block supports:

- **username** - (Required) The username with which to connect to the registry.
- **password** - (Required) The password with which to connect to the registry.
- **server** - (Required) The address to use to connect to the registry without protocol ("https"/"http"). For example: "myacr.acr.io"

A **log_analytics** block supports:

- **log_type** - (Required) The log type which should be used. Possible values are **ContainerInsights** and **ContainerInstanceLogs**.
- **workspace_id** - (Required) The Workspace ID of the Log Analytics Workspace.
- **workspace_key** - (Required) The Workspace Key of the Log Analytics Workspace.
- **metadata** - (Optional) Any metadata required for Log Analytics.

A `ports` block supports:

- `port` - (Required) The port number the container will expose.
- `protocol` - (Required) The network protocol associated with port. Possible values are `TCP` & `UDP`.

--

A `gpu` block supports:

- `count` - (Required) The number of GPUs which should be assigned to this container. Allowed values are 1, 2, or 4.
- `sku` - (Required) The Sku which should be used for the GPU. Possible values are `K80`, `P100`, or `V100`.

A `volume` block supports:

- `name` - (Required) The name of the volume mount. Changing this forces a new resource to be created.
- `mount_path` - (Required) The path on which this volume is to be mounted. Changing this forces a new resource to be created.
- `read_only` - (Optional) Specify if the volume is to be mounted as read only or not. The default value is `false`. Changing this forces a new resource to be created.
- `storage_account_name` - (Required) The Azure storage account from which the volume is to be mounted. Changing this forces a new resource to be created.
- `storage_account_key` - (Required) The access key for the Azure Storage account specified as above. Changing this forces a new resource to be created.
- `share_name` - (Required) The Azure storage share that is to be mounted as a volume. This must be created on the storage account specified as above. Changing this forces a new resource to be created.

The `readiness_probe` block supports:

- `exec` - (Optional) Commands to be run to validate container readiness. Changing this forces a new resource to be created.
- `httpget` - (Optional) The definition of the `httpget` for this container as documented in the `httpget` block below. Changing this forces a new resource to be created.

- **initial_delay_seconds** - (Optional) Number of seconds after the container has started before liveness or readiness probes are initiated. Changing this forces a new resource to be created.
- **period_seconds** - (Optional) How often (in seconds) to perform the probe. The default value is 10 and the minimum value is 1. Changing this forces a new resource to be created.
- **failure_threshold** - (Optional) How many times to try the probe before restarting the container (liveness probe) or marking the container as unhealthy (readiness probe). The default value is 3 and the minimum value is 1. Changing this forces a new resource to be created.
- **success_threshold** - (Optional) Minimum consecutive successes for the probe to be considered successful after having failed. The default value is 1 and the minimum value is 1. Changing this forces a new resource to be created.
- **timeout_seconds** - (Optional) Number of seconds after which the probe times out. The default value is 1 and the minimum value is 1. Changing this forces a new resource to be created.

The **liveness_probe** block supports:

- **exec** - (Optional) Commands to be run to validate container readiness. Changing this forces a new resource to be created.
- **httpget** - (Optional) The definition of the httpget for this container as documented in the **httpget** block below. Changing this forces a new resource to be created.
- **initial_delay_seconds** - (Optional) Number of seconds after the container has started before liveness or readiness probes are initiated. Changing this forces a new resource to be created.
- **period_seconds** - (Optional) How often (in seconds) to perform the probe. The default value is 10 and the minimum value is 1. Changing this forces a new resource to be created.
- **failure_threshold** - (Optional) How many times to try the probe before restarting the container (liveness probe) or marking the container as unhealthy (readiness probe). The default value is 3 and the minimum value is 1. Changing this forces a new resource to be created.
- **success_threshold** - (Optional) Minimum consecutive successes for the probe to be considered successful after having failed. The default value is 1 and the minimum value is 1. Changing this forces a new resource to be created.

- `timeout_seconds` - (Optional) Number of seconds after which the probe times out. The default value is 1 and the minimum value is 1. Changing this forces a new resource to be created.

The `httpget` block supports:

- `path` - (Optional) Path to access on the HTTP server. Changing this forces a new resource to be created.
- `port` - (Optional) Number of the port to access on the container. Changing this forces a new resource to be created.
- `scheme` - (Optional) Scheme to use for connecting to the host. Possible values are `Http` and `Https`. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- `id` - The container group ID.
- `ip_address` - The IP address allocated to the container group.
- `fqdn` - The FQDN of the container group derived from `dns_name_label`.

» Import

Container Group's can be imported using the `resource id`, e.g.

```
terraform import azurerm_container_group.containerGroup1 /subscriptions/00000000-0000-0000-0000-000000000000/resourceGroups/rg1/providers/Microsoft.ContainerInstance/containerGroups/containerGroup1
```

» `azurerm__container__registry`

Manages an Azure Container Registry.

Note: All arguments including the access key will be stored in the raw state as plain-text. Read more about sensitive data in state.

» Example Usage

```
resource "azurerm_resource_group" "rg" {
  name     = "resourceGroup1"
  location = "West US"
}
```



```

}

resource "azurerm_container_registry" "acr" {
  name                        = "containerRegistry1"
  resource_group_name        = "${azurerm_resource_group.rg.name}"
  location                   = "${azurerm_resource_group.rg.location}"
  sku                        = "Premium"
  admin_enabled              = false
  georeplication_locations = ["East US", "West Europe"]
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Container Registry. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the Container Registry. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **admin_enabled** - (Optional) Specifies whether the admin user is enabled. Defaults to **false**.
- **storage_account_id** - (Required for **Classic** Sku - Forbidden otherwise) The ID of a Storage Account which must be located in the same Azure Region as the Container Registry.
- **sku** - (Optional) The SKU name of the the container registry. Possible values are **Classic** (which was previously **Basic**), **Basic**, **Standard** and **Premium**.

NOTE: The **Classic** SKU is Deprecated and will no longer be available for new resources from the end of March 2019.

- **tags** - (Optional) A mapping of tags to assign to the resource.
- **georeplication_locations** - (Optional) A list of Azure locations where the container registry should be geo-replicated.

» Attributes Reference

The following attributes are exported:

- **id** - The Container Registry ID.

- `login_server` - The URL that can be used to log into the container registry.
- `admin_username` - The Username associated with the Container Registry Admin account - if the admin account is enabled.
- `admin_password` - The Password associated with the Container Registry Admin account - if the admin account is enabled.

» Import

Container Registries can be imported using the `resource id`, e.g.

```
terraform import azurerm_container_registry.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_container_service`

Manages an Azure Container Service Instance

NOTE: All arguments including the client secret will be stored in the raw state as plain-text. Read more about sensitive data in state.

DEPRECATED: Azure Container Service (ACS) has been deprecated by Azure in favour of Azure (Managed) Kubernetes Service (AKS). Support for ACS will be removed in the next major version of the AzureRM Provider (2.0) - and we **strongly recommend** you consider using Azure Kubernetes Service (AKS) for new deployments.

» Example Usage (DCOS)

```
resource "azurerm_resource_group" "test" {
  name      = "acctestRG1"
  location  = "West US"
}

resource "azurerm_container_service" "test" {
  name                        = "acctestcontservice1"
  location                   = "${azurerm_resource_group.test.location}"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  orchestration_platform     = "DCOS"

  master_profile {
    count      = 1
    dns_prefix = "acctestmaster1"
  }
}
```

```

linux_profile {
  admin_username = "acctestuser1"

  ssh_key {
    key_data = "ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAQCaZoyiz1qbd0Q8xEf6uEu1cCwYowo5FHtsF"
  }
}

agent_pool_profile {
  name      = "default"
  count     = 1
  dns_prefix = "acctestagent1"
  vm_size   = "Standard_F2"
}

diagnostics_profile {
  enabled = false
}

tags = {
  Environment = "Production"
}
}

```

» Example Usage (Kubernetes)

```

resource "azurerm_resource_group" "test" {
  name     = "acctestRG1"
  location = "West US"
}

resource "azurerm_container_service" "test" {
  name                = "acctestcontservice1"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  orchestration_platform = "Kubernetes"

  master_profile {
    count     = 1
    dns_prefix = "acctestmaster1"
  }

  linux_profile {
    admin_username = "acctestuser1"
  }
}

```

```

    ssh_key {
      key_data = "ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAQCaZoyiz1qbd0Q8xEf6uEu1cCwYowo5FHtsF
    }
  }

  agent_pool_profile {
    name      = "default"
    count     = 1
    dns_prefix = "acctestagent1"
    vm_size   = "Standard_F2"
  }

  service_principal {
    client_id     = "000000000-0000-0000-0000-000000000000"
    client_secret = "000000000000000000000000000000000000"
  }

  diagnostics_profile {
    enabled = false
  }

  tags = {
    Environment = "Production"
  }
}

```

» Example Usage (Swarm)

```

resource "azurerm_resource_group" "test" {
  name     = "acctestRG1"
  location = "West US"
}

resource "azurerm_container_service" "test" {
  name                        = "acctestcontservice1"
  location                   = "${azurerm_resource_group.test.location}"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  orchestration_platform     = "Swarm"

  master_profile {
    count     = 1
    dns_prefix = "acctestmaster1"
  }
}

```

```

linux_profile {
  admin_username = "acctestuser1"

  ssh_key {
    key_data = "ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAQCaZoyiz1qbd0Q8xEf6uEu1cCwYowo5FHtsI"
  }
}

agent_pool_profile {
  name      = "default"
  count     = 1
  dns_prefix = "acctestagent1"
  vm_size   = "Standard_F2"
}

diagnostics_profile {
  enabled = false
}

tags = {
  Environment = "Production"
}
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Container Service instance to create. Changing this forces a new resource to be created.
- **location** - (Required) The location where the Container Service instance should be created. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) Specifies the resource group where the resource exists. Changing this forces a new resource to be created.
- **orchestration_platform** - (Required) Specifies the Container Orchestration Platform to use. Currently can be either `DCOS`, `Kubernetes` or `Swarm`. Changing this forces a new resource to be created.
- **master_profile** - (Required) A Master Profile block as documented below.
- **linux_profile** - (Required) A Linux Profile block as documented below.
- **agent_pool_profile** - (Required) A Agent Pool Profile's block as documented below.

- **service_principal** - (only Required when you're using **Kubernetes** as an Orchestration Platform) A Service Principal block as documented below.
- **diagnostics_profile** - (Required) A VM Diagnostics Profile block as documented below.
- **tags** - (Optional) A mapping of tags to assign to the resource.

master_profile supports the following:

- **count** - (Required) Number of masters (VMs) in the container service cluster. Allowed values are 1, 3, and 5. The default value is 1.
- **dns_prefix** - (Required) The DNS Prefix to use for the Container Service master nodes.

linux_profile supports the following:

- **admin_username** - (Required) The Admin Username for the Cluster.
- **ssh_key** - (Required) An SSH Key block as documented below.

ssh_key supports the following:

- **key_data** - (Required) The Public SSH Key used to access the cluster.

agent_pool_profile supports the following:

- **name** - (Required) Unique name of the agent pool profile in the context of the subscription and resource group.
- **count** - (Required) Number of agents (VMs) to host docker containers. Allowed values must be in the range of 1 to 100 (inclusive). The default value is 1.
- **dns_prefix** - (Required) The DNS Prefix given to Agents in this Agent Pool.
- **vm_size** - (Required) The VM Size of each of the Agent Pool VM's (e.g. Standard_F1 / Standard_D2v2).

service_principal supports the following:

- **client_id** - (Required) The ID for the Service Principal.
- **client_secret** - (Required) The secret password associated with the service principal.

diagnostics_profile supports the following:

- **enabled** - (Required) Should VM Diagnostics be enabled for the Container Service VM's

» Attributes Reference

The following attributes are exported:

- **id** - The Container Service ID.

- `master_profile.fqdn` - FQDN for the master.
- `agent_pool_profile.fqdn` - FQDN for the agent pool.
- `diagnostics_profile.storage_uri` - The URI of the storage account where diagnostics are stored.

» `azurerm_kubernetes_cluster`

Manages a Managed Kubernetes Cluster (also known as AKS / Azure Kubernetes Service)

Note: All arguments including the client secret will be stored in the raw state as plain-text. Read more about sensitive data in state.

» Example Usage

This example provisions a basic Managed Kubernetes Cluster. Other examples of the `azurerm_kubernetes_cluster` resource can be found in the `./examples/kubernetes` directory within the Github Repository

```
resource "azurerm_resource_group" "test" {
  name     = "acctestRG1"
  location = "East US"
}

resource "azurerm_kubernetes_cluster" "test" {
  name                = "acctestaks1"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  dns_prefix          = "acctestagent1"

  agent_pool_profile {
    name          = "default"
    count         = 1
    vm_size       = "Standard_D1_v2"
    os_type       = "Linux"
    os_disk_size_gb = 30
  }

  service_principal {
    client_id     = "00000000-0000-0000-0000-000000000000"
    client_secret = "00000000000000000000000000000000"
  }
}
```

```

tags = {
  Environment = "Production"
}

output "client_certificate" {
  value = "${azurerm_kubernetes_cluster.test.kube_config.0.client_certificate}"
}

output "kube_config" {
  value = "${azurerm_kubernetes_cluster.test.kube_config_raw}"
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Managed Kubernetes Cluster to create. Changing this forces a new resource to be created.
- **location** - (Required) The location where the Managed Kubernetes Cluster should be created. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) Specifies the Resource Group where the Managed Kubernetes Cluster should exist. Changing this forces a new resource to be created.
- **agent_pool_profile** - (Required) An `agent_pool_profile` block. Currently only one agent pool can exist.
- **dns_prefix** - (Required) DNS prefix specified when creating the managed cluster. Changing this forces a new resource to be created.

NOTE: The `dns_prefix` must contain between 3 and 45 characters, and can contain only letters, numbers, and hyphens. It must start with a letter and must end with a letter or a number.

- **service_principal** - (Required) A `service_principal` block as documented below.

-
- **addon_profile** - (Optional) A `addon_profile` block.
 - **kubernetes_version** - (Optional) Version of Kubernetes specified when creating the AKS managed cluster. If not specified, the latest recommended version will be used at provisioning time (but won't auto-upgrade).
 - **linux_profile** - (Optional) A `linux_profile` block.

- `network_profile` - (Optional) A `network_profile` block.

NOTE: If `network_profile` is not defined, `kubenet` profile will be used by default.

- `role_based_access_control` - (Optional) A `role_based_access_control` block. Changing this forces a new resource to be created.
 - `tags` - (Optional) A mapping of tags to assign to the resource.
-

A `addon_profile` block supports the following:

- `aci_connector_linux` - (Optional) A `aci_connector_linux` block. For more details, please visit [Create and configure an AKS cluster to use virtual nodes](#).
 - `http_application_routing` - (Optional) A `http_application_routing` block.
 - `oms_agent` - (Optional) A `oms_agent` block. For more details, please visit [How to onboard Azure Monitor for containers](#).
-

A `agent_pool_profile` block supports the following:

- `name` - (Required) Unique name of the Agent Pool Profile in the context of the Subscription and Resource Group. Changing this forces a new resource to be created.
- `count` - (Required) Number of Agents (VMs) in the Pool. Possible values must be in the range of 1 to 100 (inclusive). Defaults to 1.
- `vm_size` - (Required) The size of each VM in the Agent Pool (e.g. `Standard_F1`). Changing this forces a new resource to be created.
- `max_pods` - (Optional) The maximum number of pods that can run on each agent.
- `os_disk_size_gb` - (Optional) The Agent Operating System disk size in GB. Changing this forces a new resource to be created.
- `os_type` - (Optional) The Operating System used for the Agents. Possible values are `Linux` and `Windows`. Changing this forces a new resource to be created. Defaults to `Linux`.
- `vnet_subnet_id` - (Optional) The ID of the Subnet where the Agents in the Pool should be provisioned. Changing this forces a new resource to be created.

NOTE: A route table should be configured on this Subnet.

A `azure_active_directory` block supports the following:

- **client_app_id** - (Required) The Client ID of an Azure Active Directory Application. Changing this forces a new resource to be created.
 - **server_app_id** - (Required) The Server ID of an Azure Active Directory Application. Changing this forces a new resource to be created.
 - **server_app_secret** - (Required) The Server Secret of an Azure Active Directory Application. Changing this forces a new resource to be created.
 - **tenant_id** - (Optional) The Tenant ID used for Azure Active Directory Application. If this isn't specified the Tenant ID of the current Subscription is used. Changing this forces a new resource to be created.
-

A **http_application_routing** block supports the following:

- **enabled** (Required) Is HTTP Application Routing Enabled? Changing this forces a new resource to be created.
-

A **linux_profile** block supports the following:

- **admin_username** - (Required) The Admin Username for the Cluster. Changing this forces a new resource to be created.
 - **ssh_key** - (Required) An **ssh_key** block. Only one is currently allowed. Changing this forces a new resource to be created.
-

A **network_profile** block supports the following:

- **network_plugin** - (Required) Network plugin to use for networking. Currently supported values are **azure** and **kubenet**. Changing this forces a new resource to be created.

NOTE: When **network_plugin** is set to **azure** - the **vnet_subnet_id** field in the **agent_pool_profile** block must be set.

- **network_policy** - (Optional) Sets up network policy to be used with Azure CNI. Network policy allows us to control the traffic flow between pods. This field can only be set when **network_plugin** is set to **azure**. Currently supported values are **calico** and **azure**. Changing this forces a new resource to be created.
- **dns_service_ip** - (Optional) IP address within the Kubernetes service address range that will be used by cluster service discovery (kube-dns). This is required when **network_plugin** is set to **azure**. Changing this forces a new resource to be created.
- **docker_bridge_cidr** - (Optional) IP address (in CIDR notation) used as the Docker bridge IP address on nodes. This is required when

`network_plugin` is set to `azure`. Changing this forces a new resource to be created.

- `pod_cidr` - (Optional) The CIDR to use for pod IP addresses. This field can only be set when `network_plugin` is set to `kubenet`. Changing this forces a new resource to be created.
- `service_cidr` - (Optional) The Network Range used by the Kubernetes service. This is required when `network_plugin` is set to `azure`. Changing this forces a new resource to be created.

NOTE: This range should not be used by any network element on or connected to this VNet. Service address CIDR must be smaller than /12.

Examples of how to use AKS with Advanced Networking can be found in the `./examples/kubernetes/` directory in the Github repository.

A `oms_agent` block supports the following:

- `enabled` - (Required) Is the OMS Agent Enabled?
- `log_analytics_workspace_id` - (Required) The ID of the Log Analytics Workspace which the OMS Agent should send data to.

A `aci_connector_linux` block supports the following:

- `enabled` - (Required) Is the virtual node addon enabled?
- `subnet_name` - (Required) The subnet name for the virtual nodes to run.

A `role_based_access_control` block supports the following:

- `azure_active_directory` - (Optional) An `azure_active_directory` block. Changing this forces a new resource to be created.
- `enabled` - (Required) Is Role Based Access Control Enabled? Changing this forces a new resource to be created.

A `service_principal` block supports the following:

- `client_id` - (Required) The Client ID for the Service Principal. Changing this forces a new resource to be created.
- `client_secret` - (Required) The Client Secret for the Service Principal. Changing this forces a new resource to be created.

A `ssh_key` block supports the following:

- **key_data** - (Required) The Public SSH Key used to access the cluster. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The Kubernetes Managed Cluster ID.
- **fqdn** - The FQDN of the Azure Kubernetes Managed Cluster.
- **kube_admin_config** - A `kube_admin_config` block as defined below. This is only available when Role Based Access Control with Azure Active Directory is enabled.
- **kube_admin_config_raw** - Raw Kubernetes config for the admin account to be used by kubectl and other compatible tools. This is only available when Role Based Access Control with Azure Active Directory is enabled.
- **kube_config** - A `kube_config` block as defined below.
- **kube_config_raw** - Raw Kubernetes config to be used by kubectl and other compatible tools
- **http_application_routing** - A `http_application_routing` block as defined below.
- **node_resource_group** - The auto-generated Resource Group which contains the resources for this Managed Kubernetes Cluster.

A `http_application_routing` block exports the following:

- **http_application_routing_zone_name** - The Zone Name of the HTTP Application Routing.
-

The `kube_admin_config` and `kube_config` blocks export the following::

- **client_key** - Base64 encoded private key used by clients to authenticate to the Kubernetes cluster.
- **client_certificate** - Base64 encoded public certificate used by clients to authenticate to the Kubernetes cluster.
- **cluster_ca_certificate** - Base64 encoded public CA certificate used as the root of trust for the Kubernetes cluster.
- **host** - The Kubernetes cluster server host.
- **username** - A username used to authenticate to the Kubernetes cluster.

- **password** - A password or token used to authenticate to the Kubernetes cluster.

NOTE: It's possible to use these credentials with the Kubernetes Provider like so:

```
provider "kubernetes" {
  host          = "${azurerm_kubernetes_cluster.main.kube_config.0.host}"
  username      = "${azurerm_kubernetes_cluster.main.kube_config.0.username}"
  password      = "${azurerm_kubernetes_cluster.main.kube_config.0.password}"
  client_certificate = "${base64decode(azurerm_kubernetes_cluster.main.kube_config.0.client_certificate)}"
  client_key      = "${base64decode(azurerm_kubernetes_cluster.main.kube_config.0.client_key)}"
  cluster_ca_certificate = "${base64decode(azurerm_kubernetes_cluster.main.kube_config.0.cluster_ca_certificate)}"
}
```

» Import

Managed Kubernetes Clusters can be imported using the `resource id`, e.g.

```
terraform import azurerm_kubernetes_cluster.cluster1 /subscriptions/00000000-0000-0000-0000-
```

» `azurerm_cognitive_account`

Manages a Cognitive Services Account.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "example-resources"
  location = "West Europe"
}

resource "azurerm_cognitive_account" "test" {
  name                = "example-account"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  kind                 = "Face"

  sku {
    name = "S0"
    tier  = "Standard"
  }
}
```

```

tags = {
  Acceptance = "Test"
}
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Cognitive Service Account. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the Cognitive Service Account is created. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **kind** - (Required) Specifies the type of Cognitive Service Account that should be created. Possible values are `Academic`, `Bing.Autosuggest`, `Bing.Autosuggest.v7`, `Bing.CustomSearch`, `Bing.Search`, `Bing.Search.v7`, `Bing.Speech`, `Bing.SpellCheck`, `Bing.SpellCheck.v7`, `ComputerVision`, `ContentModerator`, `CustomSpeech`, `Emotion`, `Face`, `LUIS`, `Recommendations`, `SpeakerRecognition`, `Speech`, `SpeechServices`, `SpeechTranslation`, `TextAnalytics`, `TextTranslation` and `WebLM`. Changing this forces a new resource to be created.
- **sku** - (Required) A sku block as defined below.
- **tags** - (Optional) A mapping of tags to assign to the resource.

A sku block supports the following:

- **name** - (Required) Specifies the Name of the Sku. Possible values are `F0`, `S0`, `S1`, `S2`, `S3`, `S4`, `S5`, `S6`, `P0`, `P1` and `P2`.
- **tier** - (Required) Specifies the Tier of the Sku. Possible values include `Free`, `Standard` and `Premium`.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Cognitive Service Account.

- **endpoint** - The endpoint used to connect to the Cognitive Service Account.
- **primary_access_key** - A primary access key which can be used to connect to the Cognitive Service Account.
- **secondary_access_key** - The secondary access key which can be used to connect to the Cognitive Service Account.

» Import

Cognitive Service Accounts can be imported using the `resource id`, e.g.

```
terraform import azurerm_cognitive_account.account1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_cosmosdb_account

Manages a CosmosDB (formally DocumentDB) Account.

» Example Usage

```
resource "azurerm_resource_group" "rg" {
  name       = "${var.resource_group_name}"
  location   = "${var.resource_group_location}"
}

resource "random_integer" "ri" {
  min = 10000
  max = 99999
}

resource "azurerm_cosmosdb_account" "db" {
  name                = "tfex-cosmos-db-${random_integer.ri.result}"
  location             = "${azurerm_resource_group.rg.location}"
  resource_group_name = "${azurerm_resource_group.rg.name}"
  offer_type          = "Standard"
  kind                 = "GlobalDocumentDB"

  enable_automatic_failover = true

  consistency_policy {
    consistency_level      = "BoundedStaleness"
    max_interval_in_seconds = 10
    max_staleness_prefix   = 200
  }
}
```

```

}

geo_location {
  location          = "${var.failover_location}"
  failover_priority = 1
}

geo_location {
  prefix            = "tfex-cosmos-db-${random_integer.ri.result}-customid"
  location          = "${azurerm_resource_group.rg.location}"
  failover_priority = 0
}
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the CosmosDB Account. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the CosmosDB Account is created. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **tags** - (Optional) A mapping of tags to assign to the resource.
- **offer_type** - (Required) Specifies the Offer Type to use for this CosmosDB Account - currently this can only be set to **Standard**.
- **kind** - (Optional) Specifies the Kind of CosmosDB to create - possible values are **GlobalDocumentDB** and **MongoDB**. Defaults to **GlobalDocumentDB**. Changing this forces a new resource to be created.
- **consistency_policy** - (Required) Specifies a **consistency_policy** resource, used to define the consistency policy for this CosmosDB account.
- **geo_location** - (Required) Specifies a **geo_location** resource, used to define where data should be replicated with the **failover_priority** 0 specifying the primary location.
- **ip_range_filter** - (Optional) CosmosDB Firewall Support: This value specifies the set of IP addresses or IP address ranges in CIDR form to be included as the allowed list of client IP's for a given database account. IP addresses/ranges must be comma separated and must not contain any spaces.

- **enable_automatic_failover** - (Optional) Enable automatic fail over for this Cosmos DB account.
- **capabilities** - (Optional) The capabilities which should be enabled for this Cosmos DB account. Possible values are **EnableAggregationPipeline**, **EnableCassandra**, **EnableGremlin**, **EnableTable**, **MongoDBv3.4**, and **mongoEnableDocLevelTTL**.
- **is_virtual_network_filter_enabled** - (Optional) Enables virtual network filtering for this Cosmos DB account.
- **virtual_network_rule** - (Optional) Specifies a **virtual_network_rules** resource, used to define which subnets are allowed to access this CosmosDB account.
- **enable_multiple_write_locations** - (Optional) Enable multi-master support for this Cosmos DB account.

consistency_policy Configures the database consistency and supports the following:

- **consistency_level** - (Required) The Consistency Level to use for this CosmosDB Account - can be either **BoundedStaleness**, **Eventual**, **Session**, **Strong** or **ConsistentPrefix**.
- **max_interval_in_seconds** - (Optional) When used with the Bounded Staleness consistency level, this value represents the time amount of staleness (in seconds) tolerated. Accepted range for this value is 5 - 86400 (1 day). Defaults to 5. Required when **consistency_level** is set to **BoundedStaleness**.
- **max_staleness_prefix** - (Optional) When used with the Bounded Staleness consistency level, this value represents the number of stale requests tolerated. Accepted range for this value is 10 – 2147483647. Defaults to 100. Required when **consistency_level** is set to **BoundedStaleness**.

Note: **max_interval_in_seconds** and **max_staleness_prefix** can only be set to custom values when **consistency_level** is set to **BoundedStaleness** - otherwise they will return the default values shown above.

geo_location Configures the geographic locations the data is replicated to and supports the following:

- **prefix** - (Optional) The string used to generate the document endpoints for this region. If not specified it defaults to **\${cosmosdb_account.name}-\${location}**. Changing this causes the location to be deleted and re-provisioned and cannot be changed for the location with failover priority 0.
- **location** - (Required) The name of the Azure region to host replicated data.
- **failover_priority** - (Required) The failover priority of the region. A failover priority of 0 indicates a write region. The maximum value for a failover priority = (total number of regions - 1). Failover priority values

must be unique for each of the regions in which the database account exists. Changing this causes the location to be re-provisioned and cannot be changed for the location with failover priority 0.

capabilities Configures the capabilities to enable for this Cosmos DB account:

- **name** - (Required) The capability to enable - Possible values are `EnableTable`, `EnableCassandra`, and `EnableGremlin`.

NOTE: The **prefix** and **failover_priority** fields of a location cannot be changed for the location with a failover priority of 0.

virtual_network_rule Configures the virtual network subnets allowed to access this Cosmos DB account and supports the following:

- **id** - (Required) The ID of the virtual network subnet.

» Attributes Reference

The following attributes are exported:

- **id** - The CosmosDB Account ID.
- **endpoint** - The endpoint used to connect to the CosmosDB account.
- **read_endpoints** - A list of read endpoints available for this CosmosDB account.
- **write_endpoints** - A list of write endpoints available for this CosmosDB account.
- **primary_master_key** - The Primary master key for the CosmosDB Account.
- **secondary_master_key** - The Secondary master key for the CosmosDB Account.
- **primary_readonly_master_key** - The Primary read-only master Key for the CosmosDB Account.
- **secondary_readonly_master_key** - The Secondary read-only master key for the CosmosDB Account.
- **connection_strings** - A list of connection strings available for this CosmosDB account. If the kind is `GlobalDocumentDB`, this will be empty.

» Import

CosmosDB Accounts can be imported using the **resource id**, e.g.

```
terraform import azurerm_cosmosdb_account.account1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_mariadb_database

Manages a MariaDB Database within a MariaDB Server

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "tfex-mariadb-database-RG"
  location = "westeurope"
}

resource "azurerm_mariadb_server" "example" {
  name                = "mariadb-svr"
  location             = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"

  sku {
    name       = "B_Gen5_2"
    capacity   = 2
    tier        = "Basic"
    family     = "Gen5"
  }

  storage_profile {
    storage_mb           = 51200
    backup_retention_days = 7
    geo_redundant_backup = "Disabled"
  }

  administrator_login          = "acctestun"
  administrator_login_password = "H@Sh1CoR3!"
  version                     = "10.2"
  ssl_enforcement              = "Enabled"
}

resource "azurerm_mariadb_database" "example" {
  name                = "mariadb_database"
  resource_group_name = "${azurerm_resource_group.example.name}"
  server_name         = "${azurerm_mariadb_server.example.name}"
  charset              = "utf8"
  collation            = "utf8_general_ci"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the MariaDB Database, which needs to be a valid MariaDB identifier. Changing this forces a new resource to be created.
- **server_name** - (Required) Specifies the name of the MariaDB Server. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the MariaDB Server exists. Changing this forces a new resource to be created.
- **charset** - (Required) Specifies the Charset for the MariaDB Database, which needs to be a valid MariaDB Charset. Changing this forces a new resource to be created.
- **collation** - (Required) Specifies the Collation for the MariaDB Database, which needs to be a valid MariaDB Collation. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the MariaDB Database.

» Import

MariaDB Database's can be imported using the **resource id**, e.g.

```
terraform import azurerm_mariadb_database.database1 /subscriptions/00000000-0000-0000-0000-000000000000/resourceGroups/test-rg/providers/Microsoft.DBforMariaDB/servers/test-server/databases/test-database
```

» azurerm__mariadb__server

Manages a MariaDB Server.

NOTE MariaDB Server is currently in Public Preview. You can find more information, including how to register for the Public Preview here.

» Example Usage

```
resource "azurerm_resource_group" "test" {
```

```

    name      = "api-rg-pro"
    location = "West Europe"
}

resource "azurerm_mariadb_server" "test" {
  name      = "mariadb-server-1"
  location  = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  sku {
    name      = "B_Gen5_2"
    capacity = 2
    tier       = "Basic"
    family    = "Gen5"
  }

  storage_profile {
    storage_mb          = 5120
    backup_retention_days = 7
    geo_redundant_backup = "Disabled"
  }

  administrator_login          = "mariadbadmin"
  administrator_login_password = "H@Sh1CoR3!"
  version                      = "10.2"
  ssl_enforcement              = "Enabled"
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the MariaDB Server. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the MariaDB Server. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **sku** - (Required) A **sku** block as defined below.
- **storage_profile** - (Required) A **storage_profile** block as defined below.

- **administrator_login** - (Required) The Administrator Login for the MariaDB Server. Changing this forces a new resource to be created.
- **administrator_login_password** - (Required) The Password associated with the **administrator_login** for the MariaDB Server.
- **version** - (Required) Specifies the version of MariaDB to use. The valid value is 10.2. Changing this forces a new resource to be created.
- **ssl_enforcement** - (Required) Specifies if SSL should be enforced on connections. Possible values are **Enabled** and **Disabled**.
- **tags** - (Optional) A mapping of tags to assign to the resource.

A **sku** block supports the following:

- **name** - (Required) Specifies the SKU Name for this MariaDB Server. The name of the SKU, follows the **tier** + **family** + **cores** pattern (e.g. **B_Gen5_1**, **GP_Gen5_8**). For more information see the product documentation.
- **capacity** - (Required) The scale up/out capacity, representing server's compute units.
- **tier** - (Required) The tier of the particular SKU. Possible values are **Basic**, **GeneralPurpose**, and **MemoryOptimized**. For more information see the product documentation.
- **family** - (Required) The **family** of the hardware (e.g. **Gen5**), before selecting your **family** check the product documentation for availability in your region.

A **storage_profile** block supports the following:

- **storage_mb** - (Required) Max storage allowed for a server. Possible values are between 5120 MB (5GB) and 1024000MB (1TB) for the Basic SKU and between 5120 MB (5GB) and 4096000 MB (4TB) for General Purpose/Memory Optimized SKUs. For more information see the product documentation.
- **backup_retention_days** - (Optional) Backup retention days for the server, supported values are between 7 and 35 days.
- **geo_redundant_backup** - (Optional) Enable Geo-redundant or not for server backup. Valid values for this property are **Enabled** or **Disabled**.

NOTE: Geo Redundant Backups cannot be configured when using the **Basic** tier.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the MariaDB Server.
- `fqdn` - The FQDN of the MariaDB Server.

» Import

MariaDB Server's can be imported using the `resource id`, e.g.

```
terraform import azurerm_mariadb_server.server1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_mysql_configuration

Sets a MySQL Configuration value on a MySQL Server.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "api-rg-pro"
  location = "West Europe"
}

resource "azurerm_mysql_server" "test" {
  name                = "mysql-server-1"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  sku {
    name       = "B_Gen4_2"
    capacity   = 2
    tier        = "Basic"
    family     = "Gen4"
  }

  storage_profile {
    storage_mb           = 5120
    backup_retention_days = 7
    geo_redundant_backup = "Disabled"
  }

  administrator_login = "psqladminun"
```

```

    administrator_login_password = "H@Sh1CoR3!"
    version                      = "5.7"
    ssl_enforcement              = "Enabled"
  }

  resource "azurerm_mysql_configuration" "test" {
    name                = "interactive_timeout"
    resource_group_name = "${azurerm_resource_group.test.name}"
    server_name         = "${azurerm_mysql_server.test.name}"
    value               = "600"
  }

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the MySQL Configuration, which needs to be a valid MySQL configuration name. Changing this forces a new resource to be created.
- **server_name** - (Required) Specifies the name of the MySQL Server. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the MySQL Server exists. Changing this forces a new resource to be created.
- **value** - (Required) Specifies the value of the MySQL Configuration. See the MySQL documentation for valid values.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the MySQL Configuration.

» Import

MySQL Configurations can be imported using the `resource id`, e.g.

```
terraform import azurerm_mysql_configuration.interactive_timeout /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm__mysql__database

Manages a MySQL Database within a MySQL Server

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "api-rg-pro"
  location  = "West Europe"
}

resource "azurerm_mysql_server" "test" {
  name                = "mysql-server-1"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  sku {
    name       = "B_Gen4_2"
    capacity  = 2
    tier       = "Basic"
    family    = "Gen4"
  }

  storage_profile {
    storage_mb          = 5120
    backup_retention_days = 7
    geo_redundant_backup = "Disabled"
  }

  administrator_login          = "mysqladminun"
  administrator_login_password = "H@Sh1CoR3!"
  version                     = "5.7"
  ssl_enforcement              = "Enabled"
}

resource "azurerm_mysql_database" "test" {
  name                = "exampleldb"
  resource_group_name = "${azurerm_resource_group.test.name}"
  server_name         = "${azurerm_mysql_server.test.name}"
  charset             = "utf8"
  collation            = "utf8_unicode_ci"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the MySQL Database, which needs to be a valid MySQL identifier. Changing this forces a new resource to be

created.

- **server_name** - (Required) Specifies the name of the MySQL Server. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the MySQL Server exists. Changing this forces a new resource to be created.
- **charset** - (Required) Specifies the Charset for the MySQL Database, which needs to be a valid MySQL Charset. Changing this forces a new resource to be created.
- **collation** - (Required) Specifies the Collation for the MySQL Database, which needs to be a valid MySQL Collation. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the MySQL Database.

» Import

MySQL Database's can be imported using the **resource id**, e.g.

```
terraform import azurerm_mysql_database.database1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm__mysql__firewall__rule

Manages a Firewall Rule for a MySQL Server

» Example Usage (Single IP Address)

```
resource "azurerm_resource_group" "test" {
  name     = "api-rg-pro"
  location = "West Europe"
}

resource "azurerm_mysql_server" "test" {
  # ...
}
```

```
resource "azurerm_mysql_firewall_rule" "test" {
  name                = "office"
  resource_group_name = "${azurerm_resource_group.test.name}"
  server_name         = "${azurerm_mysql_server.test.name}"
  start_ip_address    = "40.112.8.12"
  end_ip_address      = "40.112.8.12"
}
```

» Example Usage (IP Range)

```
resource "azurerm_resource_group" "test" {
  name     = "api-rg-pro"
  location = "West Europe"
}

resource "azurerm_mysql_server" "test" {
  # ...
}

resource "azurerm_mysql_firewall_rule" "test" {
  name                = "office"
  resource_group_name = "${azurerm_resource_group.test.name}"
  server_name         = "${azurerm_mysql_server.test.name}"
  start_ip_address    = "40.112.0.0"
  end_ip_address      = "40.112.255.255"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the MySQL Firewall Rule. Changing this forces a new resource to be created.
- **server_name** - (Required) Specifies the name of the MySQL Server. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the MySQL Server exists. Changing this forces a new resource to be created.
- **start_ip_address** - (Required) Specifies the Start IP Address associated with this Firewall Rule. Changing this forces a new resource to be created.
- **end_ip_address** - (Required) Specifies the End IP Address associated with this Firewall Rule. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the MySQL Firewall Rule.

» Import

MySQL Firewall Rule's can be imported using the `resource id`, e.g.

```
terraform import azurerm_mysql_firewall_rule.rule1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_mysql_server`

Manages a MySQL Server.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "api-rg-pro"
  location = "West Europe"
}

resource "azurerm_mysql_server" "test" {
  name                = "mysql-server-1"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  sku {
    name       = "B_Gen5_2"
    tier        = "Basic"
    family     = "Gen5"
    capacity   = "2"
  }

  storage_profile {
    storage_mb           = 5120
    backup_retention_days = 7
    geo_redundant_backup = "Disabled"
  }

  administrator_login          = "mysqladminun"
  administrator_login_password = "H@Sh1CoR3!"
}
```

```

    version                = "5.7"
    ssl_enforcement        = "Enabled"
  }

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the MySQL Server. Changing this forces a new resource to be created. This needs to be globally unique within Azure.
- **resource_group_name** - (Required) The name of the resource group in which to create the MySQL Server. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **sku** - (Required) A **sku** block as defined below.
- **storage_profile** - (Required) A **storage_profile** block as defined below.
- **administrator_login** - (Required) The Administrator Login for the MySQL Server. Changing this forces a new resource to be created.
- **administrator_login_password** - (Required) The Password associated with the **administrator_login** for the MySQL Server.
- **version** - (Required) Specifies the version of MySQL to use. Valid values are 5.6 and 5.7. Changing this forces a new resource to be created.
- **ssl_enforcement** - (Required) Specifies if SSL should be enforced on connections. Possible values are **Enabled** and **Disabled**.
- **tags** - (Optional) A mapping of tags to assign to the resource.

sku supports the following:

- **name** - (Required) Specifies the SKU Name for this MySQL Server. The name of the SKU, follows the **tier** + **family** + **cores** pattern (e.g. B_Gen4_1, GP_Gen5_8). For more information see the product documentation.
- **capacity** - (Required) The scale up/out capacity, representing server's compute units.

- **tier** - (Required) The tier of the particular SKU. Possible values are **Basic**, **GeneralPurpose**, and **MemoryOptimized**. For more information see the product documentation.
- **family** - (Required) The **family** of hardware **Gen4** or **Gen5**, before selecting your **family** check the product documentation for availability in your region.

storage_profile supports the following:

- **storage_mb** - (Required) Max storage allowed for a server. Possible values are between 5120 MB(5GB) and 1048576 MB(1TB) for the Basic SKU and between 5120 MB(5GB) and 4194304 MB(4TB) for General Purpose/Memory Optimized SKUs. For more information see the product documentation.
- **backup_retention_days** - (Optional) Backup retention days for the server, supported values are between 7 and 35 days.
- **geo_redundant_backup** - (Optional) Enable Geo-redundant or not for server backup. Valid values for this property are **Enabled** or **Disabled**, not supported for the **basic** tier.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the MySQL Server.
- **fqdn** - The FQDN of the MySQL Server.

» Import

MySQL Server's can be imported using the **resource id**, e.g.

```
terraform import azurerm_mysql_server.server1 /subscriptions/00000000-0000-0000-0000-00000000
```

» azurerm__mysql__virtual__network__rule

Manages a MySQL Virtual Network Rule.

NOTE: MySQL Virtual Network Rules can only be used with SKU Tiers of **GeneralPurpose** or **MemoryOptimized**

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "example-resources"
  location  = "West Europe"
}

resource "azurerm_virtual_network" "test" {
  name            = "example-vnet"
  address_space   = ["10.7.29.0/29"]
  location        = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_subnet" "internal" {
  name                 = "internal"
  resource_group_name  = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix        = "10.7.29.0/29"
  service_endpoints     = ["Microsoft.Sql"]
}

resource "azurerm_mysql_server" "test" {
  name                        = "mysql-server-1"
  location                   = "${azurerm_resource_group.test.location}"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  administrator_login        = "mysqladminun"
  administrator_login_password = "H@Sh1CoR3!"
  version                    = "5.7"
  ssl_enforcement            = "Enabled"

  sku {
    name       = "GP_Gen5_2"
    capacity   = 2
    tier        = "GeneralPurpose"
    family     = "Gen5"
  }

  storage_profile {
    storage_mb           = 5120
    backup_retention_days = 7
    geo_redundant_backup = "Disabled"
  }
}
```

```
resource "azurerm_mysql_virtual_network_rule" "test" {
  name                = "mysql-vnet-rule"
  resource_group_name = "${azurerm_resource_group.test.name}"
  server_name         = "${azurerm_mysql_server.test.name}"
  subnet_id           = "${azurerm_subnet.internal.id}"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the MySQL Virtual Network Rule. Cannot be empty and must only contain alphanumeric characters and hyphens. Cannot start with a number, and cannot start or end with a hyphen. Changing this forces a new resource to be created.

NOTE: **name** must be between 1-128 characters long and must satisfy all of the requirements below: 1. Contains only alphanumeric and hyphen characters 2. Cannot start with a number or hyphen 3. Cannot end with a hyphen

- **resource_group_name** - (Required) The name of the resource group where the MySQL server resides. Changing this forces a new resource to be created.
- **server_name** - (Required) The name of the SQL Server to which this MySQL virtual network rule will be applied to. Changing this forces a new resource to be created.
- **subnet_id** - (Required) The ID of the subnet that the MySQL server will be connected to.

NOTE: Due to a bug in the Azure API this resource currently doesn't expose the `ignore_missing_vnet_service_endpoint` field and defaults this to `false`. Terraform will check during the provisioning of the Virtual Network Rule that the Subnet contains the Service Rule to verify that the Virtual Network Rule can be created.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the MySQL Virtual Network Rule.

» Import

MySQL Virtual Network Rules can be imported using the `resource id`, e.g.


```
terraform import azurerm_mysql_virtual_network_rule.rule1 /subscriptions/00000000-0000-0000-
```

» azurerm_postgresql_configuration

Sets a PostgreSQL Configuration value on a PostgreSQL Server.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "api-rg-pro"
  location  = "West Europe"
}

resource "azurerm_postgresql_server" "test" {
  name                = "postgresql-server-1"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  sku {
    name      = "B_Gen4_2"
    capacity  = 2
    tier       = "Basic"
    family    = "Gen4"
  }

  storage_profile {
    storage_mb           = 5120
    backup_retention_days = 7
    geo_redundant_backup = "Disabled"
  }

  administrator_login          = "psqladminun"
  administrator_login_password = "H@Sh1CoR3!"
  version                      = "9.5"
  ssl_enforcement              = "Enabled"
}

resource "azurerm_postgresql_configuration" "test" {
  name                = "backslash_quote"
  resource_group_name = "${azurerm_resource_group.test.name}"
  server_name         = "${azurerm_postgresql_server.test.name}"
  value               = "on"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the PostgreSQL Configuration, which needs to be a valid PostgreSQL configuration name. Changing this forces a new resource to be created.
- **server_name** - (Required) Specifies the name of the PostgreSQL Server. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the PostgreSQL Server exists. Changing this forces a new resource to be created.
- **value** - (Required) Specifies the value of the PostgreSQL Configuration. See the PostgreSQL documentation for valid values.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the PostgreSQL Configuration.

» Import

PostgreSQL Configurations can be imported using the **resource id**, e.g.

```
terraform import azurerm_postgresql_configuration.backslash_quote /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_postgresql_database

Manages a PostgreSQL Database within a PostgreSQL Server

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "api-rg-pro"
  location = "West Europe"
}

resource "azurerm_postgresql_server" "test" {
  name          = "postgresql-server-1"
  location      = "${azurerm_resource_group.test.location}"
}
```

```

resource_group_name = "${azurerm_resource_group.test.name}"

sku {
  name      = "B_Gen4_2"
  capacity = 2
  tier       = "Basic"
  family    = "Gen4"
}

storage_profile {
  storage_mb           = 5120
  backup_retention_days = 7
  geo_redundant_backup = "Disabled"
}

administrator_login          = "psqladminun"
administrator_login_password = "H@Sh1CoR3!"
version                      = "9.5"
ssl_enforcement              = "Enabled"
}

resource "azurerm_postgresql_database" "test" {
  name                = "exampleldb"
  resource_group_name = "${azurerm_resource_group.test.name}"
  server_name         = "${azurerm_postgresql_server.test.name}"
  charset              = "UTF8"
  collation            = "English_United States.1252"
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the PostgreSQL Database, which needs to be a valid PostgreSQL identifier. Changing this forces a new resource to be created.
- **server_name** - (Required) Specifies the name of the PostgreSQL Server. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the PostgreSQL Server exists. Changing this forces a new resource to be created.
- **charset** - (Required) Specifies the Charset for the PostgreSQL Database, which needs to be a valid PostgreSQL Charset. Changing this forces a new resource to be created.

- **collation** - (Required) Specifies the Collation for the PostgreSQL Database, which needs to be a valid PostgreSQL Collation. Note that Microsoft uses different notation - en-US instead of en_US. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the PostgreSQL Database.

» Import

PostgreSQL Database's can be imported using the **resource id**, e.g.

```
terraform import azurerm_postgresql_database.database1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_postgresql_firewall_rule

Manages a Firewall Rule for a PostgreSQL Server

» Example Usage (Single IP Address)

```
resource "azurerm_resource_group" "test" {
  name     = "api-rg-pro"
  location = "West Europe"
}

resource "azurerm_postgresql_server" "test" {
  # ...
}

resource "azurerm_postgresql_firewall_rule" "test" {
  name                = "office"
  resource_group_name = "${azurerm_resource_group.test.name}"
  server_name         = "${azurerm_postgresql_server.test.name}"
  start_ip_address    = "40.112.8.12"
  end_ip_address      = "40.112.8.12"
}
```

» Example Usage (IP Range)

```
resource "azurerm_resource_group" "test" {
  name      = "api-rg-pro"
  location = "West Europe"
}

resource "azurerm_postgresql_server" "test" {
  # ...
}

resource "azurerm_postgresql_firewall_rule" "test" {
  name                = "office"
  resource_group_name = "${azurerm_resource_group.test.name}"
  server_name         = "${azurerm_postgresql_server.test.name}"
  start_ip_address    = "40.112.0.0"
  end_ip_address      = "40.112.255.255"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the PostgreSQL Firewall Rule. Changing this forces a new resource to be created.
- **server_name** - (Required) Specifies the name of the PostgreSQL Server. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the PostgreSQL Server exists. Changing this forces a new resource to be created.
- **start_ip_address** - (Required) Specifies the Start IP Address associated with this Firewall Rule. Changing this forces a new resource to be created.
- **end_ip_address** - (Required) Specifies the End IP Address associated with this Firewall Rule. Changing this forces a new resource to be created.

NOTE: The Azure feature `Allow access to Azure services` can be enabled by setting `start_ip_address` and `end_ip_address` to `0.0.0.0` which (is documented in the Azure API Docs).

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the PostgreSQL Firewall Rule.

» Import

PostgreSQL Firewall Rule's can be imported using the `resource id`, e.g.

```
terraform import azurearm_postgresql_firewall_rule.rule1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurearm_postgresql_server`

Manage a PostgreSQL Server.

» Example Usage

```
resource "azurearm_resource_group" "test" {
  name       = "api-rg-pro"
  location   = "West Europe"
}

resource "azurearm_postgresql_server" "test" {
  name                = "postgresql-server-1"
  location             = "${azurearm_resource_group.test.location}"
  resource_group_name = "${azurearm_resource_group.test.name}"

  sku {
    name       = "B_Gen4_2"
    capacity   = 2
    tier        = "Basic"
    family     = "Gen4"
  }

  storage_profile {
    storage_mb           = 5120
    backup_retention_days = 7
    geo_redundant_backup = "Disabled"
  }

  administrator_login          = "psqladminun"
  administrator_login_password = "H@Sh1CoR3!"
  version                     = "9.5"
  ssl_enforcement              = "Enabled"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the PostgreSQL Server. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the PostgreSQL Server. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **sku** - (Required) A **sku** block as defined below.
- **storage_profile** - (Required) A **storage_profile** block as defined below.
- **administrator_login** - (Required) The Administrator Login for the PostgreSQL Server. Changing this forces a new resource to be created.
- **administrator_login_password** - (Required) The Password associated with the **administrator_login** for the PostgreSQL Server.
- **version** - (Required) Specifies the version of PostgreSQL to use. Valid values are 9.5, 9.6, 10, 10.0, and 10.2. Changing this forces a new resource to be created.
- **ssl_enforcement** - (Required) Specifies if SSL should be enforced on connections. Possible values are **Enabled** and **Disabled**.
- **tags** - (Optional) A mapping of tags to assign to the resource.

sku supports the following:

- **name** - (Required) Specifies the SKU Name for this PostgreSQL Server. The name of the SKU, follows the **tier** + **family** + **cores** pattern (e.g. B_Gen4_1, GP_Gen5_8). For more information see the product documentation.
- **capacity** - (Required) The scale up/out capacity, representing server's compute units.
- **tier** - (Required) The tier of the particular SKU. Possible values are **Basic**, **GeneralPurpose**, and **MemoryOptimized**. For more information see the product documentation.
- **family** - (Required) The **family** of hardware **Gen4** or **Gen5**, before selecting your **family** check the product documentation for availability in your region.

`storage_profile` supports the following:

- `storage_mb` - (Required) Max storage allowed for a server. Possible values are between 5120 MB(5GB) and 1048576 MB(1TB) for the Basic SKU and between 5120 MB(5GB) and 4194304 MB(4TB) for General Purpose/Memory Optimized SKUs. For more information see the product documentation.
- `backup_retention_days` - (Optional) Backup retention days for the server, supported values are between 7 and 35 days.
- `geo_redundant_backup` - (Optional) Enable Geo-redundant or not for server backup. Valid values for this property are `Enabled` or `Disabled`, not supported for the `basic` tier.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the PostgreSQL Server.
- `fqdn` - The FQDN of the PostgreSQL Server.

» Import

PostgreSQL Server's can be imported using the `resource id`, e.g.

```
terraform import azurerm_postgresql_server.server1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_postgresql_virtual_network_rule`

Manages a PostgreSQL Virtual Network Rule.

NOTE: PostgreSQL Virtual Network Rules can only be used with SKU Tiers of `GeneralPurpose` or `MemoryOptimized`

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "example-resources"
  location = "West US"
}
```



```

resource "azurerm_virtual_network" "test" {
  name                = "example-vnet"
  address_space       = ["10.7.29.0/29"]
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_subnet" "internal" {
  name                = "internal"
  resource_group_name = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix       = "10.7.29.0/29"
  service_endpoints    = ["Microsoft.Sql"]
}

resource "azurerm_postgresql_server" "test" {
  name                = "postgresql-server-1"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  sku {
    name       = "GP_Gen5_2"
    capacity   = 2
    tier        = "GeneralPurpose"
    family     = "Gen5"
  }

  storage_profile {
    storage_mb           = 5120
    backup_retention_days = 7
    geo_redundant_backup = "Disabled"
  }

  administrator_login          = "psqladminun"
  administrator_login_password = "H0Sh1CoR3!"
  version                     = "9.5"
  ssl_enforcement              = "Enabled"
}

resource "azurerm_postgresql_virtual_network_rule" "test" {
  name                = "postgresql-vnet-rule"
  resource_group_name = "${azurerm_resource_group.test.name}"
  server_name         = "${azurerm_postgresql_server.test.name}"
  subnet_id           = "${azurerm_subnet.internal.id}"
  ignore_missing_vnet_service_endpoint = true
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the PostgreSQL virtual network rule. Cannot be empty and must only contain alphanumeric characters and hyphens. Cannot start with a number, and cannot start or end with a hyphen. Changing this forces a new resource to be created.

NOTE: **name** must be between 1-128 characters long and must satisfy all of the requirements below: 1. Contains only alphanumeric and hyphen characters 2. Cannot start with a number or hyphen 3. Cannot end with a hyphen

- **resource_group_name** - (Required) The name of the resource group where the PostgreSQL server resides. Changing this forces a new resource to be created.
- **server_name** - (Required) The name of the SQL Server to which this PostgreSQL virtual network rule will be applied to. Changing this forces a new resource to be created.
- **subnet_id** - (Required) The ID of the subnet that the PostgreSQL server will be connected to.
- **ignore_missing_vnet_service_endpoint** - (Optional) Should the Virtual Network Rule be created before the Subnet has the Virtual Network Service Endpoint enabled? Defaults to **false**.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the PostgreSQL Virtual Network Rule.

» Import

PostgreSQL Virtual Network Rules can be imported using the **resource id**, e.g.

```
terraform import azurerm_postgresql_virtual_network_rule.rule1 /subscriptions/00000000-0000-
```

» azurerm__sql__database

Allows you to manage an Azure SQL Database

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "acceptanceTestResourceGroup1"
  location  = "West US"
}

resource "azurerm_sql_server" "test" {
  name                        = "mysqlserver"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  location                   = "West US"
  version                    = "12.0"
  administrator_login        = "4dm1n157r470r"
  administrator_login_password = "4-v3ry-53cr37-p455w0rd"
}

resource "azurerm_sql_database" "test" {
  name                = "mysqldatabase"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location            = "West US"
  server_name         = "${azurerm_sql_server.test.name}"

  tags = {
    environment = "production"
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the database.
- **resource_group_name** - (Required) The name of the resource group in which to create the database. This must be the same as Database Server resource group currently.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **server_name** - (Required) The name of the SQL Server on which to create the database.
- **create_mode** - (Optional) Specifies the type of database to create. Defaults to `Default`. See below for the accepted values/
- **import** - (Optional) A Database Import block as documented below. `create_mode` must be set to `Default`.

- `source_database_id` - (Optional) The URI of the source database if `create_mode` value is not `Default`.
- `restore_point_in_time` - (Optional) The point in time for the restore. Only applies if `create_mode` is `PointInTimeRestore` e.g. 2013-11-08T22:00:40Z
- `edition` - (Optional) The edition of the database to be created. Applies only if `create_mode` is `Default`. Valid values are: `Basic`, `Standard`, `Premium`, or `DataWarehouse`. Please see Azure SQL Database Service Tiers.
- `collation` - (Optional) The name of the collation. Applies only if `create_mode` is `Default`. Azure default is `SQL_LATIN1_GENERAL_CP1_CI_AS`. Changing this forces a new resource to be created.
- `max_size_bytes` - (Optional) The maximum size that the database can grow to. Applies only if `create_mode` is `Default`. Please see Azure SQL Database Service Tiers.
- `requested_service_objective_id` - (Optional) Use `requested_service_objective_id` or `requested_service_objective_name` to set the performance level for the database. Please see Azure SQL Database Service Tiers.
- `requested_service_objective_name` - (Optional) Use `requested_service_objective_name` or `requested_service_objective_id` to set the performance level for the database. Valid values are: `S0`, `S1`, `S2`, `S3`, `P1`, `P2`, `P4`, `P6`, `P11` and `ElasticPool`. Please see Azure SQL Database Service Tiers.
- `source_database_deletion_date` - (Optional) The deletion date time of the source database. Only applies to deleted databases where `create_mode` is `PointInTimeRestore`.
- `elastic_pool_name` - (Optional) The name of the elastic database pool.
- `threat_detection_policy` - (Optional) Threat detection policy configuration. The `threat_detection_policy` block supports fields documented below.
- `tags` - (Optional) A mapping of tags to assign to the resource.

`import` supports the following:

- `storage_uri` - (Required) Specifies the blob URI of the .bacpac file.
- `storage_key` - (Required) Specifies the access key for the storage account.
- `storage_key_type` - (Required) Specifies the type of access key for the storage account. Valid values are `StorageAccessKey` or `SharedAccessKey`.
- `administrator_login` - (Required) Specifies the name of the SQL administrator.
- `administrator_login_password` - (Required) Specifies the password of the SQL administrator.

- **authentication_type** - (Required) Specifies the type of authentication used to access the server. Valid values are **SQL** or **ADPassword**.
- **operation_mode** - (Optional) Specifies the type of import operation being performed. The only allowable value is **Import**.

threat_detection_policy supports the following:

- **state** - (Required) The State of the Policy. Possible values are **Enabled**, **Disabled** or **New**.
- **disabled_alerts** - (Optional) Specifies a list of alerts which should be disabled. Possible values include **Access_Anomaly**, **Sql_Injection** and **Sql_Injection_Vulnerability**.
- **email_account_admins** - (Optional) Should the account administrators be emailed when this alert is triggered?
- **email_addresses** - (Optional) A list of email addresses which alerts should be sent to.
- **retention_days** - (Optional) Specifies the number of days to keep in the Threat Detection audit logs.
- **storage_account_access_key** - (Optional) Specifies the identifier key of the Threat Detection audit storage account. Required if **state** is **Enabled**.
- **storage_endpoint** - (Optional) Specifies the blob storage endpoint (e.g. <https://MyAccount.blob.core.windows.net>). This blob storage will hold all Threat Detection audit logs. Required if **state** is **Enabled**.
- **use_server_default** - (Optional) Should the default server policy be used? Defaults to **Disabled**.

» Attributes Reference

The following attributes are exported:

- **id** - The SQL Database ID.
- **creation_date** - The creation date of the SQL Database.
- **default_secondary_location** - The default secondary location of the SQL Database.

» Import

SQL Databases can be imported using the **resource id**, e.g.

```
terraform import azurerm_sql_database.database1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» **azurerm_sql_active_directory_administrator**

Allows you to set a user or group as the AD administrator for an Azure SQL server

» **Example Usage**

```
data "azurerm_client_config" "current" {}

resource "azurerm_resource_group" "test" {
  name     = "acceptanceTestResourceGroup1"
  location = "West US"
}

resource "azurerm_sql_server" "test" {
  name                        = "mysqlserver"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  location                   = "${azurerm_resource_group.test.location}"
  version                    = "12.0"
  administrator_login        = "4dm1n157r470r"
  administrator_login_password = "4-v3ry-53cr37-p455w0rd"
}

resource "azurerm_sql_active_directory_administrator" "test" {
  server_name      = "${azurerm_sql_server.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  login            = "sqladmin"
  tenant_id        = "${data.azurerm_client_config.current.tenant_id}"
  object_id         = "${data.azurerm_client_config.current.service_principal_object_id}"
}
```

» **Argument Reference**

The following arguments are supported:

- **server_name** - (Required) The name of the SQL Server on which to set the administrator. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group for the SQL server. Changing this forces a new resource to be created.
- **login** - (Required) The login name of the principal to set as the server administrator
- **object_id** - (Required) The ID of the principal to set as the server administrator

- `tenant_id` - (Required) The Azure Tenant ID

» Attributes Reference

The following attributes are exported:

- `id` - The SQL Active Directory Administrator ID.

» Import

A SQL Active Directory Administrator can be imported using the `resource id`, e.g.

```
terraform import azurerm_sql_active_directory_administrator.administrator /subscriptions/00000000-0000-0000-0000-000000000000/
```

» `azurerm_sql_elasticpool`

Allows you to manage an Azure SQL Elastic Pool.

NOTE: - This version of the `Elasticpool` resource is being **deprecated** and should no longer be used. Please use the `azurerm_mssql_elasticpool` version instead.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "my-resource-group"
  location = "West US"
}

resource "azurerm_sql_server" "test" {
  name                        = "my-sql-server"                                # NOTE: needs to
  resource_group_name        = "${azurerm_resource_group.test.name}"
  location                   = "${azurerm_resource_group.test.location}"
  version                    = "12.0"
  administrator_login        = "4dm1n157r470r"
  administrator_login_password = "4-v3ry-53cr37-p455w0rd"
}

resource "azurerm_sql_elasticpool" "test" {
  name                = "test"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location            = "${azurerm_resource_group.test.location}"
}
```

```

server_name      = "${azurerm_sql_server.test.name}"
edition          = "Basic"
dtu              = 50
db_dtu_min       = 0
db_dtu_max       = 5
pool_size        = 5000
}

```

NOTE on azurerm_sql_elasticpool: - The values of `edition`, `dtu`, and `pool_size` must be consistent with the Azure SQL Database Service Tiers. Any inconsistent argument configuration will be rejected.

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the elastic pool. This needs to be globally unique. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the elastic pool. This must be the same as the resource group of the underlying SQL server.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **server_name** - (Required) The name of the SQL Server on which to create the elastic pool. Changing this forces a new resource to be created.
- **edition** - (Required) The edition of the elastic pool to be created. Valid values are `Basic`, `Standard`, and `Premium`. Refer to Azure SQL Database Service Tiers for details. Changing this forces a new resource to be created.
- **dtu** - (Required) The total shared DTU for the elastic pool. Valid values depend on the `edition` which has been defined. Refer to Azure SQL Database Service Tiers for valid combinations.
- **db_dtu_min** - (Optional) The minimum DTU which will be guaranteed to all databases in the elastic pool to be created.
- **db_dtu_max** - (Optional) The maximum DTU which will be guaranteed to all databases in the elastic pool to be created.
- **pool_size** - (Optional) The maximum size in MB that all databases in the elastic pool can grow to. The maximum size must be consistent with combination of `edition` and `dtu` and the limits documented in Azure SQL Database Service Tiers. If not defined when creating an elastic pool, the value is set to the size implied by `edition` and `dtu`.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- `id` - The SQL Elastic Pool ID.
- `creation_date` - The creation date of the SQL Elastic Pool.

» `azurerm_mssql_elasticpool`

Allows you to manage an Azure SQL Elastic Pool via the 2017-10-01-preview API which allows for vCore and DTU based configurations.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "my-resource-group"
  location = "westeurope"
}

resource "azurerm_sql_server" "test" {
  name                        = "my-sql-server"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  location                   = "${azurerm_resource_group.test.location}"
  version                    = "12.0"
  administrator_login        = "4dm1n157r470r"
  administrator_login_password = "4-v3ry-53cr37-p455w0rd"
}

resource "azurerm_mssql_elasticpool" "test" {
  name                = "test-epool"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location            = "${azurerm_resource_group.test.location}"
  server_name         = "${azurerm_sql_server.test.name}"
  max_size_gb        = 756

  sku {
    name      = "GP_Gen5"
    tier       = "GeneralPurpose"
    family    = "Gen5"
    capacity = 4
  }

  per_database_settings {
```

```

    min_capacity = 0.25
    max_capacity = 4
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the elastic pool. This needs to be globally unique. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the elastic pool. This must be the same as the resource group of the underlying SQL server.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **server_name** - (Required) The name of the SQL Server on which to create the elastic pool. Changing this forces a new resource to be created.
- **sku** - (Required) A **sku** block as defined below.
- **per_database_settings** - (Required) A **per_database_settings** block as defined below.
- **max_size_gb** - (Optional) The max data size of the elastic pool in gigabytes. Conflicts with **max_size_bytes**.
- **max_size_bytes** - (Optional) The max data size of the elastic pool in bytes. Conflicts with **max_size_gb**.
- **tags** - (Optional) A mapping of tags to assign to the resource.
- **zone_redundant** - (Optional) Whether or not this elastic pool is zone redundant. **tier** needs to be **Premium** for DTU based or **BusinessCritical** for vCore based **sku**. Defaults to **false**.

sku supports the following:

- **name** - (Required) Specifies the SKU Name for this Elasticpool. The name of the SKU, will be either vCore based **tier** + **family** pattern (e.g. GP_Gen4, BC_Gen5) or the DTU based **BasicPool**, **StandardPool**, or **PremiumPool** pattern.
- **capacity** - (Required) The scale up/out capacity, representing server's compute units. For more information see the documentation for your Elasticpool configuration: vCore-based or DTU-based.

- **tier** - (Required) The tier of the particular SKU. Possible values are **GeneralPurpose**, **BusinessCritical**, **Basic**, **Standard**, or **Premium**. For more information see the documentation for your Elasticpool configuration: vCore-based or DTU-based.
- **family** - (Optional) The family of hardware **Gen4** or **Gen5**.

per_database_settings supports the following:

- **min_capacity** - (Required) The minimum capacity all databases are guaranteed.
 - **max_capacity** - (Required) The maximum capacity any one database can consume.
-

» Attributes Reference

The following attributes are exported:

- **id** - The MsSQL Elastic Pool ID.

» Import

SQL Elastic Pool can be imported using the **resource id**, e.g.

```
terraform import azurerm_mssql_elasticpool.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_sql_firewall_rule

Allows you to manage an Azure SQL Firewall Rule

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "acceptanceTestResourceGroup1"
  location = "West US"
}

resource "azurerm_sql_server" "test" {
  name                        = "mysqlserver"
  resource_group_name        = "${azurerm_resource_group.test.name}"
}
```

```

location                = "West US"
version                 = "12.0"
administrator_login     = "4dm1n157r470r"
administrator_login_password = "4-v3ry-53cr37-p455w0rd"
}

resource "azurerm_sql_firewall_rule" "test" {
  name                = "FirewallRule1"
  resource_group_name = "${azurerm_resource_group.test.name}"
  server_name         = "${azurerm_sql_server.test.name}"
  start_ip_address    = "10.0.17.62"
  end_ip_address      = "10.0.17.62"
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the firewall rule.
- **resource_group_name** - (Required) The name of the resource group in which to create the sql server.
- **server_name** - (Required) The name of the SQL Server on which to create the Firewall Rule.
- **start_ip_address** - (Required) The starting IP address to allow through the firewall for this rule.
- **end_ip_address** - (Required) The ending IP address to allow through the firewall for this rule.

NOTE: The Azure feature **Allow access to Azure services** can be enabled by setting **start_ip_address** and **end_ip_address** to 0.0.0.0 which (is documented in the Azure API Docs).

» Attributes Reference

The following attributes are exported:

- **id** - The SQL Firewall Rule ID.

» Import

SQL Firewall Rules can be imported using the **resource id**, e.g.

```
terraform import azurerm_sql_firewall_rule.rule1 /subscriptions/00000000-0000-0000-0000-0000
```

» azurerm_sql_server

Manages a SQL Azure Database Server.

Note: All arguments including the administrator login and password will be stored in the raw state as plain-text. Read more about sensitive data in state.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name       = "database-rg"
  location   = "West US"
}

resource "azurerm_sql_server" "test" {
  name                               = "mysqlserver"
  resource_group_name                = "${azurerm_resource_group.test.name}"
  location                           = "${azurerm_resource_group.test.location}"
  version                             = "12.0"
  administrator_login                = "mradministrator"
  administrator_login_password      = "thisIsDog11"

  tags = {
    environment = "production"
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the SQL Server. This needs to be globally unique within Azure.
- **resource_group_name** - (Required) The name of the resource group in which to create the SQL Server.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **version** - (Required) The version for the new server. Valid values are: 2.0 (for v11 server) and 12.0 (for v12 server).
- **administrator_login** - (Required) The administrator login name for the new server. Changing this forces a new resource to be created.

- `administrator_login_password` - (Required) The password associated with the `administrator_login` user. Needs to comply with Azure's Password Policy
- `tags` - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- `id` - The SQL Server ID.
- `fully_qualified_domain_name` - The fully qualified domain name of the Azure SQL Server (e.g. `myServerName.database.windows.net`)

» Import

SQL Servers can be imported using the `resource id`, e.g.

```
terraform import azurerm_sql_server.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_sql_virtual_network_rule`

Allows you to add, update, or remove an Azure SQL server to a subnet of a virtual network.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "example-sql-server-vnet-rule"
  location = "West US"
}

resource "azurerm_virtual_network" "vnet" {
  name            = "example-vnet"
  address_space   = ["10.7.29.0/29"]
  location        = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
}

resource "azurerm_subnet" "subnet" {
  name                 = "example-subnet"
  resource_group_name = "${azurerm_resource_group.example.name}"
  virtual_network_name = "${azurerm_virtual_network.vnet.name}"
}
```

```

    address_prefix      = "10.7.29.0/29"
    service_endpoints    = ["Microsoft.Sql"]
}

resource "azurerm_sql_server" "sqlserver" {
  name                        = "uniqueazuresqlserver"
  resource_group_name        = "${azurerm_resource_group.example.name}"
  location                   = "${azurerm_resource_group.example.location}"
  version                    = "12.0"
  administrator_login        = "4dm1n157r470r"
  administrator_login_password = "4-v3ry-53cr37-p455w0rd"
}

resource "azurerm_sql_virtual_network_rule" "sqlvnetrule" {
  name                = "sql-vnet-rule"
  resource_group_name = "${azurerm_resource_group.example.name}"
  server_name         = "${azurerm_sql_server.sqlserver.name}"
  subnet_id           = "${azurerm_subnet.subnet.id}"
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the SQL virtual network rule. Changing this forces a new resource to be created. Cannot be empty and must only contain alphanumeric characters and hyphens. Cannot start with a number, and cannot start or end with a hyphen.

NOTE: **name** must be between 1-128 characters long and must satisfy all of the requirements below: 1. Contains only alphanumeric and hyphen characters 2. Cannot start with a number or hyphen 3. Cannot end with a hyphen

- **resource_group_name** - (Required) The name of the resource group where the SQL server resides. Changing this forces a new resource to be created.
- **server_name** - (Required) The name of the SQL Server to which this SQL virtual network rule will be applied to. Changing this forces a new resource to be created.
- **subnet_id** - (Required) The ID of the subnet that the SQL server will be connected to.
- **ignore_missing_vnet_service_endpoint** - (Optional) Create the virtual network rule before the subnet has the virtual network service endpoint enabled. The default value is false.

NOTE: If `ignore_missing_vnet_service_endpoint` is false, and the target subnet does not contain the `Microsoft.SQL` endpoint in the `service_endpoints` array, the deployment will fail when it tries to create the SQL virtual network rule.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the SQL virtual network rule.

» Import

SQL Virtual Network Rules can be imported using the `resource id`, e.g.

```
terraform import azurerm_sql_virtual_network_rule.rule1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_databricks_workspace`

Manages a Databricks Workspace

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "example-resources"
  location = "West US"
}

resource "azurerm_databricks_workspace" "test" {
  name                       = "databricks-test"
  resource_group_name       = "${azurerm_resource_group.test.name}"
  location                  = "${azurerm_resource_group.test.location}"
  sku                       = "standard"

  tags = {
    Environment = "Production"
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Databricks Workspace resource. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the Resource Group in which the Databricks Workspace should exist. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource has to be created. Changing this forces a new resource to be created.
- **sku** - (Required) The **sku** to use for the Databricks Workspace. Possible values are **standard** or **premium**. Changing this forces a new resource to be created.
- **managed_resource_group_name** - (Optional) The name of the resource group where Azure should place the managed Databricks resources. Changing this forces a new resource to be created.

NOTE Azure requires that this Resource Group does not exist in this Subscription (and that the Azure API creates it) - otherwise the deployment will fail.

- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Databricks Workspace.
- **managed_resource_group_id** - The ID of the Managed Resource Group created by the Databricks Workspace.

» Import

Databrick Workspaces can be imported using the **resource id**, e.g.

```
terraform import azurerm_databrick_workspace.workspace1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm__data__factory

Manage an Azure Data Factory (Version 2).

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name      = "example"
  location  = "northeurope"
}

resource "azurerm_data_factory" "example" {
  name                = "example"
  location             = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Data Factory. Changing this forces a new resource to be created. Must be globally unique. See the Microsoft documentation for all restrictions.
- **resource_group_name** - (Required) The name of the resource group in which to create the Data Factory.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **identity** - (Optional) An `identity` block as defined below.
- **github_configuration** - (Optional) A `github_configuration` block as defined below.
- **vsts_configuration** - (Optional) A `vsts_configuration` block as defined below.
- **tags** - (Optional) A mapping of tags to assign to the resource.

A `identity` block supports the following:

- **type** - (Required) Specifies the identity type of the Data Factory. At this time the only allowed value is `SystemAssigned`.

A `github_configuration` block supports the following:

- **account_name** - (Required) Specifies the GitHub account name.
- **branch_name** - (Required) Specifies the branch of the repository to get code from.

- **git_url** - (Required) Specifies the GitHub Enterprise host name. For example: `https://github.mydomain.com`. Use `https://github.com` for open source repositories.
- **repository_name** - (Required) Specifies the name of the git repository.
- **root_folder** - (Required) Specifies the root folder within the repository. Set to `/` for the top level.

Note: You must log in to the Data Factory management UI to complete the authentication to the GitHub repository.

A `vsts_configuration` block supports the following:

- **account_name** - (Required) Specifies the VSTS account name.
- **branch_name** - (Required) Specifies the branch of the repository to get code from.
- **project_name** - (Required) Specifies the name of the VSTS project.
- **repository_name** - (Required) Specifies the name of the git repository.
- **root_folder** - (Required) Specifies the root folder within the repository. Set to `/` for the top level.
- **tenant_id** - (Required) Specifies the Tenant ID associated with the VSTS account.

» Attributes Reference

The following attributes are exported:

- **id** - The Data Factory ID.

» Import

Data Factory can be imported using the `resource id`, e.g.

```
terraform import azurerm_data_factory.example /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm__data_factory__dataset__mysql`

Manage a MySQL Dataset inside a Azure Data Factory.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name      = "example"
  location  = "northeurope"
}

resource "azurerm_data_factory" "example" {
  name                = "example"
  location            = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
}

resource "azurerm_data_factory_linked_service_mysql" "example" {
  name                = "example"
  resource_group_name = "${azurerm_resource_group.example.name}"
  data_factory_name   = "${azurerm_data_factory.example.name}"
  connection_string   = "Server=test;Port=3306;Database=test;User=test;SSLMode=1;UseSystemTr"
}

resource "azurerm_data_factory_dataset_mysql" "example" {
  name                = "example"
  resource_group_name = "${azurerm_resource_group.example.name}"
  data_factory_name   = "${azurerm_data_factory.example.name}"
  linked_service_name = "${azurerm_data_factory_linked_service_mysql.test.name}"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Data Factory Dataset MySQL. Changing this forces a new resource to be created. Must be globally unique. See the Microsoft documentation for all restrictions.
- **resource_group_name** - (Required) The name of the resource group in which to create the Data Factory Dataset MySQL. Changing this forces a new resource
- **data_factory_name** - (Required) The Data Factory name in which to associate the Dataset with. Changing this forces a new resource.
- **linked_service_name** - (Required) The Data Factory Linked Service name in which to associate the Dataset with.
- **table_name** - (Optional) The table name of the Data Factory Dataset MySQL.

- **folder** - (Optional) The folder that this Dataset is in. If not specified, the Dataset will appear at the root level.
- **schema_column** - (Optional) A **schema_column** block as defined below.
- **description** - (Optional) The description for the Data Factory Dataset MySQL.
- **annotations** - (Optional) List of tags that can be used for describing the Data Factory Dataset MySQL.
- **parameters** - (Optional) A map of parameters to associate with the Data Factory Dataset MySQL.
- **additional_properties** - (Optional) A map of additional properties to associate with the Data Factory Dataset MySQL.

A **schema_column** block supports the following:

- **name** - (Required) The name of the column.
- **type** - (Optional) Type of the column. Valid values are **Byte**, **Byte[]**, **Boolean**, **Date**, **DateTime**, **DateTimeOffset**, **Decimal**, **Double**, **Guid**, **Int16**, **Int32**, **Int64**, **Single**, **String**, **TimeSpan**. Please note these values are case sensitive.
- **description** - (Optional) The description of the column.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Data Factory Dataset.

» Import

Data Factory Dataset MySQL can be imported using the **resource id**, e.g.

```
terraform import azurerm_data_factory_dataset_mysql.example /subscriptions/00000000-0000-0000-0000-000000000000
```

» **azurerm_data_factory_dataset_postgresql**

Manage a PostgreSQL Dataset inside a Azure Data Factory.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name      = "example"
  location  = "northeurope"
}

resource "azurerm_data_factory" "example" {
  name                = "example"
  location             = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
}

resource "azurerm_data_factory_linked_service_postgresql" "example" {
  name                = "example"
  resource_group_name = "${azurerm_resource_group.example.name}"
  data_factory_name   = "${azurerm_data_factory.example.name}"
  connection_string   = "Host=example;Port=5432;Database=example;UID=example;EncryptionMethod=0"
}

resource "azurerm_data_factory_dataset_postgresql" "example" {
  name                = "example"
  resource_group_name = "${azurerm_resource_group.example.name}"
  data_factory_name   = "${azurerm_data_factory.example.name}"
  linked_service_name = "${azurerm_data_factory_linked_service_postgresql.test.name}"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Data Factory Dataset PostgreSQL. Changing this forces a new resource to be created. Must be globally unique. See the Microsoft documentation for all restrictions.
- **resource_group_name** - (Required) The name of the resource group in which to create the Data Factory Dataset PostgreSQL. Changing this forces a new resource
- **data_factory_name** - (Required) The Data Factory name in which to associate the Dataset with. Changing this forces a new resource.
- **linked_service_name** - (Required) The Data Factory Linked Service name in which to associate the Dataset with.
- **table_name** - (Optional) The table name of the Data Factory Dataset PostgreSQL.

- **folder** - (Optional) The folder that this Dataset is in. If not specified, the Dataset will appear at the root level.
- **schema_column** - (Optional) A **schema_column** block as defined below.
- **description** - (Optional) The description for the Data Factory Dataset PostgreSQL.
- **annotations** - (Optional) List of tags that can be used for describing the Data Factory Dataset PostgreSQL.
- **parameters** - (Optional) A map of parameters to associate with the Data Factory Dataset PostgreSQL.
- **additional_properties** - (Optional) A map of additional properties to associate with the Data Factory Dataset PostgreSQL.

A **schema_column** block supports the following:

- **name** - (Required) The name of the column.
- **type** - (Optional) Type of the column. Valid values are `Byte`, `Byte[]`, `Boolean`, `Date`, `DateTime`, `DateTimeOffset`, `Decimal`, `Double`, `Guid`, `Int16`, `Int32`, `Int64`, `Single`, `String`, `TimeSpan`. Please note these values are case sensitive.
- **description** - (Optional) The description of the column.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Data Factory Dataset.

» Import

Data Factory Dataset PostgreSQL can be imported using the **resource id**, e.g.

```
terraform import azurerm_data_factory_dataset_postgresql.example /subscriptions/00000000-0000-0000-0000-000000000000
```

» **azurerm__data__factory__dataset__sql__server**

Manage a SQL Server Table Dataset inside a Azure Data Factory.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name      = "example"
  location  = "northeurope"
}

resource "azurerm_data_factory" "example" {
  name                = "example"
  location             = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
}

resource "azurerm_data_factory_linked_service_sql_server" "example" {
  name                = "example"
  resource_group_name = "${azurerm_resource_group.example.name}"
  data_factory_name   = "${azurerm_data_factory.example.name}"
  connection_string   = "Integrated Security=False;Data Source=test;Initial Catalog=test;User=sa;Password=test;"
}

resource "azurerm_data_factory_dataset_sql_server_table" "example" {
  name                = "example"
  resource_group_name = "${azurerm_resource_group.example.name}"
  data_factory_name   = "${azurerm_data_factory.example.name}"
  linked_service_name = "${azurerm_data_factory_linked_service_sql_server.test.name}"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Data Factory Dataset SQL Server Table. Changing this forces a new resource to be created. Must be globally unique. See the Microsoft documentation for all restrictions.
- **resource_group_name** - (Required) The name of the resource group in which to create the Data Factory Dataset SQL Server Table. Changing this forces a new resource
- **data_factory_name** - (Required) The Data Factory name in which to associate the Dataset with. Changing this forces a new resource.
- **linked_service_name** - (Required) The Data Factory Linked Service name in which to associate the Dataset with.
- **table_name** - (Optional) The table name of the Data Factory Dataset SQL Server Table.

- **folder** - (Optional) The folder that this Dataset is in. If not specified, the Dataset will appear at the root level.
- **schema_column** - (Optional) A **schema_column** block as defined below.
- **description** - (Optional) The description for the Data Factory Dataset SQL Server Table.
- **annotations** - (Optional) List of tags that can be used for describing the Data Factory Dataset SQL Server Table.
- **parameters** - (Optional) A map of parameters to associate with the Data Factory Dataset SQL Server Table.
- **additional_properties** - (Optional) A map of additional properties to associate with the Data Factory Dataset SQL Server Table.

A **schema_column** block supports the following:

- **name** - (Required) The name of the column.
- **type** - (Optional) Type of the column. Valid values are `Byte`, `Byte[]`, `Boolean`, `Date`, `DateTime`, `DateTimeOffset`, `Decimal`, `Double`, `Guid`, `Int16`, `Int32`, `Int64`, `Single`, `String`, `TimeSpan`. Please note these values are case sensitive.
- **description** - (Optional) The description of the column.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Data Factory Dataset.

» Import

Data Factory Dataset SQL Server Table can be imported using the **resource id**, e.g.

```
terraform import azurerm_data_factory_dataset_sql_server_table.example /subscriptions/000000
```

» **azurerm_data_factory_pipeline**

Manage a Pipeline inside a Azure Data Factory.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name      = "example"
  location  = "northeurope"
}

resource "azurerm_data_factory" "example" {
  name                = "example"
  location            = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
}

resource "azurerm_data_factory_pipeline" "example" {
  name                = "example"
  resource_group_name = "${azurerm_resource_group.example.name}"
  data_factory_name   = "${azurerm_data_factory.example.name}"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Data Factory Pipeline. Changing this forces a new resource to be created. Must be globally unique. See the Microsoft documentation for all restrictions.
- **resource_group_name** - (Required) The name of the resource group in which to create the Data Factory Pipeline. Changing this forces a new resource
- **data_factory_name** - (Required) The Data Factory name in which to associate the Pipeline with. Changing this forces a new resource.
- **description** - (Optional) The description for the Data Factory Pipeline.
- **annotations** - (Optional) List of tags that can be used for describing the Data Factory Pipeline.
- **parameters** - (Optional) A map of parameters to associate with the Data Factory Pipeline.
- **variables** - (Optional) A map of variables to associate with the Data Factory Pipeline.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the Data Factory Pipeline.

» Import

Data Factory Pipeline can be imported using the `resource id`, e.g.

```
terraform import azurerm_data_factory_pipeline.example /subscriptions/00000000-0000-0000-0000-000000000000/resourceGroups/example/dataFactories/example/pipelines/example
```

» `azurerm_data_factory_linked_service_mysql`

Manage a Linked Service (connection) between MySQL and Azure Data Factory.

Note: All arguments including the `connection_string` will be stored in the raw state as plain-text. Read more about sensitive data in state.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "example"
  location = "northeurope"
}

resource "azurerm_data_factory" "example" {
  name                = "example"
  location             = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
}

resource "azurerm_data_factory_linked_service_mysql" "example" {
  name                = "example"
  resource_group_name = "${azurerm_resource_group.test.name}"
  data_factory_name   = "${azurerm_data_factory.test.name}"
  connection_string    = "Server=test;Port=3306;Database=test;User=test;SSLMode=1;UseSystemTables=1"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Data Factory Linked Service MySQL. Changing this forces a new resource to be created. Must be globally unique. See the Microsoft documentation for all restrictions.
- **resource_group_name** - (Required) The name of the resource group in which to create the Data Factory Linked Service MySQL. Changing this forces a new resource
- **data_factory_name** - (Required) The Data Factory name in which to associate the Linked Service with. Changing this forces a new resource.
- **connection_string** - (Required) The connection string in which to authenticate with MySQL.
- **description** - (Optional) The description for the Data Factory Linked Service MySQL.
- **integration_runtime_name** - (Optional) The integration runtime reference to associate with the Data Factory Linked Service MySQL.
- **annotations** - (Optional) List of tags that can be used for describing the Data Factory Linked Service MySQL.
- **parameters** - (Optional) A map of parameters to associate with the Data Factory Linked Service MySQL.
- **additional_properties** - (Optional) A map of additional properties to associate with the Data Factory Linked Service MySQL.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Data Factory Linked Service.

» Import

Data Factory Linked Service MySQL can be imported using the **resource id**, e.g.

```
terraform import azurerm_data_factory_linked_service_mysql.example /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_data_factory_linked_service_postgresql

Manage a Linked Service (connection) between PostgreSQL and Azure Data Factory.

Note: All arguments including the `connection_string` will be stored in the raw state as plain-text. Read more about sensitive data in state.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name      = "example"
  location  = "northeurope"
}

resource "azurerm_data_factory" "example" {
  name                = "example"
  location             = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
}

resource "azurerm_data_factory_linked_service_postgresql" "example" {
  name                = "example"
  resource_group_name = "${azurerm_resource_group.test.name}"
  data_factory_name   = "${azurerm_data_factory.test.name}"
  connection_string    = "Host=example;Port=5432;Database=example;UID=example;EncryptionMethod="
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Data Factory Linked Service PostgreSQL. Changing this forces a new resource to be created. Must be globally unique. See the Microsoft documentation for all restrictions.
- **resource_group_name** - (Required) The name of the resource group in which to create the Data Factory Linked Service PostgreSQL. Changing this forces a new resource
- **data_factory_name** - (Required) The Data Factory name in which to associate the Linked Service with. Changing this forces a new resource.
- **connection_string** - (Required) The connection string in which to authenticate with PostgreSQL.
- **description** - (Optional) The description for the Data Factory Linked Service PostgreSQL.
- **integration_runtime_name** - (Optional) The integration runtime reference to associate with the Data Factory Linked Service PostgreSQL.

- **annotations** - (Optional) List of tags that can be used for describing the Data Factory Linked Service PostgreSQL.
- **parameters** - (Optional) A map of parameters to associate with the Data Factory Linked Service PostgreSQL.
- **additional_properties** - (Optional) A map of additional properties to associate with the Data Factory Linked Service PostgreSQL.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Data Factory Linked Service.

» Import

Data Factory Linked Service PostgreSQL can be imported using the **resource id**, e.g.

```
terraform import azurerm_data_factory_linked_service_postgresql.example /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_data_factory_linked_service_sql_server

Manage a Linked Service (connection) between a SQL Server and Azure Data Factory.

Note: All arguments including the client secret will be stored in the raw state as plain-text. Read more about sensitive data in state.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "example"
  location = "northeurope"
}

resource "azurerm_data_factory" "example" {
  name                = "example"
  location            = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
}

resource "azurerm_data_factory_linked_service_sql_server" "example" {
```

```

name                = "example"
resource_group_name = "${azurerm_resource_group.test.name}"
data_factory_name   = "${azurerm_data_factory.test.name}"
connection_string   = "Integrated Security=False;Data Source=test;Initial Catalog=test;Use
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Data Factory Linked Service SQL Server. Changing this forces a new resource to be created. Must be globally unique. See the Microsoft documentation for all restrictions.
- **resource_group_name** - (Required) The name of the resource group in which to create the Data Factory Linked Service SQL Server. Changing this forces a new resource
- **data_factory_name** - (Required) The Data Factory name in which to associate the Linked Service with. Changing this forces a new resource.
- **connection_string** - (Required) The connection string in which to authenticate with the SQL Server.
- **description** - (Optional) The description for the Data Factory Linked Service SQL Server.
- **integration_runtime_name** - (Optional) The integration runtime reference to associate with the Data Factory Linked Service SQL Server.
- **annotations** - (Optional) List of tags that can be used for describing the Data Factory Linked Service SQL Server.
- **parameters** - (Optional) A map of parameters to associate with the Data Factory Linked Service SQL Server.
- **additional_properties** - (Optional) A map of additional properties to associate with the Data Factory Linked Service SQL Server.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Data Factory Linked Service.

» Import

Data Factory Linked Service SQL Server can be imported using the `resource` id, e.g.

```
terraform import azurerm_data_factory_linked_service_sql_server.example /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_data_lake_analytics_account

Manage an Azure Data Lake Analytics Account.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "tfex-datalake-account"
  location = "northeurope"
}

resource "azurerm_data_lake_store" "example" {
  name                = "tfexdatalakestore"
  resource_group_name = "${azurerm_resource_group.example.name}"
  location            = "${azurerm_resource_group.example.location}"
}

resource "azurerm_data_lake_analytics_account" "example" {
  name                = "tfexdatalakeaccount"
  resource_group_name = "${azurerm_resource_group.example.name}"
  location            = "${azurerm_resource_group.example.location}"

  default_store_account_name = "${azurerm_data_lake_store.example.name}"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Data Lake Analytics Account. Changing this forces a new resource to be created. Has to be between 3 to 24 characters.
- **resource_group_name** - (Required) The name of the resource group in which to create the Data Lake Analytics Account.

- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **default_store_account_name** - (Required) Specifies the data lake store to use by default. Changing this forces a new resource to be created.
- **tier** - (Optional) The monthly commitment tier for Data Lake Analytics Account. Accepted values are `Consumption`, `Commitment_100000AUHours`, `Commitment_10000AUHours`, `Commitment_1000AUHours`, `Commitment_100AUHours`, `Commitment_500000AUHours`, `Commitment_50000AUHours`, `Commitment_5000AUHours`, or `Commitment_500AUHours`.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- **id** - The Date Lake Store ID.

» Import

Date Lake Analytics Account can be imported using the **resource id**, e.g.

```
terraform import azurerm_data_lake_analytics_account.test /subscriptions/00000000-0000-0000-
```

» azurerm__data_lake_analytics__firewall__rule

Manage a Azure Data Lake Analytics Firewall Rule.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name      = "tfex_datalake_fw_rule"
  location  = "northeurope"
}

resource "azurerm_data_lake_store" "example" {
  name                = "tfexdatalakestore"
  resource_group_name = "${azurerm_resource_group.example.name}"
  location            = "${azurerm_resource_group.example.location}"
}

resource "azurerm_data_lake_analytics_account" "example" {
```

```

    name                = "tfexdatalakeaccount"
    resource_group_name = "${azurerm_resource_group.example.name}"
    location             = "${azurerm_resource_group.example.location}"

    default_store_account_name = "${azurerm_data_lake_store.example.name}"
}

resource "azurerm_data_lake_analytics_firewall_rule" "example" {
    name                = "office-ip-range"
    account_name        = "${azurerm_data_lake_analytics.example.name}"
    resource_group_name = "${azurerm_resource_group.example.name}"
    start_ip_address    = "1.2.3.4"
    end_ip_address      = "2.3.4.5"
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Data Lake Analytics. Changing this forces a new resource to be created. Has to be between 3 to 24 characters.
- **resource_group_name** - (Required) The name of the resource group in which to create the Data Lake Analytics.
- **account_name** - (Required) Specifies the name of the Data Lake Analytics for which the Firewall Rule should take effect.
- **start_ip_address** - (Required) The Start IP address for the firewall rule.
- **end_ip_address** - (Required) The End IP Address for the firewall rule.

» Attributes Reference

The following attributes are exported:

- **id** - The Date Lake Store Firewall Rule ID.

» Import

Date Lake Store Firewall Rules can be imported using the **resource id**, e.g.

```
terraform import azurerm_data_lake_analytics_firewall_rule.rule1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» **azurerm_data_lake_store**

Manage an Azure Data Lake Store.

» **Example Usage**

```
resource "azurerm_resource_group" "example" {
  name      = "example"
  location  = "northeurope"
}

resource "azurerm_data_lake_store" "example" {
  name                        = "consumptiondatalake"
  resource_group_name        = "${azurerm_resource_group.example.name}"
  location                   = "${azurerm_resource_group.example.location}"
  encryption_state           = "Enabled"
  encryption_type             = "ServiceManaged"
}
```

» **Argument Reference**

The following arguments are supported:

- **name** - (Required) Specifies the name of the Data Lake Store. Changing this forces a new resource to be created. Has to be between 3 to 24 characters.
- **resource_group_name** - (Required) The name of the resource group in which to create the Data Lake Store.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **tier** - (Optional) The monthly commitment tier for Data Lake Store. Accepted values are `Consumption`, `Commitment_1TB`, `Commitment_10TB`, `Commitment_100TB`, `Commitment_500TB`, `Commitment_1PB` or `Commitment_5PB`.
- **encryption_state** - (Optional) Is Encryption enabled on this Data Lake Store Account? Possible values are `Enabled` or `Disabled`. Defaults to `Enabled`.
- **encryption_type** - (Optional) The Encryption Type used for this Data Lake Store Account. Currently can be set to `ServiceManaged` when `encryption_state` is `Enabled` - and must be a blank string when it's `Disabled`.

NOTE: Support for User Managed encryption will be supported in the future once a bug in the API is fixed.

- `firewall_allow_azure_ips` - are Azure Service IP's allowed through the firewall? Possible values are `Enabled` and `Disabled`. Defaults to `Enabled`.
- `firewall_state` - the state of the Firewall. Possible values are `Enabled` and `Disabled`. Defaults to `Enabled`.
- `tags` - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- `id` - The Data Lake Store ID.
- `endpoint` - The Endpoint for the Data Lake Store.

» Import

Date Lake Store can be imported using the `resource id`, e.g.

```
terraform import azurerm_data_lake_store.test /subscriptions/00000000-0000-0000-0000-00000000
```

» `azurerm_data_lake_store_firewall_rule`

Manage a Azure Data Lake Store Firewall Rule.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "example"
  location = "northeurope"
}

resource "azurerm_data_lake_store" "example" {
  name                = "consumptiondatalake"
  resource_group_name = "${azurerm_resource_group.example.name}"
  location            = "${azurerm_resource_group.example.location}"
}

resource "azurerm_data_lake_store_firewall_rule" "example" {
  name = "office-ip-range"
```

```

account_name      = "${azurerm_data_lake_store.example.name}"
resource_group_name = "${azurerm_resource_group.example.name}"
start_ip_address  = "1.2.3.4"
end_ip_address    = "2.3.4.5"
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Data Lake Store. Changing this forces a new resource to be created. Has to be between 3 to 24 characters.
- **resource_group_name** - (Required) The name of the resource group in which to create the Data Lake Store.
- **account_name** - (Required) Specifies the name of the Data Lake Store for which the Firewall Rule should take effect.
- **start_ip_address** - (Required) The Start IP address for the firewall rule.
- **end_ip_address** - (Required) The End IP Address for the firewall rule.

» Attributes Reference

The following attributes are exported:

- **id** - The Date Lake Store Firewall Rule ID.

» Import

Date Lake Store Firewall Rules can be imported using the **resource id**, e.g.

```
terraform import azurerm_data_lake_store_firewall_rule.rule1 /subscriptions/00000000-0000-00
```

» azurerm__devspace__controller

Manages a DevSpace Controller.

» Example Usage

```

resource "azurerm_resource_group" "test" {
  name = "acctestRG1"
}

```

```

    location = "westeurope"
}

resource "azurerm_kubernetes_cluster" "test" {
  name           = "acctestaks1"
  location       = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  dns_prefix     = "acctestaks1"

  agent_pool_profile {
    name     = "default"
    count    = "1"
    vm_size  = "Standard_DS2_v2"
  }

  service_principal {
    client_id     = "000000000-0000-0000-0000-000000000000"
    client_secret = "00000000000000000000000000000000"
  }
}

resource "azurerm_devspace_controller" "test" {
  name           = "acctestdsc1"
  location       = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  sku {
    name = "S1"
    tier  = "Standard"
  }

  host_suffix           = "suffix"
  target_container_host_resource_id = "${azurerm_kubernetes_cluster.test.id}"
  target_container_host_credentials_base64 = "${base64encode(azurerm_kubernetes_cluster.test...)}"

  tags = {
    Environment = "Testing"
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the DevSpace Controller. Chang-

ing this forces a new resource to be created.

- **resource_group_name** - (Required) The name of the resource group under which the DevSpace Controller resource has to be created. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported location where the DevSpace Controller should exist. Changing this forces a new resource to be created.
- **sku** - (Required) A **sku** block as documented below. Changing this forces a new resource to be created.
- **host_suffix** - (Required) The host suffix for the DevSpace Controller. Changing this forces a new resource to be created.
- **target_container_host_resource_id** - (Required) The resource id of Azure Kubernetes Service cluster. Changing this forces a new resource to be created.
- **target_container_host_credentials_base64** - (Required) Base64 encoding of **kube_config_raw** of Azure Kubernetes Service cluster. Changing this forces a new resource to be created.
- **tags** - (Optional) A mapping of tags to assign to the resource.

A **sku** block supports the following:

- **name** - (Required) The name of the SKU for DevSpace Controller. Currently the only supported value is **S1**. Changing this forces a new resource to be created.
- **tier** - (Required) The tier of the SKU for DevSpace Controller. Currently the only supported value is **Standard**. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the DevSpace Controller.
- **data_plane_fqdn** - DNS name for accessing DataPlane services.

» Import

DevSpace Controller's can be imported using the **resource id**, e.g.

```
terraform import azurerm_devspace_controller.controller1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» **azurerm__dev__test__lab**

Manages a Dev Test Lab.

» **Example Usage**

```
resource "azurerm_resource_group" "test" {
  name      = "example-resources"
  location  = "West US"
}

resource "azurerm_dev_test_lab" "test" {
  name                  = "example-devtestlab"
  location              = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  tags = {
    "Sydney" = "Australia"
  }
}
```

» **Argument Reference**

The following arguments are supported:

- **name** - (Required) Specifies the name of the Dev Test Lab. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group under which the Dev Test Lab resource has to be created. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the Dev Test Lab should exist. Changing this forces a new resource to be created.
- **storage_type** - (Optional) The type of storage used by the Dev Test Lab. Possible values are **Standard** and **Premium**. Defaults to **Premium**. Changing this forces a new resource to be created.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» **Attributes Reference**

The following attributes are exported:

- `id` - The ID of the Dev Test Lab.
- `artifacts_storage_account_id` - The ID of the Storage Account used for Artifact Storage.
- `default_storage_account_id` - The ID of the Default Storage Account for this Dev Test Lab.
- `default_premium_storage_account_id` - The ID of the Default Premium Storage Account for this Dev Test Lab.
- `key_vault_id` - The ID of the Key used for this Dev Test Lab.
- `premium_data_disk_storage_account_id` - The ID of the Storage Account used for Storage of Premium Data Disk.
- `unique_identifier` - The unique immutable identifier of the Dev Test Lab.

» Import

Dev Test Labs can be imported using the `resource id`, e.g.

```
terraform import azurerm_dev_test_lab.lab1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_dev_test_linux_virtual_machine`

Manages a Linux Virtual Machine within a Dev Test Lab.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "example-resources"
  location  = "West US"
}

resource "azurerm_dev_test_lab" "test" {
  name                = "example-devtestlab"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  tags = {
    "Sydney" = "Australia"
  }
}
```

```

resource "azurerm_dev_test_virtual_network" "test" {
  name           = "example-network"
  lab_name       = "${azurerm_dev_test_lab.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  subnet {
    use_public_ip_address      = "Allow"
    use_in_virtual_machine_creation = "Allow"
  }
}

resource "azurerm_dev_test_linux_virtual_machine" "test" {
  name           = "example-vm03"
  lab_name       = "${azurerm_dev_test_lab.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location       = "${azurerm_resource_group.test.location}"
  size           = "Standard_DS2"
  username       = "exampleuser99"
  ssh_key        = "${file("~/ssh/id_rsa.pub")}"
  lab_virtual_network_id = "${azurerm_dev_test_virtual_network.test.id}"
  lab_subnet_name = "${azurerm_dev_test_virtual_network.test.subnet.0.name}"
  storage_type    = "Premium"
  notes           = "Some notes about this Virtual Machine."

  gallery_image_reference {
    offer      = "UbuntuServer"
    publisher  = "Canonical"
    sku        = "18.04-LTS"
    version    = "latest"
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Dev Test Machine. Changing this forces a new resource to be created.

NOTE: The validation requirements for the Name change based on the **os_type** used in this Virtual Machine. For a Linux VM the name must be between 1-62 characters, and for a Windows VM the name must be between 1-15 characters. It must begin and end with a letter or number, and cannot be all numbers.

- **lab_name** - (Required) Specifies the name of the Dev Test Lab in which the Virtual Machine should be created. Changing this forces a new resource

to be created.

- **resource_group_name** - (Required) The name of the resource group in which the Dev Test Lab resource exists. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the Dev Test Lab exists. Changing this forces a new resource to be created.
- **gallery_image_reference** - (Required) A **gallery_image_reference** block as defined below.
- **lab_subnet_name** - (Required) The name of a Subnet within the Dev Test Virtual Network where this machine should exist. Changing this forces a new resource to be created.
- **lab_virtual_network_id** - (Required) The ID of the Dev Test Virtual Network where this Virtual Machine should be created. Changing this forces a new resource to be created.
- **size** - (Required) The Machine Size to use for this Virtual Machine, such as **Standard_F2**. Changing this forces a new resource to be created.
- **storage_type** - (Required) The type of Storage to use on this Virtual Machine. Possible values are **Standard** and **Premium**.
- **username** - (Required) The Username associated with the local administrator on this Virtual Machine. Changing this forces a new resource to be created.

-
- **allow_claim** - (Optional) Can this Virtual Machine be claimed by users? Defaults to **true**.
 - **disallow_public_ip_address** - (Optional) Should the Virtual Machine be created without a Public IP Address? Changing this forces a new resource to be created.
 - **inbound_nat_rule** - (Optional) One or more **inbound_nat_rule** blocks as defined below. Changing this forces a new resource to be created.

NOTE: If any **inbound_nat_rule** blocks are specified then **disallow_public_ip_address** must be set to **true**.

- **notes** - (Optional) Any notes about the Virtual Machine.
- **password** - (Optional) The Password associated with the **username** used to login to this Virtual Machine. Changing this forces a new resource to be created.
- **ssh_key** - (Optional) The SSH Key associated with the **username** used to login to this Virtual Machine. Changing this forces a new resource to be created.

NOTE: One or either `password` or `ssh_key` must be specified.

- `tags` - (Optional) A mapping of tags to assign to the resource.
-

A `gallery_image_reference` block supports the following:

- `offer` - (Required) The Offer of the Gallery Image. Changing this forces a new resource to be created.
 - `publisher` - (Required) The Publisher of the Gallery Image. Changing this forces a new resource to be created.
 - `sku` - (Required) The SKU of the Gallery Image. Changing this forces a new resource to be created.
 - `version` - (Required) The Version of the Gallery Image. Changing this forces a new resource to be created.
-

A `inbound_nat_rule` block supports the following:

- `protocol` - (Required) The Protocol used for this NAT Rule. Possible values are `Tcp` and `Udp`. Changing this forces a new resource to be created.
- `backend_port` - (Required) The Backend Port associated with this NAT Rule. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the Virtual Machine.
 - `fqdn` - The FQDN of the Virtual Machine.
 - `inbound_nat_rule` - One or more `inbound_nat_rule` blocks as defined below.
 - `unique_identifier` - The unique immutable identifier of the Virtual Machine.
-

A `inbound_nat_rule` block exports the following:

- `frontend_port` - The frontend port associated with this Inbound NAT Rule.

» Import

Dev Test Linux Virtual Machines can be imported using the `resource id`, e.g.

```
terraform import azurerm_dev_test_linux_virtual_machine.machine1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm__dev__test__policy

Manages a Policy within a Dev Test Policy Set.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "example-resources"
  location = "West US"
}

resource "azurerm_dev_test_lab" "test" {
  name                = "example-devtestlab"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  tags = {
    "Sydney" = "Australia"
  }
}

resource "azurerm_dev_test_policy" "test" {
  name                = "LabVmCount"
  policy_set_name     = "default"
  lab_name            = "${azurerm_dev_test_lab.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  fact_data           = ""
  threshold           = "999"
  evaluator_type      = "MaxValuePolicy"

  tags = {
    "Acceptance" = "Test"
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Dev Test Policy. Possible values are `GalleryImage`, `LabPremiumVmCount`, `LabTargetCost`, `LabVmCount`, `LabVmSize`, `UserOwnedLabPremiumVmCount`, `UserOwnedLabVmCount` and `UserOwnedLabVmCountInSubnet`. Changing this forces a new resource to be created.
- **policy_set_name** - (Required) Specifies the name of the Policy Set within the Dev Test Lab where this policy should be created. Changing this forces a new resource to be created.
- **lab_name** - (Required) Specifies the name of the Dev Test Lab in which the Policy should be created. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the Dev Test Lab resource exists. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the Dev Test Lab exists. Changing this forces a new resource to be created.
- **description** - (Optional) A description for the Policy.
- **evaluator_type** - (Required) The Evaluation Type used for this Policy. Possible values include: `'AllowedValuesPolicy'`, `'MaxValuePolicy'`. Changing this forces a new resource to be created.
- **threshold** - (Required) The Threshold for this Policy.
- **fact_data** - (Optional) The Fact Data for this Policy.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Dev Test Policy.

» Import

Dev Test Policies can be imported using the **resource id**, e.g.

```
terraform import azurerm_dev_test_policy.policy1 /subscriptions/00000000-0000-0000-0000-0000
```

» **azurerm_dev_test_virtual_network**

Manages a Virtual Network within a Dev Test Lab.

» **Example Usage**

```
resource "azurerm_resource_group" "test" {
  name      = "example-resources"
  location  = "West US"
}

resource "azurerm_dev_test_lab" "test" {
  name                = "example-devtestlab"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  tags = {
    "Sydney" = "Australia"
  }
}

resource "azurerm_dev_test_virtual_network" "test" {
  name                = "example-network"
  lab_name            = "${azurerm_dev_test_lab.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  subnet {
    use_public_ip_address      = "Allow"
    use_in_virtual_machine_creation = "Allow"
  }
}
```

» **Argument Reference**

The following arguments are supported:

- **name** - (Required) Specifies the name of the Dev Test Virtual Network. Changing this forces a new resource to be created.
- **lab_name** - (Required) Specifies the name of the Dev Test Lab in which the Virtual Network should be created. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in

which the Dev Test Lab resource exists. Changing this forces a new resource to be created.

- **description** - (Optional) A description for the Virtual Network.
 - **subnet** - (Optional) A **subnet** block as defined below.
 - **tags** - (Optional) A mapping of tags to assign to the resource.
-

A **subnet** block supports the following:

- **use_public_ip_address** - (Required) Can Virtual Machines in this Subnet use Public IP Addresses? Possible values are **Allow**, **Default** and **Deny**.
- **use_in_virtual_machine_creation** - (Required) Can this subnet be used for creating Virtual Machines? Possible values are **Allow**, **Default** and **Deny**.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Dev Test Virtual Network.
 - **subnet** - A **subnet** block as defined below.
 - **unique_identifier** - The unique immutable identifier of the Dev Test Virtual Network.
-

A **subnet** block exports the following:

- **name** - The name of the Subnet for this Virtual Network.

» Import

Dev Test Virtual Networks can be imported using the **resource id**, e.g.

```
terraform import azurerm_dev_test_virtual_network.network1 /subscriptions/00000000-0000-0000
```

» **azurerm_dev_test_windows_virtual_machine**

Manages a Windows Virtual Machine within a Dev Test Lab.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "example-resources"
  location  = "West US"
}

resource "azurerm_dev_test_lab" "test" {
  name                = "example-devtestlab"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  tags = {
    "Sydney" = "Australia"
  }
}

resource "azurerm_dev_test_virtual_network" "test" {
  name                = "example-network"
  lab_name            = "${azurerm_dev_test_lab.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  subnet {
    use_public_ip_address      = "Allow"
    use_in_virtual_machine_creation = "Allow"
  }
}

resource "azurerm_dev_test_windows_virtual_machine" "test" {
  name                = "example-vm03"
  lab_name            = "${azurerm_dev_test_lab.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location            = "${azurerm_resource_group.test.location}"
  size                = "Standard_DS2"
  username            = "exampleuser99"
  password            = "Pa$$w0rd1234!"
  lab_virtual_network_id = "${azurerm_dev_test_virtual_network.test.id}"
  lab_subnet_name     = "${azurerm_dev_test_virtual_network.test.subnet.0.name}"
  storage_type        = "Premium"
  notes               = "Some notes about this Virtual Machine."

  gallery_image_reference {
    offer      = "UbuntuServer"
    publisher  = "Canonical"
    sku        = "18.04-LTS"
  }
}
```

```

    version    = "latest"
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Dev Test Machine. Changing this forces a new resource to be created.

NOTE: The validation requirements for the Name change based on the **os_type** used in this Virtual Machine. For a Linux VM the name must be between 1-62 characters, and for a Windows VM the name must be between 1-15 characters. It must begin and end with a letter or number, and cannot be all numbers.

- **lab_name** - (Required) Specifies the name of the Dev Test Lab in which the Virtual Machine should be created. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the Dev Test Lab resource exists. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the Dev Test Lab exists. Changing this forces a new resource to be created.
- **gallery_image_reference** - (Required) A **gallery_image_reference** block as defined below.
- **lab_subnet_name** - (Required) The name of a Subnet within the Dev Test Virtual Network where this machine should exist. Changing this forces a new resource to be created.
- **lab_virtual_network_id** - (Required) The ID of the Dev Test Virtual Network where this Virtual Machine should be created. Changing this forces a new resource to be created.
- **password** - (Required) The Password associated with the **username** used to login to this Virtual Machine. Changing this forces a new resource to be created.
- **size** - (Required) The Machine Size to use for this Virtual Machine, such as **Standard_F2**. Changing this forces a new resource to be created.
- **storage_type** - (Required) The type of Storage to use on this Virtual Machine. Possible values are **Standard** and **Premium**.
- **username** - (Required) The Username associated with the local administrator on this Virtual Machine. Changing this forces a new resource to be created.

-
- **allow_claim** - (Optional) Can this Virtual Machine be claimed by users? Defaults to **true**.
 - **disallow_public_ip_address** - (Optional) Should the Virtual Machine be created without a Public IP Address? Changing this forces a new resource to be created.
 - **inbound_nat_rule** - (Optional) One or more **inbound_nat_rule** blocks as defined below. Changing this forces a new resource to be created.

NOTE: If any **inbound_nat_rule** blocks are specified then **disallow_public_ip_address** must be set to **true**.

- **notes** - (Optional) Any notes about the Virtual Machine.
- **tags** - (Optional) A mapping of tags to assign to the resource.

A **gallery_image_reference** block supports the following:

- **offer** - (Required) The Offer of the Gallery Image. Changing this forces a new resource to be created.
- **publisher** - (Required) The Publisher of the Gallery Image. Changing this forces a new resource to be created.
- **sku** - (Required) The SKU of the Gallery Image. Changing this forces a new resource to be created.
- **version** - (Required) The Version of the Gallery Image. Changing this forces a new resource to be created.

A **inbound_nat_rule** block supports the following:

- **protocol** - (Required) The Protocol used for this NAT Rule. Possible values are **Tcp** and **Udp**. Changing this forces a new resource to be created.
- **backend_port** - (Required) The Backend Port associated with this NAT Rule. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Virtual Machine.
- **fqdn** - The FQDN of the Virtual Machine.

- `inbound_nat_rule` - One or more `inbound_nat_rule` blocks as defined below.
- `unique_identifier` - The unique immutable identifier of the Virtual Machine.

A `inbound_nat_rule` block exports the following:

- `frontend_port` - The frontend port associated with this Inbound NAT Rule.

» Import

Dev Test Windows Virtual Machines can be imported using the `resource id`, e.g.

```
terraform import azurerm_dev_test_windows_virtual_machine.machine1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_dns_a_record

Enables you to manage DNS A Records within Azure DNS.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "acceptanceTestResourceGroup1"
  location = "West US"
}

resource "azurerm_dns_zone" "test" {
  name                = "mydomain.com"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_dns_a_record" "test" {
  name           = "test"
  zone_name      = "${azurerm_dns_zone.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  ttl            = 300
  records        = ["10.0.180.17"]
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the DNS A Record.
- **resource_group_name** - (Required) Specifies the resource group where the resource exists. Changing this forces a new resource to be created.
- **zone_name** - (Required) Specifies the DNS Zone where the resource exists. Changing this forces a new resource to be created.
- **TTL** - (Required) The Time To Live (TTL) of the DNS record in seconds.
- **records** - (Required) List of IPv4 Addresses.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- **id** - The DNS A Record ID.

» Import

A records can be imported using the `resource id`, e.g.

```
terraform import azurerm_dns_a_record.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_dns_aaaa_record

Enables you to manage DNS AAAA Records within Azure DNS.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "acceptanceTestResourceGroup1"
  location = "West US"
}

resource "azurerm_dns_zone" "test" {
  name                = "mydomain.com"
  resource_group_name = "${azurerm_resource_group.test.name}"
}
```

```
resource "azurerm_dns_aaaa_record" "test" {
  name           = "test"
  zone_name      = "${azurerm_dns_zone.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  ttl           = 300
  records        = ["2607:f8b0:4009:1803::1005"]
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the DNS AAAA Record.
- **resource_group_name** - (Required) Specifies the resource group where the resource exists. Changing this forces a new resource to be created.
- **zone_name** - (Required) Specifies the DNS Zone where the resource exists. Changing this forces a new resource to be created.
- **TTL** - (Required) The Time To Live (TTL) of the DNS record in seconds.
- **records** - (Required) List of IPv6 Addresses.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- **id** - The DNS AAAA Record ID.

» Import

AAAA records can be imported using the **resource id**, e.g.

```
terraform import azurerm_dns_aaaa_record.test /subscriptions/00000000-0000-0000-0000-00000000
```

» azurerm_dns_caa_record

Enables you to manage DNS CAA Records within Azure DNS.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "acceptanceTestResourceGroup1"
  location  = "West US"
}

resource "azurerm_dns_zone" "test" {
  name              = "mydomain.com"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_dns_caa_record" "test" {
  name              = "test"
  zone_name         = "${azurerm_dns_zone.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  ttl               = 300

  record {
    flags = 0
    tag   = "issue"
    value = "example.com"
  }

  record {
    flags = 0
    tag   = "issue"
    value = "example.net"
  }

  record {
    flags = 0
    tag   = "issuewild"
    value = ";"
  }

  record {
    flags = 0
    tag   = "iodef"
    value = "mailto:terraform@nonexisting.tld"
  }

  tags = {
    Environment = "Production"
  }
}
```

}

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the DNS CAA Record.
- **resource_group_name** - (Required) Specifies the resource group where the resource exists. Changing this forces a new resource to be created.
- **zone_name** - (Required) Specifies the DNS Zone where the resource exists. Changing this forces a new resource to be created.
- **ttl** - (Required) The Time To Live (TTL) of the DNS record in seconds.
- **record** - (Required) A list of values that make up the CAA record. Each **record** block supports fields documented below.
- **tags** - (Optional) A mapping of tags to assign to the resource.

The **record** block supports:

- **flags** - (Required) Extensible CAA flags, currently only 1 is implemented to set the issuer critical flag.
- **tag** - (Required) A property tag, options are issue, issuewild and iodef.
- **value** - (Required) A property value such as a registrar domain.

» Attributes Reference

The following attributes are exported:

- **id** - The DNS CAA Record ID.

» Import

CAA records can be imported using the **resource id**, e.g.

```
terraform import azurerm_dns_caa_record.test /subscriptions/00000000-0000-0000-0000-00000000
```

» azurerm_dns_cname_record

Enables you to manage DNS CNAME Records within Azure DNS.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "acceptanceTestResourceGroup1"
  location  = "West US"
}

resource "azurerm_dns_zone" "test" {
  name              = "mydomain.com"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_dns_cname_record" "test" {
  name              = "test"
  zone_name         = "${azurerm_dns_zone.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  ttl               = 300
  record            = "contoso.com"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the DNS CNAME Record.
- **resource_group_name** - (Required) Specifies the resource group where the resource exists. Changing this forces a new resource to be created.
- **zone_name** - (Required) Specifies the DNS Zone where the resource exists. Changing this forces a new resource to be created.
- **TTL** - (Required) The Time To Live (TTL) of the DNS record in seconds.
- **record** - (Required) The target of the CNAME.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- **id** - The DNS CName Record ID.

» Import

CNAME records can be imported using the **resource id**, e.g.

```
terraform import azure_rm_dns_cname_record.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azure_rm_dns_mx_record`

Enables you to manage DNS MX Records within Azure DNS.

» Example Usage

```
resource "azure_rm_resource_group" "test" {
  name      = "acceptanceTestResourceGroup1"
  location  = "West US"
}

resource "azure_rm_dns_zone" "test" {
  name              = "mydomain.com"
  resource_group_name = "${azure_rm_resource_group.test.name}"
}

resource "azure_rm_dns_mx_record" "test" {
  name              = "test"
  zone_name         = "${azure_rm_dns_zone.test.name}"
  resource_group_name = "${azure_rm_resource_group.test.name}"
  ttl               = 300

  record {
    preference = 10
    exchange   = "mail1.contoso.com"
  }

  record {
    preference = 20
    exchange   = "mail2.contoso.com"
  }

  tags = {
    Environment = "Production"
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the DNS MX Record.
- **resource_group_name** - (Required) Specifies the resource group where the resource exists. Changing this forces a new resource to be created.
- **zone_name** - (Required) Specifies the DNS Zone where the resource exists. Changing this forces a new resource to be created.
- **ttl** - (Required) The Time To Live (TTL) of the DNS record in seconds.
- **record** - (Required) A list of values that make up the MX record. Each **record** block supports fields documented below.
- **tags** - (Optional) A mapping of tags to assign to the resource.

The **record** block supports:

- **preference** - (Required) String representing the "preference" value of the MX records. Records with lower preference value take priority.
- **exchange** - (Required) The mail server responsible for the domain covered by the MX record.

» Attributes Reference

The following attributes are exported:

- **id** - The DNS MX Record ID.

» Import

MX records can be imported using the **resource id**, e.g.

```
terraform import azurerm_dns_mx_record.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_dns_ns_record

Enables you to manage DNS NS Records within Azure DNS.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "acceptanceTestResourceGroup1"
  location = "West US"
}
```

```

resource "azurerm_dns_zone" "test" {
  name           = "mydomain.com"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_dns_ns_record" "test" {
  name           = "test"
  zone_name      = "${azurerm_dns_zone.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  ttl            = 300

  records = ["ns1.contoso.com", "ns2.contoso.com"]

  tags = {
    Environment = "Production"
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the DNS NS Record.
- **resource_group_name** - (Required) Specifies the resource group where the resource exists. Changing this forces a new resource to be created.
- **zone_name** - (Required) Specifies the DNS Zone where the resource exists. Changing this forces a new resource to be created.
- **ttl** - (Required) The Time To Live (TTL) of the DNS record in seconds.
- **records** - (Optional) A list of values that make up the NS record. *WARNING*: Either **records** or **record** is required.
- **record** - (Optional) A list of values that make up the NS record. Each **record** block supports fields documented below. This field has been deprecated and will be removed in a future release.
- **tags** - (Optional) A mapping of tags to assign to the resource.

The **record** block supports:

- **nsdname** - (Required) The value of the record.

» Attributes Reference

The following attributes are exported:

- `id` - The DNS NS Record ID.

» Import

NS records can be imported using the `resource id`, e.g.

```
terraform import azure_rm_dns_ns_record.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azure_rm_dns_ptr_record`

Enables you to manage DNS PTR Records within Azure DNS.

» Example Usage

```
resource "azure_rm_resource_group" "test" {
  name      = "acceptanceTestResourceGroup1"
  location = "West US"
}

resource "azure_rm_dns_zone" "test" {
  name                  = "mydomain.com"
  resource_group_name = "${azure_rm_resource_group.test.name}"
}

resource "azure_rm_dns_ptr_record" "test" {
  name          = "test"
  zone_name     = "${azure_rm_dns_zone.test.name}"
  resource_group_name = "${azure_rm_resource_group.test.name}"
  ttl           = 300
  records       = ["yourdomain.com"]
}
```

» Argument Reference

The following arguments are supported:

- `name` - (Required) The name of the DNS PTR Record.
- `resource_group_name` - (Required) Specifies the resource group where the resource exists. Changing this forces a new resource to be created.
- `zone_name` - (Required) Specifies the DNS Zone where the resource exists. Changing this forces a new resource to be created.

- **ttl** - (Required) The Time To Live (TTL) of the DNS record in seconds.
- **records** - (Required) List of Fully Qualified Domain Names.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- **id** - The DNS PTR Record ID.

» Import

PTR records can be imported using the **resource id**, e.g.

```
terraform import azurerm_dns_ptr_record.test /subscriptions/00000000-0000-0000-0000-00000000
```

» azurerm_dns_srv_record

Enables you to manage DNS SRV Records within Azure DNS.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "acceptanceTestResourceGroup1"
  location = "West US"
}

resource "azurerm_dns_zone" "test" {
  name          = "mydomain.com"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_dns_srv_record" "test" {
  name          = "test"
  zone_name     = "${azurerm_dns_zone.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  ttl          = 300

  record {
    priority = 1
    weight   = 5
  }
}
```

```

    port      = 8080
    target    = "target1.contoso.com"
  }

  tags = {
    Environment = "Production"
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the DNS SRV Record.
- **resource_group_name** - (Required) Specifies the resource group where the resource exists. Changing this forces a new resource to be created.
- **zone_name** - (Required) Specifies the DNS Zone where the resource exists. Changing this forces a new resource to be created.
- **ttl** - (Required) The Time To Live (TTL) of the DNS record in seconds.
- **record** - (Required) A list of values that make up the SRV record. Each **record** block supports fields documented below.
- **tags** - (Optional) A mapping of tags to assign to the resource.

The **record** block supports:

- **priority** - (Required) Priority of the SRV record.
- **weight** - (Required) Weight of the SRV record.
- **port** - (Required) Port the service is listening on.
- **target** - (Required) FQDN of the service.

» Attributes Reference

The following attributes are exported:

- **id** - The DNS SRV Record ID.

» Import

SRV records can be imported using the **resource id**, e.g.

```
terraform import azurerm_dns_srv_record.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_dns_txt_record

Enables you to manage DNS TXT Records within Azure DNS.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "acceptanceTestResourceGroup1"
  location  = "West US"
}

resource "azurerm_dns_zone" "test" {
  name              = "mydomain.com"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_dns_txt_record" "test" {
  name              = "test"
  zone_name         = "${azurerm_dns_zone.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  ttl               = 300

  record {
    value = "google-site-authenticator"
  }

  record {
    value = "more site information here"
  }

  tags = {
    Environment = "Production"
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the DNS TXT Record.
- **resource_group_name** - (Required) Specifies the resource group where the resource exists. Changing this forces a new resource to be created.

- **zone_name** - (Required) Specifies the DNS Zone where the resource exists. Changing this forces a new resource to be created.
- **ttl** - (Required) The Time To Live (TTL) of the DNS record in seconds.
- **record** - (Required) A list of values that make up the txt record. Each **record** block supports fields documented below.
- **tags** - (Optional) A mapping of tags to assign to the resource.

The **record** block supports:

- **value** - (Required) The value of the record.

» Attributes Reference

The following attributes are exported:

- **id** - The DNS TXT Record ID.

» Import

TXT records can be imported using the **resource id**, e.g.

```
terraform import azurerm_dns_txt_record.test /subscriptions/00000000-0000-0000-0000-00000000
```

» azurerm_dns_zone

Enables you to manage DNS zones within Azure DNS. These zones are hosted on Azure's name servers to which you can delegate the zone from the parent domain.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "acceptanceTestResourceGroup1"
  location = "West US"
}

resource "azurerm_dns_zone" "test" {
  name                = "mydomain.com"
  resource_group_name = "${azurerm_resource_group.test.name}"
  zone_type           = "Public"
}
```

```
resource "azurerm_dns_zone" "test_private" {
  name           = "mydomain.com"
  resource_group_name = "${azurerm_resource_group.test.name}"
  zone_type      = "Private"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the DNS Zone. Must be a valid domain name.
- **resource_group_name** - (Required) Specifies the resource group where the resource exists. Changing this forces a new resource to be created.
- **zone_type** - (Required) Specifies the type of this DNS zone. Possible values are `Public` or `Private` (Defaults to `Public`).
- **registration_virtual_network_ids** - (Optional) A list of Virtual Network ID's that register hostnames in this DNS zone. This field can only be set when **zone_type** is set to `Private`.
- **resolution_virtual_network_ids** - (Optional) A list of Virtual Network ID's that resolve records in this DNS zone. This field can only be set when **zone_type** is set to `Private`.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- **id** - The DNS Zone ID.
- **max_number_of_record_sets** - (Optional) Maximum number of Records in the zone. Defaults to 1000.
- **number_of_record_sets** - (Optional) The number of records already in the zone.
- **name_servers** - (Optional) A list of values that make up the NS record for the zone.

» Import

DNS Zones can be imported using the **resource id**, e.g.

```
terraform import azurerm_dns_zone.zone1 /subscriptions/00000000-0000-0000-0000-000000000000/
```

» azurerm_hdinsight_hadoop_cluster

Manages a HDInsight Hadoop Cluster.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "example-resources"
  location = "West Europe"
}

resource "azurerm_storage_account" "example" {
  name                     = "hdinsightstor"
  resource_group_name      = "${azurerm_resource_group.example.name}"
  location                 = "${azurerm_resource_group.example.location}"
  account_tier             = "Standard"
  account_replication_type = "LRS"
}

resource "azurerm_storage_container" "example" {
  name                     = "hdinsight"
  resource_group_name      = "${azurerm_resource_group.example.name}"
  storage_account_name     = "${azurerm_storage_account.example.name}"
  container_access_type    = "private"
}

resource "azurerm_hdinsight_hadoop_cluster" "example" {
  name                     = "example-hdicluster"
  resource_group_name      = "${azurerm_resource_group.example.name}"
  location                 = "${azurerm_resource_group.example.location}"
  cluster_version          = "3.6"
  tier                     = "Standard"

  component_version {
    hadoop = "2.7"
  }

  gateway {
    enabled = true
    username = "acctestusrgw"
    password = "TerrAform123!"
  }

  storage_account {
```

```

    storage_container_id = "${azurerm_storage_container.example.id}"
    storage_account_key  = "${azurerm_storage_account.example.primary_access_key}"
    is_default           = true
  }

  roles {
    head_node {
      vm_size   = "Standard_D3_V2"
      username  = "acctestusrvm"
      password  = "AccTestvdSC4daf986!"
    }

    worker_node {
      vm_size           = "Standard_D4_V2"
      username          = "acctestusrvm"
      password          = "AccTestvdSC4daf986!"
      target_instance_count = 3
    }

    zookeeper_node {
      vm_size   = "Standard_D3_V2"
      username  = "acctestusrvm"
      password  = "AccTestvdSC4daf986!"
    }
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name for this HDInsight Hadoop Cluster. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) Specifies the name of the Resource Group in which this HDInsight Hadoop Cluster should exist. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the Azure Region which this HDInsight Hadoop Cluster should exist. Changing this forces a new resource to be created.
- **cluster_version** - (Required) Specifies the Version of HDInsights which should be used for this Cluster. Changing this forces a new resource to be created.

- **component_version** - (Required) A **component_version** block as defined below.
- **gateway** - (Required) A **gateway** block as defined below.
- **roles** - (Required) A **roles** block as defined below.
- **storage_account** - (Required) One or more **storage_account** block as defined below.
- **tier** - (Required) Specifies the Tier which should be used for this HDInsight Hadoop Cluster. Possible values are **Standard** or **Premium**. Changing this forces a new resource to be created.

-
- **tags** - (Optional) A map of Tags which should be assigned to this HDInsight Hadoop Cluster.
-

A **component_version** block supports the following:

- **hadoop** - (Required) The version of Hadoop which should be used for this HDInsight Hadoop Cluster. Changing this forces a new resource to be created.
-

A **gateway** block supports the following:

- **enabled** - (Required) Is the Ambari portal enabled? Changing this forces a new resource to be created.
- **password** - (Required) The password used for the Ambari Portal. Changing this forces a new resource to be created.

NOTE: This password must be different from the one used for the **head_node**, **worker_node** and **zookeeper_node** roles.

- **username** - (Required) The username used for the Ambari Portal. Changing this forces a new resource to be created.
-

A **head_node** block supports the following:

- **username** - (Required) The Username of the local administrator for the Head Nodes. Changing this forces a new resource to be created.
- **vm_size** - (Required) The Size of the Virtual Machine which should be used as the Head Nodes. Changing this forces a new resource to be created.
- **password** - (Optional) The Password associated with the local administrator for the Head Nodes. Changing this forces a new resource to be created.

NOTE: If specified, this password must be at least 10 characters in length and must contain at least one digit, one uppercase and one lower case letter, one non-alphanumeric character (except characters ' " ').

- **ssh_keys** - (Optional) A list of SSH Keys which should be used for the local administrator on the Head Nodes. Changing this forces a new resource to be created.

NOTE: Either a **password** or one or more **ssh_keys** must be specified - but not both.

- **subnet_id** - (Optional) The ID of the Subnet within the Virtual Network where the Head Nodes should be provisioned within. Changing this forces a new resource to be created.
- **virtual_network_id** - (Optional) The ID of the Virtual Network where the Head Nodes should be provisioned within. Changing this forces a new resource to be created.

A **roles** block supports the following:

- **head_node** - (Required) A **head_node** block as defined above.
- **worker_node** - (Required) A **worker_node** block as defined below.
- **zookeeper_node** - (Required) A **zookeeper_node** block as defined below.

A **storage_account** block supports the following:

- **is_default** - (Required) Is this the Default Storage Account for the HDInsight Hadoop Cluster? Changing this forces a new resource to be created.

NOTE: One of the **storage_account** blocks must be marked as the default.

- **storage_account_key** - (Required) The Access Key which should be used to connect to the Storage Account. Changing this forces a new resource to be created.
- **storage_container_id** - (Required) The ID of the Storage Container. Changing this forces a new resource to be created.

NOTE: This can be obtained from the **id** of the **azurerm_storage_container** resource.

A **worker_node** block supports the following:

- **username** - (Required) The Username of the local administrator for the Worker Nodes. Changing this forces a new resource to be created.

- **vm_size** - (Required) The Size of the Virtual Machine which should be used as the Worker Nodes. Changing this forces a new resource to be created.
- **min_instance_count** - (Optional) The minimum number of instances which should be run for the Worker Nodes. Changing this forces a new resource to be created.
- **password** - (Optional) The Password associated with the local administrator for the Worker Nodes. Changing this forces a new resource to be created.

NOTE: If specified, this password must be at least 10 characters in length and must contain at least one digit, one uppercase and one lower case letter, one non-alphanumeric character (except characters ' " ').

- **ssh_keys** - (Optional) A list of SSH Keys which should be used for the local administrator on the Worker Nodes. Changing this forces a new resource to be created.

NOTE: Either a **password** or one or more **ssh_keys** must be specified - but not both.

- **subnet_id** - (Optional) The ID of the Subnet within the Virtual Network where the Worker Nodes should be provisioned within. Changing this forces a new resource to be created.
- **target_instance_count** - (Optional) The number of instances which should be run for the Worker Nodes.
- **virtual_network_id** - (Optional) The ID of the Virtual Network where the Worker Nodes should be provisioned within. Changing this forces a new resource to be created.

A **zookeeper_node** block supports the following:

- **username** - (Required) The Username of the local administrator for the Zookeeper Nodes. Changing this forces a new resource to be created.
- **vm_size** - (Required) The Size of the Virtual Machine which should be used as the Zookeeper Nodes. Changing this forces a new resource to be created.
- **password** - (Optional) The Password associated with the local administrator for the Zookeeper Nodes. Changing this forces a new resource to be created.

NOTE: If specified, this password must be at least 10 characters in length and must contain at least one digit, one uppercase and one lower case letter, one non-alphanumeric character (except characters ' " ').

- **ssh_keys** - (Optional) A list of SSH Keys which should be used for the local administrator on the Zookeeper Nodes. Changing this forces a new resource to be created.

NOTE: Either a **password** or one or more **ssh_keys** must be specified - but not both.

- **subnet_id** - (Optional) The ID of the Subnet within the Virtual Network where the Zookeeper Nodes should be provisioned within. Changing this forces a new resource to be created.
- **virtual_network_id** - (Optional) The ID of the Virtual Network where the Zookeeper Nodes should be provisioned within. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the HDInsight Hadoop Cluster.
- **https_endpoint** - The HTTPS Connectivity Endpoint for this HDInsight Hadoop Cluster.
- **ssh_endpoint** - The SSH Connectivity Endpoint for this HDInsight Hadoop Cluster.

» Import

HDInsight Hadoop Clusters can be imported using the **resource id**, e.g.

```
terraform import azurerm_hdinsight_hadoop_cluster.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm__hdinsight__hbase__cluster

Manages a HDInsight HBase Cluster.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "example-resources"
  location = "West Europe"
}

resource "azurerm_storage_account" "example" {
```



```

    name                = "hdinsightstor"
    resource_group_name  = "${azurerm_resource_group.example.name}"
    location             = "${azurerm_resource_group.example.location}"
    account_tier         = "Standard"
    account_replication_type = "LRS"
  }

  resource "azurerm_storage_container" "example" {
    name                = "hdinsight"
    resource_group_name = "${azurerm_resource_group.example.name}"
    storage_account_name = "${azurerm_storage_account.example.name}"
    container_access_type = "private"
  }

  resource "azurerm_hdinsight_hbase_cluster" "example" {
    name                = "example-hdicluster"
    resource_group_name = "${azurerm_resource_group.example.name}"
    location            = "${azurerm_resource_group.example.location}"
    cluster_version     = "3.6"
    tier                 = "Standard"

    component_version {
      hbase = "1.1"
    }

    gateway {
      enabled = true
      username = "acctestusrgw"
      password = "Terraform123!"
    }

    storage_account {
      storage_container_id = "${azurerm_storage_container.example.id}"
      storage_account_key  = "${azurerm_storage_account.example.primary_access_key}"
      is_default           = true
    }

    roles {
      head_node {
        vm_size = "Standard_D3_V2"
        username = "acctestusrv"
        password = "AccTestvdSC4daf986!"
      }

      worker_node {
        vm_size = "Standard_D3_V2"
      }
    }
  }

```

```

        username          = "acctestusrvm"
        password          = "AccTestvdSC4daf986!"
        target_instance_count = 3
    }

    zookeeper_node {
        vm_size = "Standard_D3_V2"
        username = "acctestusrvm"
        password = "AccTestvdSC4daf986!"
    }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name for this HDInsight HBase Cluster. Changing this forces a new resource to be created.
 - **resource_group_name** - (Required) Specifies the name of the Resource Group in which this HDInsight HBase Cluster should exist. Changing this forces a new resource to be created.
 - **location** - (Required) Specifies the Azure Region which this HDInsight HBase Cluster should exist. Changing this forces a new resource to be created.
 - **cluster_version** - (Required) Specifies the Version of HDInsights which should be used for this Cluster. Changing this forces a new resource to be created.
 - **component_version** - (Required) A **component_version** block as defined below.
 - **gateway** - (Required) A **gateway** block as defined below.
 - **roles** - (Required) A **roles** block as defined below.
 - **storage_account** - (Required) One or more **storage_account** block as defined below.
 - **tier** - (Required) Specifies the Tier which should be used for this HDInsight HBase Cluster. Possible values are **Standard** or **Premium**. Changing this forces a new resource to be created.
-
- **tags** - (Optional) A map of Tags which should be assigned to this HDInsight HBase Cluster.

A `component_version` block supports the following:

- **hbase** - (Required) The version of HBase which should be used for this HDInsight HBase Cluster. Changing this forces a new resource to be created.

A `gateway` block supports the following:

- **enabled** - (Required) Is the Ambari portal enabled? Changing this forces a new resource to be created.
- **password** - (Required) The password used for the Ambari Portal. Changing this forces a new resource to be created.

NOTE: This password must be different from the one used for the `head_node`, `worker_node` and `zookeeper_node` roles.

- **username** - (Required) The username used for the Ambari Portal. Changing this forces a new resource to be created.

A `head_node` block supports the following:

- **username** - (Required) The Username of the local administrator for the Head Nodes. Changing this forces a new resource to be created.
- **vm_size** - (Required) The Size of the Virtual Machine which should be used as the Head Nodes. Changing this forces a new resource to be created.
- **password** - (Optional) The Password associated with the local administrator for the Head Nodes. Changing this forces a new resource to be created.

NOTE: If specified, this password must be at least 10 characters in length and must contain at least one digit, one uppercase and one lower case letter, one non-alphanumeric character (except characters `'` `''` `'`).

- **ssh_keys** - (Optional) A list of SSH Keys which should be used for the local administrator on the Head Nodes. Changing this forces a new resource to be created.

NOTE: Either a `password` or one or more `ssh_keys` must be specified - but not both.

- **subnet_id** - (Optional) The ID of the Subnet within the Virtual Network where the Head Nodes should be provisioned within. Changing this forces a new resource to be created.

- **virtual_network_id** - (Optional) The ID of the Virtual Network where the Head Nodes should be provisioned within. Changing this forces a new resource to be created.

A **roles** block supports the following:

- **head_node** - (Required) A **head_node** block as defined above.
- **worker_node** - (Required) A **worker_node** block as defined below.
- **zookeeper_node** - (Required) A **zookeeper_node** block as defined below.

A **storage_account** block supports the following:

- **is_default** - (Required) Is this the Default Storage Account for the HDInsight HBase Cluster? Changing this forces a new resource to be created.

NOTE: One of the **storage_account** blocks must be marked as the default.

- **storage_account_key** - (Required) The Access Key which should be used to connect to the Storage Account. Changing this forces a new resource to be created.
- **storage_container_id** - (Required) The ID of the Storage Container. Changing this forces a new resource to be created.

NOTE: This can be obtained from the **id** of the **azurerm_storage_container** resource.

A **worker_node** block supports the following:

- **username** - (Required) The Username of the local administrator for the Worker Nodes. Changing this forces a new resource to be created.
- **vm_size** - (Required) The Size of the Virtual Machine which should be used as the Worker Nodes. Changing this forces a new resource to be created.
- **min_instance_count** - (Optional) The minimum number of instances which should be run for the Worker Nodes. Changing this forces a new resource to be created.
- **password** - (Optional) The Password associated with the local administrator for the Worker Nodes. Changing this forces a new resource to be created.

NOTE: If specified, this password must be at least 10 characters in length and must contain at least one digit, one uppercase and one lower case letter, one non-alphanumeric character (except characters ' " ').

- **ssh_keys** - (Optional) A list of SSH Keys which should be used for the local administrator on the Worker Nodes. Changing this forces a new resource to be created.

NOTE: Either a **password** or one or more **ssh_keys** must be specified - but not both.

- **subnet_id** - (Optional) The ID of the Subnet within the Virtual Network where the Worker Nodes should be provisioned within. Changing this forces a new resource to be created.
- **target_instance_count** - (Optional) The number of instances which should be run for the Worker Nodes.
- **virtual_network_id** - (Optional) The ID of the Virtual Network where the Worker Nodes should be provisioned within. Changing this forces a new resource to be created.

A **zookeeper_node** block supports the following:

- **username** - (Required) The Username of the local administrator for the Zookeeper Nodes. Changing this forces a new resource to be created.
- **vm_size** - (Required) The Size of the Virtual Machine which should be used as the Zookeeper Nodes. Changing this forces a new resource to be created.
- **password** - (Optional) The Password associated with the local administrator for the Zookeeper Nodes. Changing this forces a new resource to be created.

NOTE: If specified, this password must be at least 10 characters in length and must contain at least one digit, one uppercase and one lower case letter, one non-alphanumeric character (except characters ' " ').

- **ssh_keys** - (Optional) A list of SSH Keys which should be used for the local administrator on the Zookeeper Nodes. Changing this forces a new resource to be created.

NOTE: Either a **password** or one or more **ssh_keys** must be specified - but not both.

- **subnet_id** - (Optional) The ID of the Subnet within the Virtual Network where the Zookeeper Nodes should be provisioned within. Changing this forces a new resource to be created.
- **virtual_network_id** - (Optional) The ID of the Virtual Network where the Zookeeper Nodes should be provisioned within. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the HDInsight HBase Cluster.
- `https_endpoint` - The HTTPS Connectivity Endpoint for this HDInsight HBase Cluster.
- `ssh_endpoint` - The SSH Connectivity Endpoint for this HDInsight HBase Cluster.

» Import

HDInsight HBase Clusters can be imported using the `resource id`, e.g.

```
terraform import azurerm_hdinsight_hbase_cluster.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_hdinsight_kafka_cluster`

Manages a HDInsight Kafka Cluster.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "example-resources"
  location = "West Europe"
}

resource "azurerm_storage_account" "example" {
  name                     = "hdinsightstor"
  resource_group_name      = "${azurerm_resource_group.example.name}"
  location                 = "${azurerm_resource_group.example.location}"
  account_tier             = "Standard"
  account_replication_type = "LRS"
}

resource "azurerm_storage_container" "example" {
  name                     = "hdinsight"
  resource_group_name      = "${azurerm_resource_group.example.name}"
  storage_account_name     = "${azurerm_storage_account.example.name}"
  container_access_type    = "private"
}
```

```

resource "azurerm_hdinsight_kafka_cluster" "example" {
  name                = "example-hdicluster"
  resource_group_name = "${azurerm_resource_group.example.name}"
  location            = "${azurerm_resource_group.example.location}"
  cluster_version     = "3.6"
  tier                 = "Standard"

  component_version {
    kafka = "2.3"
  }

  gateway {
    enabled = true
    username = "acctestusrgw"
    password = "Terraform123!"
  }

  storage_account {
    storage_container_id = "${azurerm_storage_container.example.id}"
    storage_account_key  = "${azurerm_storage_account.example.primary_access_key}"
    is_default            = true
  }

  roles {
    head_node {
      vm_size   = "Standard_D3_V2"
      username  = "acctestusrv"
      password  = "AccTestvdSC4daf986!"
    }

    worker_node {
      vm_size           = "Standard_D3_V2"
      username          = "acctestusrv"
      password          = "AccTestvdSC4daf986!"
      number_of_disks_per_node = 3
      target_instance_count  = 3
    }

    zookeeper_node {
      vm_size   = "Standard_D3_V2"
      username  = "acctestusrv"
      password  = "AccTestvdSC4daf986!"
    }
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name for this HDInsight Kafka Cluster. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) Specifies the name of the Resource Group in which this HDInsight Kafka Cluster should exist. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the Azure Region which this HDInsight Kafka Cluster should exist. Changing this forces a new resource to be created.
- **cluster_version** - (Required) Specifies the Version of HDInsights which should be used for this Cluster. Changing this forces a new resource to be created.
- **component_version** - (Required) A **component_version** block as defined below.
- **gateway** - (Required) A **gateway** block as defined below.
- **roles** - (Required) A **roles** block as defined below.
- **storage_account** - (Required) One or more **storage_account** block as defined below.
- **tier** - (Required) Specifies the Tier which should be used for this HDInsight Kafka Cluster. Possible values are **Standard** or **Premium**. Changing this forces a new resource to be created.

-
- **tags** - (Optional) A map of Tags which should be assigned to this HDInsight Kafka Cluster.
-

A **component_version** block supports the following:

- **kafka** - (Required) The version of Kafka which should be used for this HDInsight Kafka Cluster. Changing this forces a new resource to be created.
-

A **gateway** block supports the following:

- **enabled** - (Required) Is the Ambari portal enabled? Changing this forces a new resource to be created.

- **password** - (Required) The password used for the Ambari Portal. Changing this forces a new resource to be created.

NOTE: This password must be different from the one used for the **head_node**, **worker_node** and **zookeeper_node** roles.

- **username** - (Required) The username used for the Ambari Portal. Changing this forces a new resource to be created.

A **head_node** block supports the following:

- **username** - (Required) The Username of the local administrator for the Head Nodes. Changing this forces a new resource to be created.
- **vm_size** - (Required) The Size of the Virtual Machine which should be used as the Head Nodes. Changing this forces a new resource to be created.
- **password** - (Optional) The Password associated with the local administrator for the Head Nodes. Changing this forces a new resource to be created.

NOTE: If specified, this password must be at least 10 characters in length and must contain at least one digit, one uppercase and one lower case letter, one non-alphanumeric character (except characters ' " ').

- **ssh_keys** - (Optional) A list of SSH Keys which should be used for the local administrator on the Head Nodes. Changing this forces a new resource to be created.

NOTE: Either a **password** or one or more **ssh_keys** must be specified - but not both.

- **subnet_id** - (Optional) The ID of the Subnet within the Virtual Network where the Head Nodes should be provisioned within. Changing this forces a new resource to be created.
- **virtual_network_id** - (Optional) The ID of the Virtual Network where the Head Nodes should be provisioned within. Changing this forces a new resource to be created.

A **roles** block supports the following:

- **head_node** - (Required) A **head_node** block as defined above.
- **worker_node** - (Required) A **worker_node** block as defined below.
- **zookeeper_node** - (Required) A **zookeeper_node** block as defined below.

A **storage_account** block supports the following:

- **is_default** - (Required) Is this the Default Storage Account for the HDInsight Kafka Cluster? Changing this forces a new resource to be created.

NOTE: One of the **storage_account** blocks must be marked as the default.

- **storage_account_key** - (Required) The Access Key which should be used to connect to the Storage Account. Changing this forces a new resource to be created.
- **storage_container_id** - (Required) The ID of the Storage Container. Changing this forces a new resource to be created.

NOTE: This can be obtained from the **id** of the **azurerm_storage_container** resource.

A **worker_node** block supports the following:

- **number_of_disks_per_node** - (Required) The number of Data Disks which should be assigned to each Worker Node, which can be between 1 and 8. Changing this forces a new resource to be created.
- **username** - (Required) The Username of the local administrator for the Worker Nodes. Changing this forces a new resource to be created.
- **vm_size** - (Required) The Size of the Virtual Machine which should be used as the Worker Nodes. Changing this forces a new resource to be created.
- **min_instance_count** - (Optional) The minimum number of instances which should be run for the Worker Nodes. Changing this forces a new resource to be created.
- **password** - (Optional) The Password associated with the local administrator for the Worker Nodes. Changing this forces a new resource to be created.

NOTE: If specified, this password must be at least 10 characters in length and must contain at least one digit, one uppercase and one lower case letter, one non-alphanumeric character (except characters ' " ').

- **ssh_keys** - (Optional) A list of SSH Keys which should be used for the local administrator on the Worker Nodes. Changing this forces a new resource to be created.

NOTE: Either a **password** or one or more **ssh_keys** must be specified - but not both.

- **subnet_id** - (Optional) The ID of the Subnet within the Virtual Network where the Worker Nodes should be provisioned within. Changing this forces a new resource to be created.

- **target_instance_count** - (Optional) The number of instances which should be run for the Worker Nodes.
- **virtual_network_id** - (Optional) The ID of the Virtual Network where the Worker Nodes should be provisioned within. Changing this forces a new resource to be created.

A **zookeeper_node** block supports the following:

- **username** - (Required) The Username of the local administrator for the Zookeeper Nodes. Changing this forces a new resource to be created.
- **vm_size** - (Required) The Size of the Virtual Machine which should be used as the Zookeeper Nodes. Changing this forces a new resource to be created.
- **password** - (Optional) The Password associated with the local administrator for the Zookeeper Nodes. Changing this forces a new resource to be created.

NOTE: If specified, this password must be at least 10 characters in length and must contain at least one digit, one uppercase and one lower case letter, one non-alphanumeric character (except characters ' " ').

- **ssh_keys** - (Optional) A list of SSH Keys which should be used for the local administrator on the Zookeeper Nodes. Changing this forces a new resource to be created.

NOTE: Either a **password** or one or more **ssh_keys** must be specified - but not both.

- **subnet_id** - (Optional) The ID of the Subnet within the Virtual Network where the Zookeeper Nodes should be provisioned within. Changing this forces a new resource to be created.
- **virtual_network_id** - (Optional) The ID of the Virtual Network where the Zookeeper Nodes should be provisioned within. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the HDInsight Kafka Cluster.
- **https_endpoint** - The HTTPS Connectivity Endpoint for this HDInsight Kafka Cluster.
- **ssh_endpoint** - The SSH Connectivity Endpoint for this HDInsight Kafka Cluster.

» Import

HDInsight Kafka Clusters can be imported using the `resource_id`, e.g.

```
terraform import azurerm_hdinsight_kafka_cluster.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_hdinsight_interactive_query_cluster`

Manages a HDInsight Interactive Query Cluster.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "example-resources"
  location = "West Europe"
}

resource "azurerm_storage_account" "example" {
  name                        = "hdinsightstor"
  resource_group_name        = "${azurerm_resource_group.example.name}"
  location                   = "${azurerm_resource_group.example.location}"
  account_tier                = "Standard"
  account_replication_type    = "LRS"
}

resource "azurerm_storage_container" "example" {
  name                = "hdinsight"
  resource_group_name = "${azurerm_resource_group.example.name}"
  storage_account_name = "${azurerm_storage_account.example.name}"
  container_access_type = "private"
}

resource "azurerm_hdinsight_interactive_query_cluster" "example" {
  name                = "example-hdicluster"
  resource_group_name = "${azurerm_resource_group.example.name}"
  location             = "${azurerm_resource_group.example.location}"
  cluster_version      = "3.6"
  tier                  = "Standard"

  component_version {
    interactive_hive = "2.1"
  }

  gateway {}
}
```

```

    enabled = true
    username = "acctestusrgw"
    password = "Terraform123!"
  }

  storage_account {
    storage_container_id = "${azurerm_storage_container.example.id}"
    storage_account_key  = "${azurerm_storage_account.example.primary_access_key}"
    is_default           = true
  }

  roles {
    head_node {
      vm_size = "Standard_D13_V2"
      username = "acctestusrvvm"
      password = "AccTestvdSC4daf986!"
    }

    worker_node {
      vm_size           = "Standard_D14_V2"
      username          = "acctestusrvvm"
      password          = "AccTestvdSC4daf986!"
      target_instance_count = 3
    }

    zookeeper_node {
      vm_size = "Standard_A4_V2"
      username = "acctestusrvvm"
      password = "AccTestvdSC4daf986!"
    }
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name for this HDInsight Interactive Query Cluster. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) Specifies the name of the Resource Group in which this HDInsight Interactive Query Cluster should exist. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the Azure Region which this HDInsight Interactive Query Cluster should exist. Changing this forces a new re-

source to be created.

- **cluster_version** - (Required) Specifies the Version of HDInsights which should be used for this Cluster. Changing this forces a new resource to be created.
- **component_version** - (Required) A **component_version** block as defined below.
- **gateway** - (Required) A **gateway** block as defined below.
- **roles** - (Required) A **roles** block as defined below.
- **storage_account** - (Required) One or more **storage_account** block as defined below.
- **tier** - (Required) Specifies the Tier which should be used for this HDInsight Interactive Query Cluster. Possible values are **Standard** or **Premium**. Changing this forces a new resource to be created.

-
- **tags** - (Optional) A map of Tags which should be assigned to this HDInsight Interactive Query Cluster.

A **component_version** block supports the following:

- **interactive_query** - (Required) The version of Interactive Query which should be used for this HDInsight Interactive Query Cluster. Changing this forces a new resource to be created.

A **gateway** block supports the following:

- **enabled** - (Required) Is the Ambari portal enabled? Changing this forces a new resource to be created.
- **password** - (Required) The password used for the Ambari Portal. Changing this forces a new resource to be created.

NOTE: This password must be different from the one used for the **head_node**, **worker_node** and **zookeeper_node** roles.

- **username** - (Required) The username used for the Ambari Portal. Changing this forces a new resource to be created.

A **head_node** block supports the following:

- **username** - (Required) The Username of the local administrator for the Head Nodes. Changing this forces a new resource to be created.

- **vm_size** - (Required) The Size of the Virtual Machine which should be used as the Head Nodes. Changing this forces a new resource to be created.

NOTE: High memory instances must be specified for the Head Node (Azure suggests a **Standard_D13_V2**).

- **password** - (Optional) The Password associated with the local administrator for the Head Nodes. Changing this forces a new resource to be created.

NOTE: If specified, this password must be at least 10 characters in length and must contain at least one digit, one uppercase and one lower case letter, one non-alphanumeric character (except characters ' " ').

- **ssh_keys** - (Optional) A list of SSH Keys which should be used for the local administrator on the Head Nodes. Changing this forces a new resource to be created.

NOTE: Either a **password** or one or more **ssh_keys** must be specified - but not both.

- **subnet_id** - (Optional) The ID of the Subnet within the Virtual Network where the Head Nodes should be provisioned within. Changing this forces a new resource to be created.
- **virtual_network_id** - (Optional) The ID of the Virtual Network where the Head Nodes should be provisioned within. Changing this forces a new resource to be created.

A **roles** block supports the following:

- **head_node** - (Required) A **head_node** block as defined above.
- **worker_node** - (Required) A **worker_node** block as defined below.
- **zookeeper_node** - (Required) A **zookeeper_node** block as defined below.

A **storage_account** block supports the following:

- **is_default** - (Required) Is this the Default Storage Account for the HDInsight Interactive Query Cluster? Changing this forces a new resource to be created.

NOTE: One of the **storage_account** blocks must be marked as the default.

- **storage_account_key** - (Required) The Access Key which should be used to connect to the Storage Account. Changing this forces a new resource to be created.
- **storage_container_id** - (Required) The ID of the Storage Container. Changing this forces a new resource to be created.

NOTE: This can be obtained from the `id` of the `azurerm_storage_container` resource.

A `worker_node` block supports the following:

- `username` - (Required) The Username of the local administrator for the Worker Nodes. Changing this forces a new resource to be created.
- `vm_size` - (Required) The Size of the Virtual Machine which should be used as the Worker Nodes. Changing this forces a new resource to be created.

NOTE: High memory instances must be specified for the Head Node (Azure suggests a `Standard_D14_V2`).

- `min_instance_count` - (Optional) The minimum number of instances which should be run for the Worker Nodes. Changing this forces a new resource to be created.
- `password` - (Optional) The Password associated with the local administrator for the Worker Nodes. Changing this forces a new resource to be created.

NOTE: If specified, this password must be at least 10 characters in length and must contain at least one digit, one uppercase and one lower case letter, one non-alphanumeric character (except characters `'` `"` `'` `'` `'`).

- `ssh_keys` - (Optional) A list of SSH Keys which should be used for the local administrator on the Worker Nodes. Changing this forces a new resource to be created.

NOTE: Either a `password` or one or more `ssh_keys` must be specified - but not both.

- `subnet_id` - (Optional) The ID of the Subnet within the Virtual Network where the Worker Nodes should be provisioned within. Changing this forces a new resource to be created.
- `target_instance_count` - (Optional) The number of instances which should be run for the Worker Nodes.
- `virtual_network_id` - (Optional) The ID of the Virtual Network where the Worker Nodes should be provisioned within. Changing this forces a new resource to be created.

A `zookeeper_node` block supports the following:

- `username` - (Required) The Username of the local administrator for the Zookeeper Nodes. Changing this forces a new resource to be created.

- **vm_size** - (Required) The Size of the Virtual Machine which should be used as the Zookeeper Nodes. Changing this forces a new resource to be created.
- **password** - (Optional) The Password associated with the local administrator for the Zookeeper Nodes. Changing this forces a new resource to be created.

NOTE: If specified, this password must be at least 10 characters in length and must contain at least one digit, one uppercase and one lower case letter, one non-alphanumeric character (except characters ' " ').

- **ssh_keys** - (Optional) A list of SSH Keys which should be used for the local administrator on the Zookeeper Nodes. Changing this forces a new resource to be created.

NOTE: Either a **password** or one or more **ssh_keys** must be specified - but not both.

- **subnet_id** - (Optional) The ID of the Subnet within the Virtual Network where the Zookeeper Nodes should be provisioned within. Changing this forces a new resource to be created.
- **virtual_network_id** - (Optional) The ID of the Virtual Network where the Zookeeper Nodes should be provisioned within. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the HDInsight Interactive Query Cluster.
- **https_endpoint** - The HTTPS Connectivity Endpoint for this HDInsight Interactive Query Cluster.
- **ssh_endpoint** - The SSH Connectivity Endpoint for this HDInsight Interactive Query Cluster.

» Import

HDInsight Interactive Query Clusters can be imported using the **resource id**, e.g.

```
terraform import azurerm_hdinsight_interactive_query_cluster.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_hdinsight_ml_services_cluster

Manages a HDInsight ML Services Cluster.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "example-resources"
  location = "West Europe"
}

resource "azurerm_storage_account" "example" {
  name                        = "hdinsightstor"
  resource_group_name        = "${azurerm_resource_group.example.name}"
  location                   = "${azurerm_resource_group.example.location}"
  account_tier                = "Standard"
  account_replication_type    = "LRS"
}

resource "azurerm_storage_container" "example" {
  name                = "hdinsight"
  resource_group_name = "${azurerm_resource_group.example.name}"
  storage_account_name = "${azurerm_storage_account.example.name}"
  container_access_type = "private"
}

resource "azurerm_hdinsight_ml_services_cluster" "example" {
  name                = "example-hdicluster"
  resource_group_name = "${azurerm_resource_group.example.name}"
  location            = "${azurerm_resource_group.example.location}"
  cluster_version     = "3.6"
  tier                 = "Standard"
  rstudio              = true

  gateway {
    enabled = true
    username = "acctestusrgw"
    password = "TerrAform123!"
  }

  storage_account {
    storage_container_id = "${azurerm_storage_container.example.id}"
    storage_account_key   = "${azurerm_storage_account.example.primary_access_key}"
    is_default             = true
  }
}
```

```

}

roles {
  head_node {
    vm_size   = "Standard_D3_v2"
    username  = "acctestusrvm"
    password  = "AccTestvdSC4daf986!"
  }

  worker_node {
    vm_size           = "Standard_D4_V2"
    username          = "acctestusrvm"
    password          = "AccTestvdSC4daf986!"
    target_instance_count = 3
  }

  zookeeper_node {
    vm_size   = "Standard_D3_v2"
    username  = "acctestusrvm"
    password  = "AccTestvdSC4daf986!"
  }

  edge_node {
    vm_size   = "Standard_D3_v2"
    username  = "acctestusrvm"
    password  = "AccTestvdSC4daf986!"
  }
}
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name for this HDInsight ML Services Cluster. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) Specifies the name of the Resource Group in which this HDInsight ML Services Cluster should exist. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the Azure Region which this HDInsight ML Services Cluster should exist. Changing this forces a new resource to be created.
- **cluster_version** - (Required) Specifies the Version of HDInsights which should be used for this Cluster. Changing this forces a new resource to be

created.

- **gateway** - (Required) A **gateway** block as defined below.
- **roles** - (Required) A **roles** block as defined below.
- **rstudio** - (Required) Should R Studio community edition for ML Services be installed? Changing this forces a new resource to be created.
- **storage_account** - (Required) One or more **storage_account** block as defined below.
- **tier** - (Required) Specifies the Tier which should be used for this HDInsight ML Services Cluster. Possible values are **Standard** or **Premium**. Changing this forces a new resource to be created.

-
- **tags** - (Optional) A map of Tags which should be assigned to this HDInsight ML Services Cluster.
-

A **gateway** block supports the following:

- **enabled** - (Required) Is the Ambari portal enabled? Changing this forces a new resource to be created.
- **password** - (Required) The password used for the Ambari Portal. Changing this forces a new resource to be created.

NOTE: This password must be different from the one used for the **head_node**, **worker_node** and **zookeeper_node** roles.

- **username** - (Required) The username used for the Ambari Portal. Changing this forces a new resource to be created.
-

A **edge_node** block supports the following:

- **username** - (Required) The Username of the local administrator for the Edge Node. Changing this forces a new resource to be created.
- **vm_size** - (Required) The Size of the Virtual Machine which should be used as the Edge Node. Changing this forces a new resource to be created.
- **password** - (Optional) The Password associated with the local administrator for the Edge Node. Changing this forces a new resource to be created.

NOTE: If specified, this password must be at least 10 characters in length and must contain at least one digit, one uppercase and one lower case letter, one non-alphanumeric character (except characters ' " ').

- **head_node** - (Required) A **head_node** block as defined above.
- **worker_node** - (Required) A **worker_node** block as defined below.
- **zookeeper_node** - (Required) A **zookeeper_node** block as defined below.

A **storage_account** block supports the following:

- **is_default** - (Required) Is this the Default Storage Account for the HDInsight ML Services Cluster? Changing this forces a new resource to be created.

NOTE: One of the **storage_account** blocks must be marked as the default.

- **storage_account_key** - (Required) The Access Key which should be used to connect to the Storage Account. Changing this forces a new resource to be created.
- **storage_container_id** - (Required) The ID of the Storage Container. Changing this forces a new resource to be created.

NOTE: This can be obtained from the **id** of the **azurerm_storage_container** resource.

A **worker_node** block supports the following:

- **username** - (Required) The Username of the local administrator for the Worker Nodes. Changing this forces a new resource to be created.
- **vm_size** - (Required) The Size of the Virtual Machine which should be used as the Worker Nodes. Changing this forces a new resource to be created.
- **min_instance_count** - (Optional) The minimum number of instances which should be run for the Worker Nodes. Changing this forces a new resource to be created.
- **password** - (Optional) The Password associated with the local administrator for the Worker Nodes. Changing this forces a new resource to be created.

NOTE: If specified, this password must be at least 10 characters in length and must contain at least one digit, one uppercase and one lower case letter, one non-alphanumeric character (except characters ' " ').

- **ssh_keys** - (Optional) A list of SSH Keys which should be used for the local administrator on the Worker Nodes. Changing this forces a new resource to be created.

NOTE: Either a **password** or one or more **ssh_keys** must be specified - but not both.

- **subnet_id** - (Optional) The ID of the Subnet within the Virtual Network where the Worker Nodes should be provisioned within. Changing this forces a new resource to be created.
- **target_instance_count** - (Optional) The number of instances which should be run for the Worker Nodes.
- **virtual_network_id** - (Optional) The ID of the Virtual Network where the Worker Nodes should be provisioned within. Changing this forces a new resource to be created.

A **zookeeper_node** block supports the following:

- **username** - (Required) The Username of the local administrator for the Zookeeper Nodes. Changing this forces a new resource to be created.
- **vm_size** - (Required) The Size of the Virtual Machine which should be used as the Zookeeper Nodes. Changing this forces a new resource to be created.
- **password** - (Optional) The Password associated with the local administrator for the Zookeeper Nodes. Changing this forces a new resource to be created.
- **ssh_keys** - (Optional) A list of SSH Keys which should be used for the local administrator on the Zookeeper Nodes. Changing this forces a new resource to be created.

NOTE: Either a **password** or one or more **ssh_keys** must be specified - but not both.

- **subnet_id** - (Optional) The ID of the Subnet within the Virtual Network where the Zookeeper Nodes should be provisioned within. Changing this forces a new resource to be created.
- **virtual_network_id** - (Optional) The ID of the Virtual Network where the Zookeeper Nodes should be provisioned within. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the HDInsight ML Services Cluster.
- **edge_ssh_endpoint** - The SSH Connectivity Endpoint for the Edge Node of the HDInsight ML Cluster.
- **https_endpoint** - The HTTPS Connectivity Endpoint for this HDInsight ML Services Cluster.

- `ssh_endpoint` - The SSH Connectivity Endpoint for this HDInsight ML Services Cluster.

» Import

HDInsight ML Services Clusters can be imported using the `resource_id`, e.g.

```
terraform import azurerm_hdinsight_ml_services_cluster.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_hdinsight_rserver_cluster`

Manages a HDInsight RServer Cluster.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "example-resources"
  location = "West Europe"
}

resource "azurerm_storage_account" "example" {
  name                     = "hdinsightstor"
  resource_group_name     = "${azurerm_resource_group.example.name}"
  location                 = "${azurerm_resource_group.example.location}"
  account_tier             = "Standard"
  account_replication_type = "LRS"
}

resource "azurerm_storage_container" "example" {
  name                     = "hdinsight"
  resource_group_name     = "${azurerm_resource_group.example.name}"
  storage_account_name    = "${azurerm_storage_account.example.name}"
  container_access_type   = "private"
}

resource "azurerm_hdinsight_rserver_cluster" "example" {
  name                     = "example-hdicluster"
  resource_group_name     = "${azurerm_resource_group.example.name}"
  location                 = "${azurerm_resource_group.example.location}"
  cluster_version         = "3.6"
  tier                     = "Standard"
  rstudio                  = true
}
```



```

gateway {
  enabled = true
  username = "acctestusrgw"
  password = "Terraform123!"
}

storage_account {
  storage_container_id = "${azurerm_storage_container.example.id}"
  storage_account_key = "${azurerm_storage_account.example.primary_access_key}"
  is_default           = true
}

roles {
  head_node {
    vm_size = "Standard_D3_v2"
    username = "acctestusrv"
    password = "AccTestvdSC4daf986!"
  }

  worker_node {
    vm_size           = "Standard_D4_V2"
    username          = "acctestusrv"
    password           = "AccTestvdSC4daf986!"
    target_instance_count = 3
  }

  zookeeper_node {
    vm_size = "Standard_D3_v2"
    username = "acctestusrv"
    password = "AccTestvdSC4daf986!"
  }

  edge_node {
    vm_size = "Standard_D3_v2"
    username = "acctestusrv"
    password = "AccTestvdSC4daf986!"
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name for this HDInsight RServer Cluster.

Changing this forces a new resource to be created.

- **resource_group_name** - (Required) Specifies the name of the Resource Group in which this HDInsight RServer Cluster should exist. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the Azure Region which this HDInsight RServer Cluster should exist. Changing this forces a new resource to be created.
- **cluster_version** - (Required) Specifies the Version of HDInsights which should be used for this Cluster. Changing this forces a new resource to be created.
- **gateway** - (Required) A **gateway** block as defined below.
- **roles** - (Required) A **roles** block as defined below.
- **rstudio** - (Required) Should R Studio community edition for RServer be installed? Changing this forces a new resource to be created.
- **storage_account** - (Required) One or more **storage_account** block as defined below.
- **tier** - (Required) Specifies the Tier which should be used for this HDInsight RServer Cluster. Possible values are **Standard** or **Premium**. Changing this forces a new resource to be created.

-
- **tags** - (Optional) A map of Tags which should be assigned to this HDInsight RServer Cluster.
-

A **gateway** block supports the following:

- **enabled** - (Required) Is the Ambari portal enabled? Changing this forces a new resource to be created.
- **password** - (Required) The password used for the Ambari Portal. Changing this forces a new resource to be created.

NOTE: This password must be different from the one used for the **head_node**, **worker_node** and **zookeeper_node** roles.

- **username** - (Required) The username used for the Ambari Portal. Changing this forces a new resource to be created.
-

A **edge_node** block supports the following:

- **username** - (Required) The Username of the local administrator for the Edge Node. Changing this forces a new resource to be created.

- **vm_size** - (Required) The Size of the Virtual Machine which should be used as the Edge Node. Changing this forces a new resource to be created.
- **password** - (Optional) The Password associated with the local administrator for the Edge Node. Changing this forces a new resource to be created.
- **ssh_keys** - (Optional) A list of SSH Keys which should be used for the local administrator on the Edge Node. Changing this forces a new resource to be created.

NOTE: Either a `password` or one or more `ssh_keys` must be specified - but not both.

- **subnet_id** - (Optional) The ID of the Subnet within the Virtual Network where the Edge Node should be provisioned within. Changing this forces a new resource to be created.
- **virtual_network_id** - (Optional) The ID of the Virtual Network where the Edge Node should be provisioned within. Changing this forces a new resource to be created.

A `head_node` block supports the following:

- **username** - (Required) The Username of the local administrator for the Head Nodes. Changing this forces a new resource to be created.
- **vm_size** - (Required) The Size of the Virtual Machine which should be used as the Head Nodes. Changing this forces a new resource to be created.
- **password** - (Optional) The Password associated with the local administrator for the Head Nodes. Changing this forces a new resource to be created.

NOTE: If specified, this password must be at least 10 characters in length and must contain at least one digit, one uppercase and one lower case letter, one non-alphanumeric character (except characters ' " ').

- **ssh_keys** - (Optional) A list of SSH Keys which should be used for the local administrator on the Head Nodes. Changing this forces a new resource to be created.

NOTE: Either a `password` or one or more `ssh_keys` must be specified - but not both.

- **subnet_id** - (Optional) The ID of the Subnet within the Virtual Network where the Head Nodes should be provisioned within. Changing this forces a new resource to be created.
- **virtual_network_id** - (Optional) The ID of the Virtual Network where the Head Nodes should be provisioned within. Changing this forces a new resource to be created.

A **roles** block supports the following:

- **edge_node** - (Required) A **edge_node** block as defined above.
- **head_node** - (Required) A **head_node** block as defined above.
- **worker_node** - (Required) A **worker_node** block as defined below.
- **zookeeper_node** - (Required) A **zookeeper_node** block as defined below.

A **storage_account** block supports the following:

- **is_default** - (Required) Is this the Default Storage Account for the HDInsight RServer Cluster? Changing this forces a new resource to be created.

NOTE: One of the **storage_account** blocks must be marked as the default.

- **storage_account_key** - (Required) The Access Key which should be used to connect to the Storage Account. Changing this forces a new resource to be created.
- **storage_container_id** - (Required) The ID of the Storage Container. Changing this forces a new resource to be created.

NOTE: This can be obtained from the **id** of the **azurerm_storage_container** resource.

A **worker_node** block supports the following:

- **username** - (Required) The Username of the local administrator for the Worker Nodes. Changing this forces a new resource to be created.
- **vm_size** - (Required) The Size of the Virtual Machine which should be used as the Worker Nodes. Changing this forces a new resource to be created.
- **min_instance_count** - (Optional) The minimum number of instances which should be run for the Worker Nodes. Changing this forces a new resource to be created.
- **password** - (Optional) The Password associated with the local administrator for the Worker Nodes. Changing this forces a new resource to be created.

NOTE: If specified, this password must be at least 10 characters in length and must contain at least one digit, one uppercase and one lower case letter, one non-alphanumeric character (except characters ' " ').

- **ssh_keys** - (Optional) A list of SSH Keys which should be used for the local administrator on the Worker Nodes. Changing this forces a new resource to be created.

NOTE: Either a **password** or one or more **ssh_keys** must be specified - but not both.

- **subnet_id** - (Optional) The ID of the Subnet within the Virtual Network where the Worker Nodes should be provisioned within. Changing this forces a new resource to be created.
- **target_instance_count** - (Optional) The number of instances which should be run for the Worker Nodes.
- **virtual_network_id** - (Optional) The ID of the Virtual Network where the Worker Nodes should be provisioned within. Changing this forces a new resource to be created.

A **zookeeper_node** block supports the following:

- **username** - (Required) The Username of the local administrator for the Zookeeper Nodes. Changing this forces a new resource to be created.
- **vm_size** - (Required) The Size of the Virtual Machine which should be used as the Zookeeper Nodes. Changing this forces a new resource to be created.
- **password** - (Optional) The Password associated with the local administrator for the Zookeeper Nodes. Changing this forces a new resource to be created.

NOTE: If specified, this password must be at least 10 characters in length and must contain at least one digit, one uppercase and one lower case letter, one non-alphanumeric character (except characters ' " ').

- **ssh_keys** - (Optional) A list of SSH Keys which should be used for the local administrator on the Zookeeper Nodes. Changing this forces a new resource to be created.

NOTE: Either a **password** or one or more **ssh_keys** must be specified - but not both.

- **subnet_id** - (Optional) The ID of the Subnet within the Virtual Network where the Zookeeper Nodes should be provisioned within. Changing this forces a new resource to be created.
- **virtual_network_id** - (Optional) The ID of the Virtual Network where the Zookeeper Nodes should be provisioned within. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the HDInsight RServer Cluster.
- `edge_ssh_endpoint` - The SSH Connectivity Endpoint for the Edge Node of the HDInsight RServer Cluster.
- `https_endpoint` - The HTTPS Connectivity Endpoint for this HDInsight RServer Cluster.
- `ssh_endpoint` - The SSH Connectivity Endpoint for this HDInsight RServer Cluster.

» Import

HDInsight RServer Clusters can be imported using the `resource id`, e.g.

```
terraform import azurerm_hdinsight_rserver_cluster.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_hdinsight_spark_cluster`

Manages a HDInsight Spark Cluster.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "example-resources"
  location = "West Europe"
}

resource "azurerm_storage_account" "example" {
  name                     = "hdinsightstor"
  resource_group_name      = "${azurerm_resource_group.example.name}"
  location                 = "${azurerm_resource_group.example.location}"
  account_tier             = "Standard"
  account_replication_type = "LRS"
}

resource "azurerm_storage_container" "example" {
  name                     = "hdinsight"
  resource_group_name      = "${azurerm_resource_group.example.name}"
  storage_account_name     = "${azurerm_storage_account.example.name}"
  container_access_type    = "private"
}
```

```

}

resource "azurerm_hdinsight_spark_cluster" "example" {
  name                = "example-hdicluster"
  resource_group_name = "${azurerm_resource_group.example.name}"
  location            = "${azurerm_resource_group.example.location}"
  cluster_version     = "3.6"
  tier                 = "Standard"

  component_version {
    spark = "2.3"
  }

  gateway {
    enabled = true
    username = "acctestusrgw"
    password = "Terraform123!"
  }

  storage_account {
    storage_container_id = "${azurerm_storage_container.example.id}"
    storage_account_key  = "${azurerm_storage_account.example.primary_access_key}"
    is_default            = true
  }

  roles {
    head_node {
      vm_size = "Standard_A3"
      username = "acctestusrvvm"
      password = "AccTestvdSC4daf986!"
    }

    worker_node {
      vm_size           = "Standard_A3"
      username          = "acctestusrvvm"
      password          = "AccTestvdSC4daf986!"
      target_instance_count = 3
    }

    zookeeper_node {
      vm_size = "Medium"
      username = "acctestusrvvm"
      password = "AccTestvdSC4daf986!"
    }
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name for this HDInsight Spark Cluster. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) Specifies the name of the Resource Group in which this HDInsight Spark Cluster should exist. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the Azure Region which this HDInsight Spark Cluster should exist. Changing this forces a new resource to be created.
- **cluster_version** - (Required) Specifies the Version of HDInsights which should be used for this Cluster. Changing this forces a new resource to be created.
- **component_version** - (Required) A **component_version** block as defined below.
- **gateway** - (Required) A **gateway** block as defined below.
- **roles** - (Required) A **roles** block as defined below.
- **storage_account** - (Required) One or more **storage_account** block as defined below.
- **tier** - (Required) Specifies the Tier which should be used for this HDInsight Spark Cluster. Possible values are **Standard** or **Premium**. Changing this forces a new resource to be created.

-
- **tags** - (Optional) A map of Tags which should be assigned to this HDInsight Spark Cluster.
-

A **component_version** block supports the following:

- **spark** - (Required) The version of Spark which should be used for this HDInsight Spark Cluster. Changing this forces a new resource to be created.
-

A **gateway** block supports the following:

- **enabled** - (Required) Is the Ambari portal enabled? Changing this forces a new resource to be created.

- **password** - (Required) The password used for the Ambari Portal. Changing this forces a new resource to be created.

NOTE: This password must be different from the one used for the **head_node**, **worker_node** and **zookeeper_node** roles.

- **username** - (Required) The username used for the Ambari Portal. Changing this forces a new resource to be created.

A **head_node** block supports the following:

- **username** - (Required) The Username of the local administrator for the Head Nodes. Changing this forces a new resource to be created.
- **vm_size** - (Required) The Size of the Virtual Machine which should be used as the Head Nodes. Changing this forces a new resource to be created.
- **password** - (Optional) The Password associated with the local administrator for the Head Nodes. Changing this forces a new resource to be created.

NOTE: If specified, this password must be at least 10 characters in length and must contain at least one digit, one uppercase and one lower case letter, one non-alphanumeric character (except characters ' " ').

- **ssh_keys** - (Optional) A list of SSH Keys which should be used for the local administrator on the Head Nodes. Changing this forces a new resource to be created.

NOTE: Either a **password** or one or more **ssh_keys** must be specified - but not both.

- **subnet_id** - (Optional) The ID of the Subnet within the Virtual Network where the Head Nodes should be provisioned within. Changing this forces a new resource to be created.
- **virtual_network_id** - (Optional) The ID of the Virtual Network where the Head Nodes should be provisioned within. Changing this forces a new resource to be created.

A **roles** block supports the following:

- **head_node** - (Required) A **head_node** block as defined above.
- **worker_node** - (Required) A **worker_node** block as defined below.
- **zookeeper_node** - (Required) A **zookeeper_node** block as defined below.

A **storage_account** block supports the following:

- **is_default** - (Required) Is this the Default Storage Account for the HDInsight Spark Cluster? Changing this forces a new resource to be created.

NOTE: One of the **storage_account** blocks must be marked as the default.

- **storage_account_key** - (Required) The Access Key which should be used to connect to the Storage Account. Changing this forces a new resource to be created.
- **storage_container_id** - (Required) The ID of the Storage Container. Changing this forces a new resource to be created.

NOTE: This can be obtained from the **id** of the **azurerm_storage_container** resource.

A **worker_node** block supports the following:

- **username** - (Required) The Username of the local administrator for the Worker Nodes. Changing this forces a new resource to be created.
- **vm_size** - (Required) The Size of the Virtual Machine which should be used as the Worker Nodes. Changing this forces a new resource to be created.
- **min_instance_count** - (Optional) The minimum number of instances which should be run for the Worker Nodes. Changing this forces a new resource to be created.
- **password** - (Optional) The Password associated with the local administrator for the Worker Nodes. Changing this forces a new resource to be created.

NOTE: If specified, this password must be at least 10 characters in length and must contain at least one digit, one uppercase and one lower case letter, one non-alphanumeric character (except characters ' " ').

- **ssh_keys** - (Optional) A list of SSH Keys which should be used for the local administrator on the Worker Nodes. Changing this forces a new resource to be created.

NOTE: Either a **password** or one or more **ssh_keys** must be specified - but not both.

- **subnet_id** - (Optional) The ID of the Subnet within the Virtual Network where the Worker Nodes should be provisioned within. Changing this forces a new resource to be created.
- **target_instance_count** - (Optional) The number of instances which should be run for the Worker Nodes.

- **virtual_network_id** - (Optional) The ID of the Virtual Network where the Worker Nodes should be provisioned within. Changing this forces a new resource to be created.

A **zookeeper_node** block supports the following:

- **username** - (Required) The Username of the local administrator for the Zookeeper Nodes. Changing this forces a new resource to be created.
- **vm_size** - (Required) The Size of the Virtual Machine which should be used as the Zookeeper Nodes. Changing this forces a new resource to be created.
- **password** - (Optional) The Password associated with the local administrator for the Zookeeper Nodes. Changing this forces a new resource to be created.

NOTE: If specified, this password must be at least 10 characters in length and must contain at least one digit, one uppercase and one lower case letter, one non-alphanumeric character (except characters ' " ' ').

- **ssh_keys** - (Optional) A list of SSH Keys which should be used for the local administrator on the Zookeeper Nodes. Changing this forces a new resource to be created.

NOTE: Either a **password** or one or more **ssh_keys** must be specified - but not both.

- **subnet_id** - (Optional) The ID of the Subnet within the Virtual Network where the Zookeeper Nodes should be provisioned within. Changing this forces a new resource to be created.
- **virtual_network_id** - (Optional) The ID of the Virtual Network where the Zookeeper Nodes should be provisioned within. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the HDInsight Spark Cluster.
- **https_endpoint** - The HTTPS Connectivity Endpoint for this HDInsight Spark Cluster.
- **ssh_endpoint** - The SSH Connectivity Endpoint for this HDInsight Spark Cluster.

» Import

HDInsight Spark Clusters can be imported using the `resource id`, e.g.

```
terraform import azurerm_hdinsight_spark_cluster.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_hdinsight_storm_cluster`

Manages a HDInsight Storm Cluster.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "example-resources"
  location = "West Europe"
}

resource "azurerm_storage_account" "example" {
  name                        = "hdinsightstor"
  resource_group_name        = "${azurerm_resource_group.example.name}"
  location                   = "${azurerm_resource_group.example.location}"
  account_tier                = "Standard"
  account_replication_type    = "LRS"
}

resource "azurerm_storage_container" "example" {
  name                = "hdinsight"
  resource_group_name = "${azurerm_resource_group.example.name}"
  storage_account_name = "${azurerm_storage_account.example.name}"
  container_access_type = "private"
}

resource "azurerm_hdinsight_storm_cluster" "example" {
  name                = "example-hdicluster"
  resource_group_name = "${azurerm_resource_group.example.name}"
  location            = "${azurerm_resource_group.example.location}"
  cluster_version     = "3.6"
  tier                 = "Standard"

  component_version {
    storm = "1.1"
  }

  gateway {
```

```

    enabled = true
    username = "acctestusrgw"
    password = "Terraform123!"
  }

  storage_account {
    storage_container_id = "${azurerm_storage_container.example.id}"
    storage_account_key   = "${azurerm_storage_account.example.primary_access_key}"
    is_default             = true
  }

  roles {
    head_node {
      vm_size = "Standard_A3"
      username = "acctestusrvvm"
      password = "AccTestvdSC4daf986!"
    }

    worker_node {
      vm_size           = "Standard_D3_V2"
      username          = "acctestusrvvm"
      password          = "AccTestvdSC4daf986!"
      target_instance_count = 3
    }

    zookeeper_node {
      vm_size = "Standard_A4_V2"
      username = "acctestusrvvm"
      password = "AccTestvdSC4daf986!"
    }
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name for this HDInsight Storm Cluster. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) Specifies the name of the Resource Group in which this HDInsight Storm Cluster should exist. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the Azure Region which this HDInsight Storm Cluster should exist. Changing this forces a new resource to be

created.

- **cluster_version** - (Required) Specifies the Version of HDInsights which should be used for this Cluster. Changing this forces a new resource to be created.
- **component_version** - (Required) A **component_version** block as defined below.
- **gateway** - (Required) A **gateway** block as defined below.
- **roles** - (Required) A **roles** block as defined below.
- **storage_account** - (Required) One or more **storage_account** block as defined below.
- **tier** - (Required) Specifies the Tier which should be used for this HDInsight Storm Cluster. Possible values are **Standard** or **Premium**. Changing this forces a new resource to be created.

-
- **tags** - (Optional) A map of Tags which should be assigned to this HDInsight Storm Cluster.
-

A **component_version** block supports the following:

- **storm** - (Required) The version of Storm which should be used for this HDInsight Storm Cluster. Changing this forces a new resource to be created.
-

A **gateway** block supports the following:

- **enabled** - (Required) Is the Ambari portal enabled? Changing this forces a new resource to be created.
- **password** - (Required) The password used for the Ambari Portal. Changing this forces a new resource to be created.

NOTE: This password must be different from the one used for the **head_node**, **worker_node** and **zookeeper_node** roles.

- **username** - (Required) The username used for the Ambari Portal. Changing this forces a new resource to be created.
-

A **head_node** block supports the following:

NOTE: This is also known as the **nimbus** node.

- **username** - (Required) The Username of the local administrator for the Head Nodes. Changing this forces a new resource to be created.
- **vm_size** - (Required) The Size of the Virtual Machine which should be used as the Head Nodes. Changing this forces a new resource to be created.
- **password** - (Optional) The Password associated with the local administrator for the Head Nodes. Changing this forces a new resource to be created.

NOTE: If specified, this password must be at least 10 characters in length and must contain at least one digit, one uppercase and one lower case letter, one non-alphanumeric character (except characters ' " ').

- **ssh_keys** - (Optional) A list of SSH Keys which should be used for the local administrator on the Head Nodes. Changing this forces a new resource to be created.

NOTE: Either a **password** or one or more **ssh_keys** must be specified - but not both.

- **subnet_id** - (Optional) The ID of the Subnet within the Virtual Network where the Head Nodes should be provisioned within. Changing this forces a new resource to be created.
- **virtual_network_id** - (Optional) The ID of the Virtual Network where the Head Nodes should be provisioned within. Changing this forces a new resource to be created.

A **roles** block supports the following:

- **head_node** - (Required) A **head_node** block as defined above.
- **worker_node** - (Required) A **worker_node** block as defined below.
- **zookeeper_node** - (Required) A **zookeeper_node** block as defined below.

A **storage_account** block supports the following:

- **is_default** - (Required) Is this the Default Storage Account for the HDInsight Storm Cluster? Changing this forces a new resource to be created.

NOTE: One of the **storage_account** blocks must be marked as the default.

- **storage_account_key** - (Required) The Access Key which should be used to connect to the Storage Account. Changing this forces a new resource to be created.
- **storage_container_id** - (Required) The ID of the Storage Container. Changing this forces a new resource to be created.

NOTE: This can be obtained from the `id` of the `azurerm_storage_container` resource.

A `worker_node` block supports the following:

NOTE: This is also known as the `supervisor` node.

- `username` - (Required) The Username of the local administrator for the Worker Nodes. Changing this forces a new resource to be created.
- `vm_size` - (Required) The Size of the Virtual Machine which should be used as the Worker Nodes. Changing this forces a new resource to be created.
- `min_instance_count` - (Optional) The minimum number of instances which should be run for the Worker Nodes. Changing this forces a new resource to be created.
- `password` - (Optional) The Password associated with the local administrator for the Worker Nodes. Changing this forces a new resource to be created.

NOTE: If specified, this password must be at least 10 characters in length and must contain at least one digit, one uppercase and one lower case letter, one non-alphanumeric character (except characters `'` `"` `'` `'` `'`).

- `ssh_keys` - (Optional) A list of SSH Keys which should be used for the local administrator on the Worker Nodes. Changing this forces a new resource to be created.

NOTE: Either a `password` or one or more `ssh_keys` must be specified - but not both.

- `subnet_id` - (Optional) The ID of the Subnet within the Virtual Network where the Worker Nodes should be provisioned within. Changing this forces a new resource to be created.
- `target_instance_count` - (Optional) The number of instances which should be run for the Worker Nodes.
- `virtual_network_id` - (Optional) The ID of the Virtual Network where the Worker Nodes should be provisioned within. Changing this forces a new resource to be created.

A `zookeeper_node` block supports the following:

- `username` - (Required) The Username of the local administrator for the Zookeeper Nodes. Changing this forces a new resource to be created.

- **vm_size** - (Required) The Size of the Virtual Machine which should be used as the Zookeeper Nodes. Changing this forces a new resource to be created.
- **password** - (Optional) The Password associated with the local administrator for the Zookeeper Nodes. Changing this forces a new resource to be created.

NOTE: If specified, this password must be at least 10 characters in length and must contain at least one digit, one uppercase and one lower case letter, one non-alphanumeric character (except characters ' " ').

- **ssh_keys** - (Optional) A list of SSH Keys which should be used for the local administrator on the Zookeeper Nodes. Changing this forces a new resource to be created.

NOTE: Either a **password** or one or more **ssh_keys** must be specified - but not both.

- **subnet_id** - (Optional) The ID of the Subnet within the Virtual Network where the Zookeeper Nodes should be provisioned within. Changing this forces a new resource to be created.
- **virtual_network_id** - (Optional) The ID of the Virtual Network where the Zookeeper Nodes should be provisioned within. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the HDInsight Storm Cluster.
- **https_endpoint** - The HTTPS Connectivity Endpoint for this HDInsight Storm Cluster.
- **ssh_endpoint** - The SSH Connectivity Endpoint for this HDInsight Storm Cluster.

» Import

HDInsight Storm Clusters can be imported using the **resource id**, e.g.

```
terraform import azurerm_hdinsight_storm_cluster.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm__key__vault

Manages a Key Vault.

NOTE: It's possible to define Key Vault Access Policies both within the `azurerm_key_vault` resource via the `access_policy` block and by using the `azurerm_key_vault_access_policy` resource. However it's not possible to use both methods to manage Access Policies within a KeyVault, since there'll be conflicts.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name       = "resourceGroup1"
  location   = "West US"
}

resource "azurerm_key_vault" "test" {
  name                        = "testvault"
  location                   = "${azurerm_resource_group.test.location}"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  enabled_for_disk_encryption = true
  tenant_id                  = "d6e396d0-5584-41dc-9fc0-268df99bc610"

  sku {
    name = "standard"
  }

  access_policy {
    tenant_id = "d6e396d0-5584-41dc-9fc0-268df99bc610"
    object_id = "d746815a-0433-4a21-b95d-fc437d2d475b"

    key_permissions = [
      "get",
    ]

    secret_permissions = [
      "get",
    ]

    storage_permissions = [
      "get",
    ]
  }
}
```

```

network_acls {
  default_action = "Deny"
  bypass        = "AzureServices"
}

tags = {
  environment = "Production"
}
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Key Vault. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the Key Vault. Changing this forces a new resource to be created.
- **sku** - (Required) An SKU block as described below.
- **tenant_id** - (Required) The Azure Active Directory tenant ID that should be used for authenticating requests to the key vault.
- **access_policy** - (Optional) An access policy block as described below. A maximum of 16 may be declared.

NOTE: It's possible to define Key Vault Access Policies both within the `azurerm_key_vault` resource via the `access_policy` block and by using the `azurerm_key_vault_access_policy` resource. However it's not possible to use both methods to manage Access Policies within a KeyVault, since there'll be conflicts.

- **enabled_for_deployment** - (Optional) Boolean flag to specify whether Azure Virtual Machines are permitted to retrieve certificates stored as secrets from the key vault. Defaults to **false**.
- **enabled_for_disk_encryption** - (Optional) Boolean flag to specify whether Azure Disk Encryption is permitted to retrieve secrets from the vault and unwrap keys. Defaults to **false**.
- **enabled_for_template_deployment** - (Optional) Boolean flag to specify whether Azure Resource Manager is permitted to retrieve secrets from the key vault. Defaults to **false**.

- **network_acls** - (Optional) A **network_acls** block as defined below.
- **tags** - (Optional) A mapping of tags to assign to the resource.

A **access_policy** block supports the following:

- **tenant_id** - (Required) The Azure Active Directory tenant ID that should be used for authenticating requests to the key vault. Must match the **tenant_id** used above.
- **object_id** - (Required) The object ID of a user, service principal or security group in the Azure Active Directory tenant for the vault. The object ID must be unique for the list of access policies.
- **application_id** - (Optional) The object ID of an Application in Azure Active Directory.
- **certificate_permissions** - (Optional) List of certificate permissions, must be one or more from the following: **backup**, **create**, **delete**, **deleteissuers**, **get**, **getissuers**, **import**, **list**, **listissuers**, **managecontacts**, **manageissuers**, **purge**, **recover**, **restore**, **setissuers** and **update**.
- **key_permissions** - (Optional) List of key permissions, must be one or more from the following: **backup**, **create**, **decrypt**, **delete**, **encrypt**, **get**, **import**, **list**, **purge**, **recover**, **restore**, **sign**, **unwrapKey**, **update**, **verify** and **wrapKey**.
- **secret_permissions** - (Optional) List of secret permissions, must be one or more from the following: **backup**, **delete**, **get**, **list**, **purge**, **recover**, **restore** and **set**.
- **storage_permissions** - (Optional) List of storage permissions, must be one or more from the following: **backup**, **delete**, **deletesas**, **get**, **getsas**, **list**, **listsas**, **purge**, **recover**, **regeneratekey**, **restore**, **set**, **setsas** and **update**.

A **network_acls** block supports the following:

- **bypass** - (Required) Specifies which traffic can bypass the network rules. Possible values are **AzureServices** and **None**.
- **default_action** - (Required) The Default Action to use when no rules match from **ip_rules** / **virtual_network_subnet_ids**. Possible values are **Allow** and **Deny**.
- **ip_rules** - (Optional) One or more IP Addresses, or CIDR Blocks which should be able to access this Key Vault.

- `virtual_network_subnet_ids` - (Optional) One or more Subnet ID's which should be able to access this Key Vault.

A `sku` block supports the following:

- `name` - (Required) The Name of the SKU used for this Key Vault. Possible values are `standard` and `premium`.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the Key Vault.
- `vault_uri` - The URI of the Key Vault, used for performing operations on keys and secrets.

» Import

Key Vault's can be imported using the `resource id`, e.g.

```
terraform import azurerm_key_vault.test /subscriptions/00000000-0000-0000-0000-000000000000/
```

» `azurerm_key_vault_access_policy`

Manages a Key Vault Access Policy.

NOTE: It's possible to define Key Vault Access Policies both within the `azurerm_key_vault` resource via the `access_policy` block and by using the `azurerm_key_vault_access_policy` resource. However it's not possible to use both methods to manage Access Policies within a KeyVault, since there'll be conflicts.

NOTE: Azure permits a maximum of 1024 Access Policies per Key Vault - more information can be found in this document.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "resourceGroup1"
  location = "${azurerm_resource_group.test.location}"
}
```

```

resource "azurerm_key_vault" "test" {
  name           = "testvault"
  location       = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  sku {
    name = "standard"
  }

  tenant_id = "22222222-2222-2222-2222-222222222222"

  enabled_for_disk_encryption = true

  tags = {
    environment = "Production"
  }
}

resource "azurerm_key_vault_access_policy" "test" {
  vault_name           = "${azurerm_key_vault.test.name}"
  resource_group_name = "${azurerm_key_vault.test.resource_group_name}"

  tenant_id = "00000000-0000-0000-0000-000000000000"
  object_id = "11111111-1111-1111-1111-111111111111"

  key_permissions = [
    "get",
  ]

  secret_permissions = [
    "get",
  ]
}

```

» Argument Reference

The following arguments are supported:

- **vault_name** - (Required) Specifies the name of the Key Vault resource. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the namespace. Changing this forces a new resource to be created.
- **tenant_id** - (Required) The Azure Active Directory tenant ID that should

be used for authenticating requests to the key vault. Changing this forces a new resource to be created.

- **object_id** - (Required) The object ID of a user, service principal or security group in the Azure Active Directory tenant for the vault. The object ID must be unique for the list of access policies. Changing this forces a new resource to be created.
- **application_id** - (Optional) The object ID of an Application in Azure Active Directory.
- **certificate_permissions** - (Optional) List of certificate permissions, must be one or more from the following: **backup**, **create**, **delete**, **deleteissuers**, **get**, **getissuers**, **import**, **list**, **listissuers**, **managecontacts**, **manageissuers**, **purge**, **recover**, **restore**, **setissuers** and **update**.
- **key_permissions** - (Required) List of key permissions, must be one or more from the following: **backup**, **create**, **decrypt**, **delete**, **encrypt**, **get**, **import**, **list**, **purge**, **recover**, **restore**, **sign**, **unwrapKey**, **update**, **verify** and **wrapKey**.
- **secret_permissions** - (Required) List of secret permissions, must be one or more from the following: **backup**, **delete**, **get**, **list**, **purge**, **recover**, **restore** and **set**.
- **storage_permissions** - (Optional) List of storage permissions, must be one or more from the following: **backup**, **delete**, **deletesas**, **get**, **getsas**, **list**, **listsas**, **purge**, **recover**, **regeneratekey**, **restore**, **set**, **setsas** and **update**.

» Attributes Reference

The following attributes are exported:

- **id** - Key Vault Access Policy ID.

NOTE: This Identifier is unique to Terraform and doesn't map to an existing object within Azure.

» Import

Key Vault Access Policies can be imported using the Resource ID of the Key Vault, plus some additional metadata.

If both an **object_id** and **application_id** are specified, then the Access Policy can be imported using the following code:

```
terraform import azurerm_key_vault_access_policy.test /subscriptions/00000000-0000-0000-0000-000000000000/object_id/application_id
```

where 11111111-1111-1111-1111-111111111111 is the `object_id` and 22222222-2222-2222-2222-222222222222 is the `application_id`.

Access Policies with an `object_id` but no `application_id` can be imported using the following command:

```
terraform import azurerm_key_vault_access_policy.test /subscriptions/00000000-0000-0000-0000-000000000000
```

where 11111111-1111-1111-1111-111111111111 is the `object_id`.

NOTE: Both Identifiers are unique to Terraform and don't map to an existing object within Azure.

» `azurerm_key_vault_certificate`

Manages a Key Vault Certificate.

» Example Usage (Importing a PFX)

Note: this example assumed the PFX file is located in the same directory at `certificate-to-import.pfx`.

```
data "azurerm_client_config" "current" {}

resource "azurerm_resource_group" "test" {
  name      = "key-vault-certificate-example"
  location  = "West Europe"
}

resource "azurerm_key_vault" "test" {
  name                = "keyvaultcertexample"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  tenant_id          = "${data.azurerm_client_config.current.tenant_id}"

  sku {
    name = "standard"
  }

  access_policy {
    tenant_id = "${data.azurerm_client_config.current.tenant_id}"
    object_id = "${data.azurerm_client_config.current.service_principal_object_id}"

    certificate_permissions = [
```



```

        "create",
        "delete",
        "deleteissuers",
        "get",
        "getissuers",
        "import",
        "list",
        "listissuers",
        "managecontacts",
        "manageissuers",
        "setissuers",
        "update",
    ]

    key_permissions = [
        "backup",
        "create",
        "decrypt",
        "delete",
        "encrypt",
        "get",
        "import",
        "list",
        "purge",
        "recover",
        "restore",
        "sign",
        "unwrapKey",
        "update",
        "verify",
        "wrapKey",
    ]

    secret_permissions = [
        "backup",
        "delete",
        "get",
        "list",
        "purge",
        "recover",
        "restore",
        "set",
    ]
}

tags = {

```

```

        environment = "Production"
    }
}

resource "azurerm_key_vault_certificate" "test" {
    name      = "imported-cert"
    key_vault_id = "${azurerm_key_vault.test.id}"

    certificate {
        contents = "${base64encode(file("certificate-to-import.pfx"))}"
        password = ""
    }

    certificate_policy {
        issuer_parameters {
            name = "Self"
        }
    }

    key_properties {
        exportable = true
        key_size   = 2048
        key_type   = "RSA"
        reuse_key  = false
    }

    secret_properties {
        content_type = "application/x-pkcs12"
    }
}

```

» Example Usage (Generating a new certificate)

```

data "azurerm_client_config" "current" {}

resource "azurerm_resource_group" "test" {
    name      = "key-vault-certificate-example"
    location = "West Europe"
}

resource "azurerm_key_vault" "test" {
    name                = "keyvaultcertexample"
    location            = "${azurerm_resource_group.test.location}"
    resource_group_name = "${azurerm_resource_group.test.name}"
    tenant_id          = "${data.azurerm_client_config.current.tenant_id}"
}

```

```

sku {
  name = "standard"
}

access_policy {
  tenant_id = "${data.azure_rm_client_config.current.tenant_id}"
  object_id = "${data.azure_rm_client_config.current.service_principal_object_id}"

  certificate_permissions = [
    "create", "delete", "deleteissuers",
    "get", "getissuers", "import", "list",
    "listissuers", "managecontacts", "manageissuers",
    "setissuers", "update",
  ]

  key_permissions = [
    "backup", "create", "decrypt", "delete", "encrypt", "get",
    "import", "list", "purge", "recover", "restore", "sign",
    "unwrapKey", "update", "verify", "wrapKey",
  ]

  secret_permissions = [
    "backup", "delete", "get", "list", "purge", "recover", "restore", "set",
  ]
}

tags = {
  environment = "Production"
}
}

resource "azurerm_key_vault_certificate" "test" {
  name          = "generated-cert"
  key_vault_id = "${azurerm_key_vault.test.id}"

  certificate_policy {
    issuer_parameters {
      name = "Self"
    }
  }

  key_properties {
    exportable = true
    key_size   = 2048
    key_type   = "RSA"
    reuse_key  = true
  }
}

```

```

    }

    lifetime_action {
        action {
            action_type = "AutoRenew"
        }

        trigger {
            days_before_expiry = 30
        }
    }

    secret_properties {
        content_type = "application/x-pkcs12"
    }

    x509_certificate_properties {
        # Server Authentication = 1.3.6.1.5.5.7.3.1
        # Client Authentication = 1.3.6.1.5.5.7.3.2
        extended_key_usage = ["1.3.6.1.5.5.7.3.1"]

        key_usage = [
            "cRLSign",
            "dataEncipherment",
            "digitalSignature",
            "keyAgreement",
            "keyCertSign",
            "keyEncipherment",
        ]

        subject_alternative_names {
            dns_names = ["internal.contoso.com", "domain.hello.world"]
        }

        subject = "CN=hello-world"
        validity_in_months = 12
    }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Key Vault Certificate. Chang-

ing this forces a new resource to be created.

- **key_vault_id** - (Required) The ID of the Key Vault where the Certificate should be created.
- **certificate** - (Optional) A **certificate** block as defined below, used to Import an existing certificate.
- **certificate_policy** - (Required) A **certificate_policy** block as defined below.
- **tags** - (Optional) A mapping of tags to assign to the resource.

certificate supports the following:

- **contents** - (Required) The base64-encoded certificate contents. Changing this forces a new resource to be created.
- **password** - (Optional) The password associated with the certificate. Changing this forces a new resource to be created.

certificate_policy supports the following:

- **issuer_parameters** - (Required) A **issuer_parameters** block as defined below.
- **key_properties** - (Required) A **key_properties** block as defined below.
- **lifetime_action** - (Optional) A **lifetime_action** block as defined below.
- **secret_properties** - (Required) A **secret_properties** block as defined below.
- **x509_certificate_properties** - (Optional) A **x509_certificate_properties** block as defined below.

issuer_parameters supports the following:

- **name** - (Required) The name of the Certificate Issuer. Possible values include **Self**, or the name of a certificate issuing authority supported by Azure. Changing this forces a new resource to be created.

key_properties supports the following:

- **exportable** - (Required) Is this Certificate Exportable? Changing this forces a new resource to be created.
- **key_size** - (Required) The size of the Key used in the Certificate. Possible values include 2048 and 4096. Changing this forces a new resource to be created.
- **key_type** - (Required) Specifies the Type of Key, such as **RSA**. Changing this forces a new resource to be created.
- **reuse_key** - (Required) Is the key reusable? Changing this forces a new resource to be created.

lifetime_action supports the following:

- **action** - (Required) A **action** block as defined below.
- **trigger** - (Required) A **trigger** block as defined below.

action supports the following:

- **action_type** - (Required) The Type of action to be performed when the lifetime trigger is triggered. Possible values include **AutoRenew** and **EmailContacts**. Changing this forces a new resource to be created.

trigger supports the following:

- **days_before_expiry** - (Optional) The number of days before the Certificate expires that the action associated with this Trigger should run. Changing this forces a new resource to be created. Conflicts with **lifetime_percentage**.
- **lifetime_percentage** - (Optional) The percentage at which during the Certificate's Lifetime the action associated with this Trigger should run. Changing this forces a new resource to be created. Conflicts with **days_before_expiry**.

secret_properties supports the following:

- **content_type** - (Required) The Content-Type of the Certificate, such as **application/x-pkcs12** for a PFX or **application/x-pem-file** for a PEM. Changing this forces a new resource to be created.

x509_certificate_properties supports the following:

- **extended_key_usage** - (Optional) A list of Extended/Enhanced Key Usages. Changing this forces a new resource to be created.
- **key_usage** - (Required) A list of uses associated with this Key. Possible values include **cRLSign**, **dataEncipherment**, **decipherOnly**, **digitalSignature**, **encipherOnly**, **keyAgreement**, **keyCertSign**, **keyEncipherment** and **nonRepudiation** and are case-sensitive. Changing this forces a new resource to be created.
- **subject** - (Required) The Certificate's Subject. Changing this forces a new resource to be created.
- **subject_alternative_names** - (Optional) A **subject_alternative_names** block as defined below.
- **validity_in_months** - (Required) The Certificate's Validity Period in Months. Changing this forces a new resource to be created.

subject_alternative_names supports the following:

- **dns_names** - (Optional) A list of alternative DNS names (FQDNs) identified by the Certificate. Changing this forces a new resource to be created.
- **emails** - (Optional) A list of email addresses identified by this Certificate. Changing this forces a new resource to be created.
- **upns** - (Optional) A list of User Principal Names identified by the Certificate. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- `id` - The Key Vault Certificate ID.
- `secret_id` - The ID of the associated Key Vault Secret.
- `version` - The current version of the Key Vault Certificate.
- `certificate_data` - The raw Key Vault Certificate.
- `thumbprint` - The X509 Thumbprint of the Key Vault Certificate returned as hex string.

» Import

Key Vault Certificates can be imported using the `resource id`, e.g.

```
terraform import azurerm_key_vault_certificate.test https://example-keyvault.vault.azure.net
```

» `azurerm_key_vault_key`

Manages a Key Vault Key.

» Example Usage

```
data "azurerm_client_config" "current" {}

resource "azurerm_resource_group" "test" {
  name     = "my-resource-group"
  location = "West US"
}

resource "random_id" "server" {
  keepers = {
    ami_id = 1
  }

  byte_length = 8
}

resource "azurerm_key_vault" "test" {
  name                = "keyvaultkeyexample"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  tenant_id          = "${data.azurerm_client_config.current.tenant_id}"
}
```

```

sku {
    name = "premium"
}

access_policy {
    tenant_id = "${data.azure_rm_client_config.current.tenant_id}"
    object_id = "${data.azure_rm_client_config.current.service_principal_object_id}"

    key_permissions = [
        "create",
        "get",
    ]

    secret_permissions = [
        "set",
    ]
}

tags = {
    environment = "Production"
}
}

resource "azurerm_key_vault_key" "generated" {
    name          = "generated-certificate"
    key_vault_id = "${azurerm_key_vault.test.id}"
    key_type      = "RSA"
    key_size      = 2048

    key_opts = [
        "decrypt",
        "encrypt",
        "sign",
        "unwrapKey",
        "verify",
        "wrapKey",
    ]
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Key Vault Key. Changing

this forces a new resource to be created.

- **key_vault_id** - (Required) The ID of the Key Vault where the Key should be created.
- **key_type** - (Required) Specifies the Key Type to use for this Key Vault Key. Possible values are **EC** (Elliptic Curve), **Oct** (Octet), **RSA** and **RSA-HSM**. Changing this forces a new resource to be created.
- **key_size** - (Required) Specifies the Size of the Key to create in bytes. For example, 1024 or 2048. Changing this forces a new resource to be created.
- **key_opts** - (Required) A list of JSON web key operations. Possible values include: **decrypt**, **encrypt**, **sign**, **unwrapKey**, **verify** and **wrapKey**. Please note these values are case sensitive.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- **id** - The Key Vault Key ID.
- **version** - The current version of the Key Vault Key.
- **n** - The RSA modulus of this Key Vault Key.
- **e** - The RSA public exponent of this Key Vault Key.

» Import

Key Vault Key which is Enabled can be imported using the **resource id**, e.g.

```
terraform import azurerm_key_vault_key.test https://example-keyvault.vault.azure.net/keys/example-key
```

» azurerm__key__vault__secret

Manages a Key Vault Secret.

Note: All arguments including the secret value will be stored in the raw state as plain-text. Read more about sensitive data in state.

» Example Usage

```
data "azurerm_client_config" "current" {}

resource "azurerm_resource_group" "test" {
```

```

    name      = "my-resource-group"
    location  = "West US"
}

resource "random_id" "server" {
  keepers = {
    ami_id = 1
  }

  byte_length = 8
}

resource "azurerm_key_vault" "test" {
  name                = "${format("%s%s", "kv", random_id.server.hex)}"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  tenant_id           = "${data.azurerm_client_config.current.tenant_id}"

  sku {
    name = "premium"
  }

  access_policy {
    tenant_id = "${data.azurerm_client_config.current.tenant_id}"
    object_id = "${data.azurerm_client_config.current.service_principal_object_id}"

    key_permissions = [
      "create",
      "get",
    ]

    secret_permissions = [
      "set",
      "get",
      "delete",
    ]
  }

  tags = {
    environment = "Production"
  }
}

resource "azurerm_key_vault_secret" "test" {
  name      = "secret-sauce"
  value     = "szechuan"
}

```

```

key_vault_id = "${azurerm_key_vault.test.id}"

tags = {
  environment = "Production"
}
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Key Vault Secret. Changing this forces a new resource to be created.
- **value** - (Required) Specifies the value of the Key Vault Secret.
- **key_vault_id** - (Required) The ID of the Key Vault where the Secret should be created.
- **content_type** - (Optional) Specifies the content type for the Key Vault Secret.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- **id** - The Key Vault Secret ID.
- **version** - The current version of the Key Vault Secret.

» Import

Key Vault Secrets which are Enabled can be imported using the **resource id**, e.g.

```
terraform import azurerm_key_vault_secret.test https://example-keyvault.vault.azure.net/secrets
```

» azurerm_lb

Manage a Load Balancer Resource.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "LoadBalancerRG"
  location  = "West US"
}

resource "azurerm_public_ip" "test" {
  name            = "PublicIPForLB"
  location        = "West US"
  resource_group_name = "${azurerm_resource_group.test.name}"
  allocation_method = "Static"
}

resource "azurerm_lb" "test" {
  name            = "TestLoadBalancer"
  location        = "West US"
  resource_group_name = "${azurerm_resource_group.test.name}"

  frontend_ip_configuration {
    name            = "PublicIPAddress"
    public_ip_address_id = "${azurerm_public_ip.test.id}"
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Load Balancer.
- **resource_group_name** - (Required) The name of the Resource Group in which to create the Load Balancer.
- **location** - (Required) Specifies the supported Azure Region where the Load Balancer should be created.
- **frontend_ip_configuration** - (Optional) A `frontend_ip_configuration` block as documented below.
- **sku** - (Optional) The SKU of the Azure Load Balancer. Accepted values are `Basic` and `Standard`. Defaults to `Basic`.
- **tags** - (Optional) A mapping of tags to assign to the resource.

`frontend_ip_configuration` supports the following:

- **name** - (Required) Specifies the name of the frontend ip configuration.

- **subnet_id** - The ID of the Subnet which should be associated with the IP Configuration.
- **private_ip_address** - (Optional) Private IP Address to assign to the Load Balancer. The last one and first four IPs in any range are reserved and cannot be manually assigned.
- **private_ip_address_allocation** - (Optional) The allocation method for the Private IP Address used by this Load Balancer. Possible values as **Dynamic** and **Static**.
- **public_ip_address_id** - (Optional) Th ID of a Public IP Address which should be associated with the Load Balancer.
- **zones** - (Optional) A list of Availability Zones which the Load Balancer's IP Addresses should be created in.

Please Note: Availability Zones are only supported in several regions at this time.

» Attributes Reference

The following attributes are exported:

- **id** - The Load Balancer ID.
- **private_ip_address** - The first private IP address assigned to the load balancer in **frontend_ip_configuration** blocks, if any.
- **private_ip_addresses** - The list of private IP address assigned to the load balancer in **frontend_ip_configuration** blocks, if any.

» Import

Load Balancers can be imported using the **resource id**, e.g.

```
terraform import azurerm_lb.test /subscriptions/00000000-0000-0000-0000-000000000000/resources/00000000-0000-0000-0000-000000000000
```

» azurerm_lb_backend_address_pool

Manage a Load Balancer Backend Address Pool.

NOTE: When using this resource, the Load Balancer needs to have a FrontEnd IP Configuration Attached

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name = "LoadBalancerRG"
```

```

    location = "West US"
}

resource "azurerm_public_ip" "test" {
  name                = "PublicIPForLB"
  location            = "West US"
  resource_group_name = "${azurerm_resource_group.test.name}"
  allocation_method = "Static"
}

resource "azurerm_lb" "test" {
  name                = "TestLoadBalancer"
  location            = "West US"
  resource_group_name = "${azurerm_resource_group.test.name}"

  frontend_ip_configuration {
    name                = "PublicIPAddress"
    public_ip_address_id = "${azurerm_public_ip.test.id}"
  }
}

resource "azurerm_lb_backend_address_pool" "test" {
  resource_group_name = "${azurerm_resource_group.test.name}"
  loadbalancer_id     = "${azurerm_lb.test.id}"
  name                = "BackEndAddressPool"
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Backend Address Pool.
- **resource_group_name** - (Required) The name of the resource group in which to create the resource.
- **loadbalancer_id** - (Required) The ID of the Load Balancer in which to create the Backend Address Pool.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Backend Address Pool.
- **backend_ip_configurations** - The Backend IP Configurations associated with this Backend Address Pool.

- `load_balancing_rules` - The Load Balancing Rules associated with this Backend Address Pool.

» Import

Load Balancer Backend Address Pools can be imported using the `resource id`, e.g.

```
terraform import azurerm_lb_backend_address_pool.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_lb_rule`

Manages a Load Balancer Rule.

NOTE When using this resource, the Load Balancer needs to have a FrontEnd IP Configuration Attached

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "LoadBalancerRG"
  location = "West US"
}

resource "azurerm_public_ip" "test" {
  name                = "PublicIPForLB"
  location            = "West US"
  resource_group_name = "${azurerm_resource_group.test.name}"
  allocation_method = "Static"
}

resource "azurerm_lb" "test" {
  name                = "TestLoadBalancer"
  location            = "West US"
  resource_group_name = "${azurerm_resource_group.test.name}"

  frontend_ip_configuration {
    name                = "PublicIPAddress"
    public_ip_address_id = "${azurerm_public_ip.test.id}"
  }
}

resource "azurerm_lb_rule" "test" {
```

```

resource_group_name      = "${azurerm_resource_group.test.name}"
loadbalancer_id          = "${azurerm_lb.test.id}"
name                     = "LBRule"
protocol                 = "Tcp"
frontend_port            = 3389
backend_port             = 3389
frontend_ip_configuration_name = "PublicIPAddress"
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the LB Rule.
- **resource_group_name** - (Required) The name of the resource group in which to create the resource.
- **loadbalancer_id** - (Required) The ID of the Load Balancer in which to create the Rule.
- **frontend_ip_configuration_name** - (Required) The name of the frontend IP configuration to which the rule is associated.
- **protocol** - (Required) The transport protocol for the external endpoint. Possible values are `Tcp`, `Udp` or `All`.
- **frontend_port** - (Required) The port for the external endpoint. Port numbers for each Rule must be unique within the Load Balancer. Possible values range between 0 and 65534, inclusive.
- **backend_port** - (Required) The port used for internal connections on the endpoint. Possible values range between 0 and 65535, inclusive.
- **backend_address_pool_id** - (Optional) A reference to a Backend Address Pool over which this Load Balancing Rule operates.
- **probe_id** - (Optional) A reference to a Probe used by this Load Balancing Rule.
- **enable_floating_ip** - (Optional) Floating IP is pertinent to failover scenarios: a "floating" IP is reassigned to a secondary server in case the primary server fails. Floating IP is required for SQL AlwaysOn.
- **idle_timeout_in_minutes** - (Optional) Specifies the timeout for the Tcp idle connection. The value can be set between 4 and 30 minutes. The default value is 4 minutes. This element is only used when the protocol is set to `Tcp`.
- **load_distribution** - (Optional) Specifies the load balancing distribution type to be used by the Load Balancer. Possible values are: `Default` – The load balancer is configured to use a 5 tuple hash to map traffic to available servers. `SourceIP` – The load balancer is configured to use a 2 tuple hash to map traffic to available servers. `SourceIPProtocol` – The load balancer is configured to use a 3 tuple hash to map traffic to available servers. Also known as Session Persistence, where the options are called `None`, `Client`

IP and Client IP and Protocol respectively.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the Load Balancer to which the resource is attached.

» Import

Load Balancer Rules can be imported using the `resource id`, e.g.

```
terraform import azurerm_lb_rule.test /subscriptions/00000000-0000-0000-0000-000000000000/r
```

» `azurerm_lb_outbound_rule`

Manages a Load Balancer Outbound Rule.

NOTE When using this resource, the Load Balancer needs to have a FrontEnd IP Configuration and a Backend Address Pool Attached.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "LoadBalancerRG"
  location = "West US"
}

resource "azurerm_public_ip" "test" {
  name                = "PublicIPForLB"
  location            = "West US"
  resource_group_name = "${azurerm_resource_group.test.name}"
  allocation_method   = "Static"
}

resource "azurerm_lb" "test" {
  name                = "TestLoadBalancer"
  location            = "West US"
  resource_group_name = "${azurerm_resource_group.test.name}"

  frontend_ip_configuration {
    name                = "PublicIPAddress"
    public_ip_address_id = "${azurerm_public_ip.test.id}"
  }
}
```

```

    }
}

resource "azurerm_lb_backend_address_pool" "test" {
  resource_group_name = "${azurerm_resource_group.test.name}"
  loadbalancer_id      = "${azurerm_lb.test.id}"
  name                 = "be-%d"
}

resource "azurerm_lb_outbound_rule" "test" {
  resource_group_name      = "${azurerm_resource_group.test.name}"
  loadbalancer_id          = "${azurerm_lb.test.id}"
  name                     = "OutboundRule"
  protocol                 = "Tcp"
  backend_address_pool_id  = "${azurerm_lb_backend_address_pool.test.id}"

  frontend_ip_configuration {
    name = "PublicIPAddress"
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Outbound Rule. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the resource. Changing this forces a new resource to be created.
- **loadbalancer_id** - (Required) The ID of the Load Balancer in which to create the Outbound Rule. Changing this forces a new resource to be created.
- **frontend_ip_configuration** - (Required) One or more **frontend_ip_configuration** blocks as defined below.
- **backend_address_pool_id** - (Required) The ID of the Backend Address Pool. Outbound traffic is randomly load balanced across IPs in the backend IPs.
- **protocol** - (Required) The transport protocol for the external endpoint. Possible values are `Udp`, `Tcp` or `All`.
- **enable_tcp_reset** - (Optional) Receive bidirectional TCP Reset on TCP flow idle timeout or unexpected connection termination. This element is only used when the protocol is set to `TCP`.
- **allocated_outbound_ports** - (Optional) The number of outbound ports to be used for NAT.

- `idle_timeout_in_minutes` - (Optional) The timeout for the TCP idle connection

A `frontend_ip_configuration` block supports the following:

- `name` - (Required) The name of the Frontend IP Configuration.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the Load Balancer to which the resource is attached.

» Import

Load Balancer Outbound Rules can be imported using the `resource id`, e.g.

```
terraform import azurerm_lb_outbound_rule.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_lb_nat_rule`

Manages a Load Balancer NAT Rule.

NOTE When using this resource, the Load Balancer needs to have a FrontEnd IP Configuration Attached

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "LoadBalancerRG"
  location = "West US"
}

resource "azurerm_public_ip" "test" {
  name                 = "PublicIPForLB"
  location             = "West US"
  resource_group_name = "${azurerm_resource_group.test.name}"
  allocation_method   = "Static"
}

resource "azurerm_lb" "test" {
  name = "TestLoadBalancer"
```

```

location          = "West US"
resource_group_name = "${azurerm_resource_group.test.name}"

frontend_ip_configuration {
  name          = "PublicIPAddress"
  public_ip_address_id = "${azurerm_public_ip.test.id}"
}
}

resource "azurerm_lb_nat_rule" "test" {
  resource_group_name      = "${azurerm_resource_group.test.name}"
  loadbalancer_id          = "${azurerm_lb.test.id}"
  name                     = "RDPAccess"
  protocol                 = "Tcp"
  frontend_port            = 3389
  backend_port             = 3389
  frontend_ip_configuration_name = "PublicIPAddress"
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the NAT Rule.
- **resource_group_name** - (Required) The name of the resource group in which to create the resource.
- **loadbalancer_id** - (Required) The ID of the Load Balancer in which to create the NAT Rule.
- **frontend_ip_configuration_name** - (Required) The name of the frontend IP configuration exposing this rule.
- **protocol** - (Required) The transport protocol for the external endpoint. Possible values are `Udp`, `Tcp` or `All`.
- **frontend_port** - (Required) The port for the external endpoint. Port numbers for each Rule must be unique within the Load Balancer. Possible values range between 1 and 65534, inclusive.
- **backend_port** - (Required) The port used for internal connections on the endpoint. Possible values range between 1 and 65535, inclusive.
- **enable_floating_ip** - (Optional) Enables the Floating IP Capacity, required to configure a SQL AlwaysOn Availability Group.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Load Balancer to which the resource is attached.

» Import

Load Balancer NAT Rules can be imported using the `resource id`, e.g.

```
terraform import azurerm_lb_nat_rule.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_lb_nat_pool

Manages a Load Balancer NAT pool.

NOTE When using this resource, the Load Balancer needs to have a FrontEnd IP Configuration Attached

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "LoadBalancerRG"
  location = "West US"
}

resource "azurerm_public_ip" "test" {
  name                = "PublicIPForLB"
  location             = "West US"
  resource_group_name = "${azurerm_resource_group.test.name}"
  allocation_method   = "Static"
}

resource "azurerm_lb" "test" {
  name                = "TestLoadBalancer"
  location             = "West US"
  resource_group_name = "${azurerm_resource_group.test.name}"

  frontend_ip_configuration {
    name                = "PublicIPAddress"
    public_ip_address_id = "${azurerm_public_ip.test.id}"
  }
}

resource "azurerm_lb_nat_pool" "test" {
  resource_group_name = "${azurerm_resource_group.test.name}"
  loadbalancer_id     = "${azurerm_lb.test.id}"
  name                = "SampleApplicationPool"
  protocol             = "Tcp"
  frontend_port_start = 80
}
```

```

    frontend_port_end           = 81
    backend_port                = 8080
    frontend_ip_configuration_name = "PublicIPAddress"
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the NAT pool.
- **resource_group_name** - (Required) The name of the resource group in which to create the resource.
- **loadbalancer_id** - (Required) The ID of the Load Balancer in which to create the NAT pool.
- **frontend_ip_configuration_name** - (Required) The name of the frontend IP configuration exposing this rule.
- **protocol** - (Required) The transport protocol for the external endpoint. Possible values are `Udp` or `Tcp`.
- **frontend_port_start** - (Required) The first port number in the range of external ports that will be used to provide Inbound Nat to NICs associated with this Load Balancer. Possible values range between 1 and 65534, inclusive.
- **frontend_port_end** - (Required) The last port number in the range of external ports that will be used to provide Inbound Nat to NICs associated with this Load Balancer. Possible values range between 1 and 65534, inclusive.
- **backend_port** - (Required) The port used for the internal endpoint. Possible values range between 1 and 65535, inclusive.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Load Balancer to which the resource is attached.

» Import

Load Balancer NAT Pools can be imported using the **resource id**, e.g.

```
terraform import azurerm_lb_nat_pool.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_lb_probe

Manages a LoadBalancer Probe Resource.

NOTE When using this resource, the Load Balancer needs to have a FrontEnd IP Configuration Attached

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "LoadBalancerRG"
  location  = "West US"
}

resource "azurerm_public_ip" "test" {
  name                  = "PublicIPForLB"
  location              = "West US"
  resource_group_name   = "${azurerm_resource_group.test.name}"
  allocation_method     = "Static"
}

resource "azurerm_lb" "test" {
  name                = "TestLoadBalancer"
  location            = "West US"
  resource_group_name = "${azurerm_resource_group.test.name}"

  frontend_ip_configuration {
    name                = "PublicIPAddress"
    public_ip_address_id = "${azurerm_public_ip.test.id}"
  }
}

resource "azurerm_lb_probe" "test" {
  resource_group_name = "${azurerm_resource_group.test.name}"
  loadbalancer_id     = "${azurerm_lb.test.id}"
  name                = "ssh-running-probe"
  port                = 22
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Probe.

- **resource_group_name** - (Required) The name of the resource group in which to create the resource.
- **loadbalancer_id** - (Required) The ID of the LoadBalancer in which to create the NAT Rule.
- **protocol** - (Optional) Specifies the protocol of the end point. Possible values are `Http`, `Https` or `Tcp`. If `Tcp` is specified, a received ACK is required for the probe to be successful. If `Http` is specified, a 200 OK response from the specified URI is required for the probe to be successful.
- **port** - (Required) Port on which the Probe queries the backend endpoint. Possible values range from 1 to 65535, inclusive.
- **request_path** - (Optional) The URI used for requesting health status from the backend endpoint. Required if protocol is set to `Http`. Otherwise, it is not allowed.
- **interval_in_seconds** - (Optional) The interval, in seconds between probes to the backend endpoint for health status. The default value is 15, the minimum value is 5.
- **number_of_probes** - (Optional) The number of failed probe attempts after which the backend endpoint is removed from rotation. The default value is 2. NumberOfProbes multiplied by intervalInSeconds value must be greater or equal to 10. Endpoints are returned to rotation when at least one probe is successful.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Load Balancer to which the resource is attached.

» Import

Load Balancer Probes can be imported using the **resource id**, e.g.

```
terraform import azurerm_lb_probe.test /subscriptions/00000000-0000-0000-0000-000000000000/1
```

» azurerm_logic_app_action_custom

Manages a Custom Action within a Logic App Workflow

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name = "workflow-resources"
```



```

    location = "East US"
  }

  resource "azurerm_logic_app_workflow" "test" {
    name           = "workflow1"
    location       = "${azurerm_resource_group.test.location}"
    resource_group_name = "${azurerm_resource_group.test.name}"
  }

  resource "azurerm_logic_app_action_custom" "test" {
    name           = "example-action"
    logic_app_id   = "${azurerm_logic_app_workflow.test.id}"

    body = <<BODY
  {
    "description": "A variable to configure the auto expiration age in days. Configured in r
    "inputs": {
      "variables": [
        {
          "name": "ExpirationAgeInDays",
          "type": "Integer",
          "value": -30
        }
      ]
    },
    "runAfter": {},
    "type": "InitializeVariable"
  }
  BODY
  }

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the HTTP Action to be created within the Logic App Workflow. Changing this forces a new resource to be created.

NOTE: This name must be unique across all Actions within the Logic App Workflow.

- **logic_app_id** - (Required) Specifies the ID of the Logic App Workflow. Changing this forces a new resource to be created.
- **body** - (Required) Specifies the JSON Blob defining the Body of this Custom Action.

NOTE: To make the Action more readable, you may wish to consider using HEREDOC syntax (as shown above) or the `local_file` resource to load the schema from a file on disk.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the Action within the Logic App Workflow.

» Import

Logic App Custom Actions can be imported using the `resource id`, e.g.

```
terraform import azurerm_logic_app_action_custom.custom1 /subscriptions/00000000-0000-0000-0000-000000000000/resourceGroups/rg-test/providers/Microsoft.Logic/workflows/workflow-test/actions/action-test
```

NOTE: This ID is unique to Terraform and doesn't directly match to any other resource. To compose this ID, you can take the ID Logic App Workflow and append `/actions/{name of the action}`.

» `azurerm_logic_app_action_http`

Manages an HTTP Action within a Logic App Workflow

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "workflow-resources"
  location = "East US"
}

resource "azurerm_logic_app_workflow" "test" {
  name          = "workflow1"
  location      = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_logic_app_action_http" "test" {
  name          = "webhook"
  logic_app_id = "${azurerm_logic_app_workflow.test.id}"
  method       = "GET"
  uri          = "http://example.com/some-webhook"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the HTTP Action to be created within the Logic App Workflow. Changing this forces a new resource to be created.

NOTE: This name must be unique across all Actions within the Logic App Workflow.

- **logic_app_id** - (Required) Specifies the ID of the Logic App Workflow. Changing this forces a new resource to be created.
- **method** - (Required) Specifies the HTTP Method which should be used for this HTTP Action. Possible values include **DELETE**, **GET**, **PATCH**, **POST** and **PUT**.
- **uri** - (Required) Specifies the URI which will be called when this HTTP Action is triggered.
- **body** - (Optional) Specifies the HTTP Body that should be sent to the **uri** when this HTTP Action is triggered.
- **headers** - (Optional) Specifies a Map of Key-Value Pairs that should be sent to the **uri** when this HTTP Action is triggered.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the HTTP Action within the Logic App Workflow.

» Import

Logic App HTTP Actions can be imported using the **resource id**, e.g.

```
terraform import azurerm_logic_app_action_http.webhook1 /subscriptions/00000000-0000-0000-0000-000000000000
```

NOTE: This ID is unique to Terraform and doesn't directly match to any other resource. To compose this ID, you can take the ID Logic App Workflow and append `/actions/{name of the action}`.

» `azurerm_logic_app_trigger_custom`

Manages a Custom Trigger within a Logic App Workflow

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "workflow-resources"
  location  = "East US"
}

resource "azurerm_logic_app_workflow" "test" {
  name                = "workflow1"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_logic_app_trigger_custom" "test" {
  name      = "example-trigger"
  logic_app_id = "${azurerm_logic_app_workflow.test.id}"

  body = <<BODY
{
  "recurrence": {
    "frequency": "Day",
    "interval": 1
  },
  "type": "Recurrence"
}
BODY
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the HTTP Trigger to be created within the Logic App Workflow. Changing this forces a new resource to be created.

NOTE: This name must be unique across all Triggers within the Logic App Workflow.

- **logic_app_id** - (Required) Specifies the ID of the Logic App Workflow. Changing this forces a new resource to be created.
- **body** - (Required) Specifies the JSON Blob defining the Body of this Custom Trigger.

NOTE: To make the Trigger more readable, you may wish to consider using HEREDOC syntax (as shown above) or the `local_file` resource to load the

schema from a file on disk.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the Trigger within the Logic App Workflow.

» Import

Logic App Custom Triggers can be imported using the `resource id`, e.g.

```
terraform import azurerm_logic_app_trigger_custom.custom1 /subscriptions/00000000-0000-0000-
```

NOTE: This ID is unique to Terraform and doesn't directly match to any other resource. To compose this ID, you can take the ID Logic App Workflow and append `/triggers/{name of the trigger}`.

» `azurerm_logic_app_trigger_http_request`

Manages a HTTP Request Trigger within a Logic App Workflow

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "workflow-resources"
  location = "East US"
}

resource "azurerm_logic_app_workflow" "test" {
  name          = "workflow1"
  location      = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_logic_app_trigger_http_request" "test" {
  name          = "some-http-trigger"
  logic_app_id = "${azurerm_logic_app_workflow.test.id}"

  schema = <<SCHEMA
{
  "type": "object",
  "properties": {
```

```

        "hello": {
            "type": "string"
        }
    }
}
SCHEMA
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the HTTP Request Trigger to be created within the Logic App Workflow. Changing this forces a new resource to be created.

NOTE: This name must be unique across all Triggers within the Logic App Workflow.

- **logic_app_id** - (Required) Specifies the ID of the Logic App Workflow. Changing this forces a new resource to be created.
- **schema** - (Required) A JSON Blob defining the Schema of the incoming request. This needs to be valid JSON.

NOTE: To make the Trigger more readable, you may wish to consider using HEREDOC syntax (as shown above) or the `local_file` resource to load the schema from a file on disk.

- **method** - (Optional) Specifies the HTTP Method which the request be using. Possible values include DELETE, GET, PATCH, POST or PUT.
- **relative_path** - (Optional) Specifies the Relative Path used for this Request.

NOTE: When `relative_path` is set a `method` must also be set.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the HTTP Request Trigger within the Logic App Workflow.

» Import

Logic App HTTP Request Triggers can be imported using the `resource id`, e.g.

```
terraform import azurerm_logic_app_trigger_http_request.request1 /subscriptions/00000000-0000-0000-0000-000000000000
```

NOTE: This ID is unique to Terraform and doesn't directly match to any other resource. To compose this ID, you can take the ID Logic App Workflow and append `/triggers/{name of the trigger}`.

» azurerm_logic_app_trigger_recurrence

Manages a Recurrence Trigger within a Logic App Workflow

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "workflow-resources"
  location  = "East US"
}

resource "azurerm_logic_app_workflow" "test" {
  name      = "workflow1"
  location  = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_logic_app_trigger_recurrence" "test" {
  name          = "run-every-day"
  logic_app_id = "${azurerm_logic_app_workflow.test.id}"
  frequency     = "Day"
  interval      = 1
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Recurrence Triggers to be created within the Logic App Workflow. Changing this forces a new resource to be created.

NOTE: This name must be unique across all Triggers within the Logic App Workflow.

- **logic_app_id** - (Required) Specifies the ID of the Logic App Workflow. Changing this forces a new resource to be created.

- **frequency** - (Required) Specifies the Frequency at which this Trigger should be run. Possible values include **Month**, **Week**, **Day**, **Hour**, **Minute** and **Second**.
- **interval** - (Required) Specifies interval used for the Frequency, for example a value of 4 for **interval** and **hour** for **frequency** would run the Trigger every 4 hours.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Recurrence Trigger within the Logic App Workflow.

» Import

Logic App Recurrence Triggers can be imported using the **resource id**, e.g.

```
terraform import azurerm_logic_app_trigger_recurrence.daily /subscriptions/00000000-0000-0000-0000-000000000000
```

NOTE: This ID is unique to Terraform and doesn't directly match to any other resource. To compose this ID, you can take the ID Logic App Workflow and append `/triggers/{name of the trigger}`.

» azurerm_logic_app_workflow

Manages a Logic App Workflow.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "workflow-resources"
  location = "East US"
}

resource "azurerm_logic_app_workflow" "test" {
  name          = "workflow1"
  location      = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}
```


» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Logic App Workflow. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the Resource Group in which the Logic App Workflow should be created. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the Logic App Workflow exists. Changing this forces a new resource to be created.
- **workflow_schema** - (Optional) Specifies the Schema to use for this Logic App Workflow. Defaults to `https://schema.management.azure.com/providers/Microsoft.Logic/schemas/2016-06-01/workflowdefinition.json`. Changing this forces a new resource to be created.
- **workflow_version** - (Optional) Specifies the version of the Schema used for this Logic App Workflow. Defaults to `1.0.0.0`. Changing this forces a new resource to be created.
- **parameters** - (Optional) A map of Key-Value pairs.

NOTE: Any parameters specified must exist in the Schema defined in `workflow_schema`.

- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- **id** - The Logic App Workflow ID.
- **access_endpoint** - The Access Endpoint for the Logic App Workflow

» Import

Logic App Workflows can be imported using the `resource id`, e.g.

```
terraform import azurerm_logic_app_workflow.workflow1 /subscriptions/00000000-0000-0000-0000-000000000000/resourceGroups/rg-name/providers/Microsoft.Logic/workflows/workflow-name
```

» `azurerm_management_lock`

Manages a Management Lock which is scoped to a Subscription, Resource Group or Resource.

» Example Usage (Subscription Level Lock)

```
data "azurerm_subscription" "current" {}

resource "azurerm_management_lock" "subscription-level" {
  name      = "subscription-level"
  scope     = "${data.azurerm_subscription.current.id}"
  lock_level = "CanNotDelete"
  notes     = "Items can't be deleted in this subscription!"
}
```

» Example Usage (Resource Group Level Lock)

```
resource "azurerm_resource_group" "test" {
  name      = "locked-resource-group"
  location = "West Europe"
}

resource "azurerm_management_lock" "resource-group-level" {
  name      = "resource-group-level"
  scope     = "${azurerm_resource_group.test.id}"
  lock_level = "ReadOnly"
  notes     = "This Resource Group is Read-Only"
}
```

» Example Usage (Resource Level Lock)

```
resource "azurerm_resource_group" "test" {
  name      = "locked-resource-group"
  location = "West Europe"
}

resource "azurerm_public_ip" "test" {
  name                        = "locked-publicip"
  location                   = "${azurerm_resource_group.test.location}"
  resource_group_name       = "${azurerm_resource_group.test.name}"
  allocation_method         = "Static"
  idle_timeout_in_minutes   = 30
}

resource "azurerm_management_lock" "public-ip" {
  name      = "resource-ip"
  scope     = "${azurerm_public_ip.test.id}"
  lock_level = "CanNotDelete"
}
```

```

    notes      = "Locked because it's needed by a third-party"
  }

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Management Lock. Changing this forces a new resource to be created.
- **scope** - (Required) Specifies the scope at which the Management Lock should be created. Changing this forces a new resource to be created.
- **lock_level** - (Required) Specifies the Level to be used for this Lock. Possible values are **CanNotDelete** and **ReadOnly**. Changing this forces a new resource to be created.

Note: **CanNotDelete** means authorized users are able to read and modify the resources, but not delete. **ReadOnly** means authorized users can only read from a resource, but they can't modify or delete it.

- **notes** - (Optional) Specifies some notes about the lock. Maximum of 512 characters. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Management Lock

» Import

Management Locks can be imported using the **resource id**, e.g.

```
terraform import azurerm_management_lock.lock1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_eventgrid_domain

Manages an EventGrid Domain

» Example Usage

```

resource "azurerm_resource_group" "test" {
  name     = "resourceGroup1"

```

```

    location = "West US 2"
  }

  resource "azurerm_eventgrid_domain" "test" {
    name                = "my-eventgrid-domain"
    location            = "${azurerm_resource_group.test.location}"
    resource_group_name = "${azurerm_resource_group.test.name}"

    tags = {
      environment = "Production"
    }
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the EventGrid Domain resource. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the EventGrid Domain exists. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **input_schema** - (Optional) Specifies the schema in which incoming events will be published to this domain. Allowed values are `cloudeventv01schema`, `customeventschema`, or `eventgridschema`. Defaults to `eventgridschema`. Changing this forces a new resource to be created.
- **input_mapping_fields** - (Optional) A `input_mapping_fields` block as defined below.
- **input_mapping_default_values** - (Optional) A `input_mapping_default_values` block as defined below.
- **tags** - (Optional) A mapping of tags to assign to the resource.

A `input_mapping_fields` supports the following:

- **id** - (Optional) Specifies the id of the EventGrid Event to associate with the domain. Changing this forces a new resource to be created.
- **topic** - (Optional) Specifies the topic of the EventGrid Event to associate with the domain. Changing this forces a new resource to be created.

- **event_type** - (Optional) Specifies the event type of the EventGrid Event to associate with the domain. Changing this forces a new resource to be created.
- **event_time** - (Optional) Specifies the event time of the EventGrid Event to associate with the domain. Changing this forces a new resource to be created.
- **data_version** - (Optional) Specifies the data version of the EventGrid Event to associate with the domain. Changing this forces a new resource to be created.
- **subject** - (Optional) Specifies the subject of the EventGrid Event to associate with the domain. Changing this forces a new resource to be created.

A **input_mapping_default_values** supports the following:

- **event_type** - (Optional) Specifies the default event type of the EventGrid Event to associate with the domain. Changing this forces a new resource to be created.
- **data_version** - (Optional) Specifies the default data version of the EventGrid Event to associate with the domain. Changing this forces a new resource to be created.
- **subject** - (Optional) Specifies the default subject of the EventGrid Event to associate with the domain. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the EventGrid Domain.
- **endpoint** - The Endpoint associated with the EventGrid Domain.

» Import

EventGrid Domain's can be imported using the **resource id**, e.g.

```
terraform import azurerm_eventgrid_domain.domain1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_eventgrid_event_subscription

Manages an EventGrid Event Subscription

» Example Usage

```
resource "azurerm_resource_group" "default" {
  name      = "defaultResourceGroup"
  location  = "West US 2"
}

resource "azurerm_storage_account" "default" {
  name                        = "defaultStorageAccount"
  resource_group_name        = "${azurerm_resource_group.default.name}"
  location                   = "${azurerm_resource_group.default.location}"
  account_tier                = "Standard"
  account_replication_type    = "LRS"

  tags = {
    environment = "staging"
  }
}

resource "azurerm_storage_queue" "default" {
  name                = "defaultStorageQueue"
  resource_group_name = "${azurerm_resource_group.default.name}"
  storage_account_name = "${azurerm_storage_account.default.name}"
}

resource "azurerm_eventgrid_event_subscription" "default" {
  name = "defaultEventSubscription"
  scope = "${azurerm_resource_group.default.id}"
  storage_queue_endpoint {
    storage_account_id = "${azurerm_storage_account.default.id}"
    queue_name          = "${azurerm_storage_queue.default.name}"
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the EventGrid Event Subscription resource. Changing this forces a new resource to be created.

- **scope** - (Required) Specifies the scope at which the EventGrid Event Subscription should be created. Changing this forces a new resource to be created.
- **event_delivery_schema** - (Optional) Specifies the event delivery schema for the event subscription. Possible values include: **EventGridSchema**, **CloudEventV01Schema**, **CustomInputSchema**.
- **topic_name** - (Optional) Specifies the name of the topic to associate with the event subscription.
- **storage_queue_endpoint** - (Optional) A **storage_queue_endpoint** block as defined below.
- **eventhub_endpoint** - (Optional) A **eventhub_endpoint** block as defined below.
- **hybrid_connection_endpoint** - (Optional) A **hybrid_connection_endpoint** block as defined below.
- **webhook_endpoint** - (Optional) A **webhook_endpoint** block as defined below.

NOTE: One of **storage_queue_endpoint**, **eventhub_endpoint**, **hybrid_connection_endpoint** or **webhook_endpoint** must be specified.

- **included_event_types** - (Optional) A list of applicable event types that need to be part of the event subscription.
- **subject_filter** - (Optional) A **subject_filter** block as defined below.
- **storage_blob_dead_letter_destination** - (Optional) A **storage_blob_dead_letter_destination** block as defined below.
- **retry_policy** - (Optional) A **retry_policy** block as defined below.
- **labels** - (Optional) A list of labels to assign to the event subscription.

A **storage_queue_endpoint** supports the following:

- **storage_account_id** - (Required) Specifies the id of the storage account id where the storage queue is located.
- **queue_name** - (Required) Specifies the name of the storage queue where the Event Subscription will receive events.

A **eventhub_endpoint** supports the following:

- **eventhub_id** - (Required) Specifies the id of the eventhub where the Event Subscription will receive events.
-

A `hybrid_connection_endpoint` supports the following:

- `hybrid_connection_id` - (Required) Specifies the id of the hybrid connection where the Event Subscription will receive events.

A `webhook_endpoint` supports the following:

- `url` - (Required) Specifies the url of the webhook where the Event Subscription will receive events.

A `subject_filter` supports the following:

- `subject_begins_with` - (Optional) A string to filter events for an event subscription based on a resource path prefix.
- `subject_ends_with` - (Optional) A string to filter events for an event subscription based on a resource path suffix.
- `case_sensitive` - (Optional) Specifies if `subject_begins_with` and `subject_ends_with` case sensitive. This value defaults to `false`.

A `storage_blob_dead_letter_destination` supports the following:

- `storage_account_id` - (Required) Specifies the id of the storage account id where the storage blob is located.
- `storage_blob_container_name` - (Required) Specifies the name of the Storage blob container that is the destination of the deadletter events

A `retry_policy` supports the following:

- `max_delivery_attempts` - (Required) Specifies the maximum number of delivery retry attempts for events.
- `event_time_to_live` - (Required) Specifies the time to live (in minutes) for events.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the EventGrid Event Subscription.

» Import

EventGrid Domain's can be imported using the `resource id`, e.g.


```
terraform import azurerm_eventgrid_event_subscription.eventSubscription1
/subscriptions/00000000-0000-0000-0000-000000000000/resourceGroups/group1/providers/Microsoft
```

» azurerm__eventgrid__topic

Manages an EventGrid Topic

Note: at this time EventGrid Topic's are only available in a limited number of regions.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "resourceGroup1"
  location  = "West US 2"
}

resource "azurerm_eventgrid_topic" "test" {
  name                = "my-eventgrid-topic"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  tags = {
    environment = "Production"
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the EventGrid Topic resource. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the EventGrid Topic exists. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- `id` - The EventGrid Topic ID.
- `endpoint` - The Endpoint associated with the EventGrid Topic.
- `primary_access_key` - The Primary Shared Access Key associated with the EventGrid Topic.
- `secondary_access_key` - The Secondary Shared Access Key associated with the EventGrid Topic.

» Import

EventGrid Topic's can be imported using the `resource id`, e.g.

```
terraform import azurerm_eventgrid_topic.topic1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_eventhub`

Manages a Event Hubs as a nested resource within a Event Hubs namespace.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "resourceGroup1"
  location = "West US"
}

resource "azurerm_eventhub_namespace" "test" {
  name                = "acceptanceTestEventHubNamespace"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  sku                 = "Standard"
  capacity            = 1
  kafka_enabled       = false

  tags = {
    environment = "Production"
  }
}

resource "azurerm_eventhub" "test" {
```

```

name           = "acceptanceTestEventHub"
namespace_name = "${azurerm_eventhub_namespace.test.name}"
resource_group_name = "${azurerm_resource_group.test.name}"
partition_count = 2
message_retention = 1
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the EventHub Namespace resource. Changing this forces a new resource to be created.
- **namespace_name** - (Required) Specifies the name of the EventHub Namespace. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the EventHub's parent Namespace exists. Changing this forces a new resource to be created.
- **partition_count** - (Required) Specifies the current number of shards on the Event Hub. Changing this forces a new resource to be created.
- **message_retention** - (Required) Specifies the number of days to retain the events for this Event Hub. Needs to be between 1 and 7 days; or 1 day when using a Basic SKU for the parent EventHub Namespace.
- **capture_description** - (Optional) A `capture_description` block as defined below.

A `capture_description` block supports the following:

- **enabled** - (Required) Specifies if the Capture Description is Enabled.
- **encoding** - (Required) Specifies the Encoding used for the Capture Description. Possible values are `Avro` and `AvroDeflate`.
- **interval_in_seconds** - (Optional) Specifies the time interval in seconds at which the capture will happen. Values can be between 60 and 900 seconds. Defaults to 300 seconds.
- **size_limit_in_bytes** - (Optional) Specifies the amount of data built up in your EventHub before a Capture Operation occurs. Value should be between 10485760 and 524288000 bytes. Defaults to 314572800 bytes.
- **skip_empty_archives** - (Optional) Specifies if empty files should not be emitted if no events occur during the Capture time window. Defaults to `false`.

- **destination** - (Required) A **destination** block as defined below.

A **destination** block supports the following:

- **name** - (Required) The Name of the Destination where the capture should take place. At this time the only supported value is `EventHubArchive.AzureBlockBlob`.

At this time it's only possible to Capture EventHub messages to Blob Storage. There's a Feature Request for the Azure SDK to add support for Capturing messages to Azure Data Lake here.

- **archive_name_format** - The Blob naming convention for archiving. e.g. `{Namespace}/{EventHub}/{PartitionId}/{Year}/{Month}/{Day}/{Hour}/{Minute}/{Second}`. Here all the parameters (Namespace,EventHub .. etc) are mandatory irrespective of order
- **blob_container_name** - (Required) The name of the Container within the Blob Storage Account where messages should be archived.
- **storage_account_id** - (Required) The ID of the Blob Storage Account where messages should be archived.

» Attributes Reference

The following attributes are exported:

- **id** - The EventHub ID.
- **partition_ids** - The identifiers for partitions created for Event Hubs.

» Import

EventHubs can be imported using the **resource id**, e.g.

```
terraform import azurerm_eventhub.eventhub1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» **azurerm_eventhub_authorization_rule**

Manages a Event Hubs authorization Rule within an Event Hub.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "resourceGroup1"
  location = "West US"
}
```

```

}

resource "azurerm_eventhub_namespace" "test" {
  name           = "acceptanceTestEventHubNamespace"
  location       = "West US"
  resource_group_name = "${azurerm_resource_group.test.name}"
  sku            = "Basic"
  capacity       = 2

  tags = {
    environment = "Production"
  }
}

resource "azurerm_eventhub" "test" {
  name           = "acceptanceTestEventHub"
  namespace_name = "${azurerm_eventhub_namespace.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  partition_count = 2
  message_retention = 2
}

resource "azurerm_eventhub_authorization_rule" "test" {
  name           = "navi"
  namespace_name = "${azurerm_eventhub_namespace.test.name}"
  eventhub_name  = "${azurerm_eventhub.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  listen         = true
  send           = false
  manage         = false
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the EventHub Authorization Rule resource. Changing this forces a new resource to be created.
- **namespace_name** - (Required) Specifies the name of the grandparent EventHub Namespace. Changing this forces a new resource to be created.
- **eventhub_name** - (Required) Specifies the name of the EventHub. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in

which the EventHub Namespace exists. Changing this forces a new resource to be created.

NOTE At least one of the 3 permissions below needs to be set.

- **listen** - (Optional) Does this Authorization Rule have permissions to Listen to the Event Hub? Defaults to **false**.
- **send** - (Optional) Does this Authorization Rule have permissions to Send to the Event Hub? Defaults to **false**.
- **manage** - (Optional) Does this Authorization Rule have permissions to Manage to the Event Hub? When this property is **true** - both **listen** and **send** must be too. Defaults to **false**.

» Attributes Reference

The following attributes are exported:

- **id** - The EventHub ID.
- **primary_key** - The Primary Key for the Event Hubs authorization Rule.
- **primary_connection_string** - The Primary Connection String for the Event Hubs authorization Rule.
- **secondary_key** - The Secondary Key for the Event Hubs authorization Rule.
- **secondary_connection_string** - The Secondary Connection String for the Event Hubs authorization Rule.

» Import

EventHubs can be imported using the **resource id**, e.g.

```
terraform import azurerm_eventhub_authorization_rule.rule1 /subscriptions/00000000-0000-0000
```

» azurerm_eventhub_consumer_group

Manages a Event Hubs Consumer Group as a nested resource within an Event Hub.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "resourceGroup1"
  location  = "West US"
}

resource "azurerm_eventhub_namespace" "test" {
  name                = "acceptanceTestEventHubNamespace"
  location            = "West US"
  resource_group_name = "${azurerm_resource_group.test.name}"
  sku                = "Basic"
  capacity            = 2

  tags = {
    environment = "Production"
  }
}

resource "azurerm_eventhub" "test" {
  name                = "acceptanceTestEventHub"
  namespace_name      = "${azurerm_eventhub_namespace.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  partition_count     = 2
  message_retention   = 2
}

resource "azurerm_eventhub_consumer_group" "test" {
  name                = "acceptanceTestEventHubConsumerGroup"
  namespace_name      = "${azurerm_eventhub_namespace.test.name}"
  eventhub_name       = "${azurerm_eventhub.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  user_metadata       = "some-meta-data"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the EventHub Consumer Group resource. Changing this forces a new resource to be created.
- **namespace_name** - (Required) Specifies the name of the grandparent EventHub Namespace. Changing this forces a new resource to be created.

- `eventhub_name` - (Required) Specifies the name of the EventHub. Changing this forces a new resource to be created.
- `resource_group_name` - (Required) The name of the resource group in which the EventHub Consumer Group's grandparent Namespace exists. Changing this forces a new resource to be created.
- `user_metadata` - (Optional) Specifies the user metadata.

» Attributes Reference

The following attributes are exported:

- `id` - The EventHub Consumer Group ID.

» Import

EventHub Consumer Groups can be imported using the `resource id`, e.g.

```
terraform import azurerm_eventhub_consumer_group.consumerGroup1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_eventhub_namespace`

Manage an EventHub Namespace.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "resourceGroup1"
  location = "West US"
}

resource "azurerm_eventhub_namespace" "test" {
  name                = "acceptanceTestEventHubNamespace"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  sku                 = "Standard"
  capacity             = 2

  tags = {
    environment = "Production"
  }
}
```


» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the EventHub Namespace resource. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the namespace. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **sku** - (Required) Defines which tier to use. Valid options are **Basic** and **Standard**.
- **capacity** - (Optional) Specifies the Capacity / Throughput Units for a **Standard** SKU namespace. Valid values range from 1 - 20.
- **auto_inflate_enabled** - (Optional) Is Auto Inflate enabled for the EventHub Namespace?
- **maximum_throughput_units** - (Optional) Specifies the maximum number of throughput units when Auto Inflate is Enabled. Valid values range from 1 - 20.
- **kafka_enabled** - (Optional) Is Kafka enabled for the EventHub Namespace? Defaults to **false**.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- **id** - The EventHub Namespace ID.

The following attributes are exported only if there is an authorization rule named **RootManageSharedAccessKey** which is created automatically by Azure.

- **default_primary_connection_string** - The primary connection string for the authorization rule **RootManageSharedAccessKey**.
- **default_secondary_connection_string** - The secondary connection string for the authorization rule **RootManageSharedAccessKey**.
- **default_primary_key** - The primary access key for the authorization rule **RootManageSharedAccessKey**.
- **default_secondary_key** - The secondary access key for the authorization rule **RootManageSharedAccessKey**.

» Import

EventHub Namespaces can be imported using the `resource id`, e.g.

```
terraform import azurerm_eventhub_namespace.namespace1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_eventhub_namespace_authorization_rule`

Manages an Authorization Rule for an Event Hub Namespace.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "resourceGroup1"
  location = "West US"
}

resource "azurerm_eventhub_namespace" "test" {
  name                = "acceptanceTestEventHubNamespace"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  sku                 = "Basic"
  capacity             = 2

  tags = {
    environment = "Production"
  }
}

resource "azurerm_eventhub_namespace_authorization_rule" "test" {
  name                = "navi"
  namespace_name      = "${azurerm_eventhub_namespace.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  listen = true
  send   = false
  manage = false
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Authorization Rule. Changing this forces a new resource to be created.
- **namespace_name** - (Required) Specifies the name of the EventHub Namespace. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the EventHub Namespace exists. Changing this forces a new resource to be created.

NOTE At least one of the 3 permissions below needs to be set.

- **listen** - (Optional) Grants listen access to this this Authorization Rule. Defaults to **false**.
- **send** - (Optional) Grants send access to this this Authorization Rule. Defaults to **false**.
- **manage** - (Optional) Grants manage access to this this Authorization Rule. When this property is **true** - both **listen** and **send** must be too. Defaults to **false**.

» Attributes Reference

The following attributes are exported:

- **id** - The EventHub ID.
- **primary_key** - The Primary Key for the Authorization Rule.
- **primary_connection_string** - The Primary Connection String for the Authorization Rule.
- **secondary_key** - The Secondary Key for the Authorization Rule.
- **secondary_connection_string** - The Secondary Connection String for the Authorization Rule.

» Import

EventHubs can be imported using the **resource id**, e.g.

```
terraform import azurerm_eventhub_namespace_authorization_rule.rule1 /subscriptions/00000000
```

» azurerm_iotHub

Manages an IoT Hub

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "resourceGroup1"
  location  = "West US"
}

resource "azurerm_storage_account" "test" {
  name                        = "teststa"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  location                   = "${azurerm_resource_group.test.location}"
  account_tier               = "Standard"
  account_replication_type   = "LRS"
}

resource "azurerm_storage_container" "test" {
  name                      = "test"
  resource_group_name       = "${azurerm_resource_group.test.name}"
  storage_account_name      = "${azurerm_storage_account.test.name}"
  container_access_type     = "private"
}

resource "azurerm_iothub" "test" {
  name                      = "test"
  resource_group_name       = "${azurerm_resource_group.test.name}"
  location                  = "${azurerm_resource_group.test.location}"

  sku {
    name      = "S1"
    tier       = "Standard"
    capacity  = "1"
  }

  endpoint {
    type                        = "AzureIoTHub.StorageContainer"
    connection_string          = "${azurerm_storage_account.test.primary_blob_connection_string}"
    name                       = "export"
    batch_frequency_in_seconds = 60
    max_chunk_size_in_bytes    = 10485760
    container_name              = "test"
    encoding                    = "Avro"
    file_name_format            = "${iothub}/{partition}_{YYYY}_{MM}_{DD}_{HH}_{mm}"
  }

  route {
```

```

    name          = "export"
    source         = "DeviceMessages"
    condition      = "true"
    endpoint_names = ["export"]
    enabled        = true
  }

  fallback_route {
    enabled = true
  }

  tags = {
    "purpose" = "testing"
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the IoT Hub resource. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group under which the IoT Hub resource has to be created. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource has to be created. Changing this forces a new resource to be created.
- **sku** - (Required) A **sku** block as defined below.
- **endpoint** - (Optional) An **endpoint** block as defined below.
- **ip_filter_rule** - (Optional) One or more **ip_filter_rule** blocks as defined below.
- **route** - (Optional) A **route** block as defined below.
- **fallback_route** - (Optional) A **fallback_route** block as defined below. If the fallback route is enabled, messages that don't match any of the supplied routes are automatically sent to this route. Defaults to `messages/events`.
- **tags** - (Optional) A mapping of tags to assign to the resource.

A **sku** block supports the following:

- **name** - (Required) The name of the sku. Possible values are B1, B2, B3, F1, S1, S2, and S3.
- **tier** - (Required) The billing tier for the IoT Hub. Possible values are **Basic**, **Free** or **Standard**.

NOTE: Only one IoT Hub can be on the **Free** tier per subscription.

- **capacity** - (Required) The number of provisioned IoT Hub units.

An **endpoint** block supports the following:

- **type** - (Required) The type of the endpoint. Possible values are `AzureIoTHub.StorageContainer`, `AzureIoTHub.ServiceBusQueue`, `AzureIoTHub.ServiceBusTopic` or `AzureIoTHub.EventHub`.
- **connection_string** - (Required) The connection string for the endpoint.
- **name** - (Required) The name of the endpoint. The name must be unique across endpoint types. The following names are reserved: `events`, `operationsMonitoringEvents`, `fileNotifications` and `$default`.
- **batch_frequency_in_seconds** - (Optional) Time interval at which blobs are written to storage. Value should be between 60 and 720 seconds. Default value is 300 seconds. This attribute is mandatory for endpoint type `AzureIoTHub.StorageContainer`.
- **max_chunk_size_in_bytes** - (Optional) Maximum number of bytes for each blob written to storage. Value should be between 10485760(10MB) and 524288000(500MB). Default value is 314572800(300MB). This attribute is mandatory for endpoint type `AzureIoTHub.StorageContainer`.
- **container_name** - (Optional) The name of storage container in the storage account. This attribute is mandatory for endpoint type `AzureIoTHub.StorageContainer`.
- **encoding** - (Optional) Encoding that is used to serialize messages to blobs. Supported values are 'avro' and 'avrodeflate'. Default value is 'avro'. This attribute is mandatory for endpoint type `AzureIoTHub.StorageContainer`.
- **file_name_format** - (Optional) File name format for the blob. Default format is `{iothub}/{partition}/{YYYY}/{MM}/{DD}/{HH}/{mm}`. All parameters are mandatory but can be reordered. This attribute is mandatory for endpoint type `AzureIoTHub.StorageContainer`.

An **ip_filter_rule** block supports the following:

- **name** - (Required) The name of the filter.

- **ip_mask** - (Required) The IP address range in CIDR notation for the rule.
 - **action** - (Required) The desired action for requests captured by this rule. Possible values are **Accept**, **Reject**
-

A **route** block supports the following:

- **name** - (Required) The name of the route.
 - **source** - (Required) The source that the routing rule is to be applied to, such as **DeviceMessages**. Possible values include: **RoutingSourceInvalid**, **RoutingSourceDeviceMessages**, **RoutingSourceTwinChangeEvents**, **RoutingSourceDeviceLifecycleEvents**, **RoutingSourceDeviceJobLifecycleEvents**.
 - **condition** - (Optional) The condition that is evaluated to apply the routing rule. If no condition is provided, it evaluates to true by default. For grammar, see: <https://docs.microsoft.com/azure/iot-hub/iot-hub-devguide-query-language>.
 - **endpoint_names** - (Required) The list of endpoints to which messages that satisfy the condition are routed.
 - **enabled** - (Required) Used to specify whether a route is enabled.
-

A **fallback_route** block supports the following:

- **source** - (Optional) The source that the routing rule is to be applied to, such as **DeviceMessages**. Possible values include: **RoutingSourceInvalid**, **RoutingSourceDeviceMessages**, **RoutingSourceTwinChangeEvents**, **RoutingSourceDeviceLifecycleEvents**, **RoutingSourceDeviceJobLifecycleEvents**.
- **condition** - (Optional) The condition that is evaluated to apply the routing rule. If no condition is provided, it evaluates to true by default. For grammar, see: <https://docs.microsoft.com/azure/iot-hub/iot-hub-devguide-query-language>.
- **endpoint_names** - (Optional) The endpoints to which messages that satisfy the condition are routed. Currently only 1 endpoint is allowed.
- **enabled** - (Optional) Used to specify whether the fallback route is enabled.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the IoT Hub.
- **event_hub_events_endpoint** - The EventHub compatible endpoint for events data

- `event_hub_events_path` - The EventHub compatible path for events data
- `event_hub_operations_endpoint` - The EventHub compatible endpoint for operational data
- `event_hub_operations_path` - The EventHub compatible path for operational data

NOTE: These fields can be used in conjunction with the `shared_access_policy` block to build a connection string

- `hostname` - The hostname of the IoT Hub Resource.
- `shared_access_policy` - One or more `shared_access_policy` blocks as defined below.

A `shared access policy` block contains the following:

- `key_name` - The name of the shared access policy.
- `primary_key` - The primary key.
- `secondary_key` - The secondary key.
- `permissions` - The permissions assigned to the shared access policy.

» Import

IoT Hubs can be imported using the `resource id`, e.g.

```
terraform import azurerm_iothub.hub1 /subscriptions/00000000-0000-0000-0000-000000000000/res
```

» `azurerm_iothub_consumer_group`

Manages a Consumer Group within an IoT Hub

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "resourceGroup1"
  location = "West US"
}

resource "azurerm_iothub" "test" {
  name = "test"
  resource_group_name = "${azurerm_resource_group.test.name}"
}
```



```

location          = "${azurerm_resource_group.test.location}"

sku {
  name      = "S1"
  tier       = "Standard"
  capacity  = "1"
}

tags = {
  "purpose" = "testing"
}
}

resource "azurerm_iothub_consumer_group" "test" {
  name                = "terraform"
  iothub_name         = "${azurerm_iothub.test.name}"
  eventhub_endpoint_name = "events"
  resource_group_name = "${azurerm_resource_group.foo.name}"
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of this Consumer Group. Changing this forces a new resource to be created.
- **iothub_name** - (Required) The name of the IoT Hub. Changing this forces a new resource to be created.
- **eventhub_endpoint_name** - (Required) The name of the Event Hub-compatible endpoint in the IoT hub. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group that contains the IoT hub. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the IoTHub Consumer Group.

» Import

IoTHub Consumer Groups can be imported using the **resource id**, e.g.

```
terraform import azurerm_iothub_consumer_group.group1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» **azurerm_notification_hub**

Manages a Notification Hub within a Notification Hub Namespace.

» **Example Usage**

```
resource "azurerm_resource_group" "test" {
  name      = "notificationhub-resources"
  location  = "Australia East"
}

resource "azurerm_notification_hub_namespace" "test" {
  name                = "myappnamespace"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location            = "${azurerm_resource_group.test.location}"
  namespace_type      = "NotificationHub"

  sku {
    name = "Free"
  }
}

resource "azurerm_notification_hub" "test" {
  name                = "mynotificationhub"
  namespace_name      = "${azurerm_notification_hub_namespace.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location            = "${azurerm_resource_group.test.location}"
}
```

» **Argument Reference**

The following arguments are supported:

- **name** - (Required) The name to use for this Notification Hub. Changing this forces a new resource to be created.
- **namespace_name** - (Required) The name of the Notification Hub Namespace in which to create this Notification Hub. Changing this forces a new resource to be created.

- **resource_group_name** - (Required) The name of the Resource Group in which the Notification Hub Namespace exists. Changing this forces a new resource to be created.
- **location** - (Required) The Azure Region in which this Notification Hub Namespace exists. Changing this forces a new resource to be created.
- **apns_credential** - (Optional) A **apns_credential** block as defined below.

NOTE: Removing the **apns_credential** block will currently force a recreation of this resource due to this bug in the Azure SDK for Go - we'll remove this limitation when the SDK bug is fixed.

- **gcm_credential** - (Optional) A **gcm_credential** block as defined below.

NOTE: Removing the **gcm_credential** block will currently force a recreation of this resource due to this bug in the Azure SDK for Go - we'll remove this limitation when the SDK bug is fixed.

A **apns_credential** block contains:

- **application_mode** - (Required) The Application Mode which defines which server the APNS Messages should be sent to. Possible values are **Production** and **Sandbox**.
- **bundle_id** - (Required) The Bundle ID of the iOS/macOS application to send push notifications for, such as **com.hashicorp.example**.
- **key_id** - (Required) The Apple Push Notifications Service (APNS) Key.
- **team_id** - (Required) The ID of the team the Token.
- **token** - (Required) The Push Token associated with the Apple Developer Account. This is the contents of the **key** downloaded from the Apple Developer Portal between the -----BEGIN PRIVATE KEY----- and -----END PRIVATE KEY----- blocks.

A **gcm_credential** block contains:

- **api_key** - (Required) The API Key associated with the Google Cloud Messaging service.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Notification Hub.

» Import

Notification Hubs can be imported using the `resource id`, e.g.

```
terraform import azurerm_notification_hub.hub1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_notification_hub_authorization_rule

Manages an Authorization Rule associated with a Notification Hub within a Notification Hub Namespace.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "notificationhub-resources"
  location  = "Australia East"
}

resource "azurerm_notification_hub_namespace" "test" {
  name                = "myappnamespace"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location            = "${azurerm_resource_group.test.location}"
  namespace_type      = "NotificationHub"

  sku {
    name = "Free"
  }
}

resource "azurerm_notification_hub" "test" {
  name                = "mynotificationhub"
  namespace_name      = "${azurerm_notification_hub_namespace.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location            = "${azurerm_resource_group.test.location}"
}

resource "azurerm_notification_hub_authorization_rule" "test" {
  name                = "management-auth-rule"
  notification_hub_name = "${azurerm_notification_hub.test.name}"
  namespace_name      = "${azurerm_notification_hub_namespace.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  manage              = true
  send                = true
  listen              = true
}
```

}

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name to use for this Authorization Rule. Changing this forces a new resource to be created.
- **notification_hub_name** - (Required) The name of the Notification Hub for which the Authorization Rule should be created. Changing this forces a new resource to be created.
- **namespace_name** - (Required) The name of the Notification Hub Namespace in which the Notification Hub exists. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the Resource Group in which the Notification Hub Namespace exists. Changing this forces a new resource to be created.
- **manage** - (Optional) Does this Authorization Rule have Manage access to the Notification Hub? Defaults to **false**.

NOTE: If **manage** is set to **true** then both **send** and **listen** must also be set to **true**.

- **send** - (Optional) Does this Authorization Rule have Send access to the Notification Hub? Defaults to **false**.
- **listen** - (Optional) Does this Authorization Rule have Listen access to the Notification Hub? Defaults to **false**.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Authorization Rule.
- **primary_access_key** - The Primary Access Key associated with this Authorization Rule.
- **secondary_access_key** - The Secondary Access Key associated with this Authorization Rule.

» Import

Notification Hub Authorization Rule can be imported using the **resource id**, e.g.

```
terraform import azurerm_notification_hub_authorization_rule.rule1 /subscriptions/00000000-
```

» **azurerm_notification_hub_namespace**

Manages a Notification Hub Namespace.

» **Example Usage**

```
resource "azurerm_resource_group" "test" {
  name      = "notificationhub-resources"
  location  = "Australia East"
}

resource "azurerm_notification_hub_namespace" "test" {
  name                        = "myappnamespace"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  location                   = "${azurerm_resource_group.test.location}"
  namespace_type             = "NotificationHub"

  sku {
    name = "Free"
  }
}
```

» **Argument Reference**

The following arguments are supported:

- **name** - (Required) The name to use for this Notification Hub Namespace. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the Resource Group in which the Notification Hub Namespace should exist. Changing this forces a new resource to be created.
- **location** - (Required) The Azure Region in which this Notification Hub Namespace should be created.
- **namespace_type** - (Required) The Type of Namespace - possible values are **Messaging** or **NotificationHub**. Changing this forces a new resource to be created.
- **sku** - (Required) A **sku** block as defined below.

- **enabled** - (Optional) Is this Notification Hub Namespace enabled? Defaults to **true**.

A **sku** block contains:

- **name** - (Required) The name of the SKU to use for this Notification Hub Namespace. Possible values are **Free**, **Basic** or **Standard**. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Notification Hub Namespace.
- **servicebus_endpoint** - The ServiceBus Endpoint for this Notification Hub Namespace.

» Import

Notification Hub Namespaces can be imported using the **resource id**, e.g.

```
terraform import azurerm_notification_hub_namespace.namespace1 /subscriptions/00000000-0000-
```

» azurerm__relay__namespace

Manages an Azure Relay Namespace.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "example-resources"
  location = "West Europe"
}

resource "azurerm_relay_namespace" "test" {
  name                = "example-relay"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  sku {
    name = "Standard"
  }
}
```

```

    }

    tags = {
      source = "terraform"
    }
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Azure Relay Namespace. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the Azure Relay Namespace.
- **location** - (Required) Specifies the supported Azure location where the Azure Relay Namespace exists. Changing this forces a new resource to be created.
- **sku** - (Required) A **sku** block as defined below.
- **tags** - (Optional) A mapping of tags to assign to the resource.

A **sku** block contains:

- **name** - (Required) The name of the SKU to use. At this time the only supported value is **Standard**.

» Attributes Reference

The following attributes are exported:

- **id** - The Azure Relay Namespace ID.

The following attributes are exported only if there is an authorization rule named **RootManageSharedAccessKey** which is created automatically by Azure.

- **primary_connection_string** - The primary connection string for the authorization rule **RootManageSharedAccessKey**.
- **secondary_connection_string** - The secondary connection string for the authorization rule **RootManageSharedAccessKey**.
- **primary_key** - The primary access key for the authorization rule **RootManageSharedAccessKey**.

- `secondary_key` - The secondary access key for the authorization rule `RootManageSharedAccessKey`.
- `metric_id` - The Identifier for Azure Insights metrics.

» Import

Azure Relay Namespace's can be imported using the `resource id`, e.g.

```
terraform import azurerm_relay_namespace.relay1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_servicebus_namespace`

Manage a ServiceBus Namespace.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "terraform-servicebus"
  location = "West Europe"
}

resource "azurerm_servicebus_namespace" "example" {
  name                = "tfex_sevicebus_namespace"
  location             = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
  sku                 = "Standard"

  tags = {
    source = "terraform"
  }
}
```

» Argument Reference

The following arguments are supported:

- `name` - (Required) Specifies the name of the ServiceBus Namespace resource . Changing this forces a new resource to be created.
- `resource_group_name` - (Required) The name of the resource group in which to create the namespace.

- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **sku** - (Required) Defines which tier to use. Options are basic, standard or premium.
- **capacity** - (Optional) Specifies the capacity. When **sku** is **Premium** can be 1, 2 or 4. When **sku** is **Basic** or **Standard** can be 0 only.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- **id** - The ServiceBus Namespace ID.

The following attributes are exported only if there is an authorization rule named **RootManageSharedAccessKey** which is created automatically by Azure.

- **default_primary_connection_string** - The primary connection string for the authorization rule **RootManageSharedAccessKey**.
- **default_secondary_connection_string** - The secondary connection string for the authorization rule **RootManageSharedAccessKey**.
- **default_primary_key** - The primary access key for the authorization rule **RootManageSharedAccessKey**.
- **default_secondary_key** - The secondary access key for the authorization rule **RootManageSharedAccessKey**.

» Import

Service Bus Namespace can be imported using the **resource id**, e.g.

```
terraform import azurerm_servicebus_namespace.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_servicebus_namespace_authorization_rule

Manages a ServiceBus Namespace authorization Rule within a ServiceBus.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name = "terraform-servicebus"
```

```

    location = "West US"
  }

  resource "azurerm_servicebus_namespace" "example" {
    name           = "tfex_sevicebus_namespace"
    location       = "${azurerm_resource_group.example.location}"
    resource_group_name = "${azurerm_resource_group.example.name}"
    sku            = "Standard"

    tags = {
      source = "terraform"
    }
  }

  resource "azurerm_servicebus_namespace_authorization_rule" "example" {
    name           = "examplerule"
    namespace_name = "${azurerm_servicebus_namespace.example.name}"
    resource_group_name = "${azurerm_resource_group.example.name}"

    listen = true
    send   = true
    manage = false
  }

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the ServiceBus Namespace Authorization Rule resource. Changing this forces a new resource to be created.
- **namespace_name** - (Required) Specifies the name of the ServiceBus Namespace. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the ServiceBus Namespace exists. Changing this forces a new resource to be created.

NOTE At least one of the 3 permissions below needs to be set.

- **listen** - (Optional) Grants listen access to this this Authorization Rule. Defaults to **false**.
- **send** - (Optional) Grants send access to this this Authorization Rule. Defaults to **false**.
- **manage** - (Optional) Grants manage access to this this Authorization Rule. When this property is **true** - both **listen** and **send** must be too. Defaults

to false.

» Attributes Reference

The following attributes are exported:

- `id` - The ServiceBus Topic ID.
- `primary_key` - The Primary Key for the ServiceBus Namespace authorization Rule.
- `primary_connection_string` - The Primary Connection String for the ServiceBus Namespace authorization Rule.
- `secondary_key` - The Secondary Key for the ServiceBus Namespace authorization Rule.
- `secondary_connection_string` - The Secondary Connection String for the ServiceBus Namespace authorization Rule.

» Import

ServiceBus Namespace authorization rules can be imported using the `resource` id, e.g.

```
terraform import azurerm_servicebus_namespace_authorization_rule.rule1 /subscriptions/000000
```

» `azurerm_servicebus_queue`

Manage and manage a ServiceBus Queue.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "terraform-servicebus"
  location = "West Europe"
}

resource "azurerm_servicebus_namespace" "example" {
  name                = "tfex_sevicebus_namespace"
  location            = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
  sku                 = "Standard"
}
```

```

tags = {
  source = "terraform"
}

resource "azurerm_servicebus_queue" "example" {
  name                        = "tfex_servicebus_queue"
  resource_group_name        = "${azurerm_resource_group.example.name}"
  namespace_name             = "${azurerm_servicebus_namespace.example.name}"

  enable_partitioning = true
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the ServiceBus Queue resource. Changing this forces a new resource to be created.
- **namespace_name** - (Required) The name of the ServiceBus Namespace to create this queue in. Changing this forces a new resource to be created.
- **location** - (Optional / **Deprecated**) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the namespace. Changing this forces a new resource to be created.
- **auto_delete_on_idle** - (Optional) The ISO 8601 timespan duration of the idle interval after which the Queue is automatically deleted, minimum of 5 minutes.
- **default_message_ttl** - (Optional) The ISO 8601 timespan duration of the TTL of messages sent to this queue. This is the default value used when TTL is not set on message itself.
- **duplicate_detection_history_time_window** - (Optional) The ISO 8601 timespan duration during which duplicates can be detected. Default value is 10 minutes. (PT10M)
- **enable_express** - (Optional) Boolean flag which controls whether Express Entities are enabled. An express queue holds a message in memory temporarily before writing it to persistent storage. Defaults to **false** for Basic and Standard. For Premium, it **MUST** be set to **false**.

NOTE: Service Bus Premium namespaces do not support Express Entities, so `enable_express` MUST be set to `false`.

- `enable_partitioning` - (Optional) Boolean flag which controls whether to enable the queue to be partitioned across multiple message brokers. Changing this forces a new resource to be created. Defaults to `false` for Basic and Standard. For Premium, it MUST be set to `true`.

NOTE: Partitioning is available at entity creation for all queues and topics in Basic or Standard SKUs. It is not available for the Premium messaging SKU, but any previously existing partitioned entities in Premium namespaces continue to work as expected. Please see the documentation for more information.

- `lock_duration` - (Optional) The ISO 8601 timespan duration of a peek-lock; that is, the amount of time that the message is locked for other receivers. Maximum value is 5 minutes. Defaults to 1 minute. (PT1M)
- `max_size_in_megabytes` - (Optional) Integer value which controls the size of memory allocated for the queue. For supported values see the "Queue/topic size" section of this document.
- `requires_duplicate_detection` - (Optional) Boolean flag which controls whether the Queue requires duplicate detection. Changing this forces a new resource to be created. Defaults to `false`.
- `requires_session` - (Optional) Boolean flag which controls whether the Queue requires sessions. This will allow ordered handling of unbounded sequences of related messages. With sessions enabled a queue can guarantee first-in-first-out delivery of messages. Changing this forces a new resource to be created. Defaults to `false`.
- `dead_lettering_on_message_expiration` - (Optional) Boolean flag which controls whether the Queue has dead letter support when a message expires. Defaults to `false`.
- `max_delivery_count` - (Optional) Integer value which controls when a message is automatically deadlettered. Defaults to 10.

» Attributes Reference

The following attributes are exported:

- `id` - The ServiceBus Queue ID.

» Import

Service Bus Queue can be imported using the `resource id`, e.g.

```
terraform import azurerm_servicebus_queue.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_servicebus_queue_authorization_rule`

Manages an Authorization Rule for a ServiceBus Queue.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "terraform-servicebus"
  location = "West US"
}

resource "azurerm_servicebus_namespace" "example" {
  name                = "tfex_servicebus_namespace"
  location             = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
  sku                 = "Standard"

  tags = {
    source = "terraform"
  }
}

resource "azurerm_servicebus_queue" "example" {
  name                = "tfex_servicebus_queue"
  resource_group_name = "${azurerm_resource_group.example.name}"
  namespace_name      = "${azurerm_servicebus_namespace.example.name}"

  enable_partitioning = true
}

resource "azurerm_servicebus_queue_authorization_rule" "example" {
  name                = "examplerule"
  namespace_name      = "${azurerm_servicebus_namespace.example.name}"
  queue_name          = "${azurerm_servicebus_queue.example.name}"
  resource_group_name = "${azurerm_resource_group.example.name}"

  listen = true
  send   = true
  manage = false
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Authorization Rule. Changing this forces a new resource to be created.
- **namespace_name** - (Required) Specifies the name of the ServiceBus Namespace in which the Queue exists. Changing this forces a new resource to be created.
- **queue_name** - (Required) Specifies the name of the ServiceBus Queue. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the Resource Group in which the ServiceBus Namespace exists. Changing this forces a new resource to be created.

NOTE At least one of the 3 permissions below needs to be set.

- **listen** - (Optional) Does this Authorization Rule have Listen permissions to the ServiceBus Queue? Defaults to **false**.
- **send** - (Optional) Does this Authorization Rule have Send permissions to the ServiceBus Queue? Defaults to **false**.
- **manage** - (Optional) Does this Authorization Rule have Manage permissions to the ServiceBus Queue? When this property is **true** - both **listen** and **send** must be too. Defaults to **false**.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Authorization Rule.
- **primary_key** - The Primary Key for the Authorization Rule.
- **primary_connection_string** - The Primary Connection String for the Authorization Rule.
- **secondary_key** - The Secondary Key for the Authorization Rule.
- **secondary_connection_string** - The Secondary Connection String for the Authorization Rule.

» Import

ServiceBus Queue Authorization Rules can be imported using the **resource id**, e.g.

```
terraform import azurerm_servicebus_queue_authorization_rule.rule1 /subscriptions/00000000-0000-0000-0000-000000000000
```


» azurerm_servicebus_subscription

Manage a ServiceBus Subscription.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name      = "tfex-servicebus-subscription"
  location  = "West Europe"
}

resource "azurerm_servicebus_namespace" "example" {
  name                = "tfex_sevicebus_namespace"
  location            = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
  sku                = "Standard"

  tags = {
    source = "terraform"
  }
}

resource "azurerm_servicebus_topic" "example" {
  name                = "tfex_sevicebus_topic"
  resource_group_name = "${azurerm_resource_group.example.name}"
  namespace_name      = "${azurerm_servicebus_namespace.example.name}"

  enable_partitioning = true
}

resource "azurerm_servicebus_subscription" "example" {
  name                = "tfex_sevicebus_subscription"
  resource_group_name = "${azurerm_resource_group.example.name}"
  namespace_name      = "${azurerm_servicebus_namespace.example.name}"
  topic_name          = "${azurerm_servicebus_topic.example.name}"
  max_delivery_count  = 1
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the ServiceBus Subscription resource. Changing this forces a new resource to be created.

- **namespace_name** - (Required) The name of the ServiceBus Namespace to create this Subscription in. Changing this forces a new resource to be created.
- **topic_name** - (Required) The name of the ServiceBus Topic to create this Subscription in. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the namespace. Changing this forces a new resource to be created.
- **max_delivery_count** - (Required) The maximum number of deliveries.
- **auto_delete_on_idle** - (Optional) The idle interval after which the Subscription is automatically deleted, minimum of 5 minutes. Provided in the TimeSpan format.
- **default_message_ttl** - (Optional) The TTL of messages sent to this Subscription if no TTL value is set on the message itself. Provided in the TimeSpan format.
- **lock_duration** - (Optional) The lock duration for the subscription, maximum supported value is 5 minutes. Defaults to 1 minute.
- **dead_lettering_on_message_expiration** - (Optional) Boolean flag which controls whether the Subscription has dead letter support when a message expires. Defaults to false.
- **enable_batched_operations** - (Optional) Boolean flag which controls whether the Subscription supports batched operations. Defaults to false.
- **requires_session** - (Optional) Boolean flag which controls whether this Subscription supports the concept of a session. Defaults to false. Changing this forces a new resource to be created.
- **forward_to** - (Optional) The name of a Queue or Topic to automatically forward messages to.

» TimeSpan Format

Some arguments for this resource are required in the TimeSpan format which is used to represent a length of time. The supported format is documented here

» Attributes Reference

The following attributes are exported:

- id - The ServiceBus Subscription ID.

» Import

Service Bus Subscriptions can be imported using the `resource id`, e.g.

```
terraform import azurerm_servicebus_subscription.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_servicebus_subscription_rule

Manage a ServiceBus Subscription Rule.

» Example Usage (SQL Filter)

```
resource "azurerm_resource_group" "example" {
  name     = "tfex-servicebus-subscription-rule-sql"
  location = "West Europe"
}

resource "azurerm_servicebus_namespace" "example" {
  name                = "tfex_sevicebus_namespace"
  location            = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
  sku                 = "Standard"

  tags = {
    source = "terraform"
  }
}

resource "azurerm_servicebus_topic" "example" {
  name                = "tfex_sevicebus_topic"
  resource_group_name = "${azurerm_resource_group.example.name}"
  namespace_name      = "${azurerm_servicebus_namespace.example.name}"

  enable_partitioning = true
}

resource "azurerm_servicebus_subscription" "example" {
  name                = "tfex_sevicebus_subscription"
  resource_group_name = "${azurerm_resource_group.example.name}"
  namespace_name      = "${azurerm_servicebus_namespace.example.name}"
  topic_name          = "${azurerm_servicebus_topic.example.name}"
}
```

```

    max_delivery_count = 1
  }

resource "azurerm_servicebus_subscription_rule" "example" {
  name = "tfex_sevicebus_rule"
  resource_group_name = "${azurerm_resource_group.example.name}"
  namespace_name = "${azurerm_servicebus_namespace.example.name}"
  topic_name = "${azurerm_servicebus_topic.example.name}"
  subscription_name = "${azurerm_servicebus_subscription.example.name}"
  filter_type = "SqlFilter"
  sql_filter = "color = 'red'"
}

```

» Example Usage (Correlation Filter)

```

resource "azurerm_resource_group" "example" {
  name     = "tfex-servicebus-subscription-rule-cor"
  location = "West Europe"
}

resource "azurerm_servicebus_namespace" "example" {
  name = "tfex_sevicebus_namespace"
  location = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
  sku = "Standard"

  tags = {
    source = "terraform"
  }
}

resource "azurerm_servicebus_topic" "example" {
  name = "tfex_sevicebus_topic"
  resource_group_name = "${azurerm_resource_group.example.name}"
  namespace_name = "${azurerm_servicebus_namespace.example.name}"

  enable_partitioning = true
}

resource "azurerm_servicebus_subscription" "example" {
  name = "tfex_sevicebus_subscription"
  resource_group_name = "${azurerm_resource_group.example.name}"
  namespace_name = "${azurerm_servicebus_namespace.example.name}"
  topic_name = "${azurerm_servicebus_topic.example.name}"
  max_delivery_count = 1
}

```

```

}

resource "azurerm_servicebus_subscription_rule" "example" {
  name = "tfex_sevicebus_rule"
  resource_group_name = "${azurerm_resource_group.example.name}"
  namespace_name = "${azurerm_servicebus_namespace.example.name}"
  topic_name = "${azurerm_servicebus_topic.example.name}"
  subscription_name = "${azurerm_servicebus_subscription.example.name}"
  filter_type = "CorrelationFilter"

  correlation_filter {
    correlation_id = "high"
    label = "red"
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the ServiceBus Subscription Rule. Changing this forces a new resource to be created.
- **namespace_name** - (Required) The name of the ServiceBus Namespace in which the ServiceBus Topic exists. Changing this forces a new resource to be created.
- **topic_name** - (Required) The name of the ServiceBus Topic in which the ServiceBus Subscription exists. Changing this forces a new resource to be created.
- **subscription_name** - (Required) The name of the ServiceBus Subscription in which this Rule should be created. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in the ServiceBus Namespace exists. Changing this forces a new resource to be created.
- **filter_type** - (Required) Type of filter to be applied to a BrokeredMessage. Possible values are `SqlFilter` and `CorrelationFilter`.
- **sql_filter** - (Optional) Represents a filter written in SQL language-based syntax that to be evaluated against a BrokeredMessage. Required when **filter_type** is set to `SqlFilter`.
- **correlation_filter** - (Optional) A `correlation_filter` block as documented below to be evaluated against a BrokeredMessage. Required when **filter_type** is set to `CorrelationFilter`.

- **action** - (Optional) Represents set of actions written in SQL language-based syntax that is performed against a BrokeredMessage.

correlation_filter supports the following:

- **content_type** - (Optional) Content type of the message.
- **correlation_id** - (Optional) Identifier of the correlation.
- **label** - (Optional) Application specific label.
- **message_id** - (Optional) Identifier of the message.
- **reply_to** - (Optional) Address of the queue to reply to.
- **reply_to_session_id** - (Optional) Session identifier to reply to.
- **session_id** - (Optional) Session identifier.
- **to** - (Optional) Address to send to.

NOTE: When creating a subscription rule of type **CorrelationFilter** at least one property must be set in the **correlation_filter** block.

» Attributes Reference

The following attributes are exported:

- **id** - The ServiceBus Subscription Rule ID.

» Import

Service Bus Subscription Rule can be imported using the **resource id**, e.g.

```
terraform import azurerm_servicebus_subscription.test /subscriptions/00000000-0000-0000-0000-
```

» azurerm__servicebus__topic

Manage a ServiceBus Topic.

Note Topics can only be created in Namespaces with an SKU of **standard** or higher.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "tfex-servicebus-topic"
  location = "West Europe"
```

```

}

resource "azurerm_servicebus_namespace" "example" {
  name                = "tfex_sevicebus_namespace"
  location            = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
  sku                 = "Standard"

  tags = {
    source = "terraform"
  }
}

resource "azurerm_servicebus_topic" "example" {
  name                = "tfex_sevicebus_topic"
  resource_group_name = "${azurerm_resource_group.example.name}"
  namespace_name      = "${azurerm_servicebus_namespace.example.name}"

  enable_partitioning = true
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the ServiceBus Topic resource. Changing this forces a new resource to be created.
- **namespace_name** - (Required) The name of the ServiceBus Namespace to create this topic in. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the namespace. Changing this forces a new resource to be created.
- **status** - (Optional) The Status of the Service Bus Topic. Acceptable values are `Active` or `Disabled`. Defaults to `Active`.
- **auto_delete_on_idle** - (Optional) The ISO 8601 timespan duration of the idle interval after which the Topic is automatically deleted, minimum of 5 minutes.
- **default_message_ttl** - (Optional) The ISO 8601 timespan duration of TTL of messages sent to this topic if no TTL value is set on the message itself.

- **duplicate_detection_history_time_window** - (Optional) The ISO 8601 timespan duration during which duplicates can be detected. Defaults to 10 minutes. (PT10M)
- **enable_batched_operations** - (Optional) Boolean flag which controls if server-side batched operations are enabled. Defaults to false.
- **enable_express** - (Optional) Boolean flag which controls whether Express Entities are enabled. An express topic holds a message in memory temporarily before writing it to persistent storage. Defaults to false.
- **enable_partitioning** - (Optional) Boolean flag which controls whether to enable the topic to be partitioned across multiple message brokers. Defaults to false. Changing this forces a new resource to be created.

NOTE: Partitioning is available at entity creation for all queues and topics in Basic or Standard SKUs. It is not available for the Premium messaging SKU, but any previously existing partitioned entities in Premium namespaces continue to work as expected. Please see the documentation for more information.

- **max_size_in_megabytes** - (Optional) Integer value which controls the size of memory allocated for the topic. For supported values see the "Queue/topic size" section of this document.
- **requires_duplicate_detection** - (Optional) Boolean flag which controls whether the Topic requires duplicate detection. Defaults to false. Changing this forces a new resource to be created.
- **support_ordering** - (Optional) Boolean flag which controls whether the Topic supports ordering. Defaults to false.

» Attributes Reference

The following attributes are exported:

- **id** - The ServiceBus Topic ID.

» Import

Service Bus Topics can be imported using the **resource id**, e.g.

```
terraform import azurerm_servicebus_topic.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_servicebus_topic_authorization_rule

Manages a ServiceBus Topic authorization Rule within a ServiceBus Topic.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name      = "tfex-servicebus"
  location  = "West US"
}

resource "azurerm_servicebus_namespace" "example" {
  name                = "tfex_servicebus_namespace"
  location            = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
  sku                = "Standard"

  tags = {
    source = "terraform"
  }
}

resource "azurerm_servicebus_topic" "example" {
  name                = "tfex_servicebus_topic"
  resource_group_name = "${azurerm_resource_group.example.name}"
  namespace_name      = "${azurerm_servicebus_namespace.example.name}"
}

resource "azurerm_servicebus_topic_authorization_rule" "example" {
  name                = "tfex_servicebus_topic_sasPolicy"
  namespace_name      = "${azurerm_servicebus_namespace.example.name}"
  topic_name          = "${azurerm_servicebus_topic.example.name}"
  resource_group_name = "${azurerm_resource_group.example.name}"
  listen              = true
  send                = false
  manage              = false
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the ServiceBus Topic Authorization Rule resource. Changing this forces a new resource to be created.
- **namespace_name** - (Required) Specifies the name of the ServiceBus Namespace. Changing this forces a new resource to be created.
- **topic_name** - (Required) Specifies the name of the ServiceBus Topic. Changing this forces a new resource to be created.

- **resource_group_name** - (Required) The name of the resource group in which the ServiceBus Namespace exists. Changing this forces a new resource to be created.

NOTE At least one of the 3 permissions below needs to be set.

- **listen** - (Optional) Grants listen access to this this Authorization Rule. Defaults to **false**.
- **send** - (Optional) Grants send access to this this Authorization Rule. Defaults to **false**.
- **manage** - (Optional) Grants manage access to this this Authorization Rule. When this property is **true** - both **listen** and **send** must be too. Defaults to **false**.

» Attributes Reference

The following attributes are exported:

- **id** - The ServiceBus Topic ID.
- **primary_key** - The Primary Key for the ServiceBus Topic authorization Rule.
- **primary_connection_string** - The Primary Connection String for the ServiceBus Topic authorization Rule.
- **secondary_key** - The Secondary Key for the ServiceBus Topic authorization Rule.
- **secondary_connection_string** - The Secondary Connection String for the ServiceBus Topic authorization Rule.

» Import

ServiceBus Topic authorization rules can be imported using the **resource id**, e.g.

```
terraform import azurerm_servicebus_topic_authorization_rule.rule1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm__signalr__service

Manages an Azure SignalR service.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name      = "terraform-signalr"
  location  = "West US"
}

resource "azurerm_signalr_service" "example" {
  name                = "tfex-signalr"
  location            = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"

  sku {
    name      = "Free_F1"
    capacity = 1
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the SignalR service. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the SignalR service. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the SignalR service exists. Changing this forces a new resource to be created.
- **sku** - A `sku` block as documented below.
- **tags** - (Optional) A mapping of tags to assign to the resource.

A `sku` block supports the following:

- **name** - (Required) Specifies which tier to use. Valid values are `Free_F1` and `Standard_S1`.
- **capacity** - (Required) Specifies the number of units associated with this SignalR service. Valid values are 1, 2, 5, 10, 20, 50 and 100.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the SignalR service.
- `hostname` - The FQDN of the SignalR service.
- `ip_address` - The publicly accessible IP of the SignalR service.
- `public_port` - The publicly accessible port of the SignalR service which is designed for browser/client use.
- `server_port` - The publicly accessible port of the SignalR service which is designed for customer server side use.
- `primary_access_key` - The primary access key for the SignalR service.
- `primary_connection_string` - The primary connection string for the SignalR service.
- `secondary_access_key` - The secondary access key for the SignalR service.
- `secondary_connection_string` - The secondary connection string for the SignalR service.

» Import

SignalR services can be imported using the `resource id`, e.g.

```
terraform import azurerm_signalr_service.example /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm__autoscale__setting`

Manages an AutoScale Setting which can be applied to Virtual Machine Scale Sets, App Services and other scalable resources.

NOTE: This resource has been deprecated in favour of the `azurerm_monitor_autoscale_setting` resource and will be removed in the next major version of the AzureRM Provider. The new resource shares the same fields as this one, and information on migrating across can be found in this guide.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "autoscalingTest"
  location = "West US"
```

```

}

resource "azurerm_virtual_machine_scale_set" "test" {
  # ...
}

resource "azurerm_autoscale_setting" "test" {
  name                = "myAutoscaleSetting"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location            = "${azurerm_resource_group.test.location}"
  target_resource_id  = "${azurerm_virtual_machine_scale_set.test.id}"

  profile {
    name = "defaultProfile"

    capacity {
      default = 1
      minimum = 1
      maximum = 10
    }

    rule {
      metric_trigger {
        metric_name      = "Percentage CPU"
        metric_resource_id = "${azurerm_virtual_machine_scale_set.test.id}"
        time_grain       = "PT1M"
        statistic        = "Average"
        time_window      = "PT5M"
        time_aggregation = "Average"
        operator         = "GreaterThan"
        threshold        = 75
      }

      scale_action {
        direction = "Increase"
        type      = "ChangeCount"
        value     = "1"
        cooldown  = "PT1M"
      }
    }
  }

  rule {
    metric_trigger {
      metric_name      = "Percentage CPU"
      metric_resource_id = "${azurerm_virtual_machine_scale_set.test.id}"
      time_grain       = "PT1M"
    }
  }
}

```

```

        statistic          = "Average"
        time_window        = "PT5M"
        time_aggregation   = "Average"
        operator           = "LessThan"
        threshold          = 25
    }

    scale_action {
        direction = "Decrease"
        type      = "ChangeCount"
        value     = "1"
        cooldown  = "PT1M"
    }
}

notification {
    email {
        send_to_subscription_administrator = true
        send_to_subscription_co_administrator = true
        custom_emails                      = ["admin@contoso.com"]
    }
}
}

```

» Example Usage (repeating on weekends)

```

resource "azurerm_resource_group" "test" {
    name     = "autoscalingTest"
    location = "West US"
}

resource "azurerm_virtual_machine_scale_set" "test" {
    # ...
}

resource "azurerm_autoscale_setting" "test" {
    name                = "myAutoscaleSetting"
    resource_group_name = "${azurerm_resource_group.test.name}"
    location            = "${azurerm_resource_group.test.location}"
    target_resource_id  = "${azurerm_virtual_machine_scale_set.test.id}"

    profile {
        name = "Weekends"
    }
}

```

```

capacity {
  default = 1
  minimum = 1
  maximum = 10
}

rule {
  metric_trigger {
    metric_name      = "Percentage CPU"
    metric_resource_id = "${azurerm_virtual_machine_scale_set.test.id}"
    time_grain       = "PT1M"
    statistic         = "Average"
    time_window       = "PT5M"
    time_aggregation  = "Average"
    operator          = "GreaterThan"
    threshold         = 90
  }

  scale_action {
    direction = "Increase"
    type      = "ChangeCount"
    value     = "2"
    cooldown  = "PT1M"
  }
}

rule {
  metric_trigger {
    metric_name      = "Percentage CPU"
    metric_resource_id = "${azurerm_virtual_machine_scale_set.test.id}"
    time_grain       = "PT1M"
    statistic         = "Average"
    time_window       = "PT5M"
    time_aggregation  = "Average"
    operator          = "LessThan"
    threshold         = 10
  }

  scale_action {
    direction = "Decrease"
    type      = "ChangeCount"
    value     = "2"
    cooldown  = "PT1M"
  }
}

```

```

    recurrence {
        frequency = "Week"
        timezone  = "Pacific Standard Time"
        days      = ["Saturday", "Sunday"]
        hours     = [12]
        minutes   = [0]
    }
}

notification {
    email {
        send_to_subscription_administrator = true
        send_to_subscription_co_administrator = true
        custom_emails                      = ["admin@contoso.com"]
    }
}
}

```

» Example Usage (for fixed dates)

```

resource "azurerm_resource_group" "test" {
    name      = "autoscalingTest"
    location = "West US"
}

resource "azurerm_virtual_machine_scale_set" "test" {
    # ...
}

resource "azurerm_autoscale_setting" "test" {
    name                = "myAutoscaleSetting"
    enabled              = true
    resource_group_name = "${azurerm_resource_group.test.name}"
    location             = "${azurerm_resource_group.test.location}"
    target_resource_id  = "${azurerm_virtual_machine_scale_set.test.id}"

    profile {
        name = "forJuly"

        capacity {
            default = 1
            minimum = 1
            maximum = 10
        }
    }
}

```



```

rule {
  metric_trigger {
    metric_name      = "Percentage CPU"
    metric_resource_id = "${azurerm_virtual_machine_scale_set.test.id}"
    time_grain       = "PT1M"
    statistic         = "Average"
    time_window       = "PT5M"
    time_aggregation  = "Average"
    operator          = "GreaterThan"
    threshold         = 90
  }

  scale_action {
    direction = "Increase"
    type      = "ChangeCount"
    value     = "2"
    cooldown  = "PT1M"
  }
}

rule {
  metric_trigger {
    metric_name      = "Percentage CPU"
    metric_resource_id = "${azurerm_virtual_machine_scale_set.test.id}"
    time_grain       = "PT1M"
    statistic         = "Average"
    time_window       = "PT5M"
    time_aggregation  = "Average"
    operator          = "LessThan"
    threshold         = 10
  }

  scale_action {
    direction = "Decrease"
    type      = "ChangeCount"
    value     = "2"
    cooldown  = "PT1M"
  }
}

fixed_date {
  timezone = "Pacific Standard Time"
  start    = "2020-07-01T00:00:00Z"
  end      = "2020-07-31T23:59:59Z"
}
}

```

```

notification {
  email {
    send_to_subscription_administrator = true
    send_to_subscription_co_administrator = true
    custom_emails = ["admin@contoso.com"]
  }
}
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the AutoScale Setting. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the Resource Group in the AutoScale Setting should be created. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the AutoScale Setting should exist. Changing this forces a new resource to be created.
- **profile** - (Required) Specifies one or more (up to 20) **profile** blocks as defined below.
- **target_resource_id** - (Required) Specifies the resource ID of the resource that the autoscale setting should be added to.
- **enabled** - (Optional) Specifies whether automatic scaling is enabled for the target resource. Defaults to **true**.
- **notification** - (Optional) Specifies a **notification** block as defined below.
- **tags** - (Optional) A mapping of tags to assign to the resource.

A **profile** block supports the following:

- **name** - (Required) Specifies the name of the profile.
- **capacity** - (Required) A **capacity** block as defined below.
- **rule** - (Required) One or more (up to 10) **rule** blocks as defined below.
- **fixed_date** - (Optional) A **fixed_date** block as defined below. This cannot be specified if a **recurrence** block is specified.

- **recurrence** - (Optional) A **recurrence** block as defined below. This cannot be specified if a **fixed_date** block is specified.

A **capacity** block supports the following:

- **default** - (Required) The number of instances that are available for scaling if metrics are not available for evaluation. The default is only used if the current instance count is lower than the default. Valid values are between 0 and 1000.
- **maximum** - (Required) The maximum number of instances for this resource. Valid values are between 0 and 1000.

NOTE: The maximum number of instances is also limited by the amount of Cores available in the subscription.

- **minimum** - (Required) The minimum number of instances for this resource. Valid values are between 0 and 1000.

A **rule** block supports the following:

- **metric_trigger** - (Required) A **metric_trigger** block as defined below.
- **scale_action** - (Required) A **scale_action** block as defined below.

A **metric_trigger** block supports the following:

- **metric_name** - (Required) The name of the metric that defines what the rule monitors, such as **Percentage CPU for Virtual Machine Scale Sets** and **CpuPercentage for App Service Plan**.

NOTE: The allowed value of **metric_name** highly depends on the targeting resource type, please visit Supported metrics with Azure Monitor for more details.

- **metric_resource_id** - (Required) The ID of the Resource which the Rule monitors.
- **operator** - (Required) Specifies the operator used to compare the metric data and threshold. Possible values are: **Equals**, **NotEquals**, **GreaterThan**, **GreaterThanOrEqual**, **LessThan**, **LessThanOrEqual**.
- **statistic** - (Required) Specifies how the metrics from multiple instances are combined. Possible values are **Average**, **Min** and **Max**.
- **time_aggregation** - (Required) Specifies how the data that's collected should be combined over time. Possible values include **Average**, **Count**, **Maximum**, **Minimum**, **Last** and **Total**. Defaults to **Average**.

- **time_grain** - (Required) Specifies the granularity of metrics that the rule monitors, which must be one of the pre-defined values returned from the metric definitions for the metric. This value must be between 1 minute and 12 hours and be formatted as an ISO 8601 string.
- **time_window** - (Required) Specifies the time range for which data is collected, which must be greater than the delay in metric collection (which varies from resource to resource). This value must be between 5 minutes and 12 hours and be formatted as an ISO 8601 string.
- **threshold** - (Required) Specifies the threshold of the metric that triggers the scale action.

A **scale_action** block supports the following:

- **cooldown** - (Required) The amount of time to wait since the last scaling action before this action occurs. Must be between 1 minute and 1 week and formatted as a ISO 8601 string.
- **direction** - (Required) The scale direction. Possible values are **Increase** and **Decrease**.
- **type** - (Required) The type of action that should occur. Possible values are **ChangeCount**, **ExactCount** and **PercentChangeCount**.
- **value** - (Required) The number of instances involved in the scaling action. Defaults to 1.

A **fixed_date** block supports the following:

- **end** - (Required) Specifies the end date for the profile, formatted as an RFC3339 date string.
- **start** - (Required) Specifies the start date for the profile, formatted as an RFC3339 date string.
- **timezone** (Optional) The Time Zone of the **start** and **end** times. A list of possible values can be found here. Defaults to **UTC**.

A **recurrence** block supports the following:

- **timezone** - (Required) The Time Zone used for the **hours** field. A list of possible values can be found here. Defaults to **UTC**.
- **days** - (Required) A list of days that this profile takes effect on. Possible values include **Monday**, **Tuesday**, **Wednesday**, **Thursday**, **Friday**, **Saturday** and **Sunday**.

- **hours** - (Required) A list containing a single item, which specifies the Hour interval at which this recurrence should be triggered (in 24-hour time). Possible values are from 0 to 23.
- **minutes** - (Required) A list containing a single item which specifies the Minute interval at which this recurrence should be triggered.

A **notification** block supports the following:

- **email** - (Required) A **email** block as defined below.
- **webhook** - (Optional) One or more **webhook** blocks as defined below.

A **email** block supports the following:

- **send_to_subscription_administrator** - (Optional) Should email notifications be sent to the subscription administrator? Defaults to **false**.
- **send_to_subscription_co_administrator** - (Optional) Should email notifications be sent to the subscription co-administrator? Defaults to **false**.
- **custom_emails** - (Optional) Specifies a list of custom email addresses to which the email notifications will be sent.

A **webhook** block supports the following:

- **service_uri** - (Required) The HTTPS URI which should receive scale notifications.
- **properties** - (Optional) A map of settings.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the AutoScale Setting.

» Import

AutoScale Setting can be imported using the **resource id**, e.g.

```
terraform import azurerm_autoscale_setting.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» **azurerm_monitor_action_group**

Manages an Action Group within Azure Monitor.

» **Example Usage**

```
resource "azurerm_resource_group" "test" {
  name      = "monitoring-resources"
  location = "West US"
}

resource "azurerm_monitor_action_group" "test" {
  name                        = "CriticalAlertsAction"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  short_name                  = "p0action"

  email_receiver {
    name      = "sendtoadmin"
    email_address = "admin@contoso.com"
  }

  email_receiver {
    name      = "sendtodevops"
    email_address = "devops@contoso.com"
  }

  sms_receiver {
    name      = "oncallmsg"
    country_code = "1"
    phone_number = "1231231234"
  }

  webhook_receiver {
    name      = "callmyapiaswell"
    service_uri = "http://example.com/alert"
  }
}
```

» **Argument Reference**

The following arguments are supported:

- **name** - (Required) The name of the Action Group. Changing this forces a new resource to be created.

- **resource_group_name** - (Required) The name of the resource group in which to create the Action Group instance.
- **short_name** - (Required) The short name of the action group. This will be used in SMS messages.
- **enabled** - (Optional) Whether this action group is enabled. If an action group is not enabled, then none of its receivers will receive communications. Defaults to **true**.
- **email_receiver** - (Optional) One or more **email_receiver** blocks as defined below.
- **sms_receiver** - (Optional) One or more **sms_receiver** blocks as defined below.
- **webhook_receiver** - (Optional) One or more **webhook_receiver** blocks as defined below.
- **tags** - (Optional) A mapping of tags to assign to the resource.

email_receiver supports the following:

- **name** - (Required) The name of the email receiver. Names must be unique (case-insensitive) across all receivers within an action group.
- **email_address** - (Required) The email address of this receiver.

sms_receiver supports the following:

- **name** - (Required) The name of the SMS receiver. Names must be unique (case-insensitive) across all receivers within an action group.
- **country_code** - (Required) The country code of the SMS receiver.
- **phone_number** - (Required) The phone number of the SMS receiver.

webhook_receiver supports the following:

- **name** - (Required) The name of the webhook receiver. Names must be unique (case-insensitive) across all receivers within an action group.
- **service_uri** - (Required) The URI where webhooks should be sent.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Action Group.

» Import

Action Groups can be imported using the **resource id**, e.g.

```
terraform import azurerm_monitor_action_group.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_monitor_activity_log_alert

Manages an Activity Log Alert within Azure Monitor.

» Example Usage

```
resource "azurerm_resource_group" "main" {
  name     = "example-resources"
  location = "West US"
}

resource "azurerm_monitor_action_group" "main" {
  name                = "example-actiongroup"
  resource_group_name = "${azurerm_resource_group.main.name}"
  short_name          = "p0action"

  webhook_receiver {
    name       = "callmyapi"
    service_uri = "http://example.com/alert"
  }
}

resource "azurerm_storage_account" "to_monitor" {
  name                = "examplesa"
  resource_group_name = "${azurerm_resource_group.main.name}"
  location            = "${azurerm_resource_group.main.location}"
  account_tier        = "Standard"
  account_replication_type = "GRS"
}

resource "azurerm_monitor_activity_log_alert" "main" {
  name                = "example-activitylogalert"
  resource_group_name = "${azurerm_resource_group.main.name}"
  scopes              = ["${azurerm_resource_group.main.id}"]
  description         = "This alert will monitor a specific storage account updates."

  criteria {
    resource_id      = "${azurerm_storage_account.to_monitor.id}"
    operation_name    = "Microsoft.Storage/storageAccounts/write"
    category          = "Recommendation"
  }
}
```



```

action {
  action_group_id = "${azure_rm_monitor_action_group.main.id}"

  webhook_properties = {
    from = "terraform"
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the activity log alert. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the activity log alert instance.
- **scopes** - (Required) The Scope at which the Activity Log should be applied, for example a the Resource ID of a Subscription or a Resource (such as a Storage Account).
- **criteria** - (Required) A **criteria** block as defined below.
- **action** - (Optional) One or more **action** blocks as defined below.
- **enabled** - (Optional) Should this Activity Log Alert be enabled? Defaults to **true**.
- **description** - (Optional) The description of this activity log alert.
- **tags** - (Optional) A mapping of tags to assign to the resource.

An **action** block supports the following:

- **action_group_id** - (Required) The ID of the Action Group can be sourced from the **azure_rm_monitor_action_group** resource.
- **webhook_properties** - (Optional) The map of custom string properties to include with the post operation. These data are appended to the webhook payload.

A **criteria** block supports the following:

- **category** - (Required) The category of the operation. Possible values are Administrative, Autoscale, Policy, Recommendation, Security and Service Health.
- **operation_name** - (Optional) The Resource Manager Role-Based Access Control operation name. Supported operation should be of the form: `<resourceProvider>/<resourceType>/<operation>`.

- **resource_provider** - (Optional) The name of the resource provider monitored by the activity log alert.
- **resource_type** - (Optional) The resource type monitored by the activity log alert.
- **resource_group** - (Optional) The name of resource group monitored by the activity log alert.
- **resource_id** - (Optional) The specific resource monitored by the activity log alert. It should be within one of the **scopes**.
- **caller** - (Optional) The email address or Azure Active Directory identifier of the user who performed the operation.
- **level** - (Optional) The severity level of the event. Possible values are **Verbose**, **Informational**, **Warning**, **Error**, and **Critical**.
- **status** - (Optional) The status of the event. For example, **Started**, **Failed**, or **Succeeded**.
- **sub_status** - (Optional) The sub status of the event.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the activity log alert.

» Import

Activity log alerts can be imported using the **resource id**, e.g.

```
terraform import azurerm_monitor_activity_log_alert.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_monitor_autoscale_setting

Manages a AutoScale Setting which can be applied to Virtual Machine Scale Sets, App Services and other scalable resources.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "autoscalingTest"
  location = "West US"
}

resource "azurerm_virtual_machine_scale_set" "test" {
  # ...
}
```

```

}

resource "azurerm_monitor_autoscale_setting" "test" {
  name = "myAutoscaleSetting"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location = "${azurerm_resource_group.test.location}"
  target_resource_id = "${azurerm_virtual_machine_scale_set.test.id}"

  profile {
    name = "defaultProfile"

    capacity {
      default = 1
      minimum = 1
      maximum = 10
    }

    rule {
      metric_trigger {
        metric_name = "Percentage CPU"
        metric_resource_id = "${azurerm_virtual_machine_scale_set.test.id}"
        time_grain = "PT1M"
        statistic = "Average"
        time_window = "PT5M"
        time_aggregation = "Average"
        operator = "GreaterThan"
        threshold = 75
      }

      scale_action {
        direction = "Increase"
        type = "ChangeCount"
        value = "1"
        cooldown = "PT1M"
      }
    }
  }

  rule {
    metric_trigger {
      metric_name = "Percentage CPU"
      metric_resource_id = "${azurerm_virtual_machine_scale_set.test.id}"
      time_grain = "PT1M"
      statistic = "Average"
      time_window = "PT5M"
      time_aggregation = "Average"
      operator = "LessThan"
    }
  }
}

```

```

        threshold      = 25
    }

    scale_action {
        direction = "Decrease"
        type      = "ChangeCount"
        value     = "1"
        cooldown  = "PT1M"
    }
}

notification {
    email {
        send_to_subscription_administrator = true
        send_to_subscription_co_administrator = true
        custom_emails                      = ["admin@contoso.com"]
    }
}
}

```

» Example Usage (repeating on weekends)

```

resource "azurerm_resource_group" "test" {
    name     = "autoscalingTest"
    location = "West US"
}

resource "azurerm_virtual_machine_scale_set" "test" {
    # ...
}

resource "azurerm_monitor_autoscale_setting" "test" {
    name                        = "myAutoscaleSetting"
    resource_group_name       = "${azurerm_resource_group.test.name}"
    location                  = "${azurerm_resource_group.test.location}"
    target_resource_id        = "${azurerm_virtual_machine_scale_set.test.id}"

    profile {
        name = "Weekends"

        capacity {
            default = 1
            minimum = 1
            maximum = 10
        }
    }
}

```

```

}

rule {
  metric_trigger {
    metric_name      = "Percentage CPU"
    metric_resource_id = "${azurerm_virtual_machine_scale_set.test.id}"
    time_grain       = "PT1M"
    statistic         = "Average"
    time_window       = "PT5M"
    time_aggregation  = "Average"
    operator          = "GreaterThan"
    threshold         = 90
  }

  scale_action {
    direction = "Increase"
    type      = "ChangeCount"
    value     = "2"
    cooldown  = "PT1M"
  }
}

rule {
  metric_trigger {
    metric_name      = "Percentage CPU"
    metric_resource_id = "${azurerm_virtual_machine_scale_set.test.id}"
    time_grain       = "PT1M"
    statistic         = "Average"
    time_window       = "PT5M"
    time_aggregation  = "Average"
    operator          = "LessThan"
    threshold         = 10
  }

  scale_action {
    direction = "Decrease"
    type      = "ChangeCount"
    value     = "2"
    cooldown  = "PT1M"
  }
}

recurrence {
  frequency = "Week"
  timezone  = "Pacific Standard Time"
  days      = ["Saturday", "Sunday"]
}

```

```

        hours      = [12]
        minutes    = [0]
    }
}

notification {
    email {
        send_to_subscription_administrator = true
        send_to_subscription_co_administrator = true
        custom_emails                     = ["admin@contoso.com"]
    }
}
}

```

» Example Usage (for fixed dates)

```

resource "azurerm_resource_group" "test" {
    name      = "autoscalingTest"
    location = "West US"
}

resource "azurerm_virtual_machine_scale_set" "test" {
    # ...
}

resource "azurerm_monitor_autoscale_setting" "test" {
    name                = "myAutoscaleSetting"
    enabled             = true
    resource_group_name = "${azurerm_resource_group.test.name}"
    location            = "${azurerm_resource_group.test.location}"
    target_resource_id  = "${azurerm_virtual_machine_scale_set.test.id}"
}

profile {
    name = "forJuly"

    capacity {
        default = 1
        minimum = 1
        maximum = 10
    }

    rule {
        metric_trigger {
            metric_name      = "Percentage CPU"
            metric_resource_id = "${azurerm_virtual_machine_scale_set.test.id}"

```

```

        time_grain          = "PT1M"
        statistic           = "Average"
        time_window         = "PT5M"
        time_aggregation    = "Average"
        operator            = "GreaterThan"
        threshold           = 90
    }

    scale_action {
        direction = "Increase"
        type      = "ChangeCount"
        value     = "2"
        cooldown  = "PT1M"
    }
}

rule {
    metric_trigger {
        metric_name          = "Percentage CPU"
        metric_resource_id   = "${azurerm_virtual_machine_scale_set.test.id}"
        time_grain           = "PT1M"
        statistic            = "Average"
        time_window          = "PT5M"
        time_aggregation     = "Average"
        operator             = "LessThan"
        threshold            = 10
    }

    scale_action {
        direction = "Decrease"
        type      = "ChangeCount"
        value     = "2"
        cooldown  = "PT1M"
    }
}

fixed_date {
    timezone = "Pacific Standard Time"
    start    = "2020-07-01T00:00:00Z"
    end      = "2020-07-31T23:59:59Z"
}

notification {
    email {
        send_to_subscription_administrator = true
    }
}

```

```

        send_to_subscription_co_administrator = true
        custom_emails                        = ["admin@contoso.com"]
    }
}
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the AutoScale Setting. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the Resource Group in the AutoScale Setting should be created. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the AutoScale Setting should exist. Changing this forces a new resource to be created.
- **profile** - (Required) Specifies one or more (up to 20) **profile** blocks as defined below.
- **target_resource_id** - (Required) Specifies the resource ID of the resource that the autoscale setting should be added to.
- **enabled** - (Optional) Specifies whether automatic scaling is enabled for the target resource. Defaults to **true**.
- **notification** - (Optional) Specifies a **notification** block as defined below.
- **tags** - (Optional) A mapping of tags to assign to the resource.

A **profile** block supports the following:

- **name** - (Required) Specifies the name of the profile.
 - **capacity** - (Required) A **capacity** block as defined below.
 - **rule** - (Required) One or more (up to 10) **rule** blocks as defined below.
 - **fixed_date** - (Optional) A **fixed_date** block as defined below. This cannot be specified if a **recurrence** block is specified.
 - **recurrence** - (Optional) A **recurrence** block as defined below. This cannot be specified if a **fixed_date** block is specified.
-

A **capacity** block supports the following:

- **default** - (Required) The number of instances that are available for scaling if metrics are not available for evaluation. The default is only used if the current instance count is lower than the default. Valid values are between 0 and 1000.
- **maximum** - (Required) The maximum number of instances for this resource. Valid values are between 0 and 1000.

NOTE: The maximum number of instances is also limited by the amount of Cores available in the subscription.

- **minimum** - (Required) The minimum number of instances for this resource. Valid values are between 0 and 1000.

A **rule** block supports the following:

- **metric_trigger** - (Required) A **metric_trigger** block as defined below.
- **scale_action** - (Required) A **scale_action** block as defined below.

A **metric_trigger** block supports the following:

- **metric_name** - (Required) The name of the metric that defines what the rule monitors, such as **Percentage CPU for Virtual Machine Scale Sets** and **CpuPercentage for App Service Plan**.

NOTE: The allowed value of **metric_name** highly depends on the targeting resource type, please visit [Supported metrics with Azure Monitor](#) for more details.

- **metric_resource_id** - (Required) The ID of the Resource which the Rule monitors.
- **operator** - (Required) Specifies the operator used to compare the metric data and threshold. Possible values are: **Equals**, **NotEquals**, **GreaterThan**, **GreaterThanOrEqual**, **LessThan**, **LessThanOrEqual**.
- **statistic** - (Required) Specifies how the metrics from multiple instances are combined. Possible values are **Average**, **Min** and **Max**.
- **time_aggregation** - (Required) Specifies how the data that's collected should be combined over time. Possible values include **Average**, **Count**, **Maximum**, **Minimum**, **Last** and **Total**. Defaults to **Average**.
- **time_grain** - (Required) Specifies the granularity of metrics that the rule monitors, which must be one of the pre-defined values returned from the metric definitions for the metric. This value must be between 1 minute and 12 hours and be formatted as an ISO 8601 string.

- **time_window** - (Required) Specifies the time range for which data is collected, which must be greater than the delay in metric collection (which varies from resource to resource). This value must be between 5 minutes and 12 hours and be formatted as an ISO 8601 string.
- **threshold** - (Required) Specifies the threshold of the metric that triggers the scale action.

A **scale_action** block supports the following:

- **cooldown** - (Required) The amount of time to wait since the last scaling action before this action occurs. Must be between 1 minute and 1 week and formatted as a ISO 8601 string.
- **direction** - (Required) The scale direction. Possible values are **Increase** and **Decrease**.
- **type** - (Required) The type of action that should occur. Possible values are **ChangeCount**, **ExactCount** and **PercentChangeCount**.
- **value** - (Required) The number of instances involved in the scaling action. Defaults to 1.

A **fixed_date** block supports the following:

- **end** - (Required) Specifies the end date for the profile, formatted as an RFC3339 date string.
- **start** - (Required) Specifies the start date for the profile, formatted as an RFC3339 date string.
- **timezone** (Optional) The Time Zone of the **start** and **end** times. A list of possible values can be found here. Defaults to **UTC**.

A **recurrence** block supports the following:

- **timezone** - (Required) The Time Zone used for the **hours** field. A list of possible values can be found here. Defaults to **UTC**.
- **days** - (Required) A list of days that this profile takes effect on. Possible values include **Monday**, **Tuesday**, **Wednesday**, **Thursday**, **Friday**, **Saturday** and **Sunday**.
- **hours** - (Required) A list containing a single item, which specifies the Hour interval at which this recurrence should be triggered (in 24-hour time). Possible values are from 0 to 23.
- **minutes** - (Required) A list containing a single item which specifies the Minute interval at which this recurrence should be triggered.

A **notification** block supports the following:

- **email** - (Required) A **email** block as defined below.
 - **webhook** - (Optional) One or more **webhook** blocks as defined below.
-

A **email** block supports the following:

- **send_to_subscription_administrator** - (Optional) Should email notifications be sent to the subscription administrator? Defaults to **false**.
 - **send_to_subscription_co_administrator** - (Optional) Should email notifications be sent to the subscription co-administrator? Defaults to **false**.
 - **custom_emails** - (Optional) Specifies a list of custom email addresses to which the email notifications will be sent.
-

A **webhook** block supports the following:

- **service_uri** - (Required) The HTTPS URI which should receive scale notifications.
- **properties** - (Optional) A map of settings.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the AutoScale Setting.

» Import

AutoScale Setting can be imported using the **resource id**, e.g.

```
terraform import azurerm_monitor_autoscale_setting.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm__monitor__diagnostic__setting

Manages a Diagnostic Setting for an existing Resource.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "example-resources"
  location  = "West Europe"
}

data "azurerm_storage_account" "test" {
  name                = "examplestoracc"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

data "azurerm_key_vault" "test" {
  name                = "example-vault"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_monitor_diagnostic_setting" "test" {
  name                  = "example"
  target_resource_id    = "${data.azurerm_key_vault.test.id}"
  storage_account_id    = "${data.azurerm_storage_account.test.id}"

  log {
    category = "AuditEvent"
    enabled  = false

    retention_policy {
      enabled = false
    }
  }

  metric {
    category = "AllMetrics"

    retention_policy {
      enabled = false
    }
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Diagnostic Setting. Changing

this forces a new resource to be created.

- **target_resource_id** - (Required) The ID of an existing Resource on which to configure Diagnostic Settings. Changing this forces a new resource to be created.
- **eventhub_name** - (Optional) Specifies the name of the Event Hub where Diagnostics Data should be sent. Changing this forces a new resource to be created.

NOTE: If this isn't specified then the default Event Hub will be used.

- **eventhub_authorization_rule_id** - (Optional) Specifies the ID of an Event Hub Namespace Authorization Rule used to send Diagnostics Data. Changing this forces a new resource to be created.

NOTE: This can be sourced from the `azurerm_eventhub_namespace_authorization_rule` resource and is different from a `azurerm_eventhub_authorization_rule` resource.

NOTE: One of `eventhub_authorization_rule_id`, `log_analytics_workspace_id` and `storage_account_id` must be specified.

- **log** - (Optional) One or more `log` blocks as defined below.

NOTE: At least one `log` or `metric` block must be specified.

- **log_analytics_workspace_id** - (Optional) Specifies the ID of a Log Analytics Workspace where Diagnostics Data should be sent. Changing this forces a new resource to be created.

NOTE: One of `eventhub_authorization_rule_id`, `log_analytics_workspace_id` and `storage_account_id` must be specified.

- **metric** - (Optional) One or more `metric` blocks as defined below.

NOTE: At least one `log` or `metric` block must be specified.

- **storage_account_id** - (Optional) With this parameter you can specify a storage account which should be used to send the logs to. Parameter must be a valid Azure Resource ID. Changing this forces a new resource to be created.

NOTE: One of `eventhub_authorization_rule_id`, `log_analytics_workspace_id` and `storage_account_id` must be specified.

A `log` block supports the following:

- **category** - (Required) The name of a Diagnostic Log Category for this Resource.

NOTE: The Log Categories available vary depending on the Resource being used. You may wish to use the `azurerm_monitor_diagnostic_categories` Data Source to identify which categories are available for a given Resource.

- `retention_policy` - (Required) A `retention_policy` block as defined below.
 - `enabled` - (Optional) Is this Diagnostic Log enabled? Defaults to `true`.
-

A `metric` block supports the following:

- `category` - (Required) The name of a Diagnostic Metric Category for this Resource.

NOTE: The Metric Categories available vary depending on the Resource being used. You may wish to use the `azurerm_monitor_diagnostic_categories` Data Source to identify which categories are available for a given Resource.

- `retention_policy` - (Required) A `retention_policy` block as defined below.
 - `enabled` - (Optional) Is this Diagnostic Metric enabled? Defaults to `true`.
-

A `retention_policy` block supports the following:

- `enabled` - (Required) Is this Retention Policy enabled?
- `days` - (Optional) The number of days for which this Retention Policy should apply.

NOTE: Setting this to 0 will retain the events indefinitely.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the Diagnostic Setting.

» Import

Diagnostic Settings can be imported using the `resource id`, e.g.

```
terraform import azurerm_monitor_diagnostics.test /subscriptions/XXX/resourcegroups/resource
```

NOTE: This is a Terraform specific Resource ID which uses the format `{resourceId}|{diagnosticSettingName}`

» azurerm_monitor_log_profile

Manages a Log Profile. A Log Profile configures how Activity Logs are exported.

NOTE: It's only possible to configure one Log Profile per Subscription. If you are trying to create more than one Log Profile, an error with `StatusCode=409` will occur.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name       = "logprofiletest-rg"
  location   = "eastus"
}

resource "azurerm_storage_account" "test" {
  name                         = "afscsdfytw"
  resource_group_name         = "${azurerm_resource_group.test.name}"
  location                    = "${azurerm_resource_group.test.location}"
  account_tier                 = "Standard"
  account_replication_type    = "GRS"
}

resource "azurerm_eventhub_namespace" "test" {
  name                       = "logprofileeventhub"
  location                   = "${azurerm_resource_group.test.location}"
  resource_group_name       = "${azurerm_resource_group.test.name}"
  sku                        = "Standard"
  capacity                   = 2
}

resource "azurerm_monitor_log_profile" "test" {
  name = "default"

  categories = [
    "Action",
    "Delete",
    "Write",
  ]

  locations = [
    "westus",
    "global",
  ]
}
```

```

# RootManageSharedAccessKey is created by default with listen, send, manage permissions
servicebus_rule_id = "${azurerm_eventhub_namespace.test.id}/authorizationrules/RootManageSharedAccessKey"
storage_account_id = "${azurerm_storage_account.test.id}"

retention_policy {
  enabled = true
  days    = 7
}
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Log Profile. Changing this forces a new resource to be created.
- **categories** - (Required) List of categories of the logs.
- **locations** - (Required) List of regions for which Activity Log events are stored or streamed.
- **storage_account_id** - (Optional) The resource ID of the storage account in which the Activity Log is stored. At least one of **storage_account_id** or **servicebus_rule_id** must be set.
- **servicebus_rule_id** - (Optional) The service bus (or event hub) rule ID of the service bus (or event hub) namespace in which the Activity Log is streamed to. At least one of **storage_account_id** or **servicebus_rule_id** must be set.
- **retention_policy** - (Required) A **retention_policy** block as documented below. A retention policy for how long Activity Logs are retained in the storage account.

The **retention_policy** block supports:

- **enabled** - (Required) A boolean value to indicate whether the retention policy is enabled.
- **days** - (Optional) The number of days for the retention policy. Defaults to 0.

» Attributes Reference

The following attributes are exported:

- id - The Log Profile resource ID.

» Import

A Log Profile can be imported using the `resource id`, e.g.

```
terraform import azurerm_monitor_log_profile.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_monitor_metric_alert

Manages a Metric Alert within Azure Monitor.

» Example Usage

```
resource "azurerm_resource_group" "main" {
  name     = "example-resources"
  location = "West US"
}

resource "azurerm_storage_account" "to_monitor" {
  name                        = "examplestorageaccount"
  resource_group_name        = "${azurerm_resource_group.main.name}"
  location                   = "${azurerm_resource_group.main.location}"
  account_tier               = "Standard"
  account_replication_type   = "LRS"
}

resource "azurerm_monitor_action_group" "main" {
  name                = "example-actiongroup"
  resource_group_name = "${azurerm_resource_group.main.name}"
  short_name          = "exampleact"

  webhook_receiver {
    name       = "callmyapi"
    service_uri = "http://example.com/alert"
  }
}

resource "azurerm_monitor_metric_alert" "test" {
  name                = "example-metricalert"
  resource_group_name = "${azurerm_resource_group.main.name}"
  scopes              = ["${azurerm_storage_account.to_monitor.id}"]
  description         = "Action will be triggered when Transactions count is greater than 50"
}
```

```

criteria {
  metric_namespace = "Microsoft.Storage/storageAccounts"
  metric_name       = "Transactions"
  aggregation       = "Total"
  operator          = "GreaterThan"
  threshold         = 50

  dimension {
    name      = "ApiName"
    operator  = "Include"
    values    = ["*"]
  }
}

action {
  action_group_id = "${azurerm_monitor_action_group.main.id}"
}
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Metric Alert. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the Metric Alert instance.
- **scopes** - (Required) The resource ID at which the metric criteria should be applied.
- **criteria** - (Required) One or more **criteria** blocks as defined below.
- **action** - (Optional) One or more **action** blocks as defined below.
- **enabled** - (Optional) Should this Metric Alert be enabled? Defaults to **true**.
- **auto_mitigate** - (Optional) Should the alerts in this Metric Alert be auto resolved? Defaults to **false**.
- **description** - (Optional) The description of this Metric Alert.
- **frequency** - (Optional) The evaluation frequency of this Metric Alert, represented in ISO 8601 duration format. Possible values are PT1M, PT5M, PT15M, PT30M and PT1H. Defaults to PT1M.
- **severity** - (Optional) The severity of this Metric Alert. Possible values are 0, 1, 2, 3 and 4. Defaults to 3.
- **window_size** - (Optional) The period of time that is used to monitor alert activity, represented in ISO 8601 duration format. This value must be greater than **frequency**. Possible values are PT1M, PT5M, PT15M, PT30M,

PT1H, PT6H, PT12H and P1D. Defaults to PT5M.

- **tags** - (Optional) A mapping of tags to assign to the resource.

An **action** block supports the following:

- **action_group_id** - (Required) The ID of the Action Group can be sourced from the **azurerm_monitor_action_group** resource
- **webhook_properties** - (Optional) The map of custom string properties to include with the post operation. These data are appended to the webhook payload.

A **criteria** block supports the following:

- **metric_namespace** - (Required) One of the metric namespaces to be monitored.
- **metric_name** - (Required) One of the metric names to be monitored.
- **aggregation** - (Required) The statistic that runs over the metric values. Possible values are **Average**, **Count**, **Minimum**, **Maximum** and **Total**.
- **operator** - (Required) The criteria operator. Possible values are **Equals**, **NotEquals**, **GreaterThan**, **GreaterThanOrEqual**, **LessThan** and **LessThanOrEqual**.
- **threshold** - (Required) The criteria threshold value that activates the alert.
- **dimension** - (Optional) One or more **dimension** blocks as defined below.

A **dimension** block supports the following:

- **name** - (Required) One of the dimension names.
- **operator** - (Required) The dimension operator. Possible values are **Include** and **Exclude**.
- **values** - (Required) The list of dimension values.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the metric alert.

» Import

Metric Alerts can be imported using the **resource id**, e.g.

```
terraform import azurerm_monitor_metric_alert.main /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_monitor_metric_alertrule`

Manages a metric-based alert rule in Azure Monitor.

» Example Usage (CPU Percentage of a virtual machine)

```
resource "azurerm_monitor_metric_alertrule" "test" {
  name                        = "${azurerm_virtual_machine.test.name}-cpu"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  location                   = "${azurerm_resource_group.test.location}"

  description = "An alert rule to watch the metric Percentage CPU"

  enabled = true

  resource_id = "${azurerm_virtual_machine.test.id}"
  metric_name = "Percentage CPU"
  operator    = "GreaterThan"
  threshold   = 75
  aggregation = "Average"
  period      = "PT5M"

  email_action {
    send_to_service_owners = false

    custom_emails = [
      "some.user@example.com",
    ]
  }

  webhook_action {
    service_uri = "https://example.com/some-url"

    properties = {
      severity          = "incredible"
      acceptance_test = "true"
    }
  }
}
```

» Example Usage (Storage usage of a SQL Database)

```
resource "azurerm_monitor_metric_alertrule" "test" {
  name                        = "${azurerm_sql_database.test.name}-storage"
```

```

resource_group_name = "${azurerm_resource_group.test.name}"
location            = "${azurerm_resource_group.test.location}"

description = "An alert rule to watch the metric Storage"

enabled = true

resource_id = "${azurerm_sql_database.test.id}"
metric_name = "storage"
operator     = "GreaterThan"
threshold    = 1073741824
aggregation  = "Maximum"
period       = "PT10M"

email_action {
  send_to_service_owners = false

  custom_emails = [
    "some.user@example.com",
  ]
}

webhook_action {
  service_uri = "https://example.com/some-url"

  properties = {
    severity          = "incredible"
    acceptance_test   = "true"
  }
}
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the alert rule. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the alert rule. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.

- **description** - (Optional) A verbose description of the alert rule that will be included in the alert email.
- **enabled** - (Optional) If **true**, the alert rule is enabled. Defaults to **true**.

-
- **resource_id** - (Required) The ID of the resource monitored by the alert rule.
 - **metric_name** - (Required) The metric that defines what the rule monitors.

For a comprehensive reference of supported **metric_name** values for types of **resource** refer to Supported metrics with Azure Monitor in the Azure documentation. In the referred table, the column "Metric" corresponds to supported values for **metric_name**.

- **operator** - (Required) The operator used to compare the metric data and the threshold. Possible values are **GreaterThan**, **GreaterThanOrEqual**, **LessThan**, and **LessThanOrEqual**.
- **threshold** - (Required) The threshold value that activates the alert.
- **period** - (Required) The period of time formatted in ISO 8601 duration format that is used to monitor the alert activity based on the threshold. The period must be between 5 minutes and 1 day.
- **aggregation** - (Required) Defines how the metric data is combined over time. Possible values are **Average**, **Minimum**, **Maximum**, **Total**, and **Last**.

-
- **email_action** - (Optional) A **email_action** block as defined below.
 - **webhook_action** - (Optional) A **webhook_action** block as defined below.
 - **tags** - (Optional) A mapping of tags to assign to the resource. Changing this forces a new resource to be created.

email_action supports the following:

- **send_to_service_owners** - (Optional) If **true**, the administrators (service and co-administrators) of the subscription are notified when the alert is triggered. Defaults to **false**.
- **custom_emails** - (Optional) A list of email addresses to be notified when the alert is triggered.

webhook_action supports the following:

- **service_uri** - (Required) The service uri of the webhook to POST the notification when the alert is triggered.

- **properties** - (Optional) A dictionary of custom properties to include with the webhook POST operation payload.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the alert rule.

» Import

Metric Alert Rules can be imported using the **resource id**, e.g.

```
terraform import azurerm_monitor_metric_alertrule.alertrule1 /subscriptions/00000000-0000-00
```

» azurerm__metric__alertrule

Manages a metric-based alert rule in Azure Monitor.

NOTE: This resource has been deprecated in favour of the `azurerm_monitor_metric_alertrule` resource and will be removed in the next major version of the AzureRM Provider. The new resource shares the same fields as this one, and information on migrating across can be found in this guide.

» Example Usage (CPU Percentage of a virtual machine)

```
resource "azurerm_metric_alertrule" "test" {
  name                        = "${azurerm_virtual_machine.test.name}-cpu"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  location                   = "${azurerm_resource_group.test.location}"

  description = "An alert rule to watch the metric Percentage CPU"

  enabled = true

  resource_id = "${azurerm_virtual_machine.test.id}"
  metric_name = "Percentage CPU"
  operator    = "GreaterThan"
  threshold   = 75
  aggregation = "Average"
  period      = "PT5M"

  email_action {
```

```

    send_to_service_owners = false

    custom_emails = [
        "some.user@example.com",
    ]
}

webhook_action {
    service_uri = "https://example.com/some-url"

    properties = {
        severity      = "incredible"
        acceptance_test = "true"
    }
}
}

```

» Example Usage (Storage usage of a SQL Database)

```

resource "azurerm_metric_alertrule" "test" {
    name                = "${azurerm_sql_database.test.name}-storage"
    resource_group_name = "${azurerm_resource_group.test.name}"
    location            = "${azurerm_resource_group.test.location}"

    description = "An alert rule to watch the metric Storage"

    enabled = true

    resource_id = "${azurerm_sql_database.test.id}"
    metric_name = "storage"
    operator    = "GreaterThan"
    threshold   = 1073741824
    aggregation = "Maximum"
    period      = "PT10M"

    email_action {
        send_to_service_owners = false

        custom_emails = [
            "some.user@example.com",
        ]
    }

    webhook_action {
        service_uri = "https://example.com/some-url"
    }
}

```



```

    properties = {
        severity      = "incredible"
        acceptance_test = "true"
    }
}
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the alert rule. Changing this forces a new resource to be created.
 - **resource_group_name** - (Required) The name of the resource group in which to create the alert rule. Changing this forces a new resource to be created.
 - **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
 - **description** - (Optional) A verbose description of the alert rule that will be included in the alert email.
 - **enabled** - (Optional) If **true**, the alert rule is enabled. Defaults to **true**.
-
- **resource_id** - (Required) The ID of the resource monitored by the alert rule.
 - **metric_name** - (Required) The metric that defines what the rule monitors.

For a comprehensive reference of supported **metric_name** values for types of **resource** refer to Supported metrics with Azure Monitor in the Azure documentation. In the referred table, the column "Metric" corresponds to supported values for **metric_name**.

- **operator** - (Required) The operator used to compare the metric data and the threshold. Possible values are **GreaterThan**, **GreaterThanOrEqual**, **LessThan**, and **LessThanOrEqual**.
- **threshold** - (Required) The threshold value that activates the alert.
- **period** - (Required) The period of time formatted in ISO 8601 duration format that is used to monitor the alert activity based on the threshold. The period must be between 5 minutes and 1 day.
- **aggregation** - (Required) Defines how the metric data is combined over time. Possible values are **Average**, **Minimum**, **Maximum**, **Total**, and **Last**.

-
- **email_action** - (Optional) A **email_action** block as defined below.
 - **webhook_action** - (Optional) A **webhook_action** block as defined below.
 - **tags** - (Optional) A mapping of tags to assign to the resource. Changing this forces a new resource to be created.
-

email_action supports the following:

- **send_to_service_owners** - (Optional) If **true**, the administrators (service and co-administrators) of the subscription are notified when the alert is triggered. Defaults to **false**.
 - **custom_emails** - (Optional) A list of email addresses to be notified when the alert is triggered.
-

webhook_action supports the following:

- **service_uri** - (Required) The service uri of the webhook to POST the notification when the alert is triggered.
- **properties** - (Optional) A dictionary of custom properties to include with the webhook POST operation payload.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the alert rule.

» Import

Metric Alert Rules can be imported using the **resource id**, e.g.

```
terraform import azurerm_metric_alertrule.alertrule1 /subscriptions/00000000-0000-0000-0000-
```

» **azurerm_application_gateway**

Manages an Application Gateway.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "example-resources"
  location  = "West US"
}

resource "azurerm_virtual_network" "test" {
  name                = "example-network"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location            = "${azurerm_resource_group.test.location}"
  address_space       = ["10.254.0.0/16"]
}

resource "azurerm_subnet" "frontend" {
  name                = "frontend"
  resource_group_name = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix      = "10.254.0.0/24"
}

resource "azurerm_subnet" "backend" {
  name                = "backend"
  resource_group_name = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix      = "10.254.2.0/24"
}

resource "azurerm_public_ip" "test" {
  name                = "example-pip"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location            = "${azurerm_resource_group.test.location}"
  allocation_method   = "Dynamic"
}

# since these variables are re-used - a locals block makes this more maintainable
locals {
  backend_address_pool_name = "${azurerm_virtual_network.test.name}-beap"
  frontend_port_name        = "${azurerm_virtual_network.test.name}-feport"
  frontend_ip_configuration_name = "${azurerm_virtual_network.test.name}-feip"
  http_setting_name         = "${azurerm_virtual_network.test.name}-be-htst"
  listener_name             = "${azurerm_virtual_network.test.name}-httplstn"
  request_routing_rule_name = "${azurerm_virtual_network.test.name}-rqrt"
  redirect_configuration_name = "${azurerm_virtual_network.test.name}-rdrcfg"
}
```

```

resource "azurerm_application_gateway" "network" {
  name                = "example-appgateway"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location             = "${azurerm_resource_group.test.location}"

  sku {
    name      = "Standard_Small"
    tier       = "Standard"
    capacity  = 2
  }

  gateway_ip_configuration {
    name      = "my-gateway-ip-configuration"
    subnet_id = "${azurerm_subnet.frontend.id}"
  }

  frontend_port {
    name = "${local.frontend_port_name}"
    port = 80
  }

  frontend_ip_configuration {
    name                = "${local.frontend_ip_configuration_name}"
    public_ip_address_id = "${azurerm_public_ip.test.id}"
  }

  backend_address_pool {
    name = "${local.backend_address_pool_name}"
  }

  backend_http_settings {
    name                = "${local.http_setting_name}"
    cookie_based_affinity = "Disabled"
    path                = "/path1/"
    port                = 80
    protocol             = "Http"
    request_timeout     = 1
  }

  http_listener {
    name                = "${local.listener_name}"
    frontend_ip_configuration_name = "${local.frontend_ip_configuration_name}"
    frontend_port_name = "${local.frontend_port_name}"
    protocol            = "Http"
  }
}

```

```

request_routing_rule {
  name = "${local.request_routing_rule_name}"
  rule_type = "Basic"
  http_listener_name = "${local.listener_name}"
  backend_address_pool_name = "${local.backend_address_pool_name}"
  backend_http_settings_name = "${local.http_setting_name}"
}
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Application Gateway. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to the Application Gateway should exist. Changing this forces a new resource to be created.
- **location** - (Required) The Azure region where the Application Gateway should exist. Changing this forces a new resource to be created.
- **backend_address_pool** - (Required) One or more **backend_address_pool** blocks as defined below.
- **backend_http_settings** - (Required) One or more **backend_http_settings** blocks as defined below.
- **frontend_ip_configuration** - (Required) One or more **frontend_ip_configuration** blocks as defined below.
- **frontend_port** - (Required) One or more **frontend_port** blocks as defined below.
- **gateway_ip_configuration** - (Required) One or more **gateway_ip_configuration** blocks as defined below.
- **http_listener** - (Required) One or more **http_listener** blocks as defined below.
- **request_routing_rule** - (Required) One or more **request_routing_rule** blocks as defined below.
- **sku** - (Required) A **sku** block as defined below.
- **zones** - (Optional) A collection of availability zones to spread the Application Gateway over.

Please Note: Availability Zones are only supported in several regions at this time. They are also only supported for v2 SKUs

-
- **authentication_certificate** - (Optional) One or more **authentication_certificate** blocks as defined below.
 - **disabled_ssl_protocols** - (Optional) A list of SSL Protocols which should be disabled on this Application Gateway. Possible values are TLSv1_0, TLSv1_1 and TLSv1_2.
 - **enable_http2** - (Optional) Is HTTP2 enabled on the application gateway resource? Defaults to **false**.
 - **probe** - (Optional) One or more **probe** blocks as defined below.
 - **ssl_certificate** - (Optional) One or more **ssl_certificate** blocks as defined below.
 - **tags** - (Optional) A mapping of tags to assign to the resource.
 - **url_path_map** - (Optional) One or more **url_path_map** blocks as defined below.
 - **waf_configuration** - (Optional) A **waf_configuration** block as defined below.
 - **custom_error_configuration** - (Optional) One or more **custom_error_configuration** blocks as defined below.
 - **redirect_configuration** - (Optional) A **redirect_configuration** block as defined below.

A **authentication_certificate** block supports the following:

- **name** - (Required) The Name of the Authentication Certificate to use.
- **data** - (Required) The contents of the Authentication Certificate which should be used.

A **authentication_certificate** block, within the **backend_http_settings** block supports the following:

- **name** - (Required) The name of the Authentication Certificate.

A **backend_address_pool** block supports the following:

- **name** - (Required) The name of the Backend Address Pool.

- `fqdns` - (Optional) A list of FQDN's which should be part of the Backend Address Pool.
- `fqdn_list` - (Optional **Deprecated**) A list of FQDN's which should be part of the Backend Address Pool. This field has been deprecated in favour of `fqdns` and will be removed in v2.0 of the AzureRM Provider.
- `ip_addresses` - (Optional) A list of IP Addresses which should be part of the Backend Address Pool.
- `ip_address_list` - (Optional **Deprecated**) A list of IP Addresses which should be part of the Backend Address Pool. This field has been deprecated in favour of `ip_addresses` and will be removed in v2.0 of the AzureRM Provider.

A `backend_http_settings` block supports the following:

- `cookie_based_affinity` - (Required) Is Cookie-Based Affinity enabled? Possible values are `Enabled` and `Disabled`.
- `name` - (Required) The name of the Backend HTTP Settings Collection.
- `path` - (Optional) The Path which should be used as a prefix for all HTTP requests.
- `port` - (Required) The port which should be used for this Backend HTTP Settings Collection.
- `probe_name` - (Required) The name of an associated HTTP Probe.
- `protocol` - (Required) The Protocol which should be used. Possible values are `Http` and `Https`.
- `request_timeout` - (Required) The request timeout in seconds, which must be between 1 and 86400 seconds.
- `host_name` - (Optional) Host header to be sent to the backend servers. Cannot be set if `pick_host_name_from_backend_address` is set to `true`.
- `pick_host_name_from_backend_address` - (Optional) Whether host header should be picked from the host name of the backend server. Defaults to `false`.
- `authentication_certificate` - (Optional) One or more `authentication_certificate` blocks.
- `connection_draining` - (Optional) A `connection_draining` block as defined below.

A `connection_draining` block supports the following:

- **enabled** - (Required) If connection draining is enabled or not.
 - **drain_timeout_sec** - (Required) The number of seconds connection draining is active. Acceptable values are from 1 second to 3600 seconds.
-

A **frontend_ip_configuration** block supports the following:

- **name** - (Required) The name of the Frontend IP Configuration.
- **subnet_id** - (Required) The ID of the Subnet which the Application Gateway should be connected to.
- **private_ip_address** - (Optional) The Private IP Address to use for the Application Gateway.
- **public_ip_address_id** - (Optional) The ID of a Public IP Address which the Application Gateway should use.

NOTE: The Allocation Method for this Public IP Address should be set to **Dynamic**.

- **private_ip_address_allocation** - (Optional) The Allocation Method for the Private IP Address. Possible values are **Dynamic** and **Static**.
-

A **frontend_port** block supports the following:

- **name** - (Required) The name of the Frontend Port.
 - **port** - (Required) The port used for this Frontend Port.
-

A **gateway_ip_configuration** block supports the following:

- **name** - (Required) The Name of this Gateway IP Configuration.
 - **subnet_id** - (Required) The ID of a Subnet.
-

A **http_listener** block supports the following:

- **name** - (Required) The Name of the HTTP Listener.
- **frontend_ip_configuration_name** - (Required) The Name of the Frontend IP Configuration used for this HTTP Listener.
- **frontend_port_name** - (Required) The Name of the Frontend Port use for this HTTP Listener.
- **host_name** - (Optional) The Hostname which should be used for this HTTP Listener.

- **protocol** - (Required) The Protocol to use for this HTTP Listener. Possible values are **Http** and **Https**.
 - **require_sni** - (Optional) Should Server Name Indication be Required? Defaults to **false**.
 - **ssl_certificate_name** - (Optional) The name of the associated SSL Certificate which should be used for this HTTP Listener.
 - **custom_error_configuration** - (Optional) One or more **custom_error_configuration** blocks as defined below.
-

A **match** block supports the following:

- **body** - (Optional) A snippet from the Response Body which must be present in the Response. Defaults to *****.
 - **status_code** - (Optional) A list of allowed status codes for this Health Probe.
-

A **path_rule** block supports the following:

- **name** - (Required) The Name of the Path Rule.
 - **paths** - (Required) A list of Paths used in this Path Rule.
 - **backend_address_pool_name** - (Optional) The Name of the Backend Address Pool to use for this Path Rule. Cannot be set if **redirect_configuration_name** is set.
 - **backend_http_settings_name** - (Optional) The Name of the Backend HTTP Settings Collection to use for this Path Rule. Cannot be set if **redirect_configuration_name** is set.
 - **redirect_configuration_name** - (Optional) The Name of a Redirect Configuration to use for this Path Rule. Cannot be set if **backend_address_pool_name** or **backend_http_settings_name** is set.
-

A **probe** block support the following:

- **host** - (Optional) The Hostname used for this Probe. If the Application Gateway is configured for a single site, by default the Host name should be specified as '127.0.0.1', unless otherwise configured in custom probe. Cannot be set if **pick_host_name_from_backend_http_settings** is set to **true**.
- **interval** - (Required) The Interval between two consecutive probes in seconds. Possible values range from 1 second to a maximum of 86,400 seconds.

- **name** - (Required) The Name of the Probe.
- **protocol** - (Required) The Protocol used for this Probe. Possible values are **Http** and **Https**.
- **path** - (Required) The Path used for this Probe.
- **timeout** - (Required) The Timeout used for this Probe, which indicates when a probe becomes unhealthy. Possible values range from 1 second to a maximum of 86,400 seconds.
- **unhealthy_threshold** - (Required) The Unhealthy Threshold for this Probe, which indicates the amount of retries which should be attempted before a node is deemed unhealthy. Possible values are from 1 - 20 seconds.
- **pick_host_name_from_backend_http_settings** - (Optional) Whether the host header should be picked from the backend http settings. Defaults to **false**.
- **match** - (Optional) A **match** block as defined above.
- **minimum_servers** - (Optional) The minimum number of servers that are always marked as healthy. Defaults to 0.

A **request_routing_rule** block supports the following:

- **name** - (Required) The Name of this Request Routing Rule.
- **rule_type** - (Required) The Type of Routing that should be used for this Rule. Possible values are **Basic** and **PathBasedRouting**.
- **http_listener_name** - (Required) The Name of the HTTP Listener which should be used for this Routing Rule.
- **backend_address_pool_name** - (Optional) The Name of the Backend Address Pool which should be used for this Routing Rule. Cannot be set if **redirect_configuration_name** is set.
- **backend_http_settings_name** - (Optional) The Name of the Backend HTTP Settings Collection which should be used for this Routing Rule. Cannot be set if **redirect_configuration_name** is set.
- **redirect_configuration_name** - (Optional) The Name of the Redirect Configuration which should be used for this Routing Rule. Cannot be set if either **backend_address_pool_name** or **backend_http_settings_name** is set.
- **url_path_map_name** - (Optional) The Name of the URL Path Map which should be associated with this Routing Rule.

A **sku** block supports the following:

- **name** - (Required) The Name of the SKU to use for this Application Gateway. Possible values are **Standard_Small**, **Standard_Medium**, **Standard_Large**, **Standard_v2**, **WAF_Medium**, **WAF_Large**, and **WAF_v2**.
- **tier** - (Required) The Tier of the SKU to use for this Application Gateway. Possible values are **Standard**, **Standard_v2**, **WAF** and **WAF_v2**.
- **capacity** - (Required) The Capacity of the SKU to use for this Application Gateway - which must be between 1 and 10.

A **ssl_certificate** block supports the following:

- **name** - (Required) The Name of the SSL certificate that is unique within this Application Gateway
- **data** - (Required) PFX certificate.
- **password** - (Required) Password for the pfx file specified in data.

A **url_path_map** block supports the following:

- **name** - (Required) The Name of the URL Path Map.
- **default_backend_address_pool_name** - (Optional) The Name of the Default Backend Address Pool which should be used for this URL Path Map. Cannot be set if there are **path_rules** with re-direct configurations set.
- **default_backend_http_settings_name** - (Optional) The Name of the Default Backend HTTP Settings Collection which should be used for this URL Path Map. Cannot be set if there are **path_rules** with re-direct configurations set.
- **default_redirect_configuration_name** - (Optional) The Name of the Default Redirect Configuration which should be used for this URL Path Map. Cannot be set if there are **path_rules** with Backend Address Pool or HTTP Settings set.
- **path_rule** - (Required) One or more **path_rule** blocks as defined above.

A **waf_configuration** block supports the following:

- **enabled** - (Required) Is the Web Application Firewall be enabled?
- **firewall_mode** - (Required) The Web Application Firewall Mode. Possible values are **Detection** and **Prevention**.
- **rule_set_type** - (Required) The Type of the Rule Set used for this Web Application Firewall.

- **rule_set_version** - (Required) The Version of the Rule Set used for this Web Application Firewall.
- **file_upload_limit_mb** - (Optional) The File Upload Limit in MB. Accepted values are in the range 1MB to 500MB. Defaults to 100MB.
- **request_body_check** - (Optional) Is Request Body Inspection enabled? Defaults to **true**.
- **max_request_body_size_kb** - (Optional) The Maximum Request Body Size in KB. Accepted values are in the range 1KB to 128KB. Defaults to 128KB.

A **custom_error_configuration** block supports the following:

- **status_code** - (Required) Status code of the application gateway customer error. Possible values are **HttpStatus403** and **HttpStatus502**
- **custom_error_page_url** - (Required) Error page URL of the application gateway customer error.

A **redirect_configuration** block supports the following:

- **name** - (Required) Unique name of the redirect configuration block
- **redirect_type** - (Required) The type of redirect. Possible values are **Permanent**, **Temporary**, **Found** and **SeeOther**
- **target_listener_name** - (Optional) The name of the listener to redirect to. Cannot be set if **target_url** is set.
- **target_url** - (Optional) The Url to redirect the request to. Cannot be set if **target_listener_name** is set.
- **include_path** - (Optional) Whether or not to include the path in the redirected Url. Defaults to **false**
- **include_query_string** - (Optional) Whether or not to include the query string in the redirected Url. Default to **false**

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Application Gateway.
- **authentication_certificate** - A list of **authentication_certificate** blocks as defined below.

- `backend_address_pool` - A list of `backend_address_pool` blocks as defined below.
- `backend_http_settings` - A list of `backend_http_settings` blocks as defined below.
- `frontend_ip_configuration` - A list of `frontend_ip_configuration` blocks as defined below.
- `frontend_port` - A list of `frontend_port` blocks as defined below.
- `gateway_ip_configuration` - A list of `gateway_ip_configuration` blocks as defined below.
- `enable_http2` - (Optional) Is HTTP2 enabled on the application gateway resource? Defaults to `false`.
- `http_listener` - A list of `http_listener` blocks as defined below.
- `probe` - A probe block as defined below.
- `request_routing_rule` - A list of `request_routing_rule` blocks as defined below.
- `ssl_certificate` - A list of `ssl_certificate` blocks as defined below.
- `url_path_map` - A list of `url_path_map` blocks as defined below.
- `custom_error_configuration` - A list of `custom_error_configuration` blocks as defined below.

» * **`redirect_configuration`** - A list of `redirect_configuration` blocks as defined below.

A `authentication_certificate` block exports the following:

- `id` - The ID of the Authentication Certificate.

A `authentication_certificate` block, within the `backend_http_settings` block exports the following:

- `id` - The ID of the Authentication Certificate.

A `backend_address_pool` block exports the following:

- `id` - The ID of the Backend Address Pool.

A `backend_http_settings` block exports the following:

- `id` - The ID of the Backend HTTP Settings Configuration.
 - `probe_id` - The ID of the associated Probe.
-

A `frontend_ip_configuration` block exports the following:

- `id` - The ID of the Frontend IP Configuration.
-

A `frontend_port` block exports the following:

- `id` - The ID of the Frontend Port.
-

A `gateway_ip_configuration` block exports the following:

- `id` - The ID of the Gateway IP Configuration.
-

A `http_listener` block exports the following:

- `id` - The ID of the HTTP Listener.
 - `frontend_ip_configuration_id` - The ID of the associated Frontend Configuration.
 - `frontend_port_id` - The ID of the associated Frontend Port.
 - `ssl_certificate_id` - The ID of the associated SSL Certificate.
-

A `path_rule` block exports the following:

- `id` - The ID of the Path Rule.
 - `backend_address_pool_id` - The ID of the Backend Address Pool used in this Path Rule.
 - `backend_http_settings_id` - The ID of the Backend HTTP Settings Collection used in this Path Rule.
 - `redirect_configuration_id` - The ID of the Redirect Configuration used in this Path Rule.
-

A `probe` block exports the following:

- `id` - The ID of the Probe.
-

A `request_routing_rule` block exports the following:

- `id` - The ID of the Request Routing Rule.
 - `http_listener_id` - The ID of the associated HTTP Listener.
 - `backend_address_pool_id` - The ID of the associated Backend Address Pool.
 - `backend_http_settings_id` - The ID of the associated Backend HTTP Settings Configuration.
 - `redirect_configuration_id` - The ID of the associated Redirect Configuration.
 - `url_path_map_id` - The ID of the associated URL Path Map.
-

A `ssl_certificate` block exports the following:

- `id` - The ID of the SSL Certificate.
 - `public_cert_data` - The Public Certificate Data associated with the SSL Certificate.
-

A `url_path_map` block exports the following:

- `id` - The ID of the URL Path Map.
 - `default_backend_address_pool_id` - The ID of the Default Backend Address Pool.
 - `default_backend_http_settings_id` - The ID of the Default Backend HTTP Settings Collection.
 - `default_redirect_configuration_id` - The ID of the Default Redirect Configuration.
 - `path_rule` - A list of `path_rule` blocks as defined above.
-

A `custom_error_configuration` block exports the following:

- `id` - The ID of the Custom Error Configuration.
-

A `redirect_configuration` block exports the following:

- `id` - The ID of the Redirect Configuration.

» Import

Application Gateway's can be imported using the `resource id`, e.g.

```
terraform import azurerm_application_gateway.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_application_security_group

Manage an Application Security Group.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "tf-test"
  location = "West Europe"
}

resource "azurerm_application_security_group" "test" {
  name                = "tf-appsecuritygroup"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  tags = {
    "Hello" = "World"
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Application Security Group. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the Application Security Group.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the Application Security Group.

» Import

Application Security Groups can be imported using the `resource id`, e.g.

```
terraform import azurerm_application_security_group.securitygroup1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_connection_monitor`

Configures a Connection Monitor to monitor communication between a Virtual Machine and an endpoint using a Network Watcher.

NOTE: This resource has been deprecated in favour of the `azurerm_network_connection_monitor` resource and will be removed in the next major version of the AzureRM Provider. The new resource shares the same fields as this one, and information on migrating across can be found in this guide.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "connection-monitor-rg"
  location = "West US"
}

resource "azurerm_network_watcher" "test" {
  name                = "network-watcher"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_virtual_network" "test" {
  name                = "production-network"
  address_space       = ["10.0.0.0/16"]
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_subnet" "test" {
  name = "internal"
```

```

    resource_group_name = "${azurerm_resource_group.test.name}"
    virtual_network_name = "${azurerm_virtual_network.test.name}"
    address_prefix       = "10.0.2.0/24"
}

resource "azurerm_network_interface" "test" {
    name                = "cmtest-nic"
    location             = "${azurerm_resource_group.test.location}"
    resource_group_name = "${azurerm_resource_group.test.name}"

    ip_configuration {
        name                = "testconfiguration1"
        subnet_id           = "${azurerm_subnet.test.id}"
        private_ip_address_allocation = "Dynamic"
    }
}

resource "azurerm_virtual_machine" "test" {
    name                = "cmtest-vm"
    location            = "${azurerm_resource_group.test.location}"
    resource_group_name = "${azurerm_resource_group.test.name}"
    network_interface_ids = ["${azurerm_network_interface.test.id}"]
    vm_size             = "Standard_F2"

    storage_image_reference {
        publisher = "Canonical"
        offer     = "UbuntuServer"
        sku       = "16.04-LTS"
        version   = "latest"
    }

    storage_os_disk {
        name                = "osdisk"
        caching             = "ReadWrite"
        create_option       = "FromImage"
        managed_disk_type   = "Standard_LRS"
    }

    os_profile {
        computer_name     = "cmtest-vm"
        admin_username     = "testadmin"
        admin_password     = "Password1234!"
    }

    os_profile_linux_config {
        disable_password_authentication = false
    }
}

```

```

    }
  }

  resource "azurerm_virtual_machine_extension" "test" {
    name                        = "cmtest-vm-network-watcher"
    location                  = "${azurerm_resource_group.test.location}"
    resource_group_name       = "${azurerm_resource_group.test.name}"
    virtual_machine_name      = "${azurerm_virtual_machine.test.name}"
    publisher                 = "Microsoft.Azure.NetworkWatcher"
    type                     = "NetworkWatcherAgentLinux"
    type_handler_version      = "1.4"
    auto_upgrade_minor_version = true
  }

  resource "azurerm_connection_monitor" "test" {
    name                = "cmtest-connectionmonitor"
    location            = "${azurerm_resource_group.test.location}"
    resource_group_name = "${azurerm_resource_group.test.name}"
    network_watcher_name = "${azurerm_network_watcher.test.name}"

    source {
      virtual_machine_id = "${azurerm_virtual_machine.test.id}"
    }

    destination {
      address = "terraform.io"
      port    = 80
    }

    depends_on = ["azurerm_virtual_machine_extension.test"]
  }

```

NOTE: This Resource requires that the Network Watcher Agent Virtual Machine Extension is installed on the Virtual Machine before monitoring can be started. The extension can be installed via the `azurerm_virtual_machine_extension` resource.

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Connection Monitor. Changing this forces a new resource to be created.
- **network_watcher_name** - (Required) The name of the Network Watcher. Changing this forces a new resource to be created.

- **resource_group_name** - (Required) The name of the resource group in which to create the Connection Monitor. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **auto_start** - (Optional) Specifies whether the connection monitor will start automatically once created. Defaults to **true**. Changing this forces a new resource to be created.
- **interval_in_seconds** - (Optional) Monitoring interval in seconds. Defaults to 60.
- **source** - (Required) A **source** block as defined below.
- **destination** - (Required) A **destination** block as defined below.
- **tags** - (Optional) A mapping of tags to assign to the resource.

A **source** block contains:

- **virtual_machine_id** - (Required) The ID of the Virtual Machine to monitor connectivity from.
- **port** - (Optional) The port on the Virtual Machine to monitor connectivity from. Defaults to 0 (Dynamic Port Assignment).

A **destination** block contains:

- **virtual_machine_id** - (Optional) The ID of the Virtual Machine to monitor connectivity to.
- **address** - (Optional) IP address or domain name to monitor connectivity to.
- **port** - (Required) The port on the destination to monitor connectivity to.

NOTE: One of **virtual_machine_id** or **address** must be specified.

» Attributes Reference

The following attributes are exported:

- **id** - The Connection Monitor ID.

» Import

Connection Monitors can be imported using the **resource id**, e.g.

```
terraform import azurerm_connection_monitor.monitor1 /subscriptions/00000000-0000-0000-0000-
```

» azurerm__connection__monitor

Configures a Network Connection Monitor to monitor communication between a Virtual Machine and an endpoint using a Network Watcher.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "connection-monitor-rg"
  location = "West US"
}

resource "azurerm_network_watcher" "test" {
  name                = "network-watcher"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_virtual_network" "test" {
  name                = "production-network"
  address_space       = ["10.0.0.0/16"]
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_subnet" "test" {
  name                = "internal"
  resource_group_name = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix       = "10.0.2.0/24"
}

resource "azurerm_network_interface" "test" {
  name                = "cmtest-nic"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  ip_configuration {
    name                        = "testconfiguration1"
    subnet_id                  = "${azurerm_subnet.test.id}"
    private_ip_address_allocation = "Dynamic"
  }
}
```

```

    }
}

resource "azurerm_virtual_machine" "test" {
  name                = "cmtest-vm"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  network_interface_ids = ["${azurerm_network_interface.test.id}"]
  vm_size              = "Standard_F2"

  storage_image_reference {
    publisher = "Canonical"
    offer     = "UbuntuServer"
    sku       = "16.04-LTS"
    version   = "latest"
  }

  storage_os_disk {
    name          = "osdisk"
    caching       = "ReadWrite"
    create_option = "FromImage"
    managed_disk_type = "Standard_LRS"
  }

  os_profile {
    computer_name  = "cmtest-vm"
    admin_username = "testadmin"
    admin_password = "Password1234!"
  }

  os_profile_linux_config {
    disable_password_authentication = false
  }
}

resource "azurerm_virtual_machine_extension" "test" {
  name                = "cmtest-vm-network-watcher"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  virtual_machine_name = "${azurerm_virtual_machine.test.name}"
  publisher            = "Microsoft.Azure.NetworkWatcher"
  type                 = "NetworkWatcherAgentLinux"
  type_handler_version = "1.4"
  auto_upgrade_minor_version = true
}

```

```

resource "azurerm_network_connection_monitor" "test" {
  name                = "cmtest-connectionmonitor"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  network_watcher_name = "${azurerm_network_watcher.test.name}"

  source {
    virtual_machine_id = "${azurerm_virtual_machine.test.id}"
  }

  destination {
    address = "terraform.io"
    port    = 80
  }

  depends_on = ["azurerm_virtual_machine_extension.test"]
}

```

NOTE: This Resource requires that the Network Watcher Agent Virtual Machine Extension is installed on the Virtual Machine before monitoring can be started. The extension can be installed via the `azurerm_virtual_machine_extension` resource.

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Network Connection Monitor. Changing this forces a new resource to be created.
- **network_watcher_name** - (Required) The name of the Network Watcher. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the Connection Monitor. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **auto_start** - (Optional) Specifies whether the connection monitor will start automatically once created. Defaults to `true`. Changing this forces a new resource to be created.
- **interval_in_seconds** - (Optional) Monitoring interval in seconds. Defaults to 60.
- **source** - (Required) A `source` block as defined below.

- **destination** - (Required) A **destination** block as defined below.
- **tags** - (Optional) A mapping of tags to assign to the resource.

A **source** block contains:

- **virtual_machine_id** - (Required) The ID of the Virtual Machine to monitor connectivity from.
- **port** - (Optional) The port on the Virtual Machine to monitor connectivity from. Defaults to 0 (Dynamic Port Assignment).

A **destination** block contains:

- **virtual_machine_id** - (Optional) The ID of the Virtual Machine to monitor connectivity to.
- **address** - (Optional) IP address or domain name to monitor connectivity to.
- **port** - (Required) The port on the destination to monitor connectivity to.

NOTE: One of **virtual_machine_id** or **address** must be specified.

» Attributes Reference

The following attributes are exported:

- **id** - The Connection Monitor ID.

» Import

Connection Monitors can be imported using the **resource id**, e.g.

```
terraform import azurerm_network_connection_monitor.monitor1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_ddos_protection_plan

Manages an Azure DDoS Protection Plan.

NOTE Azure only allow **one** DDoS Protection Plan per region.

NOTE: This resource has been deprecated in favour of the **azurerm_network_ddos_protection_plan** resource and will be removed in the next major version of the AzureRM Provider. The new resource shares the same fields as this one, and information on migrating across can be found in this guide.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "example-resources"
  location  = "West Europe"
}

resource "azurerm_ddos_protection_plan" "test" {
  name                = "example-protection-plan"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the DDoS Protection Plan. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the resource. Changing this forces a new resource to be created.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- **id** - The Resource ID of the DDoS Protection Plan
- **virtual_network_ids** - The Resource ID list of the Virtual Networks associated with DDoS Protection Plan.

» Import

Azure DDoS Protection Plan can be imported using the **resource id**, e.g.

```
terraform import azurerm_ddos_protection_plan.test /subscriptions/00000000-0000-0000-0000-000000000000/resourceGroups/example-rg/providers/Microsoft.Network/ddosProtectionPlans/example-protection-plan
```

» **azurerm_network_ddos_protection_plan**

Manages an AzureNetwork DDoS Protection Plan.

NOTE Azure only allow one DDoS Protection Plan per region.

» **Example Usage**

```
resource "azurerm_resource_group" "test" {
  name      = "example-resources"
  location  = "West Europe"
}

resource "azurerm_network_ddos_protection_plan" "test" {
  name                  = "example-protection-plan"
  location              = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}
```

» **Argument Reference**

The following arguments are supported:

- **name** - (Required) Specifies the name of the Network DDoS Protection Plan. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the resource. Changing this forces a new resource to be created.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» **Attributes Reference**

The following attributes are exported:

- **id** - The Resource ID of the DDoS Protection Plan
- **virtual_network_ids** - The Resource ID list of the Virtual Networks associated with DDoS Protection Plan.

» Import

Azure DDoS Protection Plan can be imported using the `resource id`, e.g.

```
terraform import azurerm_network_ddos_protection_plan.test /subscriptions/00000000-0000-0000
```

» azurerm_express_route_circuit

Manages an ExpressRoute circuit.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "exprtTest"
  location = "West US"
}

resource "azurerm_express_route_circuit" "test" {
  name                        = "expressRoute1"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  location                   = "${azurerm_resource_group.test.location}"
  service_provider_name      = "Equinix"
  peering_location            = "Silicon Valley"
  bandwidth_in_mbps          = 50

  sku {
    tier     = "Standard"
    family  = "MeteredData"
  }

  tags = {
    environment = "Production"
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the ExpressRoute circuit. Changing this forces a new resource to be created.

- **resource_group_name** - (Required) The name of the resource group in which to create the ExpressRoute circuit. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **service_provider_name** - (Required) The name of the ExpressRoute Service Provider.
- **peering_location** - (Required) The name of the peering location and **not** the Azure resource location.
- **bandwidth_in_mbps** - (Required) The bandwidth in Mbps of the circuit being created.

NOTE: Once you increase your bandwidth, you will not be able to decrease it to it's previous value.

- **sku** - (Required) A **sku** block for the ExpressRoute circuit as documented below.
- **allow_classic_operations** - (Optional) Allow the circuit to interact with classic (RDFE) resources. The default value is **false**.
- **tags** - (Optional) A mapping of tags to assign to the resource.

sku supports the following:

- **tier** - (Required) The service tier. Possible values are **Standard** or **Premium**.
- **family** - (Required) The billing mode for bandwidth. Possible values are **MeteredData** or **UnlimitedData**.

NOTE: You can migrate from **MeteredData** to **UnlimitedData**, but not the other way around.

» Attributes Reference

The following attributes are exported:

- **id** - The Resource ID of the ExpressRoute circuit.
- **service_provider_provisioning_state** - The ExpressRoute circuit provisioning state from your chosen service provider. Possible values are "NotProvisioned", "Provisioning", "Provisioned", and "Deprovisioning".
- **service_key** - The string needed by the service provider to provision the ExpressRoute circuit.

» Import

ExpressRoute circuits can be imported using the `resource id`, e.g.

```
terraform import azurerm_express_route_circuit.myExpressRoute /subscriptions/00000000-0000-0
```

» `azurerm_express_route_circuit_authorization`

Manages an ExpressRoute Circuit Authorization.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "exprtTest"
  location  = "West US"
}

resource "azurerm_express_route_circuit" "test" {
  name                        = "expressRoute1"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  location                   = "${azurerm_resource_group.test.location}"
  service_provider_name      = "Equinix"
  peering_location           = "Silicon Valley"
  bandwidth_in_mbps          = 50

  sku {
    tier      = "Standard"
    family   = "MeteredData"
  }

  allow_classic_operations = false

  tags = {
    environment = "Production"
  }
}

resource "azurerm_express_route_circuit_authorization" "test" {
  name                        = "exampleERCAuth"
  express_route_circuit_name = "${azurerm_express_route_circuit.test.name}"
  resource_group_name        = "${azurerm_resource_group.test.name}"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the ExpressRoute circuit. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the ExpressRoute circuit. Changing this forces a new resource to be created.
- **express_route_circuit_name** - (Required) The name of the Express Route Circuit in which to create the Authorization.

» Attributes Reference

The following attributes are exported:

- **id** - The Resource ID of the ExpressRoute Circuit Authorization.
- **authorization_key** - The Authorization Key.
- **authorization_use_status** - The authorization use status.

» Import

ExpressRoute Circuit Authorizations can be imported using the **resource id**, e.g.

```
terraform import azurerm_express_route_circuit_authorization.auth1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_express_route_circuit_peering

Manages an ExpressRoute Circuit Peering.

» Example Usage (Creating a Microsoft Peering)

```
resource "azurerm_resource_group" "test" {
  name     = "exprtTest"
  location = "West US"
}

resource "azurerm_express_route_circuit" "test" {
  name                = "expressRoute1"
  resource_group_name = "${azurerm_resource_group.test.name}"
}
```

```

location          = "${azurerm_resource_group.test.location}"
service_provider_name = "Equinix"
peering_location   = "Silicon Valley"
bandwidth_in_mbps   = 50

sku {
  tier      = "Standard"
  family    = "MeteredData"
}

allow_classic_operations = false

tags = {
  environment = "Production"
}
}

resource "azurerm_express_route_circuit_peering" "test" {
  peering_type      = "MicrosoftPeering"
  express_route_circuit_name = "${azurerm_express_route_circuit.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  peer_asn           = 100
  primary_peer_address_prefix = "123.0.0.0/30"
  secondary_peer_address_prefix = "123.0.0.4/30"
  vlan_id            = 300

  microsoft_peering_config {
    advertised_public_prefixes = ["123.1.0.0/24"]
  }
}

```

» Argument Reference

The following arguments are supported:

- **peering_type** - (Required) The type of the ExpressRoute Circuit Peering. Acceptable values include `AzurePrivatePeering`, `AzurePublicPeering` and `MicrosoftPeering`. Changing this forces a new resource to be created.

NOTE: only one Peering of each Type can be created. Attempting to create multiple peerings of the same type will overwrite the original peering.

- **express_route_circuit_name** - (Required) The name of the Express-Route Circuit in which to create the Peering.
- **resource_group_name** - (Required) The name of the resource group in

which to create the Express Route Circuit Peering. Changing this forces a new resource to be created.

- **primary_peer_address_prefix** - (Optional) A /30 subnet for the primary link.
- **secondary_peer_address_prefix** - (Optional) A /30 subnet for the secondary link.
- **vlan_id** - (Optional) A valid VLAN ID to establish this peering on.
- **shared_key** - (Optional) The shared key. Can be a maximum of 25 characters.
- **peer_asn** - (Optional) The Either a 16-bit or a 32-bit ASN. Can either be public or private..
- **microsoft_peering_config** - (Optional) A **microsoft_peering_config** block as defined below. Required when **peering_type** is set to **MicrosoftPeering**.

A **microsoft_peering_config** block contains:

- **advertised_public_prefixes** - (Required) A list of Advertised Public Prefixes

» Attributes Reference

The following attributes are exported:

- **id** - The Resource ID of the ExpressRoute Circuit Peering.
- **azure_asn** - The ASN used by Azure.
- **primary_azure_port** - The Primary Port used by Azure for this Peering.
- **secondary_azure_port** - The Secondary Port used by Azure for this Peering.

» Import

ExpressRoute Circuit Peerings can be imported using the **resource id**, e.g.

```
terraform import azurerm_express_route_circuit_peering.peering1 /subscriptions/00000000-0000-
```


» azurerm_firewall

Manages an Azure Firewall.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "example-resources"
  location  = "North Europe"
}

resource "azurerm_virtual_network" "test" {
  name            = "testvnet"
  address_space   = ["10.0.0.0/16"]
  location        = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_subnet" "test" {
  name                 = "AzureFirewallSubnet"
  resource_group_name = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix       = "10.0.1.0/24"
}

resource "azurerm_public_ip" "test" {
  name                = "testpip"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  allocation_method   = "Static"
  sku                 = "Standard"
}

resource "azurerm_firewall" "test" {
  name                = "testfirewall"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  ip_configuration {
    name                = "configuration"
    subnet_id          = "${azurerm_subnet.test.id}"
    public_ip_address_id = "${azurerm_public_ip.test.id}"
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Firewall. Changing this forces a new resource to be created.
 - **resource_group_name** - (Required) The name of the resource group in which to create the resource. Changing this forces a new resource to be created.
 - **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
 - **ip_configuration** - (Required) A `ip_configuration` block as documented below.
 - **tags** - (Optional) A mapping of tags to assign to the resource.
-

A `ip_configuration` block supports the following:

- **name** - (Required) Specifies the name of the IP Configuration.
- **subnet_id** - (Required) Reference to the subnet associated with the IP Configuration.

NOTE The Subnet used for the Firewall must have the name `AzureFirewallSubnet` and the subnet mask must be at least `/26`.

- **public_ip_address_id** - (Required) The Resource ID of the Public IP Address associated with the firewall.

NOTE The Public IP must have a `Static` allocation and `Standard` sku.

» Attributes Reference

The following attributes are exported:

- **id** - The Resource ID of the Azure Firewall.
 - **ip_configuration** - A `ip_configuration` block as defined below.
-

A `ip_configuration` block exports the following:

- **private_ip_address** - The private IP address of the Azure Firewall.

» Import

Azure Firewalls can be imported using the `resource id`, e.g.

```
terraform import azurerm_firewall.test /subscriptions/00000000-0000-0000-0000-000000000000/1
```

» `azurerm_firewall_application_rule_collection`

Manages an Application Rule Collection within an Azure Firewall.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "example-resources"
  location  = "North Europe"
}

resource "azurerm_virtual_network" "test" {
  name                = "testvnet"
  address_space       = ["10.0.0.0/16"]
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_subnet" "test" {
  name                 = "AzureFirewallSubnet"
  resource_group_name = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix       = "10.0.1.0/24"
}

resource "azurerm_public_ip" "test" {
  name                = "testpip"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  allocation_method   = "Static"
  sku                 = "Standard"
}

resource "azurerm_firewall" "test" {
  name                = "testfirewall"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}
```

```

    ip_configuration {
      name                = "configuration"
      subnet_id           = "${azurerm_subnet.test.id}"
      public_ip_address_id = "${azurerm_public_ip.test.id}"
    }
  }

  resource "azurerm_firewall_application_rule_collection" "test" {
    name                = "testcollection"
    azure_firewall_name = "${azurerm_firewall.test.name}"
    resource_group_name = "${azurerm_resource_group.test.name}"
    priority            = 100
    action              = "Allow"

    rule {
      name = "testrule"

      source_addresses = [
        "10.0.0.0/16",
      ]

      target_fqdns = [
        "*.google.com",
      ]

      protocol {
        port = "443"
        type = "Https"
      }
    }
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Application Rule Collection which must be unique within the Firewall. Changing this forces a new resource to be created.
- **azure_firewall_name** - (Required) Specifies the name of the Firewall in which the Application Rule Collection should be created. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) Specifies the name of the Resource Group in which the Firewall exists. Changing this forces a new resource

to be created.

- **priority** - (Required) Specifies the priority of the rule collection. Possible values are between 100 - 65000.
 - **action** - (Required) Specifies the action the rule will apply to matching traffic. Possible values are **Allow** and **Deny**.
 - **rule** - (Required) One or more **rule** blocks as defined below.
-

A **rule** block supports the following:

- **name** - (Required) Specifies the name of the rule.
 - **description** - (Optional) Specifies a description for the rule.
 - **source_addresses** - (Required) A list of source IP addresses and/or IP ranges.
 - **fqdn_tags** - (Optional) A list of FQDN tags. Possible values are **AppServiceEnvironment**, **AzureBackup**, **MicrosoftActiveProtectionService**, **WindowsDiagnostics** and **WindowsUpdate**
 - **target_fqdns** - (Optional) A list of FQDNs.
 - **protocol** - (Optional) One or more **protocol** blocks as defined below.
-

A **protocol** block supports the following:

- **port** - (Optional) Specify a port for the connection.
- **type** - (Required) Specifies the type of connection. Possible values are **Http** or **Https**.

» Import

Azure Firewall Application Rule Collections can be imported using the **resource id**, e.g.

```
terraform import azurerm_firewall_application_rule_collection.test /subscriptions/00000000-0000-0000-0000-000000000000/resourceGroups/rg-test/providers/Microsoft.Network/firewalls/firewall-test/ruleCollections/rule-test
```

» **azurerm_firewall_network_rule_collection**

Manages a Network Rule Collection within an Azure Firewall.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "example-resources"
  location  = "North Europe"
}

resource "azurerm_virtual_network" "test" {
  name            = "testvnet"
  address_space   = ["10.0.0.0/16"]
  location        = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_subnet" "test" {
  name                 = "AzureFirewallSubnet"
  resource_group_name  = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix       = "10.0.1.0/24"
}

resource "azurerm_public_ip" "test" {
  name                = "testpip"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name  = "${azurerm_resource_group.test.name}"
  allocation_method    = "Static"
  sku                 = "Standard"
}

resource "azurerm_firewall" "test" {
  name                = "testfirewall"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  ip_configuration {
    name                = "configuration"
    subnet_id           = "${azurerm_subnet.test.id}"
    public_ip_address_id = "${azurerm_public_ip.test.id}"
  }
}

resource "azurerm_firewall_network_rule_collection" "test" {
  name                = "testcollection"
  azure_firewall_name = "${azurerm_firewall.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}
```

```

priority          = 100
action            = "Allow"

rule {
  name = "testrule"

  source_addresses = [
    "10.0.0.0/16",
  ]

  destination_ports = [
    "53",
  ]

  destination_addresses = [
    "8.8.8.8",
    "8.8.4.4",
  ]

  protocols = [
    "TCP",
    "UDP",
  ]
}
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Network Rule Collection which must be unique within the Firewall. Changing this forces a new resource to be created.
- **azure_firewall_name** - (Required) Specifies the name of the Firewall in which the Network Rule Collection should be created. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) Specifies the name of the Resource Group in which the Firewall exists. Changing this forces a new resource to be created.
- **priority** - (Required) Specifies the priority of the rule collection. Possible values are between 100 - 65000.
- **action** - (Required) Specifies the action the rule will apply to matching traffic. Possible values are **Allow** and **Deny**.

- **rule** - (Required) One or more **rule** blocks as defined below.

A **rule** block supports the following:

- **name** - (Required) Specifies the name of the rule.
- **description** - (Optional) Specifies a description for the rule.
- **source_addresses** - (Required) A list of source IP addresses and/or IP ranges.
- **destination_addresses** - (Required) A list of destination IP addresses and/or IP ranges.
- **destination_ports** - (Required) A list of destination ports.
- **protocols** - (Required) A list of protocols. Possible values are **Any**, **ICMP**, **TCP** and **UDP**.

» Import

Azure Firewall Network Rule Collections can be imported using the **resource** id, e.g.

```
terraform import azurerm_firewall_network_rule_collection.test /subscriptions/00000000-0000-
```

» azurerm__local__network__gateway

Manages a local network gateway connection over which specific connections can be configured.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "localNetworkGWTest"
  location  = "West US"
}

resource "azurerm_local_network_gateway" "home" {
  name                        = "backHome"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  location                   = "${azurerm_resource_group.test.location}"
  gateway_address            = "12.13.14.15"
  address_space              = ["10.0.0.0/16"]
}
```


» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the local network gateway. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the local network gateway.
- **location** - (Required) The location/region where the local network gateway is created. Changing this forces a new resource to be created.
- **gateway_address** - (Required) The IP address of the gateway to which to connect.
- **address_space** - (Required) The list of string CIDRs representing the address spaces the gateway exposes.
- **bgp_settings** - (Optional) A **bgp_settings** block as defined below containing the Local Network Gateway's BGP speaker settings.
- **tags** - (Optional) A mapping of tags to assign to the resource.

bgp_settings supports the following:

- **asn** - (Required) The BGP speaker's ASN.
- **bgp_peering_address** - (Required) The BGP peering address and BGP identifier of this BGP speaker.
- **peer_weight** - (Optional) The weight added to routes learned from this BGP speaker.

» Attributes Reference

The following attributes are exported:

- **id** - The local network gateway unique ID within Azure.

» Import

Local Network Gateways can be imported using the **resource id**, e.g.

```
terraform import azurerm_local_network_gateway.lng1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_network_interface

Manages a Network Interface located in a Virtual Network, usually attached to a Virtual Machine.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "acceptanceTestResourceGroup1"
  location = "West US"
}

resource "azurerm_virtual_network" "test" {
  name            = "acceptanceTestVirtualNetwork1"
  address_space   = ["10.0.0.0/16"]
  location        = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_subnet" "test" {
  name                 = "testsubnet"
  resource_group_name = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix       = "10.0.2.0/24"
}

resource "azurerm_network_interface" "test" {
  name                = "acceptanceTestNetworkInterface1"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  ip_configuration {
    name                       = "testconfiguration1"
    subnet_id                 = "${azurerm_subnet.test.id}"
    private_ip_address_allocation = "Dynamic"
  }

  tags = {
    environment = "staging"
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the network interface. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the network interface. Changing this forces a new resource to be created.
- **location** - (Required) The location/region where the network interface is created. Changing this forces a new resource to be created.
- **network_security_group_id** - (Optional) The ID of the Network Security Group to associate with the network interface.
- **internal_dns_name_label** - (Optional) Relative DNS name for this NIC used for internal communications between VMs in the same VNet
- **enable_ip_forwarding** - (Optional) Enables IP Forwarding on the NIC. Defaults to `false`.
- **enable_accelerated_networking** - (Optional) Enables Azure Accelerated Networking using SR-IOV. Only certain VM instance sizes are supported. Refer to [Create a Virtual Machine with Accelerated Networking](#). Defaults to `false`.

NOTE: when using Accelerated Networking in an Availability Set - the Availability Set must be deployed on an Accelerated Networking enabled cluster.

- **dns_servers** - (Optional) List of DNS servers IP addresses to use for this NIC, overrides the VNet-level server list
- **ip_configuration** - (Required) One or more **ip_configuration** associated with this NIC as documented below.
- **tags** - (Optional) A mapping of tags to assign to the resource.

The **ip_configuration** block supports:

- **name** - (Required) User-defined name of the IP.
- **subnet_id** - (Optional) Reference to a subnet in which this NIC has been created. Required when **private_ip_address_version** is IPv4.
- **private_ip_address** - (Optional) Static IP Address.
- **private_ip_address_allocation** - (Required) Defines how a private IP address is assigned. Options are Static or Dynamic.
- **private_ip_address_version** - (Optional) The IP Version to use. Possible values are IPv4 or IPv6. Defaults to IPv4.

- `public_ip_address_id` - (Optional) Reference to a Public IP Address to associate with this NIC
- `application_gateway_backend_address_pools_ids` - (Optional / **Deprecated**) List of Application Gateway Backend Address Pool IDs references to which this NIC belongs

NOTE: At this time Network Interface <Application Gateway Backend Address Pool associations need to be configured both using this field (which is now Deprecated) and using the `azurerm_network_interface_application_gateway_backend_address_pool_association` resource. This field is deprecated and will be removed in favour of that resource in the next major version (2.0) of the AzureRM Provider.

- `load_balancer_backend_address_pools_ids` - (Optional / **Deprecated**) List of Load Balancer Backend Address Pool IDs references to which this NIC belongs

NOTE: At this time Network Interface <Load Balancer Backend Address Pool associations need to be configured both using this field (which is now Deprecated) and using the `azurerm_network_interface_backend_address_pool_association` resource. This field is deprecated and will be removed in favour of that resource in the next major version (2.0) of the AzureRM Provider.

- `load_balancer_inbound_nat_rules_ids` - (Optional / **Deprecated**) List of Load Balancer Inbound Nat Rules IDs involving this NIC

NOTE: At this time Network Interface <Load Balancer Inbound NAT Rule associations need to be configured both using this field (which is now Deprecated) and using the `azurerm_network_interface_nat_rule_association` resource. This field is deprecated and will be removed in favour of that resource in the next major version (2.0) of the AzureRM Provider.

- `application_security_group_ids` - (Optional / **Deprecated**) List of Application Security Group IDs which should be attached to this NIC

NOTE: At this time Network Interface <Application Security Group associations need to be configured both using this field (which is now Deprecated) and using the `azurerm_network_interface_application_security_group_association` resource. This field is deprecated and will be removed in favour of that resource in the next major version (2.0) of the AzureRM Provider.

- `primary` - (Optional) Is this the Primary Network Interface? If set to `true` this should be the first `ip_configuration` in the array.

» Attributes Reference

The following attributes are exported:

- `id` - The Virtual Network Interface ID.

- `mac_address` - The media access control (MAC) address of the network interface.
- `private_ip_address` - The first private IP address of the network interface.
- `private_ip_addresses` - The private IP addresses of the network interface.
- `virtual_machine_id` - Reference to a VM with which this NIC has been associated.
- `applied_dns_servers` - If the VM that uses this NIC is part of an Availability Set, then this list will have the union of all DNS servers from all NICs that are part of the Availability Set

» Import

Network Interfaces can be imported using the `resource id`, e.g.

```
terraform import azurerm_network_interface.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_network_interface_application_gateway_backend_address_pool`

Manages the association between a Network Interface and a Application Gateway's Backend Address Pool.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "example-resources"
  location = "West Europe"
}

resource "azurerm_virtual_network" "test" {
  name            = "example-network"
  address_space   = ["10.0.0.0/16"]
  location        = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_subnet" "frontend" {
  name                 = "frontend"
  resource_group_name = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix       = "10.254.0.0/24"
}
```

```

resource "azurerm_subnet" "backend" {
  name                = "backend"
  resource_group_name = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix      = "10.254.2.0/24"
}

resource "azurerm_public_ip" "test" {
  name                = "example-pip"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  allocation_method   = "Dynamic"
}

# since these variables are re-used - a locals block makes this more maintainable
locals {
  backend_address_pool_name = "${azurerm_virtual_network.test.name}-beap"
  frontend_port_name        = "${azurerm_virtual_network.test.name}-feport"
  frontend_ip_configuration_name = "${azurerm_virtual_network.test.name}-feip"
  http_setting_name         = "${azurerm_virtual_network.test.name}-be-htst"
  listener_name             = "${azurerm_virtual_network.test.name}-httplstn"
  request_routing_rule_name = "${azurerm_virtual_network.test.name}-rqrt"
}

resource "azurerm_application_gateway" "network" {
  name                = "example-appgateway"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location            = "${azurerm_resource_group.test.location}"

  sku {
    name     = "Standard_Small"
    tier      = "Standard"
    capacity = 2
  }

  gateway_ip_configuration {
    name        = "my-gateway-ip-configuration"
    subnet_id   = "${azurerm_subnet.frontend.id}"
  }

  frontend_port {
    name = "${local.frontend_port_name}"
    port = 80
  }
}

```

```

frontend_ip_configuration {
  name = "${local.frontend_ip_configuration_name}"
  public_ip_address_id = "${azurerm_public_ip.test.id}"
}

backend_address_pool {
  name = "${local.backend_address_pool_name}"
}

backend_http_settings {
  name = "${local.http_setting_name}"
  cookie_based_affinity = "Disabled"
  port = 80
  protocol = "Http"
  request_timeout = 1
}

http_listener {
  name = "${local.listener_name}"
  frontend_ip_configuration_name = "${local.frontend_ip_configuration_name}"
  frontend_port_name = "${local.frontend_port_name}"
  protocol = "Http"
}

request_routing_rule {
  name = "${local.request_routing_rule_name}"
  rule_type = "Basic"
  http_listener_name = "${local.listener_name}"
  backend_address_pool_name = "${local.backend_address_pool_name}"
  backend_http_settings_name = "${local.http_setting_name}"
}

}

resource "azurerm_network_interface" "test" {
  name = "example-nic"
  location = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  ip_configuration {
    name = "testconfiguration1"
    subnet_id = "${azurerm_subnet.frontend.id}"
    private_ip_address_allocation = "Dynamic"
  }
}

resource "azurerm_network_interface_application_gateway_backend_address_pool_association" "t

```

```

network_interface_id    = "${azurerm_network_interface.test.id}"
ip_configuration_name   = "testconfiguration1"
backend_address_pool_id = "${azurerm_application_gateway.test.backend_address_pool.0.id}"
}

```

» Argument Reference

The following arguments are supported:

- **network_interface_id** - (Required) The ID of the Network Interface. Changing this forces a new resource to be created.
- **ip_configuration_name** - (Required) The Name of the IP Configuration within the Network Interface which should be connected to the Backend Address Pool. Changing this forces a new resource to be created.
- **backend_address_pool_id** - (Required) The ID of the Application Gateway's Backend Address Pool which this Network Interface which should be connected to. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The (Terraform specific) ID of the Association between the Network Interface and the Application Gateway Backend Address Pool.

» Import

Associations between Network Interfaces and Application Gateway Backend Address Pools can be imported using the **resource id**, e.g.

```
terraform import azurerm_network_interface_application_gateway_backend_address_pool_associat
```

NOTE: This ID is specific to Terraform - and is of the format {networkInterfaceId}/ipConfigurations/{ip

» azurerm__network__interface__application__security__group__association

Manages the association between a Network Interface and a Application Security Group.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "example-resources"
  location  = "West Europe"
}

resource "azurerm_virtual_network" "test" {
  name            = "example-network"
  address_space   = ["10.0.0.0/16"]
  location        = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_subnet" "test" {
  name                 = "internal"
  resource_group_name  = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix       = "10.0.1.0/24"
}

resource "azurerm_application_security_group" "test" {
  name            = "example-asg"
  location        = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_network_interface" "test" {
  name            = "example-nic"
  location        = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  ip_configuration {
    name                       = "testconfiguration1"
    subnet_id                 = "${azurerm_subnet.test.id}"
    private_ip_address_allocation = "Dynamic"
    application_security_group_ids = [ "${azurerm_application_security_group.test.id}" ]
  }
}

resource "azurerm_network_interface_application_security_group_association" "test" {
  network_interface_id = "${azurerm_network_interface.test.id}"
  ip_configuration_name = "testconfiguration1"
  application_security_group_id = "${azurerm_application_security_group.test.id}"
}
```

» Argument Reference

The following arguments are supported:

- **network_interface_id** - (Required) The ID of the Network Interface. Changing this forces a new resource to be created.
- **ip_configuration_name** - (Required) The Name of the IP Configuration within the Network Interface which should be connected to the Application Security Group. Changing this forces a new resource to be created.
- **application_security_group_id** - (Required) The ID of the Application Security Group which this Network Interface which should be connected to. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The (Terraform specific) ID of the Association between the Network Interface and the Application Security Group.

» Import

Associations between Network Interfaces and Application Security Groups can be imported using the **resource id**, e.g.

```
terraform import azurerm_network_interface_application_security_group_association.association
```

NOTE: This ID is specific to Terraform - and is of the format {networkInterfaceId}/ipConfigurations/{ip}

» **azurerm_network_interface_backend_address_pool_association**

Manages the association between a Network Interface and a Load Balancer's Backend Address Pool.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "example-resources"
  location = "West Europe"
}
```

```
resource "azurerm_virtual_network" "test" {
```

```

    name                = "example-network"
    address_space        = ["10.0.0.0/16"]
    location              = "${azurerm_resource_group.test.location}"
    resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_subnet" "test" {
    name                = "internal"
    resource_group_name = "${azurerm_resource_group.test.name}"
    virtual_network_name = "${azurerm_virtual_network.test.name}"
    address_prefix      = "10.0.2.0/24"
}

resource "azurerm_public_ip" "test" {
    name                = "example-pip"
    location            = "${azurerm_resource_group.test.location}"
    resource_group_name = "${azurerm_resource_group.test.name}"
    allocation_method   = "Static"
}

resource "azurerm_lb" "test" {
    name                = "example-lb"
    location            = "${azurerm_resource_group.test.location}"
    resource_group_name = "${azurerm_resource_group.test.name}"

    frontend_ip_configuration {
        name                = "primary"
        public_ip_address_id = "${azurerm_public_ip.test.id}"
    }
}

resource "azurerm_lb_backend_address_pool" "test" {
    resource_group_name = "${azurerm_resource_group.test.name}"
    loadbalancer_id     = "${azurerm_lb.test.id}"
    name                = "acctestpool"
}

resource "azurerm_network_interface" "test" {
    name                = "example-nic"
    location            = "${azurerm_resource_group.test.location}"
    resource_group_name = "${azurerm_resource_group.test.name}"

    ip_configuration {
        name                = "testconfiguration1"
        subnet_id           = "${azurerm_subnet.test.id}"
        private_ip_address_allocation = "Dynamic"
    }
}

```

```

    }
  }

  resource "azurerm_network_interface_backend_address_pool_association" "test" {
    network_interface_id      = "${azurerm_network_interface.test.id}"
    ip_configuration_name    = "testconfiguration1"
    backend_address_pool_id = "${azurerm_lb_backend_address_pool.test.id}"
  }

```

» Argument Reference

The following arguments are supported:

- **network_interface_id** - (Required) The ID of the Network Interface. Changing this forces a new resource to be created.
- **ip_configuration_name** - (Required) The Name of the IP Configuration within the Network Interface which should be connected to the Backend Address Pool. Changing this forces a new resource to be created.
- **backend_address_pool_id** - (Required) The ID of the Load Balancer Backend Address Pool which this Network Interface which should be connected to. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The (Terraform specific) ID of the Association between the Network Interface and the Load Balancers Backend Address Pool.

» Import

Associations between Network Interfaces and Load Balancer Backend Address Pools can be imported using the **resource id**, e.g.

```
terraform import azurerm_network_interface_backend_address_pool_association.association1 /subscriptions/12345678-9012-3456-7890-123456789012/resourceGroups/rg1/providers/Microsoft.Network/loadBalancerBackendAddressPools/lbBackendAddressPool1/ipConfigurations/ipConfiguration1/networkInterfaces/networkInterface1
```

NOTE: This ID is specific to Terraform - and is of the format {networkInterfaceId}/ipConfigurations/{ipConfigurationName}

» azurerm_network_interface_nat_rule_association

Manages the association between a Network Interface and a Load Balancer's NAT Rule.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "example-resources"
  location = "West Europe"
}

resource "azurerm_virtual_network" "test" {
  name            = "example-network"
  address_space   = ["10.0.0.0/16"]
  location        = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_subnet" "test" {
  name                 = "internal"
  resource_group_name = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix       = "10.0.2.0/24"
}

resource "azurerm_public_ip" "test" {
  name            = "example-pip"
  location        = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  allocation_method = "Static"
}

resource "azurerm_lb" "test" {
  name            = "example-lb"
  location        = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  frontend_ip_configuration {
    name                 = "primary"
    public_ip_address_id = "${azurerm_public_ip.test.id}"
  }
}

resource "azurerm_lb_nat_rule" "test" {
  resource_group_name = "${azurerm_resource_group.test.name}"
  loadbalancer_id     = "${azurerm_lb.test.id}"
  name                = "RDPAccess"
  protocol             = "Tcp"
  frontend_port        = 3389
}
```

```

        backend_port          = 3389
        frontend_ip_configuration_name = "primary"
    }

    resource "azurerm_network_interface" "test" {
        name                = "example-nic"
        location            = "${azurerm_resource_group.test.location}"
        resource_group_name = "${azurerm_resource_group.test.name}"

        ip_configuration {
            name                = "testconfiguration1"
            subnet_id          = "${azurerm_subnet.test.id}"
            private_ip_address_allocation = "Dynamic"
        }
    }

    resource "azurerm_network_interface_nat_rule_association" "test" {
        network_interface_id = "${azurerm_network_interface.test.id}"
        ip_configuration_name = "testconfiguration1"
        nat_rule_id          = "${azurerm_lb_nat_rule.test.id}"
    }

```

» Argument Reference

The following arguments are supported:

- **network_interface_id** - (Required) The ID of the Network Interface. Changing this forces a new resource to be created.
- **ip_configuration_name** - (Required) The Name of the IP Configuration within the Network Interface which should be connected to the NAT Rule. Changing this forces a new resource to be created.
- **nat_rule_id** - (Required) The ID of the Load Balancer NAT Rule which this Network Interface which should be connected to. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The (Terraform specific) ID of the Association between the Network Interface and the Load Balancers NAT Rule.

» Import

Associations between Network Interfaces and Load Balancer NAT Rule can be imported using the `resource id`, e.g.

```
terraform import azurerm_network_interface_nat_rule_association.association1 /subscriptions/
```

NOTE: This ID is specific to Terraform - and is of the format `{networkInterfaceId}/ipConfigurations/{ip`

» azurerm__network__security__group

Manages a network security group that contains a list of network security rules. Network security groups enable inbound or outbound traffic to be enabled or denied.

NOTE on Network Security Groups and Network Security Rules: Terraform currently provides both a standalone Network Security Rule resource, and allows for Network Security Rules to be defined in-line within the Network Security Group resource. At this time you cannot use a Network Security Group with in-line Network Security Rules in conjunction with any Network Security Rule resources. Doing so will cause a conflict of rule settings and will overwrite rules.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "acceptanceTestResourceGroup1"
  location = "West US"
}

resource "azurerm_network_security_group" "test" {
  name                = "acceptanceTestSecurityGroup1"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  security_rule {
    name                = "test123"
    priority            = 100
    direction           = "Inbound"
    access              = "Allow"
    protocol            = "Tcp"
    source_port_range   = "*"
    destination_port_range = "*"
    source_address_prefix = "*"
    destination_address_prefix = "*"
  }
}
```

```

    }

    tags = {
        environment = "Production"
    }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the network security group. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the network security group. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **security_rule** - (Optional) One or more **security_rule** blocks as defined below.
- **tags** - (Optional) A mapping of tags to assign to the resource.

The **security_rule** block supports:

- **name** - (Required) The name of the security rule.
- **description** - (Optional) A description for this rule. Restricted to 140 characters.
- **protocol** - (Required) Network protocol this rule applies to. Can be **Tcp**, **Udp** or ***** to match both.
- **source_port_range** - (Optional) Source Port or Range. Integer or range between 0 and 65535 or ***** to match any. This is required if **source_port_ranges** is not specified.
- **source_port_ranges** - (Optional) List of source ports or port ranges. This is required if **source_port_range** is not specified.
- **destination_port_range** - (Optional) Destination Port or Range. Integer or range between 0 and 65535 or ***** to match any. This is required if **destination_port_ranges** is not specified.
- **destination_port_ranges** - (Optional) List of destination ports or port ranges. This is required if **destination_port_range** is not specified.
- **source_address_prefix** - (Optional) CIDR or source IP range or ***** to match any IP. Tags such as 'VirtualNetwork', 'AzureLoadBalancer' and

‘Internet’ can also be used. This is required if `source_address_prefixes` is not specified.

- `source_address_prefixes` - (Optional) List of source address prefixes. Tags may not be used. This is required if `source_address_prefix` is not specified.
- `source_application_security_group_ids` - (Optional) A List of source Application Security Group ID’s
- `destination_address_prefix` - (Optional) CIDR or destination IP range or * to match any IP. Tags such as ‘VirtualNetwork’, ‘AzureLoadBalancer’ and ‘Internet’ can also be used. This is required if `destination_address_prefixes` is not specified.
- `destination_address_prefixes` - (Optional) List of destination address prefixes. Tags may not be used. This is required if `destination_address_prefix` is not specified.
- `destination_application_security_group_ids` - (Optional) A List of destination Application Security Group ID’s
- `access` - (Required) Specifies whether network traffic is allowed or denied. Possible values are `Allow` and `Deny`.
- `priority` - (Required) Specifies the priority of the rule. The value can be between 100 and 4096. The priority number must be unique for each rule in the collection. The lower the priority number, the higher the priority of the rule.
- `direction` - (Required) The direction specifies if rule will be evaluated on incoming or outgoing traffic. Possible values are `Inbound` and `Outbound`.

» Attributes Reference

The following attributes are exported:

- `id` - The Network Security Group ID.

» Import

Network Security Groups can be imported using the `resource id`, e.g.

```
terraform import azurerm_network_security_group.group1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_network_security_rule

Manages a Network Security Rule.

NOTE on Network Security Groups and Network Security Rules:

Terraform currently provides both a standalone Network Security Rule resource, and allows for Network Security Rules to be defined in-line within the Network Security Group resource. At this time you cannot use a Network Security Group with in-line Network Security Rules in conjunction with any Network Security Rule resources. Doing so will cause a conflict of rule settings and will overwrite rules.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "acceptanceTestResourceGroup1"
  location  = "West US"
}

resource "azurerm_network_security_group" "test" {
  name                = "acceptanceTestSecurityGroup1"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_network_security_rule" "test" {
  name                = "test123"
  priority            = 100
  direction           = "Outbound"
  access              = "Allow"
  protocol            = "Tcp"
  source_port_range   = "*"
  destination_port_range = "*"
  source_address_prefix = "*"
  destination_address_prefix = "*"
  resource_group_name  = "${azurerm_resource_group.test.name}"
  network_security_group_name = "${azurerm_network_security_group.test.name}"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the security rule. This needs to be unique across all Rules in the Network Security Group. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the Network Security Rule. Changing this forces a new resource to be created.
- **network_security_group_name** - (Required) The name of the Network Security Group that we want to attach the rule to. Changing this forces a new resource to be created.
- **description** - (Optional) A description for this rule. Restricted to 140 characters.
- **protocol** - (Required) Network protocol this rule applies to. Possible values include `Tcp`, `Udp` or `*` (which matches both).
- **source_port_range** - (Optional) Source Port or Range. Integer or range between 0 and 65535 or `*` to match any. This is required if **source_port_ranges** is not specified.
- **source_port_ranges** - (Optional) List of source ports or port ranges. This is required if **source_port_range** is not specified.
- **destination_port_range** - (Optional) Destination Port or Range. Integer or range between 0 and 65535 or `*` to match any. This is required if **destination_port_ranges** is not specified.
- **destination_port_ranges** - (Optional) List of destination ports or port ranges. This is required if **destination_port_range** is not specified.
- **source_address_prefix** - (Optional) CIDR or source IP range or `*` to match any IP. Tags such as 'VirtualNetwork', 'AzureLoadBalancer' and 'Internet' can also be used. This is required if **source_address_prefixes** is not specified.
- **source_address_prefixes** - (Optional) List of source address prefixes. Tags may not be used. This is required if **source_address_prefix** is not specified.
- **source_application_security_group_ids** - (Optional) A List of source Application Security Group ID's
- **destination_address_prefix** - (Optional) CIDR or destination IP range or `*` to match any IP. Tags such as 'VirtualNetwork', 'AzureLoadBalancer' and 'Internet' can also be used. This is required if **destination_address_prefixes** is not specified.
- **destination_address_prefixes** - (Optional) List of destination address prefixes. Tags may not be used. This is required if **destination_address_prefix** is not specified.

- **destination_application_security_group_ids** - (Optional) A List of destination Application Security Group ID's
- **access** - (Required) Specifies whether network traffic is allowed or denied. Possible values are **Allow** and **Deny**.
- **priority** - (Required) Specifies the priority of the rule. The value can be between 100 and 4096. The priority number must be unique for each rule in the collection. The lower the priority number, the higher the priority of the rule.
- **direction** - (Required) The direction specifies if rule will be evaluated on incoming or outgoing traffic. Possible values are **Inbound** and **Outbound**.

» Attributes Reference

The following attributes are exported:

- **id** - The Network Security Rule ID.

» Import

Network Security Rules can be imported using the **resource id**, e.g.

```
terraform import azurerm_network_security_rule.rule1 /subscriptions/00000000-0000-0000-0000-
```

» azurerm__network__watcher

Manages a Network Watcher.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "production-nwwatcher"
  location = "West US"
}

resource "azurerm_network_watcher" "test" {
  name                = "production-nwwatcher"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Network Watcher. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the Network Watcher. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- **id** - The Network Watcher ID.

» Import

Network Watchers can be imported using the **resource id**, e.g.

```
terraform import azurerm_network_watcher.watcher1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm__packet__capture

Configures Packet Capturing against a Virtual Machine using a Network Watcher.

NOTE: This resource has been deprecated in favour of the **azurerm_network_connection_monitor** resource and will be removed in the next major version of the AzureRM Provider. The new resource shares the same fields as this one, and information on migrating across can be found in this guide.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "packet-capture-rg"
  location = "West Europe"
}
```

```

resource "azurerm_network_watcher" "test" {
  name           = "network-watcher"
  location       = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_virtual_network" "test" {
  name           = "production-network"
  address_space = ["10.0.0.0/16"]
  location       = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_subnet" "test" {
  name           = "internal"
  resource_group_name = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix   = "10.0.2.0/24"
}

resource "azurerm_network_interface" "test" {
  name           = "pctest-nic"
  location       = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  ip_configuration {
    name           = "testconfiguration1"
    subnet_id      = "${azurerm_subnet.test.id}"
    private_ip_address_allocation = "Dynamic"
  }
}

resource "azurerm_virtual_machine" "test" {
  name           = "pctest-vm"
  location       = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  network_interface_ids = ["${azurerm_network_interface.test.id}"]
  vm_size        = "Standard_F2"

  storage_image_reference {
    publisher = "Canonical"
    offer     = "UbuntuServer"
    sku       = "16.04-LTS"
    version   = "latest"
  }
}

```

```

storage_os_disk {
    name          = "osdisk"
    caching       = "ReadWrite"
    create_option = "FromImage"
    managed_disk_type = "Standard_LRS"
}

os_profile {
    computer_name = "pctest-vm"
    admin_username = "testadmin"
    admin_password = "Password1234!"
}

os_profile_linux_config {
    disable_password_authentication = false
}
}

resource "azurerm_virtual_machine_extension" "test" {
    name                = "network-watcher"
    location            = "${azurerm_resource_group.test.location}"
    resource_group_name = "${azurerm_resource_group.test.name}"
    virtual_machine_name = "${azurerm_virtual_machine.test.name}"
    publisher          = "Microsoft.Azure.NetworkWatcher"
    type               = "NetworkWatcherAgentLinux"
    type_handler_version = "1.4"
    auto_upgrade_minor_version = true
}

resource "azurerm_storage_account" "test" {
    name                = "pctestsa"
    resource_group_name = "${azurerm_resource_group.test.name}"
    location            = "${azurerm_resource_group.test.location}"
    account_tier        = "Standard"
    account_replication_type = "LRS"
}

resource "azurerm_packet_capture" "test" {
    name                = "pctestcapture"
    network_watcher_name = "${azurerm_network_watcher.test.name}"
    resource_group_name = "${azurerm_resource_group.test.name}"
    target_resource_id   = "${azurerm_virtual_machine.test.id}"

    storage_location {
        storage_account_id = "${azurerm_storage_account.test.id}"
    }
}

```

```

    depends_on = ["azurerm_virtual_machine_extension.test"]
}

```

NOTE: This Resource requires that the Network Watcher Virtual Machine Extension is installed on the Virtual Machine before capturing can be enabled which can be installed via the `azurerm_virtual_machine_extension` resource.

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name to use for this Packet Capture. Changing this forces a new resource to be created.
- **network_watcher_name** - (Required) The name of the Network Watcher. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the Network Watcher exists. Changing this forces a new resource to be created.
- **target_resource_id** - (Required) The ID of the Resource to capture packets from. Changing this forces a new resource to be created.

NOTE: Currently only Virtual Machines ID's are supported.

- **maximum_bytes_per_packet** - (Optional) The number of bytes captured per packet. The remaining bytes are truncated. Defaults to 0 (Entire Packet Captured). Changing this forces a new resource to be created.
- **maximum_bytes_per_session** - (Optional) Maximum size of the capture in Bytes. Defaults to 1073741824 (1GB). Changing this forces a new resource to be created.
- **maximum_capture_duration** - (Optional) The maximum duration of the capture session in seconds. Defaults to 18000 (5 hours). Changing this forces a new resource to be created.
- **storage_location** - (Required) A `storage_location` block as defined below. Changing this forces a new resource to be created.
- **filter** - (Optional) One or more `filter` blocks as defined below. Changing this forces a new resource to be created.

A `storage_location` block contains:

- **file_path** - (Optional) A valid local path on the targeting VM. Must include the name of the capture file (*.cap). For linux virtual machine it must start with `/var/captures`.

- **storage_account_id** - (Optional) The ID of the storage account to save the packet capture session

NOTE: At least one of **file_path** or **storage_account_id** must be specified.

A **filter** block contains:

- **local_ip_address** - (Optional) The local IP Address to be filtered on. Notation: "127.0.0.1" for single address entry. "127.0.0.1-127.0.0.255" for range. "127.0.0.1;127.0.0.5" for multiple entries. Multiple ranges not currently supported. Mixing ranges with multiple entries not currently supported. Changing this forces a new resource to be created.
- **local_port** - (Optional) The local port to be filtered on. Notation: "80" for single port entry."80-85" for range. "80;443;" for multiple entries. Multiple ranges not currently supported. Mixing ranges with multiple entries not currently supported. Changing this forces a new resource to be created.
- **protocol** - (Required) The Protocol to be filtered on. Possible values include **Any**, **TCP** and **UDP**. Changing this forces a new resource to be created.
- **remote_ip_address** - (Optional) The remote IP Address to be filtered on. Notation: "127.0.0.1" for single address entry. "127.0.0.1-127.0.0.255" for range. "127.0.0.1;127.0.0.5;" for multiple entries. Multiple ranges not currently supported. Mixing ranges with multiple entries not currently supported.. Changing this forces a new resource to be created.
- **remote_port** - (Optional) The remote port to be filtered on. Notation: "80" for single port entry."80-85" for range. "80;443;" for multiple entries. Multiple ranges not currently supported. Mixing ranges with multiple entries not currently supported. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The Packet Capture ID.
- **storage_location** - (Required) A **storage_location** block as defined below.

A **storage_location** block contains:

- **storage_path** - The URI of the storage path to save the packet capture.

» Import

Packet Captures can be imported using the `resource id`, e.g.

```
terraform import azurerm_packet_capture.capture1 /subscriptions/00000000-0000-0000-0000-0000
```

» azurerm__packet__capture

Configures Network Packet Capturing against a Virtual Machine using a Network Watcher.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "packet-capture-rg"
  location = "West Europe"
}

resource "azurerm_network_watcher" "test" {
  name                = "network-watcher"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_virtual_network" "test" {
  name                = "production-network"
  address_space       = ["10.0.0.0/16"]
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_subnet" "test" {
  name                = "internal"
  resource_group_name = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix       = "10.0.2.0/24"
}

resource "azurerm_network_interface" "test" {
  name                = "pctest-nic"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  ip_configuration {
```

```

        name = "testconfiguration1"
        subnet_id = "${azurerm_subnet.test.id}"
        private_ip_address_allocation = "Dynamic"
    }
}

resource "azurerm_virtual_machine" "test" {
    name = "pctest-vm"
    location = "${azurerm_resource_group.test.location}"
    resource_group_name = "${azurerm_resource_group.test.name}"
    network_interface_ids = ["${azurerm_network_interface.test.id}"]
    vm_size = "Standard_F2"

    storage_image_reference {
        publisher = "Canonical"
        offer = "UbuntuServer"
        sku = "16.04-LTS"
        version = "latest"
    }

    storage_os_disk {
        name = "osdisk"
        caching = "ReadWrite"
        create_option = "FromImage"
        managed_disk_type = "Standard_LRS"
    }

    os_profile {
        computer_name = "pctest-vm"
        admin_username = "testadmin"
        admin_password = "Password1234!"
    }

    os_profile_linux_config {
        disable_password_authentication = false
    }
}

resource "azurerm_virtual_machine_extension" "test" {
    name = "network-watcher"
    location = "${azurerm_resource_group.test.location}"
    resource_group_name = "${azurerm_resource_group.test.name}"
    virtual_machine_name = "${azurerm_virtual_machine.test.name}"
    publisher = "Microsoft.Azure.NetworkWatcher"
    type = "NetworkWatcherAgentLinux"
    type_handler_version = "1.4"
}

```

```

    auto_upgrade_minor_version = true
}

resource "azurerm_storage_account" "test" {
  name                = "pctestsa"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location            = "${azurerm_resource_group.test.location}"
  account_tier        = "Standard"
  account_replication_type = "LRS"
}

resource "azurerm_network_packet_capture" "test" {
  name                = "pctestcapture"
  network_watcher_name = "${azurerm_network_watcher.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  target_resource_id  = "${azurerm_virtual_machine.test.id}"

  storage_location {
    storage_account_id = "${azurerm_storage_account.test.id}"
  }

  depends_on = ["azurerm_virtual_machine_extension.test"]
}

```

NOTE: This Resource requires that the Network Watcher Virtual Machine Extension is installed on the Virtual Machine before capturing can be enabled which can be installed via the `azurerm_virtual_machine_extension` resource.

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name to use for this Network Packet Capture. Changing this forces a new resource to be created.
- **network_watcher_name** - (Required) The name of the Network Watcher. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the Network Watcher exists. Changing this forces a new resource to be created.
- **target_resource_id** - (Required) The ID of the Resource to capture packets from. Changing this forces a new resource to be created.

NOTE: Currently only Virtual Machines ID's are supported.

- **maximum_bytes_per_packet** - (Optional) The number of bytes captured per packet. The remaining bytes are truncated. Defaults to 0 (Entire Packet Captured). Changing this forces a new resource to be created.
- **maximum_bytes_per_session** - (Optional) Maximum size of the capture in Bytes. Defaults to 1073741824 (1GB). Changing this forces a new resource to be created.
- **maximum_capture_duration** - (Optional) The maximum duration of the capture session in seconds. Defaults to 18000 (5 hours). Changing this forces a new resource to be created.
- **storage_location** - (Required) A **storage_location** block as defined below. Changing this forces a new resource to be created.
- **filter** - (Optional) One or more **filter** blocks as defined below. Changing this forces a new resource to be created.

A **storage_location** block contains:

- **file_path** - (Optional) A valid local path on the targeting VM. Must include the name of the capture file (*.cap). For linux virtual machine it must start with /var/captures.
- **storage_account_id** - (Optional) The ID of the storage account to save the packet capture session

NOTE: At least one of **file_path** or **storage_account_id** must be specified.

A **filter** block contains:

- **local_ip_address** - (Optional) The local IP Address to be filtered on. Notation: "127.0.0.1" for single address entry. "127.0.0.1-127.0.0.255" for range. "127.0.0.1;127.0.0.5" for multiple entries. Multiple ranges not currently supported. Mixing ranges with multiple entries not currently supported. Changing this forces a new resource to be created.
- **local_port** - (Optional) The local port to be filtered on. Notation: "80" for single port entry. "80-85" for range. "80;443;" for multiple entries. Multiple ranges not currently supported. Mixing ranges with multiple entries not currently supported. Changing this forces a new resource to be created.
- **protocol** - (Required) The Protocol to be filtered on. Possible values include **Any**, **TCP** and **UDP**. Changing this forces a new resource to be created.
- **remote_ip_address** - (Optional) The remote IP Address to be filtered on. Notation: "127.0.0.1" for single address entry. "127.0.0.1-127.0.0.255" for range. "127.0.0.1;127.0.0.5;" for multiple entries. Multiple ranges not

currently supported. Mixing ranges with multiple entries not currently supported.. Changing this forces a new resource to be created.

- **remote_port** - (Optional) The remote port to be filtered on. Notation: "80" for single port entry."80-85" for range. "80;443;" for multiple entries. Multiple ranges not currently supported. Mixing ranges with multiple entries not currently supported. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The Packet Capture ID.
- **storage_location** - (Required) A **storage_location** block as defined below.

A **storage_location** block contains:

- **storage_path** - The URI of the storage path to save the packet capture.

» Import

Packet Captures can be imported using the **resource id**, e.g.

```
terraform import azurerm_network_packet_capture.capture1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_public_ip

Manage a Public IP Address.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "resourceGroup1"
  location = "West US"
}

resource "azurerm_public_ip" "test" {
  name                 = "acceptanceTestPublicIp1"
  location              = "West US"
  resource_group_name = "${azurerm_resource_group.test.name}"
}
```

```
allocation_method    = "Static"

tags = {
    environment = "Production"
}
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Public IP resource . Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the public ip.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **sku** - (Optional) The SKU of the Public IP. Accepted values are **Basic** and **Standard**. Defaults to **Basic**.

Note Public IP Standard SKUs require **allocation_method** to be set to **Static**.

Note: The **Standard** SKU is currently in Public Preview on an opt-in basis. More information, including how you can register for the Preview, and which regions **Standard** SKU's are available in are available [here](#)

- **allocation_method** - (Required) Defines the allocation method for this IP address. Possible values are **Static** or **Dynamic**.

Note **Dynamic** Public IP Addresses aren't allocated until they're assigned to a resource (such as a Virtual Machine or a Load Balancer) by design within Azure - more information is available below.

- **ip_version** - (Optional) The IP Version to use, IPv6 or IPv4.

Note Only **dynamic** IP address allocation is supported for IPv6.

- **idle_timeout_in_minutes** - (Optional) Specifies the timeout for the TCP idle connection. The value can be set between 4 and 30 minutes.
- **domain_name_label** - (Optional) Label for the Domain Name. Will be used to make up the FQDN. If a domain name label is specified, an A DNS record is created for the public IP in the Microsoft Azure DNS system.
- **reverse_fqdn** - (Optional) A fully qualified domain name that resolves to this public IP address. If the reverseFqdn is specified, then a PTR DNS

record is created pointing from the IP address in the in-addr.arpa domain to the reverse FQDN.

- **public_ip_prefix_id** - (Optional) If specified then public IP address allocated will be provided from the public IP prefix resource.

Please Note: Public IP Prefix are currently in Public Preview. You can find more information about Public IP Prefix Preview [here](#).

- **tags** - (Optional) A mapping of tags to assign to the resource.
- **zones** - (Optional) A collection containing the availability zone to allocate the Public IP in.

Please Note: Availability Zones are only supported in several regions at this time.

» Attributes Reference

The following attributes are exported:

- **id** - The Public IP ID.
- **ip_address** - The IP address value that was allocated.

Note Dynamic Public IP Addresses aren't allocated until they're attached to a device (e.g. a Virtual Machine/Load Balancer). Instead you can obtain the IP Address once the the Public IP has been assigned via the **azurerm_public_ip** Data Source.

- **fqdn** - Fully qualified domain name of the A DNS record associated with the public IP. This is the concatenation of the domainNameLabel and the regionalized DNS zone

» Import

Public IPs can be imported using the **resource id**, e.g.

```
terraform import azurerm_public_ip.myPublicIp /subscriptions/00000000-0000-0000-0000-00000000
```

» azurerm_route

Manages a Route within a Route Table.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "acceptanceTestResourceGroup1"
  location  = "West US"
}

resource "azurerm_route_table" "test" {
  name                = "acceptanceTestRouteTable1"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_route" "test" {
  name                = "acceptanceTestRoute1"
  resource_group_name = "${azurerm_resource_group.test.name}"
  route_table_name    = "${azurerm_route_table.test.name}"
  address_prefix      = "10.1.0.0/16"
  next_hop_type       = "vnetlocal"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the route. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the route. Changing this forces a new resource to be created.
- **route_table_name** - (Required) The name of the route table within which create the route. Changing this forces a new resource to be created.
- **address_prefix** - (Required) The destination CIDR to which the route applies, such as 10.1.0.0/16
- **next_hop_type** - (Required) The type of Azure hop the packet should be sent to. Possible values are `VirtualNetworkGateway`, `VnetLocal`, `Internet`, `VirtualAppliance` and `None`
- **next_hop_in_ip_address** - (Optional) Contains the IP address packets should be forwarded to. Next hop values are only allowed in routes where the next hop type is `VirtualAppliance`.

» Attributes Reference

The following attributes are exported:

- `id` - The Route ID.

» Import

Routes can be imported using the `resource id`, e.g.

```
terraform import azurerm_route.testRoute /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_route_table`

Manages a Route Table

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "acceptanceTestResourceGroup1"
  location = "West US"
}

resource "azurerm_route_table" "test" {
  name                        = "acceptanceTestSecurityGroup1"
  location                   = "${azurerm_resource_group.test.location}"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  disable_bgp_route_propagation = false

  route {
    name           = "route1"
    address_prefix = "10.1.0.0/16"
    next_hop_type  = "vnetlocal"
  }

  tags = {
    environment = "Production"
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the route table. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the route table. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **route** - (Optional) Can be specified multiple times to define multiple routes. Each **route** block supports fields documented below.
- **disable_bgp_route_propagation** - (Optional) Boolean flag which controls propagation of routes learned by BGP on that route table. True means disable.
- **tags** - (Optional) A mapping of tags to assign to the resource.

The **route** block supports:

- **name** - (Required) The name of the route.
- **address_prefix** - (Required) The destination CIDR to which the route applies, such as 10.1.0.0/16
- **next_hop_type** - (Required) The type of Azure hop the packet should be sent to. Possible values are **VirtualNetworkGateway**, **VnetLocal**, **Internet**, **VirtualAppliance** and **None**.
- **next_hop_in_ip_address** - (Optional) Contains the IP address packets should be forwarded to. Next hop values are only allowed in routes where the next hop type is **VirtualAppliance**.

» Attributes Reference

The following attributes are exported:

- **id** - The Route Table ID.
- **subnets** - The collection of Subnets associated with this route table.

» Import

Route Tables can be imported using the **resource id**, e.g.

```
terraform import azurerm_route_table.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm__subnet

Manages a subnet. Subnets represent network segments within the IP space defined by the virtual network.

NOTE on Virtual Networks and Subnet's: Terraform currently provides both a standalone Subnet resource, and allows for Subnets to be defined in-line within the Virtual Network resource. At this time you cannot use a Virtual Network with in-line Subnets in conjunction with any Subnet resources. Doing so will cause a conflict of Subnet configurations and will overwrite Subnet's.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "acceptanceTestResourceGroup1"
  location  = "West US"
}

resource "azurerm_virtual_network" "test" {
  name            = "acceptanceTestVirtualNetwork1"
  address_space   = ["10.0.0.0/16"]
  location        = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_subnet" "test" {
  name                = "testsubnet"
  resource_group_name = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix       = "10.0.1.0/24"
  delegation {
    name = "acctestdelegation"
    service_delegation {
      name      = "Microsoft.ContainerInstance/containerGroups"
      actions   = ["Microsoft.Network/virtualNetworks/subnets/action"]
    }
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the subnet. Changing this forces a new resource to be created.

- **resource_group_name** - (Required) The name of the resource group in which to create the subnet. Changing this forces a new resource to be created.
- **virtual_network_name** - (Required) The name of the virtual network to which to attach the subnet. Changing this forces a new resource to be created.
- **address_prefix** - (Required) The address prefix to use for the subnet.
- **network_security_group_id** - (Optional / **Deprecated**) The ID of the Network Security Group to associate with the subnet.

NOTE: At this time Subnet <=> Network Security Group associations need to be configured both using this field (which is now Deprecated) and/or using the `azurerm_subnet_network_security_group_association` resource. This field is deprecated and will be removed in favour of that resource in the next major version (2.0) of the AzureRM Provider.

- **route_table_id** - (Optional / **Deprecated**) The ID of the Route Table to associate with the subnet.

NOTE: At this time Subnet <=> Route Table associations need to be configured both using this field (which is now Deprecated) and/or using the `azurerm_subnet_route_table_association` resource. This field is deprecated and will be removed in favour of that resource in the next major version (2.0) of the AzureRM Provider.

- **service_endpoints** - (Optional) The list of Service endpoints to associate with the subnet. Possible values include: `Microsoft.AzureActiveDirectory`, `Microsoft.AzureCosmosDB`, `Microsoft.EventHub`, `Microsoft.KeyVault`, `Microsoft.ServiceBus`, `Microsoft.Sql` and `Microsoft.Storage`.
- **delegation** - (Optional) One or more `delegation` blocks as defined below.

A `delegation` block supports the following: * **name** (Required) A name for this delegation. * **service_delegation** (Required) A `service_delegation` block as defined below.

A `service_delegation` block supports the following:

NOTE: Delegating to services may not be available in all regions. Check that the service you are delegating to is available in your region using the Azure CLI

- **name** - (Required) The name of service to delegate to. Possible values include: `Microsoft.Batch/batchAccounts`, `Microsoft.ContainerInstance/containerGroups`, `Microsoft.HardwareSecurityModules/dedicatedHSMs`, `Microsoft.Logic/integrationServiceEnvironments`, `Microsoft.Netapp/volumes`, `Microsoft.ServiceFabricMesh/networks`,

Microsoft.Sql/managedInstances, Microsoft.Sql/servers or Microsoft.Web/serverFarms.

- **actions** - (Optional) A list of Actions which should be delegated. Possible values include: Microsoft.Network/virtualNetworks/subnets/action.

» Attributes Reference

The following attributes are exported:

- **id** - The subnet ID.
- **ip_configurations** - The collection of IP Configurations with IPs within this subnet.
- **name** - The name of the subnet.
- **resource_group_name** - The name of the resource group in which the subnet is created in.
- **virtual_network_name** - The name of the virtual network in which the subnet is created in
- **address_prefix** - The address prefix for the subnet

» Import

Subnets can be imported using the **resource id**, e.g.

```
terraform import azurerm_subnet.testSubnet /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_subnet_network_security_group_association

Associates a Network Security Group with a Subnet within a Virtual Network.

NOTE: Subnet <-> Network Security Group associations currently need to be configured on both this resource and using the **network_security_group_id** field on the **azurerm_subnet** resource. The next major version of the AzureRM Provider (2.0) will remove the **network_security_group_id** field from the **azurerm_subnet** resource such that this resource is used to link resources in future.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "example-resources"
  location = "West Europe"
}
```

```

resource "azurerm_virtual_network" "test" {
  name                = "example-network"
  address_space       = ["10.0.0.0/16"]
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_subnet" "test" {
  name                = "frontend"
  resource_group_name = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix       = "10.0.2.0/24"
  network_security_group_id = "${azurerm_network_security_group.test.id}"
}

resource "azurerm_network_security_group" "test" {
  name                = "example-nsg"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  security_rule {
    name                = "test123"
    priority             = 100
    direction            = "Inbound"
    access               = "Allow"
    protocol             = "Tcp"
    source_port_range    = "*"
    destination_port_range = "*"
    source_address_prefix = "*"
    destination_address_prefix = "*"
  }
}

resource "azurerm_subnet_network_security_group_association" "test" {
  subnet_id                = "${azurerm_subnet.test.id}"
  network_security_group_id = "${azurerm_network_security_group.test.id}"
}

```

» Argument Reference

The following arguments are supported:

- `network_security_group_id` - (Required) The ID of the Network Security Group which should be associated with the Subnet. Changing this forces a new resource to be created.

- `subnet_id` - (Required) The ID of the Subnet. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the Subnet.

» Import

Subnet <-> Network Security Group Associations can be imported using the `resource id` of the Subnet, e.g.

```
terraform import azurerm_subnet_network_security_group_association.association1 /subscriptions/...
```

» `azurerm_subnet_route_table_association`

Associates a Route Table with a Subnet within a Virtual Network.

NOTE: Subnet <-> Route Table associations currently need to be configured on both this resource and using the `route_table_id` field on the `azurerm_subnet` resource. The next major version of the AzureRM Provider (2.0) will remove the `route_table_id` field from the `azurerm_subnet` resource such that this resource is used to link resources in future.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "example-resources"
  location = "West Europe"
}

resource "azurerm_virtual_network" "test" {
  name            = "example-network"
  address_space   = ["10.0.0.0/16"]
  location        = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_subnet" "test" {
  name                 = "frontend"
  resource_group_name = "${azurerm_resource_group.test.name}"
}
```



```

    virtual_network_name = "${azurerm_virtual_network.test.name}"
    address_prefix        = "10.0.2.0/24"
    route_table_id        = "${azurerm_route_table.test.id}"
  }

  resource "azurerm_route_table" "test" {
    name                = "example-routetable"
    location            = "${azurerm_resource_group.test.location}"
    resource_group_name = "${azurerm_resource_group.test.name}"

    route {
      name                = "example"
      address_prefix      = "10.100.0.0/14"
      next_hop_type       = "VirtualAppliance"
      next_hop_in_ip_address = "10.10.1.1"
    }
  }

  resource "azurerm_subnet_route_table_association" "test" {
    subnet_id      = "${azurerm_subnet.test.id}"
    route_table_id = "${azurerm_route_table.test.id}"
  }

```

» Argument Reference

The following arguments are supported:

- **route_table_id** - (Required) The ID of the Route Table which should be associated with the Subnet. Changing this forces a new resource to be created.
- **subnet_id** - (Required) The ID of the Subnet. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Subnet.

» Import

Subnet Route Table Associations can be imported using the **resource id** of the Subnet, e.g.

```
terraform import azurerm_subnet_route_table_association.association1 /subscriptions/00000000
```

» azurerm__traffic__manager__endpoint

Manages a Traffic Manager Endpoint.

» Example Usage

```
resource "random_id" "server" {
  keepers = {
    azi_id = 1
  }

  byte_length = 8
}

resource "azurerm_resource_group" "test" {
  name     = "trafficmanagerendpointTest"
  location = "West US"
}

resource "azurerm_traffic_manager_profile" "test" {
  name                       = "${random_id.server.hex}"
  resource_group_name       = "${azurerm_resource_group.test.name}"

  traffic_routing_method = "Weighted"

  dns_config {
    relative_name = "${random_id.server.hex}"
    ttl           = 100
  }

  monitor_config {
    protocol = "http"
    port     = 80
    path     = "/"
  }

  tags = {
    environment = "Production"
  }
}
```

```

resource "azurerm_traffic_manager_endpoint" "test" {
  name                = "${random_id.server.hex}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  profile_name        = "${azurerm_traffic_manager_profile.test.name}"
  target              = "terraform.io"
  type                = "externalEndpoints"
  weight              = 100
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Traffic Manager endpoint. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the Traffic Manager endpoint.
- **profile_name** - (Required) The name of the Traffic Manager Profile to attach create the Traffic Manager endpoint.
- **endpoint_status** - (Optional) The status of the Endpoint, can be set to either **Enabled** or **Disabled**. Defaults to **Enabled**.
- **type** - (Required) The Endpoint type, must be one of:
 - **azureEndpoints**
 - **externalEndpoints**
 - **nestedEndpoints**
- **target** - (Optional) The FQDN DNS name of the target. This argument must be provided for an endpoint of type **externalEndpoints**, for other types it will be computed.
- **target_resource_id** - (Optional) The resource id of an Azure resource to target. This argument must be provided for an endpoint of type **azureEndpoints** or **nestedEndpoints**.
- **weight** - (Optional) Specifies how much traffic should be distributed to this endpoint, this must be specified for Profiles using the **Weighted** traffic routing method. Supports values between 1 and 1000.
- **priority** - (Optional) Specifies the priority of this Endpoint, this must be specified for Profiles using the **Priority** traffic routing method. Supports values between 1 and 1000, with no Endpoints sharing the same value. If omitted the value will be computed in order of creation.
- **endpoint_location** - (Optional) Specifies the Azure location of the Endpoint, this must be specified for Profiles using the **Performance**

routing method if the Endpoint is of either type `nestedEndpoints` or `externalEndpoints`. For Endpoints of type `azureEndpoints` the value will be taken from the location of the Azure target resource.

- `min_child_endpoints` - (Optional) This argument specifies the minimum number of endpoints that must be 'online' in the child profile in order for the parent profile to direct traffic to any of the endpoints in that child profile. This argument only applies to Endpoints of type `nestedEndpoints` and defaults to 1.
- `geo_mappings` - (Optional) A list of Geographic Regions used to distribute traffic, such as `WORLD`, `UK` or `DE`. The same location can't be specified in two endpoints. See the Geographic Hierarchies documentation for more information.

» Attributes Reference

The following attributes are exported:

- `id` - The Traffic Manager Endpoint id.

» Import

Traffic Manager Endpoints can be imported using the `resource id`, e.g.

```
terraform import azurerm_traffic_manager_endpoint.testEndpoints /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_traffic_manager_profile`

Manages a Traffic Manager Profile to which multiple endpoints can be attached.

» Example Usage

```
resource "random_id" "server" {
  keepers = {
    azi_id = 1
  }

  byte_length = 8
}

resource "azurerm_resource_group" "test" {
  name = "trafficmanagerProfile"
```

```

    location = "West US"
  }

  resource "azurerm_traffic_manager_profile" "test" {
    name                       = "${random_id.server.hex}"
    resource_group_name       = "${azurerm_resource_group.test.name}"
    traffic_routing_method    = "Weighted"

    dns_config {
      relative_name = "${random_id.server.hex}"
      ttl           = 100
    }

    monitor_config {
      protocol = "http"
      port     = 80
      path     = "/"
    }

    tags = {
      environment = "Production"
    }
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the virtual network. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the virtual network.
- **profile_status** - (Optional) The status of the profile, can be set to either `Enabled` or `Disabled`. Defaults to `Enabled`.
- **traffic_routing_method** - (Required) Specifies the algorithm used to route traffic, possible values are:
 - **Geographic** - Traffic is routed based on Geographic regions specified in the Endpoint.
 - **MultiValue** - All healthy Endpoints are returned. MultiValue routing method works only if all the endpoints of type 'External' and are specified as IPv4 or IPv6 addresses.
 - **Performance** - Traffic is routed via the User's closest Endpoint

- **Priority** - Traffic is routed to the Endpoint with the lowest **priority** value.
 - **Subnet** - Traffic is routed based on a mapping of sets of end-user IP address ranges to a specific Endpoint within a Traffic Manager profile.
 - **Weighted** - Traffic is spread across Endpoints proportional to their **weight** value.
- **dns_config** - (Required) This block specifies the DNS configuration of the Profile, it supports the fields documented below.
 - **monitor_config** - (Required) This block specifies the Endpoint monitoring configuration for the Profile, it supports the fields documented below.
 - **tags** - (Optional) A mapping of tags to assign to the resource.

The **dns_config** block supports:

- **relative_name** - (Required) The relative domain name, this is combined with the domain name used by Traffic Manager to form the FQDN which is exported as documented below. Changing this forces a new resource to be created.
- **t11** - (Required) The TTL value of the Profile used by Local DNS resolvers and clients.

The **monitor_config** block supports:

- **protocol** - (Required) The protocol used by the monitoring checks, supported values are **HTTP**, **HTTPS** and **TCP**.
- **port** - (Required) The port number used by the monitoring checks.
- **path** - (Optional) The path used by the monitoring checks. Required when **protocol** is set to **HTTP** or **HTTPS** - cannot be set when **protocol** is set to **TCP**.

» Attributes Reference

The following attributes are exported:

- **id** - The Traffic Manager Profile id.
- **fqdn** - The FQDN of the created Profile.

» Notes

The Traffic Manager is created with the location **global**.

» Import

Traffic Manager Profiles can be imported using the `resource id`, e.g.

```
terraform import azurerm_traffic_manager_profile.testProfile /subscriptions/00000000-0000-00
```

» azurerm_virtual_network

Manages a virtual network including any configured subnets. Each subnet can optionally be configured with a security group to be associated with the subnet.

NOTE on Virtual Networks and Subnet's: Terraform currently provides both a standalone Subnet resource, and allows for Subnets to be defined in-line within the Virtual Network resource. At this time you cannot use a Virtual Network with in-line Subnets in conjunction with any Subnet resources. Doing so will cause a conflict of Subnet configurations and will overwrite Subnet's.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "acceptanceTestResourceGroup1"
  location  = "West US"
}

resource "azurerm_network_security_group" "test" {
  name                = "acceptanceTestSecurityGroup1"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_ddos_protection_plan" "test" {
  name                = "ddosplan1"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}

resource "azurerm_virtual_network" "test" {
  name                = "virtualNetwork1"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  address_space       = ["10.0.0.0/16"]
  dns_servers          = ["10.0.0.4", "10.0.0.5"]

  ddos_protection_plan {
```

```

    id      = "${azurerm_ddos_protection_plan.test.id}"
    enable = true
  }

  subnet {
    name           = "subnet1"
    address_prefix = "10.0.1.0/24"
  }

  subnet {
    name           = "subnet2"
    address_prefix = "10.0.2.0/24"
  }

  subnet {
    name           = "subnet3"
    address_prefix = "10.0.3.0/24"
    security_group = "${azurerm_network_security_group.test.id}"
  }

  tags = {
    environment = "Production"
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the virtual network. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the virtual network.
- **address_space** - (Required) The address space that is used the virtual network. You can supply more than one address space. Changing this forces a new resource to be created.
- **location** - (Required) The location/region where the virtual network is created. Changing this forces a new resource to be created.
- **ddos_protection_plan** - (Optional) A `ddos_protection_plan` block as documented below.
- **dns_servers** - (Optional) List of IP addresses of DNS servers

- **subnet** - (Optional) Can be specified multiple times to define multiple subnets. Each **subnet** block supports fields documented below.
 - **tags** - (Optional) A mapping of tags to assign to the resource.
-

A **ddos_protection_plan** block supports the following:

- **id** - (Required) The Resource ID of DDoS Protection Plan.
 - **enable** - (Required) Enable/disable DDoS Protection Plan on Virtual Network.
-

The **subnet** block supports:

- **name** - (Required) The name of the subnet.
- **address_prefix** - (Required) The address prefix to use for the subnet.
- **security_group** - (Optional) The Network Security Group to associate with the subnet. (Referenced by **id**, ie. `azurerm_network_security_group.test.id`)

» Attributes Reference

The following attributes are exported:

- **id** - The virtual NetworkConfiguration ID.
 - **name** - The name of the virtual network.
 - **resource_group_name** - The name of the resource group in which to create the virtual network.
 - **location** - The location/region where the virtual network is created
 - **address_space** - The address space that is used the virtual network.
 - **subnet**- One or more **subnet** blocks as defined below.
-

The **subnet** block exports:

- **id** - The ID of this subnet.

» Import

Virtual Networks can be imported using the **resource id**, e.g.

```
terraform import azurerm_virtual_network.testNetwork /subscriptions/00000000-0000-0000-0000-
```

» azurerm_virtual_network_gateway

Manages a Virtual Network Gateway to establish secure, cross-premises connectivity.

Note: Please be aware that provisioning a Virtual Network Gateway takes a long time (between 30 minutes and 1 hour)

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "test"
  location = "West US"
}

resource "azurerm_virtual_network" "test" {
  name                = "test"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  address_space       = ["10.0.0.0/16"]
}

resource "azurerm_subnet" "test" {
  name                 = "GatewaySubnet"
  resource_group_name = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix       = "10.0.1.0/24"
}

resource "azurerm_public_ip" "test" {
  name                = "test"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  allocation_method = "Dynamic"
}

resource "azurerm_virtual_network_gateway" "test" {
  name                = "test"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  type      = "Vpn"
  vpn_type = "RouteBased"
}
```

```

active_active = false
enable_bgp    = false
sku           = "Basic"

ip_configuration {
  name                        = "vnetGatewayConfig"
  public_ip_address_id       = "${azurerm_public_ip.test.id}"
  private_ip_address_allocation = "Dynamic"
  subnet_id                  = "${azurerm_subnet.test.id}"
}

vpn_client_configuration {
  address_space = ["10.2.0.0/24"]

  root_certificate {
    name = "DigiCert-Federated-ID-Root-CA"

    public_cert_data = <<EOF
MIIDuzCCAqOgAwIBAgIQCHTZWCM+IlfFIRXIvyKSrjANBgkqhkiG9w0BAQsFADBN
MQswCQYDVQQGEwJVUzEVMBMGA1UEChMMRGlnaUNlcnQgSW5jMRkwFwYDVQQLExB3
d3cuZGlnaWlnaUNlcnQy29tMSYwJAYDVQQDEx1EaWdpQ2VydCBGZWRLcmF0ZWQgSUQg
Um9vdCBDQTAEFw0xMzAxMTUxMjAwMDBaFw0zMTUxMjAwMDBaMGBxGzAJBgNV
BAYTA1VTMRUwEwYDVQQKEwxEaWdpQ2VydCBJbmMxGTAXBgNVBAsTEHd3dy5kaWdp
Y2VydC5jb20xJjAkBgNVBAMTHURpZ21DZXJ0IEZlZGVyYXR1ZCBJRCBSb290IENB
MIIBIjANBgkqhkiG9w0BAQEFAAOCAQ8AMIIBCgKCAQEAvAEB4pcCqnNNOWE6Ur5j
QPUH+1y1F9KdHTRSza6k5iDlXq1kGS1qAkuKtw9JsiNRrjltmFnzMZRBbX8Tlfl8
zAhBmb6dDduDGED01kBsTkgwyYPxXVTKec0WxYEEF0oMn4wSYN10lt2eJAKHXjNf
GTwiibdP8CUR2ghSM2sUTI8Nt10mfc4SMHhGhYD64uJMbX98THQ/4LMGuYegou+d
GTiahfHtjn7AboSEknwAMJHCh5RlYZZ6B104QbKJ+34Q0eKgnI3X6Vc9u0zf6DH8
Dk+4zQDYRRtQtTnV03VT8jzqDlCRuNtq6YvryOWN74/dq8LQhUnXHvFyrSDMaE1X2
DwIDAQABo2MwYTAAPBgNVHRMBAf8EBTADAQH/MA4GA1UdDwEB/wQEAwIBhjAdBgNV
HQ4EFgQUGRdkFnbGt1EWjKwbUne+50aZvRYwHwYDVR0jBBgwFoAUGRdkFnbGt1EW
jKwbUne+50aZvRYwDQYJKoZIhvcNAQELBQADggEBAHcqsHkrjpeSsfuVTRiptJfP
9JbdtWqRTmOf6uJi2c8YVqI6X1KXsD8C1dUUaaHKLUIJzvKiazibVuBwMIT84AyqR
QELn3e0BtgEymEygMU569b01ZPxoFSnNXc7qDZBDef8WfqAV/sxkTi8L9BkmFYfL
uGLOhRJOFprPdoDIUBB+tmCl3oDcBy3vnUe0Eioz8zAkprcb3GHwHAK+vHmMfgcn
WsfMLH4JCLa/tRYL+Rw/N3ybCkDp00s0WUZ+AOdywS10Q/ZEnNYOMsFiw6LyIdbq
M/s/1JRt03bDSzD9TazRVzn2oBqzSa8VgIo5C1n0noAKJTlsClJKvIhnRlaLQqk=
EOF
}

revoked_certificate {
  name          = "Verizon-Global-Root-CA"
  thumbprint    = "912198EEF23DCAC40939312FEE97DD560BAE49B1"
}
}

```

}

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Virtual Network Gateway. Changing the name forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the Virtual Network Gateway. Changing the resource group name forces a new resource to be created.
- **location** - (Required) The location/region where the Virtual Network Gateway is located. Changing the location/region forces a new resource to be created.
- **type** - (Required) The type of the Virtual Network Gateway. Valid options are **Vpn** or **ExpressRoute**. Changing the type forces a new resource to be created.
- **vpn_type** - (Optional) The routing type of the Virtual Network Gateway. Valid options are **RouteBased** or **PolicyBased**. Defaults to **RouteBased**.
- **enable_bgp** - (Optional) If **true**, BGP (Border Gateway Protocol) will be enabled for this Virtual Network Gateway. Defaults to **false**.
- **active_active** - (Optional) If **true**, an active-active Virtual Network Gateway will be created. An active-active gateway requires a **HighPerformance** or an **UltraPerformance** sku. If **false**, an active-standby gateway will be created. Defaults to **false**.
- **default_local_network_gateway_id** - (Optional) The ID of the local network gateway through which outbound Internet traffic from the virtual network in which the gateway is created will be routed (*forced tunneling*). Refer to the Azure documentation on forced tunneling. If not specified, forced tunneling is disabled.
- **sku** - (Required) Configuration of the size and capacity of the virtual network gateway. Valid options are **Basic**, **Standard**, **HighPerformance**, **UltraPerformance**, **ErGw1AZ**, **ErGw2AZ**, **ErGw3AZ**, **VpnGw1**, **VpnGw2**, **VpnGw3**, **VpnGw1AZ**, **VpnGw2AZ**, and **VpnGw3AZ** and depend on the **type** and **vpn_type** arguments. A **PolicyBased** gateway only supports the **Basic** sku. Further, the **UltraPerformance** sku is only supported by an **ExpressRoute** gateway.
- **ip_configuration** (Required) One or two **ip_configuration** blocks documented below. An active-standby gateway requires exactly one **ip_configuration** block whereas an active-active gateway requires exactly two **ip_configuration** blocks.

- **vpn_client_configuration** (Optional) A **vpn_client_configuration** block which is documented below. In this block the Virtual Network Gateway can be configured to accept IPSec point-to-site connections.
- **tags** - (Optional) A mapping of tags to assign to the resource.

The **ip_configuration** block supports:

- **name** - (Optional) A user-defined name of the IP configuration. Defaults to **vnetGatewayConfig**.
- **private_ip_address_allocation** - (Optional) Defines how the private IP address of the gateways virtual interface is assigned. Valid options are **Static** or **Dynamic**. Defaults to **Dynamic**.
- **subnet_id** - (Required) The ID of the gateway subnet of a virtual network in which the virtual network gateway will be created. It is mandatory that the associated subnet is named **GatewaySubnet**. Therefore, each virtual network can contain at most a single Virtual Network Gateway.
- **public_ip_address_id** - (Optional) The ID of the public ip address to associate with the Virtual Network Gateway.

The **vpn_client_configuration** block supports:

- **address_space** - (Required) The address space out of which ip addresses for vpn clients will be taken. You can provide more than one address space, e.g. in CIDR notation.
- **root_certificate** - (Optional) One or more **root_certificate** blocks which are defined below. These root certificates are used to sign the client certificate used by the VPN clients to connect to the gateway. This setting is incompatible with the use of **radius_server_address** and **radius_server_secret**.
- **revoked_certificate** - (Optional) One or more **revoked_certificate** blocks which are defined below. This setting is incompatible with the use of **radius_server_address** and **radius_server_secret**.
- **radius_server_address** - (Optional) The address of the Radius server. This setting is incompatible with the use of **root_certificate** and **revoked_certificate**.
- **radius_server_secret** - (Optional) The secret used by the Radius server. This setting is incompatible with the use of **root_certificate** and **revoked_certificate**.
- **vpn_client_protocols** - (Optional) List of the protocols supported by the vpn client. The supported values are **SSTP**, **IkeV2** and **OpenVPN**.

NOTE: Support for **OpenVPN** as a Client Protocol is currently in Public Preview - you can register for this Preview using [this link](#).

The `bgp_settings` block supports:

- `asn` - (Optional) The Autonomous System Number (ASN) to use as part of the BGP.
- `peering_address` - (Optional) The BGP peer IP address of the virtual network gateway. This address is needed to configure the created gateway as a BGP Peer on the on-premises VPN devices. The IP address must be part of the subnet of the Virtual Network Gateway. Changing this forces a new resource to be created.
- `peer_weight` - (Optional) The weight added to routes which have been learned through BGP peering. Valid values can be between 0 and 100.

The `root_certificate` block supports:

- `name` - (Required) A user-defined name of the root certificate.
- `public_cert_data` - (Required) The public certificate of the root certificate authority. The certificate must be provided in Base-64 encoded X.509 format (PEM). In particular, this argument *must not* include the -----BEGIN CERTIFICATE----- or -----END CERTIFICATE----- markers.

The `root_revoked_certificate` block supports:

- `name` - (Required) A user-defined name of the revoked certificate.
- `public_cert_data` - (Required) The SHA1 thumbprint of the certificate to be revoked.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the Virtual Network Gateway.

» Import

Virtual Network Gateways can be imported using the `resource id`, e.g.

```
terraform import azurerm_virtual_network_gateway.testGateway /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_virtual_network_gateway_connection`

Manages a connection in an existing Virtual Network Gateway.

» Example Usage

» Site-to-Site connection

The following example shows a connection between an Azure virtual network and an on-premises VPN device and network.

```
resource "azurerm_resource_group" "test" {
  name      = "test"
  location  = "West US"
}

resource "azurerm_virtual_network" "test" {
  name                = "test"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  address_space       = ["10.0.0.0/16"]
}

resource "azurerm_subnet" "test" {
  name                 = "GatewaySubnet"
  resource_group_name  = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix       = "10.0.1.0/24"
}

resource "azurerm_local_network_gateway" "onpremise" {
  name                = "onpremise"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  gateway_address     = "168.62.225.23"
  address_space       = ["10.1.1.0/24"]
}

resource "azurerm_public_ip" "test" {
  name                = "test"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  allocation_method   = "Dynamic"
}

resource "azurerm_virtual_network_gateway" "test" {
  name                = "test"
  location            = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
}
```

```

type          = "Vpn"
vpn_type      = "RouteBased"

active_active = false
enable_bgp    = false
sku           = "Basic"

ip_configuration {
  public_ip_address_id      = "${azurerm_public_ip.test.id}"
  private_ip_address_allocation = "Dynamic"
  subnet_id                 = "${azurerm_subnet.test.id}"
}
}

resource "azurerm_virtual_network_gateway_connection" "onpremise" {
  name          = "onpremise"
  location      = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  type          = "IPsec"
  virtual_network_gateway_id = "${azurerm_virtual_network_gateway.test.id}"
  local_network_gateway_id   = "${azurerm_local_network_gateway.onpremise.id}"

  shared_key = "4-v3ry-53cr37-1p53c-5h4r3d-k3y"
}

```

» VNet-to-VNet connection

The following example shows a connection between two Azure virtual network in different locations/regions.

```

resource "azurerm_resource_group" "us" {
  name     = "us"
  location = "East US"
}

resource "azurerm_virtual_network" "us" {
  name          = "us"
  location      = "${azurerm_resource_group.us.location}"
  resource_group_name = "${azurerm_resource_group.us.name}"
  address_space = ["10.0.0.0/16"]
}

resource "azurerm_subnet" "us_gateway" {
  name          = "GatewaySubnet"

```



```

    resource_group_name = "${azurerm_resource_group.us.name}"
    virtual_network_name = "${azurerm_virtual_network.us.name}"
    address_prefix       = "10.0.1.0/24"
}

resource "azurerm_public_ip" "us" {
    name           = "us"
    location       = "${azurerm_resource_group.us.location}"
    resource_group_name = "${azurerm_resource_group.us.name}"
    allocation_method = "Dynamic"
}

resource "azurerm_virtual_network_gateway" "us" {
    name           = "us-gateway"
    location       = "${azurerm_resource_group.us.location}"
    resource_group_name = "${azurerm_resource_group.us.name}"

    type          = "Vpn"
    vpn_type      = "RouteBased"
    sku           = "Basic"

    ip_configuration {
        public_ip_address_id      = "${azurerm_public_ip.us.id}"
        private_ip_address_allocation = "Dynamic"
        subnet_id                 = "${azurerm_subnet.us_gateway.id}"
    }
}

resource "azurerm_resource_group" "europe" {
    name     = "europe"
    location = "West Europe"
}

resource "azurerm_virtual_network" "europe" {
    name           = "europe"
    location       = "${azurerm_resource_group.europe.location}"
    resource_group_name = "${azurerm_resource_group.europe.name}"
    address_space    = ["10.1.0.0/16"]
}

resource "azurerm_subnet" "europe_gateway" {
    name           = "GatewaySubnet"
    resource_group_name = "${azurerm_resource_group.europe.name}"
    virtual_network_name = "${azurerm_virtual_network.europe.name}"
    address_prefix    = "10.1.1.0/24"
}

```

```

resource "azurerm_public_ip" "europe" {
  name                = "europe"
  location             = "${azurerm_resource_group.europe.location}"
  resource_group_name = "${azurerm_resource_group.europe.name}"
  allocation_method    = "Dynamic"
}

resource "azurerm_virtual_network_gateway" "europe" {
  name                = "europe-gateway"
  location             = "${azurerm_resource_group.europe.location}"
  resource_group_name = "${azurerm_resource_group.europe.name}"

  type      = "Vpn"
  vpn_type  = "RouteBased"
  sku       = "Basic"

  ip_configuration {
    public_ip_address_id      = "${azurerm_public_ip.europe.id}"
    private_ip_address_allocation = "Dynamic"
    subnet_id                 = "${azurerm_subnet.europe_gateway.id}"
  }
}

resource "azurerm_virtual_network_gateway_connection" "us_to_europe" {
  name                = "us-to-europe"
  location             = "${azurerm_resource_group.us.location}"
  resource_group_name = "${azurerm_resource_group.us.name}"

  type                        = "Vnet2Vnet"
  virtual_network_gateway_id = "${azurerm_virtual_network_gateway.us.id}"
  peer_virtual_network_gateway_id = "${azurerm_virtual_network_gateway.europe.id}"

  shared_key = "4-v3ry-53cr37-1p53c-5h4r3d-k3y"
}

resource "azurerm_virtual_network_gateway_connection" "europe_to_us" {
  name                = "europe-to-us"
  location             = "${azurerm_resource_group.europe.location}"
  resource_group_name = "${azurerm_resource_group.europe.name}"

  type                        = "Vnet2Vnet"
  virtual_network_gateway_id = "${azurerm_virtual_network_gateway.europe.id}"
  peer_virtual_network_gateway_id = "${azurerm_virtual_network_gateway.us.id}"

  shared_key = "4-v3ry-53cr37-1p53c-5h4r3d-k3y"
}

```

}

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the connection. Changing the name forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the connection. Changing the name forces a new resource to be created.
- **location** - (Required) The location/region where the connection is located. Changing this forces a new resource to be created.
- **type** - (Required) The type of connection. Valid options are **IPsec** (Site-to-Site), **ExpressRoute** (ExpressRoute), and **Vnet2Vnet** (VNet-to-VNet). Each connection type requires different mandatory arguments (refer to the examples above). Changing the connection type will force a new connection to be created.
- **virtual_network_gateway_id** - (Required) The ID of the Virtual Network Gateway in which the connection will be created. Changing the gateway forces a new resource to be created.
- **authorization_key** - (Optional) The authorization key associated with the Express Route Circuit. This field is required only if the type is an ExpressRoute connection.
- **express_route_circuit_id** - (Optional) The ID of the Express Route Circuit when creating an ExpressRoute connection (i.e. when **type** is **ExpressRoute**). The Express Route Circuit can be in the same or in a different subscription.
- **peer_virtual_network_gateway_id** - (Optional) The ID of the peer virtual network gateway when creating a VNet-to-VNet connection (i.e. when **type** is **Vnet2Vnet**). The peer Virtual Network Gateway can be in the same or in a different subscription.
- **local_network_gateway_id** - (Optional) The ID of the local network gateway when creating Site-to-Site connection (i.e. when **type** is **IPsec**).
- **routing_weight** - (Optional) The routing weight. Defaults to 10.
- **shared_key** - (Optional) The shared IPsec key. A key must be provided if a Site-to-Site or VNet-to-VNet connection is created whereas ExpressRoute connections do not need a shared key.
- **enable_bgp** - (Optional) If **true**, BGP (Border Gateway Protocol) is enabled for this connection. Defaults to **false**.

- **express_route_gateway_bypass** - (Optional) If **true**, data packets will bypass ExpressRoute Gateway for data forwarding. This is only valid for ExpressRoute connections.
- **use_policy_based_traffic_selectors** - (Optional) If **true**, policy-based traffic selectors are enabled for this connection. Enabling policy-based traffic selectors requires an **ipsec_policy** block. Defaults to **false**.
- **ipsec_policy** (Optional) A **ipsec_policy** block which is documented below. Only a single policy can be defined for a connection. For details on custom policies refer to the relevant section in the Azure documentation.
- **tags** - (Optional) A mapping of tags to assign to the resource.

The **ipsec_policy** block supports:

- **dh_group** - (Required) The DH group used in IKE phase 1 for initial SA. Valid options are **DHGroup1**, **DHGroup14**, **DHGroup2**, **DHGroup2048**, **DHGroup24**, **ECP256**, **ECP384**, or **None**.
- **ike_encryption** - (Required) The IKE encryption algorithm. Valid options are **AES128**, **AES192**, **AES256**, **DES**, or **DES3**.
- **ike_integrity** - (Required) The IKE integrity algorithm. Valid options are **MD5**, **SHA1**, **SHA256**, or **SHA384**.
- **ipsec_encryption** - (Required) The IPSec encryption algorithm. Valid options are **AES128**, **AES192**, **AES256**, **DES**, **DES3**, **GCMAES128**, **GCMAES192**, **GCMAES256**, or **None**.
- **ipsec_integrity** - (Required) The IPSec integrity algorithm. Valid options are **GCMAES128**, **GCMAES192**, **GCMAES256**, **MD5**, **SHA1**, or **SHA256**.
- **pfs_group** - (Required) The DH group used in IKE phase 2 for new child SA. Valid options are **ECP256**, **ECP384**, **PFS1**, **PFS2**, **PFS2048**, **PFS24**, or **None**.
- **sa_datasize** - (Optional) The IPSec SA payload size in KB. Must be at least 1024 KB. Defaults to 102400000 KB.
- **sa_lifetime** - (Optional) The IPSec SA lifetime in seconds. Must be at least 300 seconds. Defaults to 27000 seconds.

» Attributes Reference

The following attributes are exported:

- **id** - The connection ID.

» Import

Virtual Network Gateway Connections can be imported using their `resource` id, e.g.

```
terraform import azurerm_virtual_network_gateway_connection.testConnection /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_virtual_network_peering

Manages a virtual network peering which allows resources to access other resources in the linked virtual network.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "peerredvnets-rg"
  location = "West US"
}

resource "azurerm_virtual_network" "test1" {
  name                = "peternetwork1"
  resource_group_name = "${azurerm_resource_group.test.name}"
  address_space       = ["10.0.1.0/24"]
  location            = "West US"
}

resource "azurerm_virtual_network" "test2" {
  name                = "peternetwork2"
  resource_group_name = "${azurerm_resource_group.test.name}"
  address_space       = ["10.0.2.0/24"]
  location            = "West US"
}

resource "azurerm_virtual_network_peering" "test1" {
  name                = "peer1to2"
  resource_group_name = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test1.name}"
  remote_virtual_network_id = "${azurerm_virtual_network.test2.id}"
}

resource "azurerm_virtual_network_peering" "test2" {
  name                = "peer2to1"
  resource_group_name = "${azurerm_resource_group.test.name}"
  virtual_network_name = "${azurerm_virtual_network.test2.name}"
}
```

```

    remote_virtual_network_id = "${azurerm_virtual_network.test1.id}"
}

```

» Example Usage (Global virtual network peering)

```

variable "location" {
  default = [
    "uksouth",
    "southeastasia",
  ]
}

variable "vnet_address_space" {
  default = [
    "10.0.0.0/16",
    "10.1.0.0/16",
  ]
}

resource "azurerm_resource_group" "vnet" {
  count      = "${length(var.location)}"
  name       = "rg-global-vnet-peering-${count.index}"
  location   = "${element(var.location, count.index)}"
}

resource "azurerm_virtual_network" "vnet" {
  count            = "${length(var.location)}"
  name             = "vnet-${count.index}"
  resource_group_name = "${element(azurerm_resource_group.vnet.*.name, count.index)}"
  address_space     = ["${element(var.vnet_address_space, count.index)}"]
  location         = "${element(azurerm_resource_group.vnet.*.location, count.index)}"
}

resource "azurerm_subnet" "nva" {
  count            = "${length(var.location)}"
  name             = "nva"
  resource_group_name = "${element(azurerm_resource_group.vnet.*.name, count.index)}"
  virtual_network_name = "${element(azurerm_virtual_network.vnet.*.name, count.index)}"
  address_prefix     = "${cidrsubnet("${element(azurerm_virtual_network.vnet.*.address_space, count.index)}", 0, 1)}"
}

# enable global peering between the two virtual network
resource "azurerm_virtual_network_peering" "peering" {
  count      = "${length(var.location)}"
  name       = "peering-to-${element(azurerm_virtual_network.vnet.*.name, count.index)}"
}

```

```

resource_group_name      = "${element(azurerm_resource_group.vnet.*.name, count.index)}"
virtual_network_name     = "${element(azurerm_virtual_network.vnet.*.name, count.index)}"
remote_virtual_network_id = "${element(azurerm_virtual_network.vnet.*.id, 1 - count.index)}"
allow_virtual_network_access = true
allow_forwarded_traffic   = true

# `allow_gateway_transit` must be set to false for vnet Global Peering
allow_gateway_transit = false
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the virtual network peering. Changing this forces a new resource to be created.
- **virtual_network_name** - (Required) The name of the virtual network. Changing this forces a new resource to be created.
- **remote_virtual_network_id** - (Required) The full Azure resource ID of the remote virtual network. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the virtual network. Changing this forces a new resource to be created.
- **allow_virtual_network_access** - (Optional) Controls if the VMs in the remote virtual network can access VMs in the local virtual network. Defaults to false.
- **allow_forwarded_traffic** - (Optional) Controls if forwarded traffic from VMs in the remote virtual network is allowed. Defaults to false.
- **allow_gateway_transit** - (Optional) Controls gatewayLinks can be used in the remote virtual network's link to the local virtual network.
- **use_remote_gateways** - (Optional) Controls if remote gateways can be used on the local virtual network. If the flag is set to **true**, and **allow_gateway_transit** on the remote peering is also **true**, virtual network will use gateways of remote virtual network for transit. Only one peering can have this flag set to **true**. This flag cannot be set if virtual network already has a gateway. Defaults to **false**.

NOTE: **use_remote_gateways** must be set to **false** if using Global Virtual Network Peerings.

» Attributes Reference

The following attributes are exported:

- `id` - The Virtual Network Peering resource ID.

» Note

Virtual Network peerings cannot be created, updated or deleted concurrently.

» Import

Virtual Network Peerings can be imported using the `resource id`, e.g.

```
terraform import azurerm_virtual_network_peering.testPeering /subscriptions/00000000-0000-00
```

» `azurerm_management_group`

Manages a Management Group.

» Example Usage

```
data "azurerm_subscription" "current" {}

resource "azurerm_management_group" "example_parent" {
  display_name      = "ParentGroup"
  subscription_ids = [
    "${data.azurerm_subscription.current.id}",
  ]
}

resource "azurerm_management_group" "example_child" {
  display_name          = "ChildGroup"
  parent_management_group_id = "${azurerm_management_group.example_parent.id}"

  subscription_ids = [
    "${data.azurerm_subscription.current.id}",
    # other subscription IDs can go here
  ]
}
```


» Argument Reference

The following arguments are supported:

- **group_id** - (Optional) The UUID for this Management Group, which needs to be unique across your tenant - which will be generated if not provided. Changing this forces a new resource to be created.
- **display_name** - (Optional) A friendly name for this Management Group. If not specified, this'll be the same as the **group_id**.
- **parent_management_group_id** - (Optional) The ID of the Parent Management Group. Changing this forces a new resource to be created.
- **subscription_ids** - (Optional) A list of Subscription ID's which should be assigned to the Management Group.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Management Group.

» Import

Management Groups can be imported using the `management_group` resource id, e.g.

```
terraform import azurerm_management_group.test /providers/Microsoft.Management/ManagementGr
```

» azurerm_log_analytics_linked_service

Links a Log Analytics (formally Operational Insights) Workspace to another resource. The (currently) only linkable service is an Azure Automation Account.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "resourcegroup-01"
  location = "West Europe"
}

resource "azurerm_automation_account" "test" {
  name = "automation-01"
```

```

location          = "${azurerm_resource_group.test.location}"
resource_group_name = "${azurerm_resource_group.test.name}"

sku {
  name = "Basic"
}

tags = {
  environment = "development"
}
}

resource "azurerm_log_analytics_workspace" "test" {
  name          = "workspace-01"
  location      = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  sku           = "PerGB2018"
  retention_in_days = 30
}

resource "azurerm_log_analytics_linked_service" "test" {
  resource_group_name = "${azurerm_resource_group.test.name}"
  workspace_name      = "${azurerm_log_analytics_workspace.test.name}"
  resource_id         = "${azurerm_automation_account.test.id}"
}

```

» Argument Reference

The following arguments are supported:

- **resource_group_name** - (Required) The name of the resource group in which the Log Analytics Linked Service is created. Changing this forces a new resource to be created.
- **workspace_name** - (Required) Name of the Log Analytics Workspace that will contain the linkedServices resource. Changing this forces a new resource to be created.
- **linked_service_name** - (Optional) Name of the type of linkedServices resource to connect to the Log Analytics Workspace specified in **workspace_name**. Currently it defaults to and only supports **automation** as a value. Changing this forces a new resource to be created.
- **resource_id** - (Optional) The ID of the Resource that will be linked to the workspace.

- **linked_service_properties** - (Optional **Deprecated**) A `linked_service_properties` block as defined below.
- **tags** - (Optional) A mapping of tags to assign to the resource.

`linked_service_properties` supports the following:

- **resource_id** - (Optional **Deprecated**) The resource id of the resource that will be linked to the workspace. This field has been deprecated in favour of the top-level `resource_id` field and will be removed in v2.0 of the AzureRM Provider.

» Attributes Reference

The following attributes are exported:

- **id** - The Log Analytics Linked Service ID.
- **name** - The automatically generated name of the Linked Service. This cannot be specified. The format is always `<workspace_name>/<linked_service_name>` e.g. `workspace1/Automation`

» Import

Log Analytics Workspaces can be imported using the `resource id`, e.g.

```
terraform import azurerm_log_analytics_linked_service.test /subscriptions/00000000-0000-0000
```

» azurerm_log_analytics_solution

Manages a Log Analytics (formally Operational Insights) Solution.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "k8s-log-analytics-test"
  location = "westeurope"
}

resource "random_id" "workspace" {
  keepers = {
    # Generate a new id each time we switch to a new resource group
    group_name = "${azurerm_resource_group.test.name}"
  }
}
```

```

    }

    byte_length = 8
  }

  resource "azurerm_log_analytics_workspace" "test" {
    name                = "k8s-workspace-${random_id.workspace.hex}"
    location            = "${azurerm_resource_group.test.location}"
    resource_group_name = "${azurerm_resource_group.test.name}"
    sku                 = "PerGB2018"
  }

  resource "azurerm_log_analytics_solution" "test" {
    solution_name      = "ContainerInsights"
    location           = "${azurerm_resource_group.test.location}"
    resource_group_name = "${azurerm_resource_group.test.name}"
    workspace_resource_id = "${azurerm_log_analytics_workspace.test.id}"
    workspace_name     = "${azurerm_log_analytics_workspace.test.name}"

    plan {
      publisher = "Microsoft"
      product   = "OMSGallery/ContainerInsights"
    }
  }
}

```

» Argument Reference

The following arguments are supported:

- **solution_name** - (Required) Specifies the name of the solution to be deployed. See here for options. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the Log Analytics solution is created. Changing this forces a new resource to be created. Note: The solution and it's related workspace can only exist in the same resource group.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **workspace_resource_id** - (Required) The full resource ID of the Log Analytics workspace with which the solution will be linked. Changing this forces a new resource to be created.
- **workspace_name** - (Required) The full name of the Log Analytics workspace with which the solution will be linked. Changing this forces a

new resource to be created.

- **plan** - (Required) A **plan** block as documented below.
-

A **plan** block includes:

- **publisher** - (Required) The publisher of the solution. For example Microsoft. Changing this forces a new resource to be created.
- **product** - (Required) The product name of the solution. For example OMSGallery/Containers. Changing this forces a new resource to be created.
- **promotion_code** - (Optional) A promotion code to be used with the solution.

» Import

Log Analytics Solutions can be imported using the **resource id**, e.g.

```
terraform import azurerm_log_analytics_solution.solution1 /subscriptions/00000000-0000-0000-
```

» azurerm_log_analytics_workspace

Manages a Log Analytics (formally Operational Insights) Workspace.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "acctestRG-01"
  location  = "East US"
}

resource "azurerm_log_analytics_workspace" "test" {
  name                = "acctest-01"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  sku                 = "PerGB2018"
  retention_in_days    = 30
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Log Analytics Workspace. Workspace name should include 4-63 letters, digits or '-'. The '-' shouldn't be the first or the last symbol. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which the Log Analytics workspace is created. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **sku** - (Required) Specifies the Sku of the Log Analytics Workspace. Possible values are **Free**, **PerNode**, **Premium**, **Standard**, **Standalone**, **Unlimited**, and **PerGB2018** (new Sku as of 2018-04-03).

NOTE: A new pricing model took effect on 2018-04-03, which requires the SKU **PerGB2018**. If you're provisioned resources before this date you have the option of remaining with the previous Pricing SKU and using the other SKU's defined above. More information about the Pricing SKU's is available at the following URL.

- **retention_in_days** - (Optional) The workspace data retention in days. Possible values range between 30 and 730.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- **id** - The Log Analytics Workspace ID.
- **primary_shared_key** - The Primary shared key for the Log Analytics Workspace.
- **secondary_shared_key** - The Secondary shared key for the Log Analytics Workspace.
- **workspace_id** - The Workspace (or Customer) ID for the Log Analytics Workspace.
- **portal_url** - The Portal URL for the Log Analytics Workspace.

» Import

Log Analytics Workspaces can be imported using the `resource id`, e.g.

```
terraform import azurerm_log_analytics_workspace.workspace1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_log_analytics_workspace_linked_service`

Links a Log Analytics (formally Operational Insights) Workspace to another resource. The (currently) only linkable service is an Azure Automation Account.

NOTE: This resource has been deprecated in favour of the `azurerm_log_analytics_linked_service` resource and will be removed in the next major version of the AzureRM Provider. The new resource shares the same fields as this one, and information on migrating across can be found in this guide.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "resourcegroup-01"
  location = "West Europe"
}

resource "azurerm_automation_account" "test" {
  name                = "automation-01"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"

  sku {
    name = "Basic"
  }

  tags = {
    environment = "development"
  }
}

resource "azurerm_log_analytics_workspace" "test" {
  name                = "workspace-01"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  sku                 = "PerGB2018"
  retention_in_days   = 30
}
```

```
resource "azurerm_log_analytics_workspace_linked_service" "test" {
  resource_group_name = "${azurerm_resource_group.test.name}"
  workspace_name      = "${azurerm_log_analytics_workspace.test.name}"
  resource_id         = "${azurerm_automation_account.test.id}"
}
```

» Argument Reference

The following arguments are supported:

- **resource_group_name** - (Required) The name of the resource group in which the Log Analytics Linked Service is created. Changing this forces a new resource to be created.
- **workspace_name** - (Required) Name of the Log Analytics Workspace that will contain the linkedServices resource. Changing this forces a new resource to be created.
- **linked_service_name** - (Optional) Name of the type of linkedServices resource to connect to the Log Analytics Workspace specified in **workspace_name**. Currently it defaults to and only supports **automation** as a value. Changing this forces a new resource to be created.
- **resource_id** - (Optional) The ID of the Resource that will be linked to the workspace.
- **linked_service_properties** - (Optional **Deprecated**) A **linked_service_properties** block as defined below.
- **tags** - (Optional) A mapping of tags to assign to the resource.

linked_service_properties supports the following:

- **resource_id** - (Optional **Deprecated**) The resource id of the resource that will be linked to the workspace. This field has been deprecated in favour of the top-level **resource_id** field and will be removed in v2.0 of the AzureRM Provider.

» Attributes Reference

The following attributes are exported:

- **id** - The Log Analytics Linked Service ID.

- **name** - The automatically generated name of the Linked Service. This cannot be specified. The format is always <workspace_name>/<linked_service_name> e.g. workspace1/Automation

» Import

Log Analytics Workspaces can be imported using the **resource id**, e.g.

```
terraform import azurerm_log_analytics_workspace_linked_service.test /subscriptions/00000000
```

» azurerm_policy_assignment

Configures the specified Policy Definition at the specified Scope. Also, Policy Set Definitions are supported.

» Example Usage

```
resource "azurerm_policy_definition" "test" {
  name           = "my-policy-definition"
  policy_type    = "Custom"
  mode           = "All"
  display_name   = "my-policy-definition"

  policy_rule = <<POLICY_RULE
  {
    "if": {
      "not": {
        "field": "location",
        "in": "[parameters('allowedLocations')]"
      }
    },
    "then": {
      "effect": "audit"
    }
  }
}
POLICY_RULE

parameters = <<PARAMETERS
{
  "allowedLocations": {
    "type": "Array",
    "metadata": {
      "description": "The list of allowed locations for resources.",
```

```

        "displayName": "Allowed locations",
        "strongType": "location"
      }
    }
  }
PARAMETERS
}

resource "azurerm_resource_group" "test" {
  name      = "test-resources"
  location  = "West Europe"
}

resource "azurerm_policy_assignment" "test" {
  name                  = "example-policy-assignment"
  scope                 = "${azurerm_resource_group.test.id}"
  policy_definition_id = "${azurerm_policy_definition.test.id}"
  description           = "Policy Assignment created via an Acceptance Test"
  display_name          = "My Example Policy Assignment"

  parameters = <<PARAMETERS
  {
    "allowedLocations": {
      "value": [ "West Europe" ]
    }
  }
}
PARAMETERS
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Policy Assignment. Changing this forces a new resource to be created.
- **scope** - (Required) The Scope at which the Policy Assignment should be applied, which must be a Resource ID (such as Subscription e.g. /subscriptions/00000000-0000-0000-000000000000 or a Resource Group e.g. /subscriptions/00000000-0000-0000-000000000000/resourceGroups/myResourceGroup). Changing this forces a new resource to be created.
- **policy_definition_id** - (Required) The ID of the Policy Definition to be applied at the specified Scope.
- **identity** - (Optional) An identity block.

- **location** - (Optional) The Azure location where this policy assignment should exist. This is required when an Identity is assigned. Changing this forces a new resource to be created.
- **description** - (Optional) A description to use for this Policy Assignment. Changing this forces a new resource to be created.
- **display_name** - (Optional) A friendly display name to use for this Policy Assignment. Changing this forces a new resource to be created.
- **parameters** - (Optional) Parameters for the policy definition. This field is a JSON object that maps to the Parameters field from the Policy Definition. Changing this forces a new resource to be created.

NOTE: This value is required when the specified Policy Definition contains the **parameters** field.

- **not_scopes** - (Optional) A list of the Policy Assignment's excluded scopes. The list must contain Resource IDs (such as Subscriptions e.g. `/subscriptions/00000000-0000-0000-000000000000` or Resource Groups e.g. `/subscriptions/00000000-0000-0000-000000000000/resourceGroups/myResourceGroup`).

An **identity** block supports the following:

- **type** - (Required) The Managed Service Identity Type of this Policy Assignment. Possible values are **SystemAssigned** (where Azure will generate a Service Principal for you), or **None** (no use of a Managed Service Identity).

NOTE: When **type** is set to **SystemAssigned**, identity the Principal ID can be retrieved after the policy has been assigned.

» Attributes Reference

The following attributes are exported:

- **id** - The Policy Assignment id.
- **identity** - An **identity** block.

An **identity** block exports the following:

- **principal_id** - The Principal ID of this Policy Assignment if **type** is **SystemAssigned**.
- **tenant_id** - The Tenant ID of this Policy Assignment if **type** is **SystemAssigned**.

» Import

Policy Assignments can be imported using the `policy_name`, e.g.

```
terraform import azurerm_policy_assignment.assignment1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_policy_definition`

Manages a policy rule definition on a management group or your provider subscription.

Policy definitions do not take effect until they are assigned to a scope using a Policy Assignment.

» Example Usage

```
resource "azurerm_policy_definition" "policy" {
  name           = "accTestPolicy"
  policy_type    = "Custom"
  mode           = "Indexed"
  display_name   = "acceptance test policy definition"

  policy_rule = <<POLICY_RULE
  {
    "if": {
      "not": {
        "field": "location",
        "in": "[parameters('allowedLocations')]"
      }
    },
    "then": {
      "effect": "audit"
    }
  }
}
POLICY_RULE

parameters = <<PARAMETERS
{
  "allowedLocations": {
    "type": "Array",
    "metadata": {
```

```

        "description": "The list of allowed locations for resources.",
        "displayName": "Allowed locations",
        "strongType": "location"
    }
}
}
PARAMETERS
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the policy definition. Changing this forces a new resource to be created.
- **policy_type** - (Required) The policy type. The value can be "BuiltIn", "Custom" or "NotSpecified". Changing this forces a new resource to be created.
- **mode** - (Required) The policy mode that allows you to specify which resource types will be evaluated. The value can be "All", "Indexed" or "NotSpecified". Changing this resource forces a new resource to be created.
- **display_name** - (Required) The display name of the policy definition.
- **description** - (Optional) The description of the policy definition.
- **management_group_id** - (Optional) The ID of the Management Group where this policy should be defined. Changing this forces a new resource to be created.

Note: if you are using `azurerm_management_group` to assign a value to `management_group_id`, be sure to use `.group_id` and not `.id`.

- **policy_rule** - (Optional) The policy rule for the policy definition. This is a json object representing the rule that contains an if and a then block.
- **metadata** - (Optional) The metadata for the policy definition. This is a json object representing additional metadata that should be stored with the policy definition.
- **parameters** - (Optional) Parameters for the policy definition. This field is a json object that allows you to parameterize your policy definition.

» Attributes Reference

The following attributes are exported:

- id - The policy definition id.

» Import

Policy Definitions can be imported using the `policy_name`, e.g.

```
terraform import azurerm_policy_definition.testPolicy /subscriptions/<SUBSCRIPTION_ID>/providers/Microsoft.Management/managementGroups/<MANAGEMENT_GROUP_ID>
```

or

```
terraform import azurerm_policy_definition.testPolicy /providers/Microsoft.Management/managementGroups/<MANAGEMENT_GROUP_ID>
```

» azurerm_policy_set_definition

Manages a policy set definition.

NOTE: Policy set definitions (also known as policy initiatives) do not take effect until they are assigned to a scope using a Policy Set Assignment.

» Example Usage

```
resource "azurerm_policy_set_definition" "test" {
  name           = "testPolicySet"
  policy_type    = "Custom"
  display_name   = "Test Policy Set"

  parameters = <<PARAMETERS
  {
    "allowedLocations": {
      "type": "Array",
      "metadata": {
        "description": "The list of allowed locations for resources.",
        "displayName": "Allowed locations",
        "strongType": "location"
      }
    }
  }
}

PARAMETERS

policy_definitions = <<POLICY_DEFINITIONS
[
  {
    "parameters": {
      "listOfAllowedLocations": {
```

```

        "value": "[parameters('allowedLocations')]"
      },
    },
    "policyDefinitionId": "/providers/Microsoft.Authorization/policyDefinitions/e768
  }
]
POLICY_DEFINITIONS
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the policy set definition. Changing this forces a new resource to be created.
- **policy_type** - (Required) The policy set type. Possible values are `BuiltIn` or `Custom`. Changing this forces a new resource to be created.
- **display_name** - (Required) The display name of the policy set definition.
- **policy_definitions** - (Required) The policy definitions for the policy set definition. This is a json object representing the bundled policy definitions .
- **description** - (Optional) The description of the policy set definition.
- **management_group_id** - (Optional) The ID of the Management Group where this policy should be defined. Changing this forces a new resource to be created.

Note: if you are using `azurerm_management_group` to assign a value to `management_group_id`, be sure to use `.group_id` and not `.id`.

- **metadata** - (Optional) The metadata for the policy set definition. This is a json object representing additional metadata that should be stored with the policy definition.
- **parameters** - (Optional) Parameters for the policy set definition. This field is a json object that allows you to parameterize your policy definition.

» Attributes Reference

The following attributes are exported:

- **id** - The policy set definition id.

» Import

Policy Set Definitions can be imported using the Resource ID, e.g.

```
terraform import azurerm_policy_set_definition.test /subscriptions/00000000-0000-0000-0000-000000000000  
or  
shell terraform import azurerm_policy_set_definition.test  
/providers/Microsoft.Management/managementgroups/my-mgmt-group-id/providers/Microsoft.Authorization/policyDefinitions/test
```

» azurerm__recovery__services__protection__policy__vm

Manages an Recovery Services VM Protection Policy.

» Example Usage

```
resource "azurerm_resource_group" "example" {  
  name      = "tfex-recovery_vault"  
  location = "West US"  
}  
  
resource "azurerm_recovery_services_vault" "example" {  
  name                = "tfex-recovery-vault"  
  location             = "${azurerm_resource_group.example.location}"  
  resource_group_name = "${azurerm_resource_group.example.name}"  
  sku                 = "Standard"  
}  
  
resource "azurerm_recovery_services_protection_policy_vm" "test" {  
  name                = "tfex-recovery-vault-policy"  
  resource_group_name = "${azurerm_resource_group.example.name}"  
  recovery_vault_name = "${azurerm_recovery_services_vault.example.name}"  
  
  timezone = "UTC"  
  
  backup {  
    frequency = "Daily"  
    time      = "23:00"  
  }  
  
  retention_daily {  
    count = 10  
  }  
  
  retention_weekly {  
    count = 10  
  }  
}
```



```

    count      = 42
    weekdays = ["Sunday", "Wednesday", "Friday", "Saturday"]
}

retention_monthly {
    count      = 7
    weekdays = ["Sunday", "Wednesday"]
    weeks      = ["First", "Last"]
}

retention_yearly {
    count      = 77
    weekdays = ["Sunday"]
    weeks      = ["Last"]
    months     = ["January"]
}
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Recovery Services Vault Policy. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the Recovery Services Protected VM. Changing this forces a new resource to be created.
- **recovery_vault_name** - (Required) Specifies the name of the Recovery Services Vault to use. Changing this forces a new resource to be created.
- **backup** - (Required) Configures the Policy backup frequency, times & days as documented in the **backup** block below.
- **timezone** - (Optional) Specifies the timezone. Defaults to UTC
- **retention_daily** - (Optional) Configures the policy daily retention as documented in the **retention_daily** block below. Required when backup frequency is **Daily**.
- **retention_weekly** - (Optional) Configures the policy weekly retention as documented in the **retention_weekly** block below. Required when backup frequency is **Weekly**.
- **retention_monthly** - (Optional) Configures the policy monthly retention as documented in the **retention_monthly** block below.

- **retention_yearly** - (Optional) Configures the policy yearly retention as documented in the **retention_yearly** block below.
 - **tags** - (Optional) A mapping of tags to assign to the resource.
-

The **backup** block supports:

- **frequency** - (Required) Sets the backup frequency. Must be either **Daily** or **Weekly**.
 - **times** - (Required) The time of day to perform the backup in 24hour format.
 - **weekdays** - (Optional) The days of the week to perform backups on. Must be one of **Sunday**, **Monday**, **Tuesday**, **Wednesday**, **Thursday**, **Friday** or **Saturday**.
-

The **retention_daily** block supports:

- **count** - (Required) The number of daily backups to keep. Must be between 1 and 9999
-

The **retention_weekly** block supports:

- **count** - (Required) The number of weekly backups to keep. Must be between 1 and 9999
 - **weekdays** - (Required) The weekday backups to retain. Must be one of **Sunday**, **Monday**, **Tuesday**, **Wednesday**, **Thursday**, **Friday** or **Saturday**.
-

The **retention_monthly** block supports:

- **count** - (Required) The number of monthly backups to keep. Must be between 1 and 9999
 - **weekdays** - (Required) The weekday backups to retain. Must be one of **Sunday**, **Monday**, **Tuesday**, **Wednesday**, **Thursday**, **Friday** or **Saturday**.
 - **weeks** - (Required) The weeks of the month to retain backups of. Must be one of **First**, **Second**, **Third**, **Fourth**, **Last**.
-

The **retention_yearly** block supports:

- **count** - (Required) The number of yearly backups to keep. Must be between 1 and 9999

- **weekdays** - (Required) The weekday backups to retain . Must be one of Sunday, Monday, Tuesday, Wednesday, Thursday, Friday or Saturday.
- **weeks** - (Required) The weeks of the month to retain backups of. Must be one of **First**, **Second**, **Third**, **Fourth**, **Last**.
- **months** - (Required) The months of the year to retain backups of. Must be one of **January**, **February**, **March**, **April**, **May**, **June**, **July**, **August**, **September**, **October**, **November** and **December**.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Recovery Services VM Protection Policy.

» Import

Recovery Services VM Protection Policy can be imported using the **resource id**, e.g.

```
terraform import azurerm_recovery_services_protection_policy_vm.policy1 /subscriptions/00000
```

» azurerm_recovery_services_protected_vm

Manages an Recovery Protected VM.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "tfex-recovery_vault"
  location = "West US"
}

resource "azurerm_recovery_services_vault" "example" {
  name                = "tfex-recovery-vault"
  location             = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
  sku                  = "Standard"
}

resource "azurerm_recovery_services_protection_policy_vm" "example" {
```

```

name                = "tfex-recovery-vault-policy"
resource_group_name = "${azurerm_resource_group.example.name}"
recovery_vault_name = "${azurerm_recovery_services_vault.example.name}"

backup {
  frequency = "Daily"
  time      = "23:00"
}
}

resource "azurerm_recovery_services_protected_vm" "vm1" {
  resource_group_name = "${azurerm_resource_group.example.name}"
  recovery_vault_name = "${azurerm_recovery_services_vault.example.name}"
  source_vm_id        = "${azurerm_virtual_machine.example.id}"
  backup_policy_id    = "${azurerm_recovery_services_protection_policy_vm.example.id}"
}

```

» Argument Reference

The following arguments are supported:

- **resource_group_name** - (Required) The name of the resource group in which to create the Recovery Services Protected VM. Changing this forces a new resource to be created.
- **recovery_vault_name** - (Required) Specifies the name of the Recovery Services Vault to use. Changing this forces a new resource to be created.
- **source_vm_id** - (Required) Specifies the ID of the VM to backup. Changing this forces a new resource to be created.
- **backup_policy_id** - (Required) Specifies the id of the backup policy to use. Changing this forces a new resource to be created.
- **tags** - (Optional) A mapping of tags to assign to the resource.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Recovery Services Vault.

» Import

Recovery Services Protected VMs can be imported using the **resource id**, e.g.

```
terraform import azurerm_recovery_services_protected_vm.item1 "/subscriptions/00000000-0000-
```

Note the ID requires quoting as there are semicolons

» **azurerm__recovery__services__vault**

Manage an Recovery Services Vault.

» **Example Usage**

```
resource "azurerm_resource_group" "rg" {
  name      = "tfex-recovery_vault"
  location  = "West US"
}

resource "azurerm_recovery_services_vault" "vault" {
  name                = "example_recovery_vault"
  location             = "${azurerm_resource_group.rg.location}"
  resource_group_name = "${azurerm_resource_group.rg.name}"
  sku                 = "Standard"
}
```

» **Argument Reference**

The following arguments are supported:

- **name** - (Required) Specifies the name of the Recovery Services Vault. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the Recovery Services Vault. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **tags** - (Optional) A mapping of tags to assign to the resource.
- **sku** - (Required) Sets the vault's SKU. Possible values include: **Standard**, **RS0**.

» **Attributes Reference**

The following attributes are exported:

- id - The ID of the Recovery Services Vault.

» Import

Recovery Services Vaults can be imported using the `resource id`, e.g.

```
terraform import azure_rm_recovery_services_vault.vault1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azure_rm_redis_cache

Manages a Redis Cache.

» Example Usage (Basic)

```
resource "azure_rm_resource_group" "test" {
  name      = "redis-resources"
  location  = "West US"
}

# NOTE: the Name used for Redis needs to be globally unique
resource "azure_rm_redis_cache" "test" {
  name                = "tf-redis-basic"
  location             = "${azure_rm_resource_group.test.location}"
  resource_group_name = "${azure_rm_resource_group.test.name}"
  capacity            = 0
  family              = "C"
  sku_name            = "Basic"
  enable_non_ssl_port = false
}
```

» Example Usage (Standard)

```
resource "azure_rm_resource_group" "test" {
  name      = "redis-resources"
  location  = "West US"
}

# NOTE: the Name used for Redis needs to be globally unique
resource "azure_rm_redis_cache" "test" {
  name                = "tf-redis-standard"
  location             = "${azure_rm_resource_group.test.location}"
  resource_group_name = "${azure_rm_resource_group.test.name}"
}
```

```

    capacity          = 2
    family             = "C"
    sku_name           = "Standard"
    enable_non_ssl_port = false
    minimum_tls_version = "1.2"
}

```

» Example Usage (Premium with Clustering)

```

resource "azurerm_resource_group" "test" {
  name     = "redis-resources"
  location = "West US"
}

# NOTE: the Name used for Redis needs to be globally unique
resource "azurerm_redis_cache" "test" {
  name                = "tf-redis-premium"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  capacity            = 1
  family              = "P"
  sku_name            = "Premium"
  enable_non_ssl_port = false
  shard_count         = 3

  redis_configuration {
    maxmemory_reserved = 2
    maxmemory_delta    = 2
    maxmemory_policy   = "allkeys-lru"
  }
}

```

» Example Usage (Premium with Backup)

```

resource "azurerm_resource_group" "test" {
  name     = "redis-resources"
  location = "West US"
}

resource "azurerm_storage_account" "test" {
  name                = "redissa"
  resource_group_name = "${azurerm_resource_group.test.name}"
  location            = "${azurerm_resource_group.test.location}"
  account_tier        = "Standard"
}

```

```

    account_replication_type = "GRS"
}

# NOTE: the Name used for Redis needs to be globally unique
resource "azurerm_redis_cache" "test" {
  name                = "tf-redis-pbkup"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  capacity            = 3
  family              = "P"
  sku_name            = "Premium"
  enable_non_ssl_port = false

  redis_configuration {
    rdb_backup_enabled      = true
    rdb_backup_frequency    = 60
    rdb_backup_max_snapshot_count = 1
    rdb_storage_connection_string = "DefaultEndpointsProtocol=https;BlobEndpoint=${azurerm_s
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Redis instance. Changing this forces a new resource to be created.
- **location** - (Required) The location of the resource group.
- **resource_group_name** - (Required) The name of the resource group in which to create the Redis instance.
- **capacity** - (Required) The size of the Redis cache to deploy. Valid values for a SKU family of C (Basic/Standard) are 0, 1, 2, 3, 4, 5, 6, and for P (Premium) family are 1, 2, 3, 4.
- **family** - (Required) The SKU family to use. Valid values are C and P, where C = Basic/Standard, P = Premium.

The pricing group for the Redis Family - either "C" or "P" at present.

- **sku_name** - (Required) The SKU of Redis to use - can be either Basic, Standard or Premium.
- **enable_non_ssl_port** - (Optional) Enable the non-SSL port (6789) - disabled by default.

- `minimum_tls_version` - (Optional) The minimum TLS version. Defaults to 1.0.
- `patch_schedule` - (Optional) A list of `patch_schedule` blocks as defined below - only available for Premium SKU's.
- `private_static_ip_address` - (Optional) The Static IP Address to assign to the Redis Cache when hosted inside the Virtual Network. Changing this forces a new resource to be created.
- `redis_configuration` - (Required) A `redis_configuration` as defined below - with some limitations by SKU - defaults/details are shown below.
- `shard_count` - (Optional) *Only available when using the Premium SKU* The number of Shards to create on the Redis Cluster.
- `subnet_id` - (Optional) The ID of the Subnet within which the Redis Cache should be deployed. Changing this forces a new resource to be created.
- `tags` - (Optional) A mapping of tags to assign to the resource.
- `zones` - (Optional) A list of a single item of the Availability Zone which the Redis Cache should be allocated in.

» -> **Please Note: Availability Zones are in Preview and only supported in several regions at this time - as such you must be opted into the Preview to use this functionality. You can opt into the Availability Zones Preview in the Azure Portal.**

- `redis_configuration` supports the following:
- `maxmemory_reserved` - (Optional) Value in megabytes reserved for non-cache usage e.g. failover. Defaults are shown below.
- `maxmemory_delta` - (Optional) The max-memory delta for this Redis instance. Defaults are shown below.
- `maxmemory_policy` - (Optional) How Redis will select what to remove when `maxmemory` is reached. Defaults are shown below.
- `maxfragmentationmemory_reserved` - (Optional) Value in megabytes reserved to accommodate for memory fragmentation. Defaults are shown below.
- `rdb_backup_enabled` - (Optional) Is Backup Enabled? Only supported on Premium SKU's.

- `rdb_backup_frequency` - (Optional) The Backup Frequency in Minutes. Only supported on Premium SKU's. Possible values are: 15, 30, 60, 360, 720 and 1440.
- `rdb_backup_max_snapshot_count` - (Optional) The maximum number of snapshots to create as a backup. Only supported for Premium SKU's.
- `rdb_storage_connection_string` - (Optional) The Connection String to the Storage Account. Only supported for Premium SKU's. In the format: `DefaultEndpointsProtocol=https;BlobEndpoint=${azurerm_storage_account.test.primary_blob`

NOTE: There's a bug in the Redis API where the original storage connection string isn't being returned, which is being tracked in this issue. In the interim you can use the `ignore_changes` attribute to ignore changes to this field e.g.:

```
resource "azurerm_redis_cache" "test" {
  # ...
  ignore_changes = ["redis_configuration.0.rdb_storage_connection_string"]
}
```

- `notify_keyspace_events` - (Optional) Keyspace notifications allows clients to subscribe to Pub/Sub channels in order to receive events affecting the Redis data set in some way. Reference

```
redis_configuration {
  maxmemory_reserved = 10
  maxmemory_delta    = 2
  maxmemory_policy    = "allkeys-lru"
}
```

» Default Redis Configuration Values

Redis Value	Basic	Standard	Premium
<code>maxmemory_reserved</code>	2	50	200
<code>maxfragmentationmemory_reserved</code>	2	50	200
<code>maxmemory_delta</code>	2	50	200
<code>maxmemory_policy</code>	volatile-lru	volatile-lru	volatile-lru

Important: *The `maxmemory_reserved`, `maxmemory_delta` and `maxfragmentationmemory_reserved` settings are only available for Standard and Premium caches. More details are available in the Relevant Links section below.*

- `patch_schedule` supports the following:
- `day_of_week` (Required) the Weekday name - possible values include Monday, Tuesday, Wednesday etc.

- `start_hour_utc` - (Optional) the Start Hour for maintenance in UTC - possible values range from 0 - 23.

Note: The Patch Window lasts for 5 hours from the `start_hour_utc`.

» Attributes Reference

The following attributes are exported:

- `id` - The Route ID.
- `hostname` - The Hostname of the Redis Instance
- `ssl_port` - The SSL Port of the Redis Instance
- `port` - The non-SSL Port of the Redis Instance
- `primary_access_key` - The Primary Access Key for the Redis Instance
- `secondary_access_key` - The Secondary Access Key for the Redis Instance
- `redis_configuration` - A `redis_configuration` block as defined below:

A `redis_configuration` block exports the following:

- `maxclients` - Returns the max number of connected clients at the same time.

» Relevant Links

- Azure Redis Cache: SKU specific configuration limitations
- Redis: Available Configuration Settings

» Import

Redis Cache's can be imported using the `resource id`, e.g.

```
terraform import azurerm_redis_cache.cache1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_redis_firewall_rule`

Manages a Firewall Rule associated with a Redis Cache.

» Example Usage

```
resource "random_id" "server" {
  keepers = {
    azi_id = 1
  }

  byte_length = 8
}

resource "azurerm_resource_group" "test" {
  name     = "redis-resourcegroup"
  location = "West Europe"
}

resource "azurerm_redis_cache" "test" {
  name                = "redis${random_id.server.hex}"
  location             = "${azurerm_resource_group.test.location}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  capacity            = 1
  family              = "P"
  sku_name            = "Premium"
  enable_non_ssl_port = false

  redis_configuration {
    maxclients      = 256
    maxmemory_reserved = 2
    maxmemory_delta  = 2
    maxmemory_policy  = "allkeys-lru"
  }
}

resource "azurerm_redis_firewall_rule" "test" {
  name                = "someIPrange"
  redis_cache_name    = "${azurerm_redis_cache.test.name}"
  resource_group_name = "${azurerm_resource_group.test.name}"
  start_ip            = "1.2.3.4"
  end_ip              = "2.3.4.5"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Firewall Rule. Changing this forces a

new resource to be created.

- **redis_cache_name** - (Required) The name of the Redis Cache. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which this Redis Cache exists.
- **start_ip** - (Required) The lowest IP address included in the range
- **end_ip** - (Required) The highest IP address included in the range.

» Attributes Reference

The following attributes are exported:

- **id** - The Redis Firewall Rule ID.

» Import

Redis Firewall Rules can be imported using the **resource id**, e.g.

```
terraform import azurerm_redis_firewall_rule.rule1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_search_service

Allows you to manage an Azure Search Service.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "acceptanceTestResourceGroup1"
  location = "West US"
}

resource "azurerm_search_service" "test" {
  name                        = "acceptanceTestSearchService1"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  location                   = "${azurerm_resource_group.test.location}"
  sku                        = "standard"

  tags = {
    environment = "staging"
    database    = "test"
  }
}
```

```
}  
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Search Service. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the Search Service. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **sku** - (Required) Valid values are **free** and **standard**. **standard2** and **standard3** are also valid, but can only be used when it's enabled on the backend by Microsoft support. **free** provisions the service in shared clusters. **standard** provisions the service in dedicated clusters. Changing this forces a new resource to be created.
- **replica_count** - (Optional) Default is 1. Valid values include 1 through 12. Valid only when **sku** is **standard**. Changing this forces a new resource to be created.
- **partition_count** - (Optional) Default is 1. Valid values include 1, 2, 3, 4, 6, or 12. Valid only when **sku** is **standard**. Changing this forces a new resource to be created.
- **tags** - (Optional) A mapping of tags to assign to the resource. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported:

- **id** - The Search Service ID.
- **primary_key** - The Search Service Administration primary key.
- **secondary_key** - The Search Service Administration secondary key.

» Import

Search Services can be imported using the **resource id**, e.g.

```
terraform import azurerm_search_service.service1 /subscriptions/00000000-0000-0000-0000-0000
```

» **azurerm__security__center__contact**

Manages the subscription's Security Center Contact.

NOTE: Owner access permission is required.

» **Example Usage**

```
resource "azurerm_security_center_contact" "example" {  
  email = "contact@example.com"  
  phone = "+1-555-555-5555"  
  
  alert_notifications = true  
  alerts_to_admins    = true  
}
```

» **Argument Reference**

The following arguments are supported:

- **email** - (Required) The email of the Security Center Contact.
- **phone** - (Required) The phone number of the Security Center Contact.
- **alert_notifications** - (Required) Whether to send security alerts notifications to the security contact.
- **alerts_to_admins** - (Required) Whether to send security alerts notifications to subscription admins.

» **Attributes Reference**

The following attributes are exported:

- **id** - The Security Center Contact ID.

» **Import**

The contact can be imported using the **resource id**, e.g.

```
terraform import azurerm_security_center_contact.example /subscriptions/00000000-0000-0000-0000-0000
```

» **azurerm_security_center_subscription_pricing**

Manages the Pricing Tier for Azure Security Center in the current subscription.

NOTE: This resource requires the **Owner** permission on the Subscription.

NOTE: Deletion of this resource does not change or reset the pricing tier to **Free**

» **Example Usage**

```
resource "azurerm_security_center_subscription_pricing" "example" {  
  tier = "Standard"  
}
```

» **Argument Reference**

The following arguments are supported:

- **tier** - (Required) The pricing tier to use. Possible values are **Free** and **Standard**.

NOTE: Changing the pricing tier to **Standard** affects all resources in the subscription and could be quite costly.

» **Attributes Reference**

The following attributes are exported:

- **id** - The subscription pricing ID.

» **Import**

The pricing tier can be imported using the **resource id**, e.g.

```
terraform import azurerm_security_center_subscription_pricing.example /subscriptions/00000000-0000-0000-0000-000000000000
```

» **azurerm_security_center_workspace**

Manages the subscription's Security Center Workspace.

NOTE: Owner access permission is required.

NOTE: The subscription's pricing model can not be **Free** for this to have any affect.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name      = "tfex-security-workspace"
  location  = "westus"
}

resource "azurerm_log_analytics_workspace" "example" {
  name                = "tfex-security-workspace"
  location             = "${azurerm_resource_group.example.location}"
  resource_group_name = "${azurerm_resource_group.example.name}"
  sku                 = "PerGB2018"
}

resource "azurerm_security_center_workspace" "example" {
  scope          = "/subscriptions/00000000-0000-0000-0000-000000000000"
  workspace_id   = "${azurerm_log_analytics_workspace.example.id}"
}
```

» Argument Reference

The following arguments are supported:

- **scope** - (Required) The scope of VMs to send their security data to the desired workspace, unless overridden by a setting with more specific scope.
- **workspace_id** - (Required) The ID of the Log Analytics Workspace to save the data in.

» Attributes Reference

The following attributes are exported:

- **id** - The Security Center Workspace ID.

» Import

The contact can be imported using the **resource id**, e.g.

```
terraform import azurerm_security_center_workspace.example /subscriptions/00000000-0000-0000-0000-000000000000
```

» **azurerm_scheduler_job_collection**

Manages a Scheduler Job Collection.

NOTE: Support for Scheduler Job Collections has been deprecated by Microsoft in favour of Logic Apps (more information can be found at [this link](#)) - as such we plan to remove support for this resource as a part of version 2.0 of the AzureRM Provider.

» **Example Usage**

```
resource "azurerm_resource_group" "rg" {
  name      = "tfex-job_collection"
  location  = "West US"
}

resource "azurerm_scheduler_job_collection" "jobs" {
  name                  = "example_job_collection"
  location              = "${azurerm_resource_group.rg.location}"
  resource_group_name   = "${azurerm_resource_group.rg.name}"
  sku                  = "free"
  state                = "enabled"

  quota {
    max_job_count          = 5
    max_recurrence_interval = 24
    max_recurrence_frequency = "hour"
  }
}
```

» **Argument Reference**

The following arguments are supported:

- **name** - (Required) Specifies the name of the Scheduler Job Collection. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the Scheduler Job Collection. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **tags** - (Optional) A mapping of tags to assign to the resource.

- **sku** - (Required) Sets the Job Collection's pricing level's SKU. Possible values include: **Standard**, **Free**, **P10Premium**, **P20Premium**.
- **state** - (Optional) Sets Job Collection's state. Possible values include: **Enabled**, **Disabled**, **Suspended**.
- **quota** - (Optional) Configures the Job collection quotas as documented in the **quota** block below.

The **quota** block supports:

- **max_job_count** - (Optional) Sets the maximum number of jobs in the collection.
- **max_recurrence_frequency** - (Required) The maximum frequency of recurrence. Possible values include: **Minute**, **Hour**, **Day**, **Week**, **Month**
- **max_recurrence_interval** - (Optional) The maximum interval between recurrence.

» Attributes Reference

The following attributes are exported:

- **id** - The ID of the Scheduler Job Collection.

» Import

Scheduler Job Collections can be imported using the **resource id**, e.g.

```
terraform import azurerm_scheduler_job_collection.jobcollection1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_scheduler_job

Manages a Scheduler Job.

NOTE: Support for Scheduler Job has been deprecated by Microsoft in favour of Logic Apps (more information can be found at [this link](#)) - as such we plan to remove support for this resource as a part of version 2.0 of the AzureRM Provider.

» Example Usage (single web get now)

```
resource "azurerm_scheduler_job" "web-once-now" {
  name                = "tfex-web-once-now"
  resource_group_name = "${azurerm_resource_group.example.name}"
}
```

```

job_collection_name = "${azurerm_scheduler_job_collection.example.name}"

# re-enable it each run
state = "enabled"

action_web {
    # defaults to get
    url = "http://this.url.fails"
}

# default start time is now
}

```

» Example Usage (recurring daily with retry and basic authentication)

```

resource "azurerm_scheduler_job" "web-recurring-daily" {
    name = "tfex-web-recurring-daily"
    resource_group_name = "${azurerm_resource_group.example.name}"
    job_collection_name = "${azurerm_scheduler_job_collection.example.name}"

    action_web {
        url      = "https://this.url.fails"
        method   = "put"
        body     = "this is some text"

        headers = {
            Content-Type = "text"
        }

        authentication_basic {
            username = "login"
            password = "apassword"
        }
    }

    retry {
        # retry every 5 min a maximum of 10 times
        interval = "00:05:00"
        count    = 10
    }

    recurrence {
        frequency = "day"
        count     = 1000
    }
}

```

```

    # run 4 times an hour every 12 hours
    hours    = [0, 12]
    minutes  = [0, 15, 30, 45]
  }

  start_time = "2018-07-07T07:07:07-07:00"
}

```

» Example Usage (recurring monthly with an error action and client certificate authentication)

```

resource "azurerm_scheduler_job" "web-recurring-daily" {
  name                        = "tfex-web-recurring-daily"
  resource_group_name        = "${azurerm_resource_group.example.name}"
  job_collection_name        = "${azurerm_scheduler_job_collection.example.name}"

  action_web {
    url = "https://this.url.fails"

    authentication_certificate {
      pfx      = "${base64encode(file("your_cert.pfx"))}"
      password = "cert_password"
    }
  }

  error_action_web {
    url      = "https://this.url.fails"
    method   = "put"
    body     = "The job failed"

    headers = {
      "Content-Type" = "text"
    }

    authentication_basic {
      username = "login"
      password = "apassword"
    }
  }

  recurrence {
    frequency = "monthly"
    count     = 1000
  }
}

```

```

monthly_occurrences = [
  {
    # first Sunday
    day      = "Sunday"
    occurrence = 1
  },
  {
    # third Sunday
    day      = "Sunday"
    occurrence = 3
  },
  {
    # last Sunday
    day      = "Sunday"
    occurrence = -1
  },
]
}

start_time = "2018-07-07T07:07:07-07:00"
}

```

» Example Usage (storage queue action)

```

resource "azurerm_storage_account" "example" {
  name                        = "tfexstorageaccount"
  resource_group_name        = "${azurerm_resource_group.example.name}"
  location                   = "${azurerm_resource_group.example.location}"
  account_tier               = "Standard"
  account_replication_type   = "LRS"
}

resource "azurerm_storage_queue" "example" {
  name                        = "tfex-schedulerjob-storagequeue"
  resource_group_name        = "${azurerm_resource_group.example.name}"
  storage_account_name       = "${azurerm_storage_account.example.name}"
}

resource "azurerm_scheduler_job" "storage-once-now" {
  name                        = "tfex-storage-once-now"
  resource_group_name        = "${azurerm_resource_group.example.name}"
  job_collection_name        = "${azurerm_scheduler_job_collection.example.name}"

  action_storage_queue {
    storage_account_name = "${azurerm_storage_account.example.name}"
  }
}

```

```

        storage_queue_name    = "${azurerm_storage_queue.example.name}"
        sas_token              = "${azurerm_storage_account.example.primary_access_key}"
        message                = "storage message"
    }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Scheduler Job. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the Scheduler Job. Changing this forces a new resource to be created.
- **job_collection_name** - (Required) Specifies the name of the Scheduler Job Collection in which the Job should exist. Changing this forces a new resource to be created.
- **action_web** - (Optional) A **action_web** block defining the job action as described below. Note this is identical to an **error_action_web** block.

NOTE At least one of **error_action_web** or **action_storage_queue** needs to be set.

- **action_storage_queue** - (Optional) A **action_storage_queue** block defining a storage queue job action as described below. Note this is identical to an **error_action_storage_queue** block.
- **error_action_web** - (Optional) A **error_action_web** block defining the action to take on an error as described below. Note this is identical to an **action_web** block.
- **error_action_storage_queue** - (Optional) A **error_action_storage_queue** block defining the a web action to take on an error as described below. Note this is identical to an **action_storage_queue** block.
- **retry** - (Optional) A **retry** block defining how to retry as described below.
- **recurrence** - (Optional) A **recurrence** block defining a job occurrence schedule.
- **start_time** - (Optional) The time the first instance of the job is to start running at.
- **state** - (Optional) The sets or gets the current state of the job. Can be set to either **Enabled** or **Completed**

web_action & **error_web_action** block supports the following:

- **url** - (Required) Specifies the URL of the web request. Must be HTTPS for authenticated requests.
- **method** - (Optional) Specifies the method of the request. Defaults to **Get** and must be one of **Get**, **Put**, **Post**, **Delete**.
- **body** - (Optional) Specifies the request body.
- **headers** - (Optional) A map specifying the headers sent with the request.
- **authentication_basic** - (Optional) An **authentication_active_directory** block which defines the Active Directory oauth configuration to use.
- **authentication_certificate** - (Optional) An **authentication_certificate** block which defines the client certificate information to be use.
- **authentication_active_directory** - (Optional) An **authentication_active_directory** block which defines the OAUTH Active Directory information to use.

authentication_basic block supports the following:

- **username** - (Required) Specifies the username to use.
- **password** - (Required) Specifies the password to use.

authentication_certificate block supports the following:

- **pfx** - (Required) Specifies the pfx certificate in base-64 format.
- **password** - (Required) Specifies the certificate password.

authentication_active_directory block supports the following:

- **client_id** - (Required) Specifies the client ID to use.
- **tenant_id** - (Required) Specifies the tenant ID to use.
- **client_secret** - (Required) Specifies the secret to use.
- **audience** - (Optional) Specifies the audience.

action_storage_queue & **error_action_storage_queue** block supports the following:

- **storage_account_name** - (Required) Specifies the the storage account name.
- **storage_queue_name** - (Required) Specifies the the storage account queue.
- **sas_token** - (Required) Specifies a SAS token/key to authenticate with.
- **message** - (Required) The message to send into the queue.

retry block supports the following:

- **interval** - (Required) Specifies the duration between retries.
- **count** - (Required) Specifies the number of times a retry should be attempted.

recurrence block supports the following:

- **frequency** - (Required) Specifies the frequency of recurrence. Must be one of **Minute**, **Hour**, **Day**, **Week**, **Month**.
- **interval** - (Optional) Specifies the interval between executions. Defaults to 1.

- **count** - (Optional) Specifies the maximum number of times that the job should run.
- **end_time** - (Optional) Specifies the time at which the job will cease running. Must be less than 500 days into the future.
- **minutes** - (Optional) Specifies the minutes of the hour that the job should execute at. Must be between 0 and 59
- **hours** - (Optional) Specifies the hours of the day that the job should execute at. Must be between 0 and 23
- **week_days** - (Optional) Specifies the days of the week that the job should execute on. Must be one of **Monday**, **Tuesday**, **Wednesday**, **Thursday**, **Friday**, **Saturday**, **Sunday**. Only applies when **Week** is used for frequency.
- **month_days** - (Optional) Specifies the days of the month that the job should execute on. Must be non zero and between -1 and 31. Only applies when **Month** is used for frequency.
- **monthly_occurrences** - (Optional) Specifies specific monthly occurrences like "last sunday of the month" with **monthly_occurrences** blocks. Only applies when **Month** is used for frequency.

monthly_occurrences block supports the following:

- **day** - (Optional) Specifies the day of the week that the job should execute on. Must be one of **Monday**, **Tuesday**, **Wednesday**, **Thursday**, **Friday**, **Saturday**, **Sunday**.
- **occurrence** - (Optional) Specifies the week the job should run on. For example 1 for the first week, -1 for the last week of the month. Must be between -5 and 5.

» Attributes Reference

The following attributes are exported:

- **id** - The Scheduler Job ID.

authentication_certificate block exports the following:

- **thumbprint** - (Computed) The certificate thumbprint.
- **expiration** - (Computed) The certificate expiration date.
- **subject_name** - (Computed) The certificate's certificate subject name.

» Import

Scheduler Job can be imported using a **resource id**, e.g.

```
terraform import azurerm_scheduler_job.job1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» azurerm_service_fabric_cluster

Manage a Service Fabric Cluster.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "example-resources"
  location  = "West Europe"
}

resource "azurerm_service_fabric_cluster" "test" {
  name                        = "example-servicefabric"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  location                   = "${azurerm_resource_group.test.location}"
  reliability_level           = "Bronze"
  upgrade_mode                = "Manual"
  cluster_code_version        = "6.3.176.9494"
  vm_image                   = "Windows"
  management_endpoint         = "https://example:80"

  node_type {
    name                = "first"
    instance_count      = 3
    is_primary          = true
    client_endpoint_port = 2020
    http_endpoint_port  = 80
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Service Fabric Cluster. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the Resource Group in which the Service Fabric Cluster exists. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the Azure Region where the Service Fabric Cluster should exist. Changing this forces a new resource to be created.

- **reliability_level** - (Required) Specifies the Reliability Level of the Cluster. Possible values include **None**, **Bronze**, **Silver**, **Gold** and **Platinum**.

NOTE: The Reliability Level of the Cluster depends on the number of nodes in the Cluster: **Platinum** requires at least 9 VM's, **Gold** requires at least 7 VM's, **Silver** requires at least 5 VM's, **Bronze** requires at least 3 VM's.

- **management_endpoint** - (Required) Specifies the Management Endpoint of the cluster such as **http://example.com**. Changing this forces a new resource to be created.
- **node_type** - (Required) One or more **node_type** blocks as defined below.
- **upgrade_mode** - (Required) Specifies the Upgrade Mode of the cluster. Possible values are **Automatic** or **Manual**.
- **vm_image** - (Required) Specifies the Image expected for the Service Fabric Cluster, such as **Windows**. Changing this forces a new resource to be created.

-
- **cluster_code_version** - (Optional) Required if Upgrade Mode set to **Manual**, Specifies the Version of the Cluster Code of the cluster.
 - **add_on_features** - (Optional) A List of one or more features which should be enabled, such as **DnsService**.
 - **azure_active_directory** - (Optional) An **azure_active_directory** block as defined below. Changing this forces a new resource to be created.
 - **certificate** - (Optional) A **certificate** block as defined below.
 - **reverse_proxy_certificate** - (Optional) A **reverse_proxy_certificate** block as defined below.
 - **client_certificate_thumbprint** - (Optional) One or two **client_certificate_thumbprint** blocks as defined below.

NOTE: If Client Certificates are enabled then at a Certificate must be configured on the cluster.

- **diagnostics_config** - (Optional) A **diagnostics_config** block as defined below. Changing this forces a new resource to be created.
- **fabric_settings** - (Optional) One or more **fabric_settings** blocks as defined below.
- **tags** - (Optional) A mapping of tags to assign to the resource.

A **azure_active_directory** block supports the following:

- **tenant_id** - (Required) The Azure Active Directory Tenant ID. Changing this forces a new resource to be created.
- **cluster_application_id** - (Required) The Azure Active Directory Cluster Application ID. Changing this forces a new resource to be created.
- **client_application_id** - (Required) The Azure Active Directory Client ID which should be used for the Client Application. Changing this forces a new resource to be created.

A **certificate** block supports the following:

- **thumbprint** - (Required) The Thumbprint of the Certificate.
- **thumbprint_secondary** - (Required) The Secondary Thumbprint of the Certificate.
- **x509_store_name** - (Required) The X509 Store where the Certificate Exists, such as My.

A **reverse_proxy_certificate** block supports the following:

- **thumbprint** - (Required) The Thumbprint of the Certificate.
- **thumbprint_secondary** - (Required) The Secondary Thumbprint of the Certificate.
- **x509_store_name** - (Required) The X509 Store where the Certificate Exists, such as My.

A **client_certificate_thumbprint** block supports the following:

- **thumbprint** - (Required) The Thumbprint associated with the Client Certificate.
- **is_admin** - (Required) Does the Client Certificate have Admin Access to the cluster? Non-admin clients can only perform read only operations on the cluster.

A **diagnostics_config** block supports the following:

- **storage_account_name** - (Required) The name of the Storage Account where the Diagnostics should be sent to.
- **protected_account_key_name** - (Required) The protected diagnostics storage key name, such as **StorageAccountKey1**.
- **blob_endpoint** - (Required) The Blob Endpoint of the Storage Account.

- `queue_endpoint` - (Required) The Queue Endpoint of the Storage Account.
 - `table_endpoint` - (Required) The Table Endpoint of the Storage Account.
-

A `fabric_settings` block supports the following:

- `name` - (Required) The name of the Fabric Setting, such as **Security** or **Federation**.
 - `parameters` - (Optional) A map containing settings for the specified Fabric Setting.
-

A `node_type` block supports the following:

- `name` - (Required) The name of the Node Type. Changing this forces a new resource to be created.
- `placement_properties` - (Optional) The placement tags applied to nodes in the node type, which can be used to indicate where certain services (workload) should run.
- `capacities` - (Optional) The capacity tags applied to the nodes in the node type, the cluster resource manager uses these tags to understand how much resource a node has.
- `instance_count` - (Required) The number of nodes for this Node Type.
- `is_primary` - (Required) Is this the Primary Node Type? Changing this forces a new resource to be created.
- `client_endpoint_port` - (Required) The Port used for the Client Endpoint for this Node Type. Changing this forces a new resource to be created.
- `http_endpoint_port` - (Required) The Port used for the HTTP Endpoint for this Node Type. Changing this forces a new resource to be created.
- `durability_level` - (Optional) The Durability Level for this Node Type. Possible values include **Bronze**, **Gold** and **Silver**. Defaults to **Bronze**. Changing this forces a new resource to be created.
- `application_ports` - (Optional) A `application_ports` block as defined below.
- `ephemeral_ports` - (Optional) A `ephemeral_ports` block as defined below.
- `reverse_proxy_endpoint_port` - (Optional) The Port used for the Reverse Proxy Endpoint for this Node Type. Changing this will upgrade the cluster.

A `application_ports` block supports the following:

- `start_port` - (Required) The start of the Application Port Range on this Node Type.
 - `end_port` - (Required) The end of the Application Port Range on this Node Type.
-

A `ephemeral_ports` block supports the following:

- `start_port` - (Required) The start of the Ephemeral Port Range on this Node Type.
- `end_port` - (Required) The end of the Ephemeral Port Range on this Node Type.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the Service Fabric Cluster.
- `cluster_endpoint` - The Cluster Endpoint for this Service Fabric Cluster.

» Import

Service Fabric Clusters can be imported using the `resource id`, e.g.

```
terraform import azurerm_service_fabric_cluster.cluster1 /subscriptions/00000000-0000-0000-0000-000000000000
```

» `azurerm_stream_analytics_job`

Manages a Stream Analytics Job.

» Example Usage

```
resource "azurerm_resource_group" "example" {
  name     = "example-resources"
  location = "West Europe"
}

resource "azurerm_stream_analytics_job" "example" {
```

```

name = "example-job"
resource_group_name = "${azurerm_resource_group.example.name}"
location = "${azurerm_resource_group.example.location}"
compatibility_level = "1.1"
data_locale = "en-GB"
events_late_arrival_max_delay_in_seconds = 60
events_out_of_order_max_delay_in_seconds = 50
events_out_of_order_policy = "Adjust"
output_error_policy = "Drop"
streaming_units = 3

transformation_query = <<QUERY
    SELECT *
    INTO [YourOutputAlias]
    FROM [YourInputAlias]
QUERY
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Stream Analytics Job. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the Resource Group where the Stream Analytics Job should exist. Changing this forces a new resource to be created.
- **location** - (Required) The Azure Region in which the Resource Group exists. Changing this forces a new resource to be created.
- **compatibility_level** - (Required) Specifies the compatibility level for this job - which controls certain runtime behaviors of the streaming job. Possible values are 1.0 and 1.1.

NOTE: Support for Compatibility Level 1.2 is dependent on a new version of the Stream Analytics API, which being tracked in this issue.

- **data_locale** - (Required) Specifies the Data Locale of the Job, which should be a supported .NET Culture.
- **events_late_arrival_max_delay_in_seconds** - (Required) Specifies the maximum tolerable delay in seconds where events arriving late could be included. Supported range is -1 (indefinite) to 1814399 (20d 23h 59m 59s).
- **events_out_of_order_max_delay_in_seconds** - (Required) Specifies the maximum tolerable delay in seconds where out-of-order events can be

adjusted to be back in order. Supported range is 0 to 599 (9m 59s).

- **events_out_of_order_policy** - (Required) Specifies the policy which should be applied to events which arrive out of order in the input event stream. Possible values are **Adjust** and **Drop**.
- **output_error_policy** - (Required) Specifies the policy which should be applied to events which arrive at the output and cannot be written to the external storage due to being malformed (such as missing column values, column values of wrong type or size). Possible values are **Drop** and **Stop**.
- **streaming_units** - (Required) Specifies the number of streaming units that the streaming job uses. Supported values are 1, 3, 6 and multiples of 6 up to 120.
- **transformation_query** - (Required) Specifies the query that will be run in the streaming job, written in Stream Analytics Query Language (SAQL).

» Attributes Reference

The following attributes are exported in addition to the arguments listed above:

- **id** - The ID of the Stream Analytics Job.
- **job_id** - The Job ID assigned by the Stream Analytics Job.

» Import

Stream Analytics Job's can be imported using the **resource id**, e.g.

```
terraform import azurerm_stream_analytics_job.test /subscriptions/00000000-0000-0000-0000-000000000000
```

» **azurerm_stream_analytics_function_javascript_udf**

Manages a JavaScript UDF Function within Stream Analytics Streaming Job.

» Example Usage

```
data "azurerm_resource_group" "example" {
  name = "example-resources"
}

data "azurerm_stream_analytics_job" "example" {
  name = "example-job"
```



```

    resource_group_name = "${azurerm_resource_group.example.name}"
  }

  resource "azurerm_stream_analytics_function_javascript_udf" "test" {
    name = "example-javascript-function"
    stream_analytics_job_name = "${data.azurerm_stream_analytics_job.example.name}"
    resource_group_name = "${data.azurerm_stream_analytics_job.example.resource_group_name}"
    script = <<SCRIPT
function getRandomNumber(in) {
  return in;
}
SCRIPT

    input {
      type = "bigint"
    }

    output {
      type = "bigint"
    }
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the JavaScript UDF Function. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the Resource Group where the Stream Analytics Job exists. Changing this forces a new resource to be created.
- **stream_analytics_job_name** - (Required) The name of the Stream Analytics Job where this Function should be created. Changing this forces a new resource to be created.
- **input** - (Required) One or more **input** blocks as defined below.
- **output** - (Required) An **output** blocks as defined below.
- **script** - (Required) The JavaScript of this UDF Function.

A **input** block supports the following:

- **type** - The Data Type for the Input Argument of this JavaScript Function. Possible values include **array**, **any**, **bigint**, **datetime**, **float**,

`nvarchar(max)` and `record`.

A output block supports the following:

- `type` - The Data Type output from this JavaScript Function. Possible values include `array`, `any`, `bigint`, `datetime`, `float`, `nvarchar(max)` and `record`.

» Attributes Reference

The following attributes are exported in addition to the arguments listed above:

- `id` - The ID of the Stream Analytics JavaScript UDF Function.

» Import

Stream Analytics JavaScript UDF Functions can be imported using the `resource id`, e.g.

```
terraform import azurerm_stream_analytics_output_blob.test /subscriptions/00000000-0000-0000
```

» `azurerm_stream_analytics_output_blob`

Manages a Stream Analytics Output to Blob Storage.

» Example Usage

```
data "azurerm_resource_group" "example" {
  name = "example-resources"
}

data "azurerm_stream_analytics_job" "example" {
  name                        = "example-job"
  resource_group_name        = "${azurerm_resource_group.example.name}"
}

resource "azurerm_storage_account" "example" {
  name                        = "examplesa"
  resource_group_name        = "${data.azurerm_resource_group.example.name}"
  location                   = "${data.azurerm_resource_group.example.location}"
  account_tier                = "Standard"
  account_replication_type    = "LRS"
}
```

```

}

resource "azurerm_storage_container" "example" {
  name = "example"
  resource_group_name = "${data.azurerm_resource_group.example.name}"
  storage_account_name = "${azurerm_storage_account.example.name}"
  container_access_type = "private"
}

resource "azurerm_stream_analytics_output_blob" "example" {
  name = "output-to-blob-storage"
  stream_analytics_job_name = "${data.azurerm_stream_analytics_job.example.name}"
  resource_group_name = "${data.azurerm_stream_analytics_job.example.resource_group_name}"
  storage_account_name = "${azurerm_storage_account.example.name}"
  storage_account_key = "${azurerm_storage_account.example.primary_access_key}"
  storage_container_name = "${azurerm_storage_container.example.name}"
  path_pattern = "some-pattern"
  date_format = "yyyy-MM-dd"
  time_format = "HH"

  serialization {
    type = "Csv"
    encoding = "UTF8"
    field_delimiter = ","
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Stream Output. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the Resource Group where the Stream Analytics Job exists. Changing this forces a new resource to be created.
- **stream_analytics_job_name** - (Required) The name of the Stream Analytics Job. Changing this forces a new resource to be created.
- **date_format** - (Required) The date format. Wherever `{date}` appears in **path_pattern**, the value of this property is used as the date format instead.
- **path_pattern** - (Required) The blob path pattern. Not a regular expression. It represents a pattern against which blob names will be matched to

determine whether or not they should be included as input or output to the job.

- **storage_account_name** - (Required) The name of the Storage Account.
- **storage_account_key** - (Required) The Access Key which should be used to connect to this Storage Account.
- **storage_container_name** - (Required) The name of the Container within the Storage Account.
- **time_format** - (Required) The time format. Wherever **{time}** appears in **path_pattern**, the value of this property is used as the time format instead.
- **serialization** - (Required) A **serialization** block as defined below.

A **serialization** block supports the following:

- **type** - (Required) The serialization format used for outgoing data streams. Possible values are **Avro**, **Csv** and **Json**.
- **encoding** - (Optional) The encoding of the incoming data in the case of input and the encoding of outgoing data in the case of output.

NOTE: This is required when **type** is set to **Csv** or **Json**.

- **field_delimiter** - (Optional) The delimiter that will be used to separate comma-separated value (CSV) records. Possible values are (space), , (comma), (tab), | (pipe) and ;.

NOTE: This is required when **type** is set to **Csv**.

- **format** - (Optional) Specifies the format of the JSON the output will be written in.

NOTE: This is Required and can only be specified when **type** is set to **Json**.

» Attributes Reference

The following attributes are exported in addition to the arguments listed above:

- **id** - The ID of the Stream Analytics Output Blob Storage.

» Import

Stream Analytics Outputs to Blob Storage can be imported using the **resource id**, e.g.

```
terraform import azurerm_stream_analytics_output_blob.test /subscriptions/00000000-0000-0000
```

» azurerm_stream_analytics_output_eventhub

Manages a Stream Analytics Output to an EventHub.

» Example Usage

```
data "azurerm_resource_group" "example" {
  name = "example-resources"
}

data "azurerm_stream_analytics_job" "example" {
  name                        = "example-job"
  resource_group_name       = "${azurerm_resource_group.example.name}"
}

resource "azurerm_eventhub_namespace" "test" {
  name                = "example-ehnamespace"
  location            = "${data.azurerm_resource_group.example.location}"
  resource_group_name = "${data.azurerm_resource_group.example.name}"
  sku                 = "Standard"
  capacity            = 1
  kafka_enabled       = false
}

resource "azurerm_eventhub" "test" {
  name                = "example-eventhub"
  namespace_name      = "${azurerm_eventhub_namespace.example.name}"
  resource_group_name = "${data.azurerm_resource_group.example.name}"
  partition_count     = 2
  message_retention   = 1
}

resource "azurerm_stream_analytics_output_eventhub" "example" {
  name                = "output-to-eventhub"
  stream_analytics_job_name = "${data.azurerm_stream_analytics_job.example.name}"
  resource_group_name   = "${data.azurerm_stream_analytics_job.example.resource_group_name}"
  eventhub_name         = "${azurerm_eventhub.example.name}"
  servicebus_namespace = "${azurerm_eventhub_namespace.example.name}"
  shared_access_policy_key = "${azurerm_eventhub_namespace.example.default_primary_key}"
  shared_access_policy_name = "RootManageSharedAccessKey"

  serialization {
    type = "Avro"
  }
}
```

}

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Stream Output. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the Resource Group where the Stream Analytics Job exists. Changing this forces a new resource to be created.
- **stream_analytics_job_name** - (Required) The name of the Stream Analytics Job. Changing this forces a new resource to be created.
- **eventhub_name** - (Required) The name of the Event Hub.
- **servicebus_namespace** - (Required) The namespace that is associated with the desired Event Hub, Service Bus Queue, Service Bus Topic, etc.
- **shared_access_policy_key** - (Required) The shared access policy key for the specified shared access policy.
- **shared_access_policy_name** - (Required) The shared access policy name for the Event Hub, Service Bus Queue, Service Bus Topic, etc.
- **serialization** - (Required) A **serialization** block as defined below.

A **serialization** block supports the following:

- **type** - (Required) The serialization format used for outgoing data streams. Possible values are **Avro**, **Csv** and **Json**.
- **encoding** - (Optional) The encoding of the incoming data in the case of input and the encoding of outgoing data in the case of output.

NOTE: This is required when **type** is set to **Csv** or **Json**.

- **field_delimiter** - (Optional) The delimiter that will be used to separate comma-separated value (CSV) records. Possible values are (space), , (comma), (tab), | (pipe) and ;.

NOTE: This is required when **type** is set to **Csv**.

- **format** - (Optional) Specifies the format of the JSON the output will be written in.

NOTE: This is Required and can only be specified when **type** is set to **Json**.

» Attributes Reference

The following attributes are exported in addition to the arguments listed above:

- `id` - The ID of the Stream Analytics Output EventHub.

» Import

Stream Analytics Outputs to an EventHub can be imported using the `resource id`, e.g.

```
terraform import azurerm_stream_analytics_output_eventhub.test /subscriptions/00000000-0000-
```

» `azurerm_stream_analytics_output_servicebus_queue`

Manages a Stream Analytics Output to a ServiceBus Queue.

» Example Usage

```
data "azurerm_resource_group" "example" {
  name = "example-resources"
}

data "azurerm_stream_analytics_job" "example" {
  name                = "example-job"
  resource_group_name = "${azurerm_resource_group.example.name}"
}

resource "azurerm_servicebus_namespace" "example" {
  name                = "example-namespace"
  location            = "${data.azurerm_resource_group.example.location}"
  resource_group_name = "${data.azurerm_resource_group.example.name}"
  sku                 = "Standard"
}

resource "azurerm_servicebus_queue" "example" {
  name                = "example-queue"
  resource_group_name = "${data.azurerm_resource_group.example.name}"
  namespace_name      = "${azurerm_servicebus_namespace.example.name}"
  enable_partitioning = true
}

resource "azurerm_stream_analytics_output_servicebus_queue" "test" {
```

```

name = "blob-storage-output"
stream_analytics_job_name = "${data.azure_rm_stream_analytics_job.example.name}"
resource_group_name = "${data.azure_rm_stream_analytics_job.example.resource_group_name}"
queue_name = "${azurerm_servicebus_queue.example.name}"
servicebus_namespace = "${azurerm_servicebus_namespace.example.name}"
shared_access_policy_key = "${azurerm_servicebus_namespace.example.default_primary_key}"
shared_access_policy_name = "RootManageSharedAccessKey"

serialization {
  format = "Avro"
}
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Stream Output. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the Resource Group where the Stream Analytics Job exists. Changing this forces a new resource to be created.
- **stream_analytics_job_name** - (Required) The name of the Stream Analytics Job. Changing this forces a new resource to be created.
- **queue_name** - (Required) The name of the Service Bus Queue.
- **servicebus_namespace** - (Required) The namespace that is associated with the desired Event Hub, Service Bus Queue, Service Bus Topic, etc.
- **shared_access_policy_key** - (Required) The shared access policy key for the specified shared access policy.
- **shared_access_policy_name** - (Required) The shared access policy name for the Event Hub, Service Bus Queue, Service Bus Topic, etc.
- **serialization** - (Required) A `serialization` block as defined below.

A `serialization` block supports the following:

- **type** - (Required) The serialization format used for outgoing data streams. Possible values are `Avro`, `Csv` and `Json`.
- **encoding** - (Optional) The encoding of the incoming data in the case of input and the encoding of outgoing data in the case of output.

NOTE: This is required when `type` is set to `Csv` or `Json`.

- `field_delimiter` - (Optional) The delimiter that will be used to separate comma-separated value (CSV) records. Possible values are (space), , (comma), (tab), | (pipe) and ;.

NOTE: This is required when `type` is set to `Csv`.

- `format` - (Optional) Specifies the format of the JSON the output will be written in.

NOTE: This is Required and can only be specified when `type` is set to `Json`.

» Attributes Reference

The following attributes are exported in addition to the arguments listed above:

- `id` - The ID of the Stream Analytics Output ServiceBus Queue.

» Import

Stream Analytics Output ServiceBus Queue's can be imported using the resource `id`, e.g.

```
terraform import azurerm_stream_analytics_output_servicebus_queue.test /subscriptions/000000
```

» azurerm_stream_analytics_stream_input_blob

Manages a Stream Analytics Stream Input Blob.

» Example Usage

```
data "azurerm_resource_group" "example" {
  name = "example-resources"
}

data "azurerm_stream_analytics_job" "example" {
  name                  = "example-job"
  resource_group_name = "${azurerm_resource_group.example.name}"
}

resource "azurerm_storage_account" "example" {
  name                = "examplestoracc"
  resource_group_name = "${azurerm_resource_group.example.name}"
  location            = "${azurerm_resource_group.example.location}"
  account_tier        = "Standard"
}
```

```

    account_replication_type = "LRS"
}

resource "azurerm_storage_container" "example" {
  name                        = "example"
  resource_group_name        = "${azurerm_resource_group.example.name}"
  storage_account_name       = "${azurerm_storage_account.example.name}"
  container_access_type      = "private"
}

resource "azurerm_stream_analytics_stream_input_eventhub" "test" {
  name                        = "eventhub-stream-input"
  stream_analytics_job_name  = "${data.azurerm_stream_analytics_job.example.name}"
  resource_group_name        = "${data.azurerm_stream_analytics_job.example.resource_group}"
  storage_account_name       = "${azurerm_storage_account.example.name}"
  storage_account_key        = "${azurerm_storage_account.example.primary_access_key}"
  storage_container_name     = "${azurerm_storage_container.example.name}"
  path_pattern               = "some-random-pattern"
  date_format                = "yyyy/MM/dd"
  time_format                = "HH"

  serialization {
    type      = "Json"
    encoding  = "UTF8"
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Stream Input Blob. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the Resource Group where the Stream Analytics Job exists. Changing this forces a new resource to be created.
- **stream_analytics_job_name** - (Required) The name of the Stream Analytics Job. Changing this forces a new resource to be created.
- **date_format** - (Required) The date format. Wherever `{date}` appears in **path_pattern**, the value of this property is used as the date format instead.
- **path_pattern** - (Required) The blob path pattern. Not a regular expression. It represents a pattern against which blob names will be matched to

determine whether or not they should be included as input or output to the job.

- **storage_account_name** - (Required) The name of the Storage Account.
- **storage_account_key** - (Required) The Access Key which should be used to connect to this Storage Account.
- **storage_container_name** - (Required) The name of the Container within the Storage Account.
- **time_format** - (Required) The time format. Wherever **{time}** appears in **path_pattern**, the value of this property is used as the time format instead.
- **serialization** - (Required) A **serialization** block as defined below.

A **serialization** block supports the following:

- **type** - (Required) The serialization format used for incoming data streams. Possible values are **Avro**, **Csv** and **Json**.
- **encoding** - (Optional) The encoding of the incoming data in the case of input and the encoding of outgoing data in the case of output.

NOTE: This is required when **type** is set to **Csv** or **Json**.

- **field_delimiter** - (Optional) The delimiter that will be used to separate comma-separated value (CSV) records. Possible values are (space), , (comma), (tab), | (pipe) and ;.

NOTE: This is required when **type** is set to **Csv**.

» Attributes Reference

The following attributes are exported in addition to the arguments listed above:

- **id** - The ID of the Stream Analytics Stream Input Blob.

» Import

Stream Analytics Stream Input Blob's can be imported using the **resource id**, e.g.

```
terraform import azurerm_stream_analytics_stream_input_blob.test /subscriptions/00000000-0000-0000-0000-000000000000/resourceGroups/rg-test/providers/Microsoft.StreamAnalytics/streamingjobs/job-test/streaminputblobs/blob-test
```

» azurerm_stream_analytics_stream_input_eventhub

Manages a Stream Analytics Stream Input EventHub.

» Example Usage

```
data "azurerm_resource_group" "example" {
  name = "example-resources"
}

data "azurerm_stream_analytics_job" "example" {
  name                        = "example-job"
  resource_group_name = "${azurerm_resource_group.example.name}"
}

resource "azurerm_eventhub_namespace" "example" {
  name                = "example-namespace"
  location            = "${data.azurerm_resource_group.example.location}"
  resource_group_name = "${data.azurerm_resource_group.example.name}"
  sku                 = "Standard"
  capacity            = 1
  kafka_enabled       = false
}

resource "azurerm_eventhub" "example" {
  name                = "example-eventhub"
  namespace_name      = "${azurerm_eventhub_namespace.example.name}"
  resource_group_name = "${data.azurerm_resource_group.example.name}"
  partition_count     = 2
  message_retention    = 1
}

resource "azurerm_eventhub_consumer_group" "test" {
  name                = "example-consumergroup"
  namespace_name      = "${azurerm_eventhub_namespace.test.name}"
  eventhub_name       = "${azurerm_eventhub.test.name}"
  resource_group_name = "${data.azurerm_resource_group.test.name}"
}

resource "azurerm_stream_analytics_stream_input_eventhub" "test" {
  name                = "eventhub-stream-input"
  stream_analytics_job_name = "${data.azurerm_stream_analytics_job.example.name}"
  resource_group_name    = "${data.azurerm_stream_analytics_job.example.resource_group}"
  eventhub_consumer_group_name = "${azurerm_eventhub_consumer_group.example.name}"
}
```

```

eventhub_name          = "${azurerm_eventhub.example.name}"
servicebus_namespace   = "${azurerm_eventhub_namespace.example.name}"
shared_access_policy_key = "${azurerm_eventhub_namespace.example.default_primary_key}"
shared_access_policy_name = "RootManageSharedAccessKey"

serialization {
  type      = "Json"
  encoding  = "UTF8"
}
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Stream Input EventHub. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the Resource Group where the Stream Analytics Job exists. Changing this forces a new resource to be created.
- **stream_analytics_job_name** - (Required) The name of the Stream Analytics Job. Changing this forces a new resource to be created.
- **eventhub_consumer_group_name** - (Required) The name of an Event Hub Consumer Group that should be used to read events from the Event Hub. Specifying distinct consumer group names for multiple inputs allows each of those inputs to receive the same events from the Event Hub.
- **eventhub_name** - (Required) The name of the Event Hub.
- **servicebus_namespace** - (Required) The namespace that is associated with the desired Event Hub, Service Bus Queue, Service Bus Topic, etc.
- **shared_access_policy_key** - (Required) The shared access policy key for the specified shared access policy.
- **shared_access_policy_name** - (Required) The shared access policy name for the Event Hub, Service Bus Queue, Service Bus Topic, etc.
- **serialization** - (Required) A `serialization` block as defined below.

A `serialization` block supports the following:

- **type** - (Required) The serialization format used for incoming data streams. Possible values are `Avro`, `Csv` and `Json`.

- **encoding** - (Optional) The encoding of the incoming data in the case of input and the encoding of outgoing data in the case of output.

NOTE: This is required when **type** is set to **Csv** or **Json**.

- **field_delimiter** - (Optional) The delimiter that will be used to separate comma-separated value (CSV) records. Possible values are (space), , (comma), (tab), | (pipe) and ;.

NOTE: This is required when **type** is set to **Csv**.

» Attributes Reference

The following attributes are exported in addition to the arguments listed above:

- **id** - The ID of the Stream Analytics Stream Input EventHub.

» Import

Stream Analytics Stream Input EventHub's can be imported using the **resource id**, e.g.

```
terraform import azurerm_stream_analytics_stream_input_eventhub.test /subscriptions/00000000
```

» azurerm_stream_analytics_stream_input_iothub

Manages a Stream Analytics Stream Input IoT Hub.

» Example Usage

```
data "azurerm_resource_group" "example" {
  name = "example-resources"
}

data "azurerm_stream_analytics_job" "example" {
  name                        = "example-job"
  resource_group_name       = "${azurerm_resource_group.example.name}"
}

resource "azurerm_iothub" "test" {
  name                        = "example-iothub"
  resource_group_name       = "${azurerm_resource_group.example.name}"
  location                   = "${azurerm_resource_group.example.location}"
}
```

```

sku {
  name      = "S1"
  tier       = "Standard"
  capacity  = "1"
}

resource "azurerm_stream_analytics_stream_input_iothub" "test" {
  name                        = "example-iothub-input"
  stream_analytics_job_name  = "${data.azurerm_stream_analytics_job.example.name}"
  resource_group_name        = "${data.azurerm_stream_analytics_job.example.resource_group}"
  endpoint                   = "messages/events"
  eventhub_consumer_group_name = "$Default"
  iothub_namespace           = "${azurerm_iothub.example.name}"
  shared_access_policy_key    = "${azurerm_iothub.example.shared_access_policy.0.primary_key}"
  shared_access_policy_name   = "iothubowner"

  serialization {
    type      = "Json"
    encoding  = "UTF8"
  }
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the Stream Input IoT Hub. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the Resource Group where the Stream Analytics Job exists. Changing this forces a new resource to be created.
- **stream_analytics_job_name** - (Required) The name of the Stream Analytics Job. Changing this forces a new resource to be created.
- **eventhub_consumer_group_name** - (Required) The name of an Event Hub Consumer Group that should be used to read events from the Event Hub. Specifying distinct consumer group names for multiple inputs allows each of those inputs to receive the same events from the Event Hub.
- **endpoint** - (Required) The IoT Hub endpoint to connect to (ie. messages/events, messages/operationsMonitoringEvents, etc.).
- **iothub_namespace** - (Required) The name or the URI of the IoT Hub.

- **serialization** - (Required) A **serialization** block as defined below.
- **shared_access_policy_key** - (Required) The shared access policy key for the specified shared access policy.
- **shared_access_policy_name** - (Required) The shared access policy name for the Event Hub, Service Bus Queue, Service Bus Topic, etc.

A **serialization** block supports the following:

- **type** - (Required) The serialization format used for incoming data streams. Possible values are **Avro**, **Csv** and **Json**.
- **encoding** - (Optional) The encoding of the incoming data in the case of input and the encoding of outgoing data in the case of output.

NOTE: This is required when **type** is set to **Csv** or **Json**.

- **field_delimiter** - (Optional) The delimiter that will be used to separate comma-separated value (CSV) records. Possible values are (space), , (comma), (tab), | (pipe) and ;.

NOTE: This is required when **type** is set to **Csv**.

» Attributes Reference

The following attributes are exported in addition to the arguments listed above:

- **id** - The ID of the Stream Analytics Stream Input IoTHub.

» Import

Stream Analytics Stream Input IoTHub's can be imported using the **resource id**, e.g.

```
terraform import azurerm_stream_analytics_stream_input_iotHub.test /subscriptions/00000000-0000-0000-0000-000000000000/resourceGroups/rg-test/providers/Microsoft.StreamAnalytics/streamingjobs/instance1
```

» azurerm_storage_account

Manage an Azure Storage Account.

» Example Usage

```
resource "azurerm_resource_group" "testrg" {
  name     = "resourceGroupName"
  location = "West US"
}
```



```

    location = "westus"
  }

  resource "azurerm_storage_account" "testsa" {
    name                        = "storageaccountname"
    resource_group_name       = "${azurerm_resource_group.testrg.name}"
    location                  = "westus"
    account_tier              = "Standard"
    account_replication_type = "GRS"

    tags = {
      environment = "staging"
    }
  }
}

```

» Example Usage with Network Rules

```

resource "azurerm_resource_group" "testrg" {
  name      = "resourceGroupName"
  location = "westus"
}

resource "azurerm_virtual_network" "test" {
  name            = "virtnetname"
  address_space   = ["10.0.0.0/16"]
  location        = "${azurerm_resource_group.testrg.location}"
  resource_group_name = "${azurerm_resource_group.testrg.name}"
}

resource "azurerm_subnet" "test" {
  name                = "subnetname"
  resource_group_name = "${azurerm_resource_group.testrg.name}"
  virtual_network_name = "${azurerm_virtual_network.test.name}"
  address_prefix       = "10.0.2.0/24"
  service_endpoints    = ["Microsoft.Sql", "Microsoft.Storage"]
}

resource "azurerm_storage_account" "testsa" {
  name                        = "storageaccountname"
  resource_group_name       = "${azurerm_resource_group.testrg.name}"

  location                  = "${azurerm_resource_group.testrg.location}"
  account_tier              = "Standard"
  account_replication_type = "LRS"
}

```

```

network_rules {
  ip_rules          = ["100.0.0.1"]
  virtual_network_subnet_ids = ["${azurerm_subnet.test.id}"]
}

tags = {
  environment = "staging"
}
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the storage account. Changing this forces a new resource to be created. This must be unique across the entire Azure service, not just within the resource group.
- **resource_group_name** - (Required) The name of the resource group in which to create the storage account. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **account_kind** - (Optional) Defines the Kind of account. Valid options are **Storage**, **StorageV2** and **BlobStorage**. Changing this forces a new resource to be created. Defaults to **Storage**.
- **account_tier** - (Required) Defines the Tier to use for this storage account. Valid options are **Standard** and **Premium**. Changing this forces a new resource to be created
- **account_replication_type** - (Required) Defines the type of replication to use for this storage account. Valid options are **LRS**, **GRS**, **RAGRS** and **ZRS**.
- **access_tier** - (Optional) Defines the access tier for **BlobStorage** and **StorageV2** accounts. Valid options are **Hot** and **Cool**, defaults to **Hot**.
- **enable_blob_encryption** - (Optional) Boolean flag which controls if Encryption Services are enabled for Blob storage, see here for more information. Defaults to **true**.
- **enable_file_encryption** - (Optional) Boolean flag which controls if Encryption Services are enabled for File storage, see here for more information. Defaults to **true**.
- **enable_https_traffic_only** - (Optional) Boolean flag which forces HTTPS if enabled, see here for more information.

- `is_hns_enabled` - (Optional) Is Hierarchical Namespace enabled? This can be used with Azure Data Lake Storage Gen 2 (see here for more information). Changing this forces a new resource to be created.
- `account_encryption_source` - (Optional) The Encryption Source for this Storage Account. Possible values are `Microsoft.Keyvault` and `Microsoft.Storage`. Defaults to `Microsoft.Storage`.
- `custom_domain` - (Optional) A `custom_domain` block as documented below.
- `network_rules` - (Optional) A `network_rules` block as documented below.
- `tags` - (Optional) A mapping of tags to assign to the resource.
- `identity` - (Optional) A Managed Service Identity block as defined below.

-
- `custom_domain` supports the following:
 - `name` - (Optional) The Custom Domain Name to use for the Storage Account, which will be validated by Azure.
 - `use_subdomain` - (Optional) Should the Custom Domain Name be validated by using indirect CNAME validation?

-
- `network_rules` supports the following:
 - `bypass` - (Optional) Specifies whether traffic is bypassed for Logging/Metrics/AzureServices. Valid options are any combination of `Logging`, `Metrics`, `AzureServices`, or `None`.
 - `ip_rules` - (Optional) List of public IP or IP ranges in CIDR Format. Only IPV4 addresses are allowed. Private IP address ranges (as defined in RFC 1918) are not allowed.
 - `virtual_network_subnet_ids` - (Optional) A list of resource ids for subnets.

Note: If specifying `network_rules`, one of either `ip_rules` or `virtual_network_subnet_ids` must be specified.

Note: More information on Validation is available here

`identity` supports the following:

- `type` - (Required) Specifies the identity type of the Storage Account. At this time the only allowed value is `SystemAssigned`.

The assigned `principal_id` and `tenant_id` can be retrieved after the identity type has been set to `SystemAssigned` and Storage Account has been created. More details are available below.

» Attributes Reference

The following attributes are exported in addition to the arguments listed above:

- `id` - The storage account Resource ID.
- `primary_location` - The primary location of the storage account.
- `secondary_location` - The secondary location of the storage account.
- `primary_blob_endpoint` - The endpoint URL for blob storage in the primary location.
- `primary_blob_host` - The hostname with port if applicable for blob storage in the primary location.
- `secondary_blob_endpoint` - The endpoint URL for blob storage in the secondary location.
- `secondary_blob_host` - The hostname with port if applicable for blob storage in the secondary location.
- `primary_queue_endpoint` - The endpoint URL for queue storage in the primary location.
- `primary_queue_host` - The hostname with port if applicable for queue storage in the primary location.
- `secondary_queue_endpoint` - The endpoint URL for queue storage in the secondary location.
- `secondary_queue_host` - The hostname with port if applicable for queue storage in the secondary location.
- `primary_table_endpoint` - The endpoint URL for table storage in the primary location.
- `primary_table_host` - The hostname with port if applicable for table storage in the primary location.
- `secondary_table_endpoint` - The endpoint URL for table storage in the secondary location.
- `secondary_table_host` - The hostname with port if applicable for table storage in the secondary location.
- `primary_file_endpoint` - The endpoint URL for file storage in the primary location.

- **primary_file_host** - The hostname with port if applicable for file storage in the primary location.
- **secondary_file_endpoint** - The endpoint URL for file storage in the secondary location.
- **secondary_file_host** - The hostname with port if applicable for file storage in the secondary location.
- **primary_dfs_endpoint** - The endpoint URL for DFS storage in the primary location.
- **primary_dfs_host** - The hostname with port if applicable for DFS storage in the primary location.
- **secondary_dfs_endpoint** - The endpoint URL for DFS storage in the secondary location.
- **secondary_dfs_host** - The hostname with port if applicable for DFS storage in the secondary location.
- **primary_web_endpoint** - The endpoint URL for web storage in the primary location.
- **primary_web_host** - The hostname with port if applicable for web storage in the primary location.
- **secondary_web_endpoint** - The endpoint URL for web storage in the secondary location.
- **secondary_web_host** - The hostname with port if applicable for web storage in the secondary location.
- **primary_access_key** - The primary access key for the storage account.
- **secondary_access_key** - The secondary access key for the storage account.
- **primary_connection_string** - The connection string associated with the primary location.
- **secondary_connection_string** - The connection string associated with the secondary location.
- **primary_blob_connection_string** - The connection string associated with the primary blob location.
- **secondary_blob_connection_string** - The connection string associated with the secondary blob location.
- **identity** - An **identity** block as defined below, which contains the Identity information for this Storage Account.

identity exports the following:

- **principal_id** - The Principal ID for the Service Principal associated with the Identity of this Storage Account.
- **tenant_id** - The Tenant ID for the Service Principal associated with the Identity of this Storage Account.

You can access the Principal ID via `${azurerm_storage_account.test.identity.0.principal_id}` and the Tenant ID via `${azurerm_storage_account.test.identity.0.tenant_id}`

» Import

Storage Accounts can be imported using the `resource id`, e.g.

```
terraform import azurerm_storage_account.storageAcc1 /subscriptions/00000000-0000-0000-0000-
```

» azurerm__storage__blob

Manage an Azure Storage Blob.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "acctestRG-d"
  location = "westus"
}

resource "azurerm_storage_account" "test" {
  name                        = "acctestaccs"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  location                   = "westus"
  account_tier               = "Standard"
  account_replication_type   = "LRS"
}

resource "azurerm_storage_container" "test" {
  name                = "vhds"
  resource_group_name = "${azurerm_resource_group.test.name}"
  storage_account_name = "${azurerm_storage_account.test.name}"
  container_access_type = "private"
}

resource "azurerm_storage_blob" "testsb" {
```

```

name = "sample.vhd"

resource_group_name    = "${azurerm_resource_group.test.name}"
storage_account_name   = "${azurerm_storage_account.test.name}"
storage_container_name = "${azurerm_storage_container.test.name}"

type = "page"
size = 5120
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the storage blob. Must be unique within the storage container the blob is located.
- **resource_group_name** - (Required) The name of the resource group in which to create the storage container. Changing this forces a new resource to be created.
- **storage_account_name** - (Required) Specifies the storage account in which to create the storage container. Changing this forces a new resource to be created.
- **storage_container_name** - (Required) The name of the storage container in which this blob should be created.
- **type** - (Optional) The type of the storage blob to be created. One of either **block** or **page**. When not copying from an existing blob, this becomes required.
- **size** - (Optional) Used only for **page** blobs to specify the size in bytes of the blob to be created. Must be a multiple of 512. Defaults to 0.
- **content_type** - (Optional) The content type of the storage blob. Cannot be defined if **source_uri** is defined. Defaults to **application/octet-stream**.
- **source** - (Optional) An absolute path to a file on the local system. Cannot be defined if **source_uri** is defined.
- **source_uri** - (Optional) The URI of an existing blob, or a file in the Azure File service, to use as the source contents for the blob to be created. Changing this forces a new resource to be created. Cannot be defined if **source** is defined.
- **parallelism** - (Optional) The number of workers per CPU core to run for concurrent uploads. Defaults to 8.

- **attempts** - (Optional) The number of attempts to make per page or block when uploading. Defaults to 1.
- **metadata** - (Optional) A map of custom blob metadata.

» Attributes Reference

The following attributes are exported in addition to the arguments listed above:

- **id** - The ID of the Storage Blob.
- **url** - The URL of the blob

» Import

Storage Blob's can be imported using the **resource id**, e.g.

```
terraform import azurerm_storage_blob.blob1 https://example.blob.core.windows.net/container/
```

» azurerm_storage_container

Manage an Azure Storage Container.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "acctestRG"
  location = "westus"
}

resource "azurerm_storage_account" "test" {
  name                        = "accteststorageaccount"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  location                   = "westus"
  account_tier                = "Standard"
  account_replication_type    = "LRS"

  tags = {
    environment = "staging"
  }
}

resource "azurerm_storage_container" "test" {
  name = "vhds"
```



```

resource_group_name    = "${azurerm_resource_group.test.name}"
storage_account_name   = "${azurerm_storage_account.test.name}"
container_access_type = "private"
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the storage container. Must be unique within the storage service the container is located.
- **resource_group_name** - (Required) The name of the resource group in which to create the storage container. Changing this forces a new resource to be created.
- **storage_account_name** - (Required) Specifies the storage account in which to create the storage container. Changing this forces a new resource to be created.
- **container_access_type** - (Optional) The 'interface' for access the container provides. Can be either `blob`, `container` or `private`. Defaults to `private`.

» Attributes Reference

The following attributes are exported in addition to the arguments listed above:

- **id** - The ID of the Storage Container.
- **properties** - Key-value definition of additional properties associated to the storage container

» Import

Storage Containers can be imported using the `resource id`, e.g.

```
terraform import azurerm_storage_container.container1 https://example.blob.core.windows.net/
```

» azurerm__storage__queue

Manage an Azure Storage Queue.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "example-resources"
  location  = "westus"
}

resource "azurerm_storage_account" "test" {
  name                        = "examplestorageacc"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  location                   = "${azurerm_resource_group.test.location}"
  account_tier               = "Standard"
  account_replication_type   = "LRS"
}

resource "azurerm_storage_queue" "test" {
  name                      = "mysamplequeue"
  resource_group_name       = "${azurerm_resource_group.test.name}"
  storage_account_name      = "${azurerm_storage_account.test.name}"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the storage queue. Must be unique within the storage account the queue is located.
- **resource_group_name** - (Required) The name of the resource group in which to create the storage queue. Changing this forces a new resource to be created.
- **storage_account_name** - (Required) Specifies the storage account in which to create the storage queue. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported in addition to the arguments listed above:

- **id** - The ID of the Storage Queue.

» Import

Storage Queue's can be imported using the **resource id**, e.g.

```
terraform import azurerm_storage_queue.queue1 https://example.queue.core.windows.net/queue1
```

» azurerm__storage__share

Manage an Azure Storage File Share.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "azuretest"
  location = "westus"
}

resource "azurerm_storage_account" "test" {
  name                        = "azureteststorage"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  location                   = "westus"
  account_tier               = "Standard"
  account_replication_type   = "LRS"
}

resource "azurerm_storage_share" "testshare" {
  name = "sharename"

  resource_group_name = "${azurerm_resource_group.test.name}"
  storage_account_name = "${azurerm_storage_account.test.name}"

  quota = 50
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the share. Must be unique within the storage account where the share is located.
- **resource_group_name** - (Required) The name of the resource group in which to create the share. Changing this forces a new resource to be created.
- **storage_account_name** - (Required) Specifies the storage account in which to create the share. Changing this forces a new resource to be created.

- **quota** - (Optional) The maximum size of the share, in gigabytes. Must be greater than 0, and less than or equal to 5 TB (5120 GB). Default is 5120.

» Attributes Reference

The following attributes are exported in addition to the arguments listed above:

- **id** - The storage share Resource ID.
- **url** - The URL of the share

» Import

Storage Shares can be imported using the `resource id`, e.g.

```
terraform import azurerm_storage_share.testShare storageShareName/resourceGroupName/storageAccountName
```

NOTE: This identifier is unique to Terraform

» azurerm__storage__table

Manage an Azure Storage Table.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name     = "azuretest"
  location = "westus"
}

resource "azurerm_storage_account" "test" {
  name                     = "azureteststorage1"
  resource_group_name     = "${azurerm_resource_group.test.name}"
  location                 = "westus"
  account_tier             = "Standard"
  account_replication_type = "LRS"
}

resource "azurerm_storage_table" "test" {
  name                = "mysampletable"
  resource_group_name = "${azurerm_resource_group.test.name}"
  storage_account_name = "${azurerm_storage_account.test.name}"
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the storage table. Must be unique within the storage account the table is located.
- **resource_group_name** - (Required) The name of the resource group in which to create the storage table. Changing this forces a new resource to be created.
- **storage_account_name** - (Required) Specifies the storage account in which to create the storage table. Changing this forces a new resource to be created.

» Attributes Reference

The following attributes are exported in addition to the arguments listed above:

- **id** - The ID of the Storage Table.

» Import

Storage Table's can be imported using the **resource id**, e.g.

```
terraform import azurerm_storage_table.table1 https://example.table.core.windows.net/table1
```

» azurerm__template__deployment

Manage a template deployment of resources

Note on ARM Template Deployments: Due to the way the underlying Azure API is designed, Terraform can only manage the deployment of the ARM Template - and not any resources which are created by it. This means that when deleting the **azurerm_template_deployment** resource, Terraform will only remove the reference to the deployment, whilst leaving any resources created by that ARM Template Deployment. One workaround for this is to use a unique Resource Group for each ARM Template Deployment, which means deleting the Resource Group would contain any resources created within it - however this isn't ideal. More information.

» Example Usage

Note: This example uses Storage Accounts and Public IP's which are natively supported by Terraform - we'd highly recommend using the Native Resources

where possible instead rather than an ARM Template, for the reasons outlined above.

```
resource "azurerm_resource_group" "test" {
  name      = "acctestRG-01"
  location  = "West US"
}

resource "azurerm_template_deployment" "test" {
  name                = "acctesttemplate-01"
  resource_group_name = "${azurerm_resource_group.test.name}"

  template_body = <<DEPLOY
{
  "$schema": "https://schema.management.azure.com/schemas/2015-01-01/deploymentTemplate.json",
  "contentVersion": "1.0.0.0",
  "parameters": {
    "storageAccountType": {
      "type": "string",
      "defaultValue": "Standard_LRS",
      "allowedValues": [
        "Standard_LRS",
        "Standard_GRS",
        "Standard_ZRS"
      ],
      "metadata": {
        "description": "Storage Account type"
      }
    }
  },
  "variables": {
    "location": "[resourceGroup().location]",
    "storageAccountName": "[concat(uniquestring(resourceGroup().id), 'storage')]",
    "publicIPAddressName": "[concat('myPublicIp', uniquestring(resourceGroup().id))]",
    "publicIPAddressType": "Dynamic",
    "apiVersion": "2015-06-15",
    "dnsLabelPrefix": "terraform-acctest"
  },
  "resources": [
    {
      "type": "Microsoft.Storage/storageAccounts",
      "name": "[variables('storageAccountName')]",
      "apiVersion": "[variables('apiVersion')]",
      "location": "[variables('location')]",
      "properties": {
        "accountType": "[parameters('storageAccountType')]"
      }
    }
  ]
}
```

```

    }
  },
  {
    "type": "Microsoft.Network/publicIPAddresses",
    "apiVersion": "[variables('apiVersion')]",
    "name": "[variables('publicIPAddressName')]",
    "location": "[variables('location')]",
    "properties": {
      "publicIPAllocationMethod": "[variables('publicIPAllocationType')]",
      "dnsSettings": {
        "domainNameLabel": "[variables('dnsLabelPrefix')]"
      }
    }
  }
],
"outputs": {
  "storageAccountName": {
    "type": "string",
    "value": "[variables('storageAccountName')]"
  }
}
}
DEPLOY

# these key-value pairs are passed into the ARM Template's `parameters` block
parameters = {
  "storageAccountType" = "Standard_GRS"
}

deployment_mode = "Incremental"
}

output "storageAccountName" {
  value = "${lookup(azurerm_template_deployment.test.outputs, "storageAccountName")}"
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the template deployment. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the template deployment.
- **deployment_mode** - (Required) Specifies the mode that is used to deploy

resources. This value could be either **Incremental** or **Complete**. Note that you will almost *always* want this to be set to **Incremental** otherwise the deployment will destroy all infrastructure not specified within the template, and Terraform will not be aware of this.

- **template_body** - (Optional) Specifies the JSON definition for the template.

Note: There's a **file** function available which allows you to read this from an external file, which helps makes this more resource more readable.

- **parameters** - (Optional) Specifies the name and value pairs that define the deployment parameters for the template.
- **parameters_body** - (Optional) Specifies a valid Azure JSON parameters file that define the deployment parameters. It can contain KeyVault references

Note: There's a **file** function available which allows you to read this from an external file, which helps makes this more resource more readable.

» Attributes Reference

The following attributes are exported:

- **id** - The Template Deployment ID.
- **outputs** - A map of supported scalar output types returned from the deployment (currently, Azure Template Deployment outputs of type String, Int and Bool are supported, and are converted to strings - others will be ignored) and can be accessed using **.outputs["name"]**.

» Note

Terraform does not know about the individual resources created by Azure using a deployment template and therefore cannot delete these resources during a destroy. Destroying a template deployment removes the associated deployment operations, but will not delete the Azure resources created by the deployment. In order to delete these resources, the containing resource group must also be destroyed. More information.

» **azurerm__media__services__account**

Manages a Media Services Account.

» Example Usage

```
resource "azurerm_resource_group" "test" {
  name      = "media-resources"
  location  = "West Europe"
}

resource "azurerm_storage_account" "test" {
  name                        = "examplestoracc"
  resource_group_name        = "${azurerm_resource_group.test.name}"
  location                   = "${azurerm_resource_group.test.location}"
  account_tier               = "Standard"
  account_replication_type   = "GRS"
}

resource "azurerm_media_services_account" "test" {
  name                      = "examplemediaacc"
  location                  = "${azurerm_resource_group.test.location}"
  resource_group_name       = "${azurerm_resource_group.test.name}"

  storage_account {
    id      = "${azurerm_storage_account.test.id}"
    is_primary = true
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) Specifies the name of the Media Services Account. Changing this forces a new resource to be created.
- **resource_group_name** - (Required) The name of the resource group in which to create the Media Services Account. Changing this forces a new resource to be created.
- **location** - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- **storage_account** - (Required) One or more **storage_account** blocks as defined below.

A **storage_account** block supports the following:

- **id** - (Required) Specifies the ID of the Storage Account that will be associated with the Media Services instance.
- **is_primary** - (Required) Specifies whether the storage account should be the primary account or not. Defaults to **false**.

NOTE: Whilst multiple **storage_account** blocks can be specified - one of them must be set to the primary

» Attributes Reference

The following attributes are exported:

- **id** - The Resource ID of the Media Services Account.

» Import

Media Services Accounts can be imported using the **resource id**, e.g.

```
terraform import azure_rm_media_services_account.account /subscriptions/00000000-0000-0000-0000-000000000000
```