

» Data Source: newrelic__alert__channel

Use this data source to get information about a specific alert channel in New Relic that already exists. More information on Terraform's data sources can be found [here](#).

» Example Usage

```
# Data source
data "newrelic_alert_channel" "foo" {
  name = "foo@example.com"
}

# Resource
resource "newrelic_alert_policy" "foo" {
  name = "foo"
}

# Using the data source and resource together
resource "newrelic_alert_policy_channel" "foo" {
  policy_id = newrelic_alert_policy.foo.id
  channel_id = data.newrelic_alert_channel.foo.id
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the alert channel in New Relic.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- **id** - The ID of the alert channel.
- **type** - Alert channel type, either: `email`, `opsgenie`, `pagerduty`, `slack`, `victorops`, or `webhook`.
- **config** - Alert channel configuration.
- **policy_ids** - A list of policy IDs associated with the alert channel.

» Data Source: newrelic__alert__policy

Use this data source to get information about a specific alert policy in New Relic that already exists. More information on Terraform's data sources can be found [here](#).

» Example Usage

```
data "newrelic_alert_channel" "foo" {
  name = "foo@example.com"
}

data "newrelic_alert_policy" "foo" {
  name = "foo policy"
}

resource "newrelic_alert_policy_channel" "foo" {
  policy_id = data.newrelic_alert_policy.foo.id
  channel_id = data.newrelic_alert_channel.foo.id
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the alert policy in New Relic.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- **id** - The ID of the alert policy.
- **incident_preference** - The rollup strategy for the policy. Options include: PER_POLICY, PER_CONDITION, or PER_CONDITION_AND_TARGET. The default is PER_POLICY.
- **created_at** - The time the policy was created.
- **updated_at** - The time the policy was last updated.

» Data Source: newrelic__application

Use this data source to get information about a specific application in New Relic that already exists. More information on Terraform's data sources can be found

here.

» Example Usage

```
data "newrelic_application" "app" {
  name = "my-app"
}

resource "newrelic_alert_policy" "foo" {
  name = "foo"
}

resource "newrelic_alert_condition" "foo" {
  policy_id = newrelic_alert_policy.foo.id

  name          = "foo"
  type          = "apm_app_metric"
  entities      = [data.newrelic_application.app.id]
  metric        = "apdex"
  runbook_url   = "https://www.example.com"

  term {
    duration      = 5
    operator      = "below"
    priority      = "critical"
    threshold     = "0.75"
    time_function = "all"
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the application in New Relic.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- **id** - The ID of the application.
- **instance_ids** - A list of instance IDs associated with the application.
- **host_ids** - A list of host IDs associated with the application.

» Data Source: newrelic__key__transaction

Use this data source to get information about a specific key transaction in New Relic that already exists. More information on Terraform's data sources can be found [here](#).

» Example Usage

```
data "newrelic_key_transaction" "txn" {
  name = "txn"
}

resource "newrelic_alert_policy" "foo" {
  name = "foo"
}

resource "newrelic_alert_condition" "foo" {
  policy_id = newrelic_alert_policy.foo.id

  name          = "foo"
  type          = "apm_kt_metric"
  entities      = [data.newrelic_key_transaction.txn.id]
  metric        = "error_percentage"
  runbook_url   = "https://www.example.com"

  term {
    duration      = 5
    operator      = "below"
    priority      = "critical"
    threshold     = "0.75"
    time_function = "all"
  }
}
```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the key transaction in New Relic.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- `id` - The ID of the application.

» Data Source: `newrelic_plugin`

Use this data source to get information about a specific installed plugin in New Relic. More information on Terraform's data sources can be found [here](#).

Each plugin published to New Relic's Plugin Central is assigned a GUID. Once you have installed a plugin into your account it is assigned an ID. This account-specific ID is required when creating Plugins alert conditions.

» Example Usage

```
data "newrelic_plugin" "foo" {
  guid = "com.example.my-plugin"
}

resource "newrelic_alert_policy" "foo" {
  name = "foo"
}

resource "newrelic_plugins_alert_condition" "foo" {
  policy_id      = newrelic_alert_policy.foo.id
  name           = "foo"
  metric         = "Component/Summary/Consumers[consumers]"
  plugin_id      = data.newrelic_plugin.foo.id
  plugin_guid    = data.newrelic_plugin.foo.guid
  value_function = "average"
  metric_description = "Queue consumers"

  term {
    duration      = 5
    operator       = "below"
    priority      = "critical"
    threshold     = "0.75"
    time_function = "all"
  }
}
```

» Argument Reference

The following arguments are supported:

- `guid` - (Required) The GUID of the plugin in New Relic.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- `id` - The ID of the installed plugin instance.

» Data Source: `newrelic_plugin_component`

Use this data source to get information about a single plugin component in New Relic that already exists. More information on Terraform's data sources can be found [here](#).

Each plugin component reporting into to New Relic is assigned a unique ID. Once you have a plugin component reporting data into your account, its component ID can be used to create Plugins alert conditions.

» Example Usage

```
data "newrelic_plugin" "foo" {
  guid = "com.example.my-plugin"
}

data "newrelic_plugin_component" "foo" {
  plugin_id = data.newrelic_plugin.foo.id
  name      = "My Plugin Component"
}

resource "newrelic_alert_policy" "foo" {
  name = "foo"
}

resource "newrelic_plugins_alert_condition" "foo" {
  policy_id      = newrelic_alert_policy.foo.id
  name           = "foo"
  metric         = "Component/Summary/Consumers[consumers]"
  plugin_id      = data.newrelic_plugin.foo.id
  plugin_guid    = data.newrelic_plugin.foo.guid
  entities       = [data.newrelic_plugin_component.foo.id]
  value_function = "average"
  metric_description = "Queue consumers"
```

```

term {
  duration      = 5
  operator      = "below"
  priority      = "critical"
  threshold     = "0.75"
  time_function = "all"
}
}

```

» Argument Reference

The following arguments are supported:

- `plugin_id` - (Required) The ID of the plugin instance this component belongs to.
- `name` - (Required) The name of the plugin component.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- `id` - The ID of the plugin component.
- `health_status` - The health status of the plugin component.

» Data Source: `newrelic__synthetics__monitor`

Use this data source to get information about a specific synthetics monitor in New Relic that already exists. This can be used to set up a Synthetics alert condition.

» Example Usage

```

data "newrelic_synthetics_monitor" "bar" {
  name = "bar"
}

resource "newrelic_synthetics_alert_condition" "baz" {
  policy_id = newrelic_alert_policy.foo.id

  name          = "baz"
  monitor_id    = data.newrelic_synthetics_monitor.bar.id
  runbook_url   = "https://www.example.com"
}

```

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the synthetics monitor in New Relic.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- **monitor_id** - The ID of the synthetics monitor.

» Data Source: `newrelic__synthetics__secure__credential`

Use this data source to get information about a specific Synthetics secure credential in New Relic that already exists.

Note that the secure credential's value is not returned as an attribute for security reasons.

» Example Usage

```
data "newrelic__synthetics__secure__credential" "foo" {  
  key = "MY_KEY"  
}
```

» Argument Reference

The following arguments are supported:

- **key** - (Required) The secure credential's key name. Regardless of the case used in the configuration, the provider will provide an upcased key to the underlying API.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- **description** - The secure credential's description.
- **created_at** - The time the secure credential was created.
- **last_updated** - The time the secure credential was last updated.

» Resource: `newrelic__alert__channel`

Use this resource to create and manage New Relic alert policies.

» Example Usage

» Email

```
resource "newrelic_alert_channel" "foo" {
  name = "foo"
  type = "email"

  config {
    recipients          = "foo@example.com"
    include_json_attachment = "1"
  }
}
```

See additional examples.

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the channel.
- **type** - (Required) The type of channel. One of: `email`, `slack`, `opsgenie`, `pagerduty`, `victorops`, or `webhook`.
- **config** - (Optional) A nested block that describes an alert channel configuration. Only one config block is permitted per alert channel definition. See Nested config blocks below for details.
- **configuration** - **Deprecated** (Optional) A map of key/value pairs with channel type specific values. This argument is deprecated. Use the **config** argument instead.

» Nested config blocks

Each alert channel type supports a specific set of arguments for the **config** block:

- **email**
 - **recipients** - (Required) Comma delimited list of email addresses.
 - **include_json_attachment** - (Optional) 0 or 1. Flag for whether or not to attach a JSON document containing information about the associated alert to the email that is sent to recipients.
- **webhook**

- **base_url** - (Required) The base URL of the webhook destination.
- **auth_password** - (Optional) Specifies an authentication password for use with a channel. Supported by the **webhook** channel type.
- **auth_type** - (Optional) Specifies an authentication method for use with a channel. Supported by the **webhook** channel type. Only HTTP basic authentication is currently supported via the value **BASIC**.
- **auth_username** - (Optional) Specifies an authentication username for use with a channel. Supported by the **webhook** channel type.
- **headers** - (Optional) A map of key/value pairs that represents extra HTTP headers to be sent along with the webhook payload.
- **headers_string** - (Optional) Use instead of **headers** if the desired payload is more complex than a list of key/value pairs (e.g. a set of headers that makes use of nested objects). The value provided should be a valid JSON string with escaped double quotes. Conflicts with **headers**.
- **payload** - (Optional) A map of key/value pairs that represents the webhook payload. Must provide **payload_type** if setting this argument.
- **payload_string** - (Optional) Use instead of **payload** if the desired payload is more complex than a list of key/value pairs (e.g. a payload that makes use of nested objects). The value provided should be a valid JSON string with escaped double quotes. Conflicts with **payload**.
- **payload_type** - (Optional) Can either be **application/json** or **application/x-www-form-urlencoded**. The **payload_type** argument is *required* if **payload** is set.
- **pagerduty**
 - **service_key** - (Required) Specifies the service key for integrating with Pagerduty.
- **victorops**
 - **key** - (Required) The key for integrating with VictorOps.
 - **route_key** - (Required) The route key for integrating with VictorOps.
- **slack**
 - **url** - (Required) Your organization's Slack URL.
 - **channel** - (Optional) The Slack channel to send notifications to.
- **opsgenie**
 - **api_key** - (Required) The API key for integrating with OpsGenie.
 - **region** - (Required) The data center region to store your data. Valid values are **US** and **EU**. Default is **US**.
 - **teams** - (Optional) A set of teams for targeting notifications. Multiple values are comma separated.
 - **tags** - (Optional) A set of tags for targeting notifications. Multiple values are comma separated.
 - **recipients** - (Optional) A set of recipients for targeting notifications. Multiple values are comma separated.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- `id` - The ID of the channel.

» Additional Examples

» Slack

```
resource "newrelic_alert_channel" "foo" {
  name = "slack-example"
  type = "slack"

  config {
    url      = "https://<YourOrganization>.slack.com"
    channel = "example-alerts-channel"
  }
}
```

» OpsGenie

```
resource "newrelic_alert_channel" "foo" {
  name = "opsgenie-example"
  type = "opsgenie"

  config {
    api_key      = "abc123"
    teams        = "team1, team2"
    tags          = "tag1, tag2"
    recipients   = "user1@domain.com, user2@domain.com"
  }
}
```

» PagerDuty

```
resource "newrelic_alert_channel" "foo" {
  name = "pagerduty-example"
  type = "pagerduty"

  config {
    service_key = "abc123"
  }
}
```

» VictorOps

```
resource "newrelic_alert_channel" "foo" {
  name = "victorops-example"
  type = "victorops"

  config {
    key          = "abc123"
    route_key    = "/example"
  }
}
```

» Webhook

```
resource "newrelic_alert_channel" "foo" {
  name = "webhook-example"
  type = "webhook"

  config {
    base_url = "http://www.test.com"
    payload_type = "application/json"
    payload = {
      condition_name = "$CONDITION_NAME"
      policy_name = "$POLICY_NAME"
    }

    headers = {
      header1 = value1
      header2 = value2
    }
  }
}
```

» Webhook with complex payload

```
resource "newrelic_alert_channel" "foo" {
  name = "webhook-example"
  type = "webhook"

  config {
    base_url = "http://www.test.com"
    payload_type = "application/json"
    payload_string = <<EOF
  {
    "my_custom_values": {
      "condition_name": "$CONDITION_NAME",
```

```

    "policy_name": "$POLICY_NAME"
  }
}
EOF
}
}

```

» Import

Alert channels can be imported using the id, e.g.

```
$ terraform import newrelic_alert_channel.main <id>
```

» Resource: newrelic_alert_condition

Use this resource to create and manage alert conditions for APM, Browser, and Mobile in New Relic.

» Example Usage

```

data "newrelic_application" "app" {
  name = "my-app"
}

resource "newrelic_alert_policy" "foo" {
  name = "foo"
}

resource "newrelic_alert_condition" "foo" {
  policy_id = newrelic_alert_policy.foo.id

  name          = "foo"
  type          = "apm_app_metric"
  entities      = [data.newrelic_application.app.id]
  metric        = "apdex"
  runbook_url   = "https://www.example.com"
  condition_scope = "application"

  term {
    duration      = 5
    operator       = "below"
    priority      = "critical"
    threshold     = "0.75"
  }
}

```

```

    time_function = "all"
  }
}

```

» Argument Reference

The following arguments are supported:

- **policy_id** - (Required) The ID of the policy where this condition should be used.
- **name** - (Required) The title of the condition. Must be between 1 and 64 characters, inclusive.
- **type** - (Required) The type of condition. One of: `apm_app_metric`, `apm_kt_metric`, `browser_metric`, `mobile_metric`
- **entities** - (Required) The instance IDs associated with this condition.
- **metric** - (Required) The metric field accepts parameters based on the type set. One of these metrics based on type:
 - `apm_app_metric`
 - * `apdex`
 - * `error_percentage`
 - * `response_time_background`
 - * `response_time_web`
 - * `throughput_background`
 - * `throughput_web`
 - * `user_defined`
 - `apm_kt_metric`
 - * `apdex`
 - * `error_count`
 - * `error_percentage`
 - * `response_time`
 - * `throughput`
 - `browser_metric`
 - * `ajax_response_time`
 - * `ajax_throughput`
 - * `dom_processing`
 - * `end_user_apdex`
 - * `network`
 - * `page_rendering`
 - * `page_view_throughput`
 - * `page_views_with_js_errors`
 - * `request_queuing`
 - * `total_page_load`
 - * `user_defined`
 - * `web_application`
 - `mobile_metric`

- * database
 - * images
 - * json
 - * mobile_crash_rate
 - * network_error_percentage
 - * network
 - * status_error_percentage
 - * user_defined
 - * view_loading
- **condition_scope** - (Required for some types) **application** or **instance**. Choose **application** for most scenarios. If you are using the JVM plugin in New Relic, the **instance** setting allows your condition to trigger for specific app instances.
 - **enabled** - (Optional) Whether the condition is enabled or not. Defaults to true.
 - **gc_metric** - (Optional) A valid Garbage Collection metric e.g. **GC/G1 Young Generation**.
 - **violation_close_timer** - (Optional) Automatically close instance-based violations, including JVM health metric violations, after the number of hours specified. Must be: 1, 2, 4, 8, 12 or 24.
 - **runbook_url** - (Optional) Runbook URL to display in notifications.
 - **term** - (Required) A list of terms for this condition. See Terms below for details.
 - **user_defined_metric** - (Optional) A custom metric to be evaluated.
 - **user_defined_value_function** - (Optional) One of: **average**, **min**, **max**, **total**, or **sample_size**.

» Terms

The **term** mapping supports the following arguments:

- **duration** - (Required) In minutes, must be in the range of 5 to 120, inclusive.
- **operator** - (Optional) **above**, **below**, or **equal**. Defaults to **equal**.
- **priority** - (Optional) **critical** or **warning**. Defaults to **critical**. Terms must include at least one **critical** priority term
- **threshold** - (Required) Must be 0 or greater.
- **time_function** - (Required) **all** or **any**.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- **id** - The ID of the alert condition.

» Import

Alert conditions can be imported using notation `alert_policy_id:alert_condition_id`, e.g.

```
$ terraform import newrelic_alert_condition.main 123456:6789012345
```

» Resource: newrelic__alert__policy

Use this resource to create and manage New Relic alert policies.

» Example Usage

```
resource "newrelic_alert_policy" "foo" {  
  name = "example"  
  incident_preference = "PER_POLICY" # PER_POLICY is default  
}
```

See additional examples.

» Argument Reference

The following arguments are supported:

- **name** - (Required) The name of the policy.
- **incident_preference** - (Optional) The rollup strategy for the policy. Options include: `PER_POLICY`, `PER_CONDITION`, or `PER_CONDITION_AND_TARGET`. The default is `PER_POLICY`.
- **channel_ids** - (Optional) An array of channel IDs (integers) to assign to the policy. Adding or removing channel IDs from this array will result in a new alert policy resource being created and the old one being destroyed. Also note that channel IDs *cannot* be imported via `terraform import` (see Import for info).

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- **id** - The ID of the policy.
- **created_at** - The time the policy was created.
- **updated_at** - The time the policy was last updated.

» Additional Examples

» Provision multiple notification channels and add those channels to a policy

```
# Provision a Slack notification channel.
resource "newrelic_alert_channel" "slack_channel" {
  name = "slack-example"
  type = "slack"

  config {
    url      = "https://hooks.slack.com/services/<*****>/<*****>"
    channel = "example-alerts-channel"
  }
}

# Provision an email notification channel.
resource "newrelic_alert_channel" "email_channel" {
  name = "email-example"
  type = "email"

  config {
    recipients          = "example@testing.com"
    include_json_attachment = "1"
  }
}

# Provision the alert policy.
resource "newrelic_alert_policy" "policy_with_channels" {
  name              = "example-with-channels"
  incident_preference = "PER_CONDITION"

  # Add the provisioned channels to the policy.
  channel_ids = [
    newrelic_alert_channel.slack_channel.id,
    newrelic_alert_channel.email_channel.id,
  ]
}
```

» Reference existing notification channels and add those channel to a policy

```
# Reference an existing Slack notification channel.
data "newrelic_alert_channel" "slack_channel" {
  name = "slack-channel-notification"
}
```

```
# Reference an existing email notification channel.
data "newrelic_alert_channel" "email_channel" {
  name = "test@example.com"
}

# Provision the alert policy.
resource "newrelic_alert_policy" "policy_with_channels" {
  name          = "example-with-channels"
  incident_preference = "PER_CONDITION"

  # Add the referenced channels to the policy.
  channel_ids = [
    data.newrelic_alert_channel.slack_channel.id,
    data.newrelic_alert_channel.email_channel.id,
  ]
}
```

» Import

Alert policies can be imported using the resource's id.

Example import:

```
$ terraform import newrelic_alert_policy.policy_with_channels 23423556
```

Please note that channel IDs (`channel_ids`) *cannot* be imported due channels being a separate resource. However, to add channels to an imported alert policy, you can import the policy, add the `channel_ids` attribute with the associated channel IDs, then run `terraform apply`. This will result in the original alert policy being destroyed and a new alert policy being created along with the channels being added to the policy.

» Resource: `newrelic__alert__policy__channel`

Use this resource to map alert policies to alert channels in New Relic.

» Example Usage

The example below will apply multiple alert channels to an existing New Relic alert policy.

```

# Fetches the data for this policy from your New Relic account
# and is referenced in the newrelic_alert_policy_channel block below.
data "newrelic_alert_policy" "example_policy" {
  name = "my-alert-policy"
}

# Creates an email alert channel.
resource "newrelic_alert_channel" "email_channel" {
  name = "bar"
  type = "email"

  config {
    recipients          = "foo@example.com"
    include_json_attachment = "1"
  }
}

# Creates a Slack alert channel.
resource "newrelic_alert_channel" "slack_channel" {
  name = "slack-channel-example"
  type = "slack"

  config {
    channel = "#example-channel"
    url     = "http://example-org.slack.com"
  }
}

# Applies the created channels above to the alert policy
# referenced at the top of the config.
resource "newrelic_alert_policy_channel" "foo" {
  policy_id = newrelic_alert_policy.example_policy.id
  channel_ids = [
    data.newrelic_alert_channel.email_channel.id,
    data.newrelic_alert_channel.slack_channel.id
  ]
}

```

» Argument Reference

The following arguments are supported:

- `policy_id` - (Required) The ID of the policy.
- `channel_ids` - (Optional*) Array of channel IDs to apply to the specified policy. We recommend sorting channel IDs in ascending order to avoid

drift your Terraform state.

- **channel_id** - **Deprecated!** (Optional*) The ID of the channel. Please use the **channel_ids** argument instead.

*Note: Even though **channel_id** and **channel_ids** are optional, at least one of those arguments must be used for this resource to work.

» Import

Alert policy channels can be imported using the following notation:
<policyID>:<channelID>:<channelID>, e.g.

```
$ terraform import newrelic_alert_policy_channel.foo 123456:3462754:2938324
```

When importing `newrelic_alert_policy_channel` resource, the attribute `channel_ids*` will be set in your Terraform state. You can import multiple channels as long as those channel IDs are included as part of the import ID hash.

*Note: The attribute `channel_id` is *deprecated* and will not be set when importing this resource.

» Resource: newrelic_application_label

» **DEPRECATED!** Use at your own risk. This feature may be removed in the next major release.

Use this resource to create, update, and delete an Application label in New Relic.

» Example Usage

```
data "newrelic_application" "app1" {
  name="myapp1"
}

data "newrelic_application" "app2" {
  name="myapp2"
}

resource "newrelic_application_label" "foo" {
  category = "Team"
  name = "MyTeam"
  links {
    applications = [
      data.newrelic_application.app1.id,
      data.newrelic_application.app2.id
    ]
  }
}
```

```

    ]
    servers = []
  }
}

```

» Argument Reference

The following arguments are supported:

- **category** - (Required) A string representing the label key/category.
- **name** - (Required) A string that will be assigned to the label.
- **links** - (Required) The resources to which label should be assigned to. At least one of the following attributes must be set.
 - **applications** - An array of application IDs.
 - **servers** - An array of server IDs.

» Import

Application labels can be imported using a concatenated **category** and **name**, e.g.

```
$ terraform import newrelic_application_label.my_label <CATEGORY>:<NAME>
```

» Resource: newrelic__dashboard

Use this resource to create and manage New Relic dashboards.

» Example Usage: Create a New Relic Dashboard

```

data "newrelic_application" "my_application" {
  name = "My Application"
}

resource "newrelic_dashboard" "exempldash" {
  title = "New Relic Terraform Example"

  filter {
    event_types = [
      "Transaction"
    ]
    attributes = [
      "appName",

```

```

        "name"
    ]
}

widget {
    title = "Requests per minute"
    visualization = "billboard"
    nrql = "SELECT rate(count(*), 1 minute) FROM Transaction"
    row = 1
    column = 1
}

widget {
    title = "Error rate"
    visualization = "gauge"
    nrql = "SELECT percentage(count(*), WHERE error IS True) FROM Transaction"
    threshold_red = 2.5
    row = 1
    column = 2
}

widget {
    title = "Average transaction duration, by application"
    visualization = "facet_bar_chart"
    nrql = "SELECT average(duration) FROM Transaction FACET appName"
    row = 1
    column = 3
}

widget {
    title = "Apdex, top 5 by host"
    duration = 1800000
    visualization = "metric_line_chart"
    entity_ids = [
        data.newrelic_application.my_application.id,
    ]
    metric {
        name = "Apdex"
        values = [ "score" ]
    }
    facet = "host"
    limit = 5
    row = 2
    column = 1
}

```

```

    widget {
      title = "Requests per minute, by transaction"
      visualization = "facet_table"
      nrql = "SELECT rate(count(*), 1 minute) FROM Transaction FACET name"
      row = 2
      column = 2
    }

    widget {
      title = "Dashboard Note"
      visualization = "markdown"
      source = "### Helpful Links\n\n* [New Relic One](https://one.newrelic.com)\n* [Developer"
      row = 2
      column = 3
    }
  }
}

```

» Argument Reference

The following arguments are supported:

- **title** - (Required) The title of the dashboard.
- **icon** - (Optional) The icon for the dashboard. Valid values are `adjust`, `archive`, `bar-chart`, `bell`, `bolt`, `bug`, `bullhorn`, `bullseye`, `clock-o`, `cloud`, `cog`, `comments-o`, `crosshairs`, `dashboard`, `envelope`, `fire`, `flag`, `flask`, `globe`, `heart`, `leaf`, `legal`, `life-ring`, `line-chart`, `magic`, `mobile`, `money`, `none`, `paper-plane`, `pie-chart`, `puzzle-piece`, `road`, `rocket`, `shopping-cart`, `sitemap`, `sliders`, `tablet`, `thumbs-down`, `thumbs-up`, `trophy`, `usd`, `user`, and `users`. Defaults to `bar-chart`.
- **visibility** - (Optional) Determines who can see the dashboard in an account. Valid values are `all` or `owner`. Defaults to `all`.
- **editable** - (Optional) Determines who can edit the dashboard in an account. Valid values are `all`, `editable_by_all`, `editable_by_owner`, or `read_only`. Defaults to `editable_by_all`.
- **widget** - (Optional) A nested block that describes a visualization. Up to 300 `widget` blocks are allowed in a dashboard definition. See Nested widget blocks below for details.
- **filter** - (Optional) A nested block that describes a dashboard filter. Exactly one nested `filter` block is allowed. See Nested filter block below for details.

» Attribute Refence

In addition to all arguments above, the following attributes are exported:

- `dashboard_url` - The URL for viewing the dashboard.

» Nested widget blocks

All nested `widget` blocks support the following common arguments:

- `title` - (Required) A title for the widget.
- `visualization` - (Required) How the widget visualizes data. Valid values are `billboard`, `gauge`, `billboard_comparison`, `facet_bar_chart`, `faceted_line_chart`, `facet_pie_chart`, `facet_table`, `faceted_area_chart`, `heatmap`, `attribute_sheet`, `single_event`, `histogram`, `funnel`, `raw_json`, `event_feed`, `event_table`, `uniques_list`, `line_chart`, `comparison_line_chart`, `markdown`, and `metric_line_chart`.
- `row` - (Required) Row position of widget from top left, starting at 1.
- `column` - (Required) Column position of widget from top left, starting at 1.
- `width` - (Optional) Width of the widget. Valid values are 1 to 3 inclusive. Defaults to 1.
- `height` - (Optional) Height of the widget. Valid values are 1 to 3 inclusive. Defaults to 1.
- `notes` - (Optional) Description of the widget.

Each `visualization` type supports an additional set of arguments:

- `billboard`, `billboard_comparison`:
 - `nrql` - (Required) Valid NRQL query string. See Writing NRQL Queries for help.
 - `threshold_red` - (Optional) Threshold above which the displayed value will be styled with a red color.
 - `threshold_yellow` - (Optional) Threshold above which the displayed value will be styled with a yellow color.
- `gauge`:
 - `nrql` - (Required) Valid NRQL query string. See Writing NRQL Queries for help.
 - `threshold_red` - (Required) Threshold above which the displayed value will be styled with a red color.
 - `threshold_yellow` - (Optional) Threshold above which the displayed value will be styled with a yellow color.
- `facet_bar_chart`, `facet_pie_chart`, `facet_table`, `faceted_area_chart`, `faceted_line_chart`, or `heatmap`:
 - `nrql` - (Required) Valid NRQL query string. See Writing NRQL Queries for help.
 - `drilldown_dashboard_id` - (Optional) The ID of a dashboard to link to from the widget's facets.
- `attribute_sheet`, `comparison_line_chart`, `event_feed`, `event_table`, `funnel`, `histogram`, `line_chart`, `raw_json`, `single_event`, or

- `uniques_list:`
 - `nrql` - (Required) Valid NRQL query string. See Writing NRQL Queries for help.
- `markdown:`
 - `source` - (Required) The markdown source to be rendered in the widget.
- `metric_line_chart:`
 - `entity_ids` - (Required) A collection of entity ids to display data for. These are typically application IDs.
 - `metric` - (Required) A nested block that describes a metric. Nested `metric` blocks support the following arguments:
 - * `name` - (Required) The metric name to display.
 - * `values` - (Required) The metric values to display.
 - `duration` - (Required) The duration, in ms, of the time window represented in the chart.
 - `end_time` - (Optional) The end time of the time window represented in the chart in epoch time. When not set, the time window will end at the current time.
 - `facet` - (Optional) Can be set to "host" to facet the metric data by host.
 - `limit` - (Optional) The limit of distinct data series to display.
- `application_breakdown:`
 - `entity_ids` - (Required) A collection of entity IDs to display data. These are typically application IDs.

» Nested filter block

The optional filter block supports the following arguments: * `event_types` - (Optional) A list of event types to enable filtering for. * `attributes` - (Optional) A list of attributes belonging to the specified event types to enable filtering for.

» Import

New Relic dashboards can be imported using their ID, e.g.

```
$ terraform import newrelic_dashboard.my_dashboard 8675309
```

» Resource: `newrelic_infra_alert_condition`

Use this resource to create and manage Infrastructure alert conditions in New Relic.

» Example Usage

```
resource "newrelic_alert_policy" "foo" {
  name = "foo"
}

resource "newrelic_infra_alert_condition" "high_disk_usage" {
  policy_id = newrelic_alert_policy.foo.id

  name      = "High disk usage"
  type      = "infra_metric"
  event     = "StorageSample"
  select    = "diskUsedPercent"
  comparison = "above"
  where     = "(`hostname` LIKE '%frontend%')"

  critical {
    duration      = 25
    value         = 90
    time_function = "all"
  }

  warning {
    duration      = 10
    value         = 90
    time_function = "all"
  }
}

resource "newrelic_infra_alert_condition" "high_db_conn_count" {
  policy_id = newrelic_alert_policy.foo.id

  name      = "High database connection count"
  type      = "infra_metric"
  event     = "DatastoreSample"
  select    = "provider.databaseConnections.Average"
  comparison = "above"
  where     = "(`hostname` LIKE '%db%')"
  integration_provider = "RdsDbInstance"

  critical {
    duration      = 25
    value         = 90
    time_function = "all"
  }
}
```

```

}

resource "newrelic_infra_alert_condition" "process_not_running" {
  policy_id = newrelic_alert_policy.foo.id

  name          = "Process not running (/usr/bin/ruby)"
  type          = "infra_process_running"
  comparison    = "equal"
  process_where = "`commandName` = '/usr/bin/ruby'"

  critical {
    duration = 5
    value    = 0
  }
}

resource "newrelic_infra_alert_condition" "host_not_reporting" {
  policy_id = newrelic_alert_policy.foo.id

  name      = "Host not reporting"
  type      = "infra_host_not_reporting"
  event     = "StorageSample"
  select    = "diskUsedPercent"
  where     = "(`hostname` LIKE '%frontend%')"

  critical {
    duration = 5
  }
}

```

» Argument Reference

The following arguments are supported:

- **policy_id** - (Required) The ID of the alert policy where this condition should be used.
- **name** - (Required) The Infrastructure alert condition's name.
- **type** - (Required) The type of Infrastructure alert condition. Valid values are `infra_process_running`, `infra_metric`, and `infra_host_not_reporting`.
- **event** - (Required) The metric event; for example, `SystemSample` or `StorageSample`. Supported by the `infra_metric` condition type.
- **select** - (Required) The attribute name to identify the metric being targeted; for example, `cpuPercent`, `diskFreePercent`, or `memoryResidentSizeBytes`. The underlying API will automatically populate this value for Infrastructure integrations (for example

`diskFreePercent`), so make sure to explicitly include this value to avoid diff issues. Supported by the `infra_metric` condition type.

- **comparison** - (Required) The operator used to evaluate the threshold value. Valid values are `above`, `below`, and `equal`. Supported by the `infra_metric` and `infra_process_running` condition types.
- **critical** - (Required) Identifies the threshold parameters for opening a critical alert violation. See Thresholds below for details.
- **warning** - (Optional) Identifies the threshold parameters for opening a warning alert violation. See Thresholds below for details.
- **enabled** - (Optional) Whether the condition is turned on or off. Valid values are `true` and `false`. Defaults to `true`.
- **where** - (Optional) If applicable, this identifies any Infrastructure host filters used; for example: `hostname LIKE '%cassandra%'`.
- **process_where** - (Optional) Any filters applied to processes; for example: `commandName = 'java'`. Supported by the `infra_process_running` condition type.
- **integration_provider** - (Optional) For alerts on integrations, use this instead of `event`. Supported by the `infra_metric` condition type.
- **runbook_url** - (Optional) Runbook URL to display in notifications.
- **violation_close_timer** - (Optional) Determines how much time will pass before a violation is automatically closed. Setting the time limit to 0 prevents a violation from being force-closed.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- **id** - The ID of the Infrastructure alert condition.
- **created_at** - The timestamp the alert condition was created.
- **updated_at** - The timestamp the alert condition was last updated.

» Thresholds

The **critical** and **warning** threshold mapping supports the following arguments:

- **duration** - (Required) Identifies the number of minutes the threshold must be passed or met for the alert to trigger. Threshold durations must be between 1 and 60 minutes (inclusive).
- **value** - (Optional) Threshold value, computed against the **comparison** operator. Supported by `infra_metric` and `infra_process_running` alert condition types.
- **time_function** - (Optional) Indicates if the condition needs to be sustained or to just break the threshold once; `all` or `any`. Supported by the `infra_metric` alert condition type.

» Import

Infrastructure alert conditions can be imported using a composite ID of <policy_id>:<condition_id>, e.g.

```
$ terraform import newrelic_infra_alert_condition.main 12345:67890
```

» Resource: newrelic__insights__event

Use this resource to create one or more Insights events during a terraform run.

» Example Usage

```
resource "newrelic_insights_event" "foo" {
  event {
    type = "MyEvent"

    timestamp = 1232471100

    attribute {
      key   = "a_string_attribute"
      value = "a string"
    }
    attribute {
      key   = "an_integer_attribute"
      value = 42
      type  = "int"
    }
    attribute {
      key   = "a_float_attribute"
      value = 101.1
      type  = "float"
    }
  }
}
```

» Argument Reference

The following arguments are supported:

- **event** - (Required) An event to insert into Insights. Multiple event blocks can be defined. See Events below for details.

» Events

The `event` mapping supports the following arguments:

- `type` - (Required) The event's name. Can be a combination of alphanumeric characters, underscores, and colons.
- `timestamp` - (Optional) Must be a Unix epoch timestamp. You can define timestamps either in seconds or in milliseconds.
- `attribute` - (Required) An attribute to include in your event payload. Multiple attribute blocks can be defined for an event. See Attributes below for details.

» Attributes

The `attribute` mapping supports the following arguments:

- `key` - (Required) The name of the attribute.
- `value` - (Required) The value of the attribute.
- `type` - (Optional) Specify the type for the attribute value. This is useful when passing integer or float values to Insights. Allowed values are `string`, `int`, or `float`. Defaults to `string`.

» Resource: `newrelic_nrql_alert_condition`

Use this resource to create and manage NRQL alert conditions in New Relic.

» Example Usage

» Type: `static` (default)

```
resource "newrelic_alert_policy" "foo" {
  name = "foo"
}

resource "newrelic_nrql_alert_condition" "foo" {
  policy_id = newrelic_alert_policy.foo.id

  name          = "foo"
  type          = "static"
  runbook_url   = "https://www.example.com"
  enabled       = true

  term {
    duration     = 5
  }
}
```

```

        operator      = "below"
        priority      = "critical"
        threshold     = "1"
        time_function = "all"
    }

    nrql {
        query          = "SELECT count(*) FROM SyntheticCheck WHERE monitorId = '<monitorId>'"
        since_value    = "3"
    }

    value_function = "single_value"
}

```

See additional examples.

» Argument Reference

The following arguments are supported:

- **policy_id** - (Required) The ID of the policy where this condition should be used.
- **name** - (Required) The title of the condition.
- **type** - (Optional) The type of the condition. Valid values are **static** or **outlier**. Defaults to **static**.
- **runbook_url** - (Optional) Runbook URL to display in notifications.
- **enabled** - (Optional) Whether to enable the alert condition. Valid values are **true** and **false**. Defaults to **true**.
- **term** - (Required) A list of terms for this condition. See Terms below for details.
- **nrql** - (Required) A NRQL query. See NRQL below for details.
- **value_function** - (Optional) Possible values are **single_value**, **sum**.
- **expected_groups** - (Optional) Number of expected groups when using outlier detection.
- **ignore_overlap** - (Optional) Whether to look for a convergence of groups when using outlier detection.
- **violation_time_limit_seconds** - (Optional) Sets a time limit, in seconds, that will automatically force-close a long-lasting violation after the time limit you select. Possible values are 3600, 7200, 14400, 28800, 43200, and 86400.

» Terms

The **term** mapping supports the following arguments:

- **duration** - (Required) In minutes, must be in the range of 1 to 120, inclusive.
- **operator** - (Optional) **above**, **below**, or **equal**. Defaults to **equal**.
- **priority** - (Optional) **critical** or **warning**. Defaults to **critical**.
- **threshold** - (Required) Must be 0 or greater.
- **time_function** - (Required) **all** or **any**.

» NRQL

The `nrql` attribute supports the following arguments:

- **query** - (Required) The NRQL query to execute for the condition.
- **since_value** - (Required) The value to be used in the `SINCE <X> MINUTES AGO` clause for the NRQL query. Must be between 1 and 20.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- **id** - The ID of the NRQL alert condition.

» Additional Examples

» Type: `outlier`

```
resource "newrelic_alert_policy" "foo" {
  name = "foo"
}

resource "newrelic_nrql_alert_condition" "foo" {
  policy_id = newrelic_alert_policy.foo.id

  name          = "outlier-example"
  runbook_url   = "https://bar.example.com"
  enabled       = true

  term {
    duration      = 10
    operator      = "above"
    priority      = "critical"
    threshold     = "0.65"
    time_function = "all"
  }
  nrql {
```



```

        query          = "SELECT percentile(duration, 99) FROM Transaction FACET remote_ip"
        since_value    = "3"
    }
    type                = "outlier"
    expected_groups     = 2
    ignore_overlap      = true
}

```

» Import

Alert conditions can be imported using a composite ID of <policy_id>:<condition_id>, e.g.

```
$ terraform import newrelic_nrql_alert_condition.main 12345:67890
```

The actual values for `policy_id` and `condition_id` can be retrieved from the following URL when looking at the alert condition:

https://alerts.newrelic.com/accounts/<account_id>/policies/<policy_id>/conditions/<condition_id>

» Resource: newrelic_plugins_alert_condition

Use this resource to create and manage plugins alert conditions in New Relic.

» Example Usage

```

data "newrelic_plugin" "foo" {
    guid = "com.example.my-plugin"
}

data "newrelic_plugin_component" "foo" {
    plugin_id = data.newrelic_plugin.foo.id
    name      = "MyPlugin"
}

resource "newrelic_alert_policy" "foo" {
    name = "foo"
}

resource "newrelic_plugins_alert_condition" "foo" {
    policy_id      = newrelic_alert_policy.foo.id
    name           = "foo"
    entities       = [data.newrelic_plugin_component.foo.id]
    metric         = "Component/Summary/Consumers[consumers]"
}

```

```

plugin_id      = data.newrelic_plugin.foo.id
plugin_guid    = data.newrelic_plugin.foo.guid
value_function = "average"
metric_description = "Queue consumers"

term {
  duration      = 5
  operator      = "below"
  priority      = "critical"
  threshold     = "0.75"
  time_function = "all"
}
}

```

» Argument Reference

The following arguments are supported:

- **policy_id** - (Required) The ID of the policy where this condition should be used.
- **name** - (Required) The title of the condition. Must be between 1 and 64 characters, inclusive.
- **metric** - (Required) The plugin metric to evaluate.
- **entities** - (Required) The plugin component IDs to target.
- **plugin_id** - (Required) The ID of the installed plugin instance which produces the metric.
- **plugin_guid** - (Required) The GUID of the plugin which produces the metric.
- **metric_description** - (Required) The metric description.
- **value_function** - (Required) The value function to apply to the metric data. One of `min`, `max`, `average`, `sample_size`, `total`, or `percent`.
- **runbook_url** - (Optional) Runbook URL to display in notifications.
- **enabled** - (Optional) Whether or not this condition is enabled.
- **term** - (Required) A list of terms for this condition. See Terms below for details.

» Terms

The **term** mapping supports the following arguments:

- **duration** - (Required) In minutes, must be in the range of 5 to 120, inclusive.
- **operator** - (Optional) `above`, `below`, or `equal`. Defaults to `equal`.
- **priority** - (Optional) `critical` or `warning`. Defaults to `critical`.
- **threshold** - (Required) Must be 0 or greater.

- `time_function` - (Required) `all` or `any`.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- `id` - The ID of the alert condition.

» Import

Alert conditions can be imported using the `id`, e.g.

```
$ terraform import newrelic_plugins_alert_condition.main 12345
```

» Resource: `newrelic_synthetics_alert_condition`

Use this resource to create and manage synthetics alert conditions in New Relic.

» Example Usage

```
data "newrelic_synthetics_monitor" "foo" {
  name = "foo"
}

resource "newrelic_synthetics_alert_condition" "foo" {
  policy_id = newrelic_alert_policy.foo.id

  name           = "foo"
  monitor_id     = data.newrelic_synthetics_monitor.foo.id
  runbook_url    = "https://www.example.com"
}
```

» Argument Reference

The following arguments are supported:

- `policy_id` - (Required) The ID of the policy where this condition should be used.
- `name` - (Required) The title of this condition.
- `monitor_id` - (Required) The ID of the Synthetics monitor to be referenced in the alert condition.
- `runbook_url` - (Optional) Runbook URL to display in notifications.

- **enabled** - (Optional) Set whether to enable the alert condition. Defaults to `true`.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- **id** - The ID of the Synthetics alert condition.

» Import

Synthetics alert conditions can be imported using a composite ID of `<policy_id>:<condition_id>`, e.g.

```
$ terraform import newrelic_synthetics_alert_condition.main 12345:67890
```

» Resource: `newrelic__synthetics__label`

Use this resource to create, update, and delete a Synthetics label in New Relic.

» Example Usage

```
resource "newrelic_synthetics_label" "foo" {
  monitor_id = newrelic_synthetics_monitor.foo.id
  type       = "MyCategory"
  value      = "MyValue"
}
```

» Argument Reference

The following arguments are supported:

- **monitor_id** - (Required) The ID of the monitor that will be assigned the label.
- **type** - (Required) A string representing the label key/category.
- **value** - (Required) A string representing the label value.

» Attributes Reference

The following attributes are exported:

- **href** - The URL of the Synthetics label.

» Import

Synthetics labels can be imported using an ID in the format `<monitor_id>:<type>:<value>`, e.g.

```
$ terraform import newrelic_synthetics_labels.foo 1a272364-f204-4cd3-ae2a-2d15a2bedadd:MyCat
```

» Resource: newrelic__synthetics__monitor

Use this resource to create, update, and delete a synthetics monitor in New Relic.

» Example Usage

» Type: SIMPLE

```
resource "newrelic_synthetics_monitor" "foo" {
  name = "foo"
  type = "SIMPLE"
  frequency = 5
  status = "ENABLED"
  locations = ["AWS_US_EAST_1", "AWS_US_EAST_2"]

  uri = "https://example.com" # Required for type "SIMPLE"
  validation_string = "add example validation check here" # Optional for type "SIMPLE"
  verify_ssl = true # Optional for type "SIMPLE"
}
```

See additional examples.

» Argument Reference

The following arguments are supported:

- **name** - (Required) The title of this monitor.
- **type** - (Required) The monitor type. Valid values are `SIMPLE`, `BROWSER`, `SCRIPT_BROWSER`, and `SCRIPT_API`.
- **frequency** - (Required) The interval (in minutes) at which this monitor should run.
- **status** - (Required) The monitor status (i.e. `ENABLED`, `MUTED`, `DISABLED`).
- **locations** - (Required) The locations in which this monitor should be run.
- **sla_threshold** - (Optional) The base threshold for the SLA report.

The `SIMPLE` monitor type supports the following additional arguments:

- `uri` - (Required) The URI for the monitor to hit.
- `validation_string` - (Optional) The string to validate against in the response.
- `verify_ssl` - (Optional) Verify SSL.
- `bypass_head_request` - (Optional) Bypass HEAD request.
- `treat_redirect_as_failure` - (Optional) Fail the monitor check if redirected.

The **BROWSER** monitor type supports the following additional arguments:

- `uri` - (Required) The URI for the monitor to hit.
- `validation_string` - (Optional) The string to validate against in the response.
- `verify_ssl` - (Optional) Verify SSL.

» Attributes Reference

The following attributes are exported:

- `id` - The ID of the Synthetics monitor.

» Additional Examples

Type: **BROWSER**

```
resource "newrelic_synthetics_monitor" "foo" {
  name = "foo"
  type = "BROWSER"
  frequency = 5
  status = "ENABLED"
  locations = ["AWS_US_EAST_1"]

  uri = "https://example.com" # required for type "SIMPLE"
  validation_string = "add example validation check here" # optional for type "SIMPLE"
  verify_ssl = true # optional for type "SIMPLE"
  bypass_head_request = true # Note: optional for type "SIMPLE"
  treat_redirect_as_failure = true # Note: optional for type "SIMPLE"
}
```

Type: **SCRIPT_BROWSER**

```
resource "newrelic_synthetics_monitor" "foo" {
  name = "foo"
  type = "SCRIPT_BROWSER"
  frequency = 5
  status = "ENABLED"
  locations = ["AWS_US_EAST_1"]
}
```

```

}

Type: SCRIPT_API

resource "newrelic_synthetic_monitor" "foo" {
  name = "foo"
  type = "SCRIPT_API"
  frequency = 5
  status = "ENABLED"
  locations = ["AWS_US_EAST_1"]
}

```

» Import

Synthetic monitors can be imported using the id, e.g.

```
$ terraform import newrelic_synthetic_monitor.main <id>
```

» Resource: newrelic_synthetic_monitor_script

Use this resource to update a synthetic monitor script in New Relic.

» Example Usage

```

resource "newrelic_synthetic_monitor" "foo" {
  name = "foo"
  type = "SCRIPT_BROWSER"
  frequency = 5
  status = "ENABLED"
  locations = ["AWS_US_EAST_1"]
}

data "template_file" "foo_script" {
  template = file("${path.module}/foo_script.tpl")
}

resource "newrelic_synthetic_monitor_script" "foo_script" {
  monitor_id = newrelic_synthetic_monitor.foo.id
  text = data.template_file.foo_script.rendered
}

```

» Argument Reference

The following arguments are supported:

- `monitor_id` - (Required) The ID of the monitor to attach the script to.
- `text` - (Required) The plaintext representing the monitor script.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- `id` - The ID of the Synthetics monitor that the script is attached to.

» Import

Synthetics monitor scripts can be imported using the `id`, e.g.

```
$ terraform import newrelic_synthetics_monitor_script.main <id>
```

» Resource: `newrelic__synthetics__secure__credential`

Use this resource to create and manage New Relic Synthetic secure credentials.

» Example Usage

```
resource "newrelic_synthetics_secure_credential" "foo" {  
  key = "MY_KEY"  
  value = "My value"  
  description = "My description"  
}
```

» Argument Reference

The following arguments are supported:

- `key` - (Required) The secure credential's key name. Regardless of the case used in the configuration, the provider will provide an upcased key to the underlying API.
- `value` - (Required) The secure credential's value.
- `description` - (Optional) The secure credential's description.

» Attributes Reference

In addition to all arguments above, the following attributes are exported:

- `created_at` - The time the secure credential was created.
- `updated_at` - The time the secure credential was last updated.

» Import

A Synthetics secure credential can be imported using its `key`:

```
$ terraform import newrelic_synthetics_secure_credential.foo MY_KEY
```

» Resource: `newrelic_workload`

Use this resource to create, update, and delete a New Relic One workload.

A New Relic Personal API key is required to provision this resource. Set the `personal_api_key` attribute in the `provider` block or the `NEWRELIC_PERSONAL_API_KEY` environment variable with your Personal API key.

» Example Usage

```
resource "newrelic_workload" "foo" {
  name = "Example workload"
  account_id = 12345678

  entity_guids = ["MjUyMDUyOHxBUE18QVBQTE1DQVRJT058MjE1MDM3Nzk1"]

  entity_search_query {
    query = "name like 'Example application'"
  }

  scope_account_ids = [12345678]
}
```

» Argument Reference

The following arguments are supported:

- `name` - (Required) The workload's name.

- **account_id** - (Required) The New Relic account ID where you want to create the workload.
- **entity_guids** - (Optional) A list of entity GUIDs manually assigned to this workload.
- **entity_search_query** - (Optional) A list of search queries that define a dynamic workload. See Nested `entity_search_query` blocks below for details.
- **scope_account_ids** - (Optional) A list of account IDs that will be used to get entities from.

» Nested `entity_search_query` blocks

All nested `entity_search_query` blocks support the following common arguments:

- **query** - (Required) The query.

» Attributes Reference

The following attributes are exported:

- **guid** - The unique entity identifier of the workload in New Relic.
- **workload_id** - The unique entity identifier of the workload.
- **permalink** - The URL of the workload.
- **composite_entity_search_query** - The composite query used to compose a dynamic workload.

» Import

New Relic One workloads can be imported using a concatenated string of the format `<account_id>:<workload_id>:<guid>`, e.g.

```
$ terraform import newrelic_workload.foo 12345678:1456:MjUyMDUyOHxBUE18QVBRTE1DQVRJT058MjE1
```